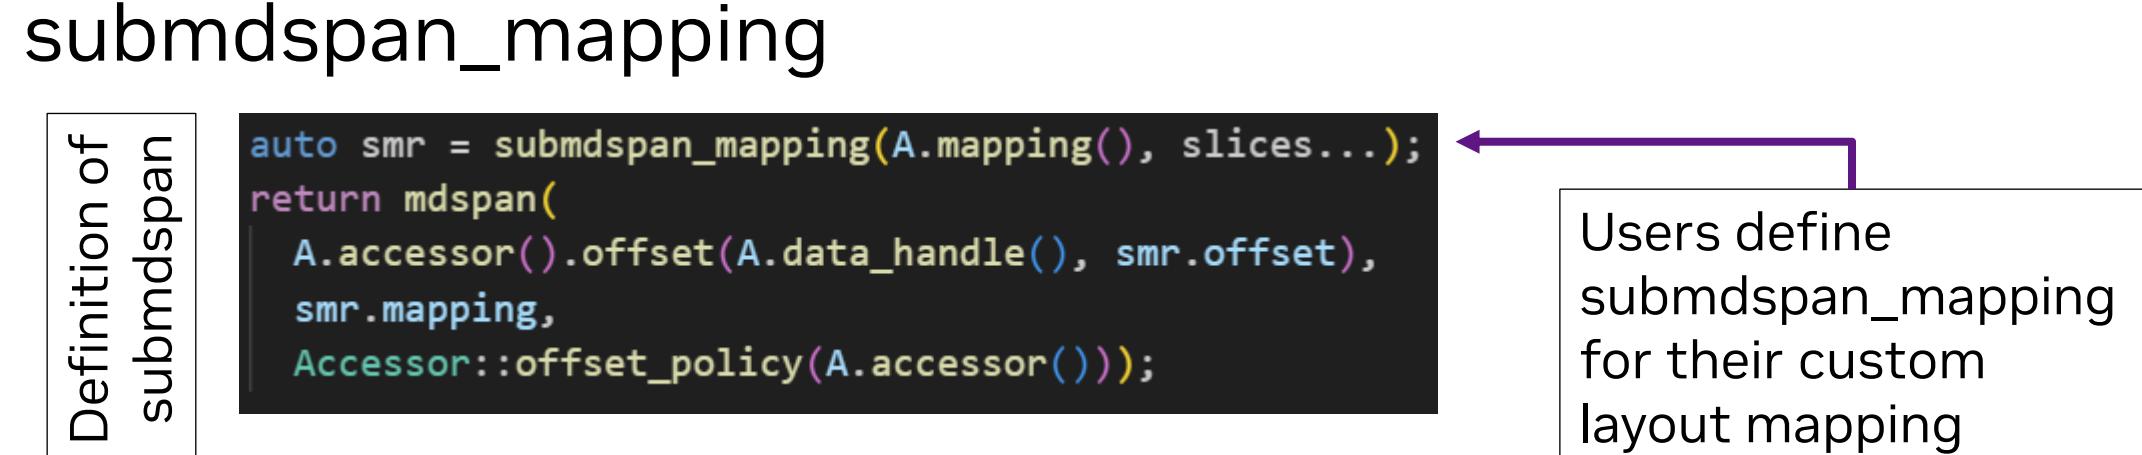
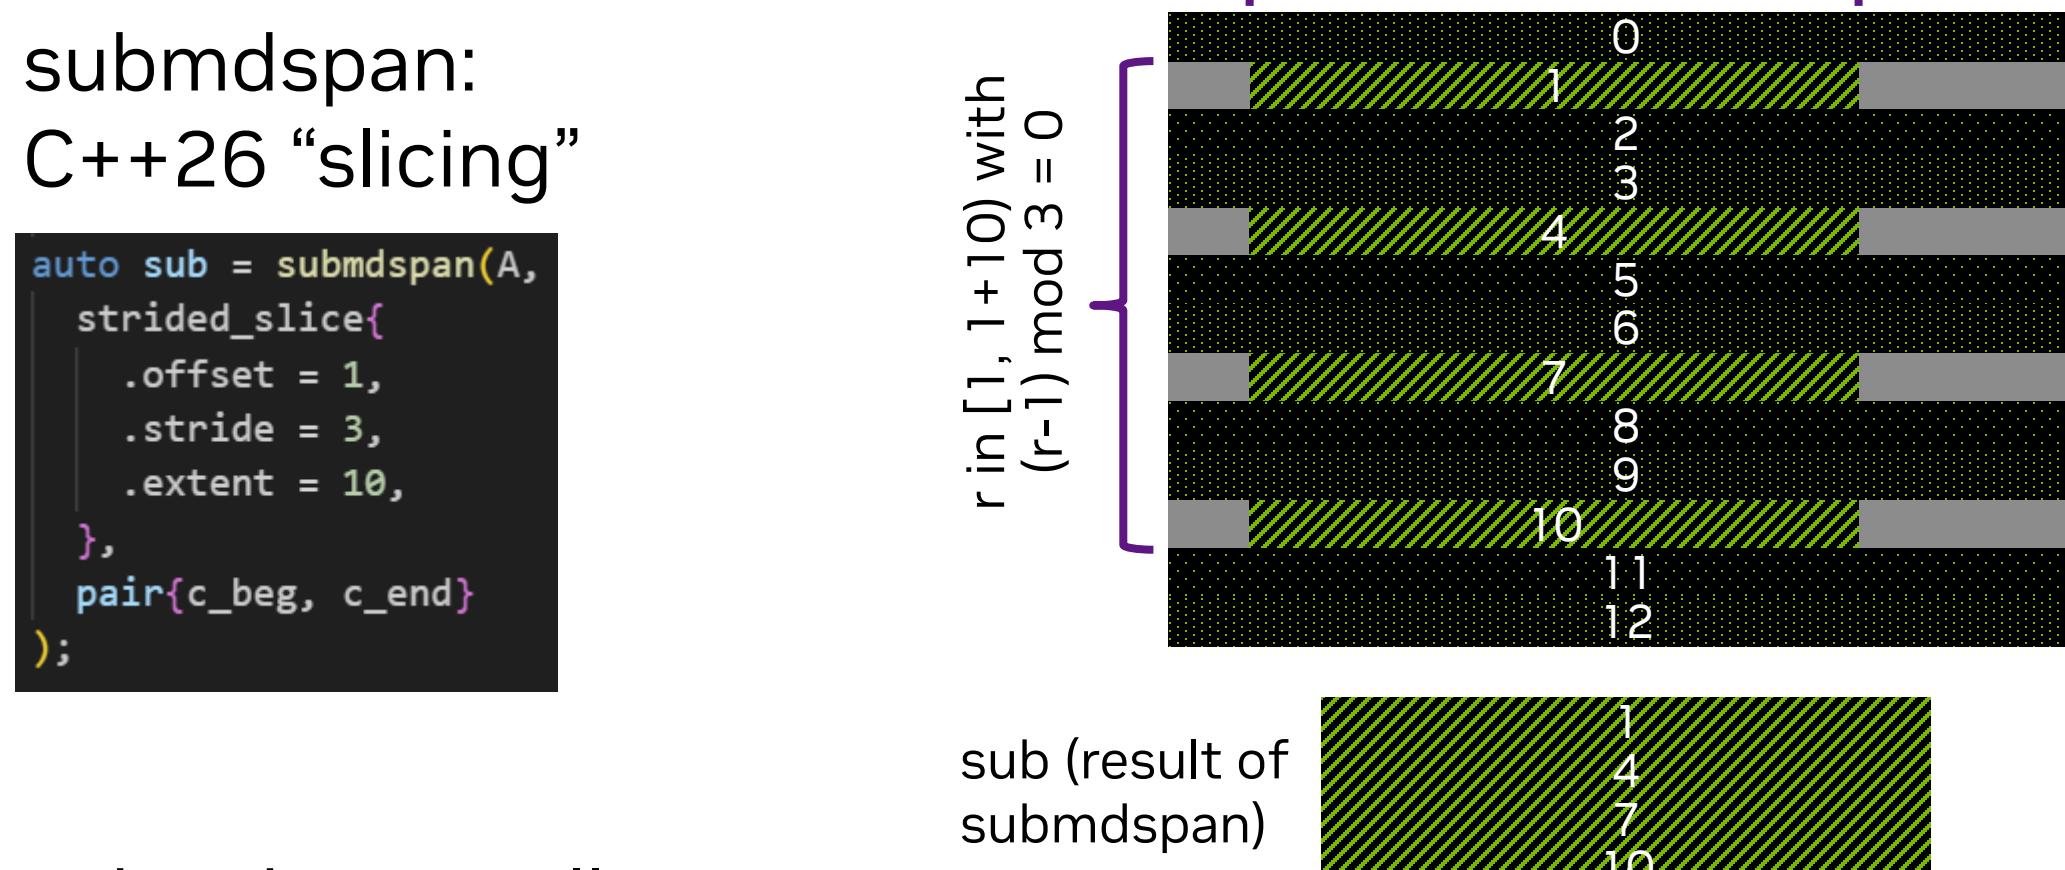
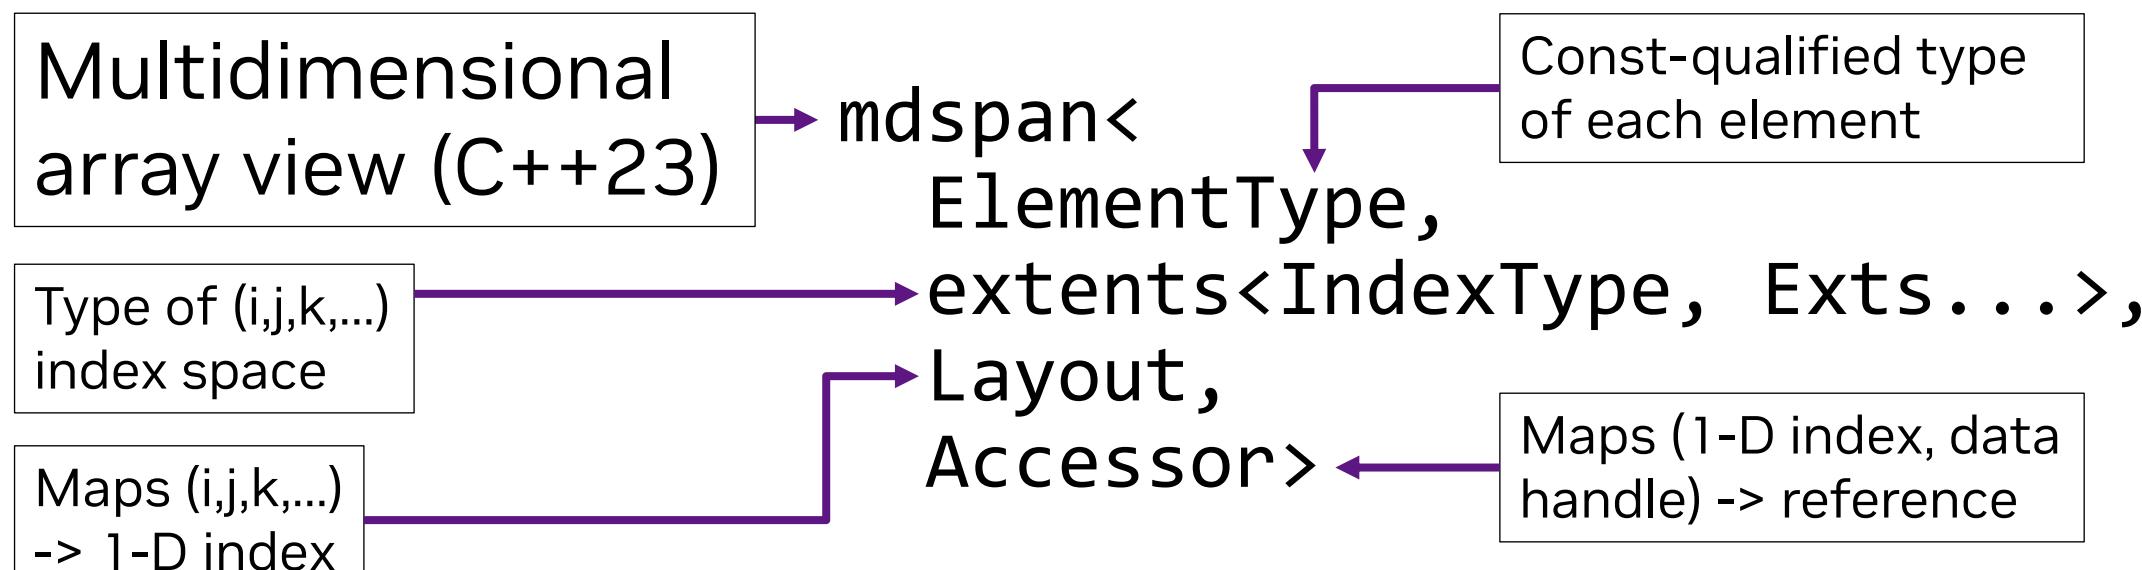


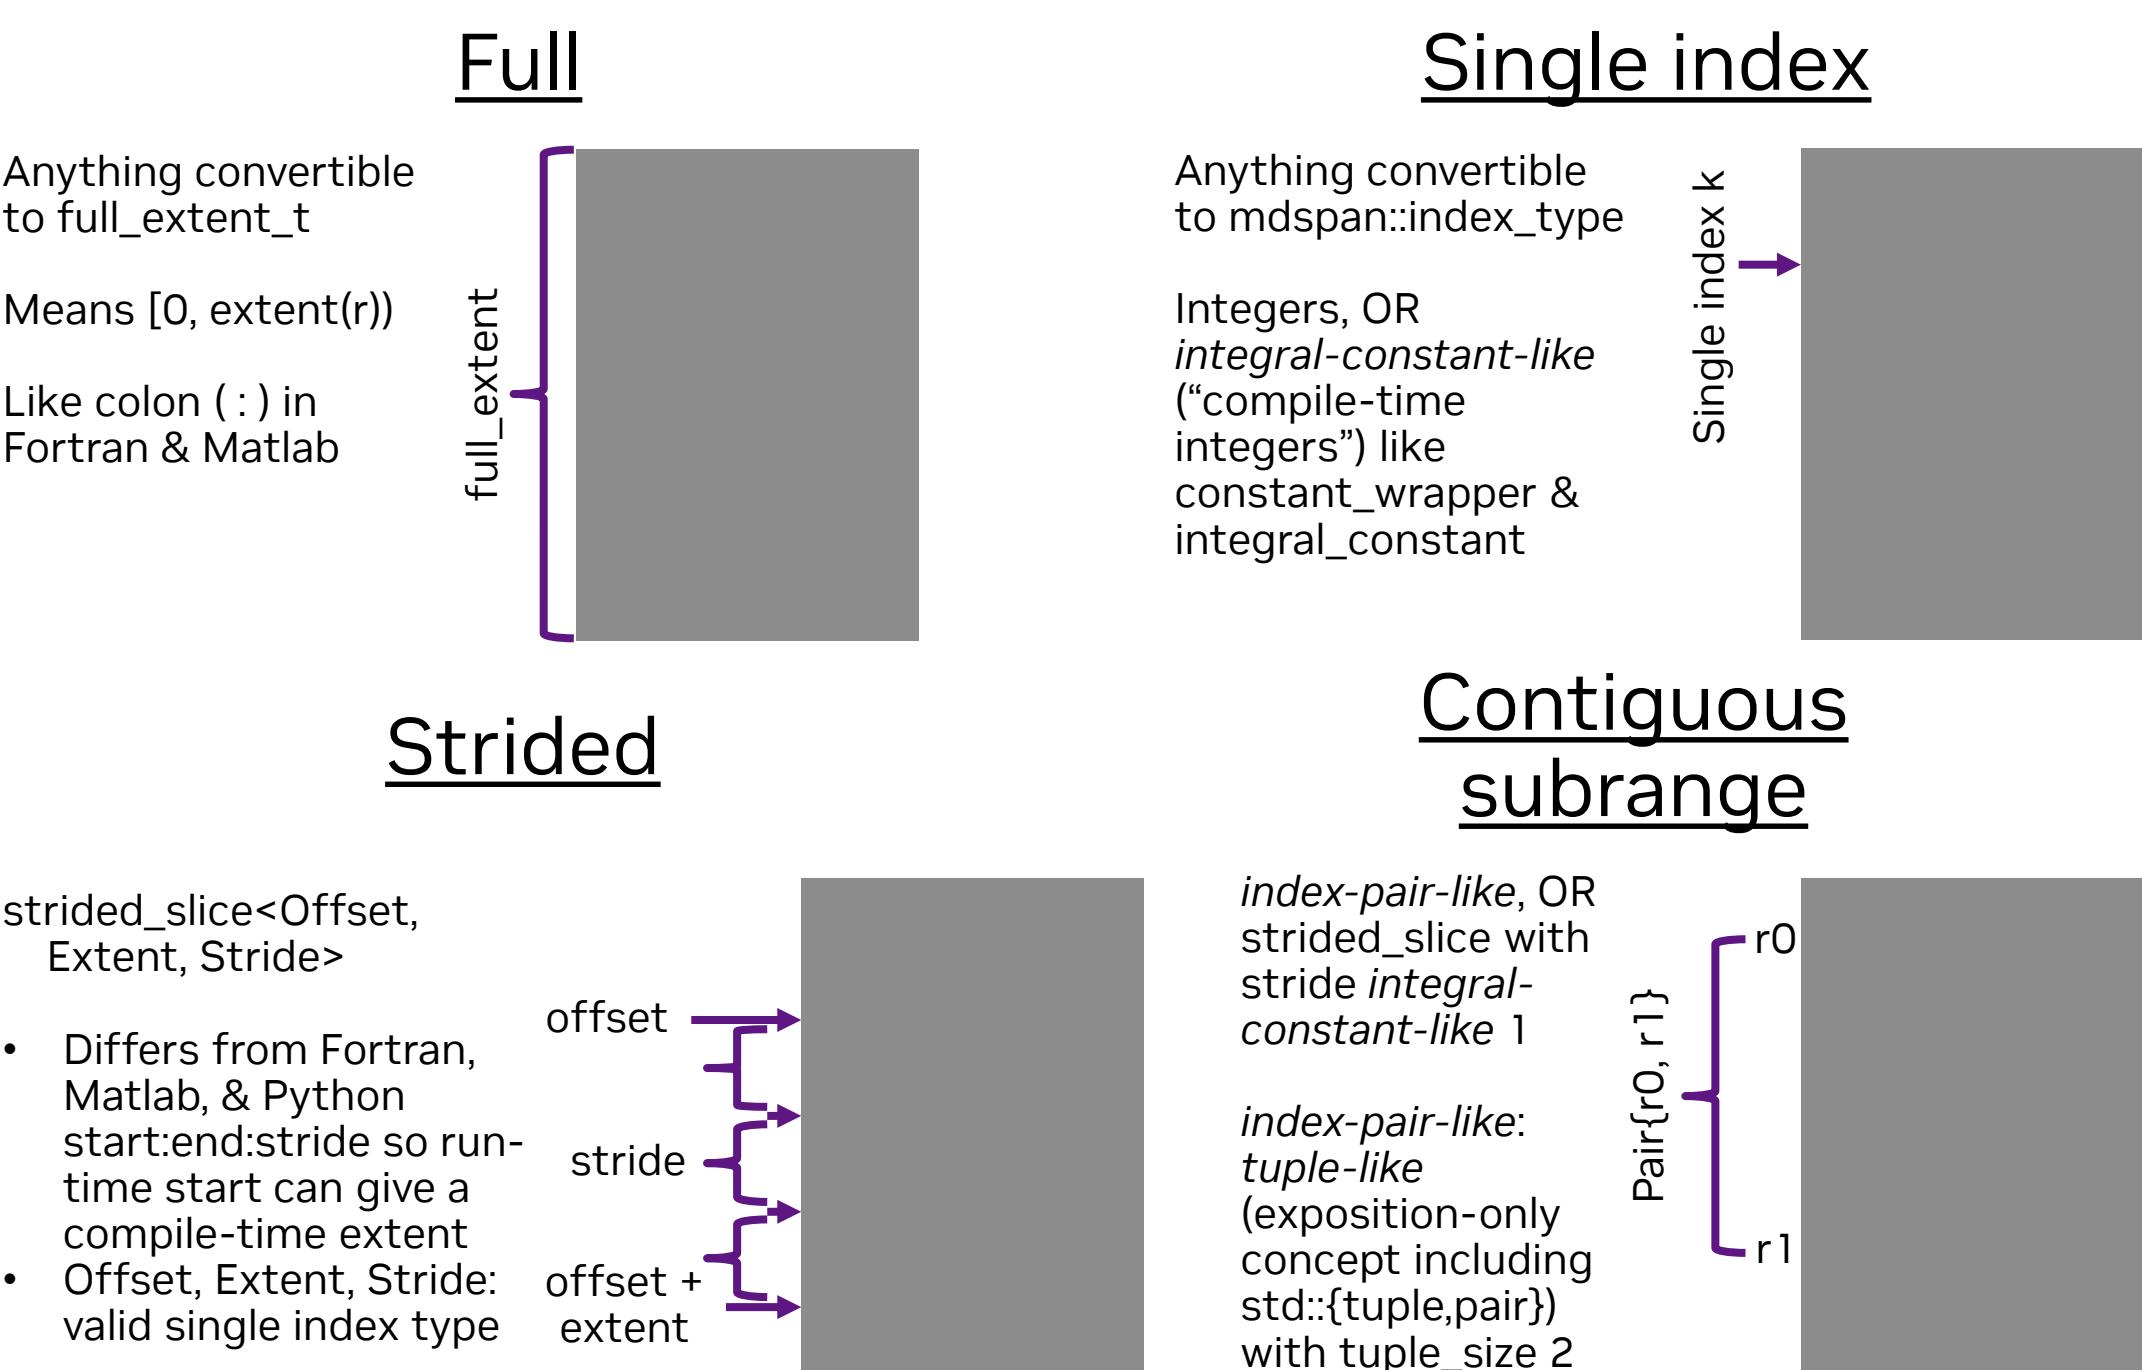
Performance effects of future-proofing submdspan_mapping

Mark Hoemmen, mhoemmen@nvidia.com

Background



Valid slice types: 4 categories



Design issue & fix

FEATURE: Valid slice types are an open set

- Permit user-defined pair or tuple types (e.g., cuda::std::tuple)
- Index wrappers (to prevent mixing up indices in nested for loops)

Later C++ versions may add new slice types

Accidentally, by expanding tuple-like

- Current definition: std::array<T,2>, std::complex, std::pair, std::tuple<X, Y>, std::ranges::subrange
- C++26 added std::complex, making it a valid slice type (unintended by submdspan authors)
- Proposals in flight (e.g., P2769) to redefine tuple-like as open set

Deliberately, e.g., by adding Fortran-style “start : end : stride”

BUG: Adding new slice type breaks user-defined submdspan_mapping customizations

- submdspan can call submdspan_mapping with any valid slice type
- Undefined behavior (UB) if submdspan calls it with invalid slice type
- submdspan_mapping NOT required to be ill-formed for invalid slice types
- Bad for custom layout mappings written to earlier C++ versions
- Users might exploit this: e.g., “Fake slice type” that makes custom submdspan_mapping print debug info & returns original mapping

FIX: submdspan “canonicalizes” slices

- Full → full_extent_t
- Single index → index_type or cw<Value>
- Strided → strided_slice (apply single index to members)
- Contiguous subrange → Strided with stride_type = cw<1>

cw<Value>: constant_wrapper, P2781 (in C++26)

User-defined submdspan_mapping must fail to compile with anything else

Fixed definition of submdspan

```
auto [...canonical_slices] =
    submdspan_canonicalize_slices(src.extents(), slices...);
auto smr = submdspan_mapping(A.mapping(), canonical_slices...);
return mdspan(
    A.accessor().offset(A.data_handle(), smr.offset),
    smr.mapping,
    Accessor::offset_policy(A.accessor()));
```

User code only sees canonical slices

P3663 proposes this fix

- Passed C++ Standard Committee design review; awaits wording review
- Hope to apply to C++26 via either National Body comment or Defect Report

Performance

submdspan may be called in tight loops

```
using space_3d_t = extents<int, dynamic_extent, 3>;
mdspan<float, space_3d_t> points_in_3d_space = /* ... */;
for (int r = 0; r < points_in_3d_space.extents(0); ++r) {
    auto point_r = submdspan(points_in_3d_space, r, full_extent);
    for (int c = r+1; c < points_in_3d_space.extents(1); ++c) {
        auto point_c = submdspan(points_in_3d_space, c, full_extent);
        forces[r, c] += force(point_r, point_c);
    }
}
```

Loop over N points in 3-D space

$O(N^2)$ calls to submdspan

submdspan MUST be fast or users will bypass it

- Point of mdspan is to write layout- & accessor-generic code
- Bypass → users get raw pointers & do error-prone index arithmetic

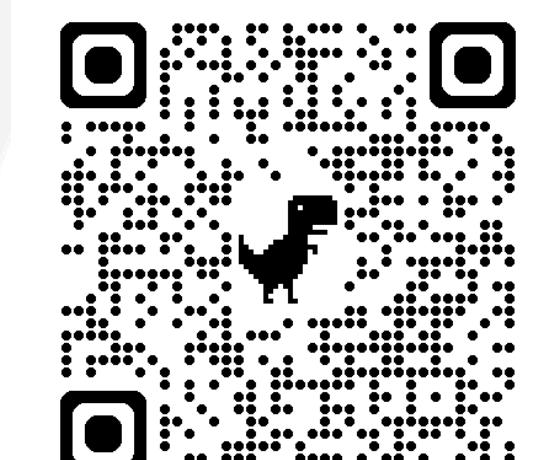
Canonicalization copies slices

- May consume hardware resources like GPU registers
- Increased code complexity may hinder inlining
- But fix may even improve run time by simplifying code
 - Fewer types → fewer function instantiations
 - strided_slice is an aggregate (unlike pair or tuple)

2 implementations

- Fully C++26: pack indexing, structured bindings can introduce a pack
- C++20 back-port

Both available in a pull request of the reference mdspan implementation:
<https://github.com/kokkos/mdspan>



Benchmarks

- Iterate over all elements by slicing mdspan recursively until rank 1
- submdspan(x, k, uss, full, ..., full, uss) for every row k
 - uss (“unit-stride slice”): either std::pair or cuda::std::pair (for CUDA builds)
 - Force conversions: k → index_type, full → full_extent_t, pair → strided_slice
 - Stay layout_right_padded to avoid layout_stride (which computes & stores std::array)
- element_type is uint8_t: minimize memory; overflow-free integer arithmetic
- index_type = int, size_t t x all {static, dynamic} extents: {2, 2, 2, 2, 2}
- Time in seconds, averaged over many iterations
- Hardware
 - NVIDIA H100 GPU
 - AMD EPYC 7232 CPU (8 cores, 2 threads/core)

Results

	int with static extents				int with dynamic extents			
	GCC 15.1	Clang 21.25.5	nvc++ 14.3	CUDA 13 + GCC 2.71E-08	GCC 15.1	Clang 21.25.5	nvc++ 14.3	CUDA 13 + GCC 6.83E-08
Baseline	1.13E-09 1.21E-07	2.25E-08			Baseline	1.17E-07 2.09E-07	1.05E-07	
C++26	1.16E-09 3.14E-09				C++26	1.91E-07 7.01E-07		
C++20	1.16E-09 3.14E-09	5.28E-08	2.75E-08		C++20	1.93E-07 7.03E-07	5.37E-07	3.33E-08
	size_t with static extents				size_t with dynamic extents			
	GCC 15.1	Clang 21.25.5	nvc++ 14.3	CUDA 13 + GCC 3.22E-08	GCC 15.1	Clang 21.25.5	nvc++ 14.3	CUDA 13 + GCC 8.04E-08
Baseline	1.32E-09 3.14E-09	3.45E-09			Baseline	1.23E-07 2.34E-08	9.70E-08	
C++26	1.16E-09 3.14E-09				C++26	1.50E-07 7.04E-08		
C++20	1.16E-09 3.14E-09	4.76E-08	3.14E-08		C++20	1.50E-07 7.06E-08	5.45E-07	3.31E-08