

+ 25

# Building a High-Performance Binary Serialization Format with In-Place Modification

HAMISH MORRISON

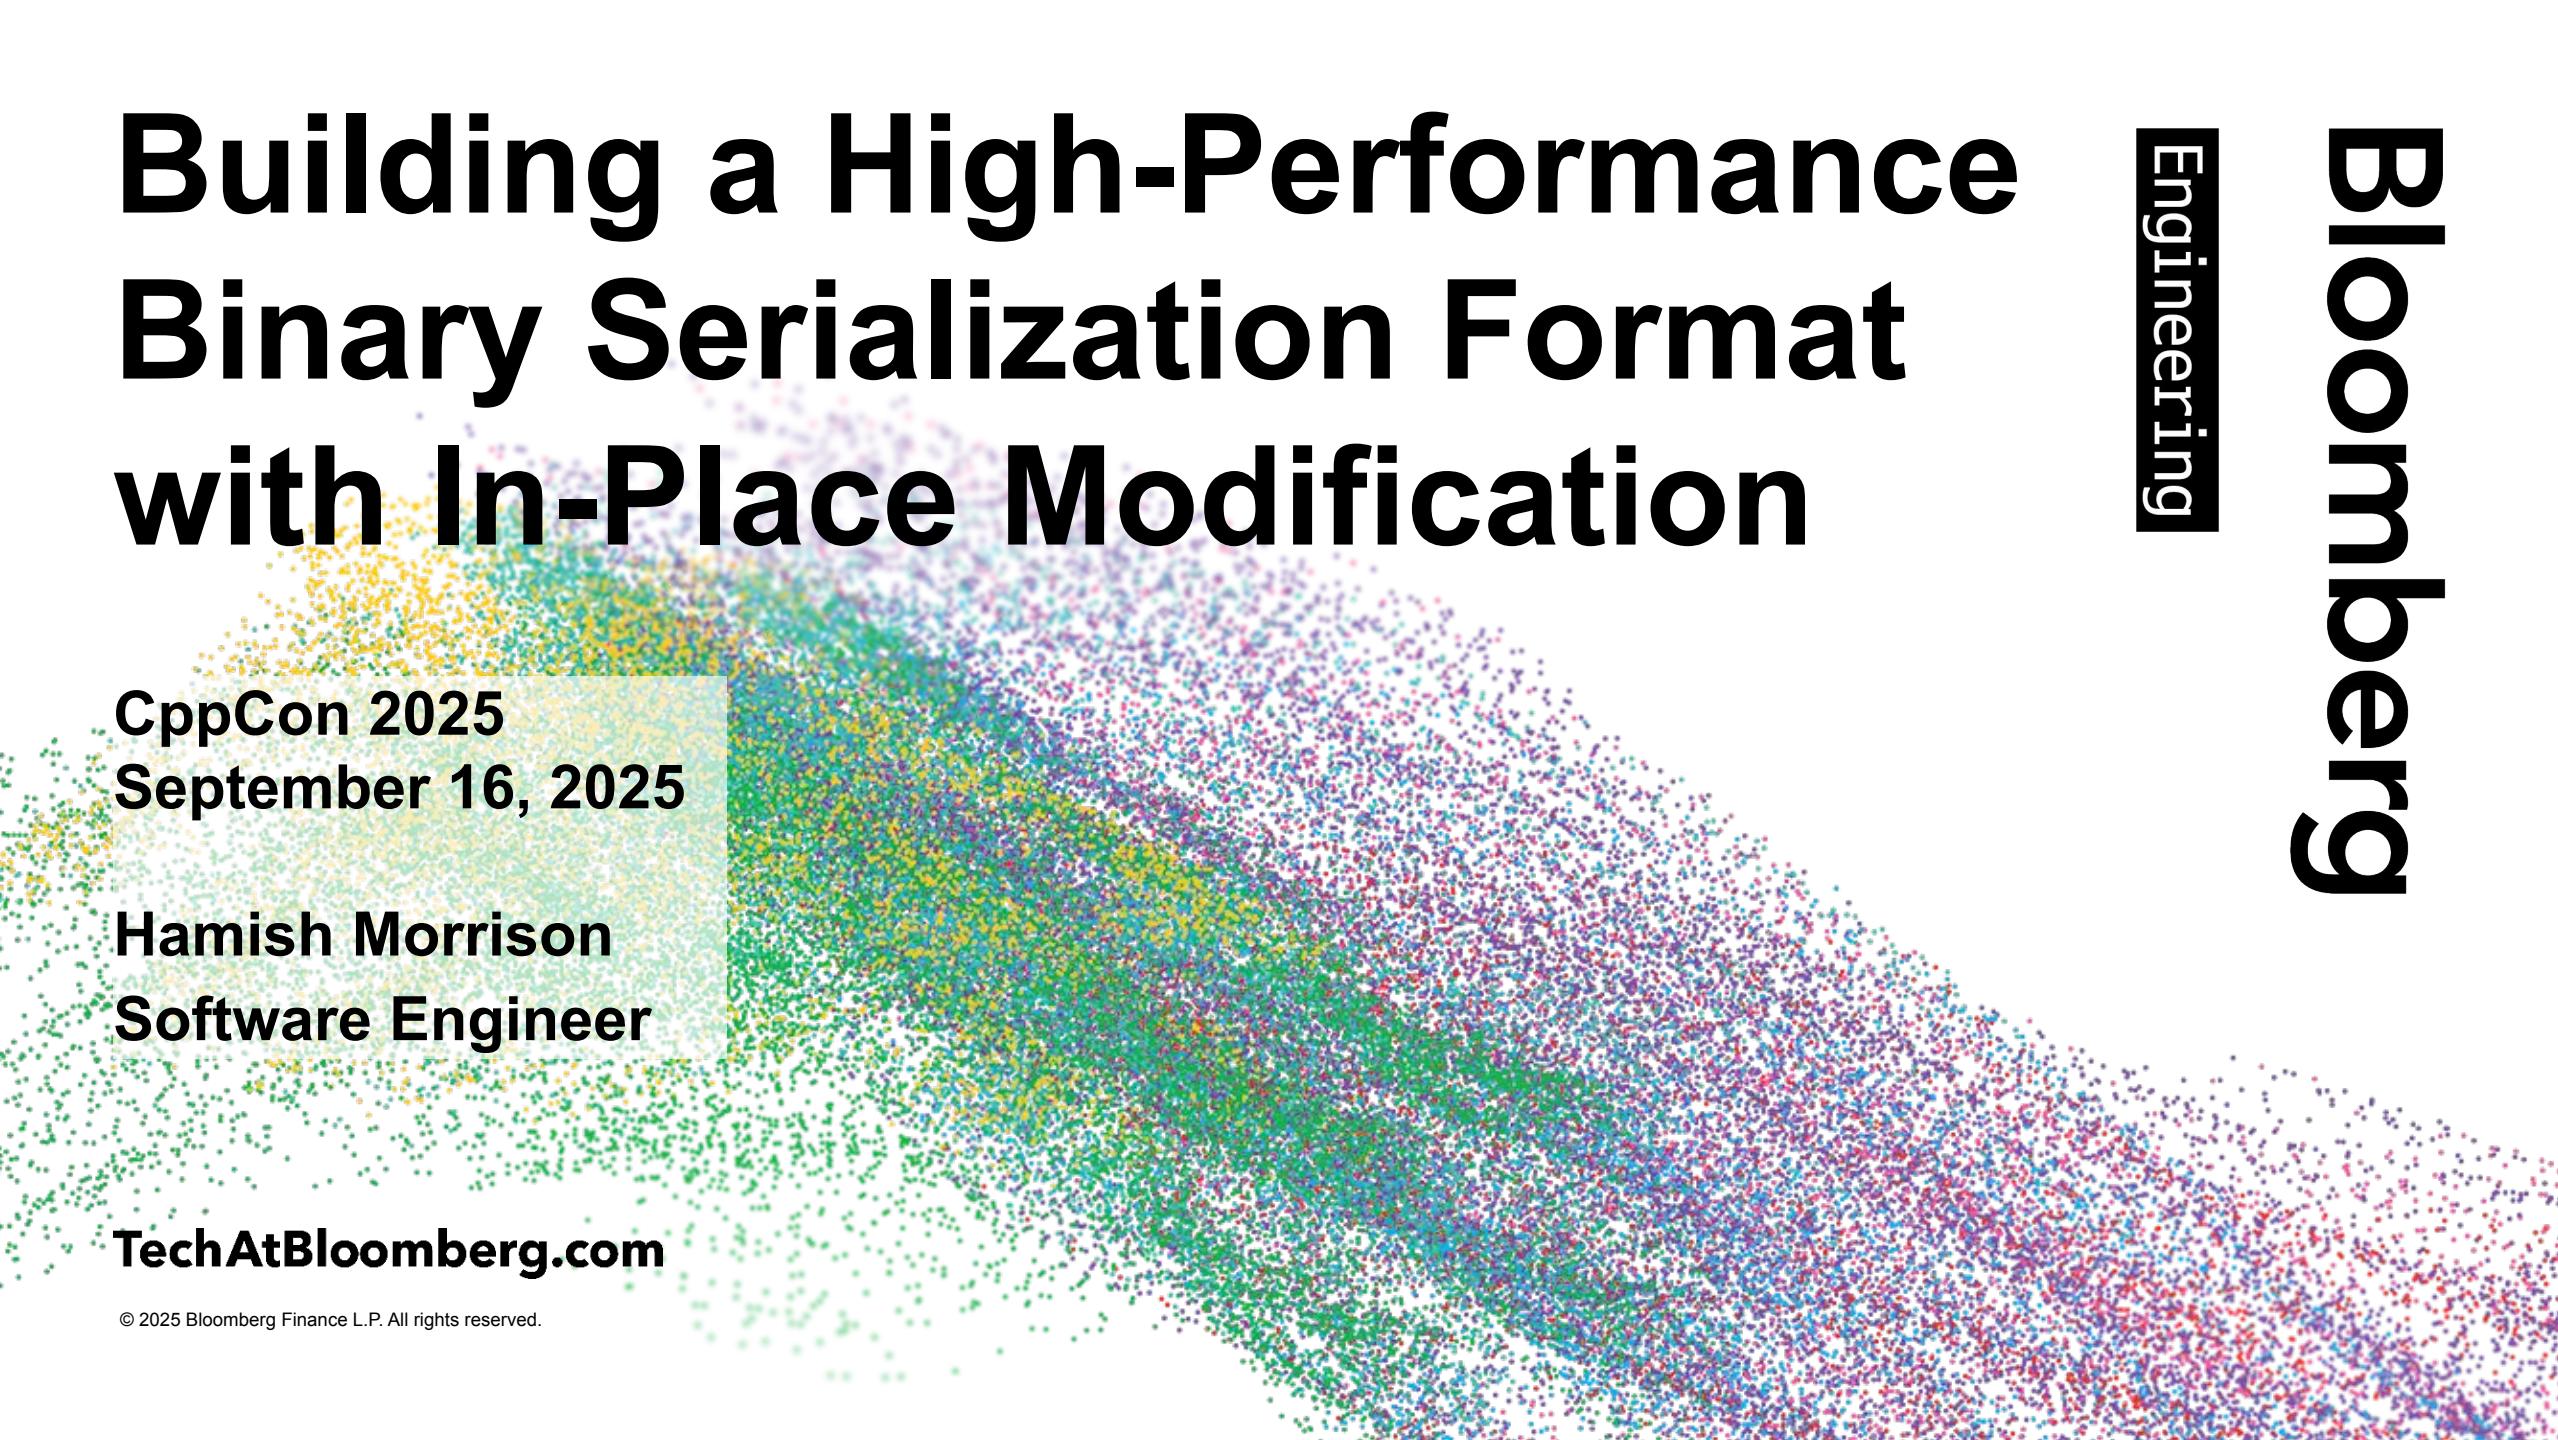


**Cppcon**  
The C++ Conference

20  
25



# Building a High-Performance Binary Serialization Format with In-Place Modification



Engineering

Bloomberg

CppCon 2025  
September 16, 2025

Hamish Morrison  
Software Engineer

TechAtBloomberg.com

# The Basics – What is Serialization?

- I have some data structure and I want to...
  - Send it over the network
  - Persist it somewhere
  - Share it between processes
- To do this, I need to turn it into a sequence of bytes
- And I need to be able to turn those bytes back into original data structure

# The Basics – What is Serialization?

- Example data structure:

```
struct Trade {  
    unsigned          price;  
    unsigned          volume;  
    unsigned long    timestamp;  
    char              tradeId[20];  
    char              buyBroker[5];  
    char              sellBroker[5];  
    unsigned          flags;  
};
```

# The Basics – What is Serialization?

- How do I serialize this?

```
struct Trade {  
    unsigned          price;  
    unsigned          volume;  
    unsigned long    timestamp;  
    char              tradeId[20];  
    char              buyBroker[5];  
    char              sellBroker[5];  
    unsigned          flags;  
};
```

# Serializing our Trade struct

- How about this:

```
size_t serialize(char *buffer, const Trade& trade)
{
    memcpy(buffer, &trade, sizeof(trade));
    return sizeof(Trade);
}
```

# Serializing our Trade struct

- How about this:

```
size_t serialize(char *buffer, const Trade& trade)
{
    memcpy(buffer, &trade, sizeof(trade));
    return sizeof(Trade);
}
```

- This is no good

# Serializing our Trade struct

- How about this:

```
size_t serialize(char *buffer, const Trade& trade)
{
    memcpy(buffer, &trade, sizeof(trade));
    return sizeof(Trade);
}
```

- This is no good
- Why?

# Serializing our Trade struct

- The problems with using memcpy:

```
struct Trade {  
    unsigned price;  
    unsigned volume;  
    unsigned long timestamp;  
    char tradeId[20];  
    char buyBroker[5];  
    char sellBroker[5];  
    unsigned flags;  
};
```

The diagram illustrates the memory layout of the `Trade` struct. It shows the fields: `price`, `volume`, `timestamp`, `tradeId`, `buyBroker`, `sellBroker`, and `flags`. A dashed line at the top spans the first four fields and is labeled "Endianness". Another dashed line spans the entire row of fields (from `timestamp` to `flags`) and is labeled "Size of types". A third dashed line at the bottom spans the last three fields (`buyBroker`, `sellBroker`, and `flags`) and is labeled "Padding between types".

# Serializing our Trade struct

- A better version:

```
void serialize(char *buffer, const Trade& trade)
{
    uint32_t price = htonl(trade.price);
    memcpy(buffer, &price, sizeof(price));
    buffer += sizeof(price);
    uint32_t volume = htonl(trade.size);
    memcpy(buffer, &size, sizeof(size));
    buffer += sizeof(volume);
    ...
}
```

- Types with defined sizes
- Swap everything to network order
- Consistent padding

# Serializing our Trade struct

- It's better, but it's still not great
  - We have to write a lot of boilerplate
  - The sending and receiving code have to agree exactly on this format
  - Evolving the schema is not possible
  - It's not easy to see how to encode more complex nested structures

# How Do We Improve on This?

- Two broad approaches:
  - Schema-based
  - Schemaless

In a schema-based format, the receiver has to know something about the shape of the data before it can decode it meaningfully.

In schemaless formats, it doesn't need to know anything.

# Schema-Based Formats

- In schema-based formats, the user defines a schema, usually in a DSL
- Often, code is generated from the schema to encode and decode messages
- e.g., in protobufs:

```
message Trade {  
    optional uint32 price      = 1;  
    optional uint32 volume     = 2;  
    optional uint64 timestamp  = 3;  
    optional string tradeId   = 4;  
    ...  
};
```

# Schema-Based Formats

- Sender and receiver need to agree on the schema, to an extent

```
message Trade {  
    optional uint32 price      = 1;  
    optional uint32 volume     = 2;  
    optional uint64 timestamp  = 3;  
    optional string tradeId   = 4;  
    ...  
};
```

# Schema-Based Formats

- Sender and receiver need to agree on the schema, to an extent
- But the schema is allowed to evolve in some ways

```
message Trade {  
    optional uint32 price      = 1;  
    optional uint32 volume     = 2;  
    optional uint64 timestamp  = 3;  
    optional string tradeId   = 4;  
    ...  
};
```

# Schema-Based Formats

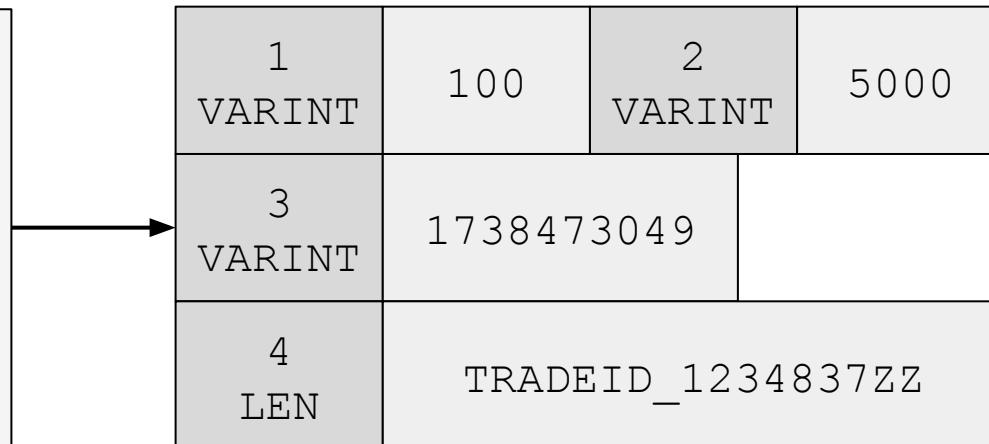
- Sender and receiver need to agree on the schema, to an extent
- But the schema is allowed to evolve in some ways
- e.g., a new field can be added

```
message Trade {  
    optional uint32 price      = 1;  
    optional uint32 volume     = 2;  
    optional uint64 timestamp  = 3;  
    optional string tradeId   = 4;  
    optional string brokerId  = 5;  
    ...  
};
```

# Schema-Based Formats

- On the wire, the protobuf message is encoded as a sequence of pairs
- Pairs of:
  - Tag – which is a field number and wire type
  - And value

```
message Trade {  
    optional uint32 price      = 1;  
    optional uint32 volume     = 2;  
    optional uint64 timestamp  = 3;  
    optional string tradeId   = 4;  
    ...  
};
```

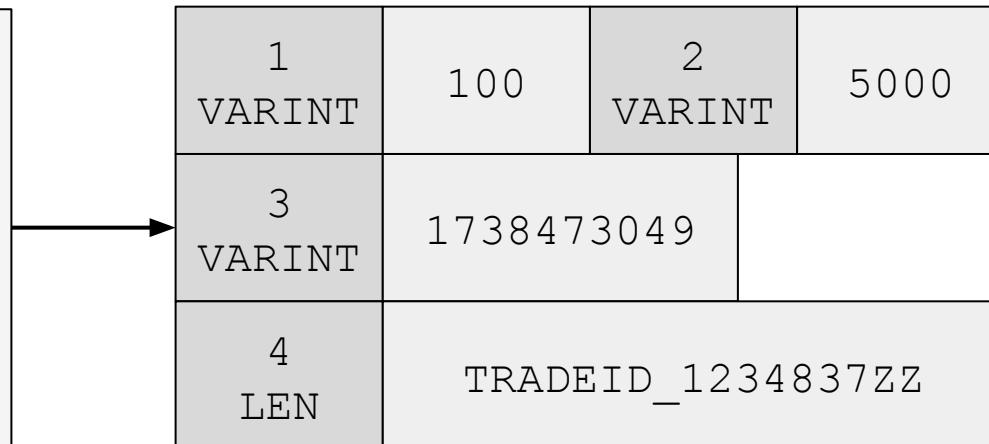


1 VARINT	100	2 VARINT	5000
3 VARINT	1738473049		
4 LEN	TRADEID_1234837ZZ		

# Schema-Based Formats

- This encoding is somewhat self-describing
- I can traverse through it, even if I don't know the schema
- But I need the schema to make sense of it (e.g. to know field names, etc)

```
message Trade {  
    optional uint32 price      = 1;  
    optional uint32 volume     = 2;  
    optional uint64 timestamp  = 3;  
    optional string tradeId   = 4;  
    ...  
};
```

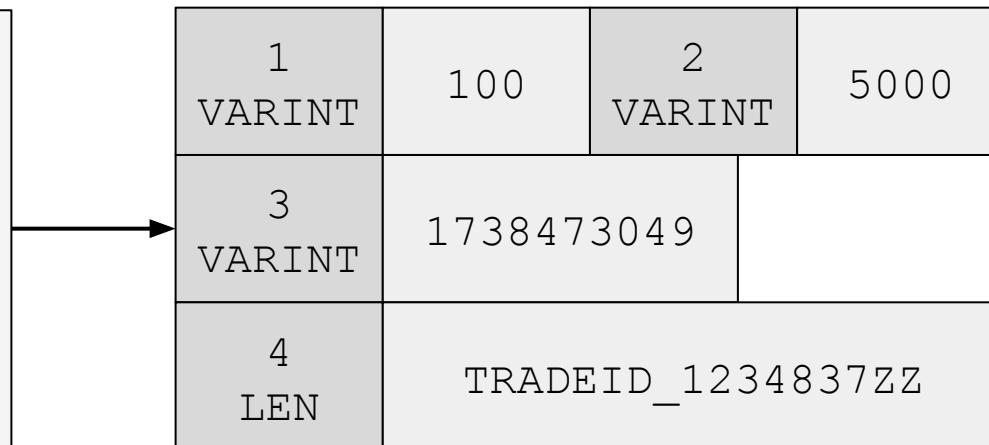


1 VARINT	100	2 VARINT	5000
3 VARINT	1738473049		
4 LEN	TRADEID_1234837ZZ		

# Schema-Based Formats

- What if I don't know the shape of my data up front?
- What if I want to encode more complex, arbitrarily nested documents of maps and arrays?

```
message Trade {  
    optional uint32 price      = 1;  
    optional uint32 volume     = 2;  
    optional uint64 timestamp  = 3;  
    optional string tradeId   = 4;  
    ...  
};
```



1 VARINT	100	2 VARINT	5000
3 VARINT	1738473049		
4 LEN	TRADEID_1234837ZZ		

# Schemaless Formats

- These formats generally let you encode anything, without having to agree on a schema between sender and receiver
- The user often gets a dictionary-style API to interact with, something like:

```
class Value {  
    int     asInt() const;  
    string asString() const;  
    Value   get(std::string key);  
    void    set(std::string key, int value);  
    void    set(std::string key, string value);  
};
```

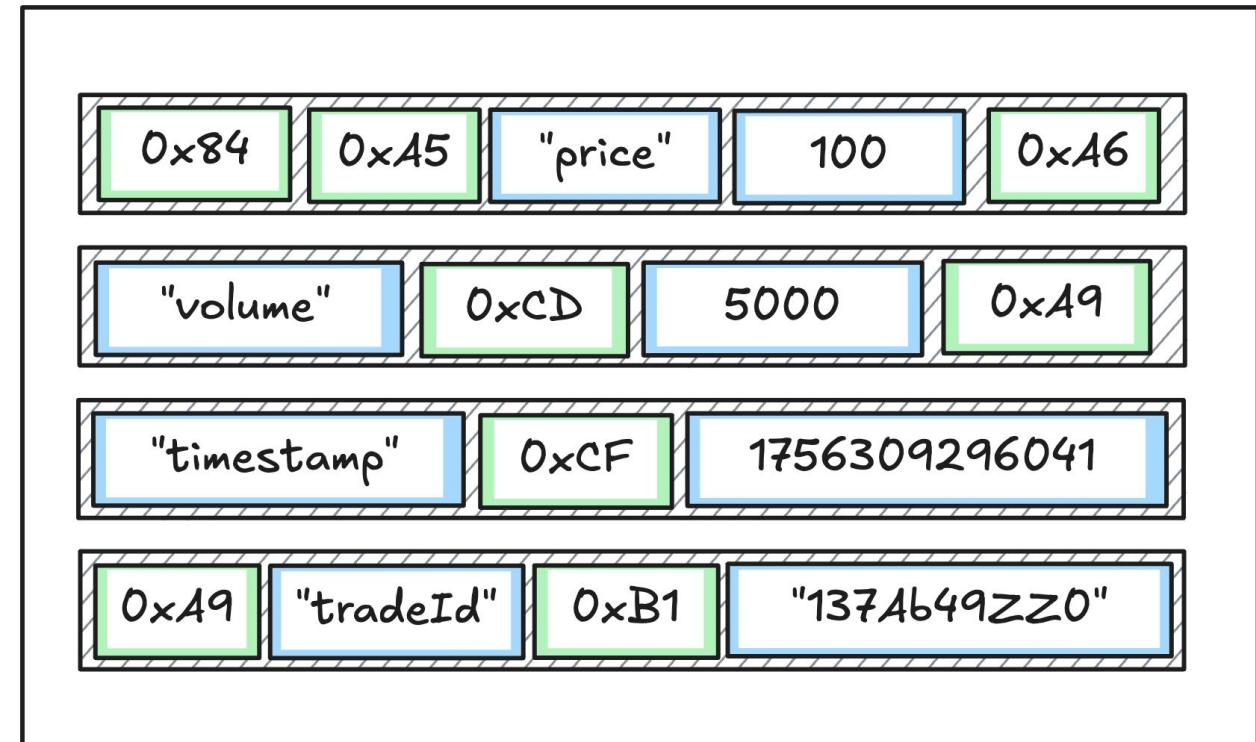
# Schemaless Formats

- Examples of schemaless formats include:
  - JSON
  - BSON
  - MsgPack
  - CBOR
  - FlexBuffers
- The binary formats are usually encoded as a sequence of tag, value pairs

# Schemaless Formats

- An example document, encoded in both JSON and MsgPack:

```
{  
  "price": 100,  
  "volume": 5000,  
  "timestamp": 1756309296041,  
  "tradeId": "137A649ZZ0"  
}
```



# Schemaless Formats

- An example document, encoded in both JSON and MsgPack:

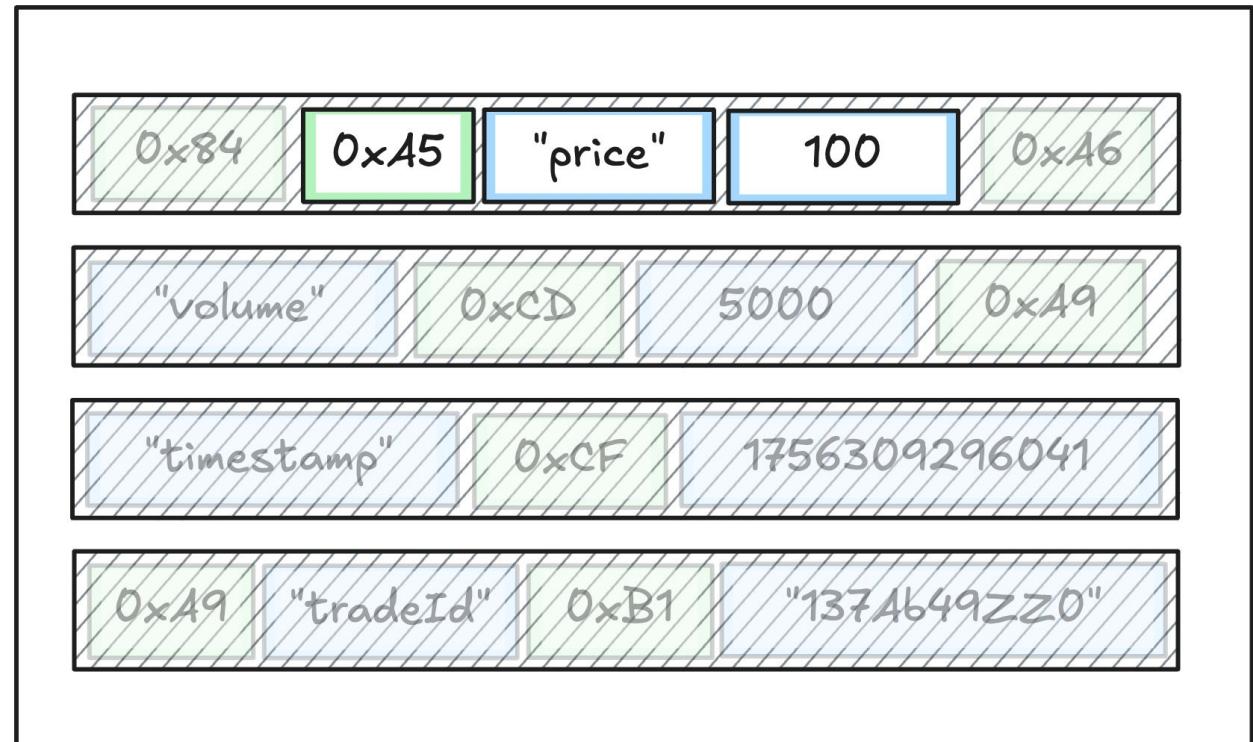
```
{  
  "price": 100,  
  "volume": 5000,  
  "timestamp": 1756309296041,  
  "tradeId": "137A649ZZ0"  
}
```



# Schemaless Formats

- An example document, encoded in both JSON and MsgPack:

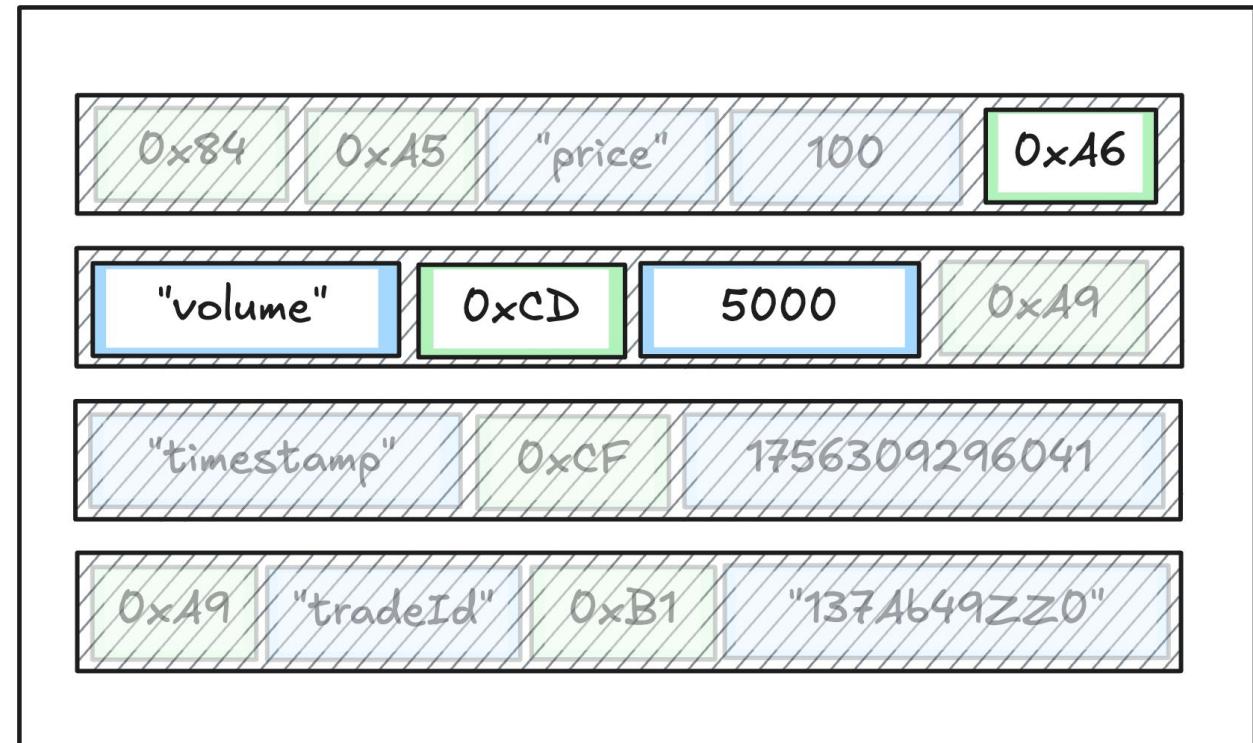
```
{  
    "price": 100,  
    "volume": 5000,  
    "timestamp": 1756309296041,  
    "tradeId": "137A649ZZ0"  
}
```



# Schemaless Formats

- An example document, encoded in both JSON and MsgPack:

```
{  
    "price": 100,  
    "volume    "timestamp": 1756309296041,  
    "tradeId": "137A649ZZ0"  
}
```



# Schemaless Formats

- The tagged formats often get creative about packing small values
- In MsgPack, there are various values that do not need a tag byte:
  - true is encoded as 0xc2 and false as 0xc3
  - Small integer values (less than 128) are simply encoded as themselves, without a tag
- As a result, binary schemaless formats are compact and elegant

# Schemaless Formats

- What is the cost of using a schemaless format?
- In order to read an encoded message, it typically gets decoded into some in-memory data structure
- Processing is then done on the in-memory data structure
- It is then encoded back into the serialized form

# Schemaless Formats

- What if I want to avoid the cost of always re-serializing the data?
- Maybe I want to pull a document out of a database and make some small, targeted modifications to it, and then write it back
- Or maybe I have some messages I'm receiving from the network, and I just want to pull one field out of each of them

# Schemaless Formats

- Can we build a schemaless serialization format that supports fast traversal and modification in-place?
- **Let's give it a try!**

# Baseline Serialization Format

- First of all, let's create a simple, unoptimized tagged serialization format
  - Similar to MsgPack, CBOR, etc.
- This will serve as a baseline for our benchmarks
- We'll build on top of it to create our more advanced versions

# Baseline Serialization Format

- Baseline serialization format:
  - Supports strings, ints, floats, maps, and arrays
  - Encoded as a sequence of tag, value pairs
  - Tag will encode the type (and the size, if necessary)
  - Values will be in little-endian (to be friendly to modern CPUs)
  - Maps will be a sequence of key, value pairs
  - Arrays will be a sequence of values

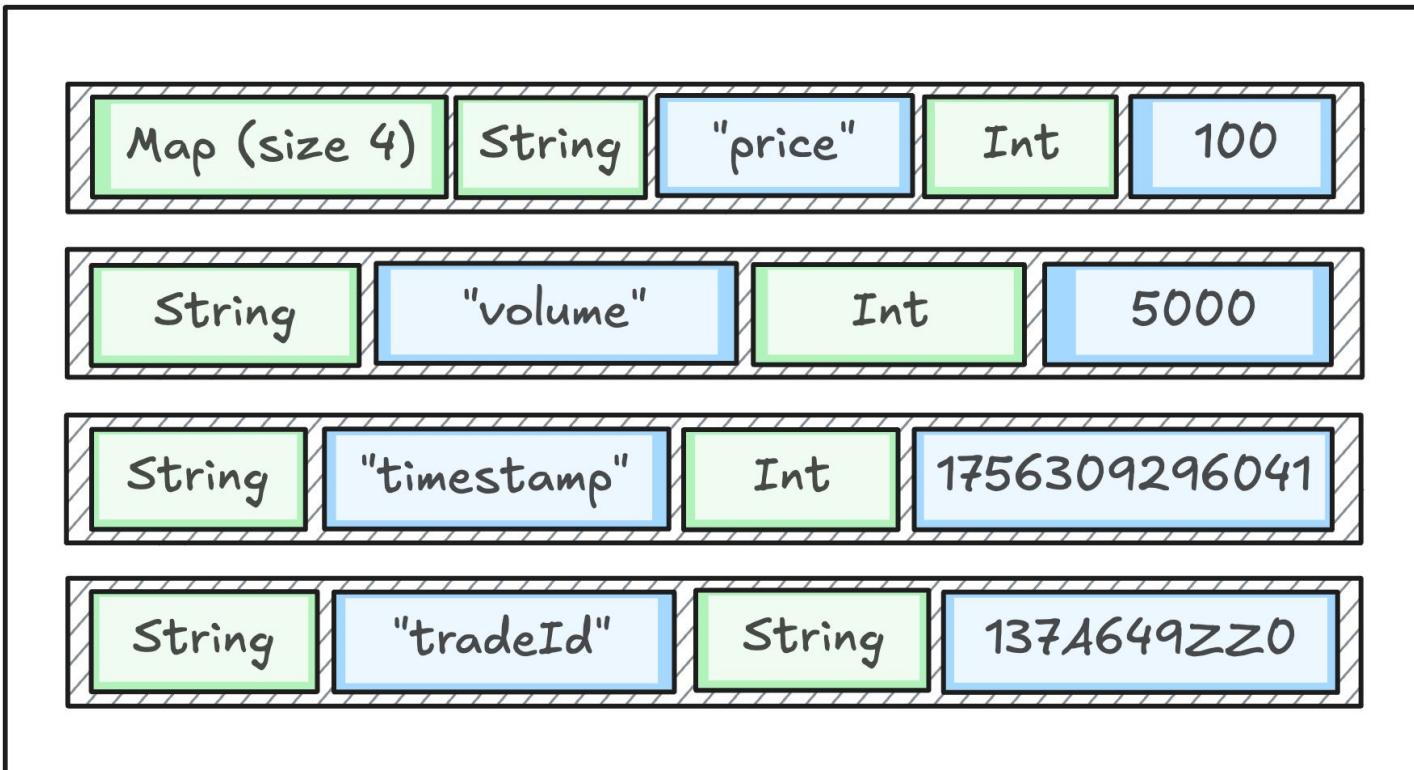
# Baseline Serialization Format

- Our interface will look something like this:

```
class Value {  
    float asFloat() const;                      // get this value as a float  
    int   asInt() const;                         // get this value as an int  
    ...                                           // etc...  
  
    void makeArray();                           // make this value an array  
    void makeMap();                            // make this value a map  
  
    Value operator=(auto value);                // assign the given value  
    Value append(auto value);                  // append a value to this array  
  
    Value set(std::string_view key, auto value); // set key to value in map  
  
    Value operator[](std::string_view key) const; // get value for key in map  
    Value operator[](int index) const;           // get value at index in array  
};
```

# Baseline Serialization Format

- To implement read operations, we will do a linear search through the document, looking for the right key



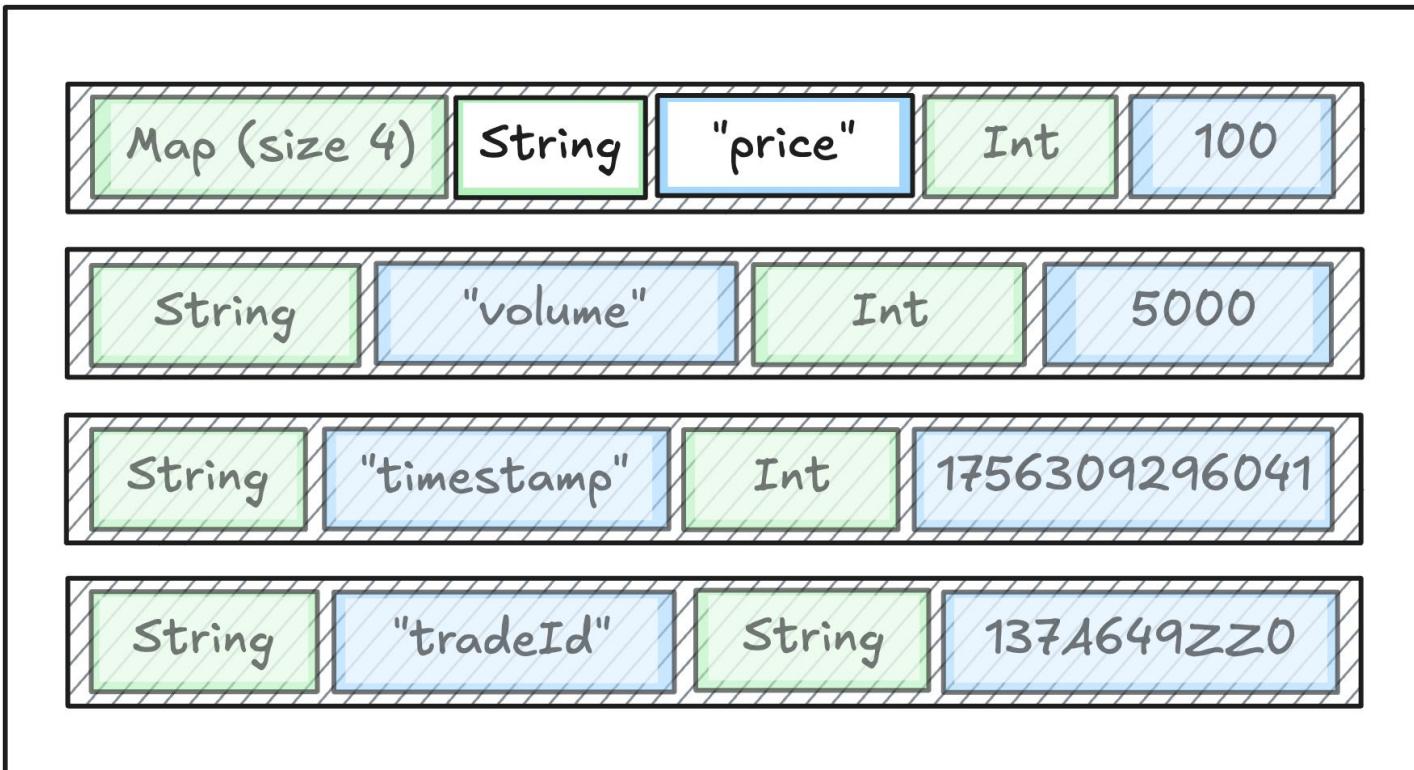
```
Value doc(buf);  
Value ts = doc["timestamp"];  
int64_t value = ts.toInt();
```

Bloomberg

Engineering

# Baseline Serialization Format

- To implement read operations, we will do a linear search through the document, looking for the right key



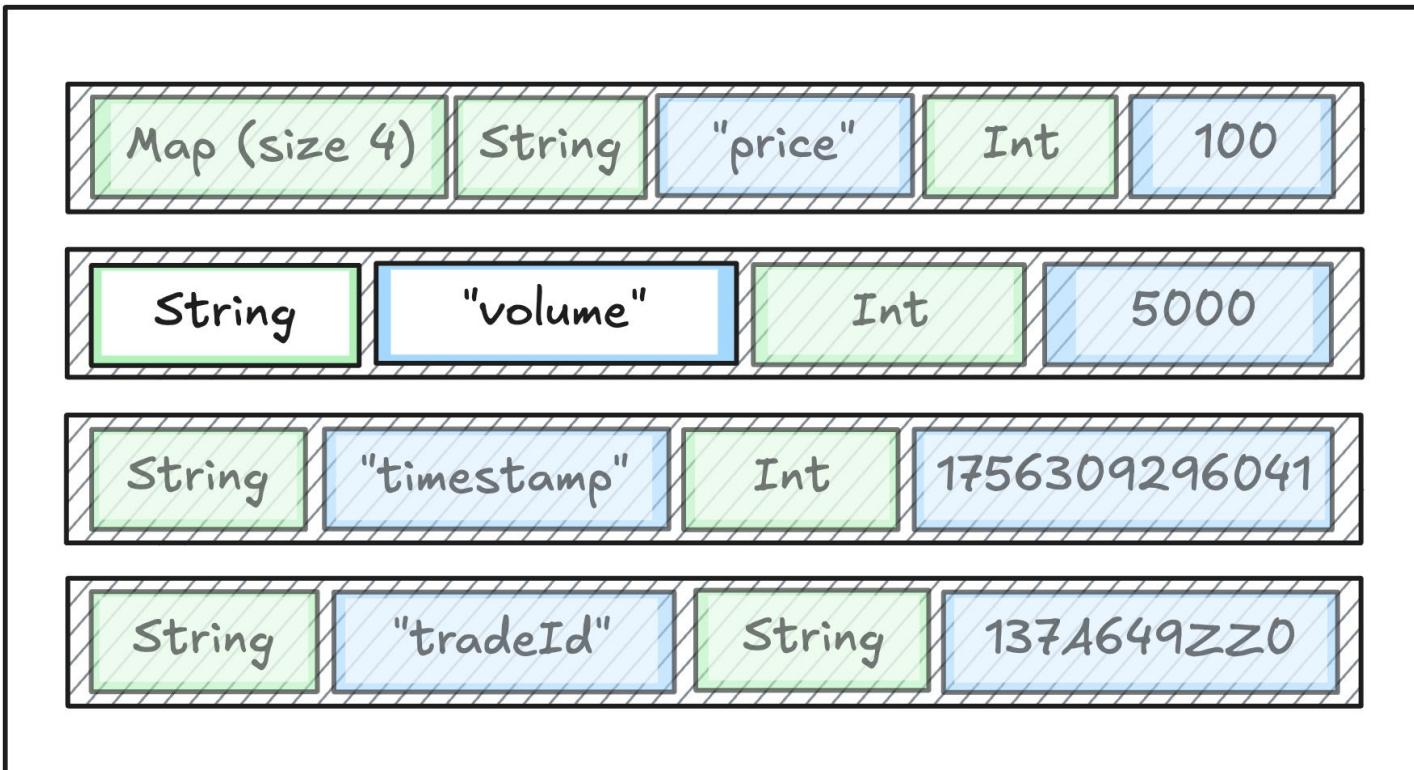
```
Value doc(buf);  
Value ts = doc["timestamp"];  
int64_t value = ts.toInt();
```

Bloomberg

Engineering

# Baseline Serialization Format

- To implement read operations, we will do a linear search through the document, looking for the right key



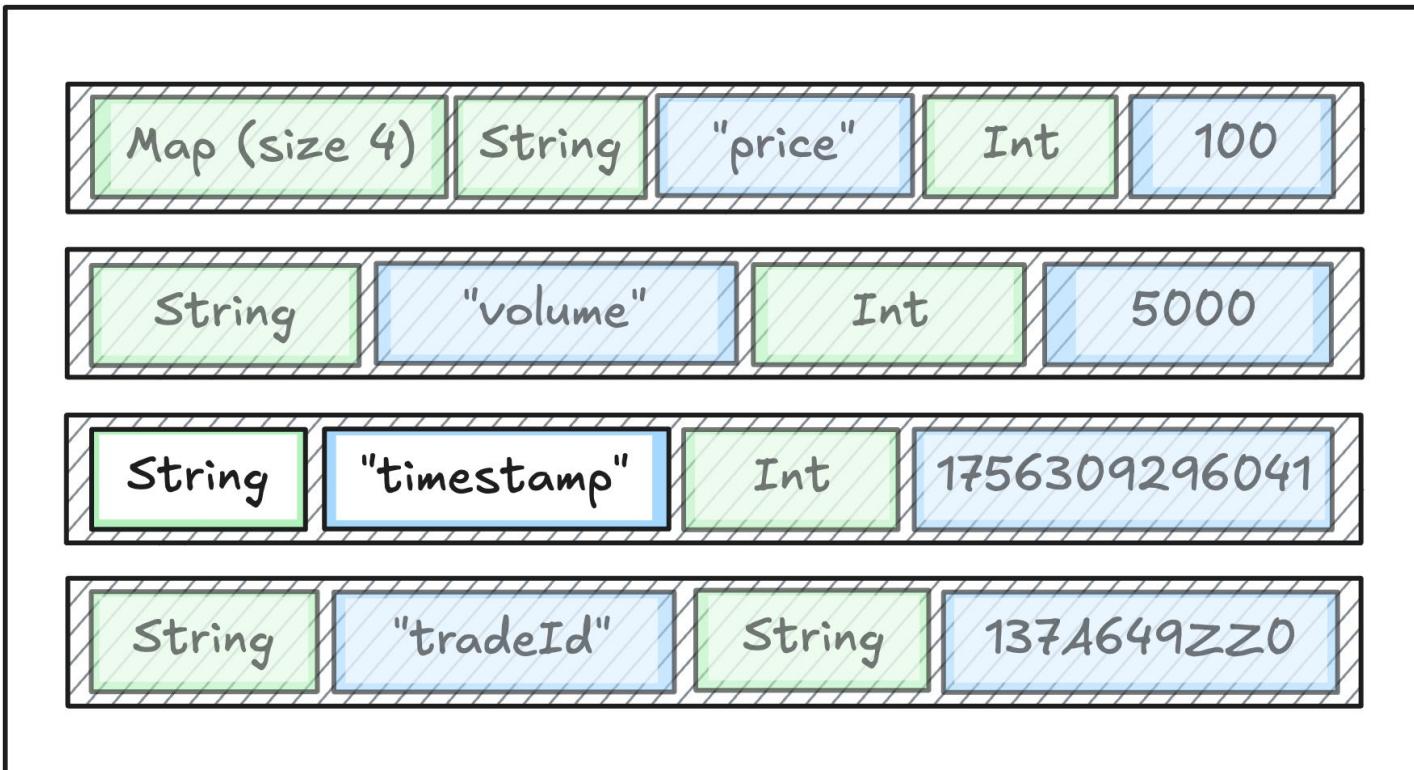
```
Value doc(buf);  
Value ts = doc["timestamp"];  
int64_t value = ts.toInt();
```

Bloomberg

Engineering

# Baseline Serialization Format

- To implement read operations, we will do a linear search through the document, looking for the right key



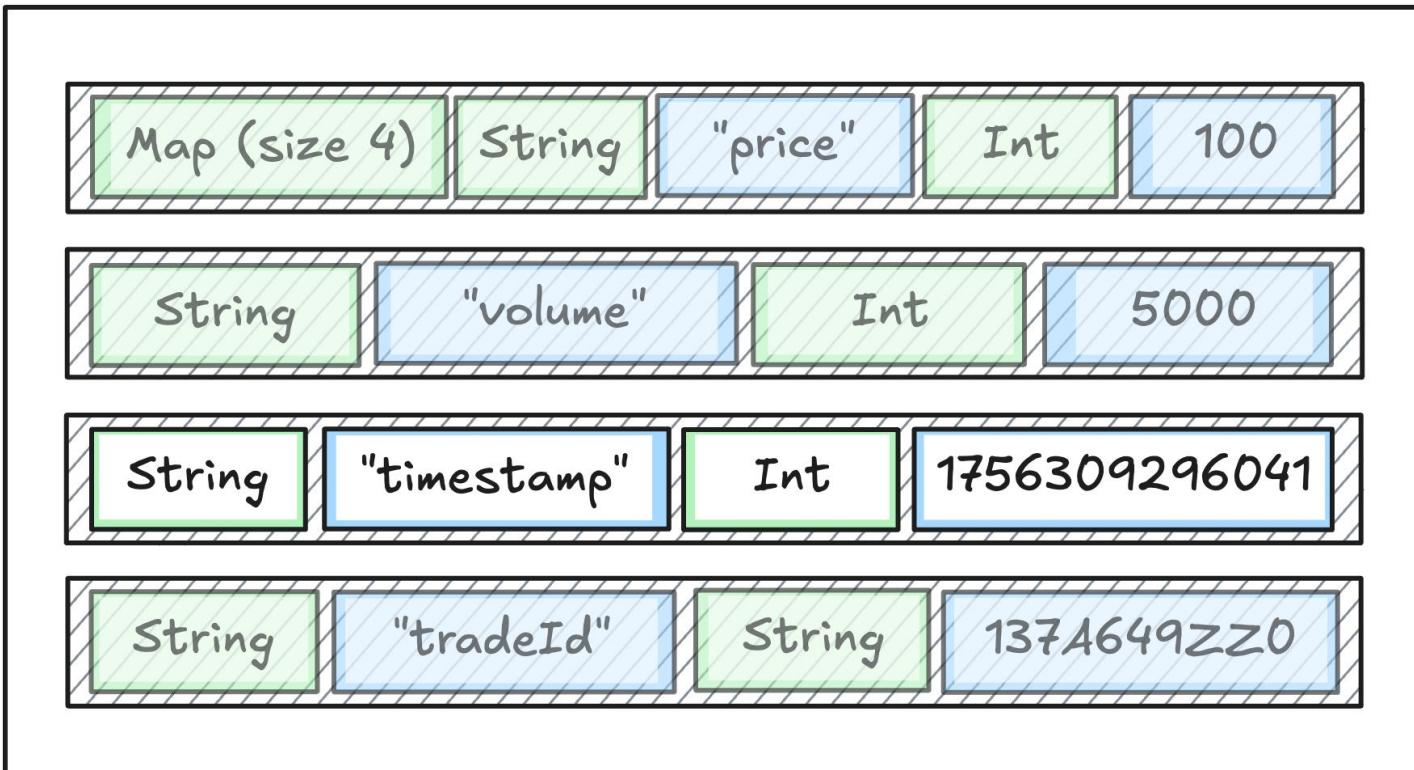
```
Value doc(buf);  
Value ts = doc["timestamp"];  
int64_t value = ts.toInt();
```

Bloomberg

Engineering

# Baseline Serialization Format

- To implement read operations, we will do a linear search through the document, looking for the right key



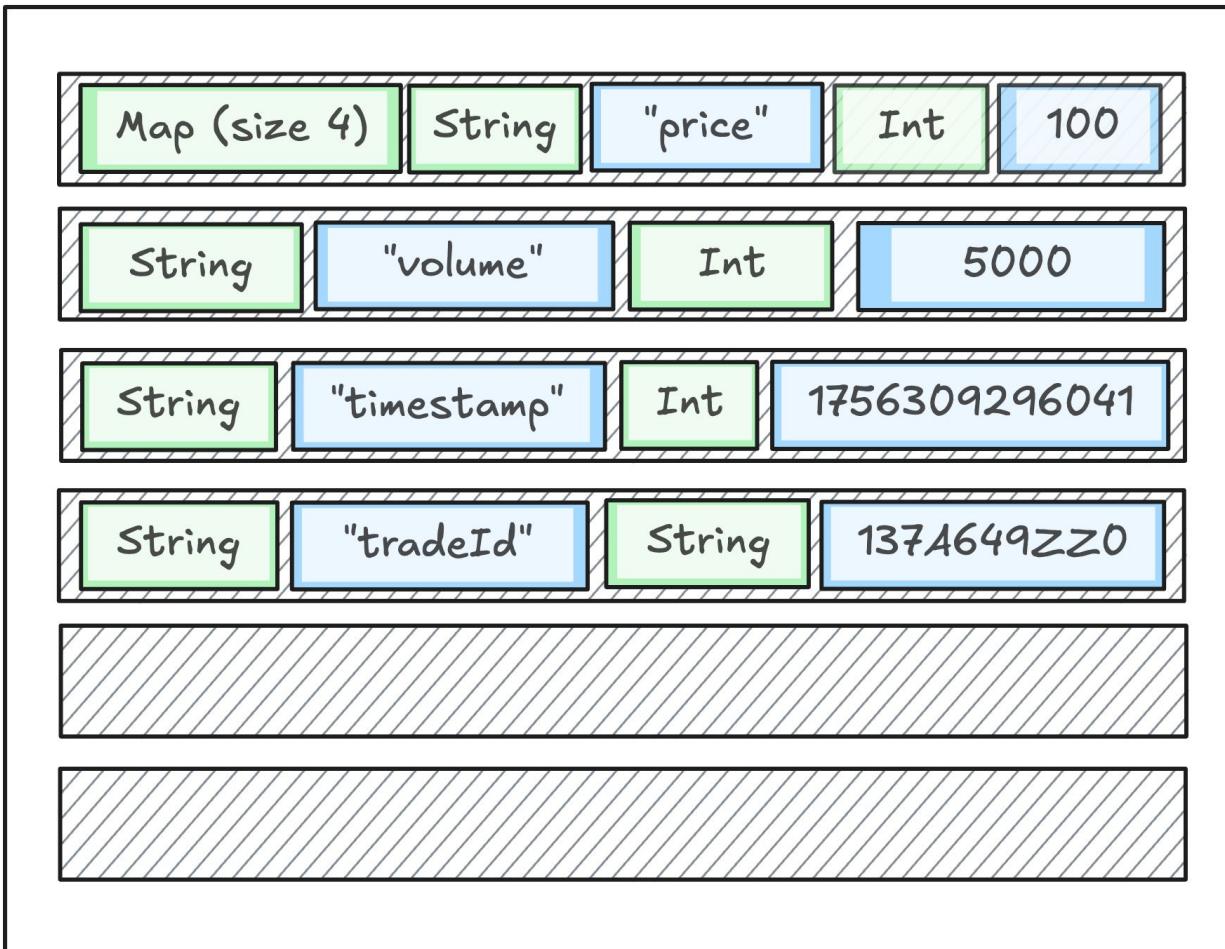
```
Value doc(buf);  
Value ts = doc["timestamp"];  
int64_t value = ts.toInt();
```

Bloomberg

Engineering

# Baseline Serialization Format

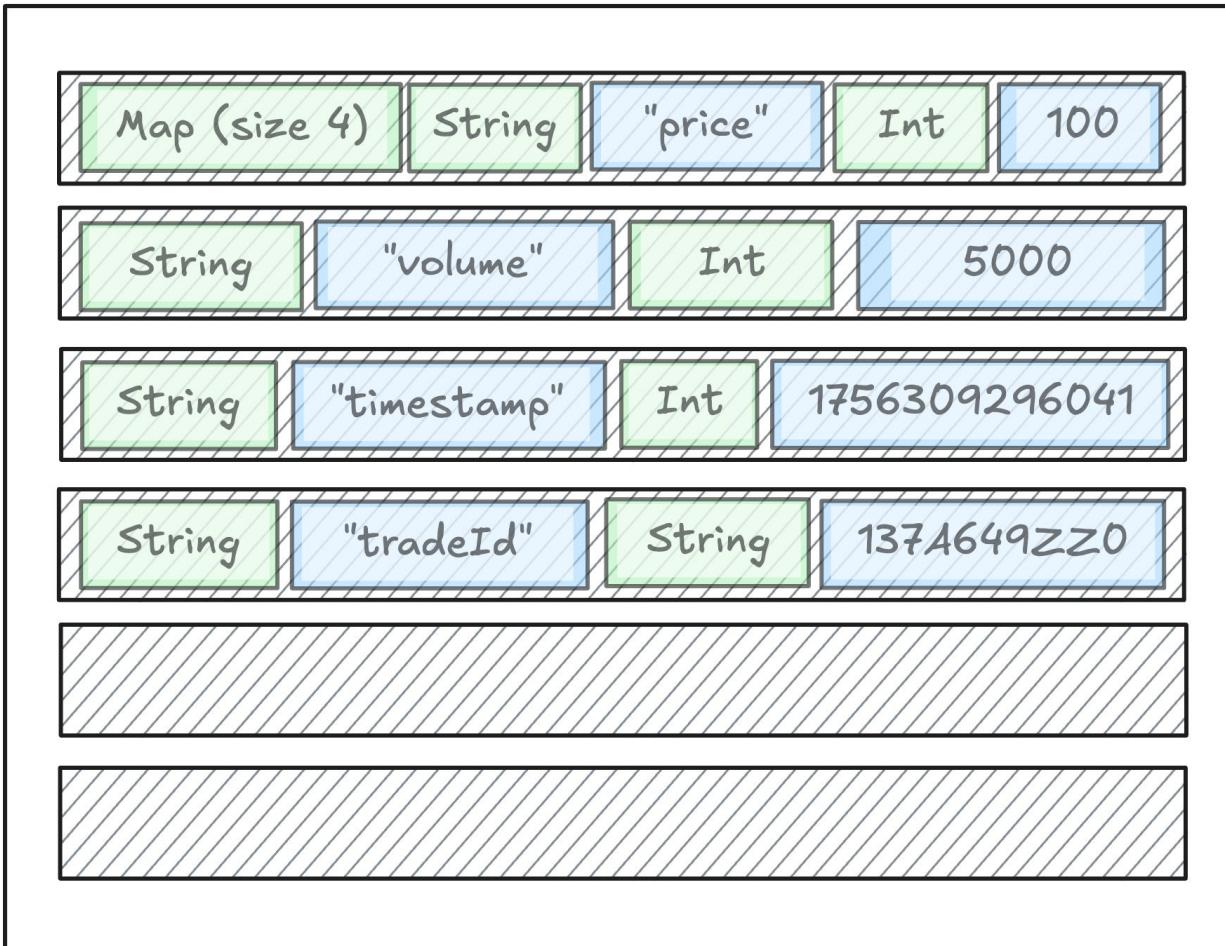
- To make modifications, we will modify in-place, and move everything that comes after to make space



```
Value doc(buf);
Value price = doc["price"];
price.makeMap();
price.set("mantissa", 100);
price.set("exponent", -2);
```

# Baseline Serialization Format

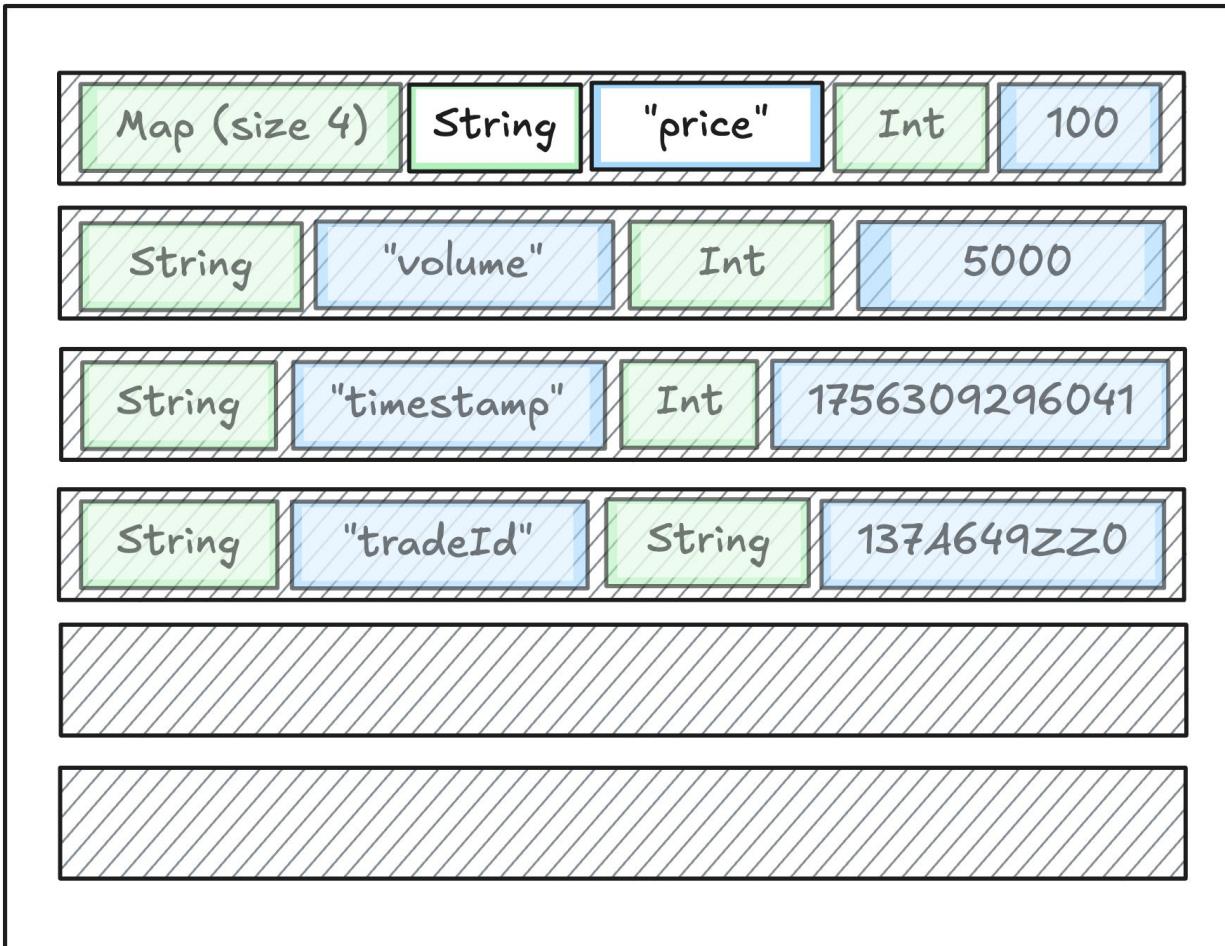
- To make modifications, we will modify in-place, and move everything that comes after to make space



```
Value doc(buf);
Value price = doc["price"];
price.makeMap();
price.set("mantissa", 100);
price.set("exponent", -2);
```

# Baseline Serialization Format

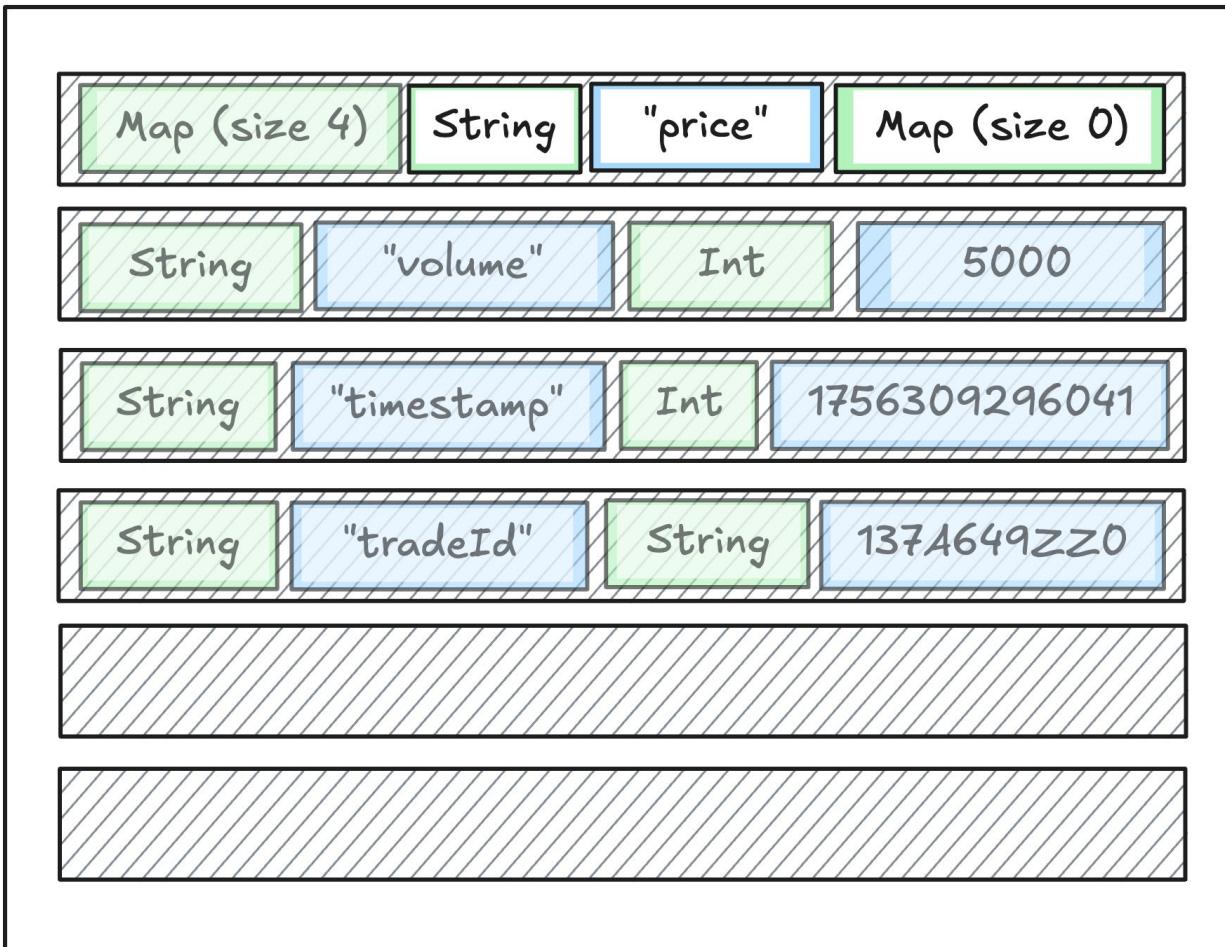
- To make modifications, we will modify in-place, and move everything that comes after to make space



```
Value doc(buf);
Value price = doc["price"];
price.makeMap();
price.set("mantissa", 100);
price.set("exponent", -2);
```

# Baseline Serialization Format

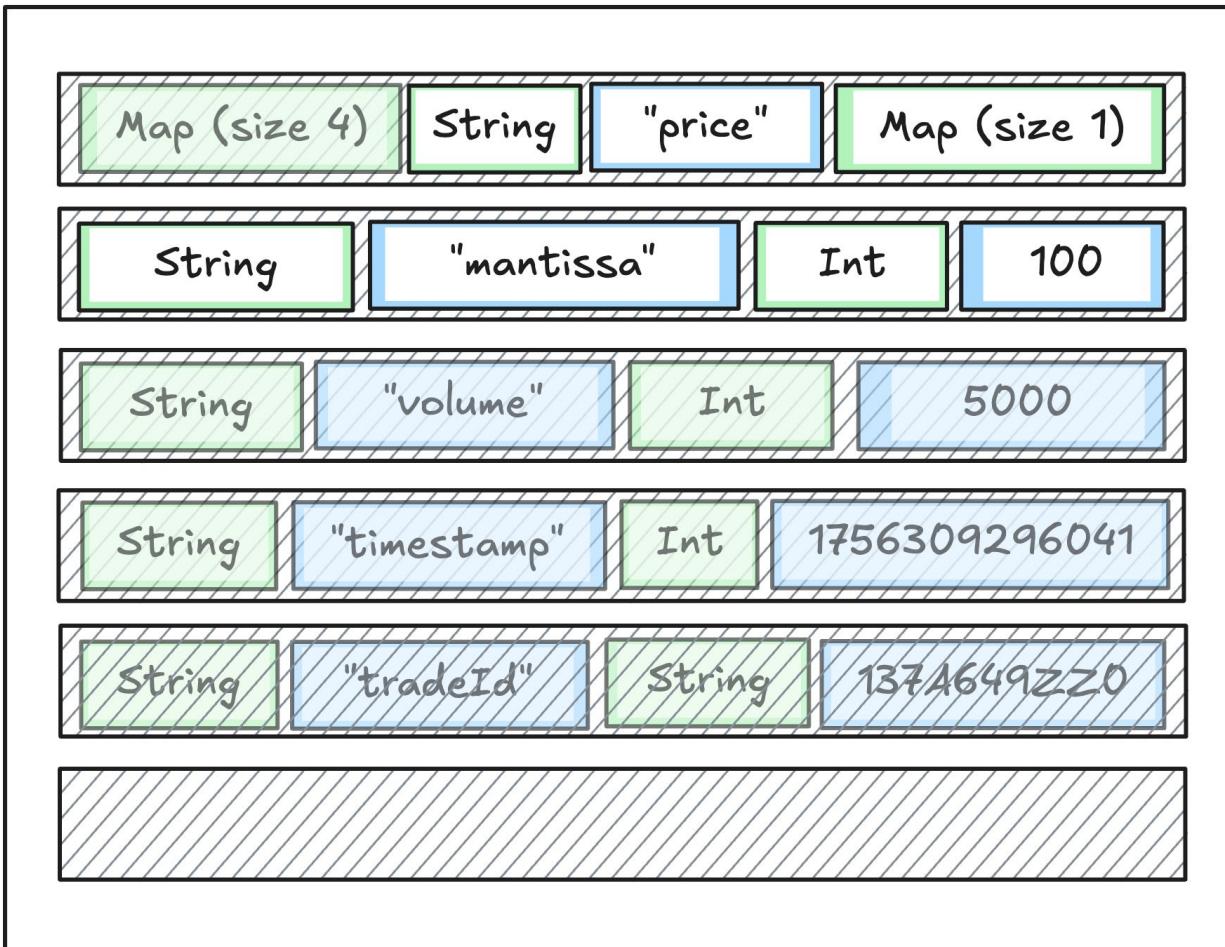
- To make modifications, we will modify in-place, and move everything that comes after to make space



```
Value doc(buf);
Value price = doc["price"];
price.makeMap();
price.set("mantissa", 100);
price.set("exponent", -2);
```

# Baseline Serialization Format

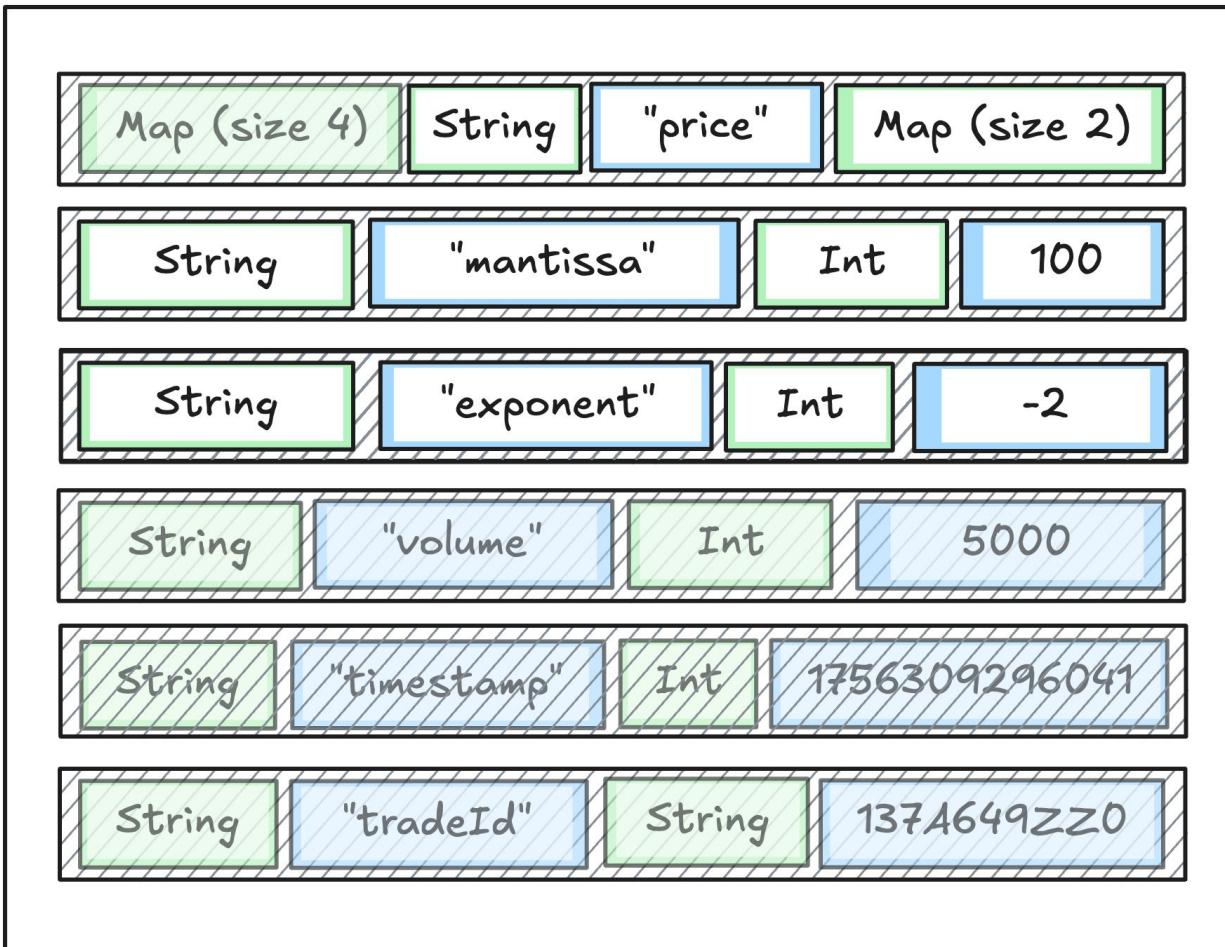
- To make modifications, we will modify in-place, and move everything that comes after to make space



```
Value doc(buf);
Value price = doc["price"];
price.makeMap();
price.set("mantissa", 100);
price.set("exponent", -2);
```

# Baseline Serialization Format

- To make modifications, we will modify in-place, and move everything that comes after to make space



```
Value doc(buf);
Value price = doc["price"];
price.makeMap();
price.set("mantissa", 100);
price.set("exponent", -2);
```

# Baseline Serialization Format – Benchmarks

- So, how does this perform?
- Let's create a simple benchmark
- We'll parameterize our tests on:
  - Number of fields in document
  - Number of operations to perform
  - Percentage of operations that are:
    - Reads
    - Writes to existing values
    - Structural modifications (add or remove value)

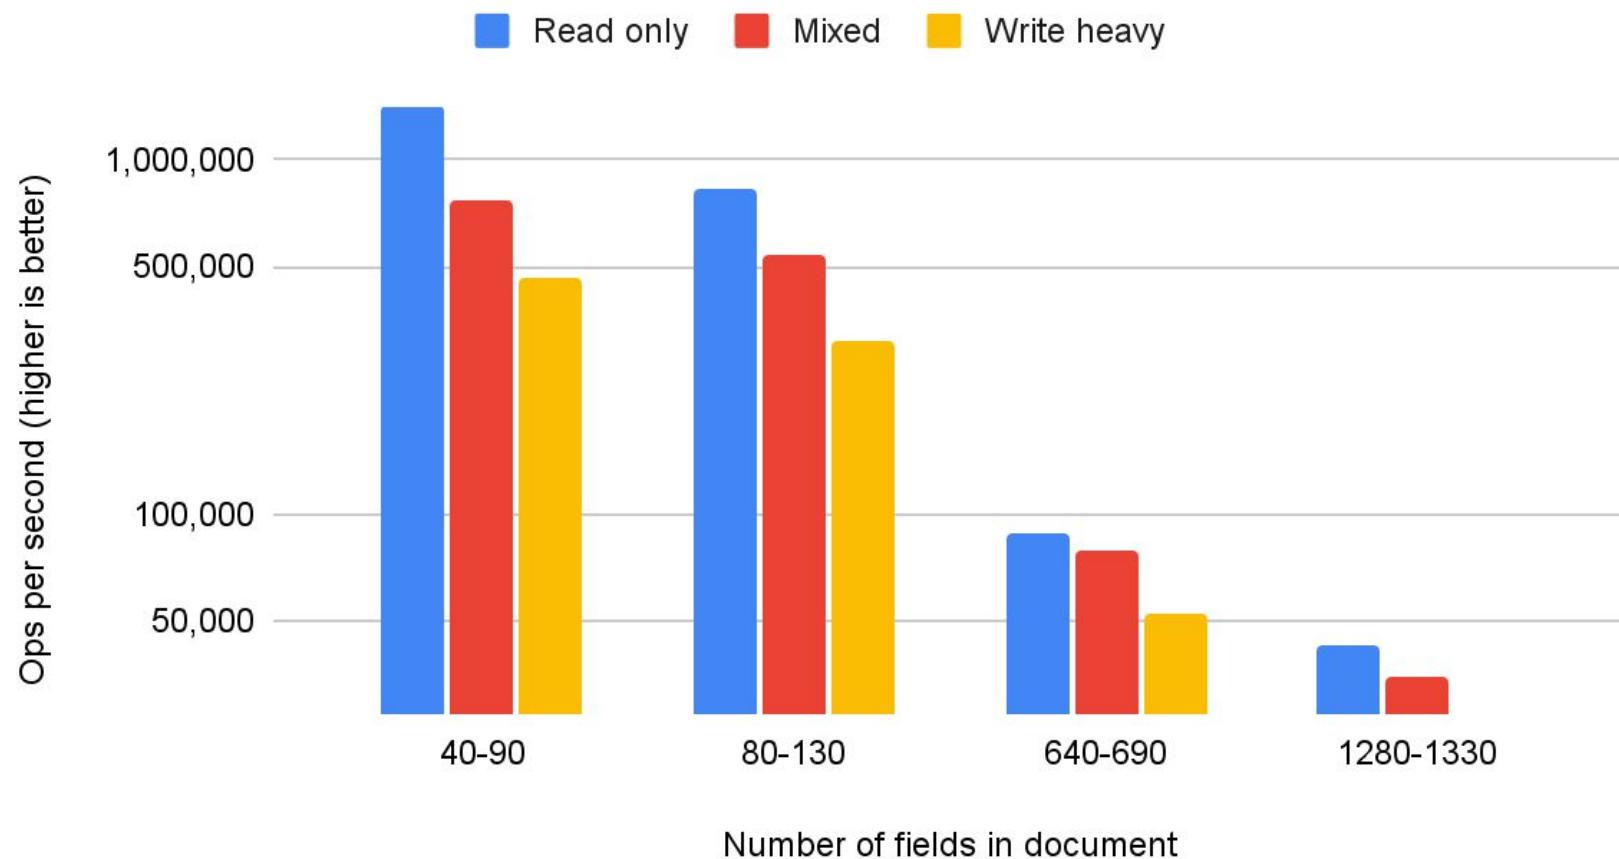
# Benchmarking Setup

- Benchmark environment:
  - Linux 3.10
  - 72-core Intel Xeon Gold 6240 @ 2.6 GHz (Cascade Lake)
  - gcc 13.1.1 20230614 (Red Hat 13.1.1-4)
  - -O2 -march=native -std=c++23
- Machine load < 1%
- These benchmarks are synthetic and don't necessarily represent the results you'd get in a production workload

# Baseline Serialization Format – Benchmarks

- Results

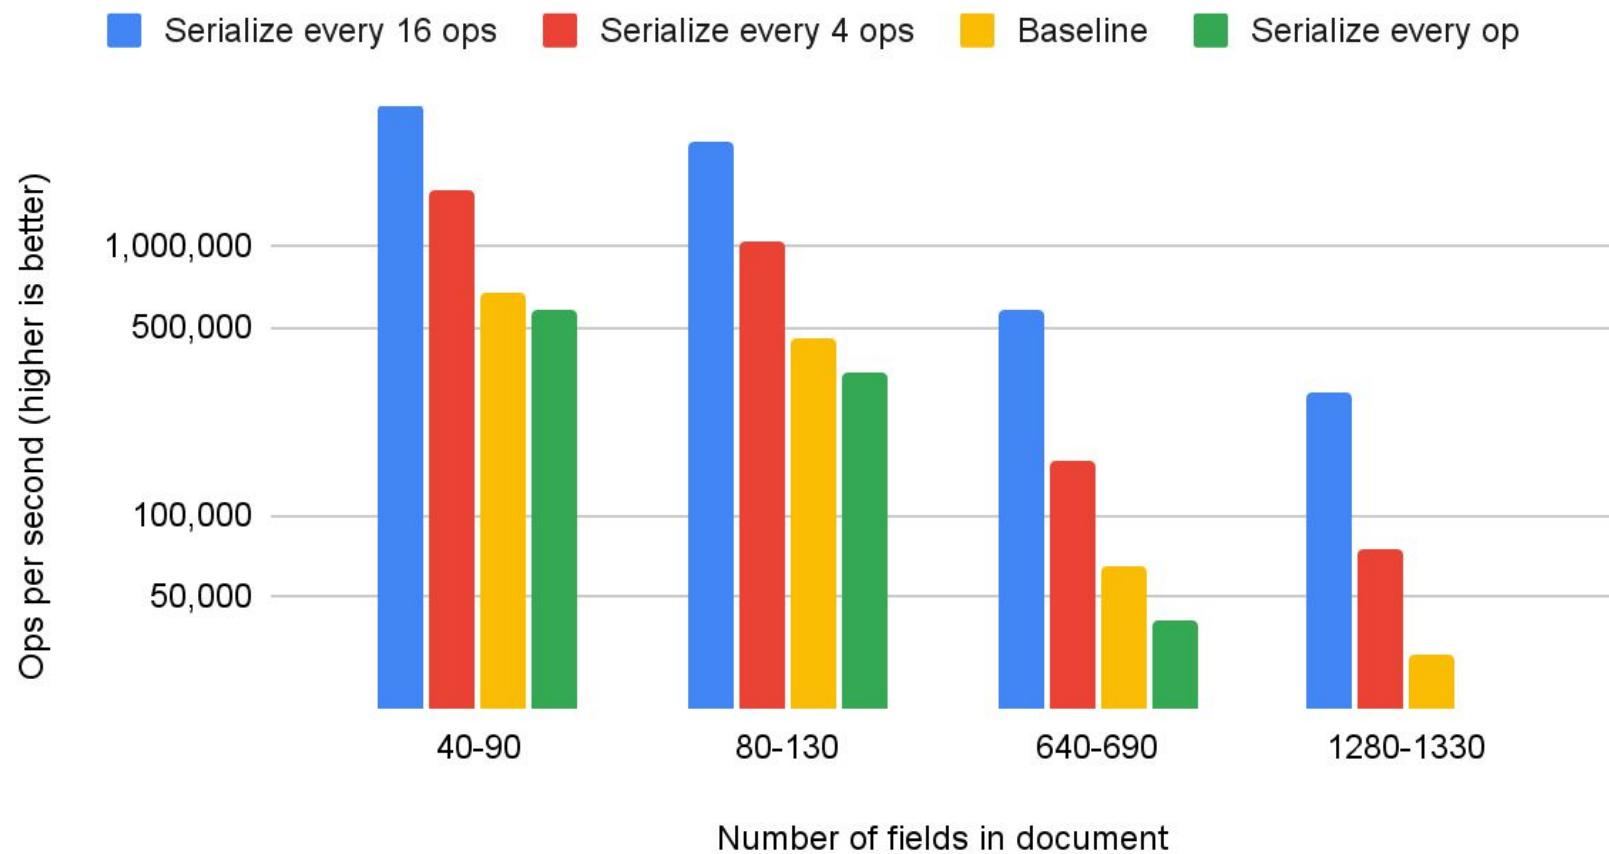
Performance on different work loads



# Baseline Serialization Format – Benchmarks

- Results vs. operating on an unordered\_map and re-serializing

Performance (mixed workload)



# Faster Searches

How do we make searches faster?

[TechAtBloomberg.com](https://TechAtBloomberg.com)

© 2025 Bloomberg Finance L.P. All rights reserved.

Bloomberg  
Engineering

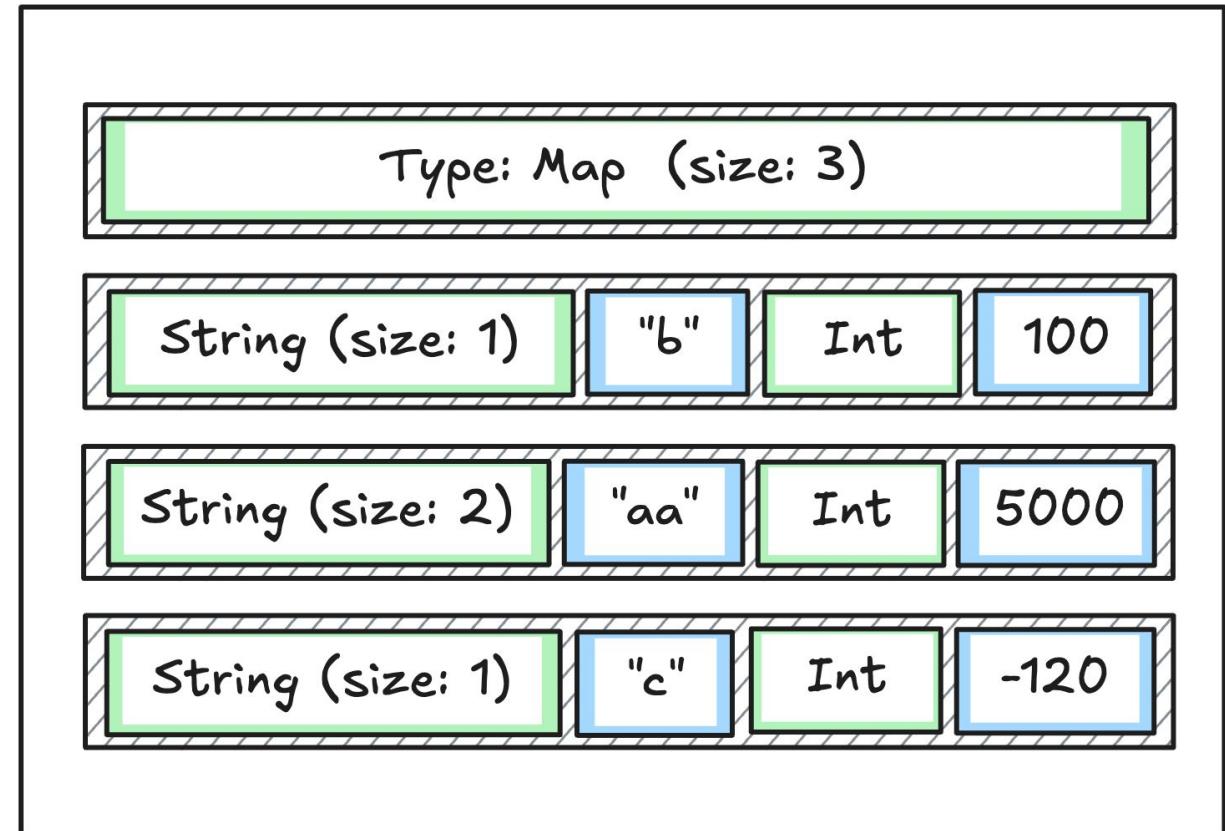
# Profiling

- Profiling reveals that most of the time is spent checking the size of fields while searching for elements of interest:

Samples: 18K of event 'cycles:uppp', Event count (approx.): 15146430238			
Overhead	Command	Shared Object	Symbol
41.85%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::tagged::TaggedUtil::sizeOfMap
29.72%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::tagged::TaggedUtil::sizeOfValue
7.89%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::tagged::TaggedUtil::mapElementAt
3.15%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::tagged::Value::operator[]
3.05%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::adv::container::performOperationC
1.82%	adv_container_b	libc-2.17.so	[.] __memcmp_sse4_1
1.75%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::adv::ContainerBenchmark::prepareC
1.67%	adv_container_b	adv_container_benchmarks	[.] std::mersenne_twister_engine<unsigned long, 32ul, 62
1.16%	adv_container_b	adv_container_benchmarks	[.] std::uniform_int_distribution<char>::operator()<std::char_traits<char>, std::random_device::result_type>
0.92%	adv_container_b	libc-2.17.so	[.] __memmove_ssse3_back
0.82%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::adv::ContainerBenchmark::setupIni
0.53%	adv_container_b	libc-2.17.so	[.] _int_malloc
0.49%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::adv::ContainerBenchmark::getUnuse
0.48%	adv_container_b	adv_container_benchmarks	[.] experimental::ser::adv::ContainerBenchmark::generate

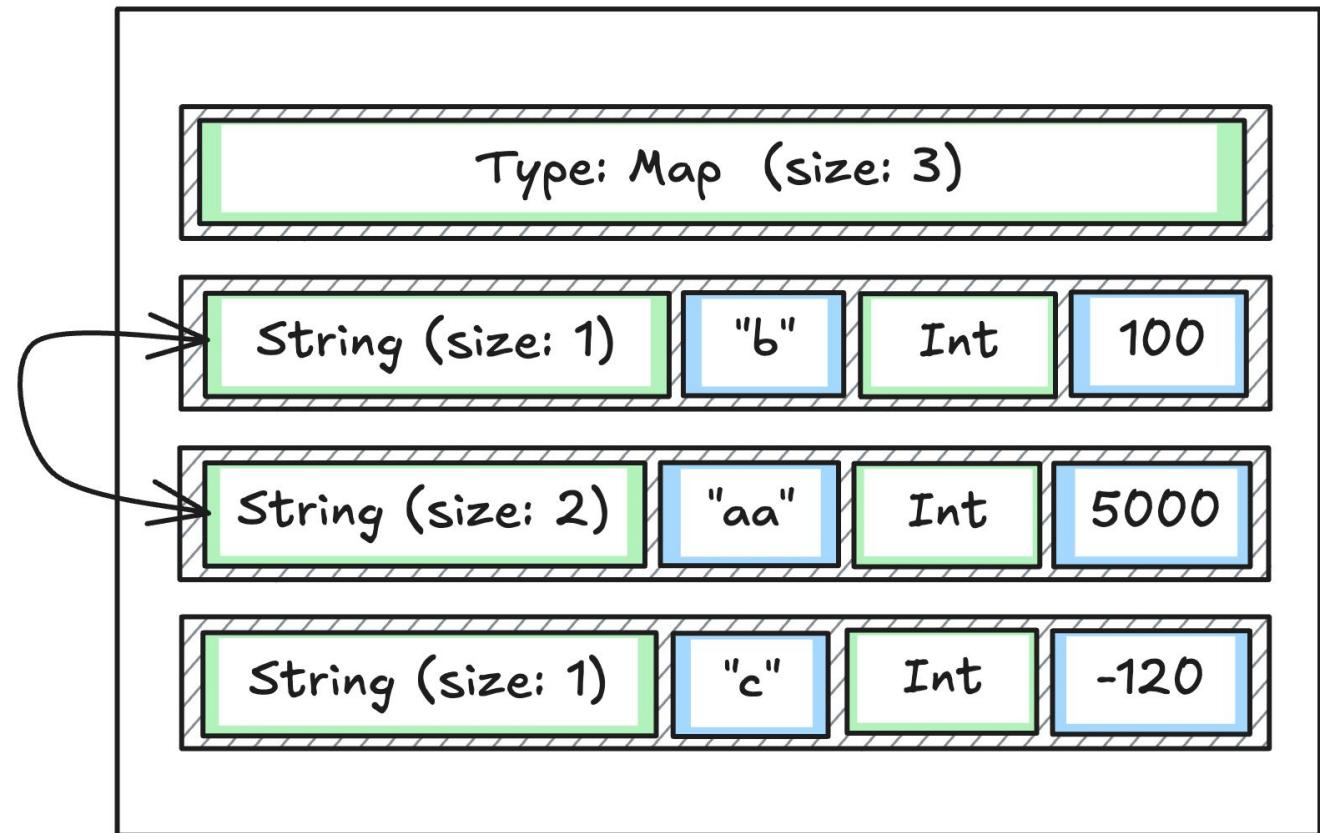
# Faster Searches

- A lot of time is spent skipping over elements
- The key we are searching for could appear anywhere in the map
- What if we put our keys in sorted order?



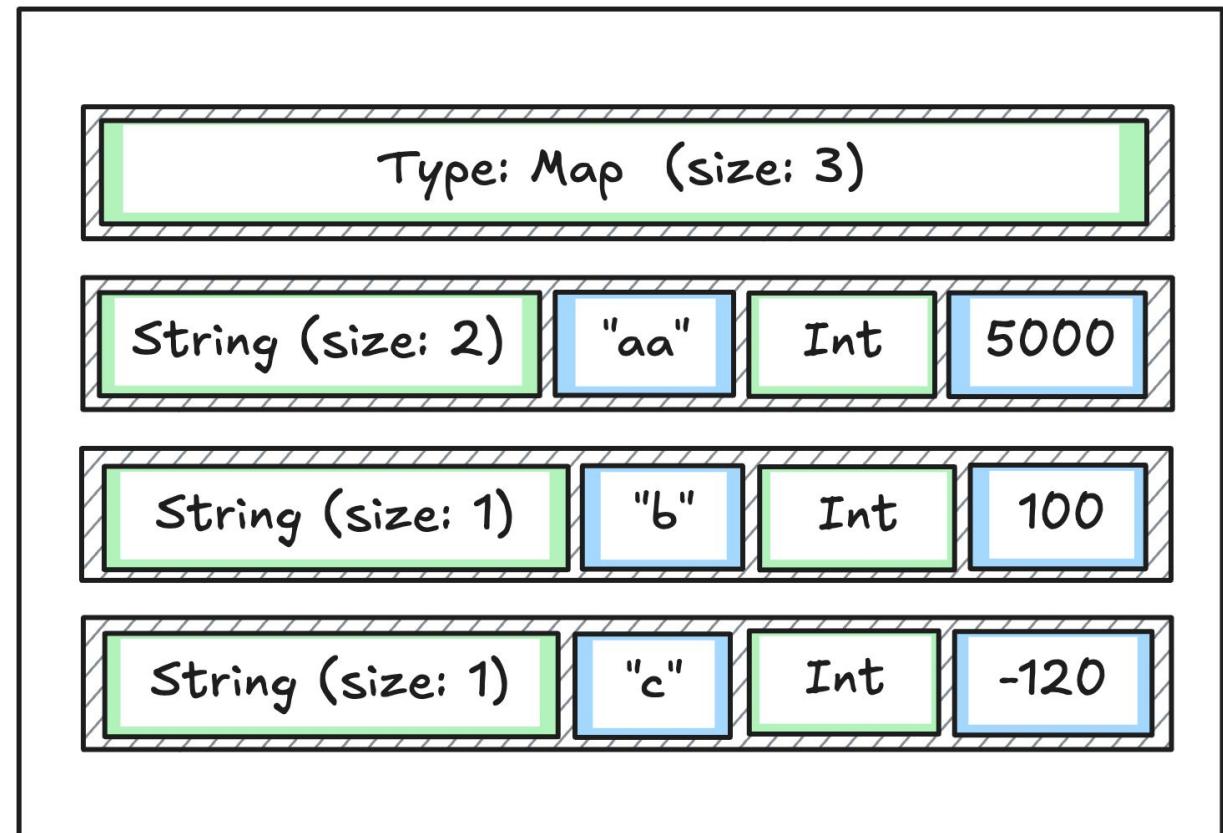
# Faster Searches

- A lot of time is spent skipping over elements
- The key we are searching for could appear anywhere in the map
- What if we put our keys in sorted order?



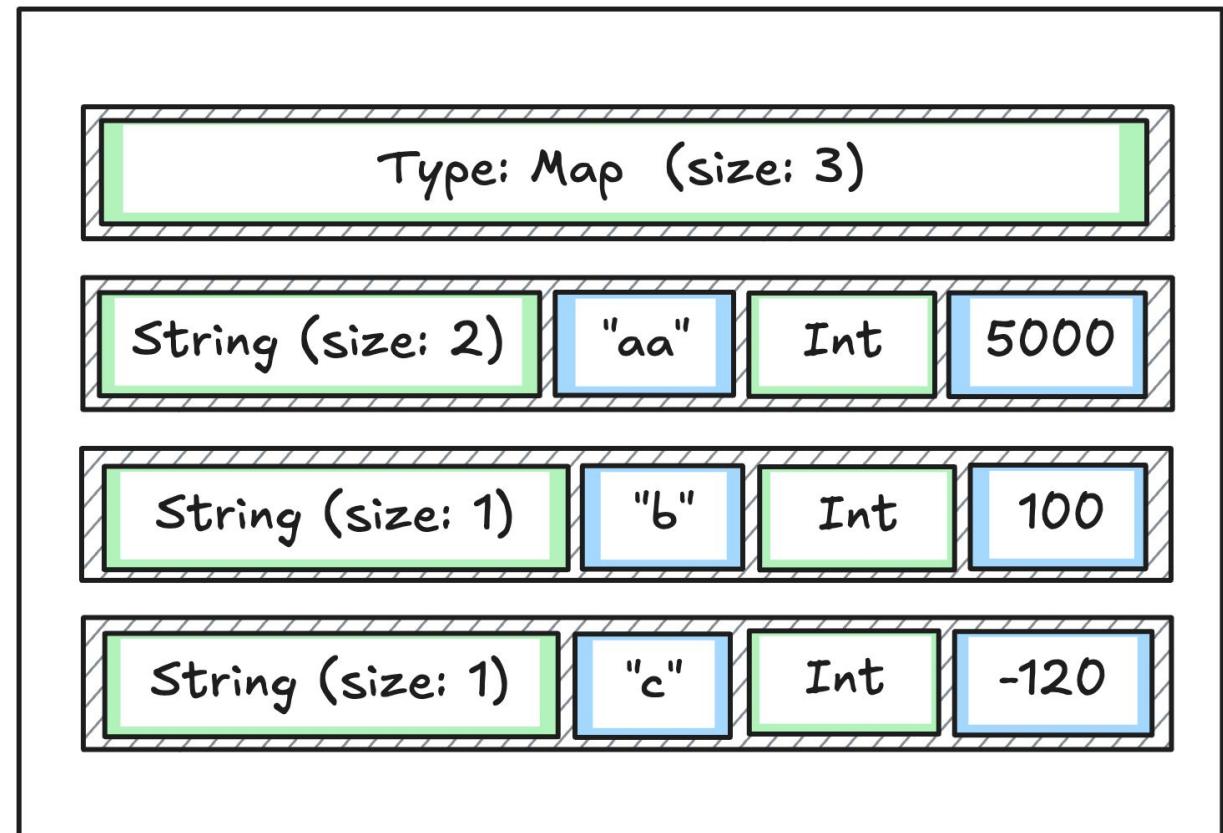
# Faster Searches

- A lot of time is spent skipping over elements
- The key we are searching for could appear anywhere in the map
- What if we put our keys in sorted order?



# Faster Searches

- A lot of time is spent skipping over elements
- The key we are searching for could appear anywhere in the map
- What if we put our keys in sorted order?
- Now can we do binary search?
  - Not quite, because our keys and values are variable sized
  - This makes it hard to jump to the middle key

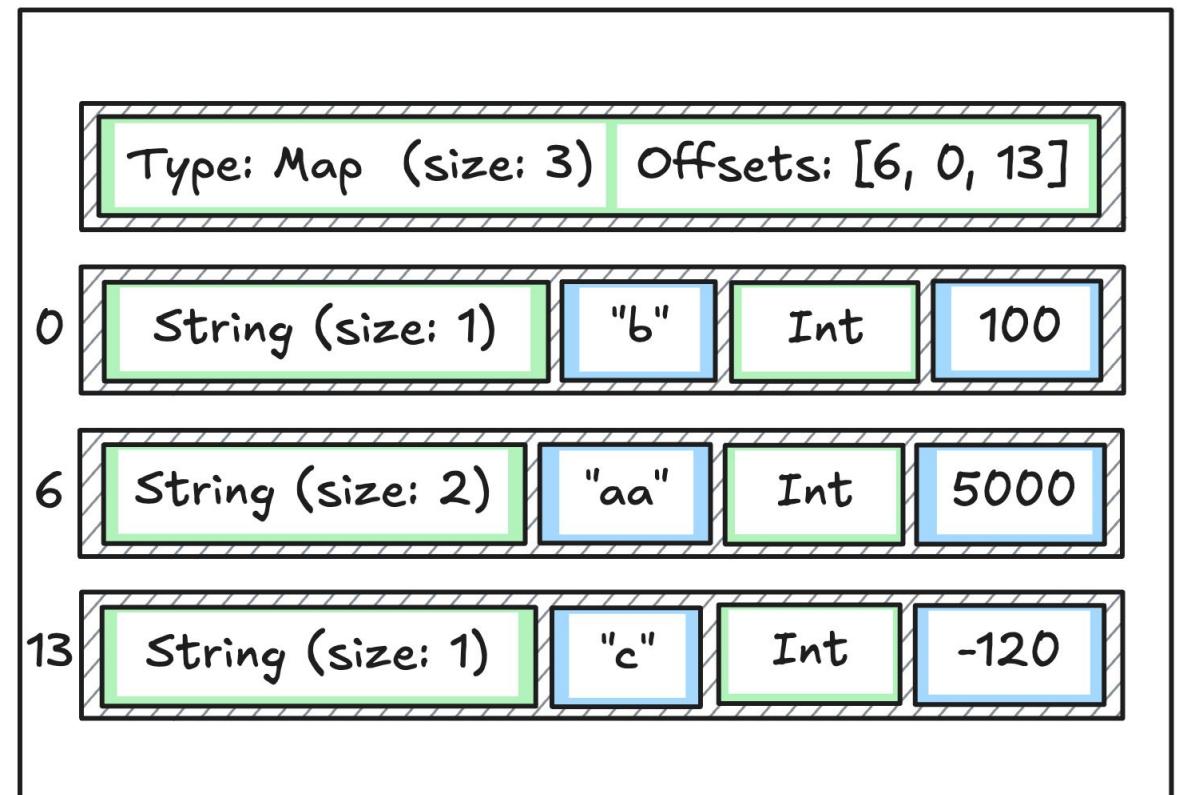


# Faster Searches

**Idea 1: Store a table of offsets at the start of the map**

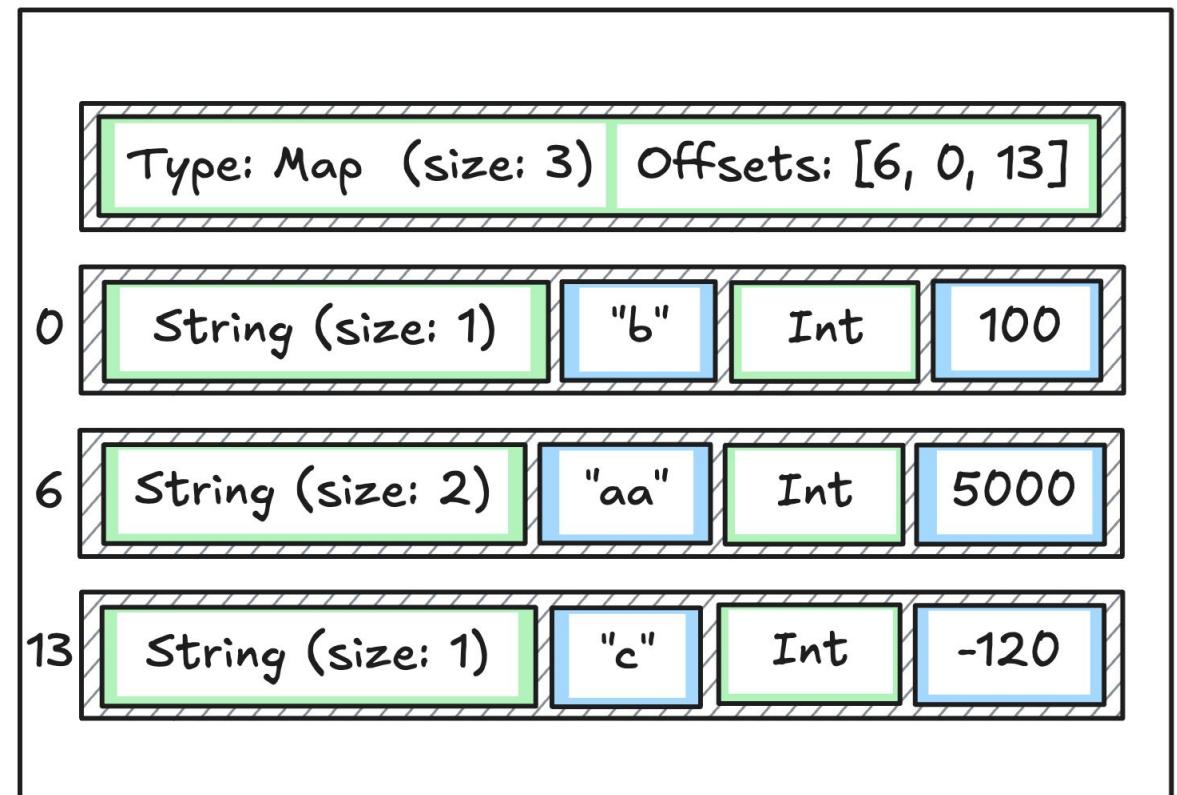
# Faster Searches

- Insert an offset table at the start of the map
- This makes it easy to jump to certain elements
- The actual elements don't need to be in sorted order, only the offsets



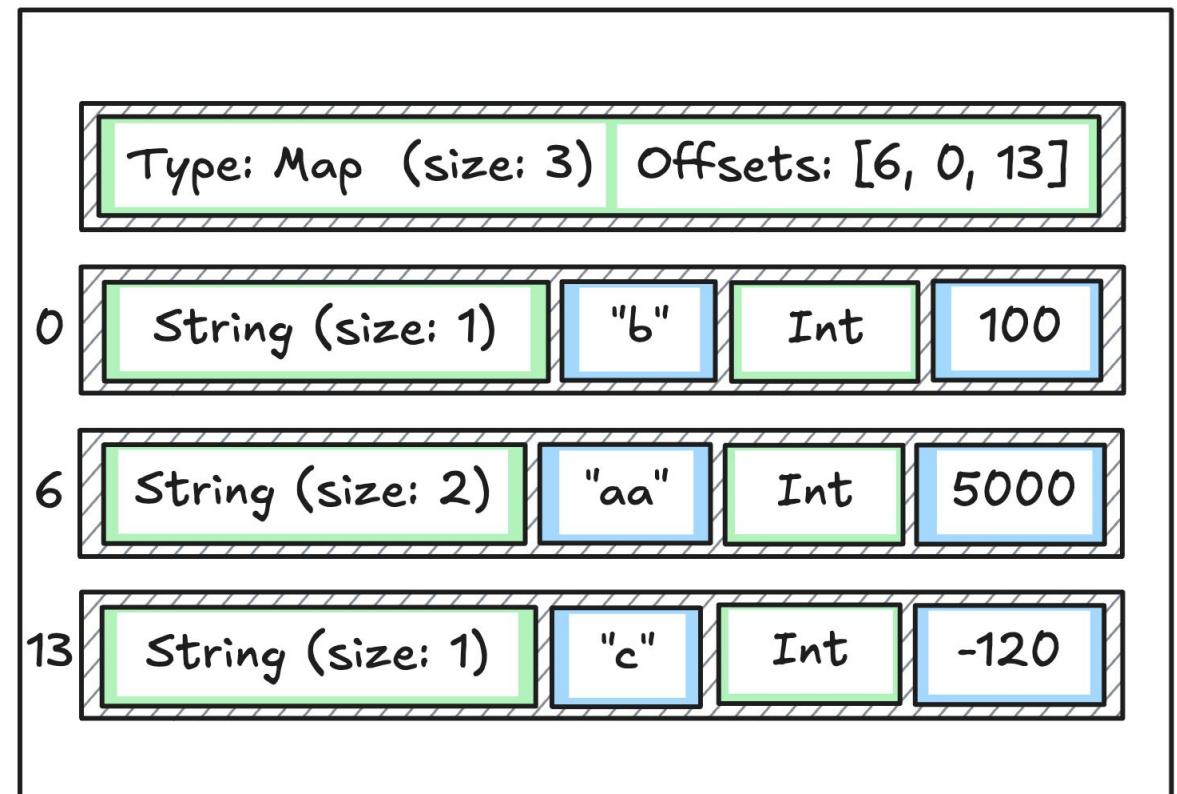
# Faster Searches

- Insert an offset table at the start of the map
- This makes it easy to jump to certain elements
- The actual elements don't need to be in sorted order, only the offsets
- In this example, the offsets tell us that key "aa" comes first, followed by "b", followed by "c"



# Faster Searches

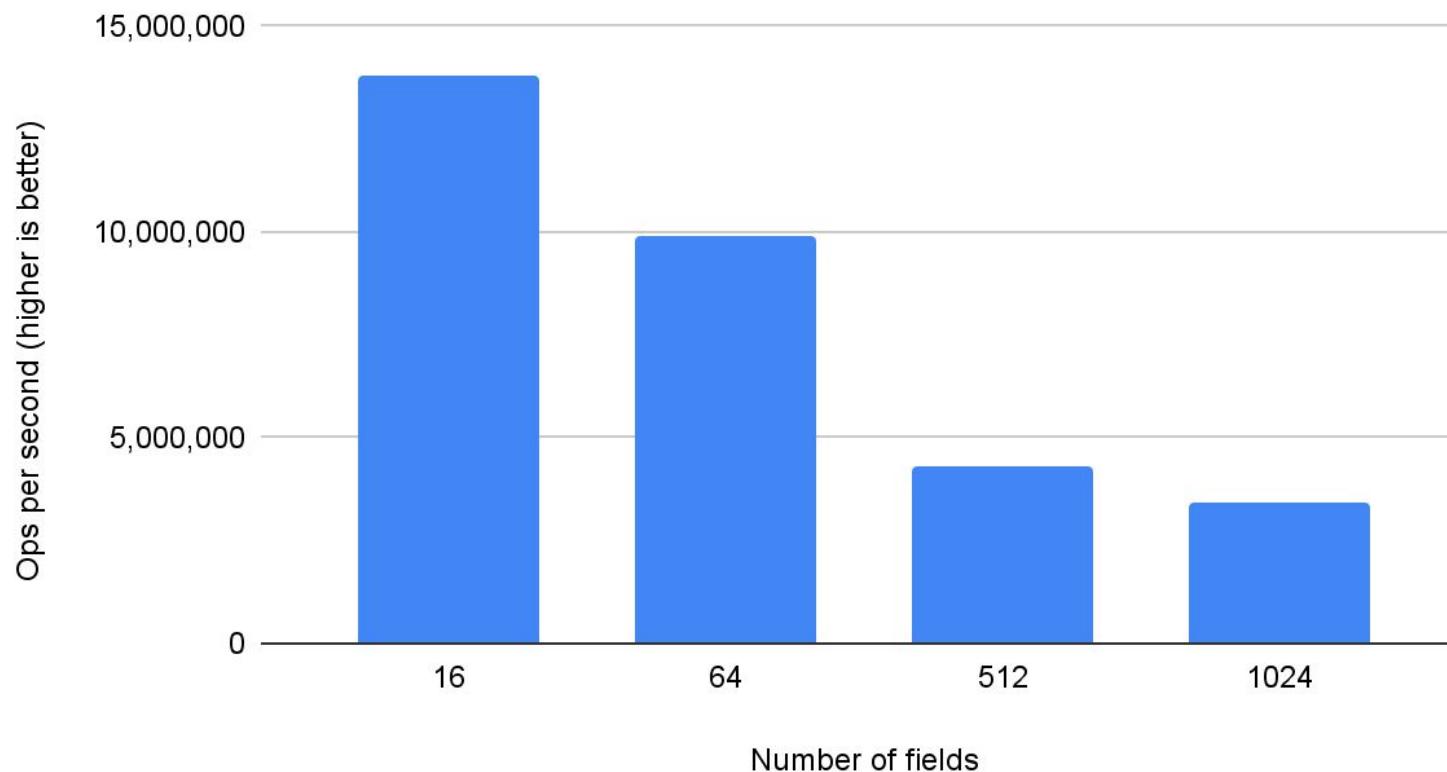
- Insert an offset table at the start of the map
- This makes it easy to jump to certain elements
- The actual elements don't need to be in sorted order, only the offsets
- In this example, the offsets tell us that key "aa" comes first, followed by "b", followed by "c"
- Now it's easier to do binary search



# Faster Searches – Benchmark Results

- Benchmark results for binary search on offsets:

Read-only workload with offset-based binary search



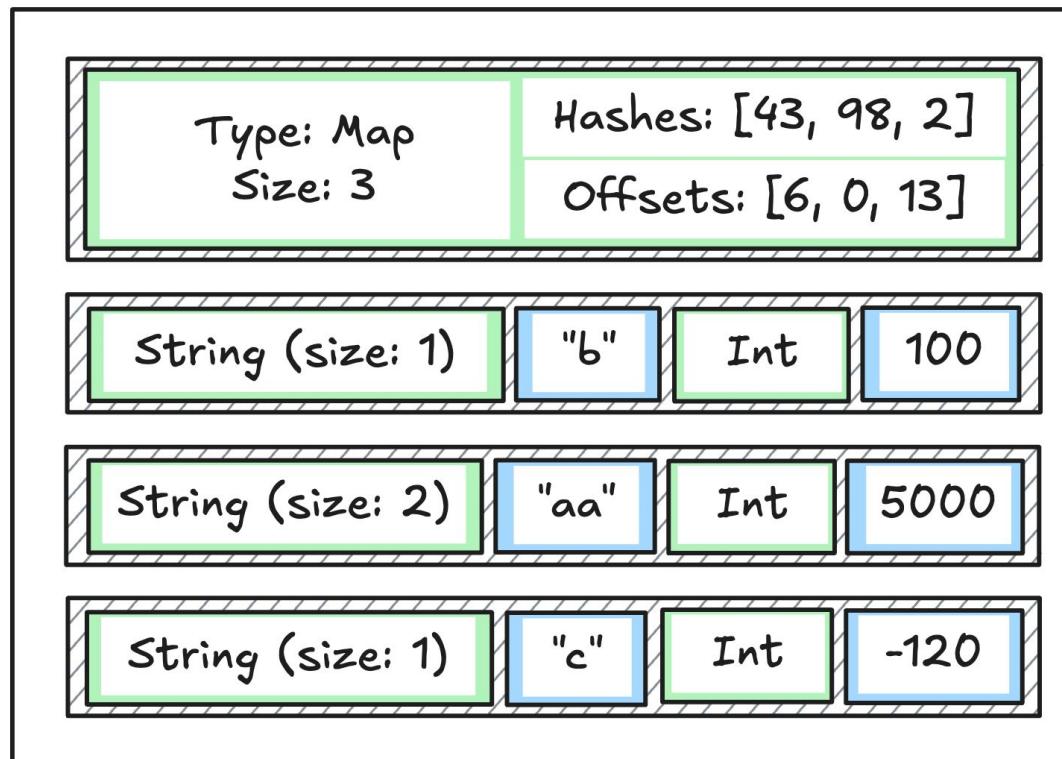
# Faster Searches

Idea 1: Store a table of offsets at the start of the map

**Idea 2: Store hashes of keys at the start of the map too**

# Faster Searches

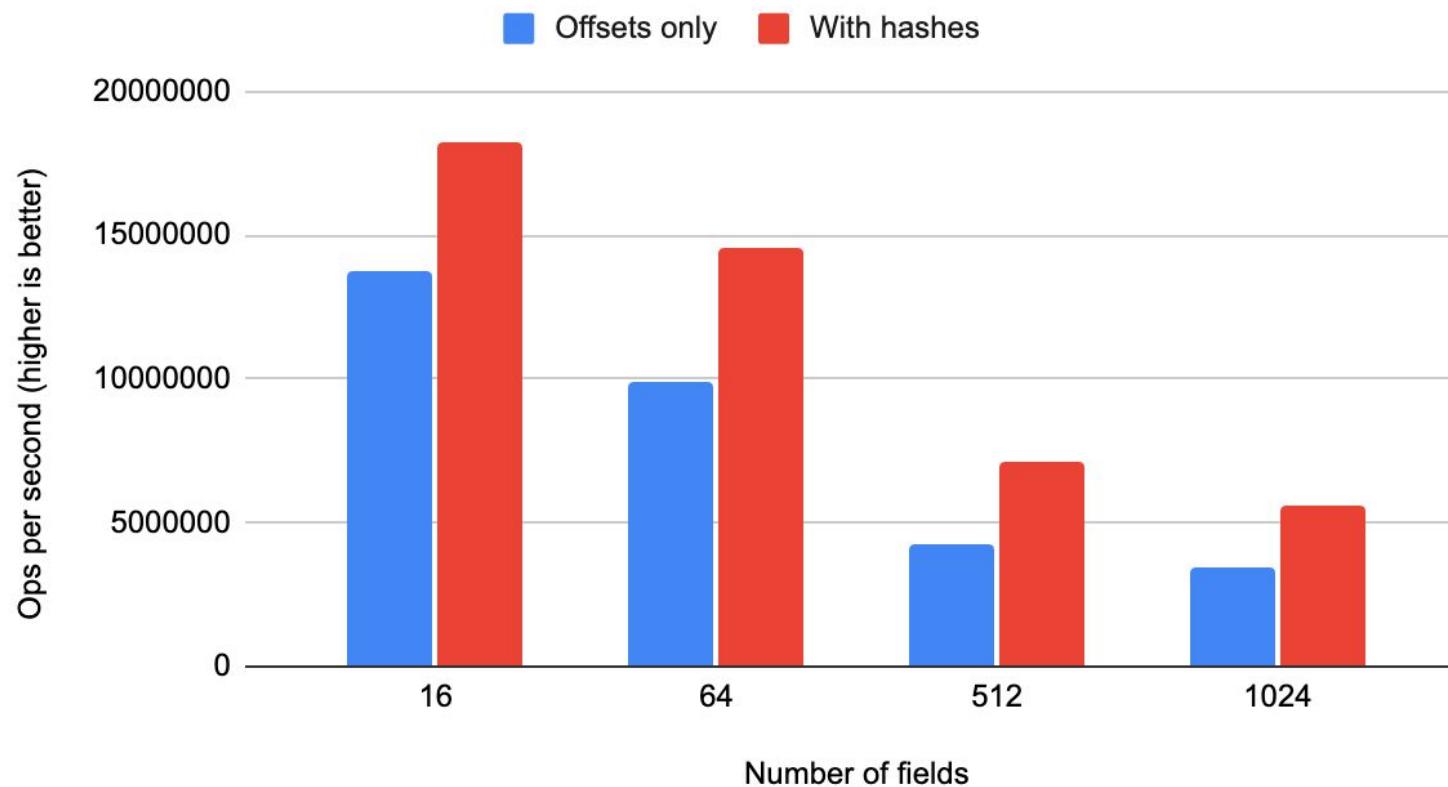
- A further enhancement:
  - Store an array of hashes at the start of our map
  - Search these first, before checking actual keys
    - Comparing small integers is faster than comparing strings
    - Hashes are contiguous in memory, which is more cache-friendly



# Faster Searches – Benchmark Results

- How does this scheme compare to the previous one?

Read-only workload



# Faster Searches – Benchmark Results

- Can we do even faster than that?
- perf shows that a lot of the cost is in this if condition in std::lower\_bound
- **sub** is cheap, so this probably indicates the previous branch is mispredicted

		_ZSt13__lower_boundIPKtiN9__gnu_cxx5__ops14_Iter_less_valEET_S5_S5_RKT0_T1_():
2.14		_forwardIterator __middle = __first;
2.63		std::advance(__middle, __half);
		if (__comp(__middle, __val))
	2.63	cmp %edi,%r15d
	22.64	→ jle 239838 <experimental::ser::adv::sortedstring::Container::search(std::basic_s
		{
		__first = __middle;
		++__first;
		__len = __len - __half - 1;
	0.05	< <b>sub</b> %rsi,%rax>
		++__first;
	0.09	lea 0x2(%rdx),%rbx
		__len = __len - __half - 1;
	0.33	dec %rax
		while (__len > 0)
		test %rax,%rax

# Binary Search – Mispredicted Branches

- The mispredicted branch is the one where we decide whether to look at the lower or upper half of the sub-array
- This makes sense – we really can't predict whether the value will be in the upper or lower half each time around the loop

```
const T *lower_bound(const T *begin, const T *end, T value) {  
    int64_t len = end - begin;  
    while (len > 0) {  
        const int64_t step = count / 2;  
        if (begin[step] < value) { ←  
            begin = &begin[step + 1];  
            len -= step + 1;  
        }  
        else len = step;  
    }  
    return begin;  
}
```

This branch is often  
mispredicted

Bloomberg

Engineering

# Binary Search – Mispredicted Branches

- To fix this, we can remove the branching and use a conditional move instruction instead
- This is sometimes called “branchless” programming

# Faster Searches

Idea 1: Store a table of offsets at the start of the map

Idea 2: Store hashes of keys at the start of the map too

**Idea 3: Use branchless binary search**

# Branchless Binary Search

- “Array Layouts for Comparison-Based Searching” by Khuong and Morin [1] has a good implementation of branchless binary search
- It looks something like this:

```
const T *search(const T *begin, const T *end, T key)
{
    const T *base = begin;
    int64_t len = end - begin;
    while (len > 1) {
        const int64_t half = len / 2;
        if (base[half] < key) {
            base += half; ←
        }
        len -= half;
    }
    return *base < key ? base + 1 : base;
}
```

Doing only one assignment in the if statement makes the compiler more likely to generate a CMOV instruction

[1]: <https://arxiv.org/abs/1509.05053>

# Branchless Binary Search – Perf Report

- Our comparison and assignment in the loop is done using cmov:

```
ZNK12experimental3ser3adv12sortedstring9Container10binsearch3EPKtS5_t():
    cmp    $0x2,%r9
→ jbe   2490b0 <int experimental::ser::adv::performOperationGeneric<experimental::sortedstring>()
    mov    $0xfffffffffffffff,%rax
    nop
7.47   mov    %r13,%rdx
2.91   sar    %rdx
0.03   lea    (%rdx,%rax,1),%rsi
18.68  cmp    %r11w,(%r14,%rsi,2)
4.33   cmoveb %rsi,%rax
    sub    %rdx,%r13
    cmp    $0x1,%r13
→ jg    248dc8 <int experimental::ser::adv::performOperationGeneric<experimental::sortedstring>()
    lea    0x2(%r14,%rax,2),%r13
_ZNK12experimental3ser3adv12sortedstring9Container6searchESt17basic_string<
    cmp    %r13,%r8
→ je    249130 <int experimental::ser::adv::performOperationGeneric<experimental::sortedstring>()
```

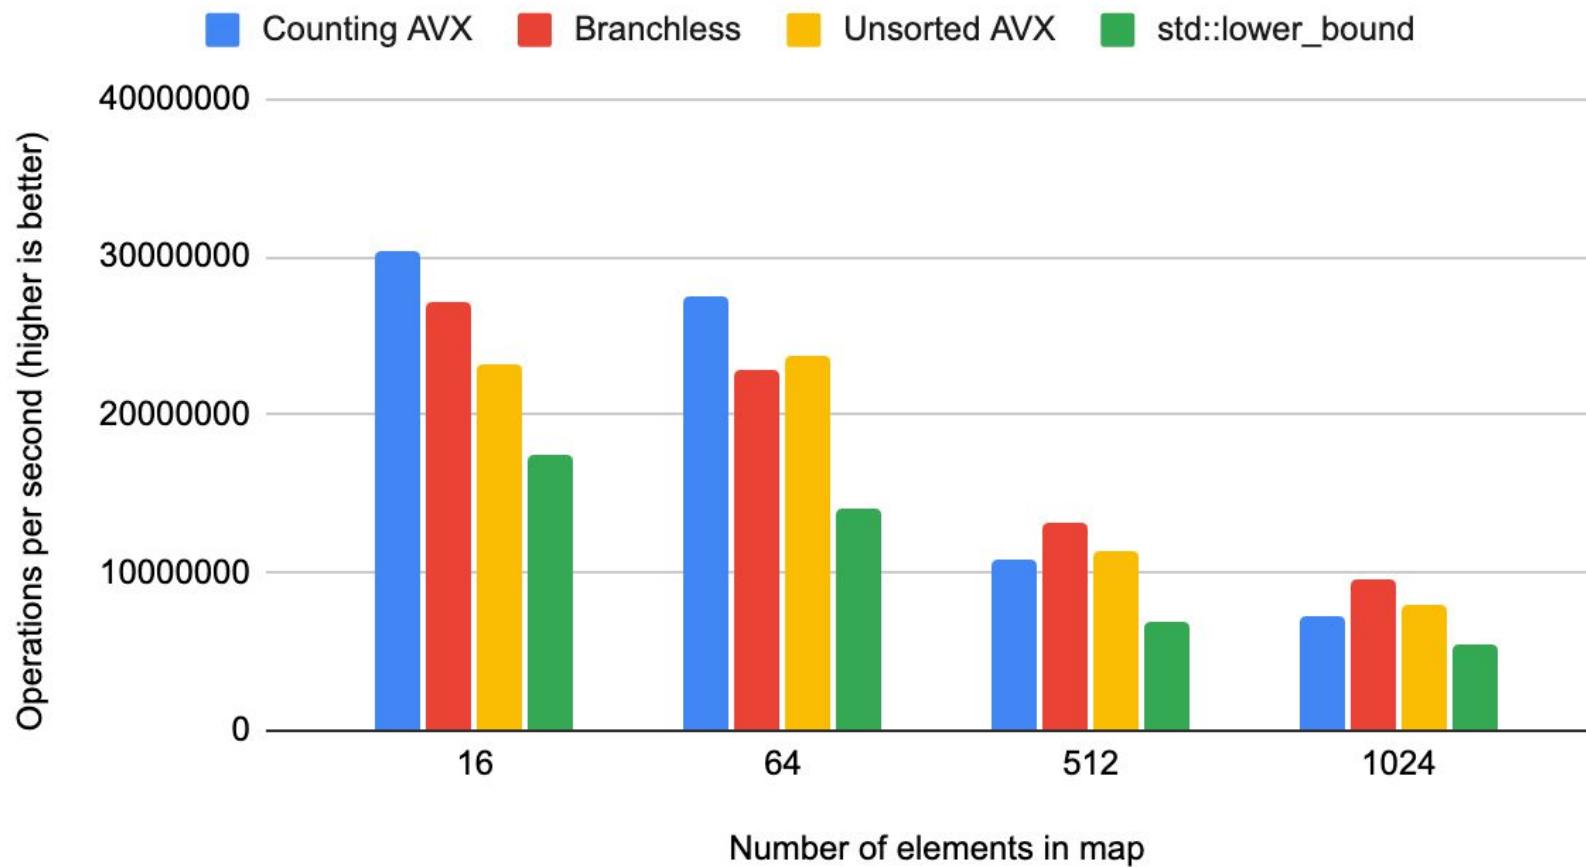
# Other Search Algorithms – Side Note

- I went down a bit of a rabbit hole and tried other accelerated search algorithms:
  - Vectorized counting linear search on a sorted array
  - Vectorized standard linear search on an unsorted array
- For a good article on these and more, see this blog post:  
<https://dirtyhandscoding.github.io/posts/performance-comparison-linear-search-vs-binary-search.html>

# Faster Searches – Benchmark Results

- Comparison of all hash-search methods:

Read-only workload



# Faster Searches

Idea 1: Store a table of offsets at the start of the map

Idea 2: Store hashes of keys at the start of the map too

Idea 3: Use branchless binary search

**I'm pretty happy with search performance now!**

# Faster Modifications

How do we make modifications faster?

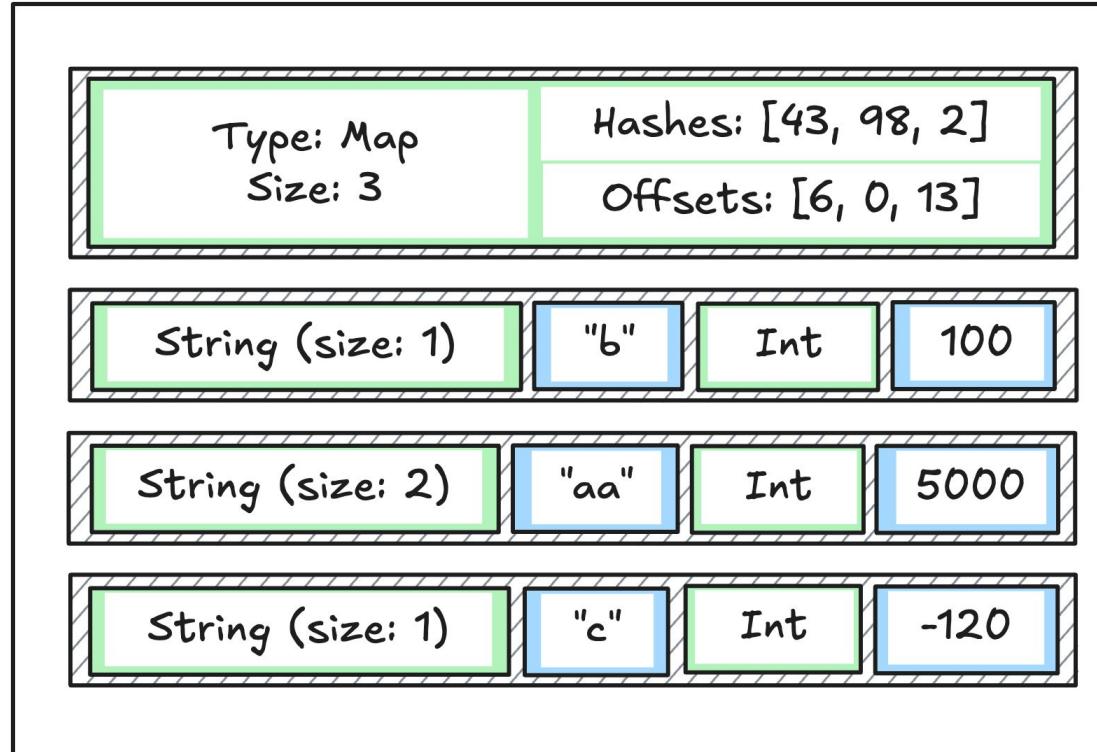
**TechAtBloomberg.com**

© 2025 Bloomberg Finance L.P. All rights reserved.

**Bloomberg**  
Engineering

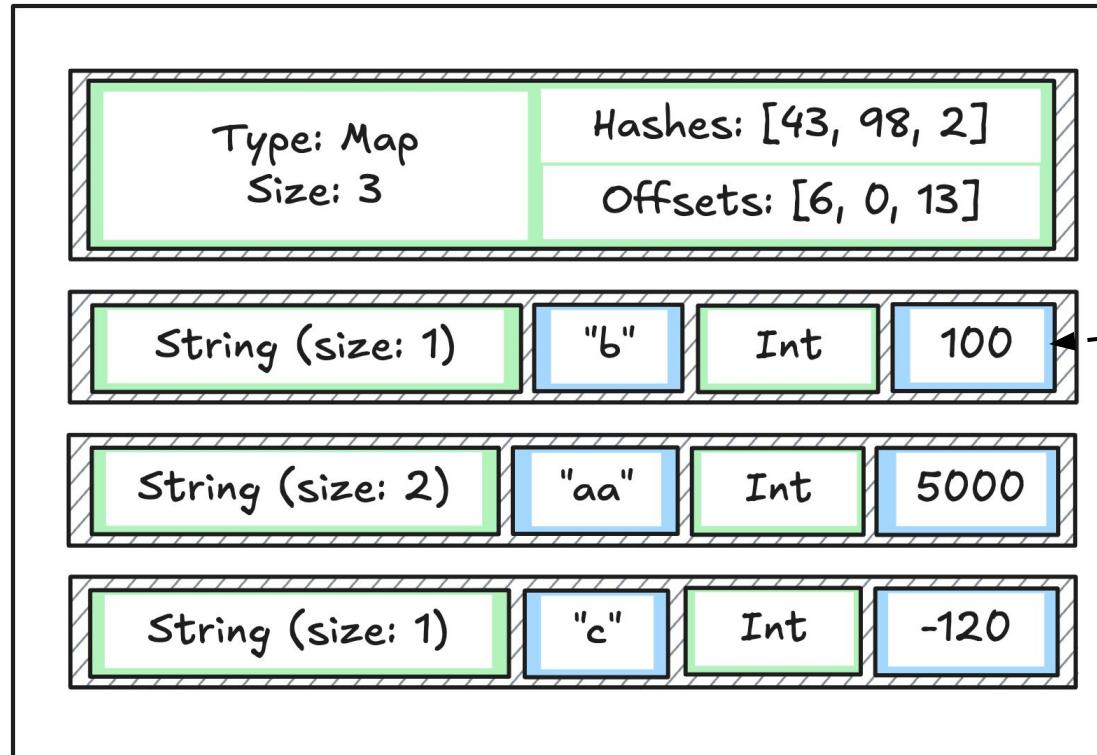
# Making our Format Modifiable

- In our baseline format, we performed modifications by just moving everything around
- Now that we've introduced offsets, this is problematic
- Why?



# Making our Format Modifiable

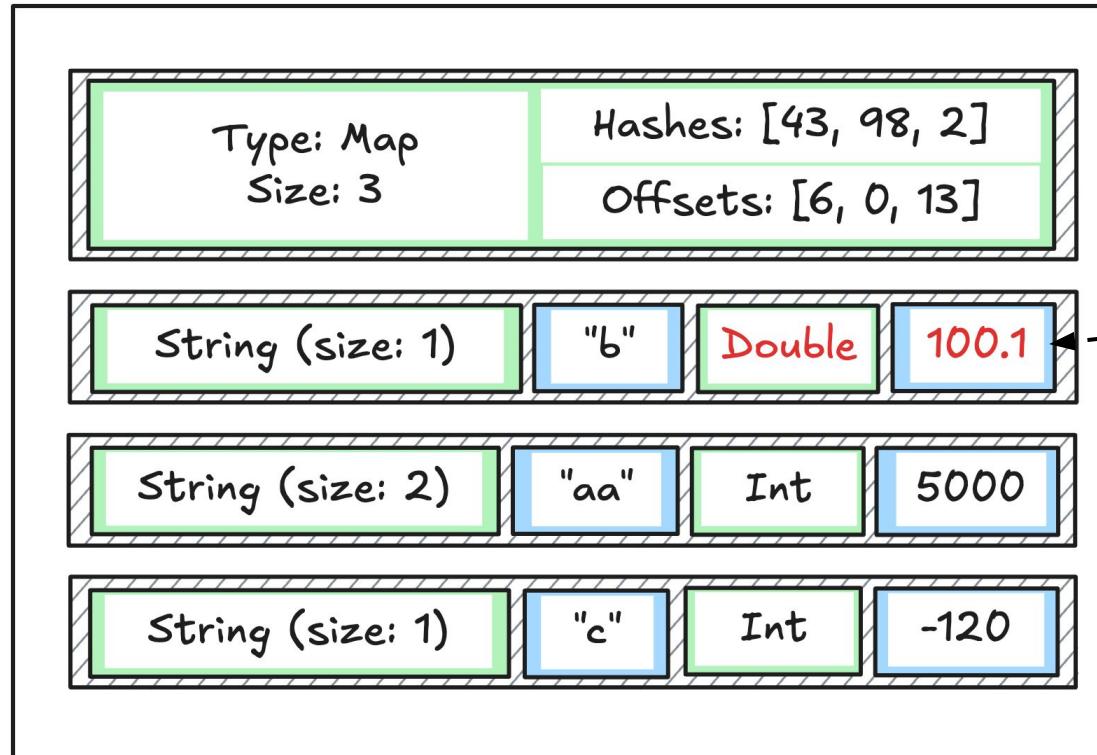
- In our baseline format, we performed modifications by just moving everything around
- Now that we've introduced offsets, this is problematic
- Why?



If, for example, I make '100' into a double, I'll change its size

# Making our Format Modifiable

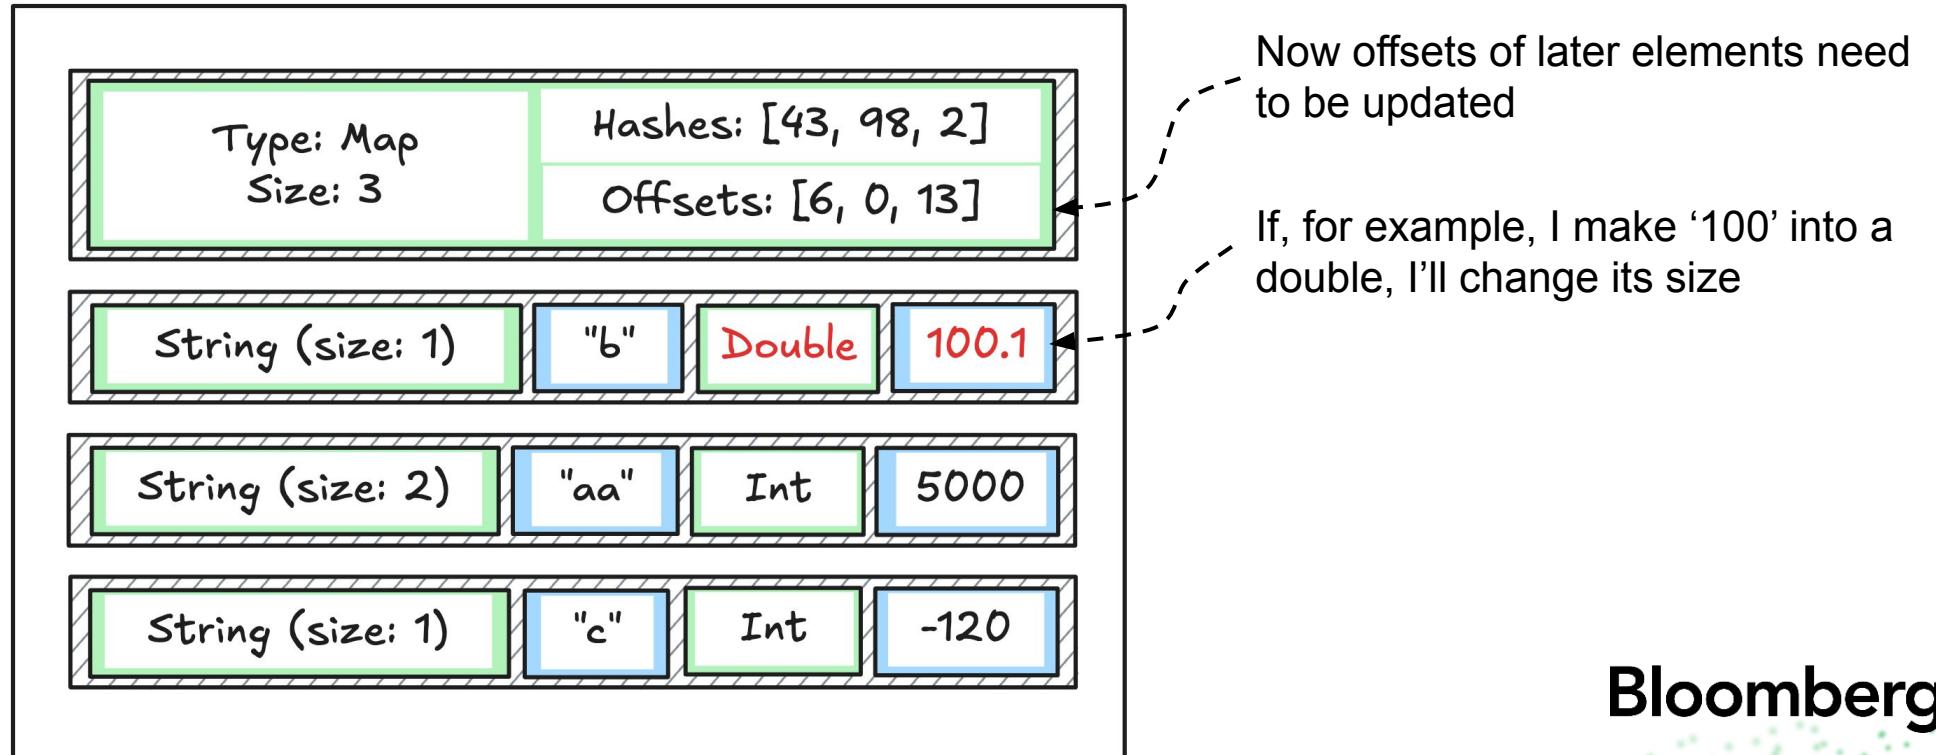
- In our baseline format, we performed modifications by just moving everything around
- Now that we've introduced offsets, this is problematic
- Why?



If, for example, I make '100' into a double, I'll change its size

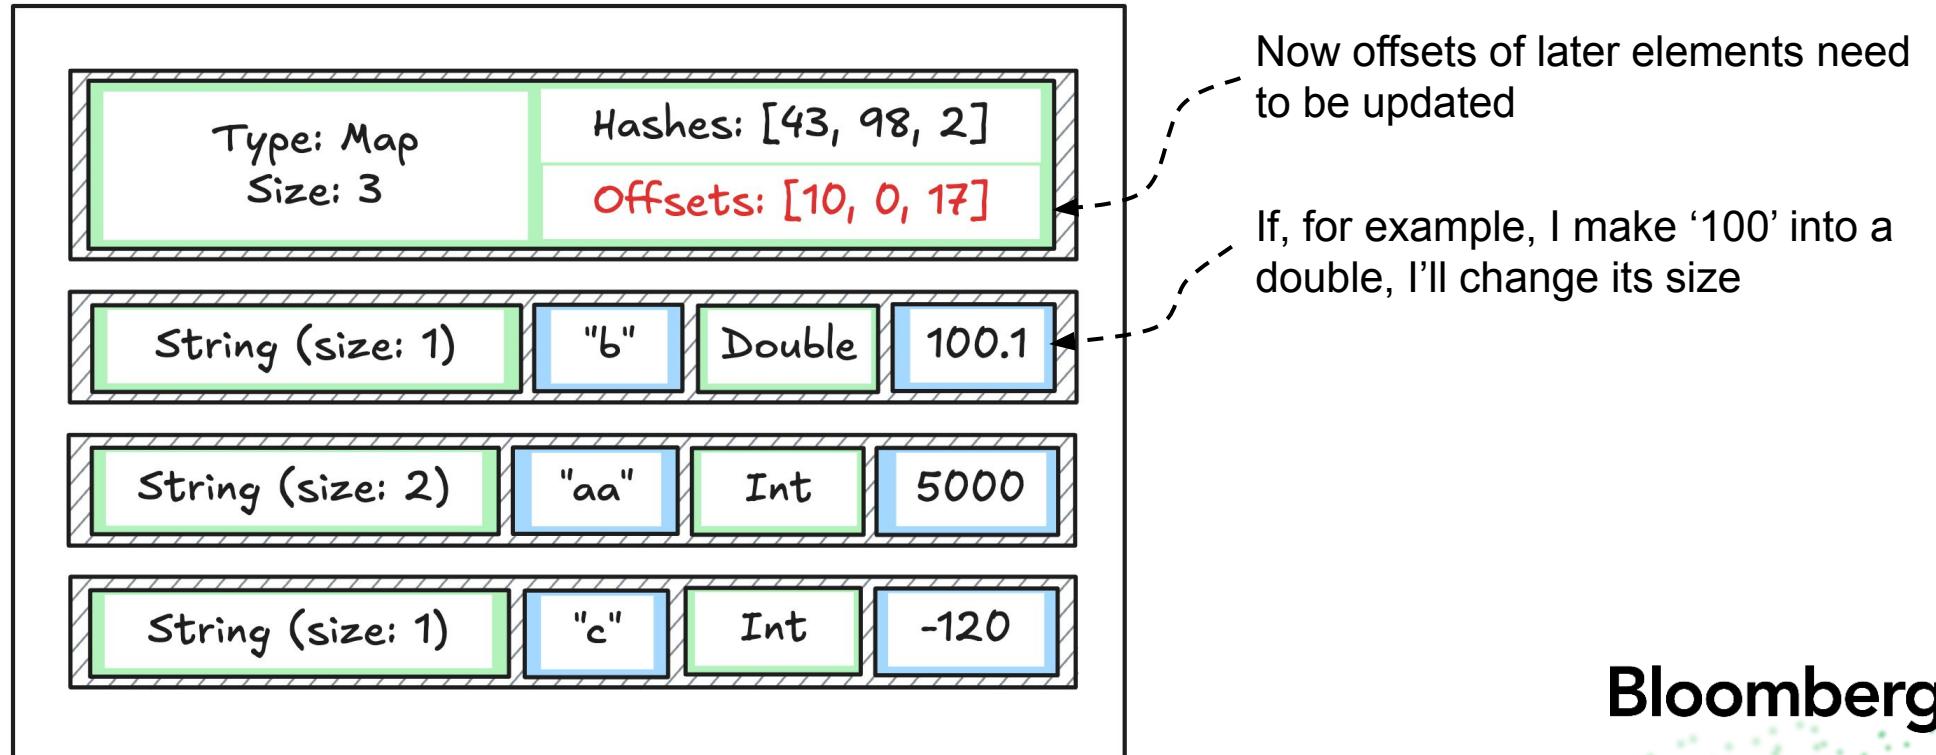
# Making our Format Modifiable

- In our baseline format, we performed modifications by just moving everything around
- Now that we've introduced offsets, this is problematic
- Why?



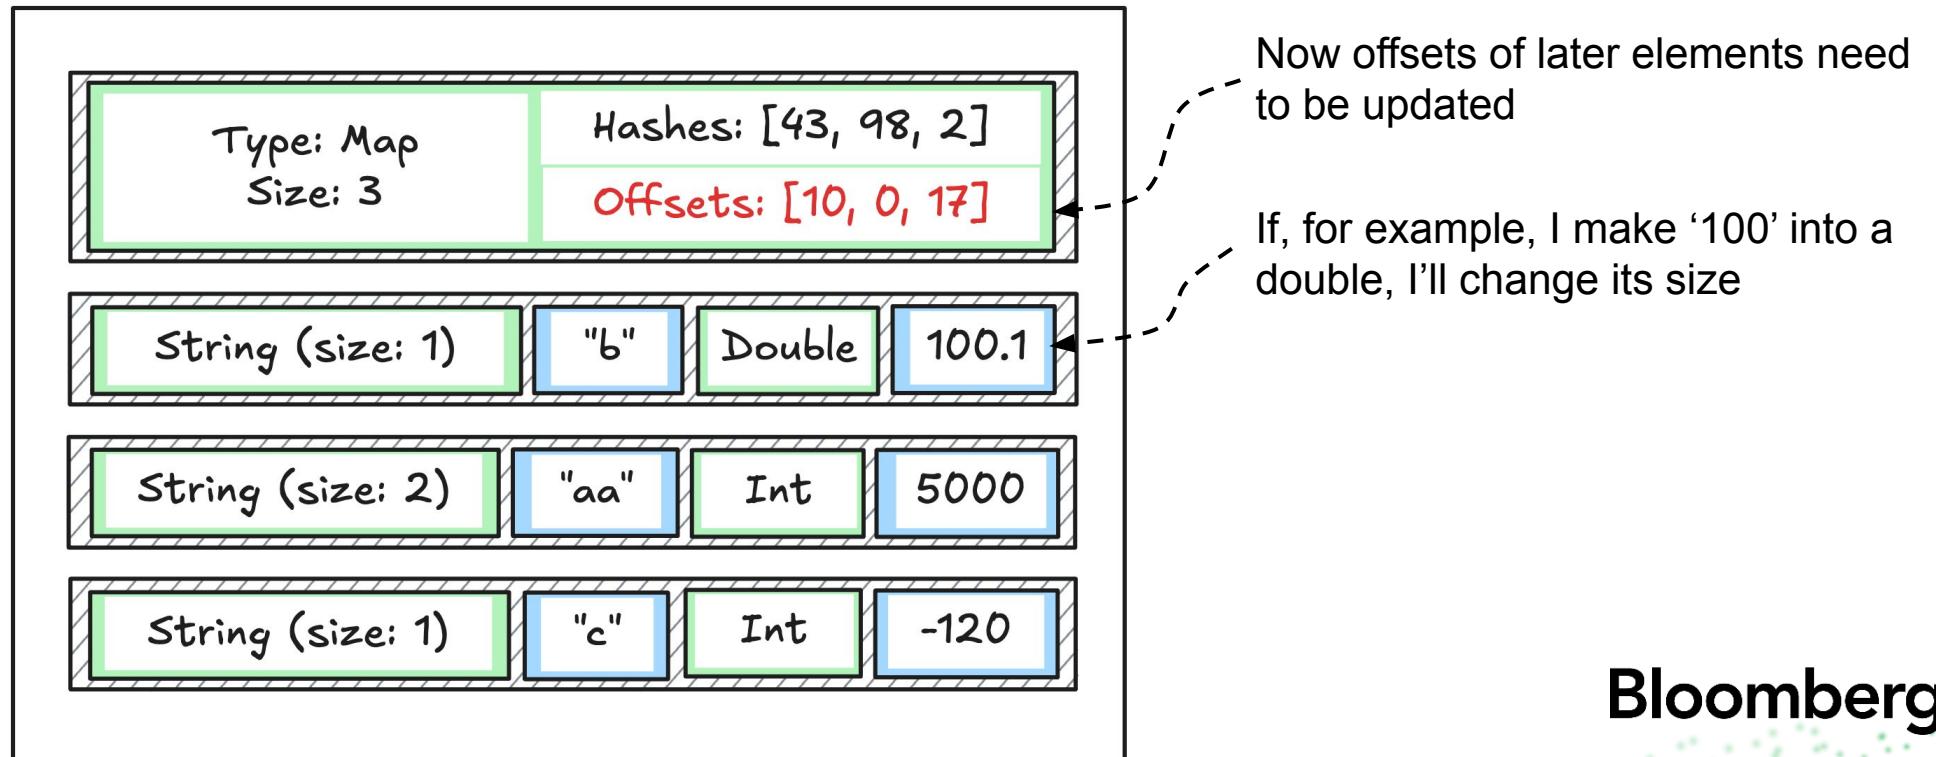
# Making our Format Modifiable

- In our baseline format, we performed modifications by just moving everything around
- Now that we've introduced offsets, this is problematic
- Why?



# Making our Format Modifiable

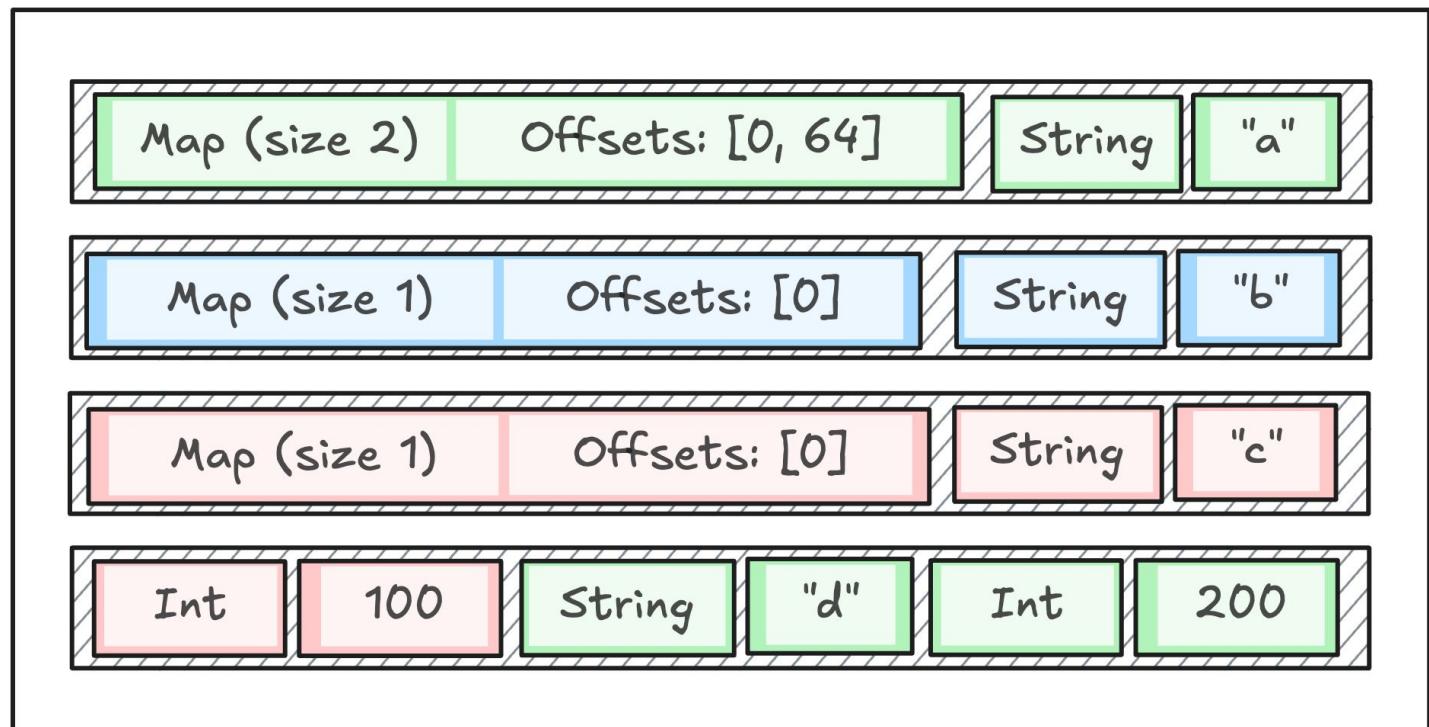
- Is this a big deal?
- On the face of it, no
- But the real problem comes when this map is nested within others



# The Problem of Nested Maps / Arrays

- If this map is nested in others, then moving elements within it will also move elements that come after this in *containing* maps
- To understand why this is a problem, consider the following document:

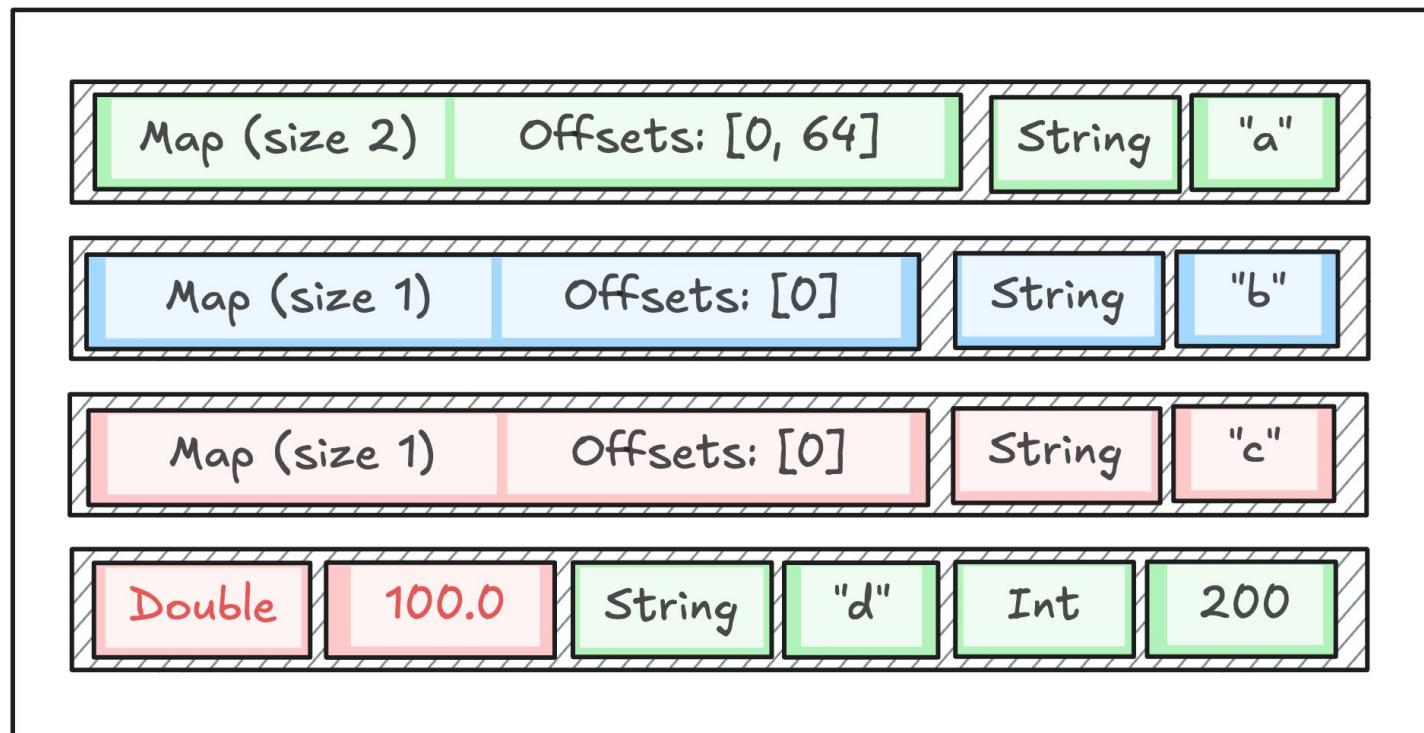
```
{  
  "a":  
  {  
    "b":  
    {  
      "c": 100  
    }  
  }  
  "d": 200  
}
```



# The Problem of Nested Maps / Arrays

- If this map is nested in others, then moving elements within it will also move elements that come after this in *containing* maps
- To understand why this is a problem, consider the following document:

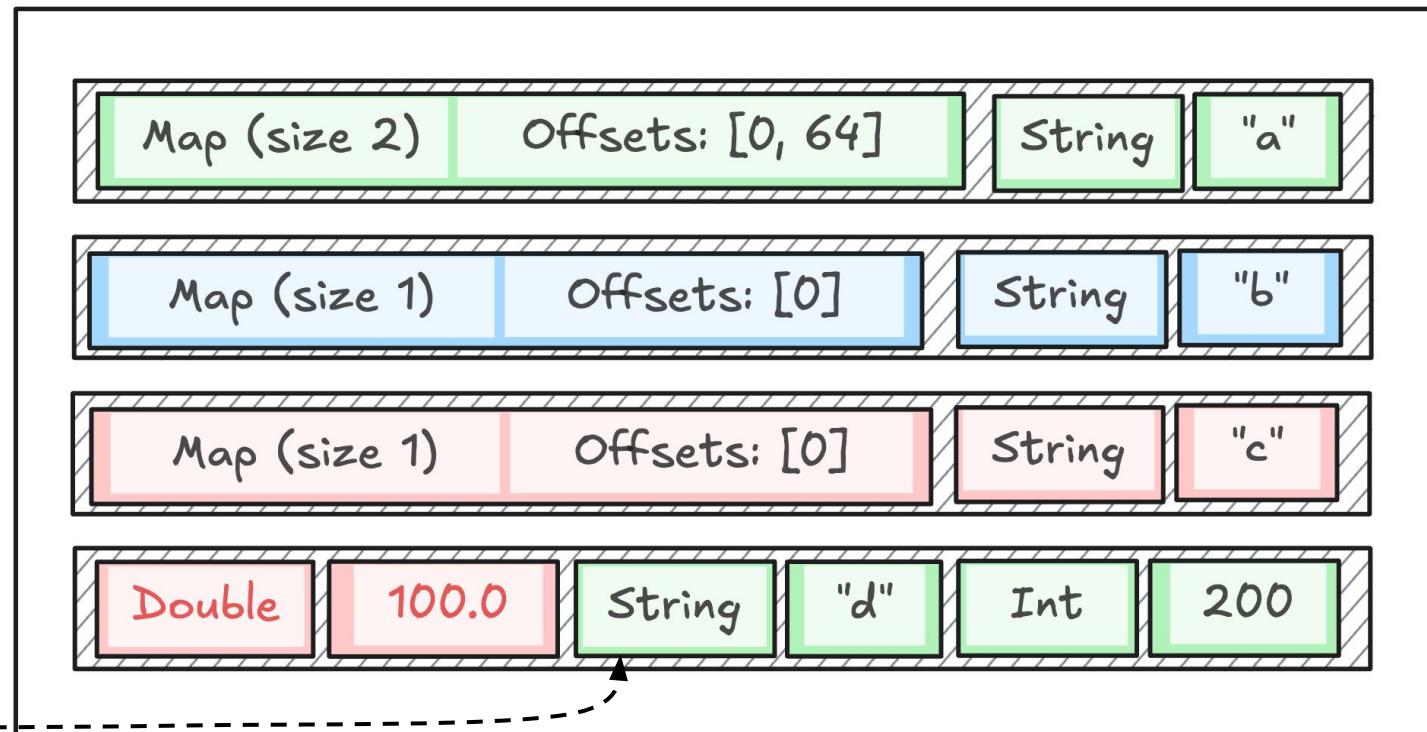
```
{  
  "a":  
  {  
    "b":  
    {  
      "c": 100.0  
    }  
  }  
  "d": 200  
}
```



# The Problem of Nested Maps / Arrays

- If this map is nested in others, then moving elements within it will also move elements that come after this in *containing* maps
- To understand why this is a problem, consider the following document:

```
{  
  "a":  
  {  
    "b":  
    {  
      "c": 100.0  
    }  
  }  
  "d": 200  
}
```

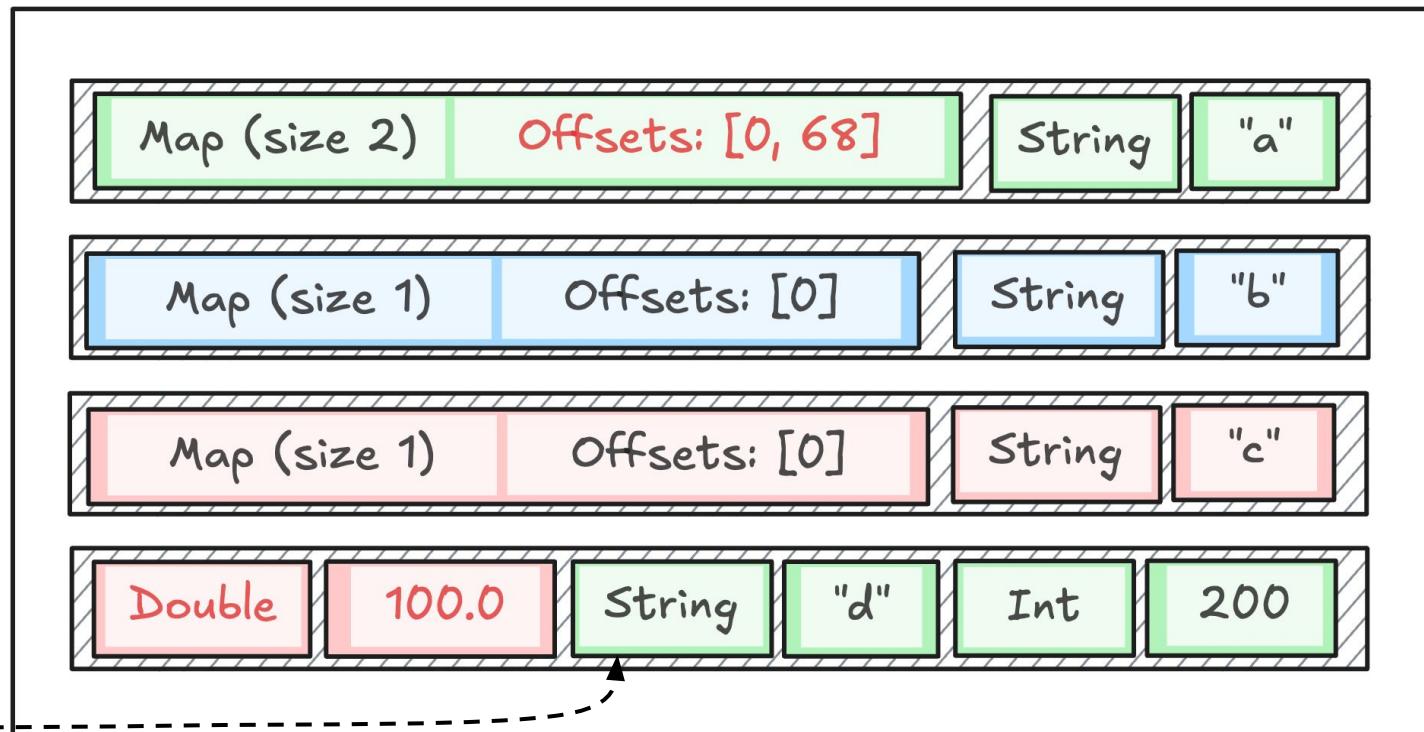


The offset of "d" now changes!

# The Problem of Nested Maps / Arrays

- If this map is nested in others, then moving elements within it will also move elements that come after this in *containing* maps
- To understand why this is a problem, consider the following document:

```
{  
  "a":  
  {  
    "b":  
    {  
      "c": 100.0  
    }  
  }  
  "d": 200  
}
```

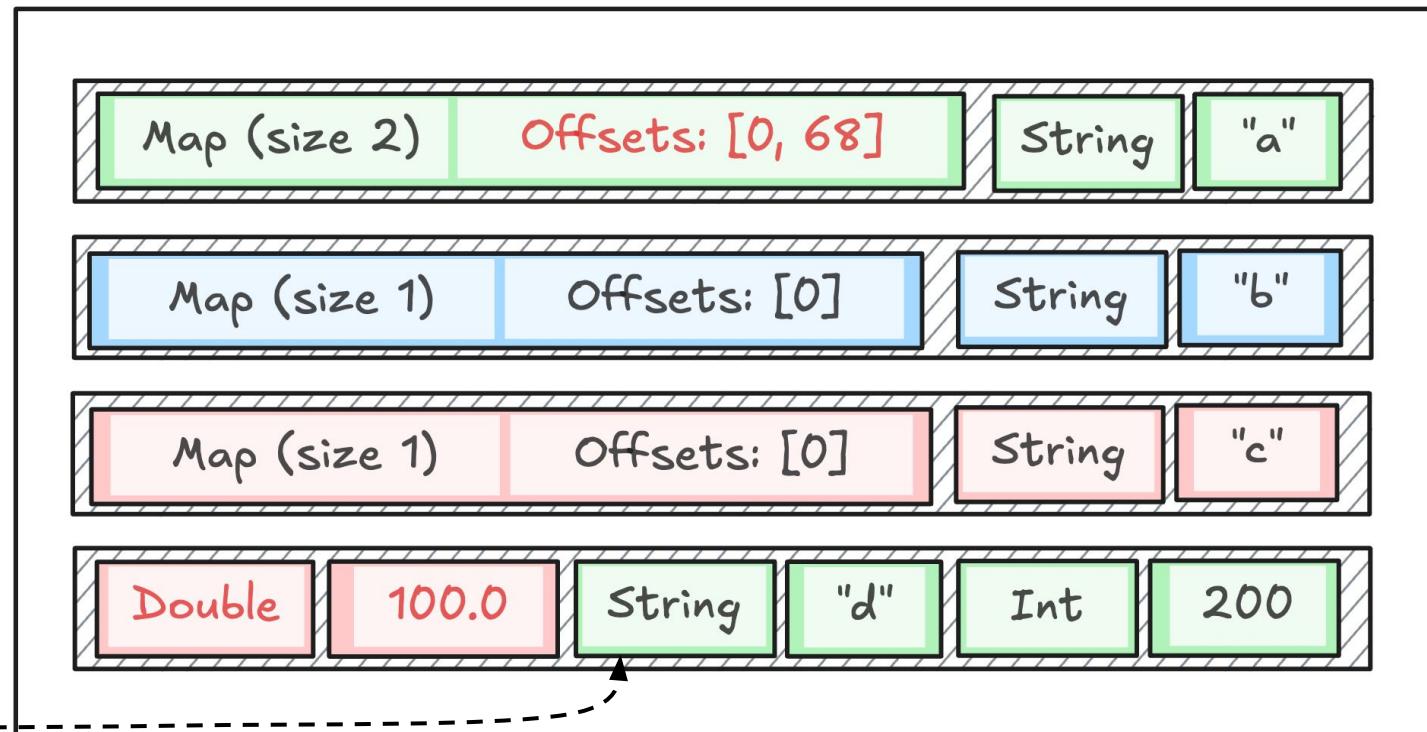


The offset of "d" now changes!

# The Problem of Nested Maps / Arrays

- If I have a large amount of nesting, this will become quite expensive
- I need to track which parents I've traversed through, or store parent pointers, both of which add complexity

```
{  
    "a":  
    {  
        "b":  
        {  
            "c": 100.0  
        }  
    }  
    "d": 200  
}
```



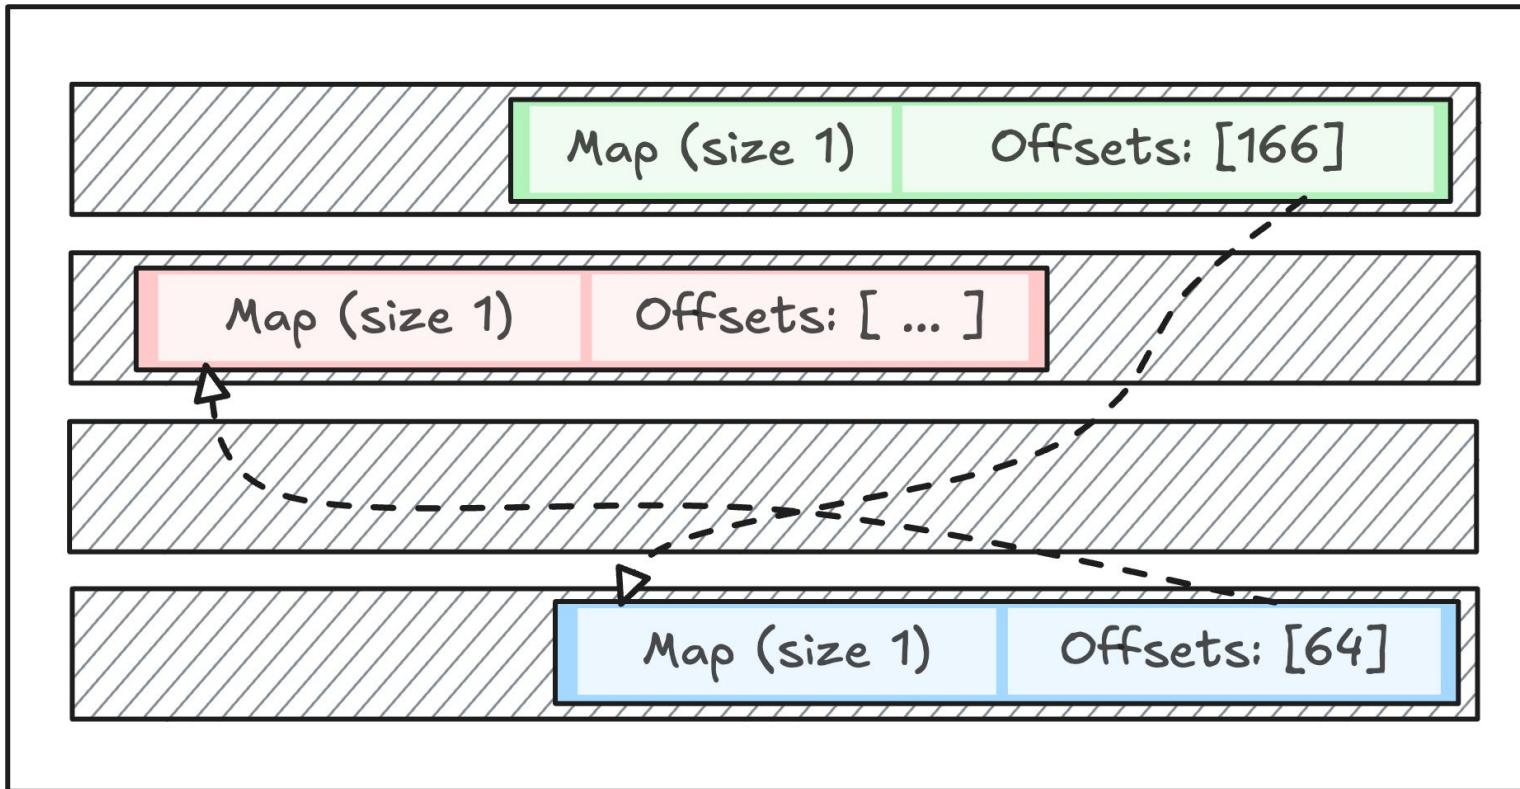
The offset of "d" now changes!

# Faster Modifications

**Idea 1: Move items “out-of-line” of their containing object**

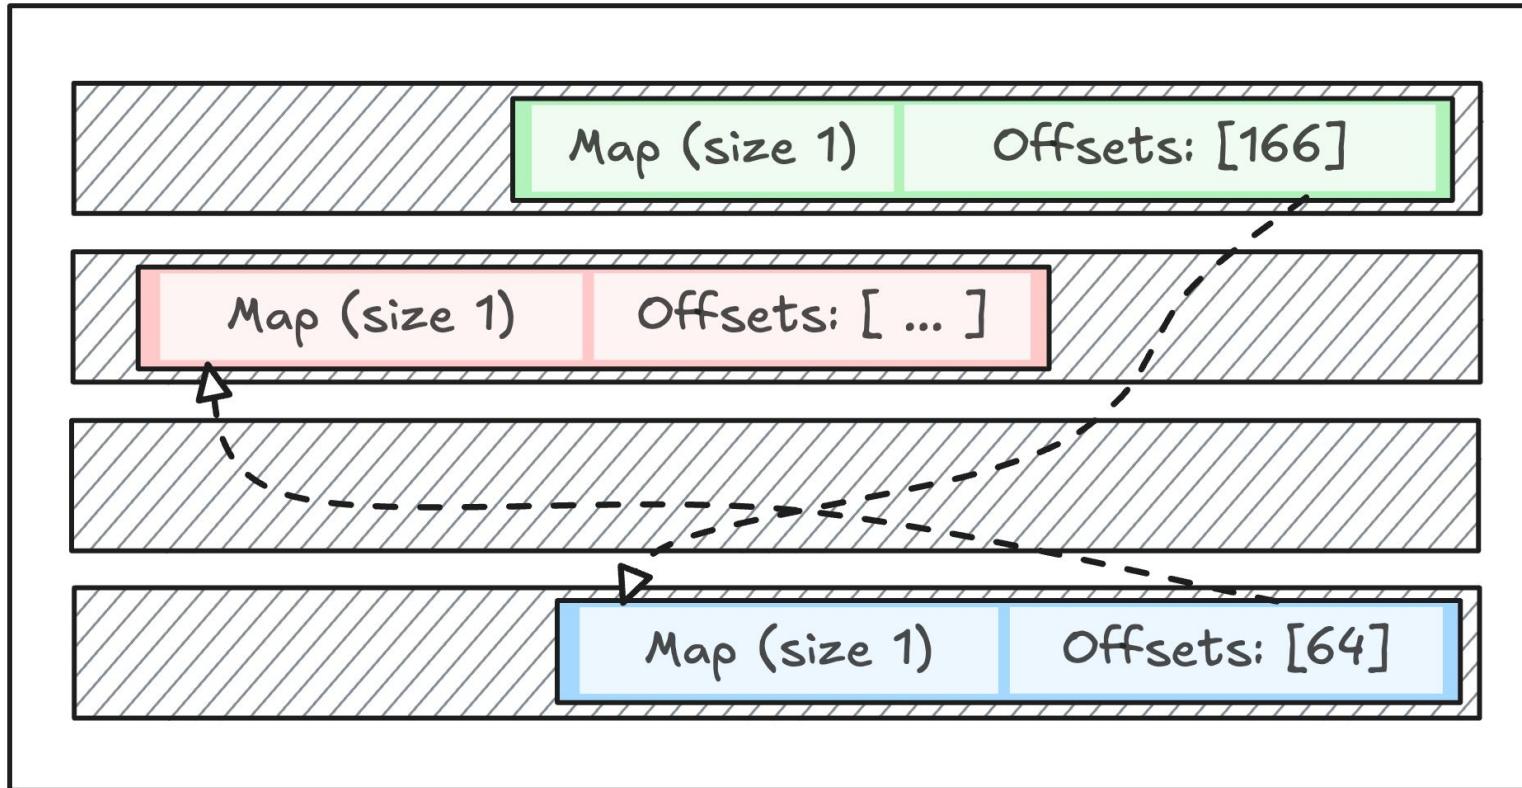
# Move Items “Out-of-Line” of Containing Object

- That means that offsets can now point anywhere else in the buffer
- Our nested maps could now look like this:



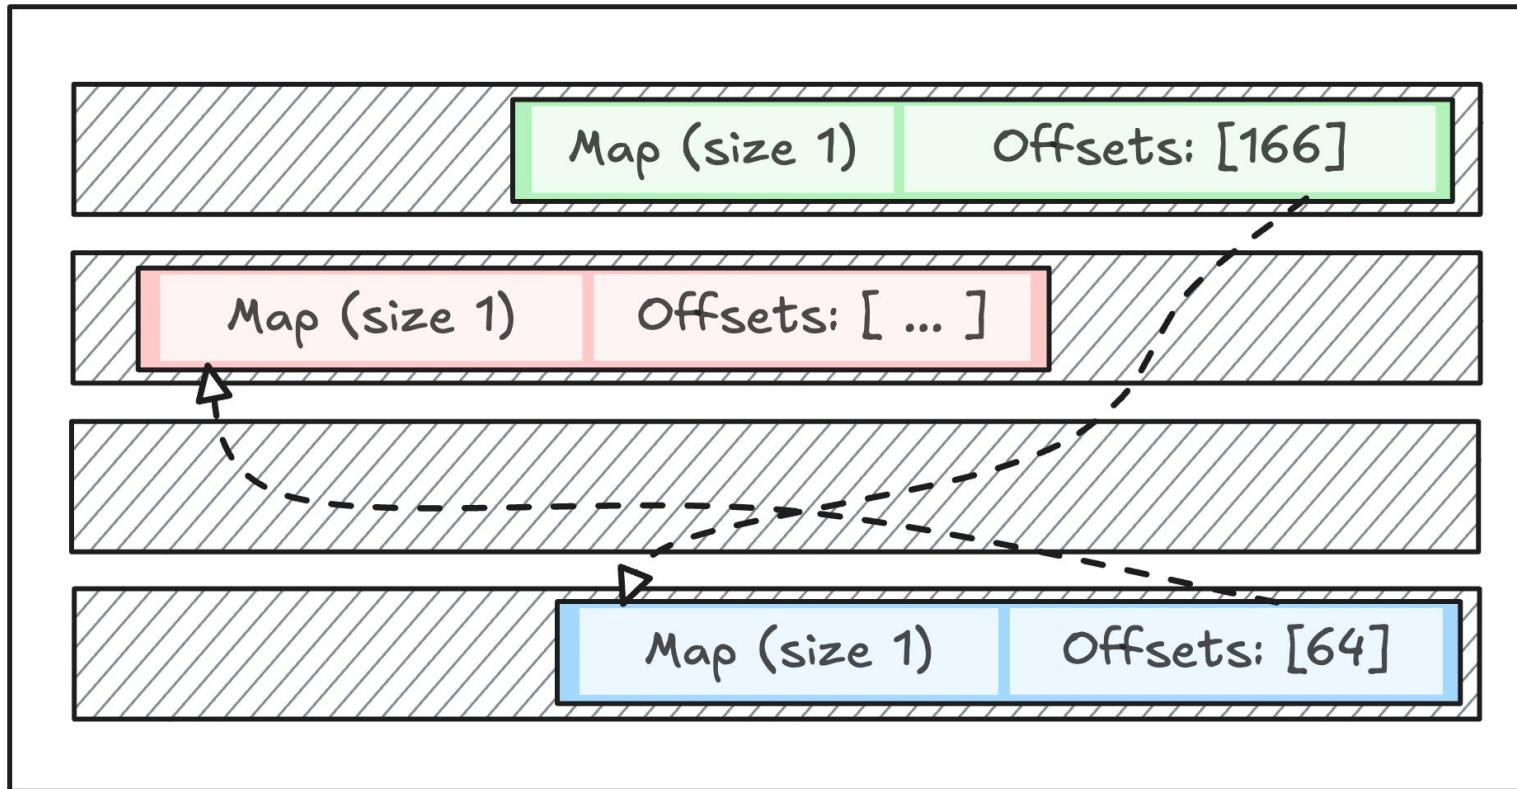
# Move Items “Out-of-Line” of Containing Object

- Our offsets have essentially become small pointers
- The nesting is now expressed by maps “pointing to” other maps (and keys/values/etc.) via offsets



# Allocating Space for Elements

- Now we have another problem:
- How do we allocate space for elements within our buffer?
  - ... And how do we free up space?



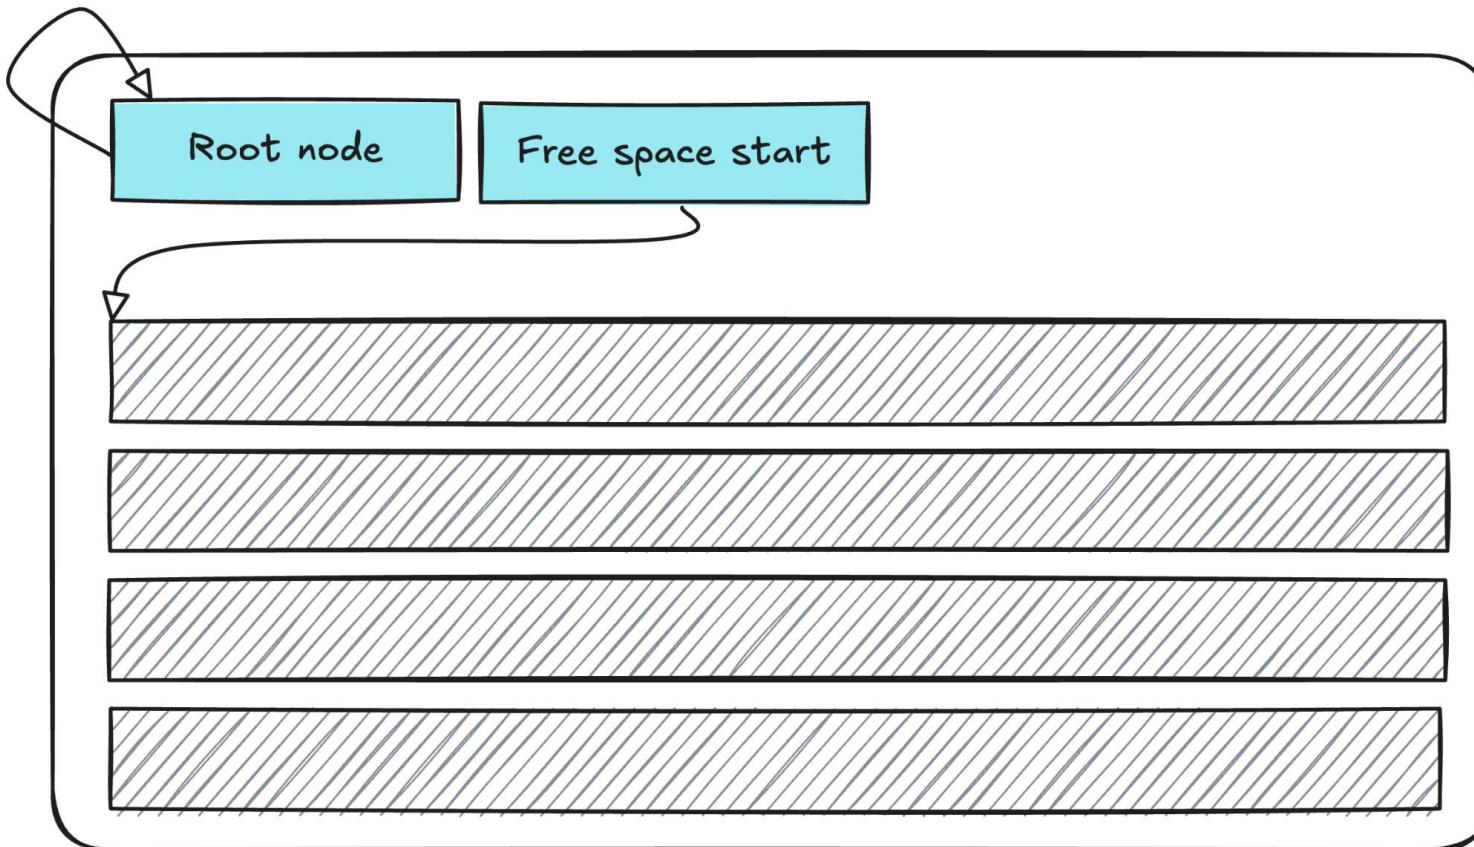
# Faster Modifications

Idea 1: Move items “out-of-line” of their containing object

**Idea 2: Implement a bump allocator**

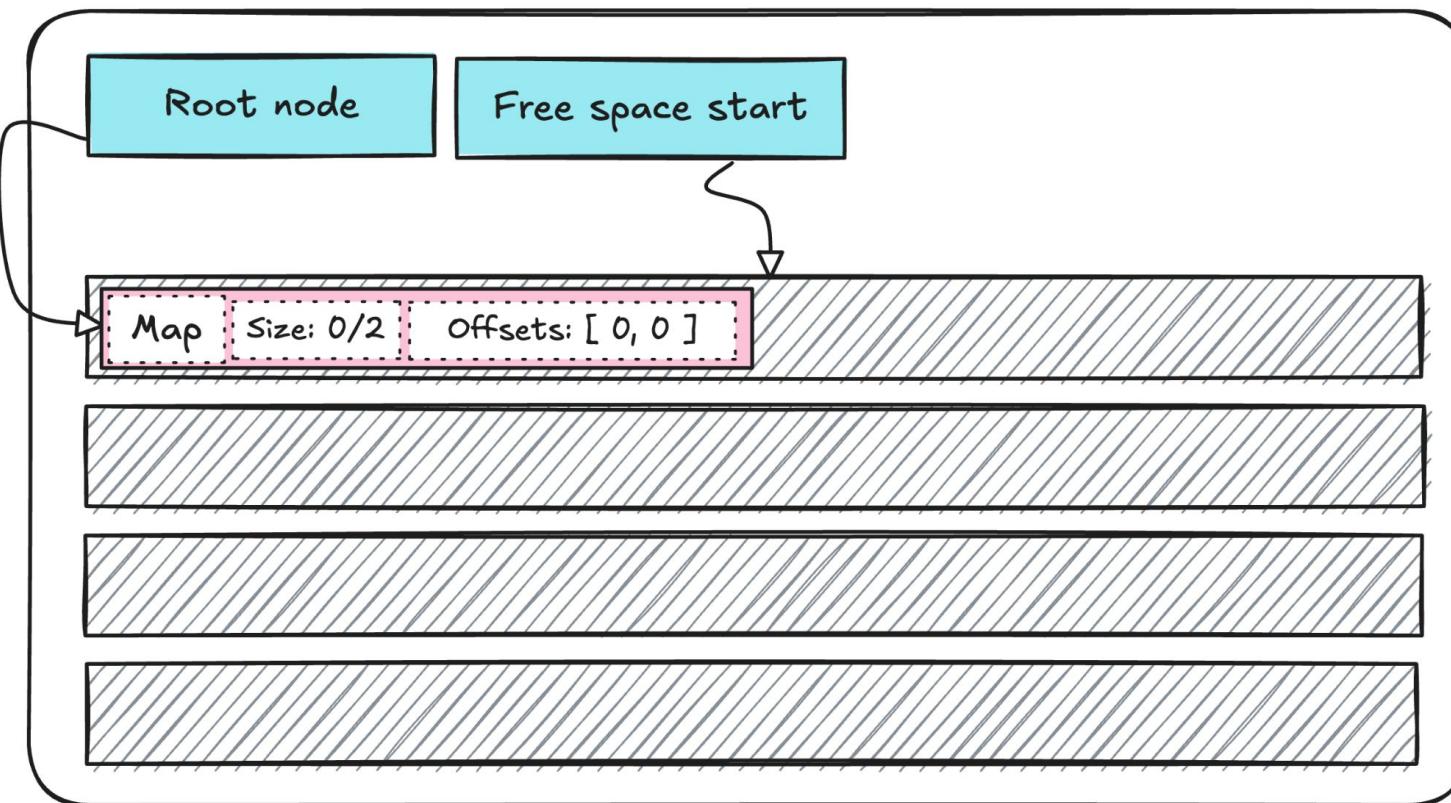
# Bump Allocator

- The most simple option is a “bump allocator”
- In this, we keep an offset to the start of the free space within our buffer
- When we want to allocate something, we just bump this offset



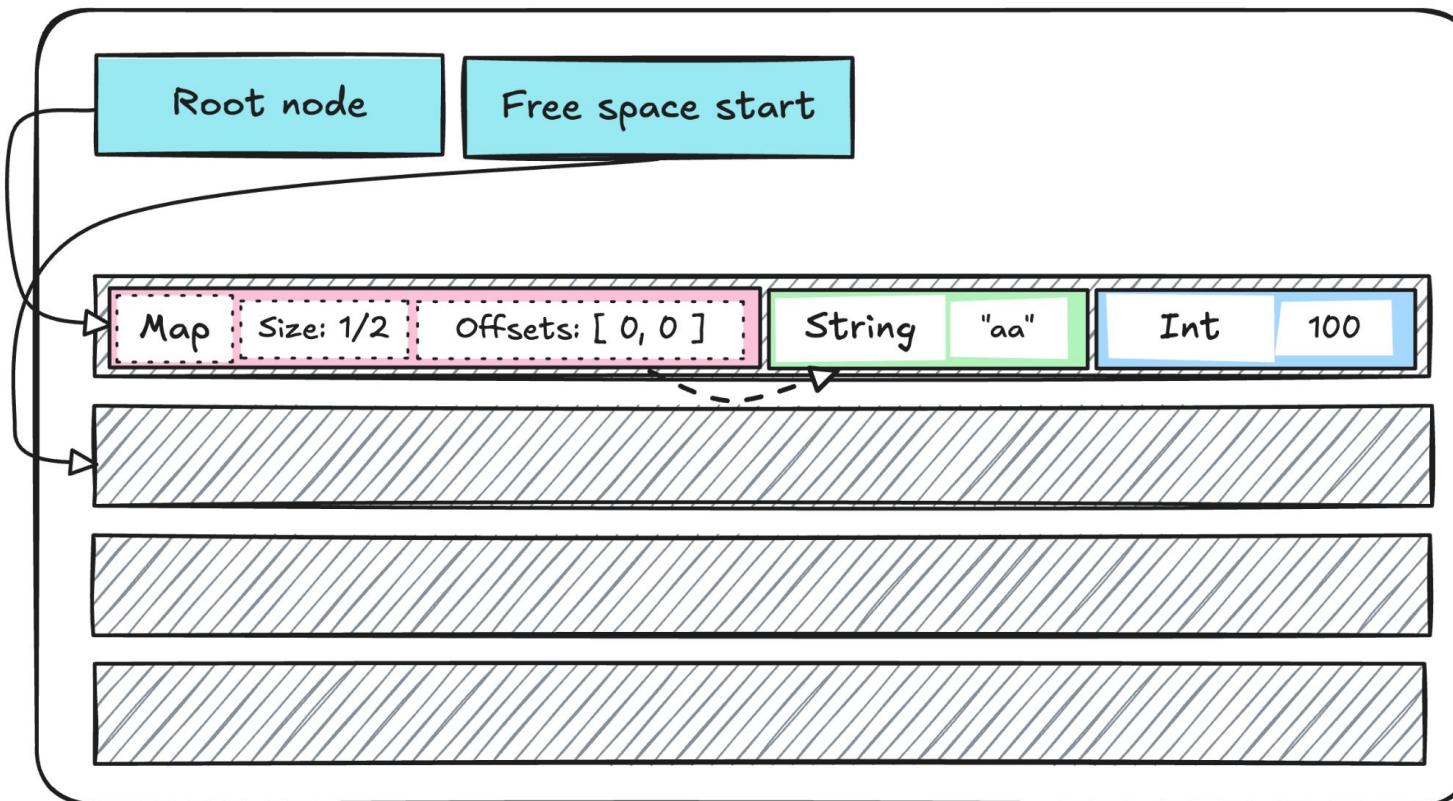
# Bump Allocator

- First, we'll turn our root node into a map



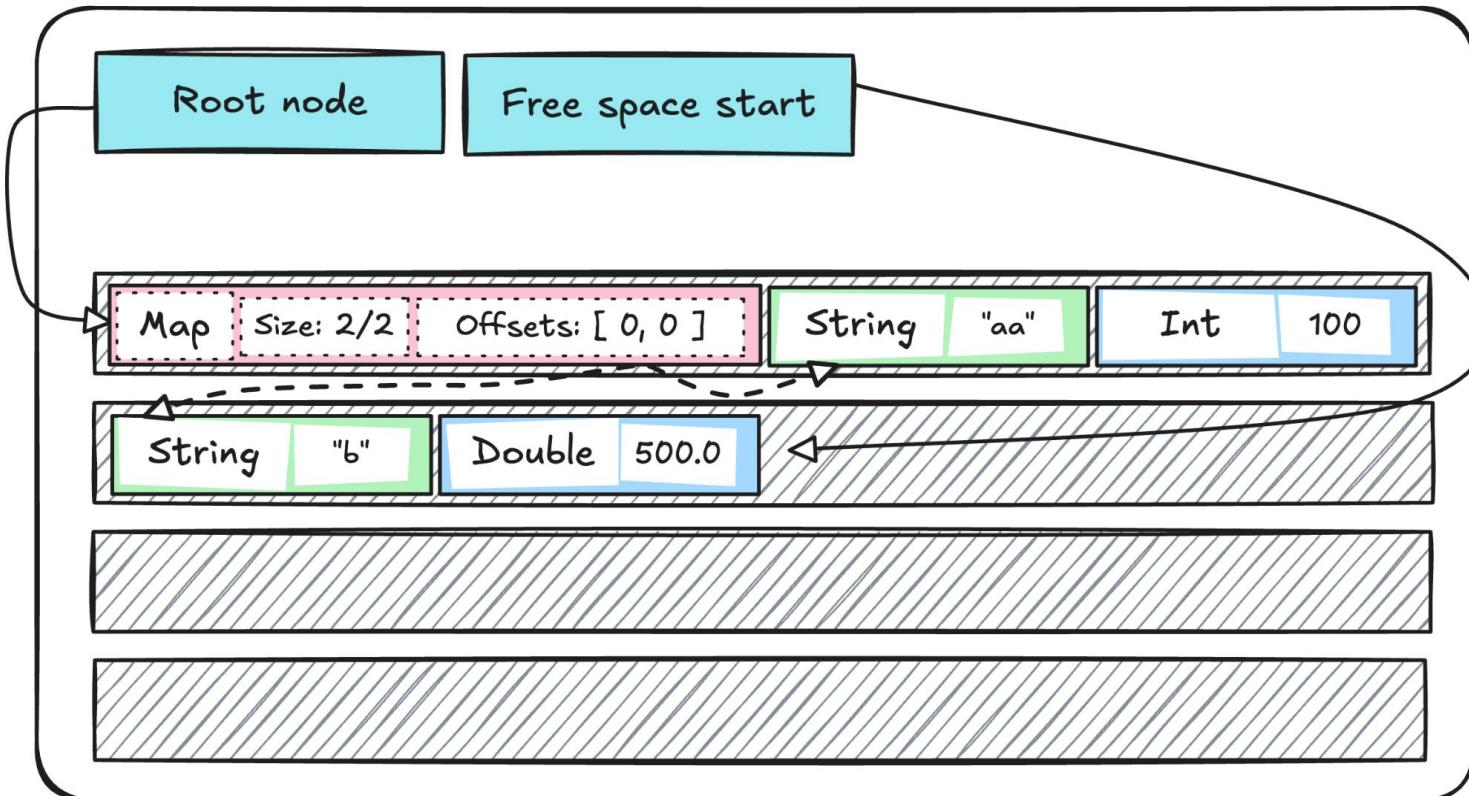
# Bump Allocator

- Now, let's add some keys and values to our map



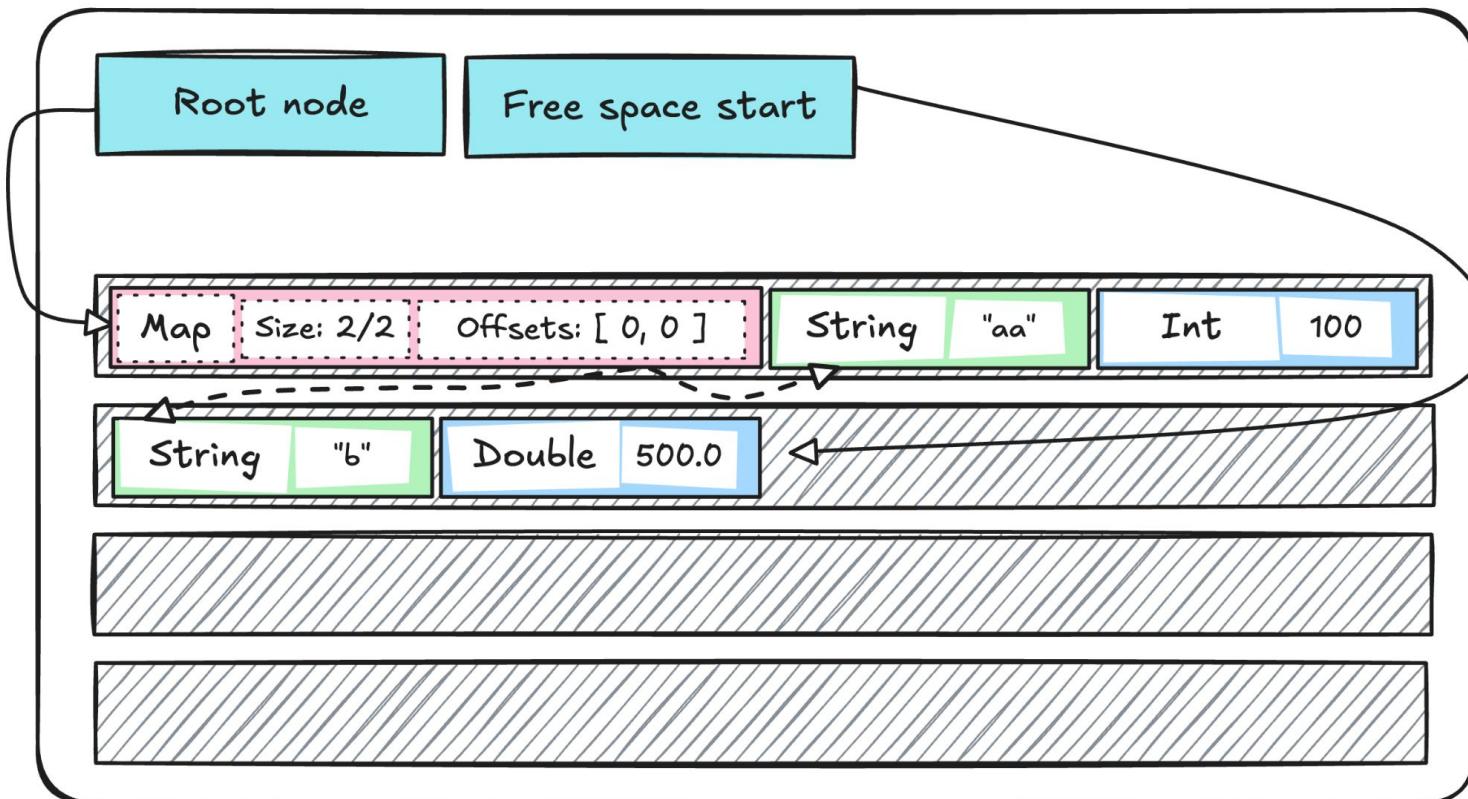
# Bump Allocator

- Now, let's add some keys and values to our map



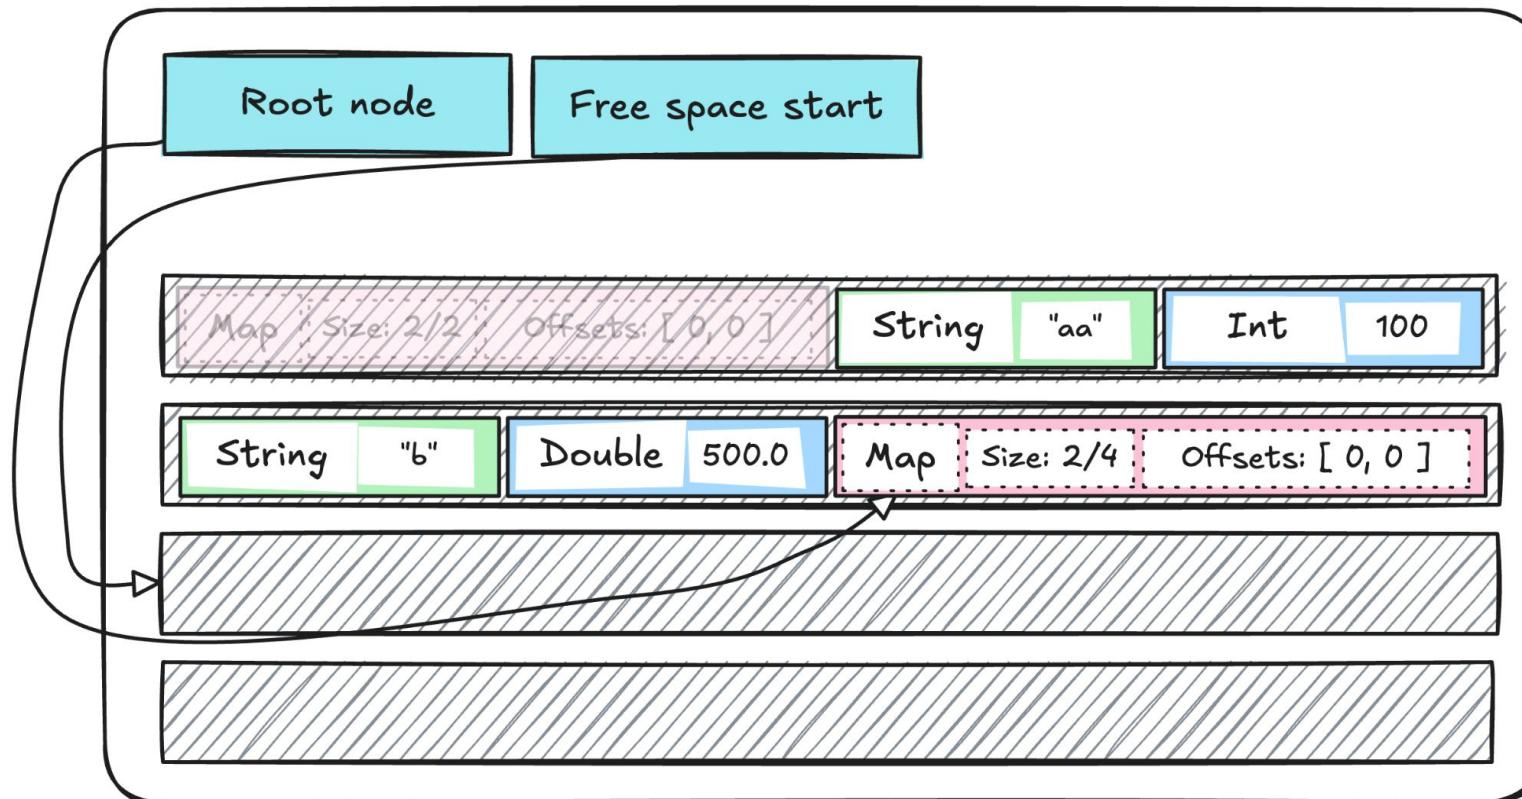
# Bump Allocator

- Now I want to add another key and value, but I'm out of space in my map
- So I need to allocate more space for a larger map



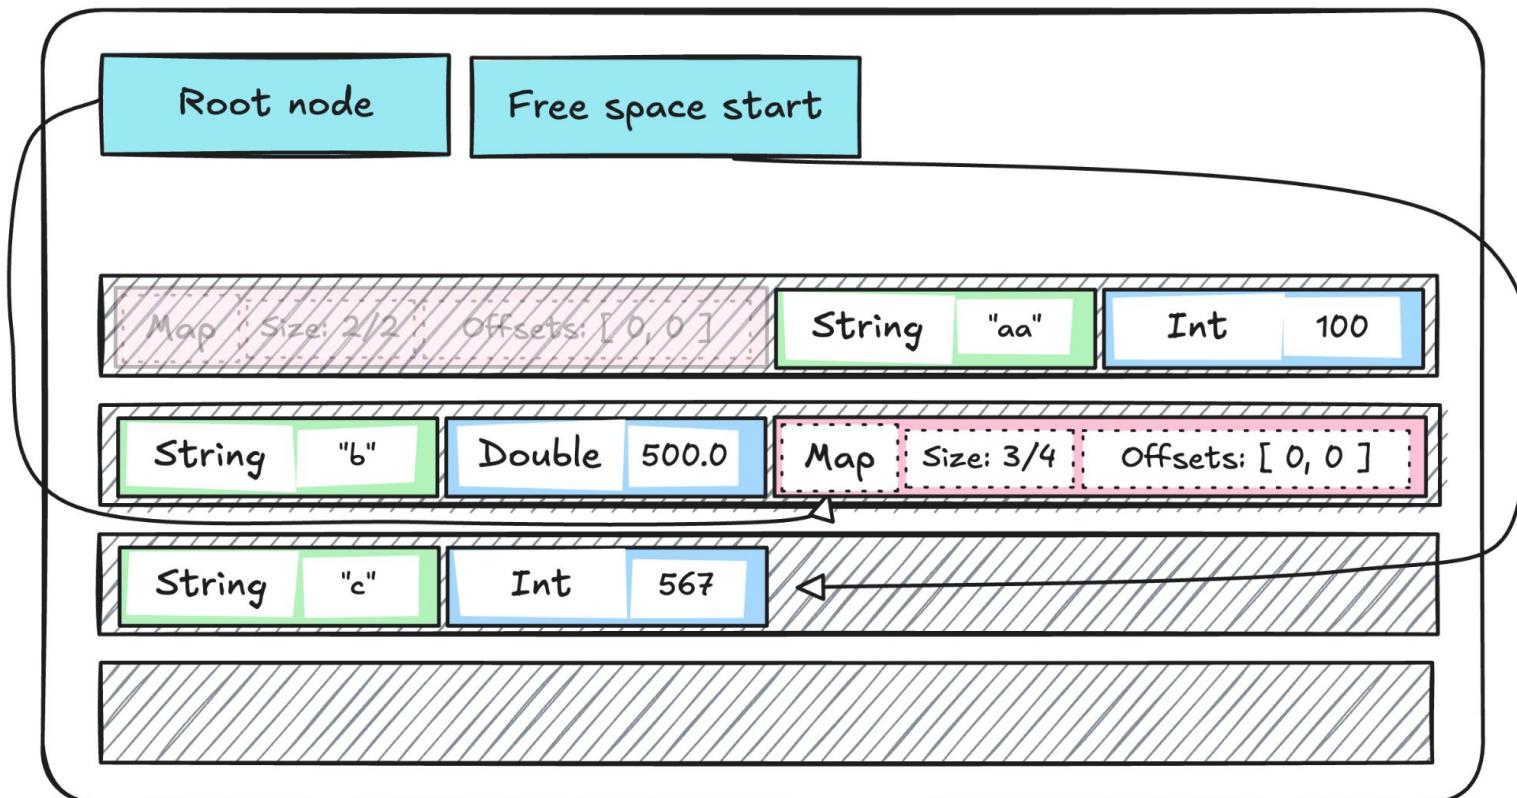
# Bump Allocator

- Now I want to add another key and value, but I'm out of space in my map
- So I need to allocate more space for a larger map



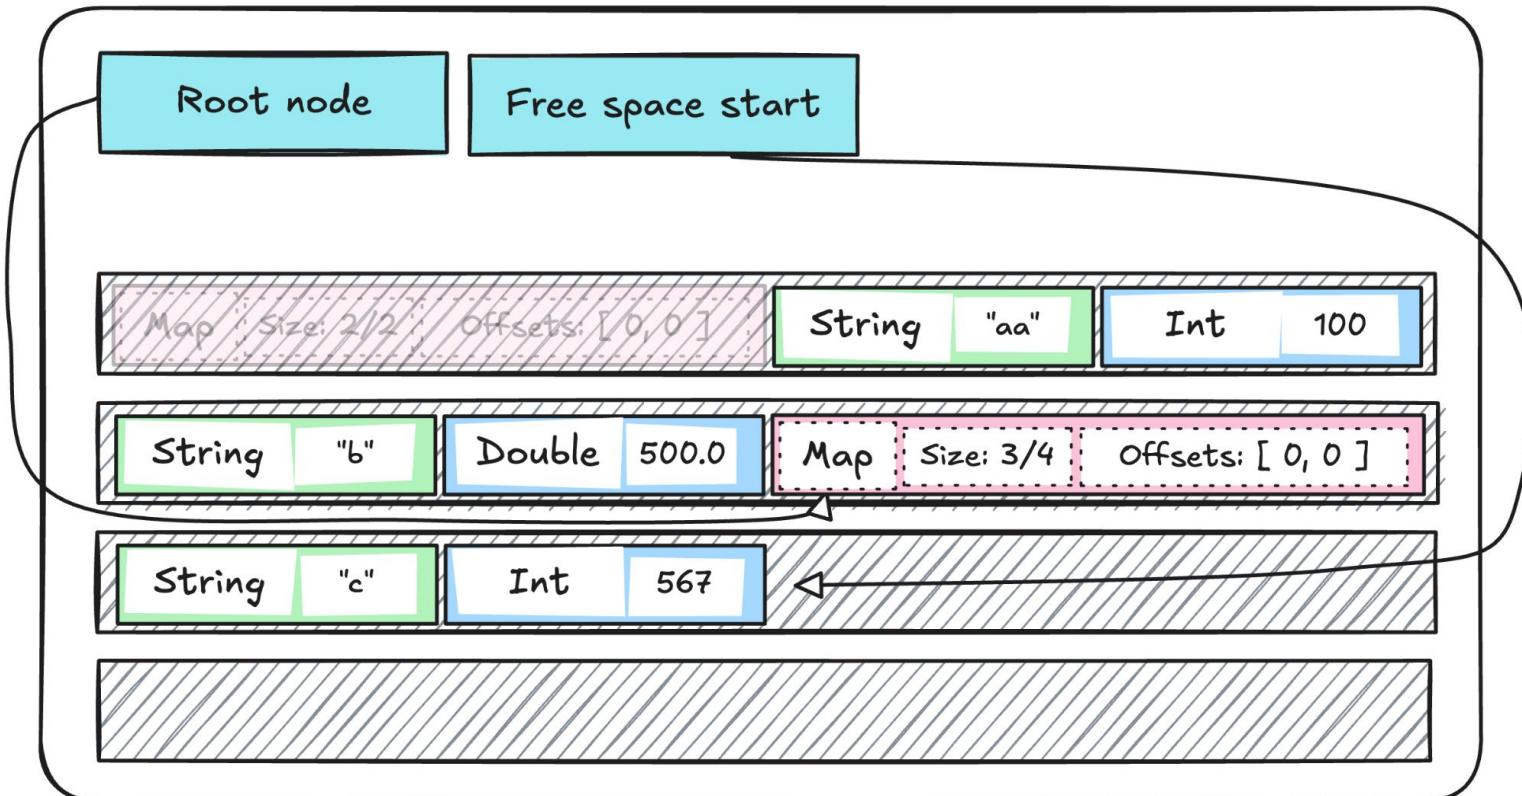
# Bump Allocator

- Now I want to add another key and value, but I'm out of space in my map
- So I need to allocate more space for a larger map
- And now I can add my value



# Bump Allocator

- Clearly, this strategy wastes a lot of space
- I will need to compact the whole buffer once I run out of space



# How can I Improve on This?

- We could keep track of free space by maintaining a free-list through the buffer
- We would then try to satisfy allocations from the free-list before attempting to resize or compact

# Faster Modifications

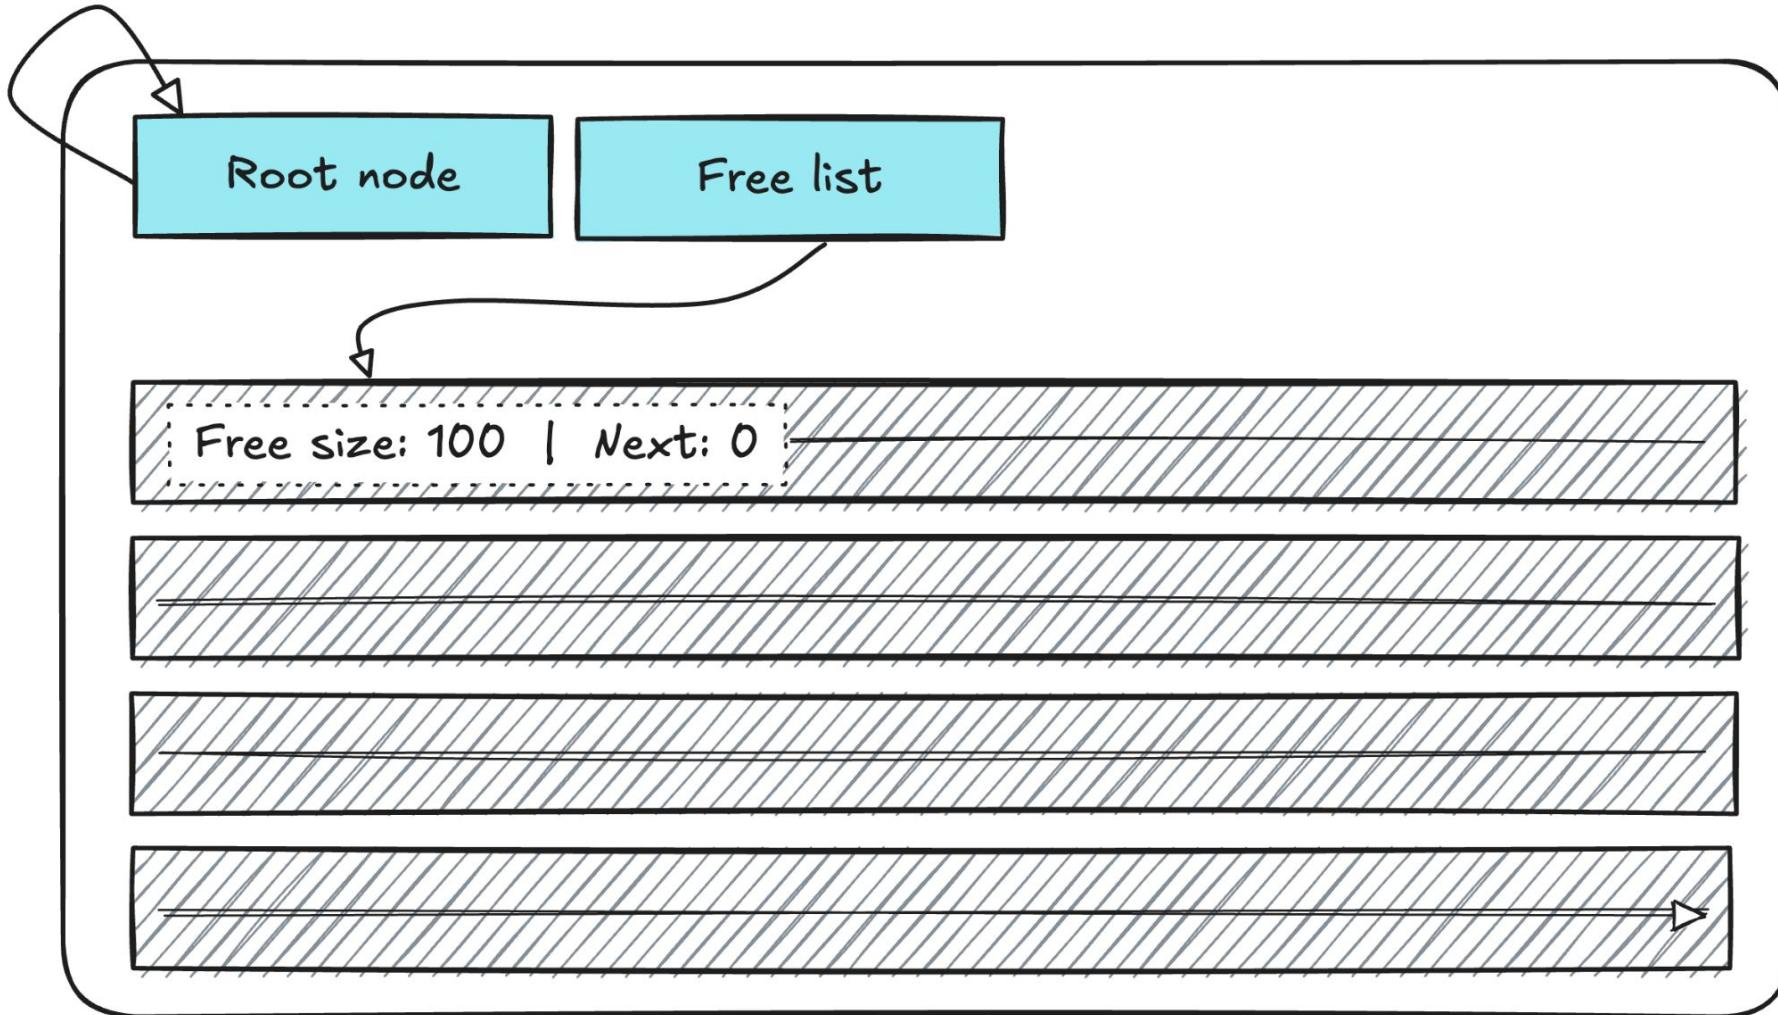
Idea 1: Move items “out-of-line” of their containing object

~~Idea 2: Implement a bump allocator~~

**Idea 3: Implement a free-list allocator**

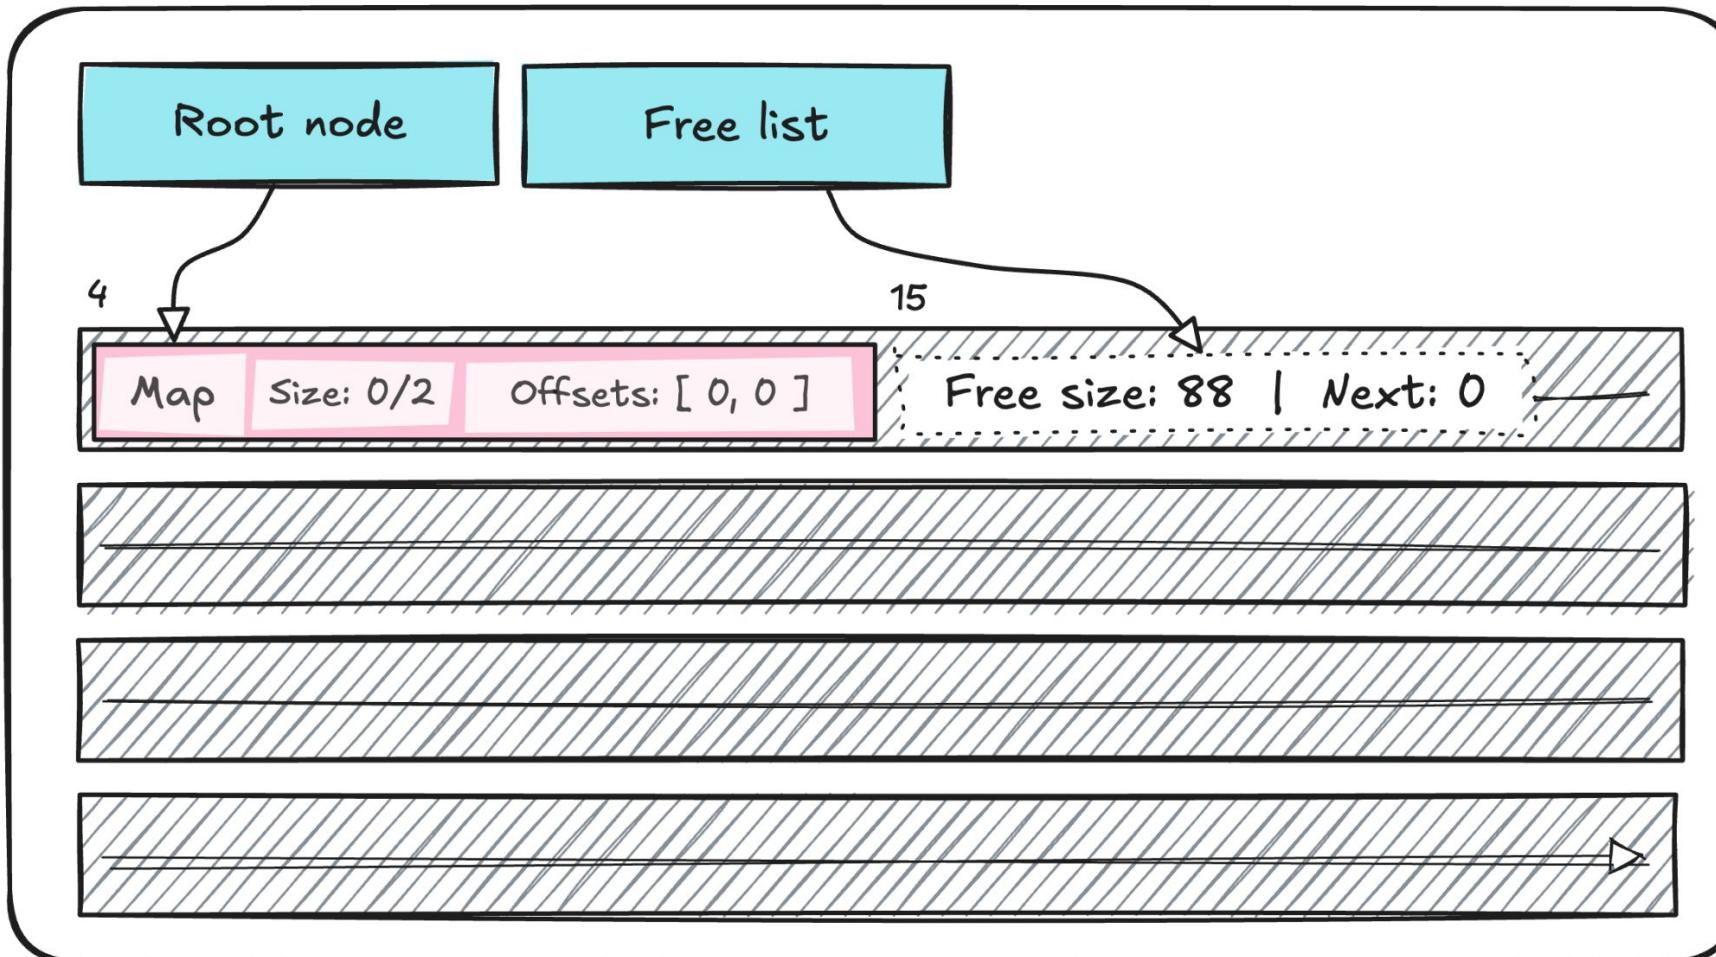
# Free List Allocator

- This would look like this:



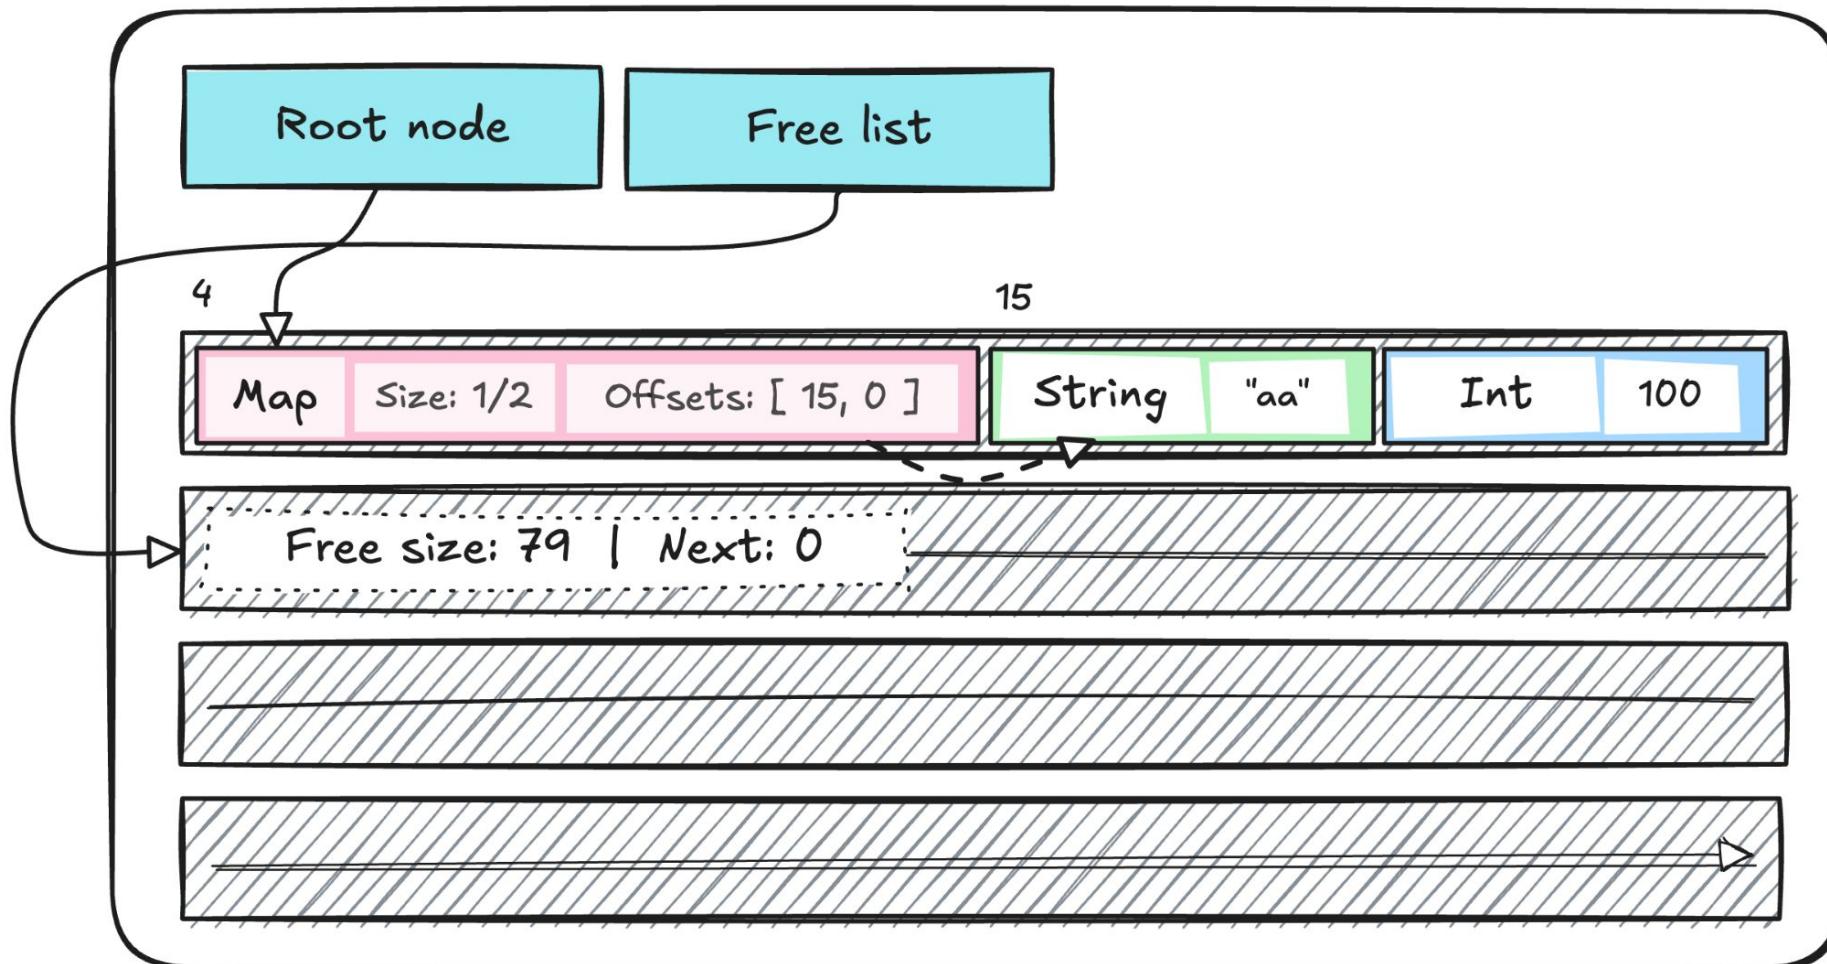
# Free List Allocator

- This would look like this:



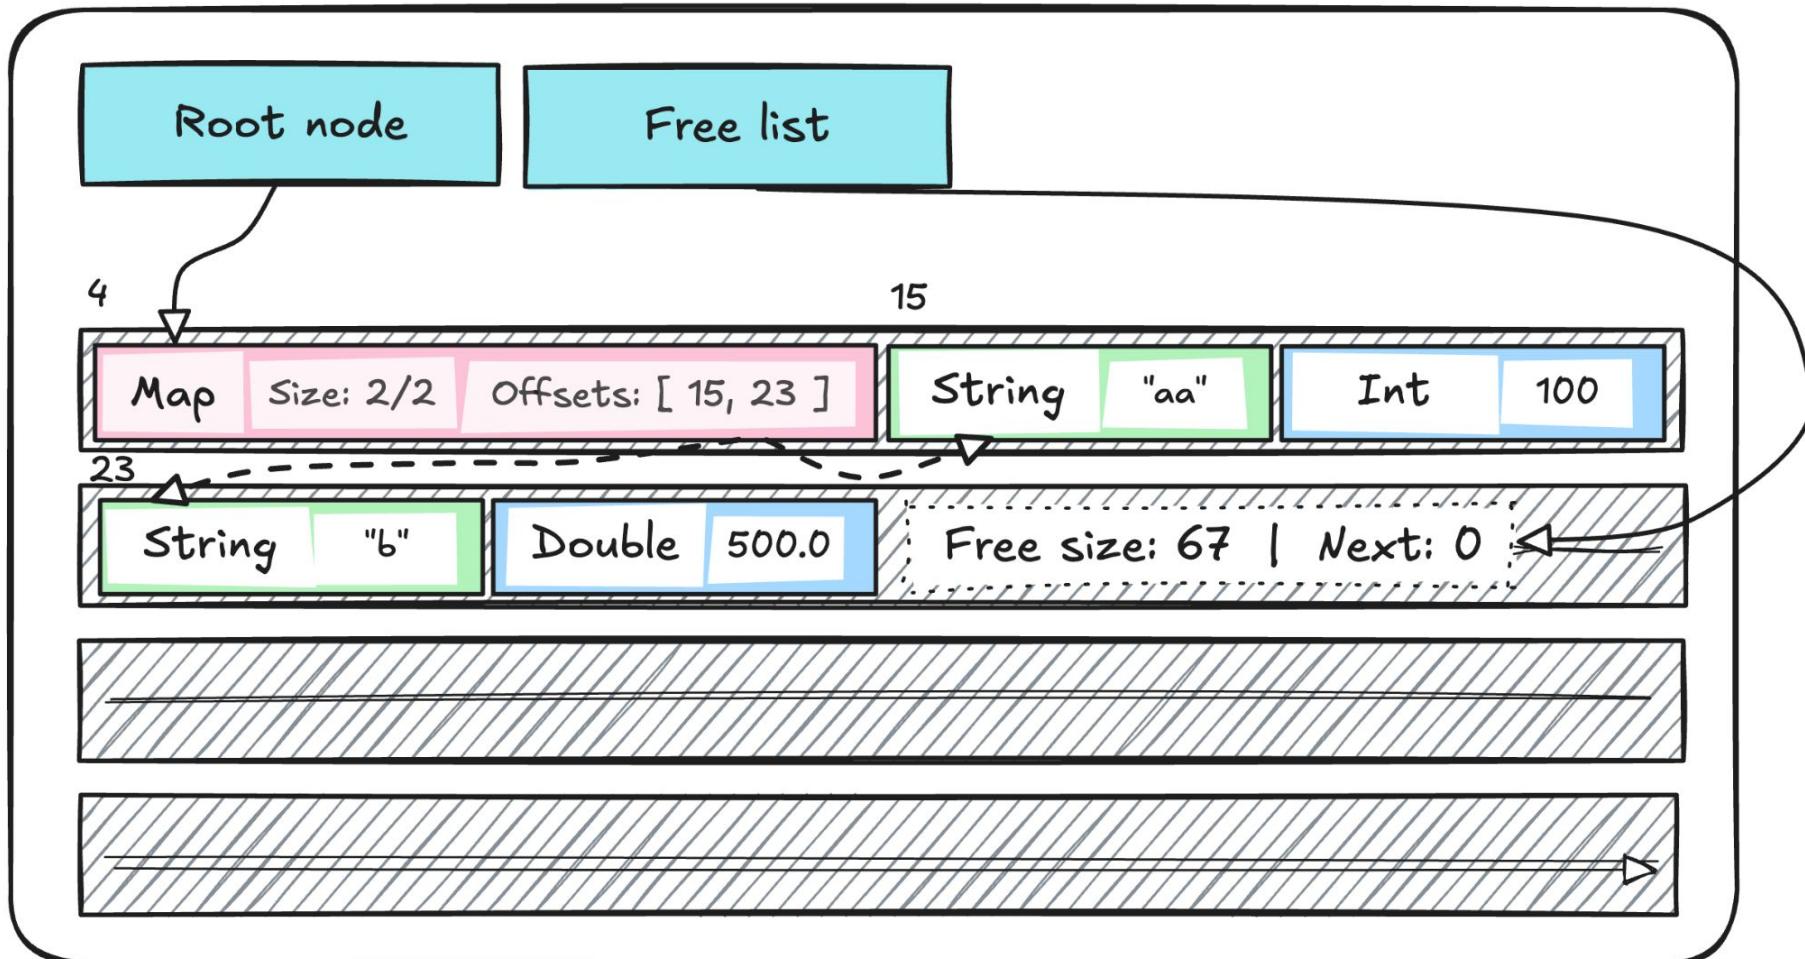
# Free List Allocator

- This would look like this:



# Free List Allocator

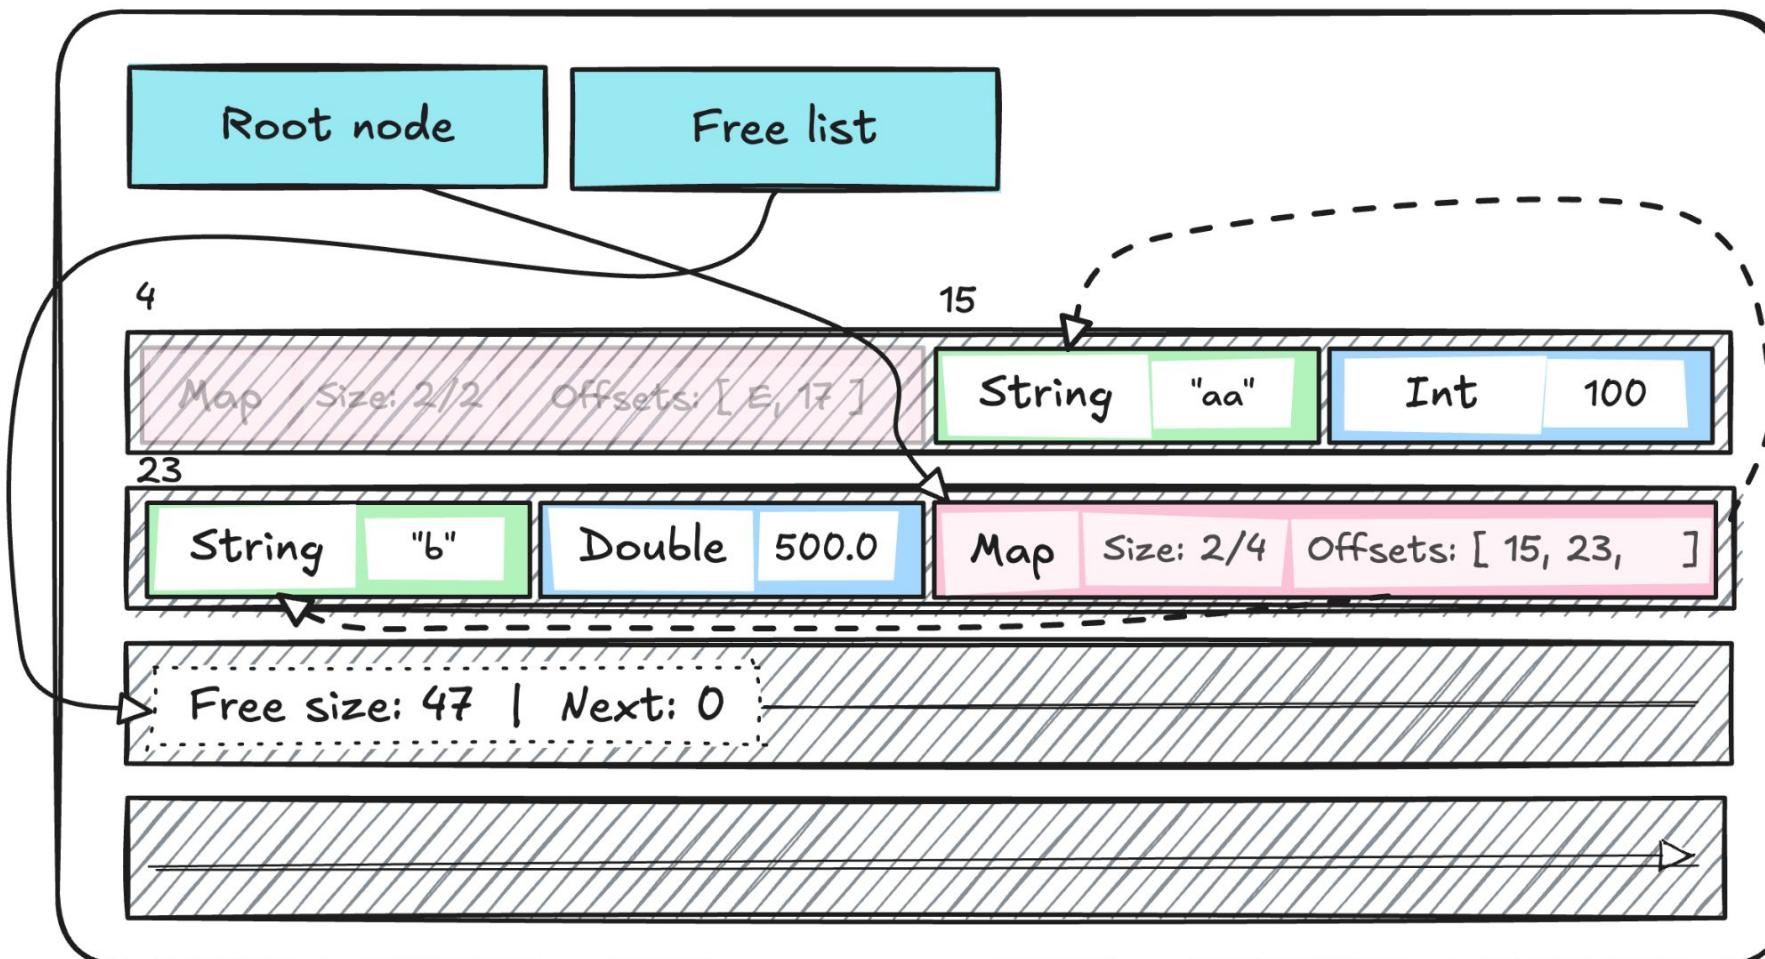
- This would look like this:



Bloomberg  
Engineering

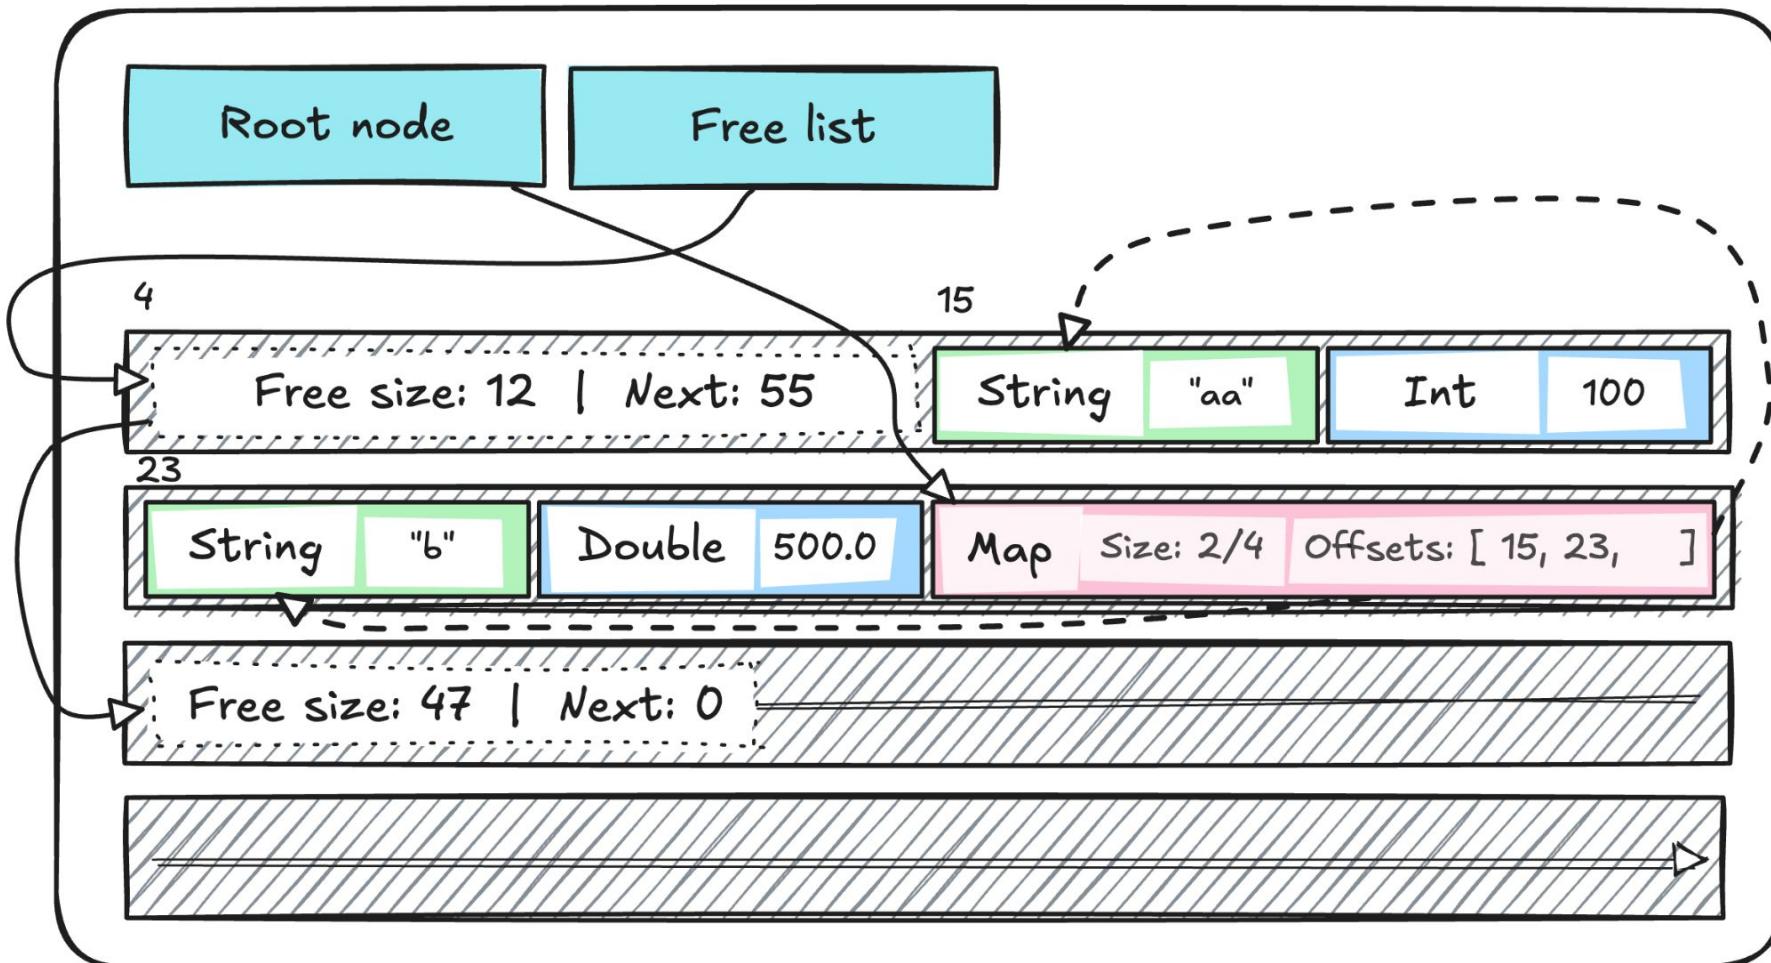
# Free List Allocator

- This would look like this:



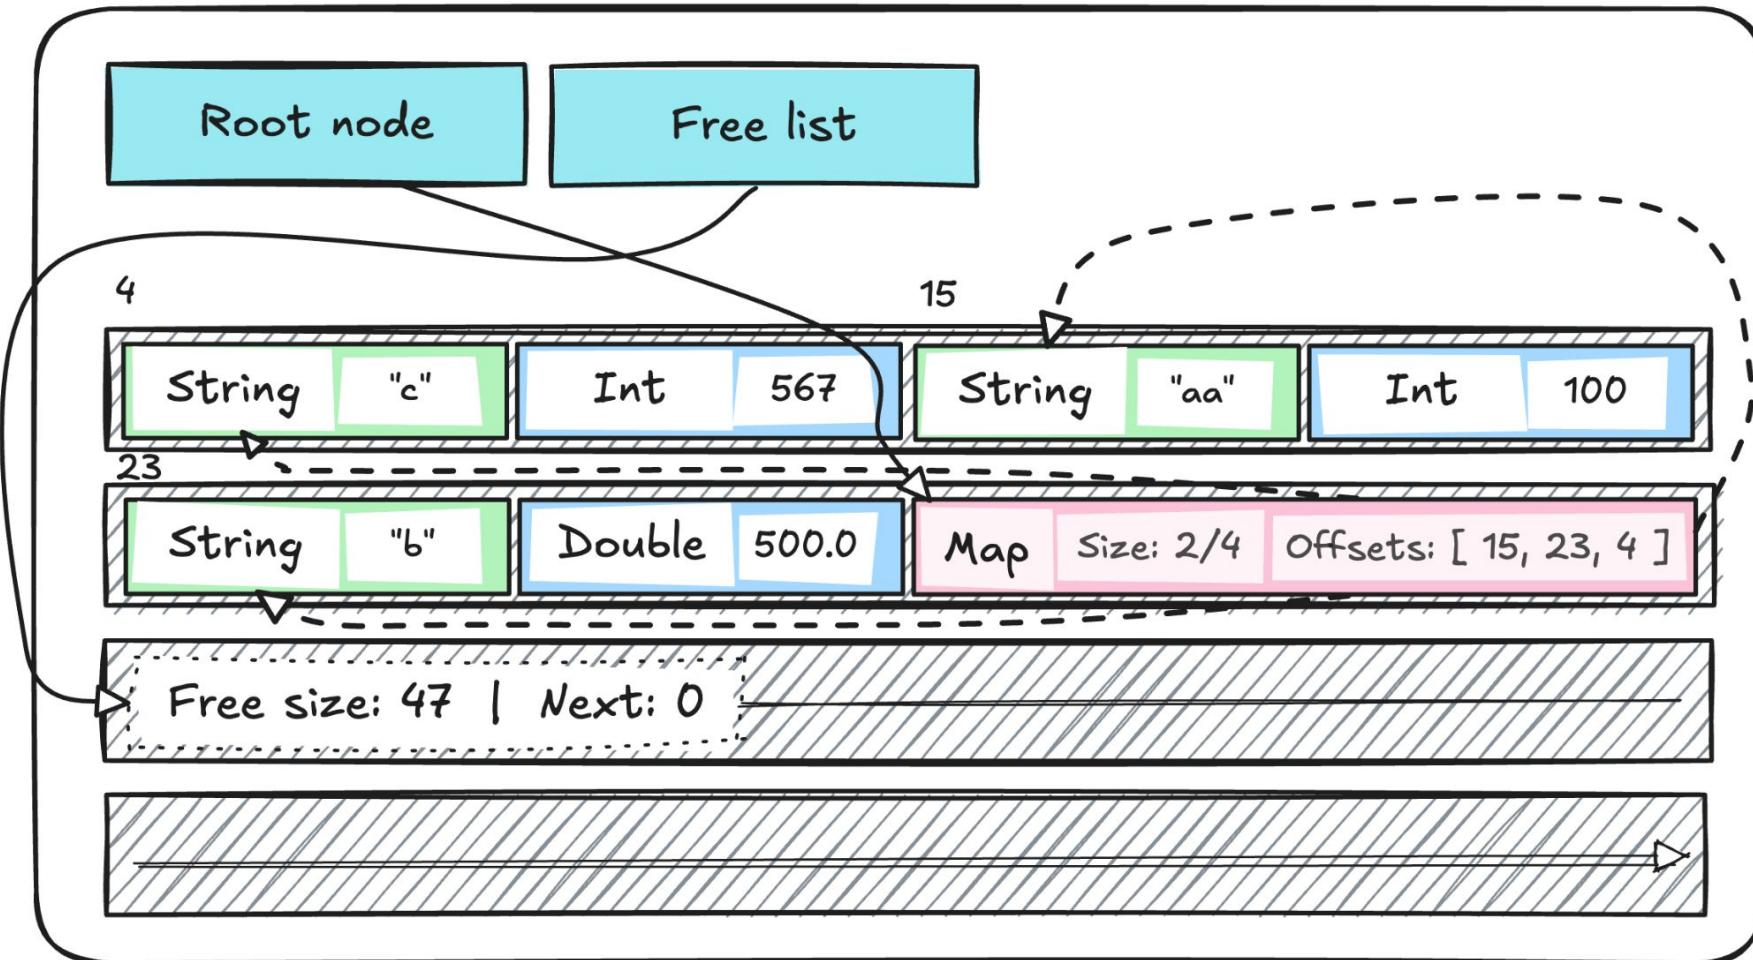
# Free List Allocator

- This would look like this:



# Free List Allocator

- This would look like this:



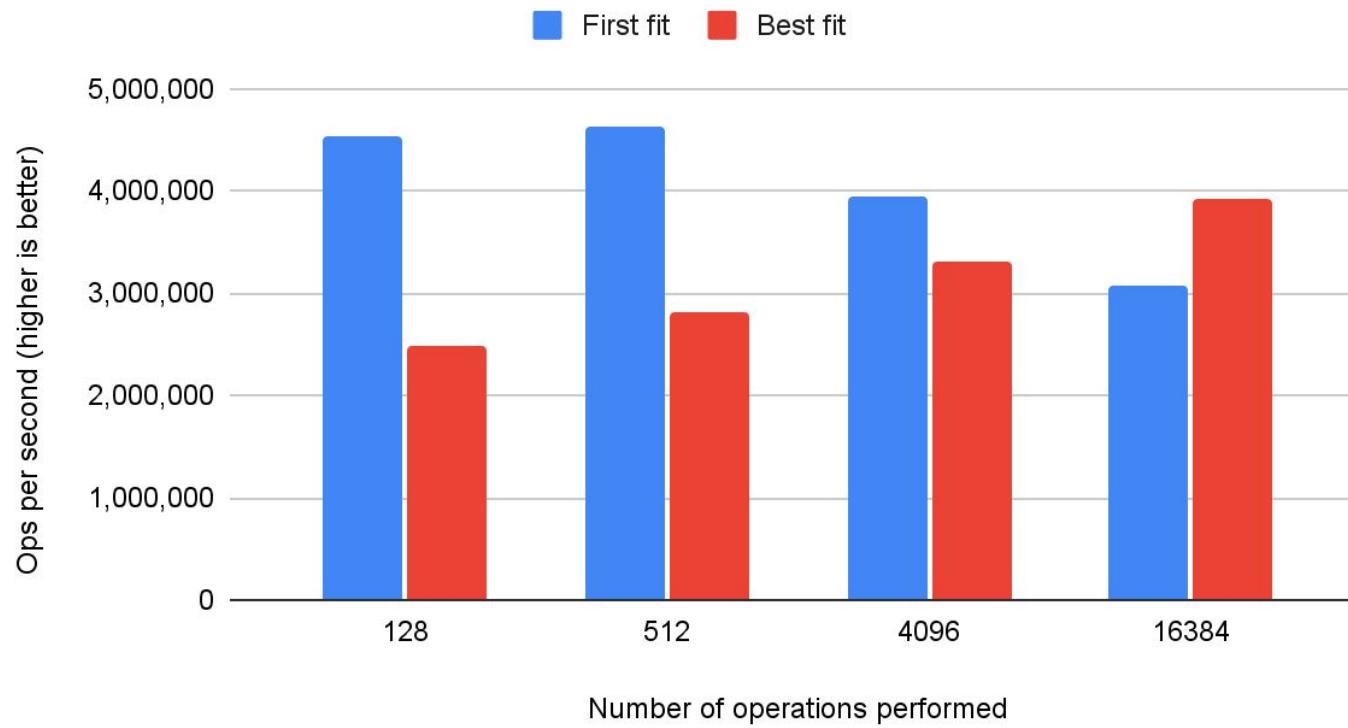
# Free List Allocator

- We can satisfy new allocations using a few schemes:
  - First fit – return the first block I find that fits our allocation
  - Best fit – return the smallest block I can find that fits our allocation
- I benchmarked both schemes, in terms of both wasted space and performance

# Free List Allocator

- First-fit performs better initially, but performance degrades as more operations are performed on the buffer
- After many operations, best-fit wins

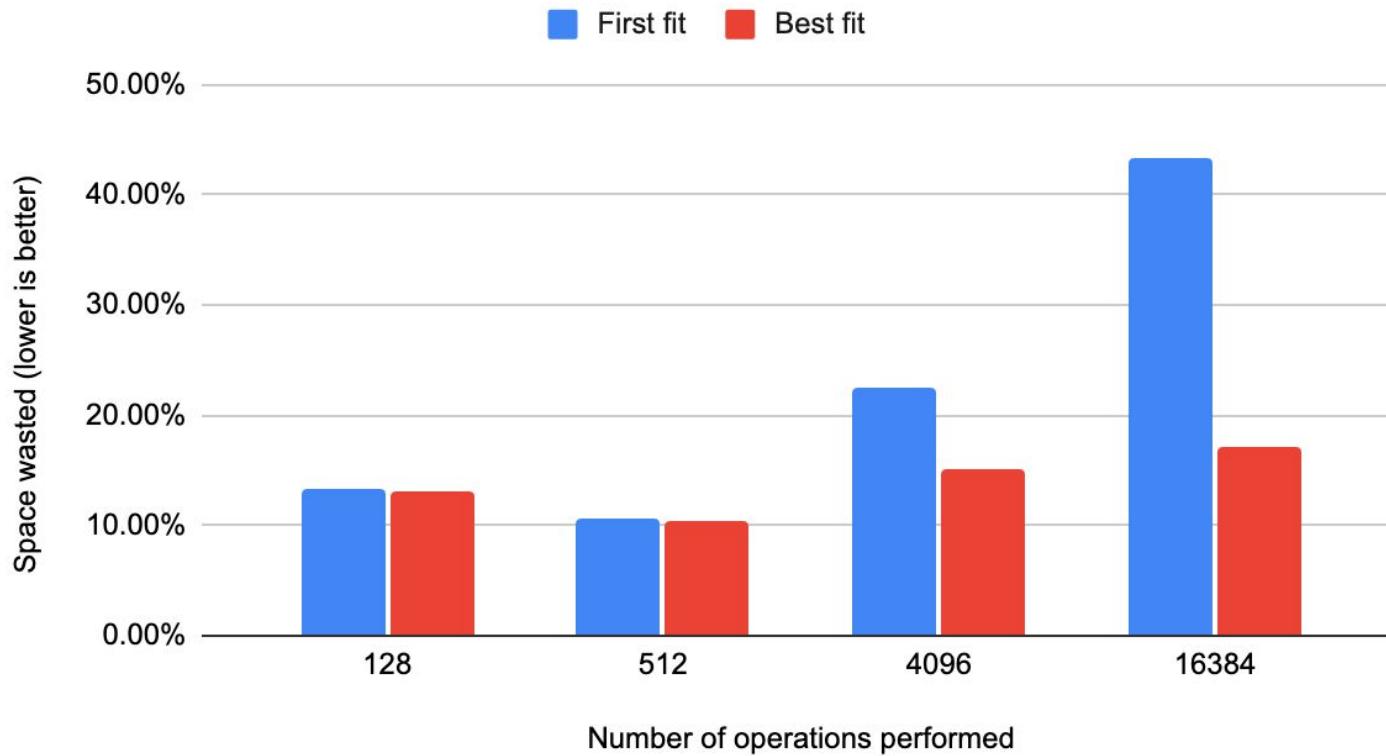
Write-heavy workflow



# Free List Allocator

- With first-fit, space wastage increases considerably the more operations are performed:

Write-heavy workflow



# Free List Allocator – Profiling

- Profiling allocation shows a lot of time is spent reading the size of each free-list chunk
- The free-list chunks are scattered throughout the buffer – so they might not be in cache



```
size > TaggedUtil::freeSpaceSize(&d_buffer[freeListOff])) {  
    freeListPtr = TaggedUtil::freeSpacePtr(&d_buffer[freeListOff]);  
    freeListOff = *freeListPtr;  
    movzwl 0x3(%rcx),%ecx  
    while (freeListOff != 0 &&  
    test  %rcx,%rcx  
→ je    2438f8 <experimental::ser::adv::container::Container::allocate(int,  
_ZNSt6vectorIhSaIhEEixEm():  
    return *(this->_M_impl._M_start + __n);  
    add   %r11,%rcx  
_ZN12experimental3ser6tagged10TaggedUtil13freeSpaceSizeEPKh():  
    return (data[2] << 8) | data[1];  
[44.66]  
    movzwl 0x1(%rcx),%esi  
_ZN12experimental3ser3adv9container9Container8allocateEiPPhS5_S5_():  
    cmp   %r9,%rsi  
→ jb    243828 <experimental::ser::adv::container::Container::allocate(int,
```

# Faster Modifications

Idea 1: Move items “out-of-line” of their containing object

~~Idea 2: Implement a bump allocator~~

Idea 3: Implement a free-list allocator

**Idea 4: Prefetch free-list entries while iterating**

# Free List Allocator – Prefetching

- We can issue a prefetch for the next item on the free list, before we know if `freeListOff != 0`

```
uint16_t *freeListPtr = &hdr->freeListHead;
uintptr_t freeListOff = *freeListPtr;

__builtin_prefetch(&d_buffer[freeListOff]);

while (freeListOff != 0 &&
       size > freeSpaceSize(&d_buffer[freeListOff])) {

    freeListPtr = freeSpacePtr(&d_buffer[freeListOff]);
    freeListOff = *freeListPtr;

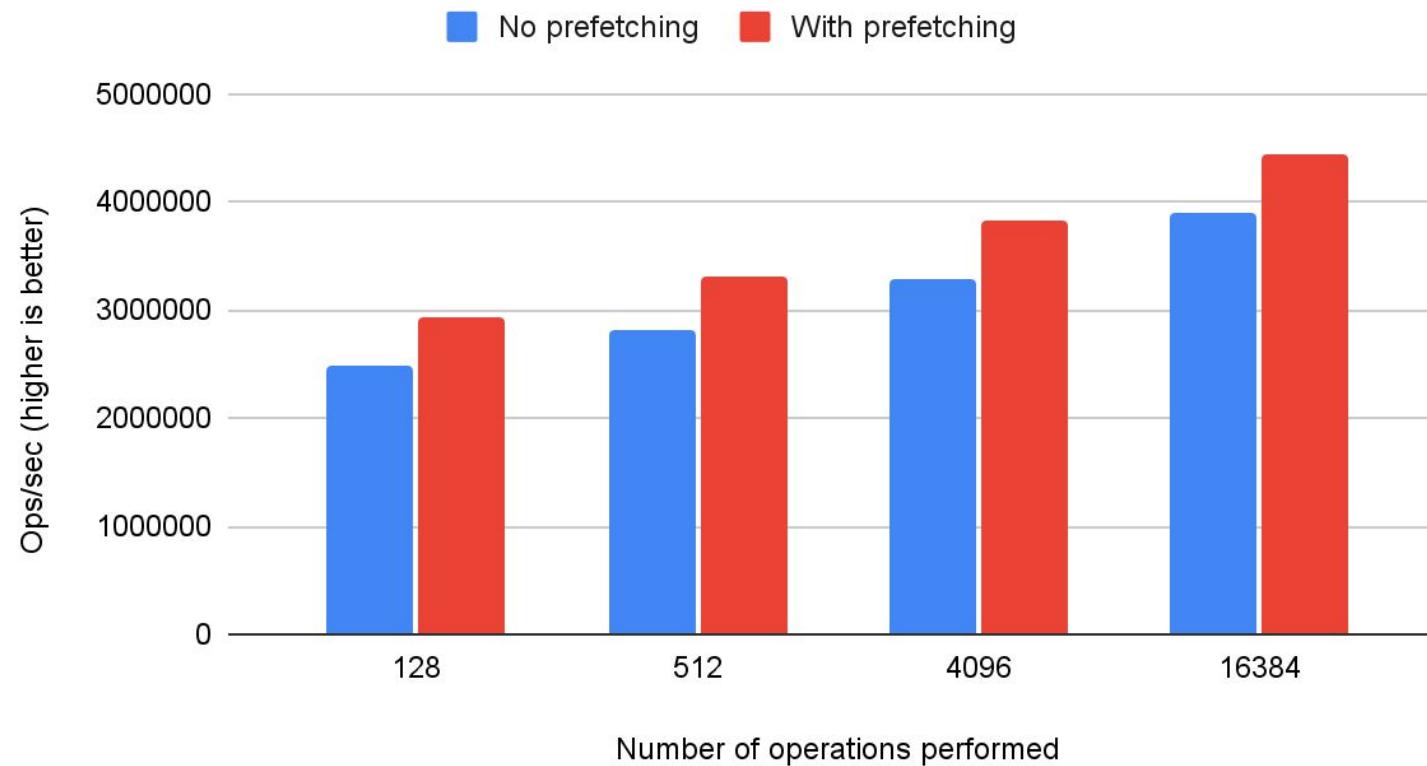
    __builtin_prefetch(&d_buffer[freeListOff]);
}
```

If `freeListOff` is zero, we do a harmless prefetch of our buffer header, which is very likely already in cache anyway

# Free List Allocator – Prefetching

- This gives us a ~15% performance boost with best-fit
- No improvement for first-fit

Write heavy-workflow, best-fit allocation



# Free List Allocator – Prefetching

- A word of warning on prefetching:
  - After getting this result, I experimented with adding prefetching in other parts of the code
  - ... But it uniformly made things worse!
  - So the hardware prefetcher is generally pretty good
  - Prefetching is only likely to be beneficial if you have genuinely unpredictable access patterns

# Faster Modifications

Idea 1: Move items “out-of-line” of their containing object

~~Idea 2: Implement a bump allocator~~

Idea 3: Implement a free-list allocator

Idea 4: Prefetch free-list entries while iterating

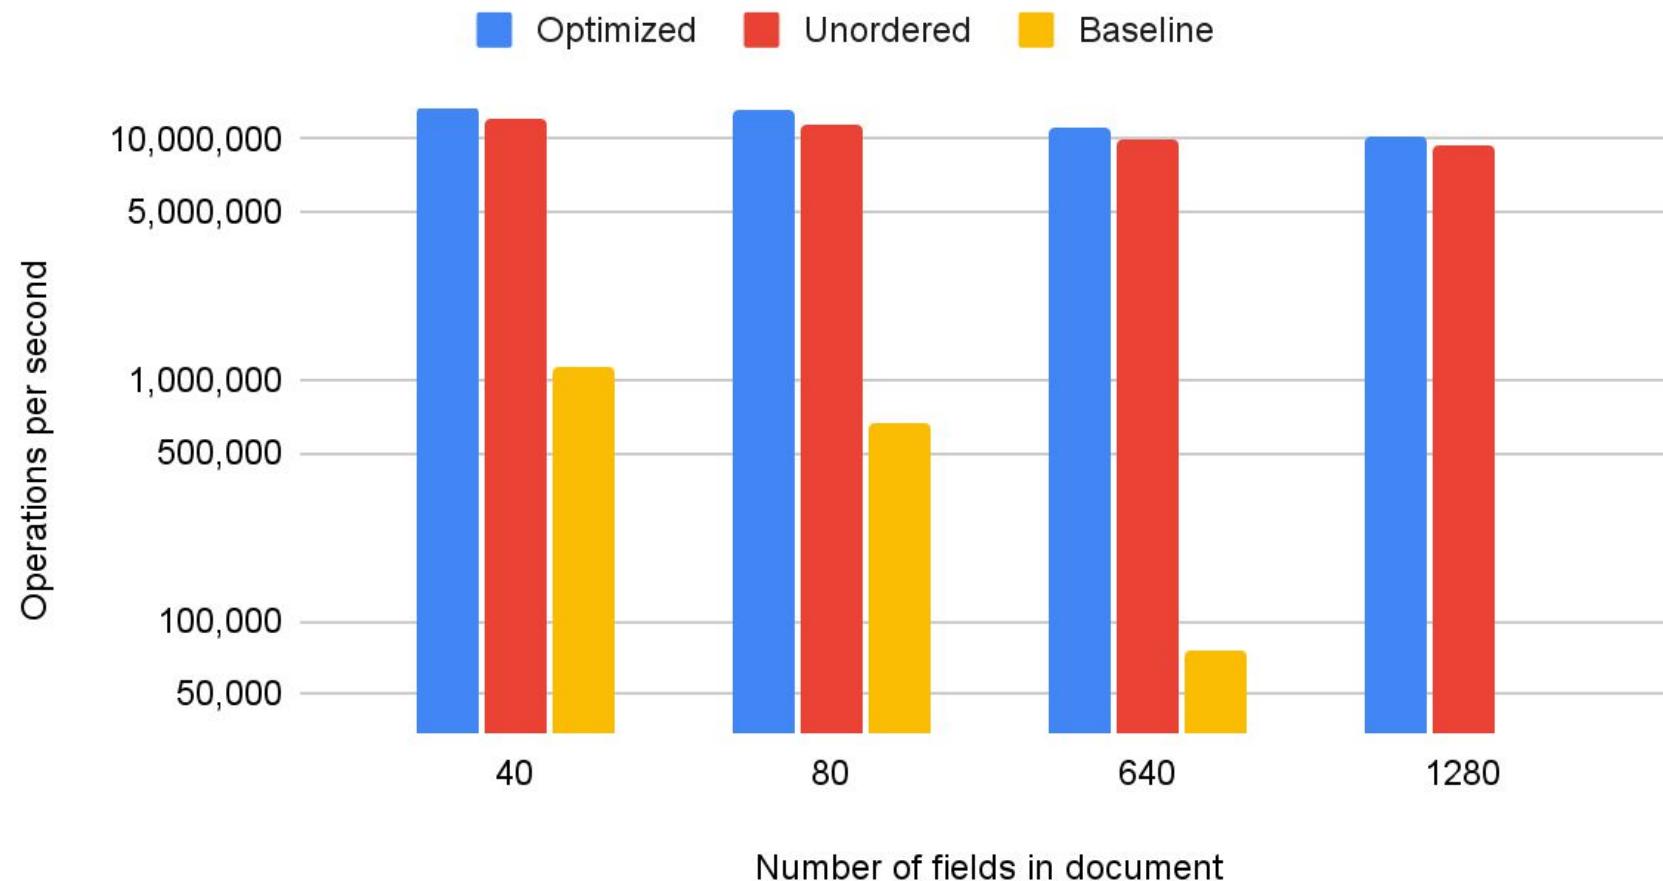
**I'm happy with modification performance now!**

# Final Benchmarks

- For my final benchmarks I decided to compare the format against the baseline format, and against unordered\_map
- In the unordered\_map version:
  - Our unordered\_maps have two levels of nesting
  - Our leaf values are variants of int, double, string, and vector of int
  - I used heterogeneous lookup to prevent excessive string allocations
  - I use the same fast hash function as my format

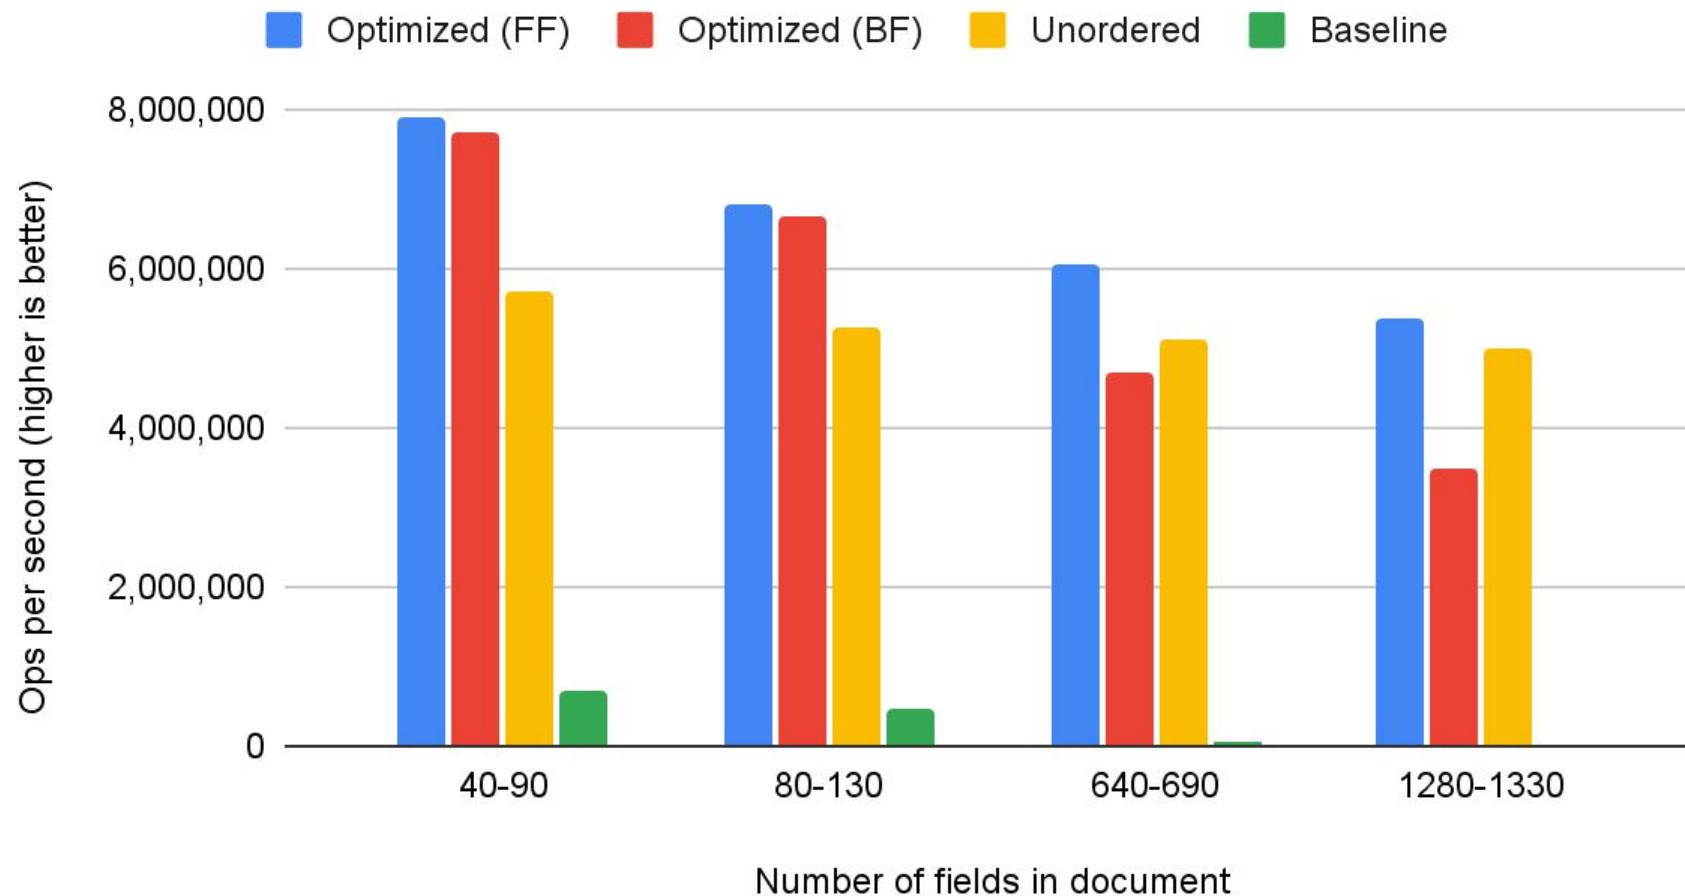
# Final Benchmarks

Read operations only



# Final Benchmarks

## Mixed operations



# Limitations

- Fragmentation is a problem with first-fit
  - Can we achieve a middle-ground between first- and best-fit?
- Because of the added metadata necessary for performing searches, our format is not as compact as other schemaless formats
  - Our overhead is ~4 bytes per key/value pair in a map
  - Or ~2 bytes per value in an array
- Our format is full of offsets, so if we load a corrupted buffer, we could easily make out-of-bounds accesses
  - Our current implementation assumes it deals with a valid input buffer, so the results are based on that

# Conclusions

- We've created a fast, schemaless serialization format which we can edit in place
- It compares favourably to using an unordered\_map in memory
- By creating our own memory allocator and replacing pointers with offsets, we can put in-memory data structures into a small buffer

# Conclusions

- When trying to improve performance:
  - Create highly parameterized benchmarks
  - Profile
  - Make algorithmic/data structure improvements
  - Explore vectorization, branchless code, prefetching

# Future Work

- Explore how to compact the buffer efficiently
  - Compact in-place, or copy into a new buffer?
  - Could recent SIMD scatter/gather instructions help here?
- Explore building an offset-free serialization format
  - Could we build indexes of our data as needed, instead of storing them in the buffer?
  - Taking inspiration from things like `simdjson`
- Explore compressing repeated keys or fragments of keys to save space

# Resources

- “Algorithms for Modern Hardware” by Sergey Slotin, available online:  
<https://en.algorithmica.org/hpc/>
  - A comprehensive review of techniques for writing high performance code
- “The Garbage Collection Handbook” by Richard Jones, Antony Hosking and Eliot B. Moss
  - Covers a lot of memory allocation and buffer compaction techniques
- “Going Nowhere Faster” by Chandler Carruth at CppCon 2017:  
<https://www.youtube.com/watch?v=2EWejmKlxs>
  - Good intro to benchmarking and profiling tools, including perf
- Fedor Pikus’s “Performance and Efficiency in C++” training
  - Covers all this stuff in a lot of depth!

# Thank you!

We are hiring: [bloomberg.com/engineering](https://bloomberg.com/engineering)

Engineering

Bloomberg

TechAtBloomberg.com

© 2025 Bloomberg Finance L.P. All rights reserved.