

+ 25

# Clean Code! Horrible Performance?

SANDOR DARGO



20  
25 | A graphic of two white mountain peaks with a yellow square at the top of the taller one, positioned next to the year '2025'.  
September 13 - 19

# Who Am I?


Sándor DARGÓ

Senior Engineer at **Spotify**

Enthusiastic blogger: [sandordargo.com](http://sandordargo.com)

Curious oenophile

Fortunate father of two



Have you ever seen engineers  
prematurely optimizing their  
code?

Have you ever prematurely  
optimized your code?

Did it feel great?

Did it bring more value  
to the business?

# Why this talk?

My earlier talk - *Why clean code is not the norm* - lead to discussions about performance

I don't like absolutes but too many people talk in absolutes

I think engineering is more nuanced

# Agenda

What is clean code?

What is software quality?

Where and why do we use C++?

What is software performance?

Is clean code actually slower?

Should we drop clean code for faster code?

Does clean code imply horrible performance?

# What is clean code?

What is clean code?

What is software quality?

Where and why do we use C++?

What is software performance?

Is clean code actually slower?

Should we drop clean code for faster code?

Does clean code imply horrible performance?

# Clean code is *NOT* a closed set of rules

The book Clean Code != The concept of clean code

Depends on the language


Depends on the context

Changes over time

Not a silver bullet

# Then what is clean code?

Code that is easy to understand  
and easy to change



Clean code is part of software  
quality!


But is high quality software  
important?

# Quality isn't so important when you

Build a proof of concept

Search a product market fit

Must deliver (fast) to survive



# What is software quality?

---

What is clean code?

What is software quality?

Where and why do we use C++?

What is software performance?

Is clean code actually slower?

Should we drop clean code for faster code?

Does clean code imply horrible performance?

# What is quality?

Quality is undefinable  
You recognize it when you see it

# Software quality has several definitions

It might even be a meaningless term

According to Derek Jones ([The aura of software quality](#))

*“People in industry are very interested in software quality, and sometimes they have the confusing experience of talking to me about it. My first response, on being asked about software quality, is to **ask what the questioner means by software quality**. After letting them fumble around for 10 seconds or so, trying to articulate an answer, I **offer several possibilities (which they are often not happy with)**. Then I explain how “software quality” is a meaningless marketing term. This leaves them confused and unhappy. People have a yearning for software quality which makes them easy prey for the snake-oil salesmen.”*

# CISQ defines the 4 pillars of structural quality

Consortium for Information & Software Quality™

Security

Reliability

Maintainability

Performance efficiency

# *Clean code is part of software quality!*

Helps reducing the number of bugs (**Security, Reliability, Maintainability**)

Helps reducing the time to fix bugs (**Reliability, Maintainability**)

Decreases maintenance costs (**Maintainability**)

# CISQ defines the 4 pillars of structural quality

Consortium for Information & Software Quality™

Security

Reliability

Maintainability

Performance efficiency



# Where do we use C++?

What is clean code?

What is software quality?

Where and why do we use C++?

What is software performance?

Is clean code actually slower?

Should we drop clean code for faster code?

Does clean code imply horrible performance?

# C++ is just everywhere



# Why do we use C++?

What is clean code?

What is software quality?

Where and why do we use C++?

What is software performance?

Is clean code actually slower?

Should we drop clean code for faster code?

Does clean code imply horrible performance?

# Because it's legacy!

It's been with us since the 80s

Tons of (often complex and critical) existing code

Rewrites rarely make sense even if C++ is not needed

Loss of time

Problem of existing ecosystem

Problem of knowledge

# C++ is also evolving!

Since C++11, stable and steady evolution

A new standard every 3 years

Safer and more expressive code

**Not owned by a  
single corporation!**

# C++ has an economic advantage!

It's fast!

Good for performance critical apps!


It uses little space!

Good for embedded

For anything that has to be distributed

Consumes less energy!

Good for everyone!



	Energy
(c) C	1.00
(c) Rust	1.03
(c) C++	1.34
(c) Ada	1.70
	1.98
	2.14
	2.18
	2.27
	2.40
	2.52
	2.79
	3.10
	3.14
	3.23
	3.83
	4.13
	4.45
	7.91
(i) TypeScript	21.50
(i) Hack	24.02
(i) PHP	29.30
(v) Erlang	42.23
(i) Lua	45.98
(i) JRuby	46.54
(i) Ruby	69.91
(i) Python	75.88
(i) Perl	79.58

	Time
(c) C	1.00
(c) Rust	1.04
(c) C++	1.56
(c) Ada	1.85
(v) Java	1.89
(c) Chapel	2.14
(c) Go	2.83
(c) Pascal	3.02
(c) Ocaml	3.09
(v) C#	3.14
(v) Lisp	3.40
(c) Haskell	3.55
(c) Swift	4.20
(c) Fortran	4.20
(v) F#	6.30
(i) JavaScript	6.52
(i) Dart	6.67
(v) Racket	11.27
(i) Hack	26.99
(i) PHP	27.64
(v) Erlang	36.71
(i) JRuby	43.44
(i) TypeScript	46.20
(i) Ruby	59.34
(i) Perl	65.79
(i) Python	71.90
(i) Lua	82.91

	Mb
(c) Pascal	1.00
(c) Go	1.05
(c) C	1.17
(c) Fortran	1.24
(c) C++	1.34
(c) Ada	1.47
(c) Rust	1.54
(v) Lisp	1.92
(c) Haskell	2.45
(i) PHP	2.57
(c) Swift	2.71
(i) Python	2.80
(c) Ocaml	2.82
(v) C#	2.85
(i) Hack	3.34
(v) Racket	3.52
(i) Ruby	3.97
(c) Chapel	4.00
(v) F#	4.25
(i) JavaScript	4.59
(i) TypeScript	4.69
(v) Java	6.01
(i) Perl	6.62
(i) Lua	6.72
(v) Erlang	7.20
(i) Dart	8.64
(i) JRuby	19.84

# What is software performance?

What is clean code?

What is software quality?

Where and why do we use C++?

What is software performance?

Is clean code actually slower?

Should we drop clean code for faster code?

Does clean code imply horrible performance?

# Performance is...

A non-functional property, like reliability or usability

Often listed as a requirement — but rarely defined

A measure of efficiency — like cost, time to deliver

Like quality, it's context-dependent and multi-dimensional

Sufficient performance ensures  
that the software can achieve its  
functional requirements in a  
**reasonable** timespan

# Performance != Big-O notation

Language is important

Hardware is important

I/O can - and will - dominate

Moreover, it's not only about execution speed

# Dimensions That Actually Matter

**Response time:** how fast it responds

**Throughput:** how many requests it can handle

**Resource utilization:** how efficiently it uses system resources

**Scalability:** ability to maintain performance levels as the workload increases

**Capacity:** maximum load or number of users before severe performance degradation

**Stability:** ability to perform consistently over time without crashes, memory leaks, or other issues

# The Real Enemies of Performance

Inefficient algorithms

Unmeasured assumptions

Abstraction overload

Bad data access patterns

Network & I/O bottlenecks

# Is clean code actually slower?

What is clean code?

What is software quality?

Where and why do we use C++?

What is software performance?

Is clean code actually slower?

Should we drop clean code for faster code?

Does clean code imply horrible performance?

# Clean code is actually slower

Clean code is actually slower  
than what?

Clean code is not the fastest of  
all possible codes



# Clean code is optimization for maintainability

It's not the most efficient / fastest code

It's not the smallest code either

But will it **harm** performance?

Optimization is always about  
compromises

# Optimization is always about compromises

## Optimize for readability

You will hurt runtime performance

You will hurt binary size

## Optimize for runtime performance

You will hurt readability

You might hurt portability and binary size

## Optimize for binary size

You might hurt readability and portability

You will hurt runtime performance

# Should we drop clean code for faster code?

---

What is clean code?

What is software quality?

Where and why do we use  
C++?

What is software  
performance?

Is clean code actually slower?

Should we drop clean code  
for faster code?

Does clean code imply  
horrible performance?

# Sometimes every bit of performance matters

Every 100ms increase in page load time costs Amazon 1% in sales

For every additional second of page load time, 10% of BBC users leave

A broker could lose \$4 million/ms in revenue, if a competitor is 5 ms ahead

An extra 500ms in search page generation time, dropped Google's traffic by 20%

Akamai found that every 100ms delay in a website load time drops the conversion rate by 7%

# Performance is crucial for Spotify too when you

Open the app

Start playing a track

Tap on any button

# Otherwise, there are more important aspects

Small binary

Thinking about low bandwidth

And embedded devices

Maintainable code

We're almost 20 years old!

With >7k employees!

But even when performance matters, the first answer is not low-level code optimization

Optimize for maintainability first,  
then for performance if needed


Start by writing clean code, clean architecture and later deal with performance - if needed

# Algorithmic complexity isn't always so important

My very first project back in 2013 suffered from

Slow speed

Immense memory leak



# Some said speed, culprit was an $O(n^4)$ “algo”

I wasn't really bothered by that piece of nice code...

I measured and measured

It took a fraction of the execution time



*"The overall performance improvement gained by optimizing a single part of a system is limited by the fraction of time that the improved part is actually used"* - Amdahl's law

You can often ignore slow code if  
it's not on the hot path

# The key was to use a 3rd party library properly

Some startup / cleanup functions were called too often

Calling them at the right time solved all the imminent issues



**Reading documentation and  
using libraries as they are  
intended to is more important  
than writing fast code...**

# How to support 15x traffic?

Instead of making the code “faster”:

Make your system more scalable

Improve operability

Improve resource utilization

Finetune network settings

# Improve operability

Remove memory leaks

Fix core dumps

Remove as much UB as possible

# Consolidate external requests

Optimize orders of magnitude slower network/db calls first

If the API offers some batching option, use it!

If not, try to batch it on call site

Think about (local) caching

Think about anything that reduces the number of calls you have to make

But what about all those costly  
language features?

# Will the common enemies ruin your performance?

Unnecessary heap allocations?


Copying large objects unnecessarily?

Unnecessary virtual functions?

You won't have many of those if your code is clean

“Some” usage will not ruin your code

The ONLY VALID MEASUREMENT  
OF CODE QUALITY: WTFs/MINUTE



# Optimization efforts must be concentrated

Overall performance improvement gain in a single part is limited!

Follow best practices, keep your code clean you'll have a reasonable performance

Start optimizing the hot path

Leave the non-hot path clean as long as you prove else

Performance is not l'art pour l'art



*Many devs in a  
code review*



# Software engineering is not about performance

Engineering is in the crossroads of defining and recognizing what is good enough


Just because we **can** do something, we don't have to

Software engineering is not competitive programming


Our goal is to help achieve business goals - for the long term

# But why maintainability is so important?


## Cost per CPU Clock Cycle (1970-2020)



## Software Engineer Salaries (1970-2020)



# Software Engineer Salaries vs CPU Cost per Clock Cycle (1970-2020)



# Performance is not l'art pour l'art

We have “cheap” and strong hardware

We also have expensive developers

Unless you have other strict requirements, optimize for maintainability

# Does clean code imply horrible performance?

---

What is clean code?

What is software quality?

Where and why do we use  
C++?

What is software  
performance?

Is clean code actually slower?

Should we drop clean code  
for faster code?

Does clean code imply  
horrible performance?

# Sometimes

Sometimes  
If requirements are extreme...

Sometimes  
If requirements are extreme...  
In most cases, it won't

Sometimes  
If requirements are extreme...  
In most cases, it won't  
Not even in the realm of C++

# Optimize for maintainability for better ROI

Keep your code **clean, understandable** and **extensible**

Aim for **sufficient** performance

**Optimize** your hot path **if needed**

Focus on **delivering business value** for the long run

+ 25

# Clean Code! Horrible Performance?

SANDOR DARGO



20  
25 | A graphic of two white mountain peaks with a yellow triangle at the top of the taller one, positioned next to the year '2025'.  
September 13 - 19