

25

# Connecting C++ Tools to AI Agents Using the Model Context Protocol (MCP)

BEN MCMORRAN



**Cppcon**  
The C++ Conference

20  
25




September 13 - 19

A brief, incomplete history  
of LLMs in developer tools



# Chat is born

## ChatGPT



November 2022

Chat interface

LLM


Type “Help me fix this test.”  
Paste unit test source code  
Press send

Compute completion text

Display result

November 2022


# IDE integration and better context



Context providers

Chat interface

LLM



Early 2023

# Automatic context (aka RAG)

Help me fix this test.

✓ Read [test\\_calculator.cpp](#)

The issue with your test lies in the logic of the

Context providers

Chat interface

LLM

User sends “Fix this test.”

Request relevant context

Search files for keywords

Build final prompt text

Compute completion text


Display result

Mid 2023

Context providers

Chat interface

LLM



Mid 2023

# Agents and tool calling

Now let me compile and run the test to verify it works:

Run command in terminal

```
g++ -o test_calculator test_calculator.cpp
```

Compiling the C++ test file

Continue

Cancel

# Tools provide **information** or take **actions**

```
{  
  "name": "add",  
  "description": "Computes the sum of two numbers",  
  "parameters": {  
    "type": "object",  
    "properties": {  
      "a": {  
        "type": "number",  
        "description": "The first number to add"  
      },  
      "b": {  
        "type": "number",  
        "description": "The second number to add"  
      },  
    },  
    "required": ["a", "b"]  
  }  
}
```




- Client** Send What is 37 + 94? (plus tool metadata)
- LLM** 37 + 94 is <tool>add(37,94)</tool>
- Client** Locally run add() function. Result is 131.
- Client** Send 37 + 94 is <tool>add(37,94)</tool>  
<tool\_response>131</tool\_response>
- LLM** 131
- Client** Render response as “37 + 94 is 131”

|

Extensible?  
Tool providers

Chat interface

LLM





Late 2024 – early 2025

# Interoperable tools

Chat interfaces



Tool providers




November 2024

# Interoperable tools

MCP clients



MCP servers




Model Context  
Protocol




November 2024

# MCP under the hood



# MCP under the hood

initialize  
notifications/initialized  
tools/list  
tools/call  
  
prompts/...  
resources/...  
sampling/createMessage  
elicitation/create  
completion/complete  
logging/setLevel  
notifications/...  
roots/list  
ping



Demo: write an MCP server

Best practices for MCP

# You might not need MCP if...

...you want to add a small amount of information to every request.

```
1 ## Functions  
2  
3 - Write short functions with a single purpose. Less than 20 instructions.  
4 - Name functions with a verb and something else.  
5 - If it returns a boolean, use isX or hasX, canX, etc.  
6 - If it doesn't return anything (void), use executeX or saveX, etc.
```

Custom instructions/rules and prompt files are commonly available

You might not need MCP if...

...you can leverage built-in tools like run\_in\_terminal or web search.

# You might not need MCP if...


...you only need to support one IDE or chat interface.

```
export interface LanguageModelTool<T> {
    /**
     * Invoke the tool with the given input and return a result.
     *
     * The provided {@link LanguageModelToolInvocationOptions.input} has been validated against the declared schema.
     */
    invoke(options: LanguageModelToolInvocationOptions<T>, token: CancellationToken): ProviderResult<LanguageModelToolResult>;
    
    /**
     * Called once before a tool is invoked. It's recommended to implement this to customize the progress message that appears
     * while the tool is running, and to provide a more useful message with context from the invocation input. Can also
     * signal that a tool needs user confirmation before running, if appropriate.
     *
     * *Note 1:* Must be free of side-effects.
     * *Note 2:* A call to `prepareInvocation` is not necessarily followed by a call to `invoke`.
     */
    prepareInvocation?(options: LanguageModelToolInvocationPrepareOptions<T>,
        token: CancellationToken): ProviderResult<PreparedToolInvocation>;
}
```

Platform-specific APIs to provide tools are often richer and easier


# You might not need MCP if...

...you have a highly structured workflow.



Better to write a normal program and call the LLM as needed.

# Security and the "lethal trifecta"



# Security and the "lethal trifecta"

Private data:  
**Internal  
documents**

CVE-2025-32711

Communication:  
**Pre-fetching  
image links**

Untrusted content:  
**Email from  
attacker**

# Help the LLM use your tools

## Write **great** tool descriptions

- What does it do?
- What information does it return?
- When should and shouldn't it be used?
- Any important limitations or constraints?

The LLM itself can help write these!

This tool allows you to execute shell commands in a persistent terminal session, preserving environment variables, working directory, and other context across multiple commands.

### Command Execution:

- Supports multi-line commands

### Directory Management:

- Must use absolute paths to avoid navigation issues.

### Program Execution:

- Supports Python, Node.js, and other executables.
- Install dependencies via pip, npm, etc.

### Background Processes:

- For long-running tasks (e.g., servers), set isBackground=true.
- Returns a terminal ID for checking status and runtime later.

### Output Management:

- Output is automatically truncated if longer than 60KB to prevent context overflow
- Use filters like 'head', 'tail', 'grep' to limit output size
- For pager commands, disable paging: use 'git --no-pager' or add '| cat'

### Best Practices:

- Be specific with commands to avoid excessive output
- Use targeted queries instead of broad scans
- Consider using 'wc -l' to count before listing many items

# Help the LLM use your tools

Robustly handle invalid input and try LLM-friendly schema changes

```
{  
    "file": "/src/app.cpp",  
    "operation": "definition",  
    "offset": 5294  
}  
  
{  
    "file": "/src/app.cpp",  
    "operation": "definition",  
    "symbol": "App::run",  
    "context": "\tif (ready) {\n\t\tApp::run();\n\t}\n"}  
}
```

# Advanced MCP

Try **elicitation** and **sampling** messages for advanced tools

**Elicitation:** Server requests more information from the user *during tool invocation*

**Sampling:** Server sends messages directly to the LLM, proxied through the client

# Take our survey to win LEGO™ prizes!



[aka.ms/cppcon/mcp](http://aka.ms/cppcon/mcp)

## Microsoft at CppCon

### Today

- 16:45 **Welcome to v1.0 of the meta::[[verse]]!** by Inbal Levi

### Tomorrow

- 14:00 **MSVC C++ Dynamic Debugging: How We Enabled Full Debuggability of Optimized Code** by Eric Brumer
- 16:45 **It's Dangerous to Go Alone: A Game Developer Tutorial** by Michael Price

### Friday

- 9:00 **Reflection-based JSON in C++ at Gigabytes per Second** by Daniel Lemire & Francisco Geiman Thiesen
- 13:30 **Duck-Tape Chronicles: Rust/C++ Interop** by Victor Ciura