

+ 25

C++ Parallel Programming Models

ERAN GILAD



20
25 |  September 13 - 19

Nice to meet you!

Day job:

Database internals @



Parallel Programming Models in C++

Community activity:



Conference



Meetup



Chit-chat

Why parallel programming?

1. Performance
2. Hide I/O
3. Responsiveness

Why is it hard?

non-deterministic execution, data races, deadlocks

Having *structure* makes things easier

Programming model

Model → structure → abstractions



User can reason about correctness
Library can provide useful features

Runtime can optimize scheduling

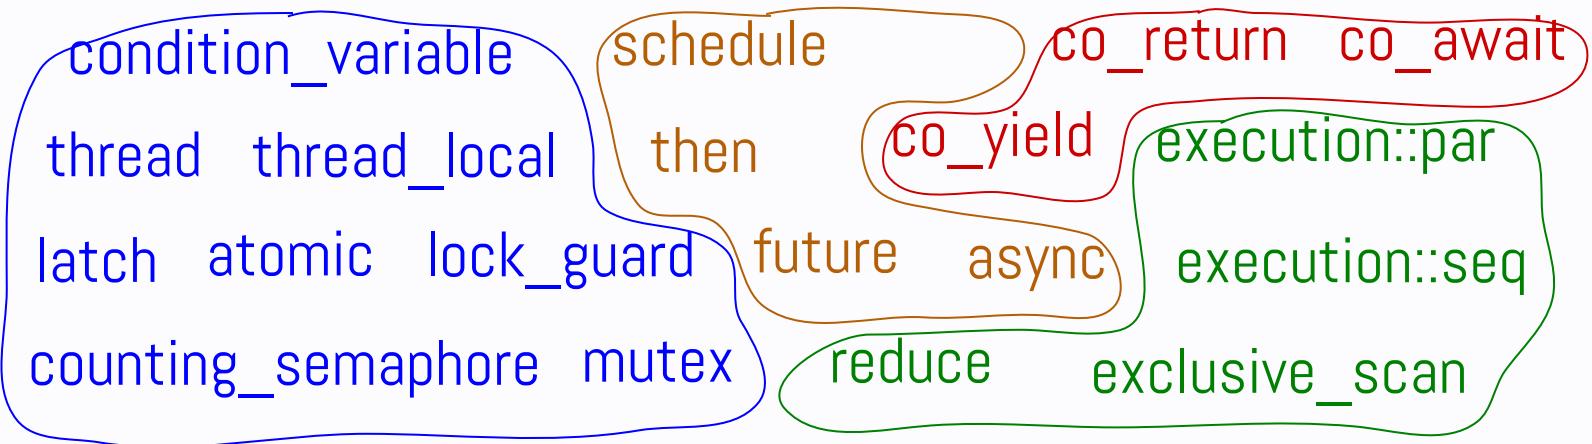


Crucial in parallel programming!

C++ parallel programming facilities

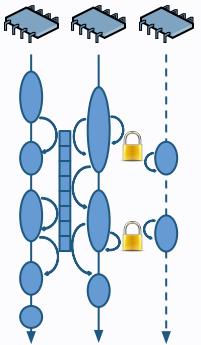
condition_variable schedule co_return co_await
thread thread_local then co_yield execution::par
latch atomic lock_guard future async execution::seq
counting_semaphore mutex reduce exclusive_scan

C++ parallel programming facilities



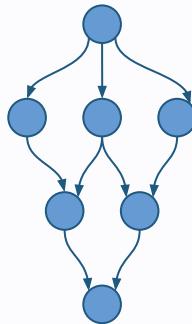
Not random facilities, not a single API:
4 parallel programming models

C++ parallel programming models



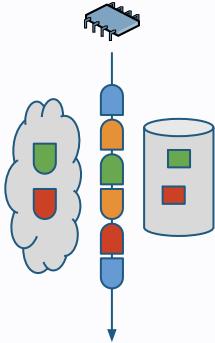
Unstructured

thread, mutex, atomics
(C++11)



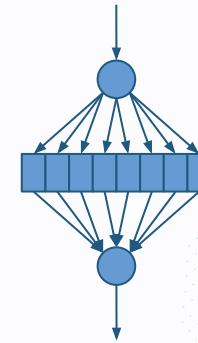
Task parallelism

async, senders
(C++11, C++26)



Cooperative tasks

coroutines
(C++20)



Data parallelism

parallel algorithms
(C++17)

Not different abstraction levels - different program structuring

Talk outline

Intro

Unstructured parallelism

Task-based parallelism

Cooperative multitasking

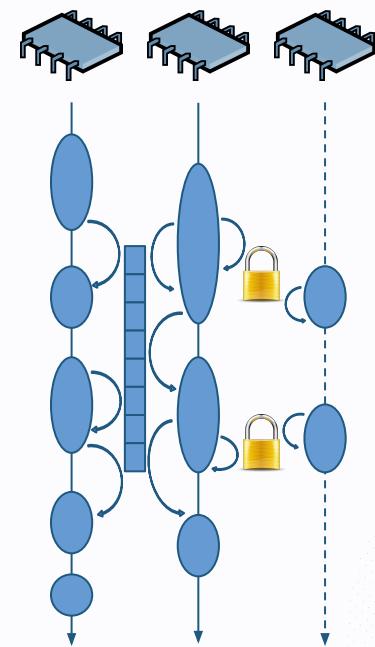
Data parallelism

Mixing models

*This talk is about the forest,
not about the trees*

Model #1: Unstructured parallelism

- Ad hoc use of parallelization facilities
 - Solve a particular need
- Relies on lower-level abstractions
- Min overhead, max hardware utilization
- Min safety, so requires max proficiency
- Central C++11 feature



Major components (partial list)

C++11/14

Convenience utils

RAI lock wrappers, call_once, `jthread`

C++20

Threading and sync primitives

`thread`, `mutex`, `condition_variable`, `semaphore`, `latch`, `barrier`

Memory model

`atomic` and friends, `thread_local`

Exclusive use cases

- Construct higher level facilities
 - E.g., thread pool, spin lock etc.
- Concurrent data structures
 - At least thread safe, usually better if lock free
- Long running services

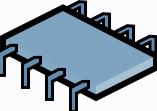
Missing part - safe shared state

- Threads communicate via shared state
 - atomics are too fine-grained
 - locking complete containers is too coarse grained
- We need *concurrent data structures!*
 - Nothing as of C++23
 - RCU and hazard pointers coming up in C++26 🤝
 - Concurrent queues in the works

Unstructured parallelism pros and cons

- Maximal control
- Usually maximal performance (when done right)
- Complicated memory model, hard to make ideal use
- Data races, deadlocks, non-determinism

Thread-level API shortcomings

Thread ≈ 

Get execution entry point,
no functional semantics

Synchronization, communication,
scheduling interleaved in code

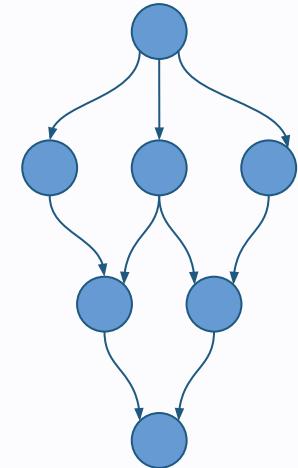
Higher abstraction level + clever runtime



Less work, less bugs, hopefully better performance

Model #2: Task-based parallelism

- Task: limited computation providing single result
 - Function, lambda, loop iteration etc.
 - *Ideally* has inputs and/or output, no side effects
- Decouples functionality from execution
 - One thread can run many tasks
 - One task can migrate among many threads (?)



Tasks are asynchronous

Execution:

- Somewhere, sometime
- Creator can proceed until the outcome is needed

Task results:

- *Future*: A handle to the task outcome
- Or passed directly to a following task

C++11 tasks - std::async et al.

```
future<int> answer = async(..., []{ return 42; });
do_caller_stuff();
print("found the answer: {}", answer.get());
```

- Spawn tasks using std::async
- Obtain results using std::future
- The runtime assigns tasks to worker threads
 - Yes, the C++ runtime can create and manage threads without user intervention

std::async execution

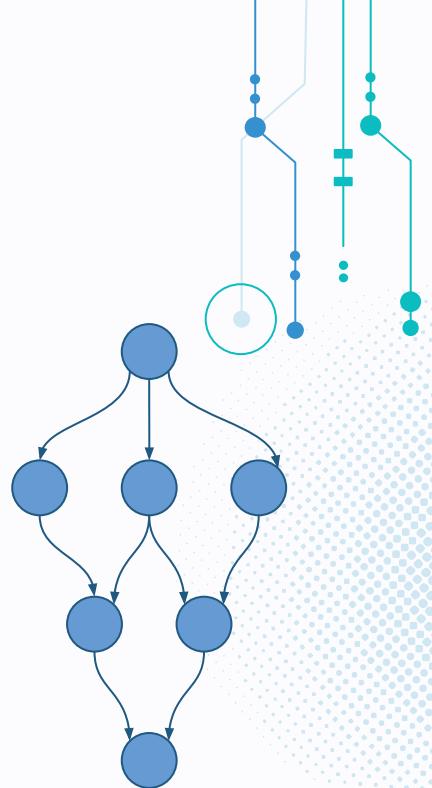
```
future<int> answer = async(when, []{ return 42; });
do_caller_stuff();
print("found the answer: {}", answer.get());
```

Task execution determined by **when**:

1. **launch::async** - on a new thread ("as if")
2. **launch::deferred** - right before the results are needed
3. **async | deferred** - implementation dependent (default)

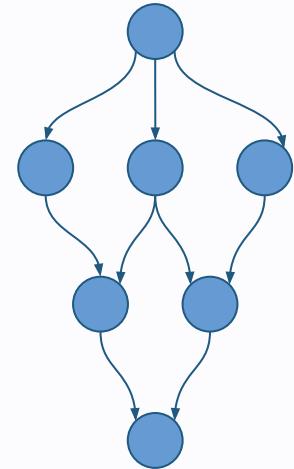
`std::async` tasks limitations

- No composition - follow-up, await many, etc.
- No execution config - e.g., thread pool
- No scheduling options - e.g., yield, stop
- No changes since C++11 



C++26 Sender/receivers tasks

- Composable by design
 - Chaining, waiting, splitting
- Configurable scheduler
 - Thread pools, coroutines, GPU, etc.
- Support stopping and error reporting
- Coming up in C++26!



Senders as tasks

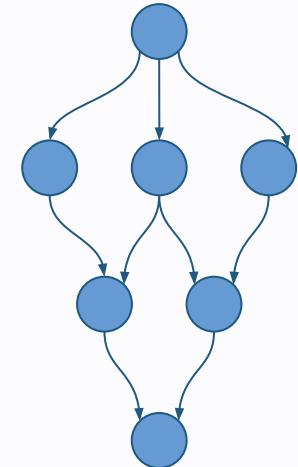
```
auto tasks =  
    just("world")           // task 1  
    | then([](string who) { // task 2  
        return format("Hello {}", who); })  
    | then([](string greet) { // task 3  
        println("{}!", greet); });  
  
sync_wait(tasks);
```



Senders define a graph of tasks, with input-output channeling

The future of async C++

- Senders/receivers are extendable
 - Schedulers, thread pools, processor type
- Allow flexible task graphs
 - With clean error handling
- Common infra for async operations
 - Sender performs the async work
 - Receiver called when done



Model #3: Cooperative Multitasking

- Long running tasks should not block
 - Assuming $\# \text{tasks} > \# \text{cores}$
 - Breaking looping sender tasks can be cumbersome
- Tasks are a user-mode concept
 - The kernel can't help

Enter coroutines!

Coroutines

C++20 coroutine: a function that can be explicitly suspended and resumed by the C++ code

- No operating system involvement
- No assembly required
- Full language (and compiler) support
- Well defined state management and error handling

Example

```
int main() {  
    coro_t c = coro_foo();  
    cout << "c suspended";  
    c.resume();  
}  
  
coro_t coro_foo() {  
    cout << "Before await";  
    co_await suspend_always{};  
    cout << "After await";  
}
```

The diagram illustrates the execution flow between the `main()` function and the `coro_foo()` coroutine. It shows the sequence of events and the points where control is transferred between them.

- Execution Flow:**
 - Control starts in `main()`.
 - It calls `coro_foo()`, which returns a `coro_t` handle.
 - It prints "c suspended".
 - It resumes the coroutine using `c.resume()`.
- Coroutine Behavior:**
 - The coroutine body begins with `cout << "Before await";`
 - It then performs a `co_await suspend_always{}` operation, which suspends the coroutine.
 - After the await, it continues with `cout << "After await";`
- Control Transfers:**
 - A red arrow points from `main()` to the start of `coro_foo()`.
 - A red arrow points from the end of `coro_foo()` back to `main()`.
 - A red arrow points from `main()` back to the `co_await` point in `coro_foo()`.

coro_t: define coroutine behavior and resume handle

co_await *expr*

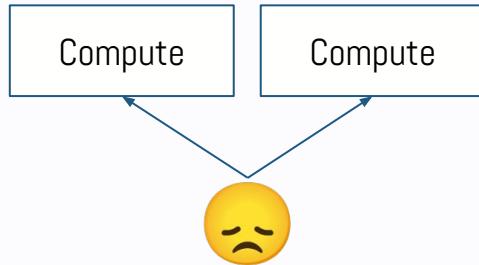
- *expr* should evaluate to an *awaitable object*
- The awaitable is a scheduling point:
 - `await_ready()` - should the coroutine be suspended
 - `await_suspend()` - get the suspended coroutine handle
 - `await_resume()` - after coroutine resumed
- `co_await foo():`
 - execute `foo`, evaluate its return value (*awaitable*)
 - suspend the current coroutine if needed

Coroutine scheduling

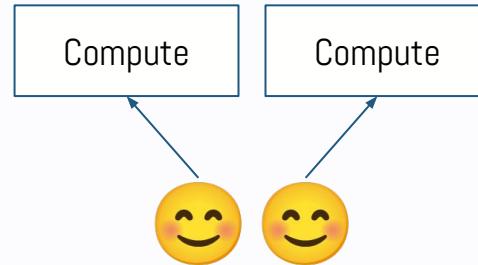
- `co_await` returns control to caller or another coroutine
 - It *does not* create any parallelism!
- The program must:
 - Track suspended coroutines
 - Resume ready coroutines
- Requirements:
 - Some source of parallelism (async I/O, worker threads, etc.)
 - Scheduler

Concurrency and parallelism

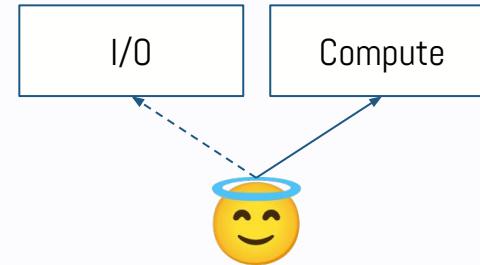
Task concurrency,
no parallelism



Task concurrency,
multi-CPU parallelism

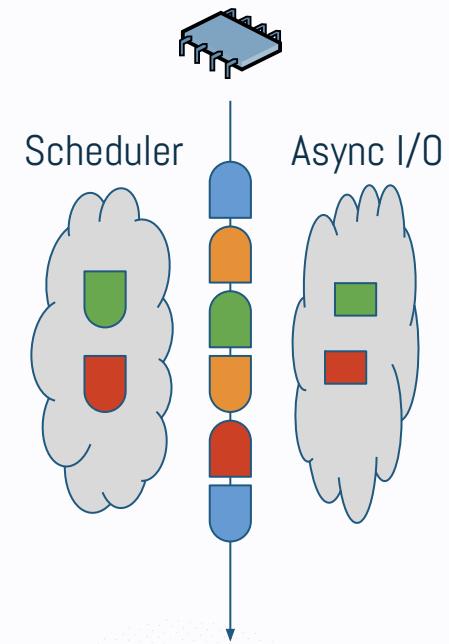


Task concurrency,
CPU/I/O parallelism



Coroutine-based multitasking example

1. The scheduler starts a coroutine task
2. The task invokes **async I/O**
3. The task **suspends itself**
4. The scheduler switches to another task
5. The async I/O completes
6. The awaiting task marked as ready
7. The scheduler **resumes it** when possible



Parallel coroutines execution

- Commonly one thread per core
- Suspended tasks can move among cores
 - Can you use `thread_local` in this scenario?
- Sync primitives must be coroutine-aware
 - `co_await coro_lock()` instead of `std::mutex.lock()`
 - A coroutine can't switch if its thread is blocked!
- Expected scheduler complexity - races, load balancing, task affinity, etc.

Threads vs. tasks

Threads use cases	Tasks use cases
Parallelism abstraction (CPU)	Concurrency abstraction (function)
Background services/workers	Many short and independent ops
Long complex parallel operations	"local"/ad-hoc (async I/O)
Equal amount of per-thread work	Dynamic load balancing

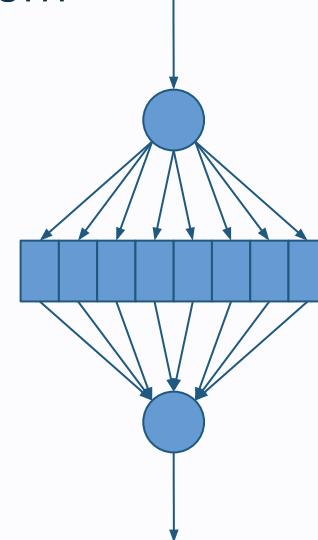
Fibers / user threads / green threads

- Stackful coroutines
 - A full stack must be allocated for each fiber
- Any function running on a fiber can be suspended
 - Not just the coroutine body
- Decoupled executor (fiber) and work (function)
 - Allows fiber pools etc.
- Requires an "ecosystem" - scheduler and sync objects
- Not supported in C++ (yet?)

Model #4: Data parallelism

A programming model in which parallelism stems from the individual computations associated with each element in a collection.

- Requires a lot of non-dependent data
- Procedure:
 - Divide element processing among processors
 - Short and simple operation on each element



C++ Parallel algorithms

```
res = algorithm(  
    exec_policy,  
    container);
```

Sequential execution context

Declarative API

Accelerate data processing

- High level abstraction
 - Few customization points for user
 - No control over parallelism, scheduling, work distribution etc.
 - Library + runtime can be very efficient
 - When used properly: no data races, deadlocks etc.

Using parallel algorithms

```
vector<int> v = {1, 2, 3, 5, 11, 20};  
int res = reduce(execution::par, v.begin(), v.end());  
assert(res == 42);
```

- Most standard algorithms have a parallel overload
 - First parameter: *ExecutionPolicy*
- Complexity requirements more lax
- Implementation isn't specified – can also use GPUs

sequenced_policy

- Forces execution to take place on the calling thread
- Minor differences from no-policy call
- `execution::seq` is an *instance of* `sequence_policy`
 - Algorithms can be overloaded by policy
 - Policy is a *compile time* decision!

```
reduce(execution::seq, ..., ...);
```

parallel_policy

- Execution on caller or another thread (runtime pool)
- Per thread, semantics are similar to sequenced_policy - unspecified order, no interleaving
- Data races are now possible if multiple operations write to unprotected data

```
reduce(execution::par, ..., ...);
```

parallel_unsequenced_policy

- A single algorithm thread can now process multiple elements at the same time
 - Operations must not use any locks
- More user restrictions => more library options
 - Vectorization can now be used
 - finer grained scheduling

```
reduce(execution::par_unseq, ... , ...);
```

unsequenced_policy

- Operations can be interleaved on a single thread
 - Not a multithreaded context
 - But vectorization can still be used!
- C++20 addition

```
reduce(execution::unseq, ..., ...);
```

Non-standard policies

- Vendor-specific
- Can allow the use of accelerators
 - GPU
 - FPGA
 - ASIC
 - ...

Parallel algorithms != parallel containers

- C++ separates algorithms from containers
 - Thread-safe containers by default? No - zero-overhead!
- Parallel algorithms can modify data but not structure
 - Unlike sequential algorithms, which can modify both
- Accessing a container processed by a parallel algorithm from another thread? Possible race!

Mixing models - unstructured context

Most cores are used:

- E.g., server
- Use coroutines for async I/O
- No resources for parallel algorithms or async compute

Most cores are unused:

- E.g., background service
- Tasks and parallel algorithms can be used
 - Impact? depends
- Careful with shared state!

Reminder: runtime schedulers aren't aware of user threads

Mixing models - `async/sender` tasks context

Most cores are used:

- Namely, many tasks created
- No point in creating threads or using parallel algorithms
- Awkward context for coroutine execution
 - Even if tasks are coroutines

Most cores are unused:

- E.g., ad-hoc work, UI worker
- Can use parallel algorithms from tasks
- Creating threads less suitable - spawn tasks instead
- Using sync mechanism (mutex etc.) doesn't fit the model

Mixing models - coroutine tasks context

Most cores are used:

- Namely, the scheduler uses many cores
- No point in using other models from within coroutines

Most cores are unused:

- E.g., limited I/O needs
- Parallel algorithms from coroutine will block it
- So will async/sender tasks
- Tasks too short for thread creation

Mixing models - parallel algorithms

- Within a parallel algorithm (e.g., user lambda):
 - No point in spawning async tasks or creating threads
 - Inappropriate context for coroutines
 - Using sync mechanisms to access external state possible but might kill concurrency

Summary



Unstructured

Low-level building blocks



Cooperative multitasking

great for async I/O,
missing scheduler



Task parallelism

functional decomposition,
much better with senders



Data parallelism

declarative, no executors
control (yet!)

Mixing models hardly works. Parallel programming is hard.

