

+ 25

Cutting C++ Exception Time by 93.4%

KHALIL ESTELL



20
25 | A graphic of three white mountain peaks with a yellow square at the top of the tallest one.
September 13 - 19

Cutting C++ Exception time by 90%

Part 2

@kammce





Chapters

Chapter 0. The beginning

Chapter 1. Benchmarking

Chapter 2. Optimizing GCC Exceptions

Chapter 3. Estell Exception Runtime

Chapter 4. Nearpoint Search

Chapter 5. Future of Exceptions

Chapter 0.



CppCon 2023 SG 14 ISO Meeting
October 4th



Conclusion

From the data: $\beta = 190$, $\sigma = 129$, Thus: $\beta > \sigma$ ✓

With 150 functions: 783 x `tl::unexpected()` calls, which corresponds to if/else checks.

What does this tell us?

- If you only have **4kB**, **8kB**, **16kB**, and even **32kB** flash, then exceptions may not be the right choice for your application because you cannot afford the +8kB flash cost for exceptions.
- With flash sizes **64kB** and above, exceptions seem to be the best option for code size, *IF you can handle the performance cost of throws*.
- Improving unwinding performance, removing the "need" for dynamic memory allocation, and removing TLS requirements, would make exceptions an extremely attractive options for embedded systems.

Improving **exceptions** would
make them an **extremely**
attractive options for
embedded systems...



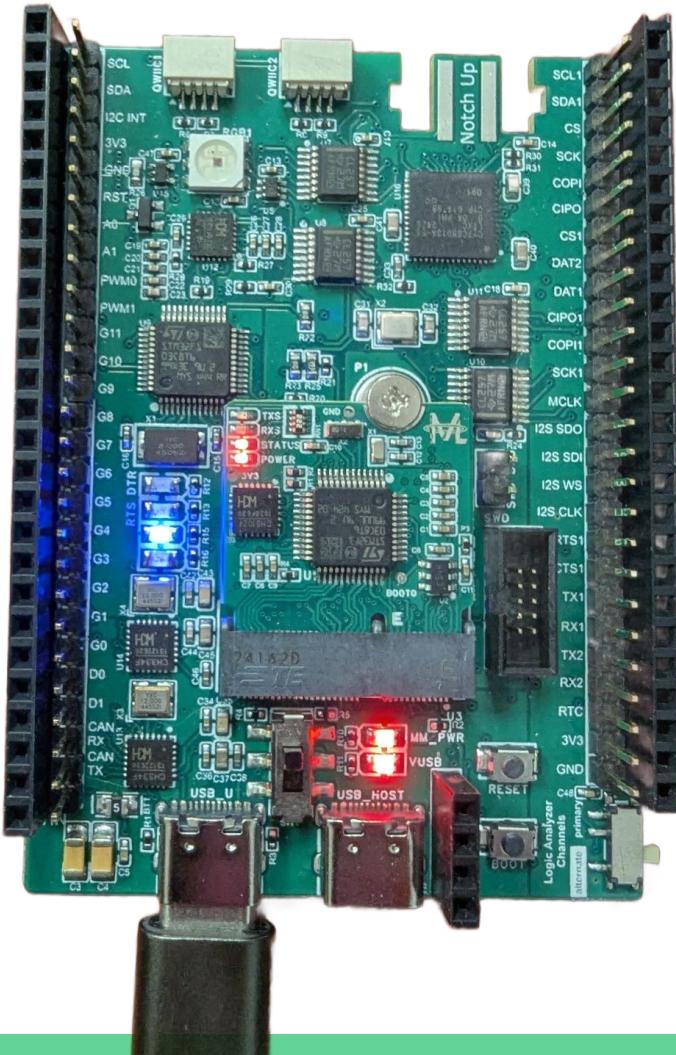
Chapter 1.



Measuring Performance
Benchmark Design

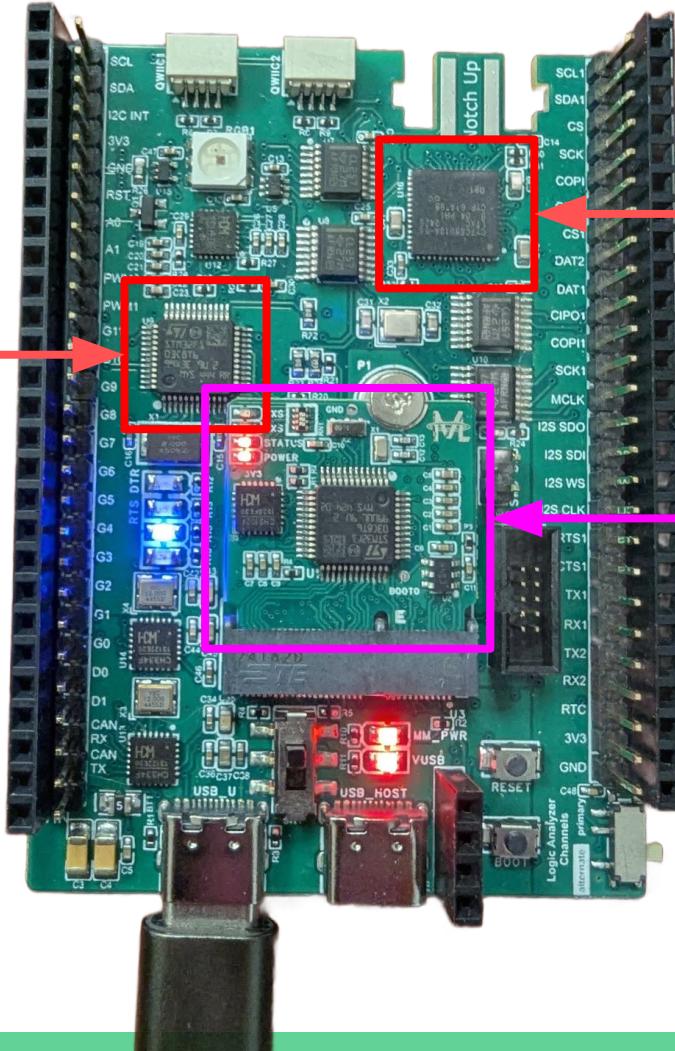
Test Setup & Hardware

Testing Hardware: HALbORD P3



Testing Hardware: HALbORD P3

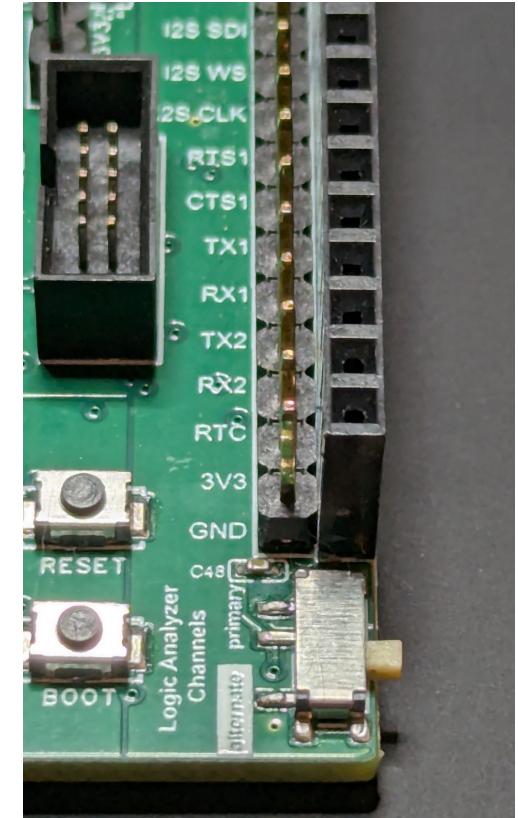
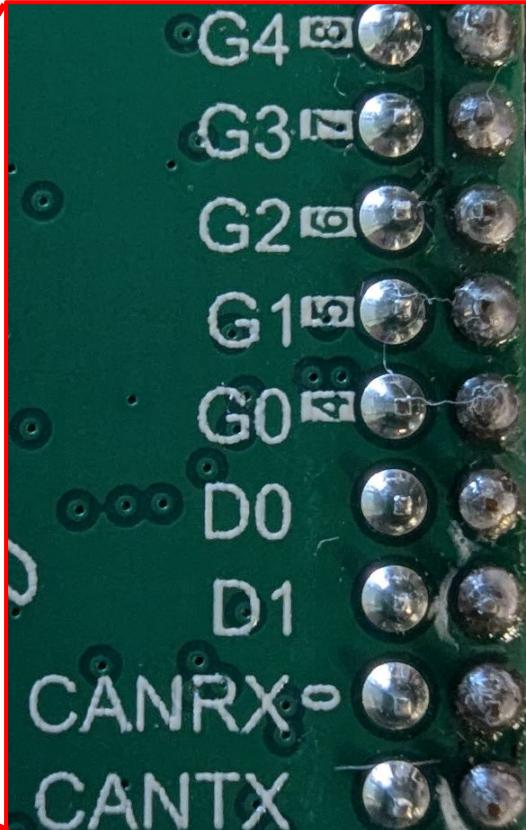
SWD/JTAG
Debugger
(gdb)



2x16-Channel
Logic Analyzer

MCU

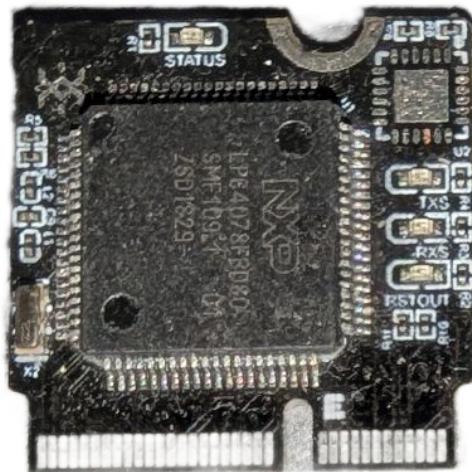
Logic Analyzer Channels & Switch





MicroMod

512kB of ROM



LPC4078

64kB of ROM

64MHz
Cortex-M3



STM32F103C8

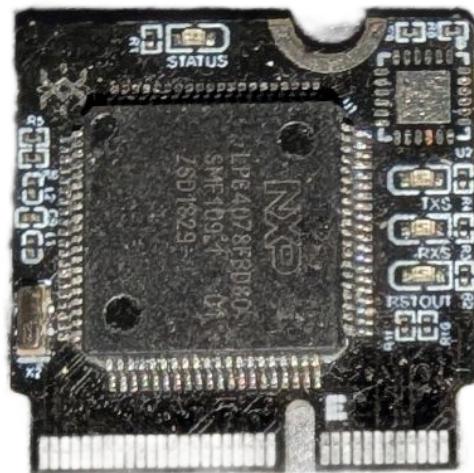
16MB of ROM



RP2040



512kB of ROM



LPC4078

MicroMod

64kB of ROM

64MHz

Cortex-M3



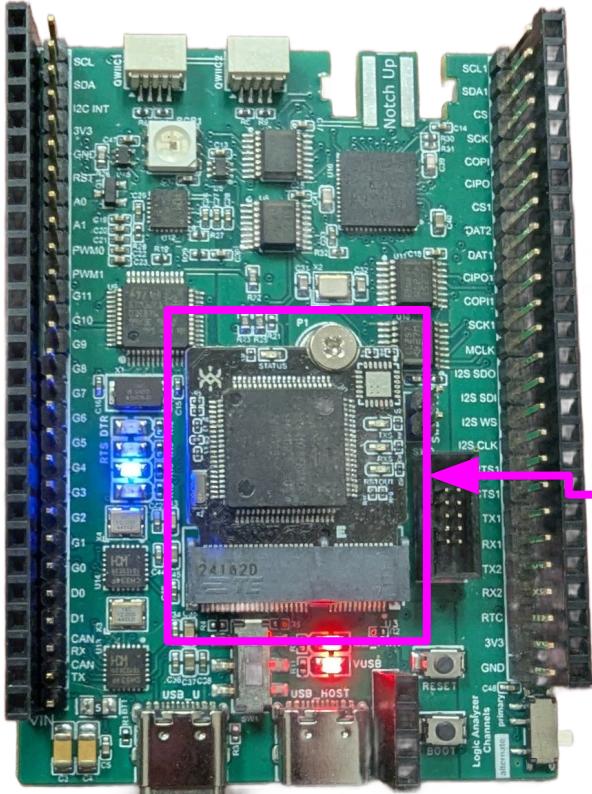
STM32F103C8

16MB of ROM



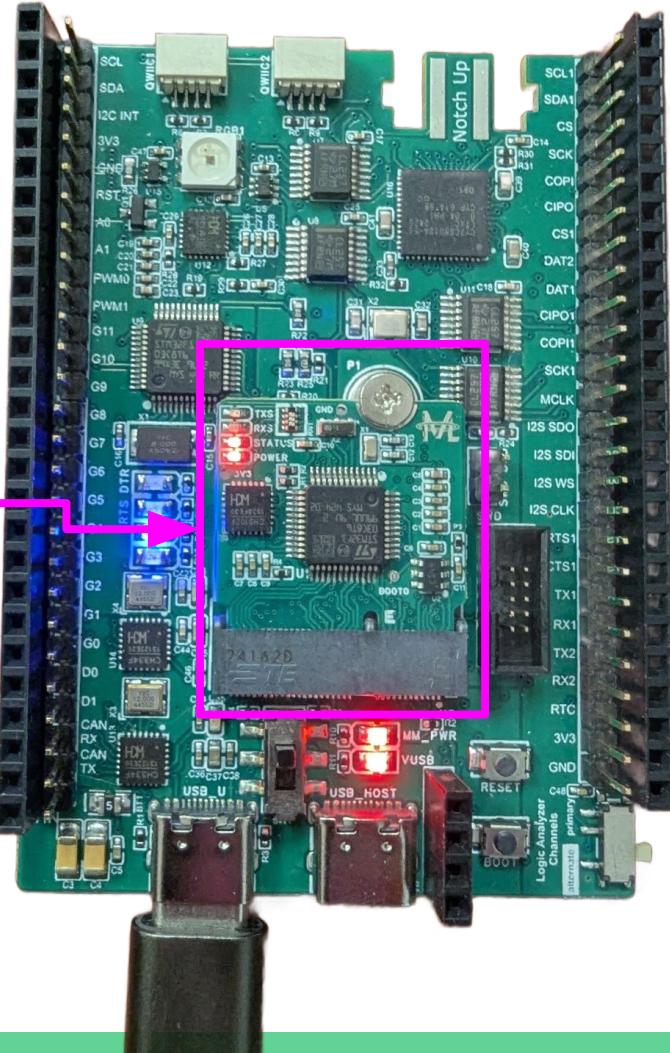
RP2040

Testing Hardware: HALbORD P3

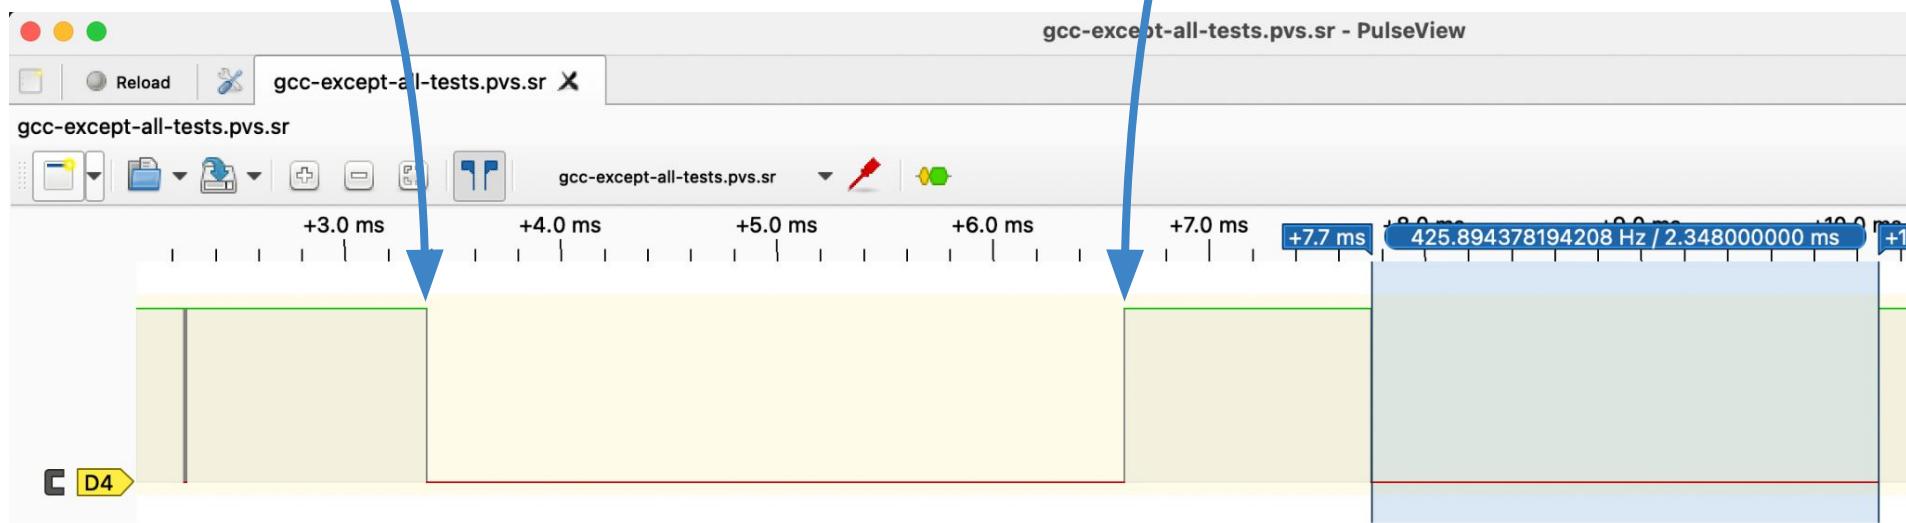


stm32f103c8
MicroMod

Ip4078
MicroMod



Using the Logic Analyzer for Measurement



Microcontroller Benchmark Benefits: No OS



~~fopen()~~



~~networking~~



~~std::print()~~



~~std::chrono::steady_clock::now();~~



~~std::thread~~



Benchmark Design

My Toolchain

14.2



1.8.10

C⁻¹²

Benchmark Specification



expected



exceptions

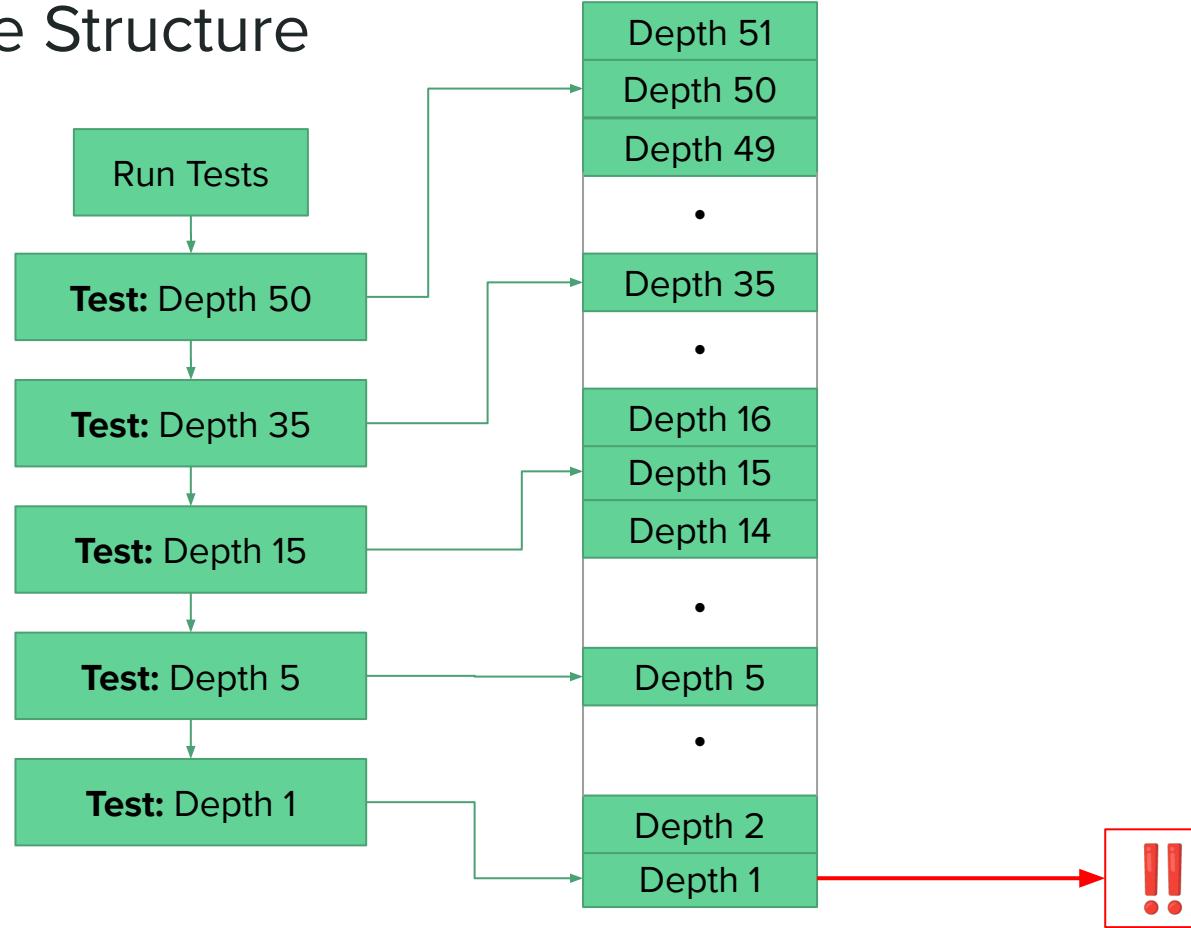
Goal: Measure propagation time up the stack from ***error detection*** to ***handler***.

- Should compare idiomatic `std::expected<T, E>` & `exceptions`.
- Programs and functions must be **isomorphic** to each other
- Test on Release (-O3) mode for highest performance

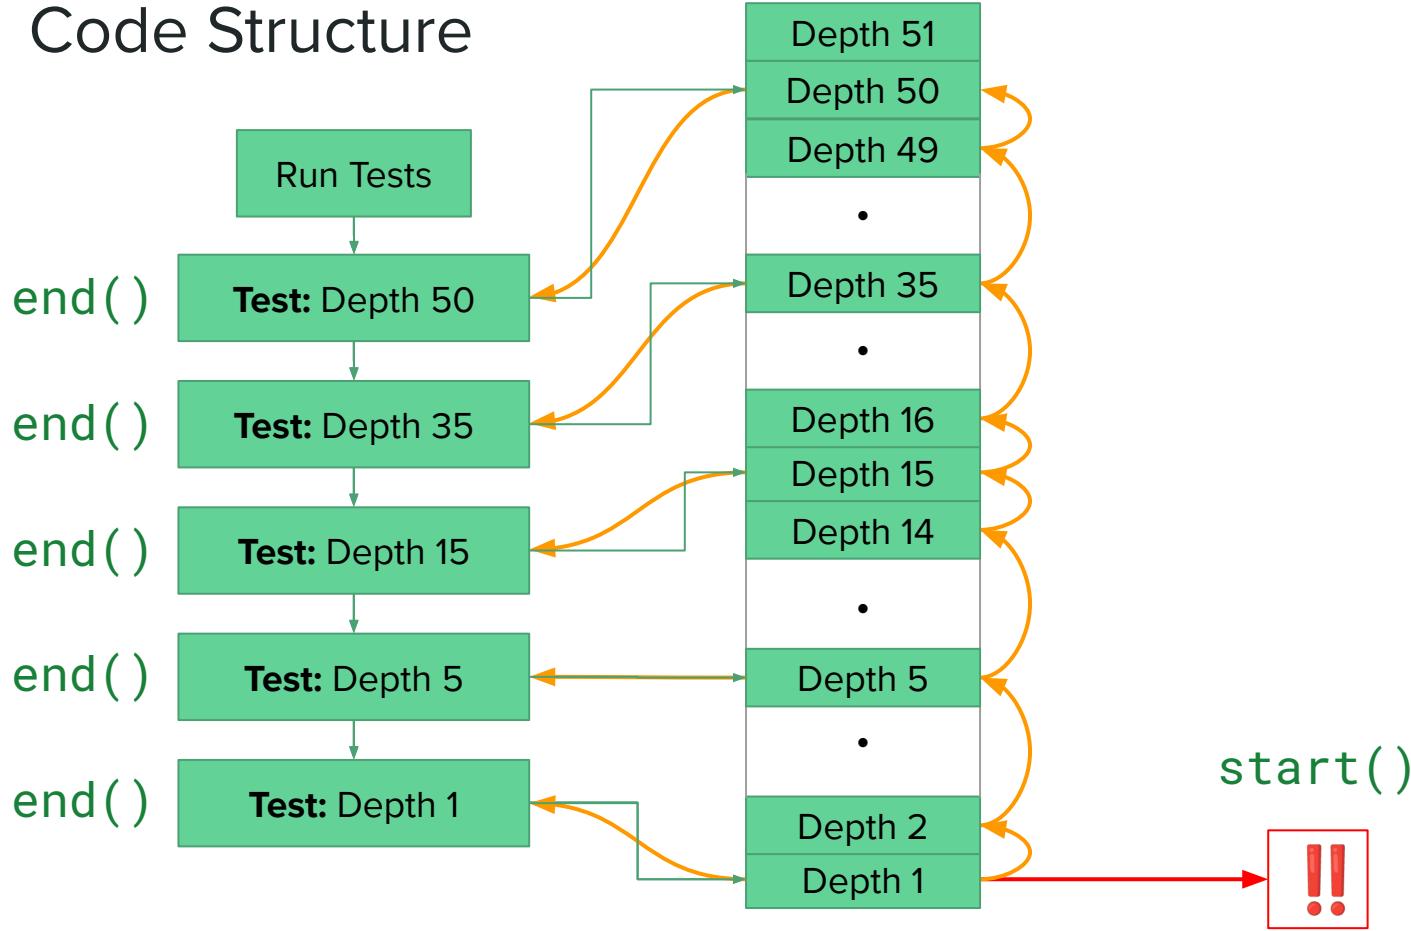
What we are **NOT** Testing

- Allocation Speed

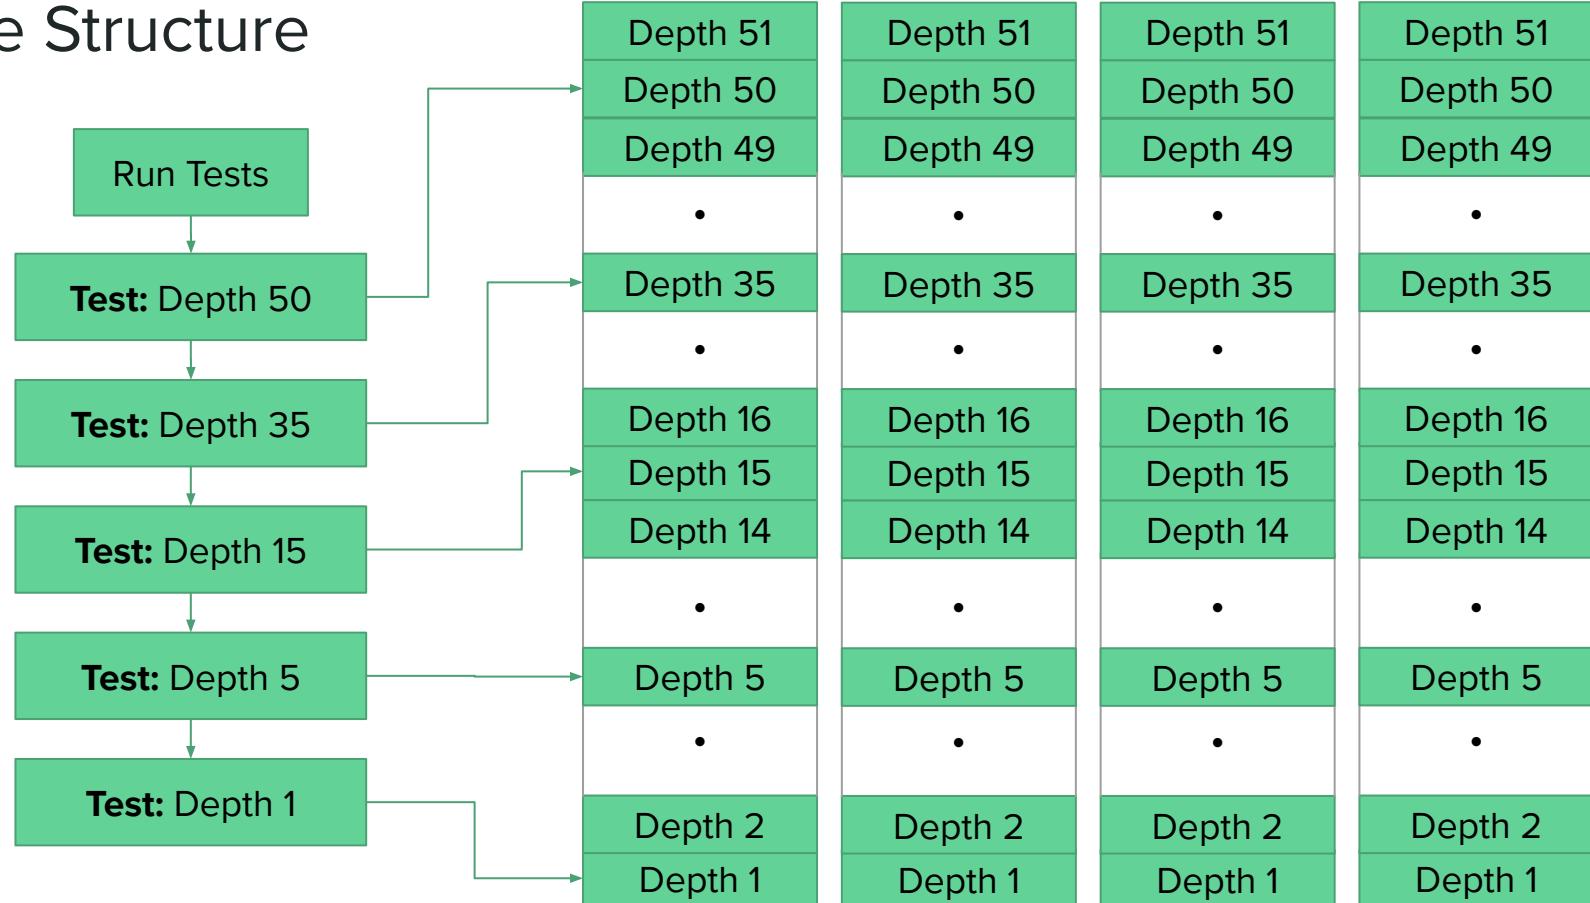
Code Structure



Code Structure



Code Structure



```
// Global side effect
std::int32_t volatile side_effect = 0;

class simple_object
{
public:
    explicit simple_object(std::int32_t value);
    void do_work();

private:
    std::int32_t m_value;
};

class destructor_object
{
public:
    explicit destructor_object(std::int32_t value);
    void do_work();

    ~destructor_object()
    {
        side_effect = side_effect + 1;
    }

private:
    std::int32_t m_value;
};
```

```
// Global side effect
std::int32_t volatile side_effect = 0;

class simple_object
{
public:
    explicit simple_object(std::int32_t value);
    void do_work();

private:
    std::int32_t m_value;
};

class destructor_object
{
public:
    explicit destructor_object(std::int32_t value);
    void do_work();

    ~destructor_object()
    {
        side_effect = side_effect + 1;
    }

private:
    std::int32_t m_value;
};
```

Benchmark depth function

```
[[gnu::noinline]]
int depth_50_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, 10> preserve_frame = { 8 };

    int const result = depth_49_percent_050();

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return -1;
    }
    return result + side_effect;
}
```

Benchmark depth function

```
[[gnu::noinline]]
int depth_50_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, 10> preserve_frame = { 8 };

    int const result = depth_49_percent_050();

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return -1;
    }
    return result + side_effect;
}
```

Benchmark depth function

```
[[gnu::noinline]]
int depth_50_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, 10> preserve_frame = { 8 };

    int const result = depth_49_percent_050();

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return -1;
    }
    return result + side_effect;
}
```

Preventing Tail Call Optimizations

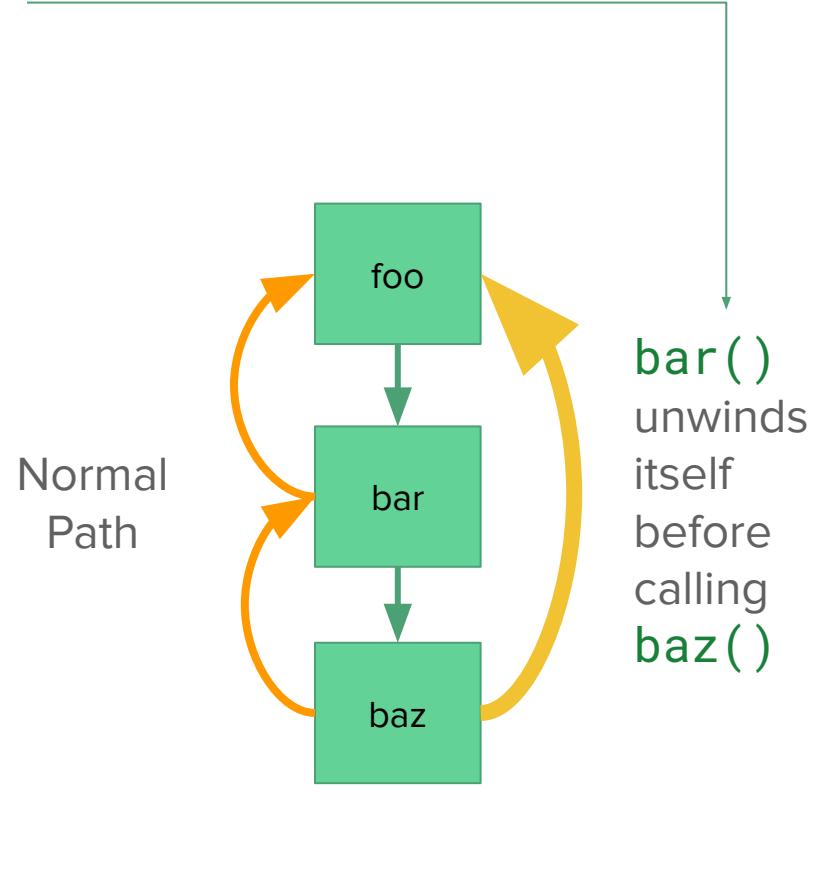
```
[[gnu::noinline]]
int depth_50_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, 10> preserve_frame = { 8 };

    int const result = depth_49_percent_050();

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return -1;
    }

    return result + side_effect;
}
```



Benchmark depth function

```
[[gnu::noinline]]
int depth_50_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, [10>]preserve_frame = { 8 };

    int const result = depth_49_percent_050();

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return -1;
    }
    return result + side_effect;
}
```

```
[[gnu::noinline]]
int depth_49_percent_050()
{
    destructor_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, [4>]preserve_frame = { 8 };

    int const result = depth_48_percent_050();

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return -1;
    }
    return result + side_effect;
}
```

Error types (4B, 16B, 65B)

```
struct test_error_4
{
    std::array<std::uint8_t, 4> data;
};

struct test_error_16
{
    std::array<std::uint8_t, 16> data;
};

struct test_error_65
{
    std::array<std::uint8_t, 65> data;
};
```

- Trivially relocatable (can be **memcpy**)
- Trivially constructible
- Trivially destructible
- Array of 4B has the **same performance** as **int** or **std::errc**

```

[[gnu::noinline]]
std::expected<int, test_error_04>
depth_50_error_04_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, 10> preserve_frame = { 8 };

    auto const result = depth_49_error_04_percent_050();

    if (!result) {
        return std::unexpected(result.error());
    }

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return std::unexpected(result.error());
    }

    return result.value() + side_effect;
}

```

```

[[gnu::noinline]]
int depth_50_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, 10> preserve_frame = { 8 };

    int const result = depth_49_percent_050();

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return -1;
    }

    return result + side_effect;
}

```

std::expected

```
[[gnu::noinline]]
std::expected<int, test_error_04>
depth_50_error_04_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, [10]> preserve_frame = { 8 };

    auto const result = depth_49_error_04_percent_050();

    if (!result) {
        return std::unexpected(result.error());
    }

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return std::unexpected(result.error());
    }

    return result.value() + side_effect;
}
```

```
[[gnu::noinline]]
int depth_50_percent_050()
{
    simple_object obj(side_effect >> 8);
    obj.do_work();

    std::array<std::int8_t volatile, [10]> preserve_frame = { 8 };

    int const result = depth_49_percent_050();

    // Use the variable after the call
    if (preserve_frame[0] < 0) {
        // Never executed but prevents tail-call optimization
        return -1;
    }

    return result + side_effect;
}
```

std::expected

Emitting the Error

```
[[gnu::noinline]]  
  
std::expected<int, test_error_04> depth_01_error_04_percent_000()  
{  
  
    simple_object obj(side_effect >> 8);  
  
    obj.do_work();  
  
    std::array<std::int8_t volatile, 14> preserve_frame = { 8 };  
  
    // This is where the error originates  
  
    if (side_effect > 0) {  
  
        [start();]  
  
        return std::unexpected(test_error_04{  
            .data = { 0xDE, 0xAD, 0xBE, 0xEF } });  
    }  
  
    // Use the variable after the call  
  
    if (preserve_frame[0] < 0) {  
  
        // Never executed but prevents tail-call optimization  
  
        return -1;  
    }  
  
    return side_effect + 1;  
}
```

```
[[gnu::noinline]]  
  
int depth_01_percent_000() {  
  
    simple_object obj(side_effect >> 8);  
  
    obj.do_work();  
  
    std::array<std::int8_t volatile, 14> preserve_frame = { 8 };  
  
    // Use the variable after the call  
  
    if (preserve_frame[0] < 0) {  
  
        // Never executed but prevents tail-call optimization  
  
        return -1;  
    }  
  
    if (error_size_select == 4) {  
  
        [start();]  
  
        throw test_error_04{.data = {0xDE, 0xAD, 0xBE, 0xEF}};  
    }  
  
    if (error_size_select == 16) {  
  
        [start();]  
  
        throw test_error_16{.data = {0xDE, 0xAD, 0xBE, 0xEF}};  
    }  
  
    if (error_size_select == 65) {  
  
        [start();]  
  
        throw test_error_65{.data = {0xDE, 0xAD, 0xBE, 0xEF}};  
    }  
  
    return side_effect;  
}
```

Reaching the Handler

```
[[gnu::noinline]]  
  
void run_test_depth_50_error_04_cleanup_100 ()  
{  
    // Ensure we will return error  
    side_effect = 1;  
  
    auto const result =  
        depth_50_error_04_percent_100 ();  
    if (!result) {  
        end ();  
    }  
  
    pause ();  
}
```

```
// Test runner  
[[gnu::noinline]]  
  
void  
run_test_depth_50_error_04_cleanup_000 ()  
{  
    // Ensure we will throw  
    side_effect = 1;  
    error_size_select = 4;  
  
    try {  
        depth_50_percent_000 ();  
    } catch (test_error_04 const& e) {  
        end ();  
    }  
  
    pause ();  
}
```

How tests are executed

```
// Test runner  
void run_test()  
{  
    run_test_error_04_cleanup_000();  
    run_test_error_04_cleanup_025();  
    run_test_error_04_cleanup_050();  
    run_test_error_04_cleanup_100();  
    run_test_error_16_cleanup_000();  
    run_test_error_16_cleanup_025();  
    run_test_error_16_cleanup_050();  
    run_test_error_16_cleanup_100();  
    run_test_error_65_cleanup_000();  
    run_test_error_65_cleanup_025();  
    run_test_error_65_cleanup_050();  
    run_test_error_65_cleanup_100();  
}
```

```
void run_test_error_04_cleanup_000()  
{  
    run_test_depth_50_error_04_cleanup_000();  
    run_test_depth_35_error_04_cleanup_000();  
    run_test_depth_15_error_04_cleanup_000();  
    run_test_depth_05_error_04_cleanup_000();  
    run_test_depth_01_error_04_cleanup_000();  
}
```

```
void run_test_error_65_cleanup_100()  
{  
    run_test_depth_50_error_65_cleanup_100();  
    run_test_depth_35_error_65_cleanup_100();  
    run_test_depth_15_error_65_cleanup_100();  
    run_test_depth_05_error_65_cleanup_100();  
    run_test_depth_01_error_65_cleanup_100();  
}
```

Testing Script Steps

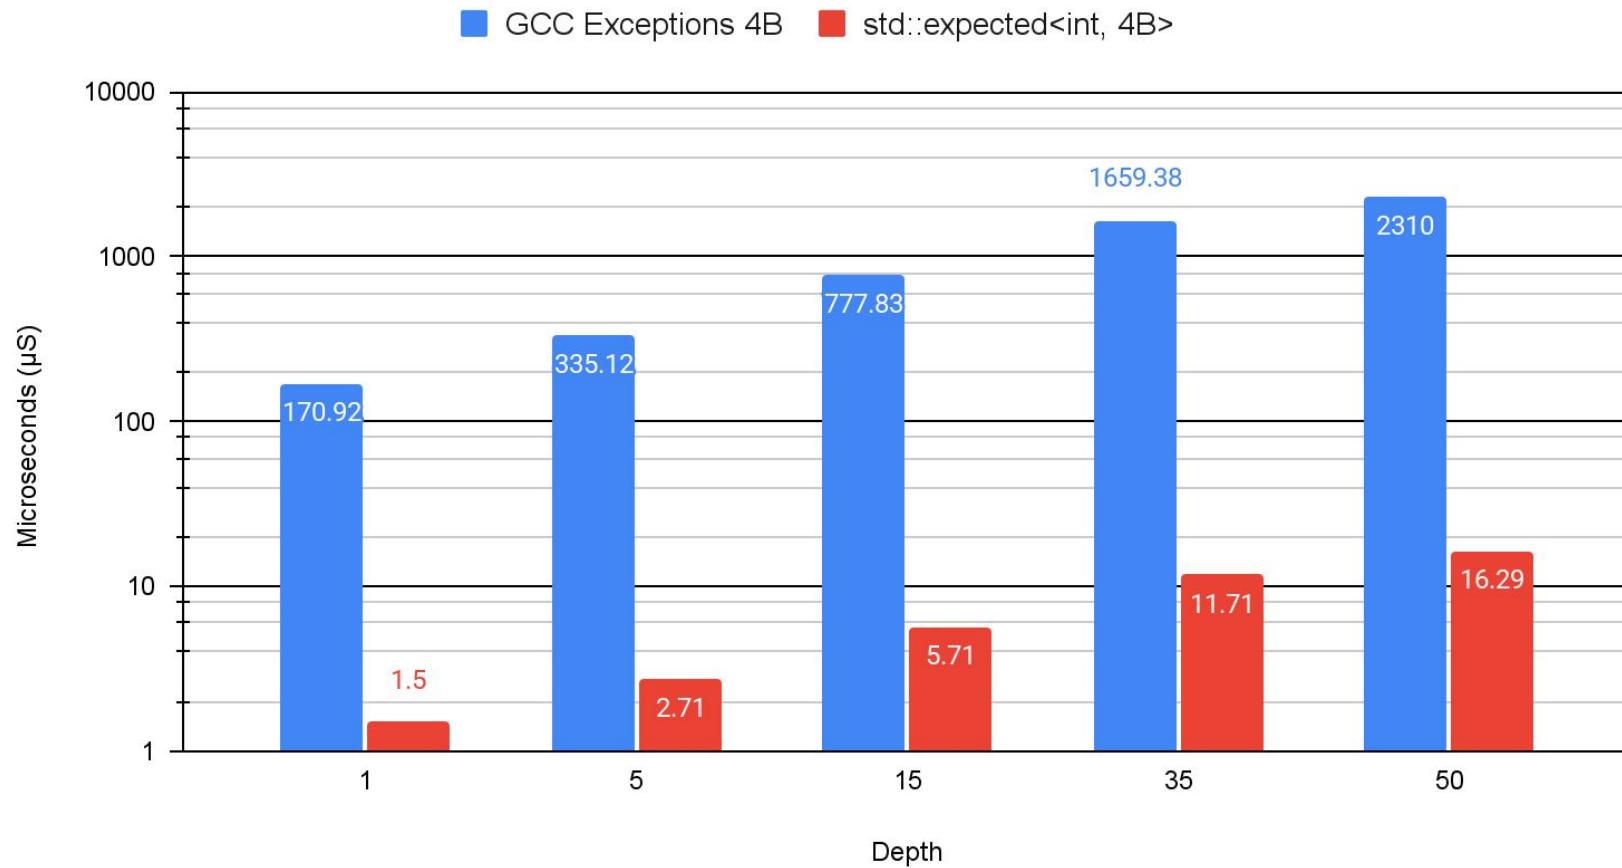




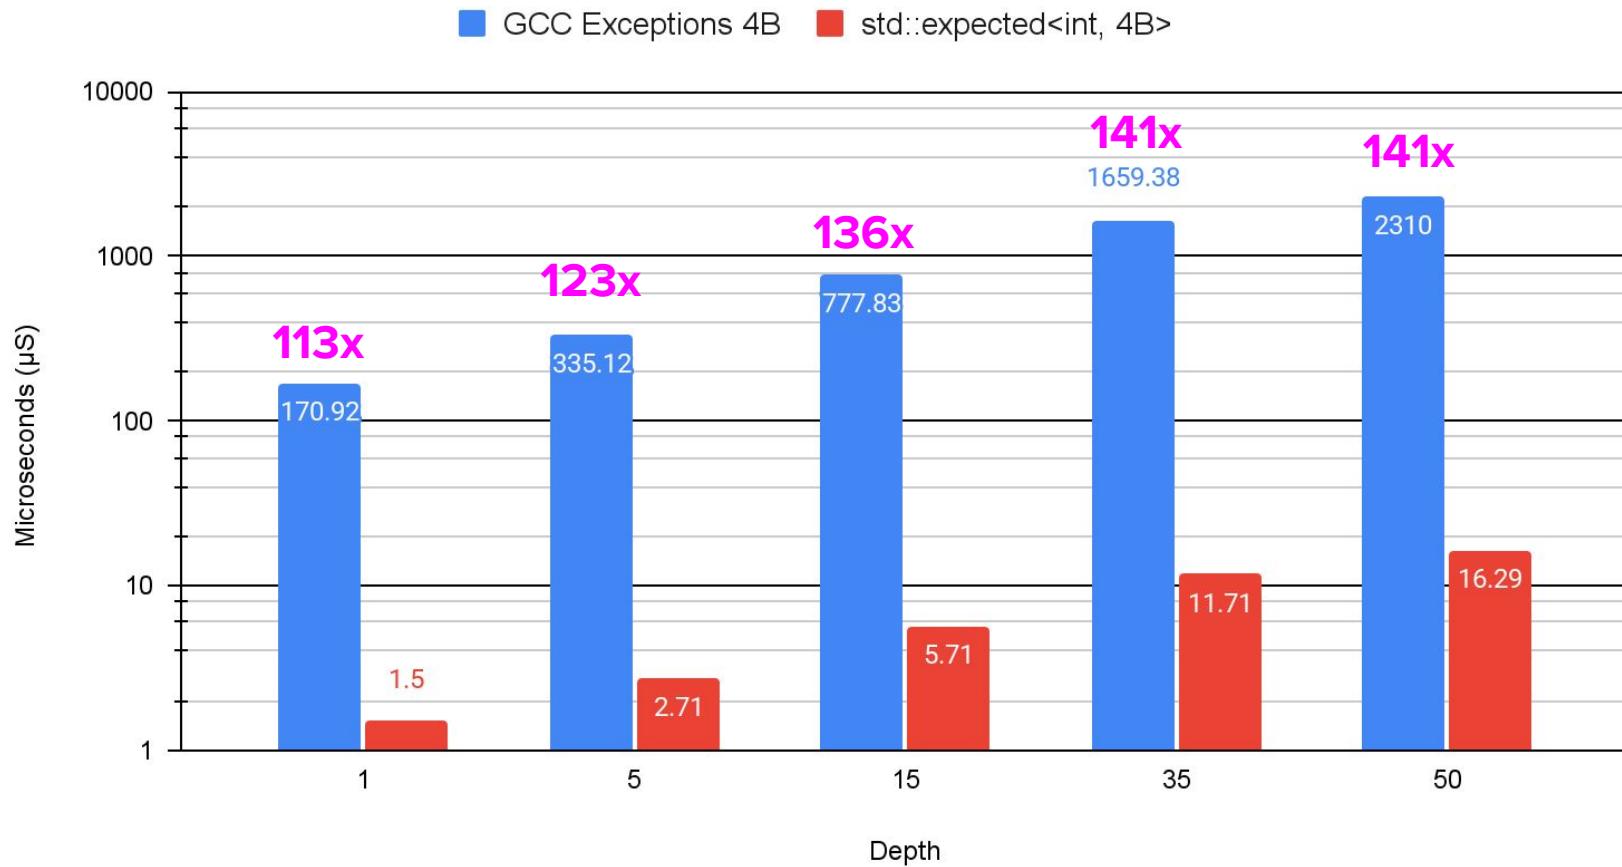
Results!

*LOG SCALE

Error Propagation Time - 0% Cleanup

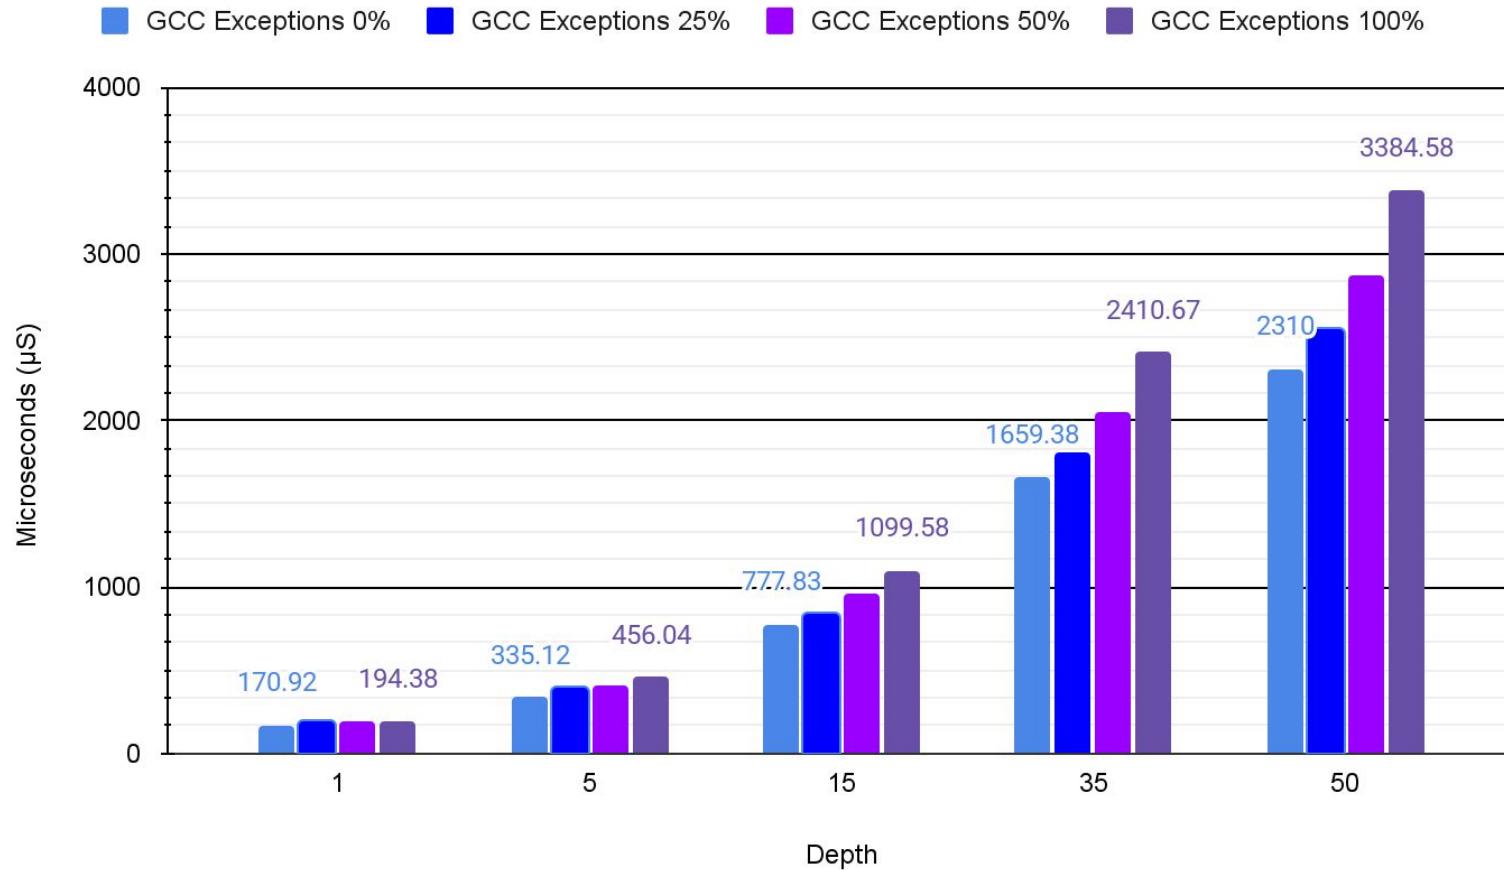


Error Propagation Time - 0% Cleanup

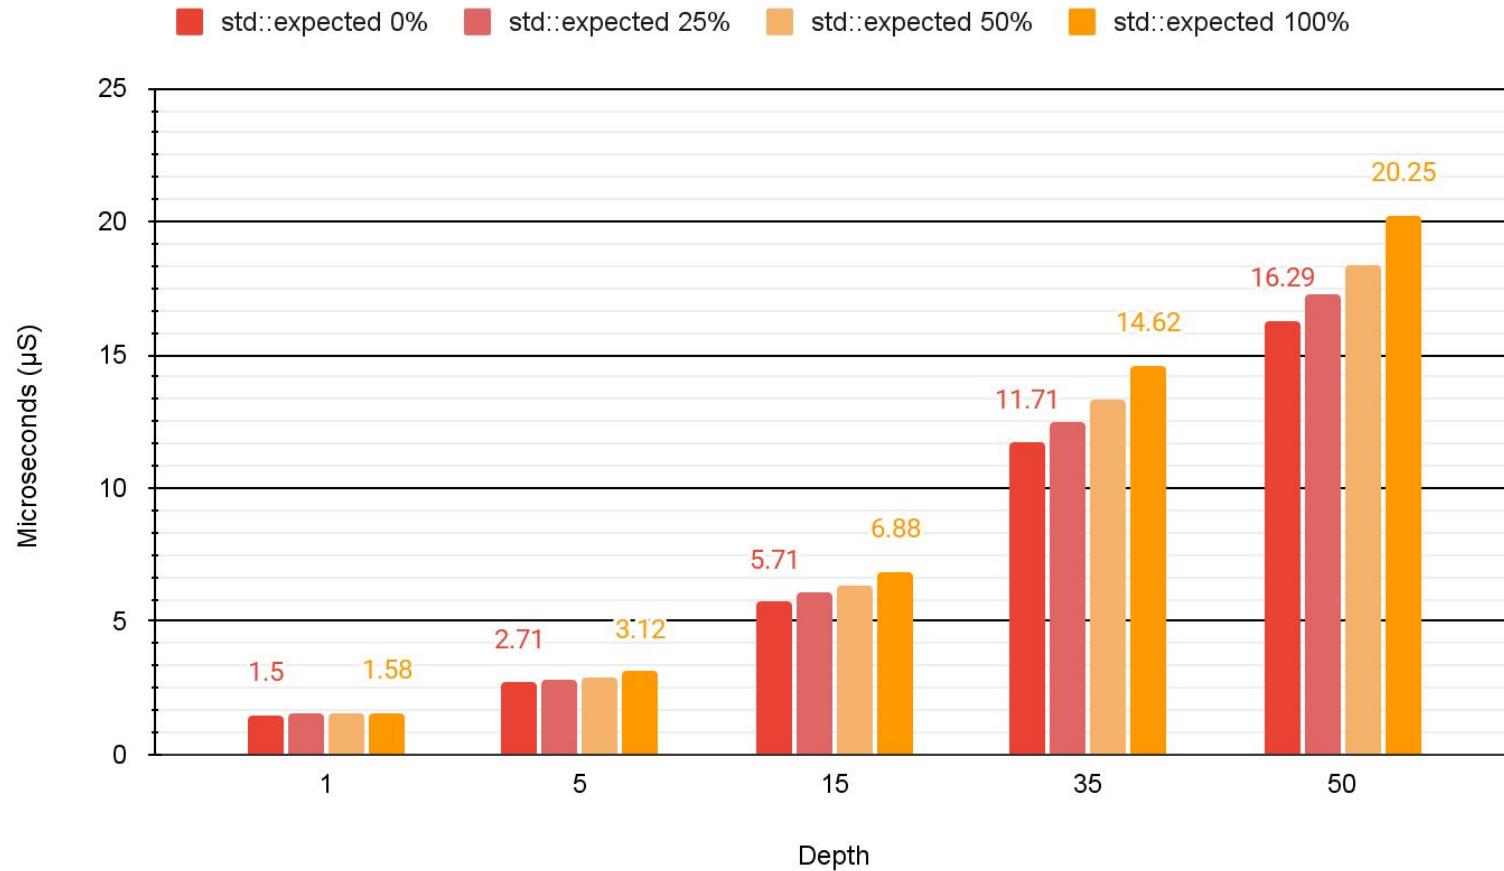


How does cleanup % effect
things?

Exceptions propagation time, 4B error size, Cleanup -Variable-

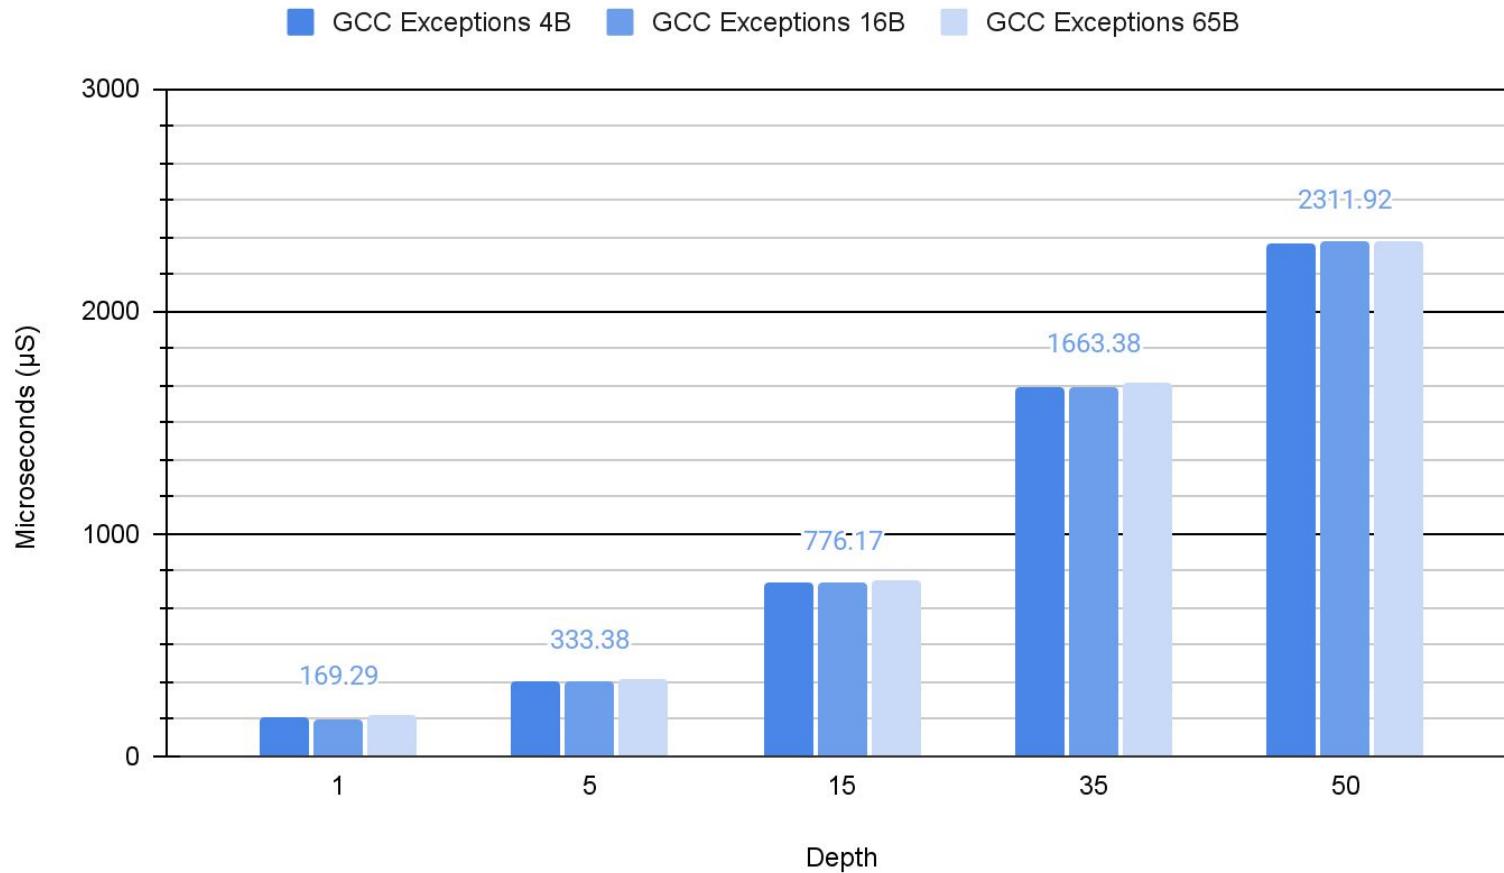


std::expected propagation time, 4B error size, Cleanup -Variable-

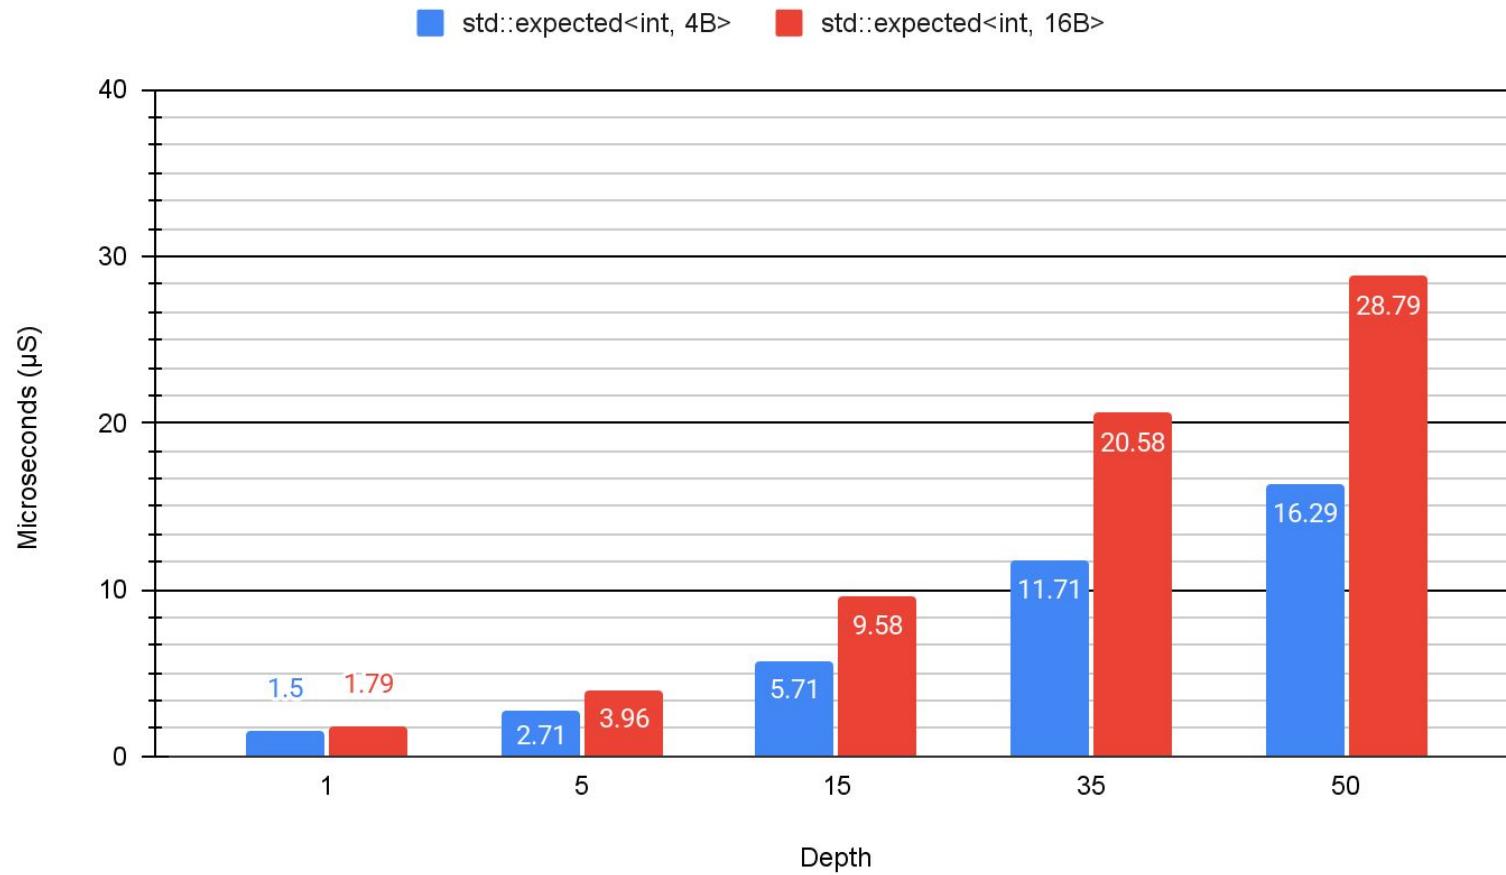


Error object size?

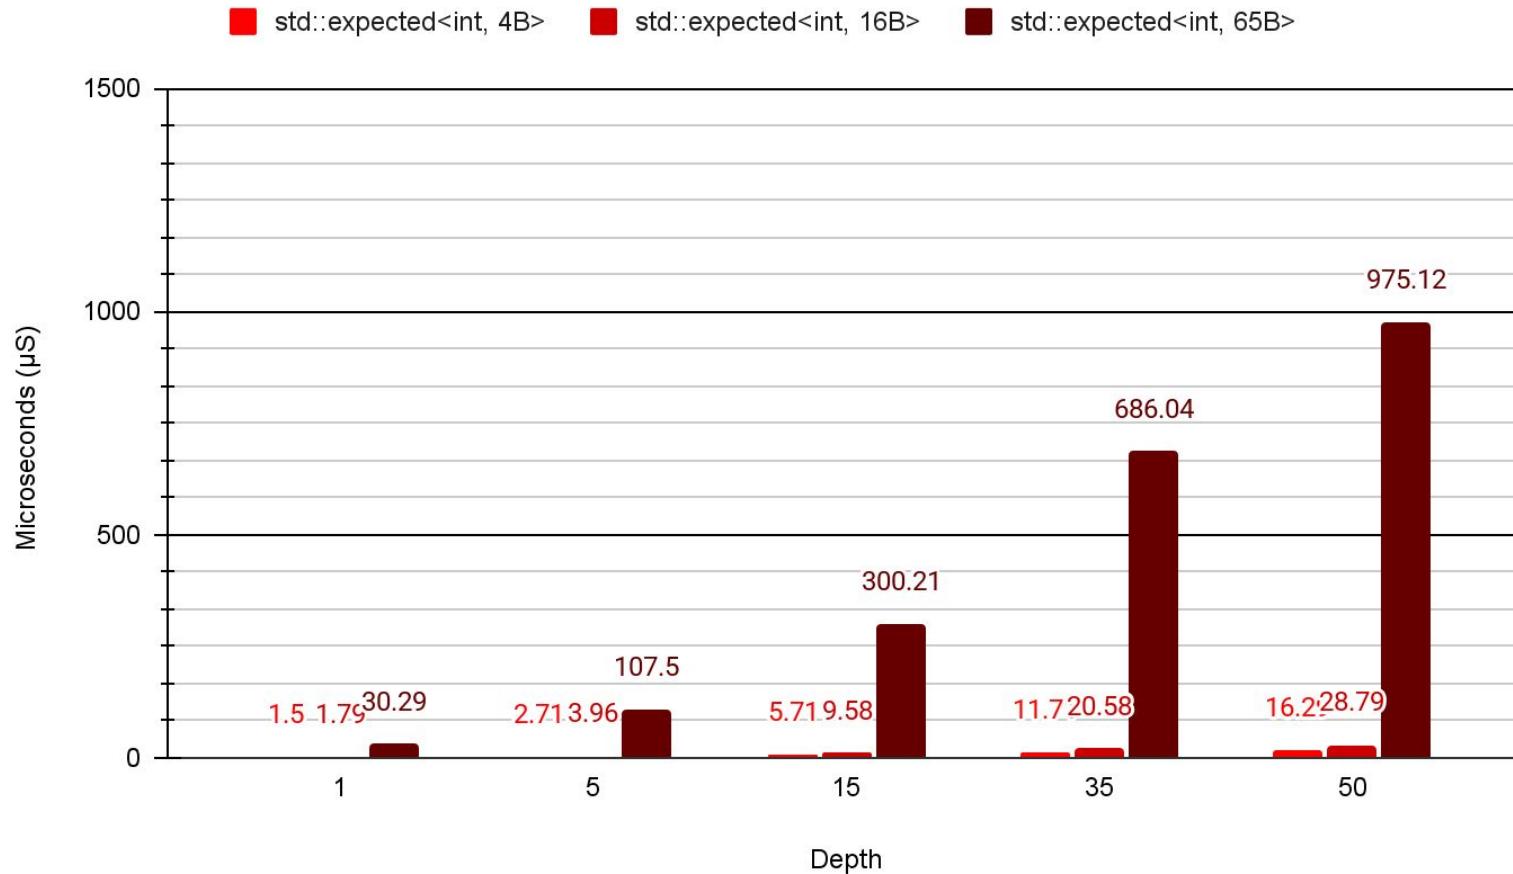
GCC Exceptions 0% Cleanup, Variable Error Sizes



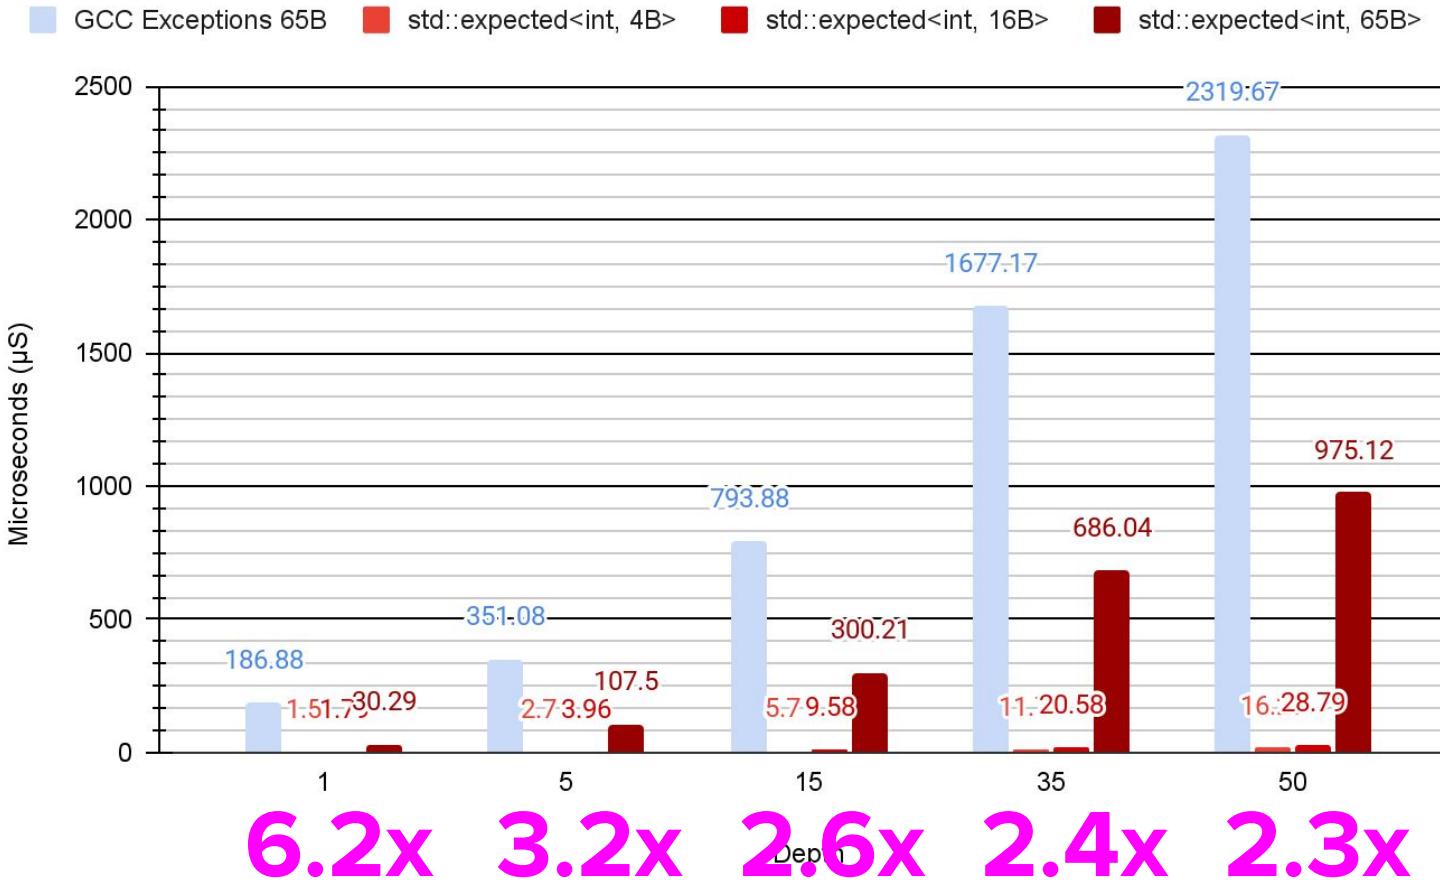
*LINEAR SCALE std::expected<T,E> 0% Cleanup, Variable Error Sizes



std::expected<T,E> 0% Cleanup, Variable Error Sizes



Propagation Time: 0% Cleanup, errors size [4B, 16B, 65B]



Why the slow down?

```
08000f24 <depth_13_error_65_percent_000 ()>:  
@ Other stuff...  
8000f44: f7ff ffa8    bl  8000e98  
<depth_12_error_65_percent_000 ()>  
8000f48: f89d 608c    ldrb.w  r6, [sp, #140]  @ 0x8c  
8000f4c: b976          cbnz  r6, 8000f6c  
<depth_13_error_65_percent_000 ()+0x48>  
8000f4e: 2241          movs  r2, #65 @ 0x41  
8000f50: a912          add   r1, sp, #72 @ 0x48  
8000f52: a801          add   r0, sp, #4  
8000f54: f00a f99c    bl  800b290 <memcpy>  
8000f58: 2241          movs  r2, #65 @ 0x41  
8000f5a: 4628          mov   r0, r5  
8000f5c: a901          add   r1, sp, #4  
8000f5e: f00a f997    bl  800b290 <memcpy>
```

```
08000e80 <depth_13_error_16_percent_000 ()>:  
@ Other stuff...  
8000ea0: f7ff ffae    bl  8000e00  
<depth_12_error_16_percent_000 ()>  
8000ea4: f89d 3024    ldrb.w  r3, [sp, #36] @ 0x24  
8000ea8: b95b          cbnz  r3, 8000ec2  
<depth_13_error_16_percent_000 ()+0x42>  
8000eaa: 7423          strb  r3, [r4, #16]  
8000eac: f10d 0c04    add.w ip, sp, #4  
8000eb0: ab05          add   r3, sp, #20  
8000eb2: cb0f          ldmia r3, {r0, r1, r2, r3}  
8000eb4: e88c 000f    stmia.w ip, {r0, r1, r2, r3}  
8000eb8: e884 000f    stmia.w r4, {r0, r1, r2, r3}  
8000ebc: 4620          mov   r0, r4  
8000ebe: b00b          add   sp, #44 @ 0x2c  
8000ec0: bd30          pop   {r4, r5, pc}
```

What we've learned

- Result types are VERY fast for errors $\leq 64B$
- Cleanup for exceptions is slow
- All runtime is linear for all parameters
- No non-determinism found with exceptions
 - Same numbers $\pm 1\text{us}$

All benchmarks going forward will use 4B (`sizeof(std::errc)`) as the error size.

Where do I begin?



What takes up the most time?

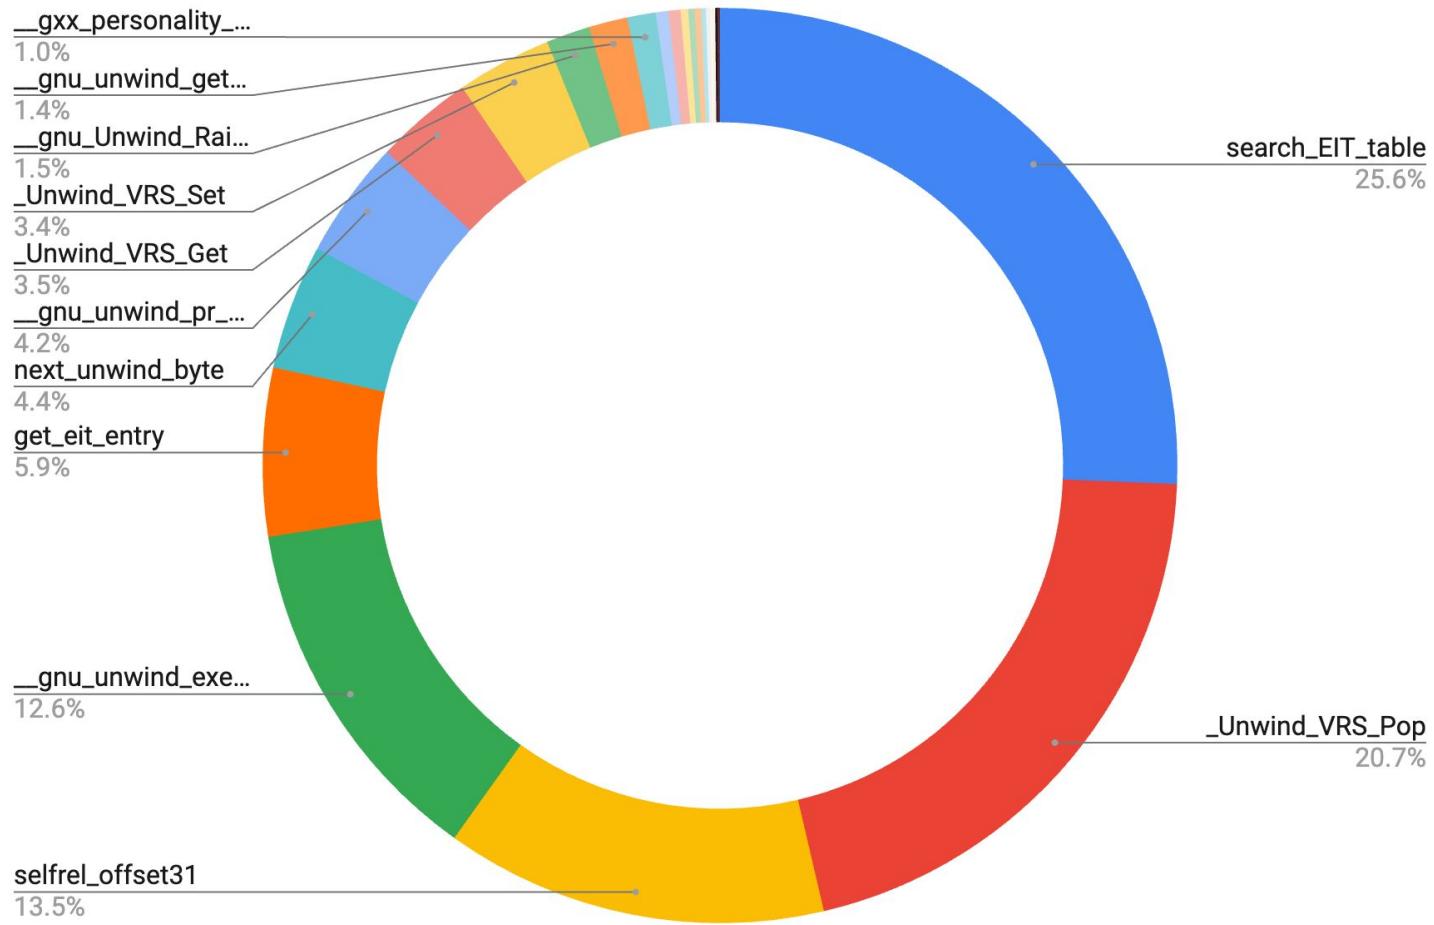


Python: control GDB & record instructions per function

```
my_stdout.out("step,function,address")
for step in track(range(instruction_steps),
                  console=my_stderr,
                  description='Processing...'):
    response = gdbmi.write('stepi')
    for message in response:
        payload = message['payload']
        if payload and isinstance(payload, dict) and 'frame' in payload:
            if payload['frame']['func'] == "start":
                # Continue from here in order to reach the next call to throw
                gdbmi.write('continue', read_response=False)
                continue
            my_stdout.out("{},{},{}".format(
                step,
                payload['frame']['func'],
                payload['frame']['addr']
            )))

```

V0

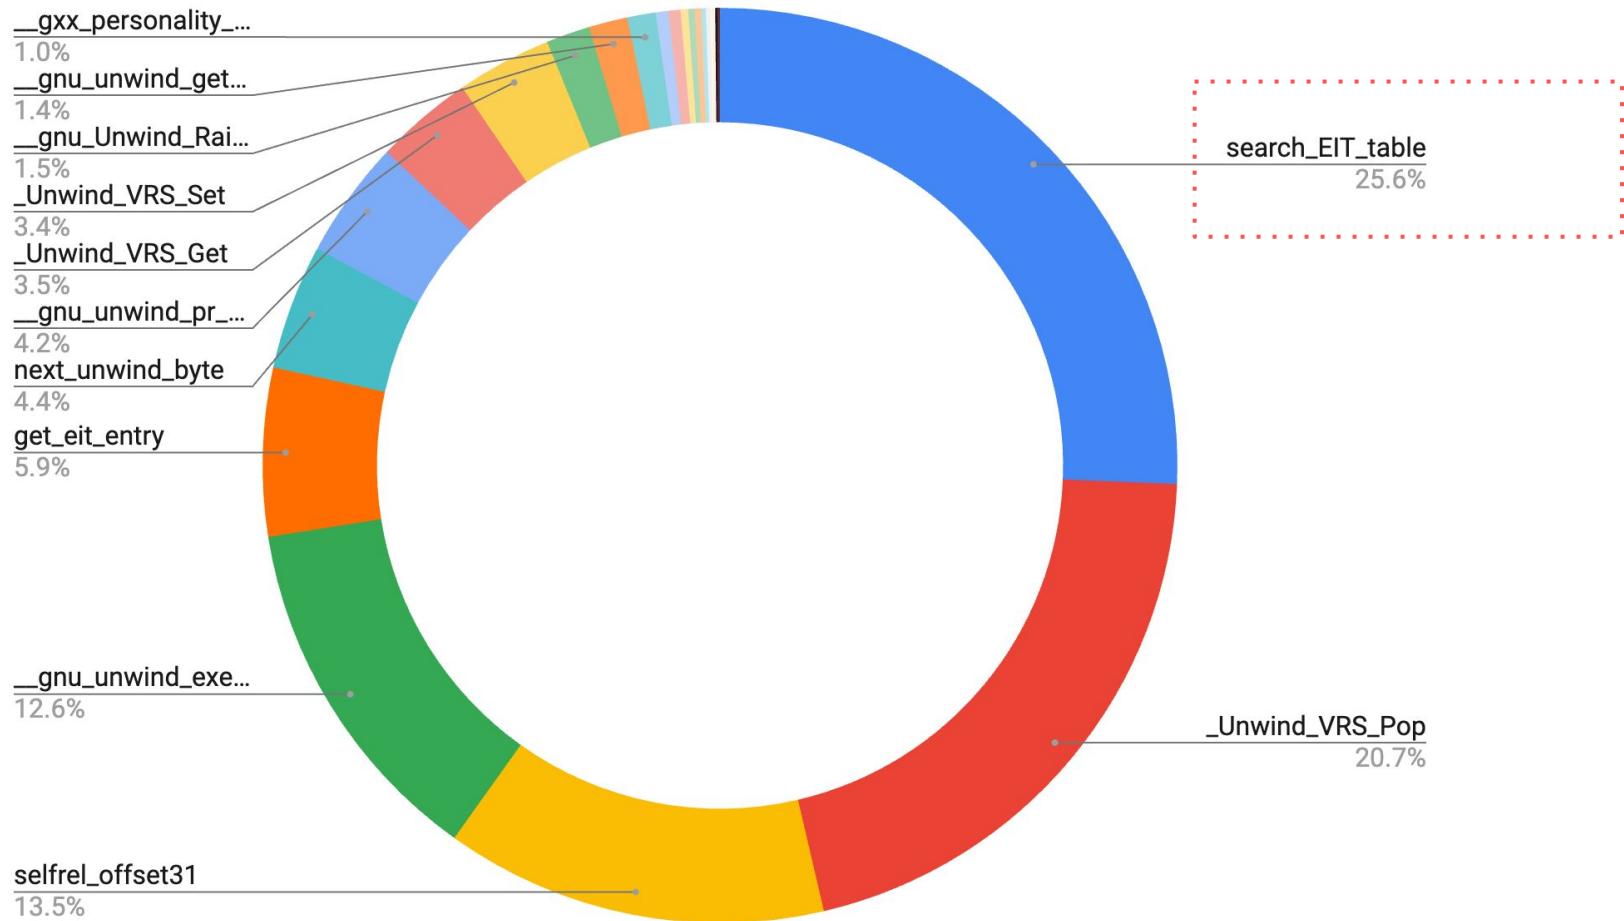


Chapter 2.

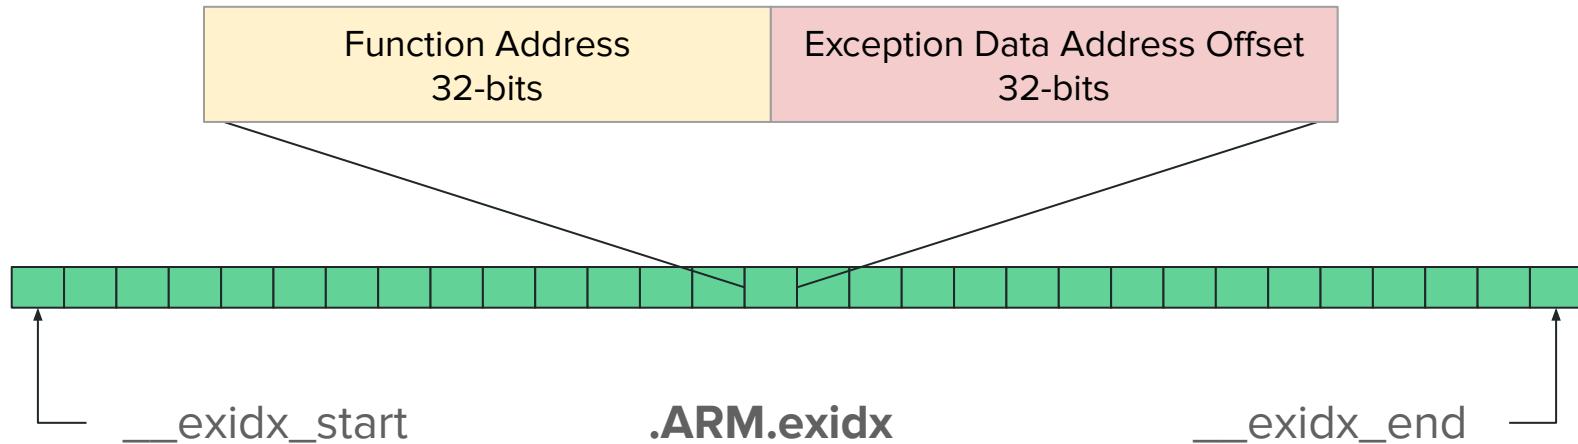


Optimizing GCC exceptions

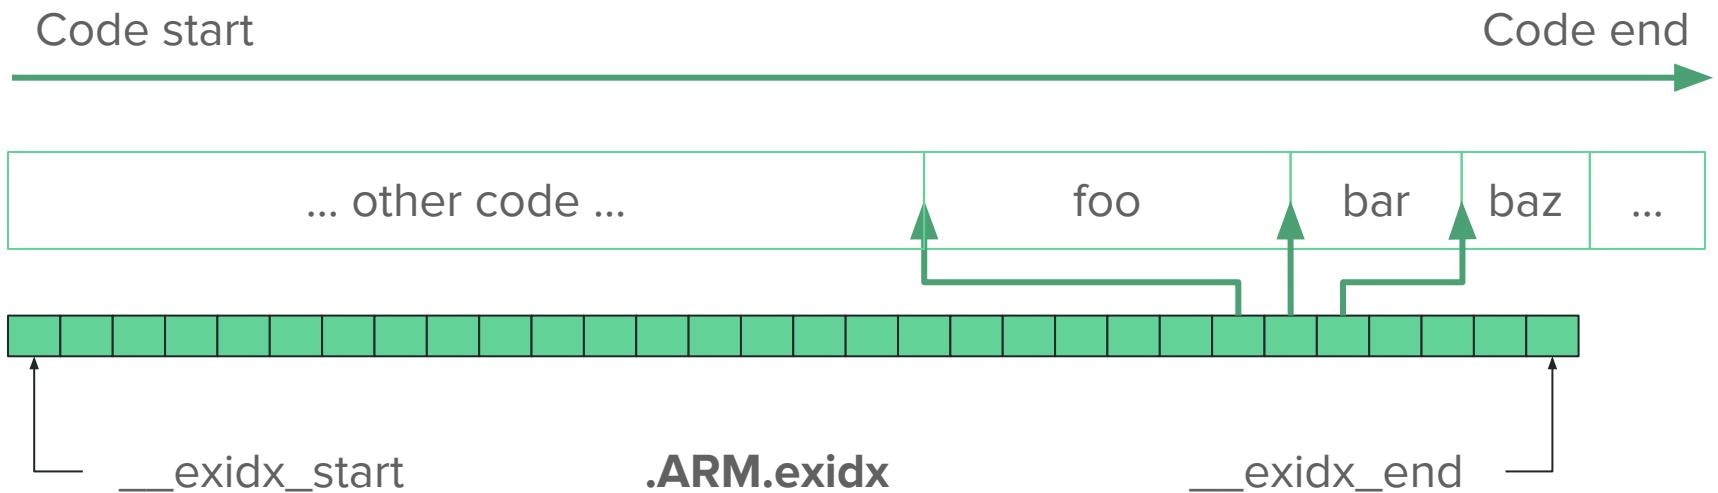
V0



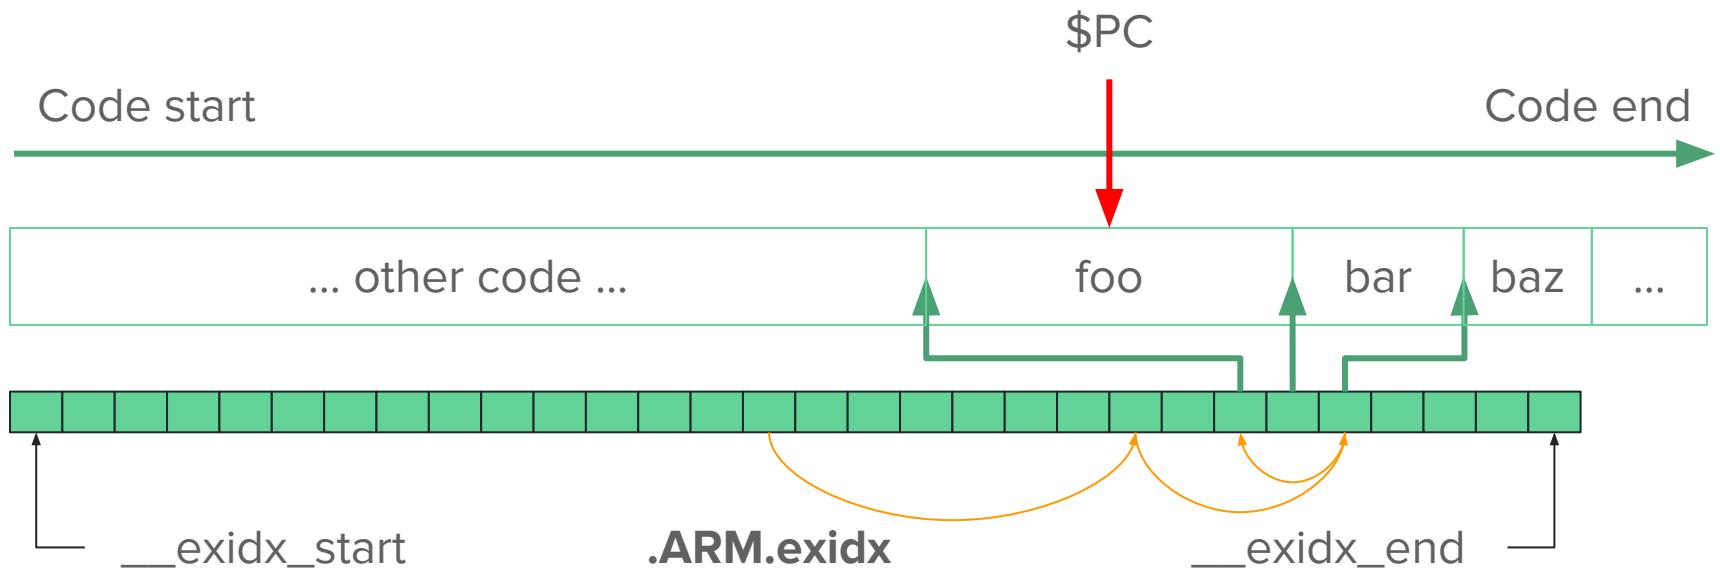
ARM Exception Index: Index Entry



ARM Exception Index: Ordering



ARM Exception Index: Ordering



```
static __EIT_entry const* search_EIT_table(
    __EIT_entry const* table, int nrec, _uw return_address)
{
    _uw next_fn;
    _uw this_fn;
    int n, left, right;
    if (nrec == 0) return (__EIT_entry*)0;
    left = 0;
    right = nrec - 1;
    while (1) {
        n = (left + right) / 2;
        this_fn = selfrel_offset31(&table[n].fnoffset);
        if (n != nrec - 1) next_fn = selfrel_offset31(&table[n + 1].fnoffset) - 1;
        else next_fn = (_uw)0 - 1;
        if (return_address < this_fn) {
            if (n == left) return (__EIT_entry*)0;
            right = n - 1;
        } else if (return_address <= next_fn) return &table[n];
        else left = n + 1;
    }
}
```



TO EXIT FULL SCREEN, PRESS **ESC**



+ 23



Andrei Alexandrescu

Robots Are after Your Job
(or at Least the Boring Parts
of It): Exploring Generative
AI for C++

Commonly Taught `binary_search`

The "violet eyes" version

- First comparison in the middle of the range
- Each pass reduces searched space in half
- 1-2 tests per pass (some consider it 1)
- May return early (as soon as element found)
- Not found decision made after range collapsed

```
template <typename I, typename T>
bool
binary_search(I b, I e, const T& v) {
    while (b < e) {
        auto m = b + (e - b) / 2;
        if (v < *m) {
            e = m;
        } else if (*m < v) {
            b = m + 1;
        } else {
            return true;
        }
    }
    return false;
}
```



7

Video Sponsorship Provided By:

ansatz

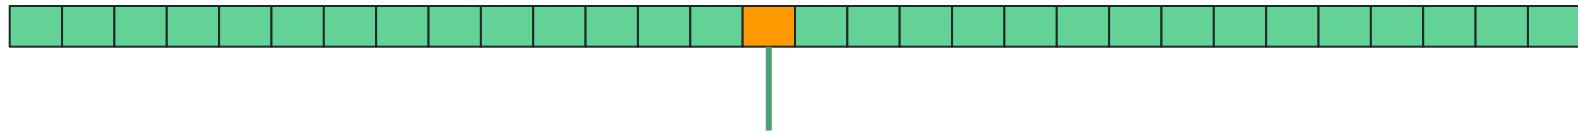
think-cell



13:10 / 1:33:08



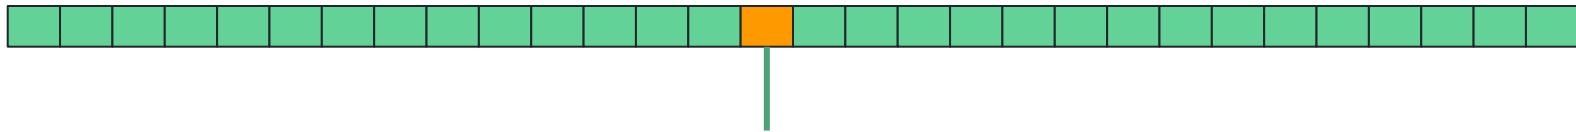
Imagine a list with 1024 entries



???

chance

Imagine a list with 1024 entries



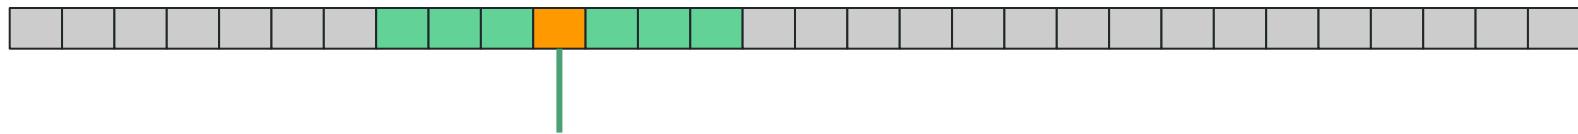
1/1024
chance

Imagine a list with 1024 entries



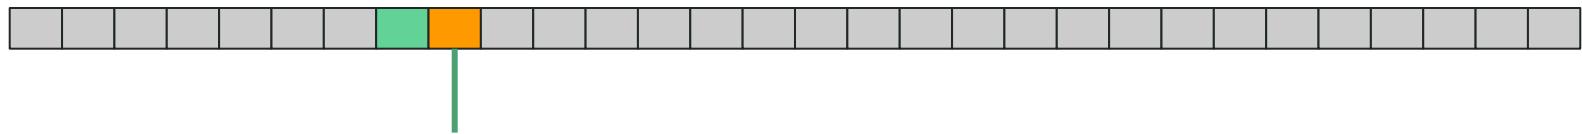
1/512
chance

Imagine a list with 1024 entries



1/256 chance

Imagine a list with 1024 entries

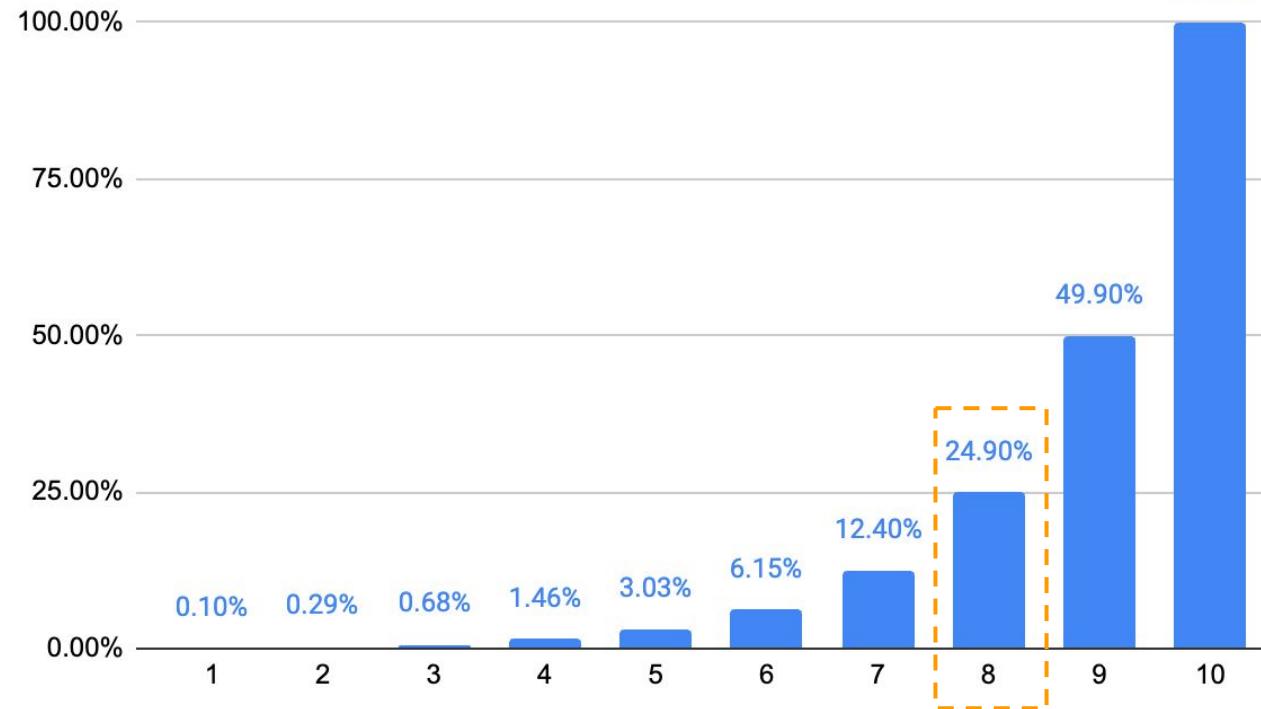


50-50
chance

24.90% chance this guess early returns



Binary Search Early Return Probability (set size = 1024) 99.90%





Andrei Alexandrescu

Robots Are after Your Job
(or at Least the Boring Parts
of It): Exploring Generative
AI for C++

"Industrial" std::lower_bound

- Uses indexing
- No early return
- Always performs $\log_2(n)$ steps (worst case)
 - However, one call to less per pass
 - Less branchy
- Everything *left* of the result is *strictly less*
- More informative/actionable than just a bool
- Similar std::upper_bound from the other side
- Finally, std::equal_range returns both

NVIDIA 9

```
template <class T, class I, class P>
I
lower_bound(I b, I e, const T& v, P p) {
    for (size_t size = e - b; size > 0; ) {
        size_t cut = size / 2;
        if (p(b[cut], v)) {
            b += ++cut;
            size -= cut;
        } else {
            size = cut;
        }
    }
    return b;
}
```

Video Sponsorship Provided By:

ansatz

think-cell

```
static __EIT_entry const* search_EIT_table(
    __EIT_entry const* table, int nrec, _uw return_address)
{
    _uw next_fn;
    _uw this_fn;
    int n, left, right;
    if (nrec == 0) return (__EIT_entry*)0;
    left = 0;
    right = nrec - 1;
    while (1) {
        n = (left + right) / 2;
        this_fn = selfrel_offset31(&table[n].fnoffset);
        if (n != nrec - 1) next_fn = selfrel_offset31(&table[n + 1].fnoffset) - 1;
        else next_fn = (_uw)0 - 1;
        if (return_address < this_fn) {
            if (n == left) return (__EIT_entry*)0;
            right = n - 1;
        } else if (return_address <= next_fn) return &table[n];
        else left = n + 1;
    }
}
```

```
static __EIT_entry const* search_EIT_table(
    __EIT_entry const* table, int nrec, _uw return_address)
{
    _uw next_fn;
    _uw this_fn;
    int n, left, right;
    if (nrec == 0) return (__EIT_entry*)0;
    left = 0;
    right = nrec - 1;
    while (1) {
        n = (left + right) / 2;
        this_fn = selfrel_offset31(&table[n].fnoffset);
        if (n != nrec - 1) next_fn = selfrel_offset31(&table[n + 1].fnoffset) - 1;
        else next_fn = (_uw)0 - 1;
        if (return_address < this_fn) {
            if (n == left) return (__EIT_entry*)0;
            right = n - 1;
        } else if (return_address <= next_fn) return &table[n];
        else left = n + 1;
    }
}
```



```
#include <span>
#include <algorithm>

__EIT_entry const* search_EIT_table(__EIT_entry const* table,
                                    int nrec,
                                    std::uint32_t return_address)

{
    if (nrec == 0) {
        return nullptr;
    }

    auto const& next_entry = std::ranges::upper_bound(
        std::span(table, nrec), return_address, eit_entry_less_than{});
    auto const& actual_entry = *(next_entry - 1);

    return &actual_entry;
}
```

```
#include <span>
#include <algorithm>

__EIT_entry const* search_EIT_table(__EIT_entry const* table,
                                    int nrec,
                                    std::uint32_t return_address)

{
    if (nrec == 0) {
        return nullptr;
    }

    start_cycles = uptime();

    auto const& next_entry = std::ranges::upper_bound(
        std::span(table, nrec), return_address, eit_entry_less_than{});
    auto const& actual_entry = *(next_entry - 1);

    end_cycles = uptime();

    return &actual_entry;
}
```

RESULT! For 710 functions

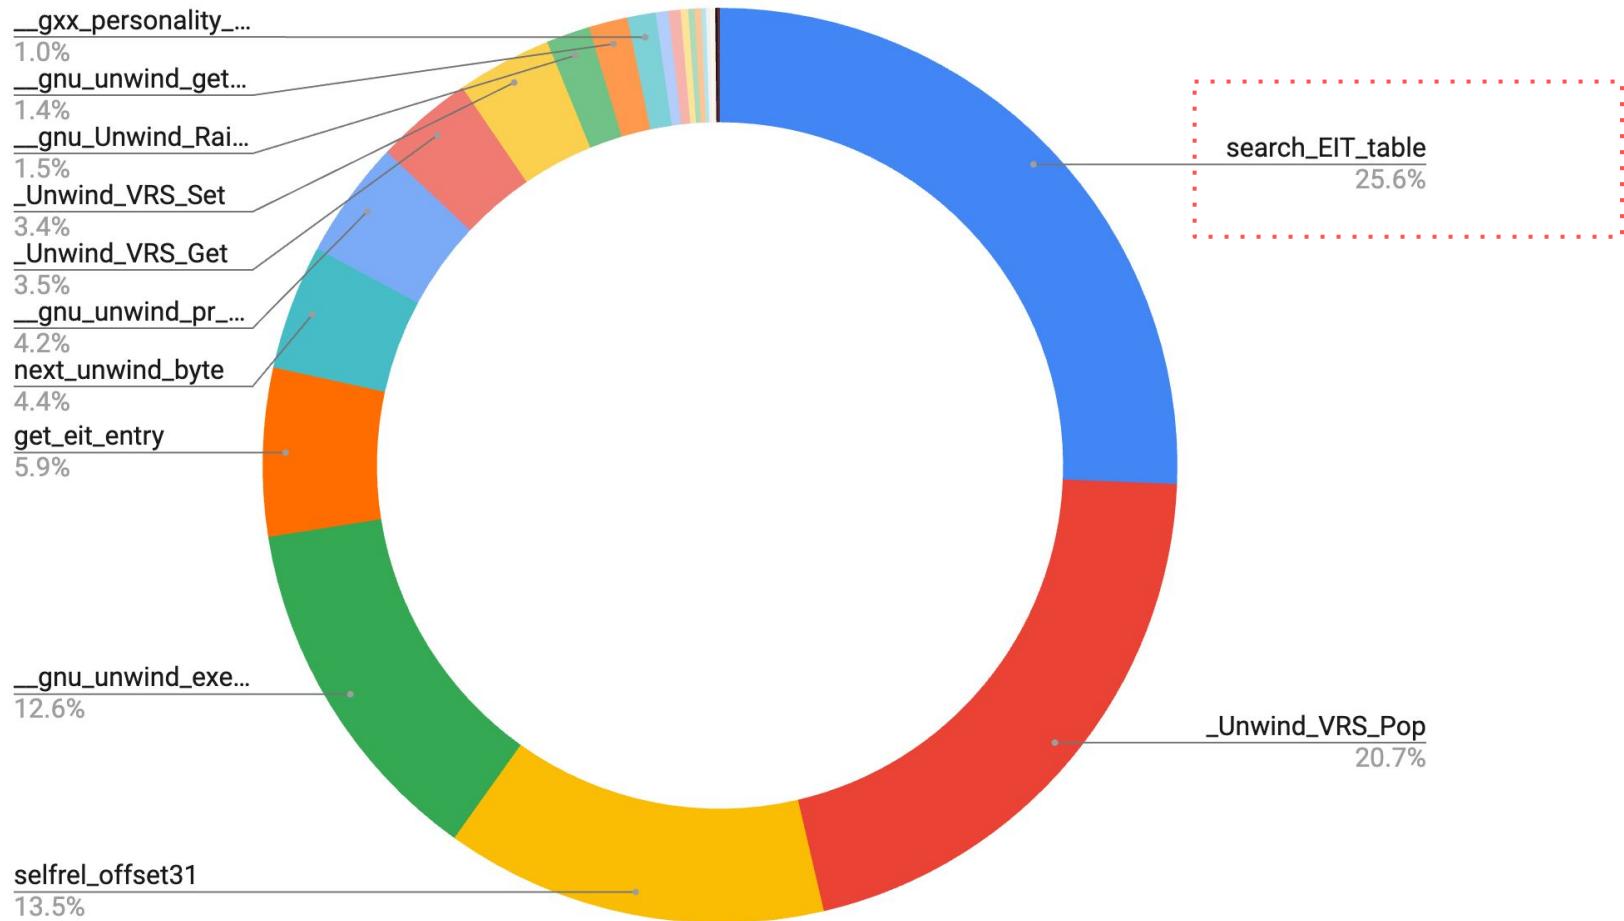
Implementation	Cycles		
	Average	Minimum	Maximum
GCC	379	148	477

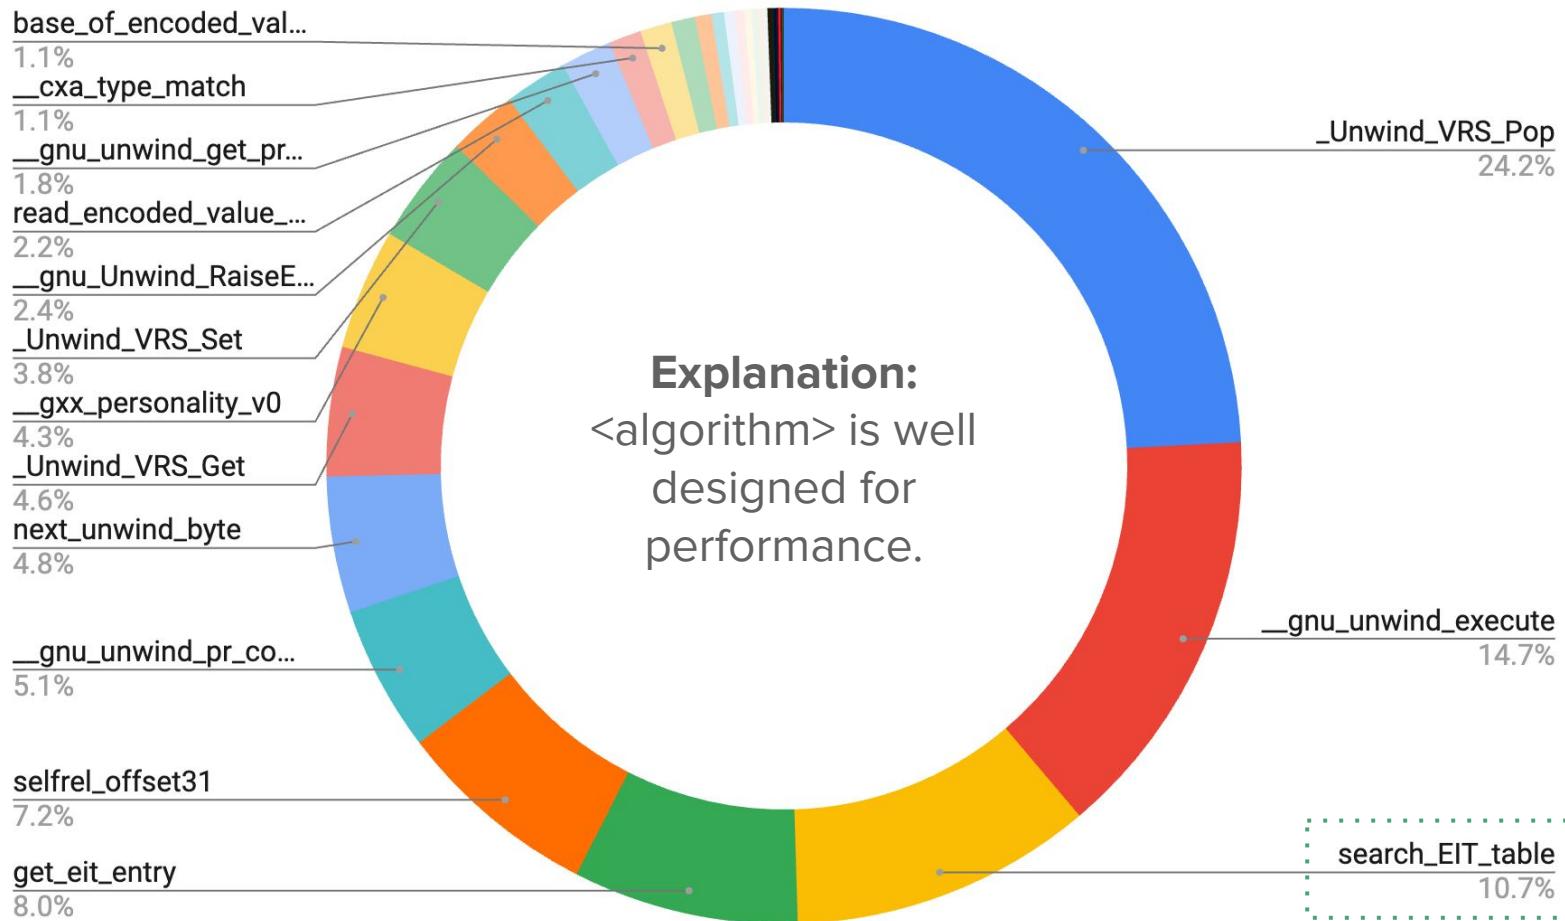
RESULT! For 710 functions

Implementation	Cycles		
	Average	Minimum	Maximum
GCC 	379	148	477
std::upper_bound 	270	261	288

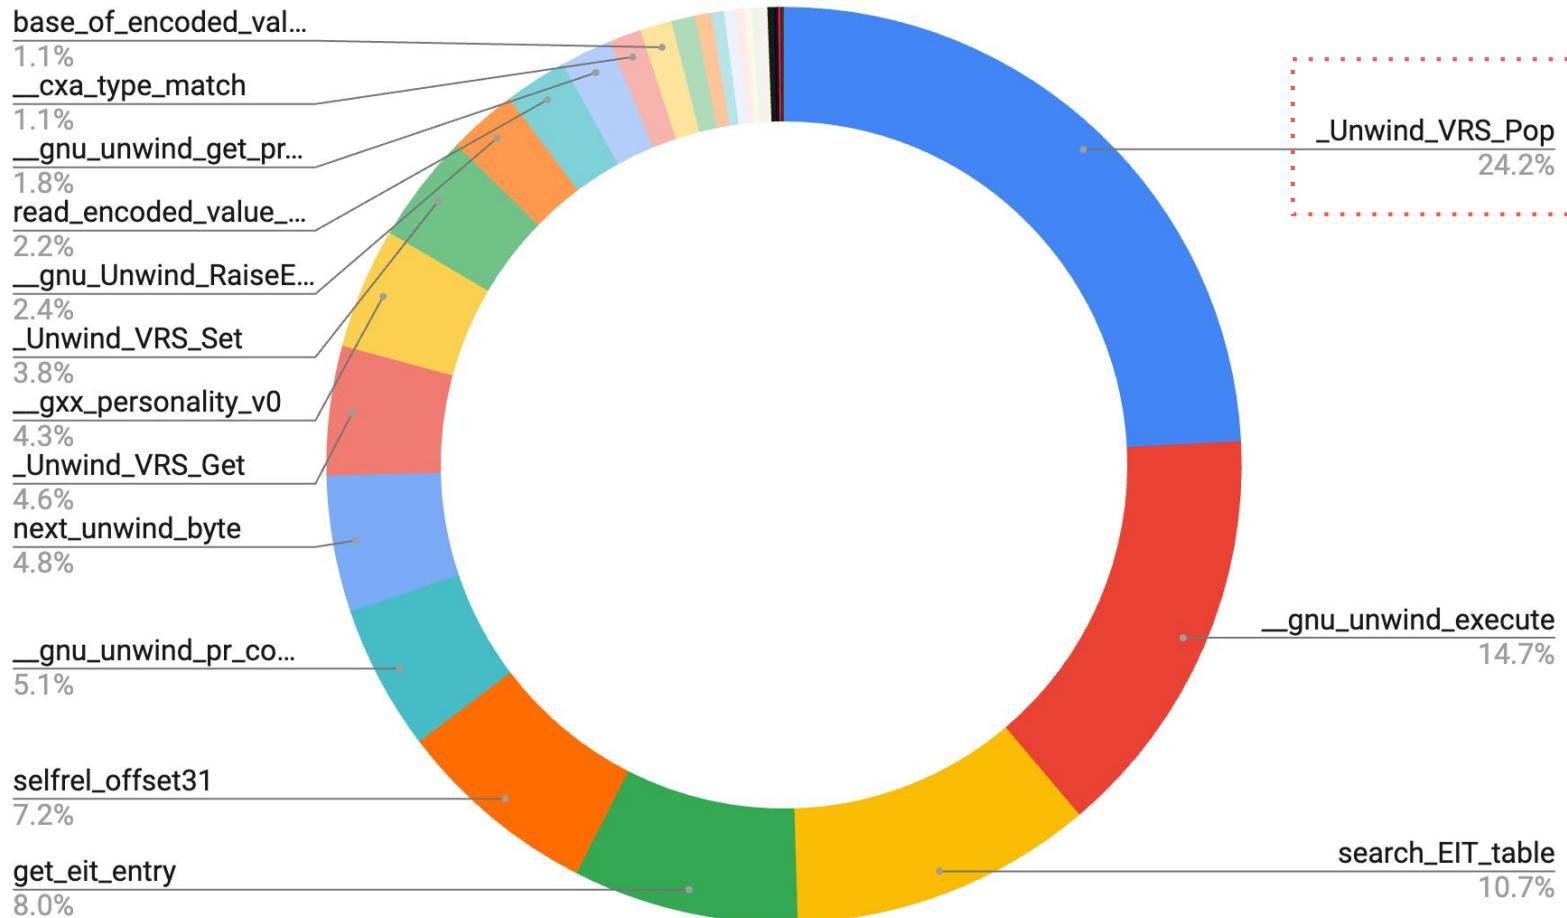
-28.8% the needed cycles on average

V0





V1

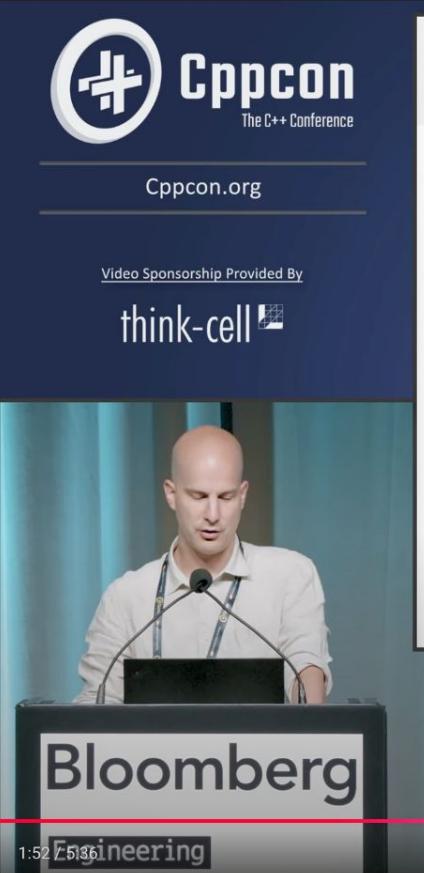


```
/* ABI defined function to pop registers off the stack. */
_Unwind_VRS_Result _Unwind_VRS_Pop (_Unwind_Context *context,
                                     _Unwind_VRS_RegClass regclass,
                                     _uw discriminator,
                                     _Unwind_VRS_DataRepresentation representation)
{
    phasel_vrs *vrs = (phasel_vrs *) context;
    switch (regclass) {
        case _UVRSC_CORE: {
            _uw *ptr;
            _uw mask;
            int i;
            if (representation != _UVRSD_UINT32)
                return _UVRSR_FAILED;

            mask = discriminator & 0xffff;
            ptr = (_uw *) vrs->core.r[R_SP];
            /* Pop the requested registers. */
            for (i = 0; i < 16; i++)
            {
                if (mask & (1 << i))
                    vrs->core.r[i] = *(ptr++);
            }
            /* The rest ... */
        }
    }
}
```

```
/* ABI defined function to pop registers off the stack. */  
  
_Unwind_VRS_Result _Unwind_VRS_Pop (_Unwind_Context *context,  
    _Unwind_VRS_RegClass regclass,  
    _uw discriminator,  
    _Unwind_VRS_DataRepresentation representation)  
{  
  
    phasel_vrs *vrs = (phasel_vrs *) context;  
    switch (regclass) {  
        case _UVRSC_CORE: {  
            _uw *ptr;  
            _uw mask;  
            int i;  
            if (representation != _UVRSD_UINT32)  
                return _UVRSP_FAILED;  
  
            mask = discriminator & 0xffff;  
            ptr = (_uw *) vrs->core.r[R_SP];  
            /* Pop the requested registers. */  
            for (i = 0; i < 16; i++)  
            {  
                if (mask & (1 << i))  
                    vrs->core.r[i] = *(ptr++);  
            }  
            /* The rest ... */  
        }  
    }  
}
```

```
mask = discriminator & 0xffff;  
  
ptr = (_uw *) vrs->core.r[R_SP];  
/* Pop the requested registers. */  
for (i = 0; i < 16; i++)  
{  
    if (mask & (1 << i))  
        vrs->core.r[i] = *(ptr++);  
}
```



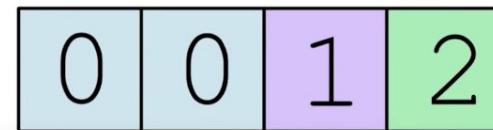
`push_back` with an $\mathcal{O}(n)$ budget

```
v.push_back(2)
```

Budget



Cost



Amortized O(1) Complexity

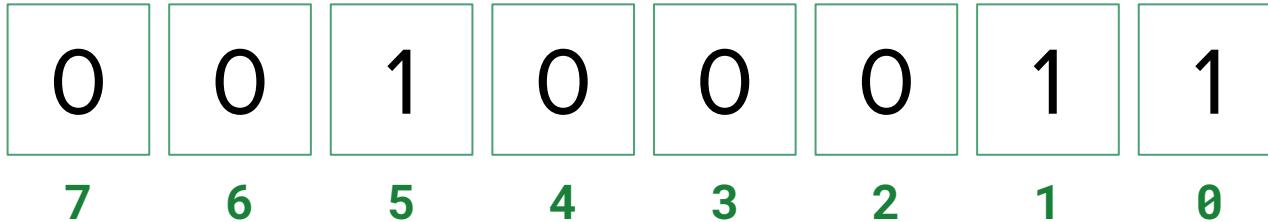
Andreas Weis

```
/* ABI defined function to pop registers off the stack. */  
  
_Unwind_VRS_Result _Unwind_VRS_Pop (_Unwind_Context *context,  
    _Unwind_VRS_RegClass regclass,  
    _uw discriminator,  
    _Unwind_VRS_DataRepresentation representation)  
{  
  
    phasel_vrs *vrs = (phasel_vrs *) context;  
    switch (regclass) {  
        case _UVRSC_CORE: {  
            _uw *ptr;  
            _uw mask;  
            int i;  
            if (representation != _UVRSD_UINT32)  
                return _UVRSR_FAILED;  
  
            mask = discriminator & 0xffff;  
            ptr = (_uw *) vrs->core.r[R_SP];  
            /* Pop the requested registers. */  
            for (i = 0; i < 16; i++)  
            {  
                if (mask & (1 << i))  
                    vrs->core.r[i] = *(ptr++);  
            }  
            /* The rest ... */  
    }  
}
```

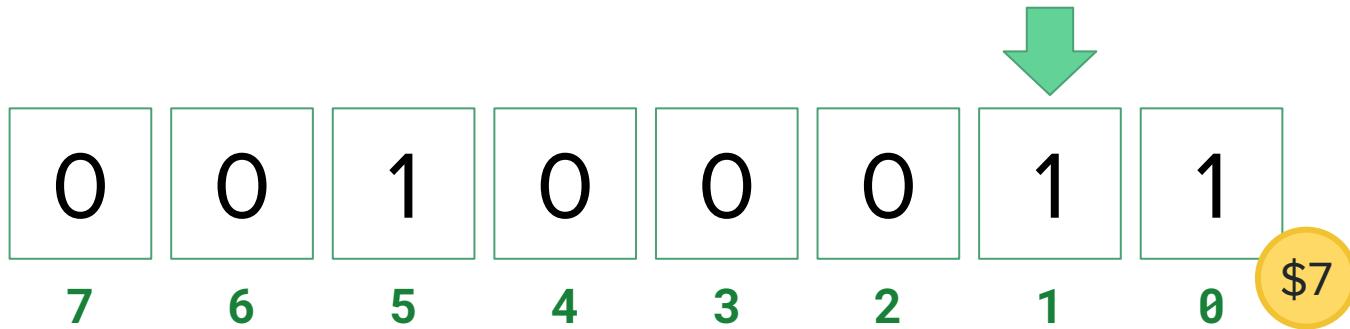


= 1x operation

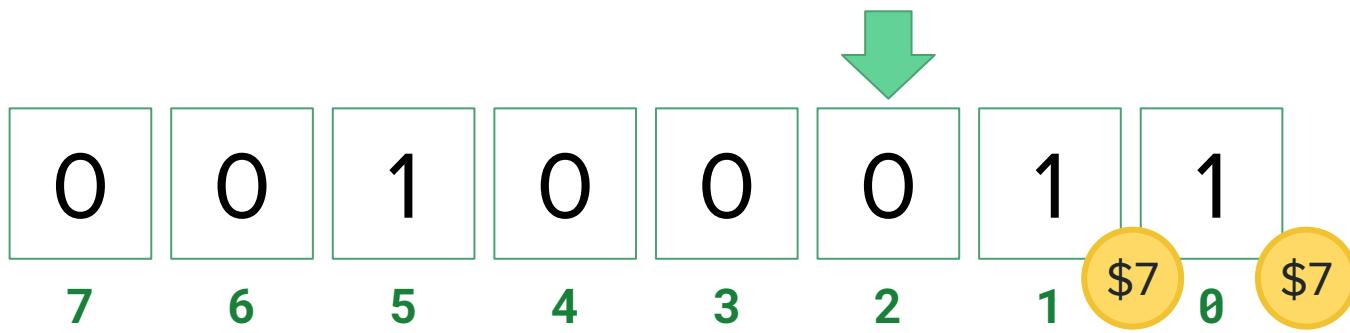
```
mask = discriminator & 0xffff;  
  
ptr = (_uw *) vrs->core.r[R_SP];  
  
/* Pop the requested registers. */  
  
for (i = 0; i < 16; i++)  
{  
  
    if (mask & (1 << i))  
  
        vrs->core.r[i] = *(ptr++);  
}  
}
```



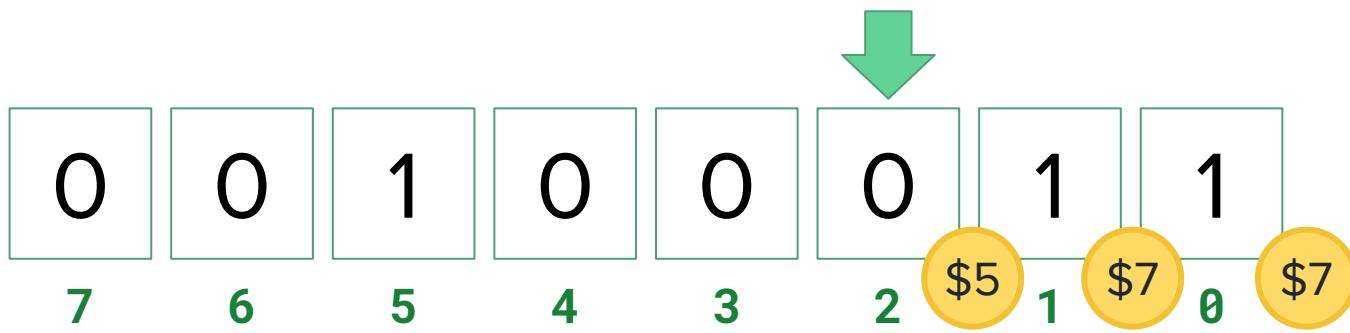
```
mask = discriminator & 0xffff;
ptr = (_uw *) vrs->core.r[R_SP];
/* Pop the requested registers. */
for (i = 0; i < 16; i++)
{
    if (mask & (1 << i))
        vrs->core.r[i] = * (ptr++);
}
```



```
mask = discriminator & 0xffff;
ptr = (_uw *) vrs->core.r[R_SP];
/* Pop the requested registers. */
for (i = 0; i < 16; i++)
{
    if (mask & (1 << i))
        vrs->core.r[i] = * (ptr++);
}
```



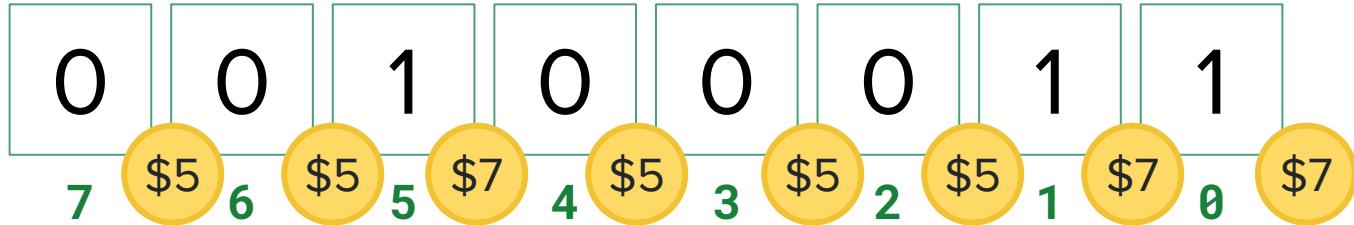
```
mask = discriminator & 0xffff;
ptr = (_uw *) vrs->core.r[R_SP];
/* Pop the requested registers. */
for (i = 0; i < 16; i++)
{
    if (mask & (1 << i))
        vrs->core.r[i] = * (ptr++);
}
```



```

mask = discriminator & 0xffff;
ptr = (_uw *) vrs->core.r[R_SP];
/* Pop the requested registers. */
for (i = 0; i < 16; i++)
{
    if (mask & (1 << i))
        vrs->core.r[i] = * (ptr++);
}

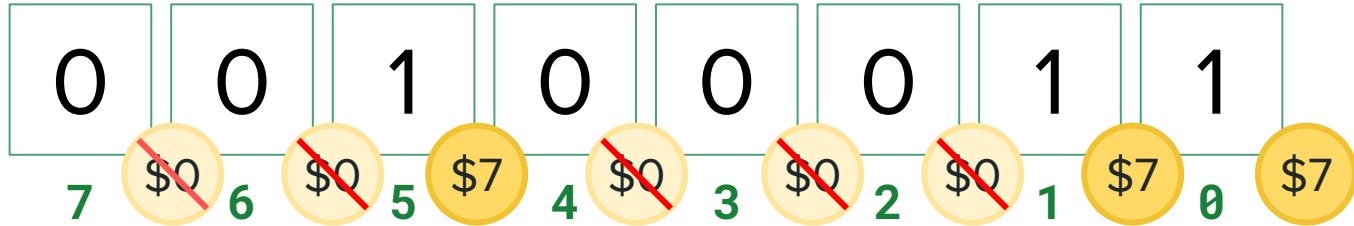
```



```
mask = discriminator & 0xffff;
ptr = (_uw *) vrs->core.r[R_SP];
/* Pop the requested registers. */
for (i = 0; i < 16; i++)
{
    if (mask & (1 << i))
        vrs->core.r[i] = * (ptr++);
}
```

TOTAL

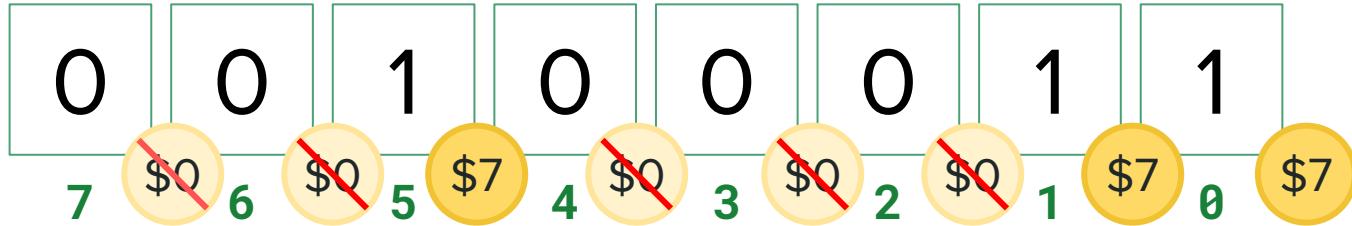
\$46



```
mask = discriminator & 0xffff;
ptr = (_uw *) vrs->core.r[R_SP];
/* Pop the requested registers. */
for (i = 0; i < 16; i++)
{
    if (mask & (1 << i))
        vrs->core.r[i] = * (ptr++);
}
```

TOTAL



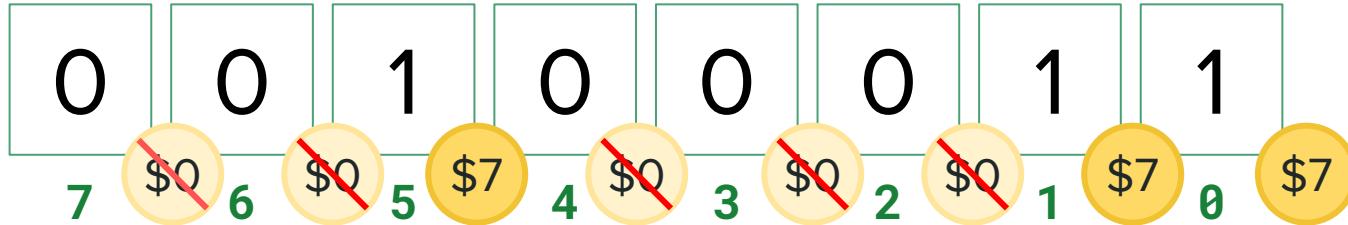


```
mask = discriminator & 0xffff;
ptr = (_uw *) vrs->core.r[R_SP];
/* Pop the requested registers. */
for (i = 0; i < 16; i++)
{
    if (mask & (1 << i))
        vrs->core.r[i] = * (ptr++);
}
```

TOTAL

\$21

-54% off



```
mask = discriminator & 0xffff;
ptr = (_uw *) vrs->core.r[R_SP];
/* Pop the requested registers. */
for (i = 0; i < 16; i++)
{
    if (mask & (1 << i))
        vrs->core.r[i] = * (ptr++);
}
```

BARGAIN DEAL!

\$21

-54% off

7	6	5	4	3	2	1	0
---	---	---	---	---	---	---	---

std::countz_zero

$$\left(\begin{array}{ccccccc} 0 & 0 & 1 & 0 & 0 & 0 & 1 \\ \end{array} \right) = 0$$

std::countz_zero

$$\left(\begin{array}{ccccccc} 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ \end{array} \right) = 6$$

std::countz_zero

Defined in header `<bit>`

```
template< class T >
constexpr int countz_zero( T x ) noexcept; (since C++20)
```

Returns the number of consecutive `0` bits in the value of `x`, starting from the least significant bit ("right").

ARM Cortex M4

```
rbit    r2, r3  
clz     r2, r2
```



X86-64

```
bsf edx, ecx
```



RISC-V

```
call    __ctzsi2
```



<https://godbolt.org/z/ezKEGGhGx>

[cppreference.com](#)

Create account

Page Discussion

Standard revision: Diff ▾ View Edit

C++ Utilities library Bit manipulation

std::countz_zero

Defined in header `<bit>`

```
template< class T >  
constexpr int countz_zero( T x ) noexcept;      (since C++20)
```

Returns the number of consecutive `0` bits in the value of `x`, starting from the least significant bit ("right").

```
while (discriminator) {  
    auto const reg_index = std::countz(discriminator);  
    vrs->core.r[reg_index] = * (ptr++);  
    discriminator &= ~ (1 << reg_index);  
}
```

std::countz

Defined in header `<bit>`

```
template< class T >  
constexpr int countz( T x ) noexcept; (since C++20)
```

Returns the number of consecutive `0` bits in the value of `x`, starting from the least significant bit ("right").

```
$ while (discriminator) {  
$     auto const reg_index = std::countz(discriminator);  
$     vrs->core.r[reg_index] = * (ptr++);  
$     discriminator &= ~ (1 << reg_index);  
$ }
```

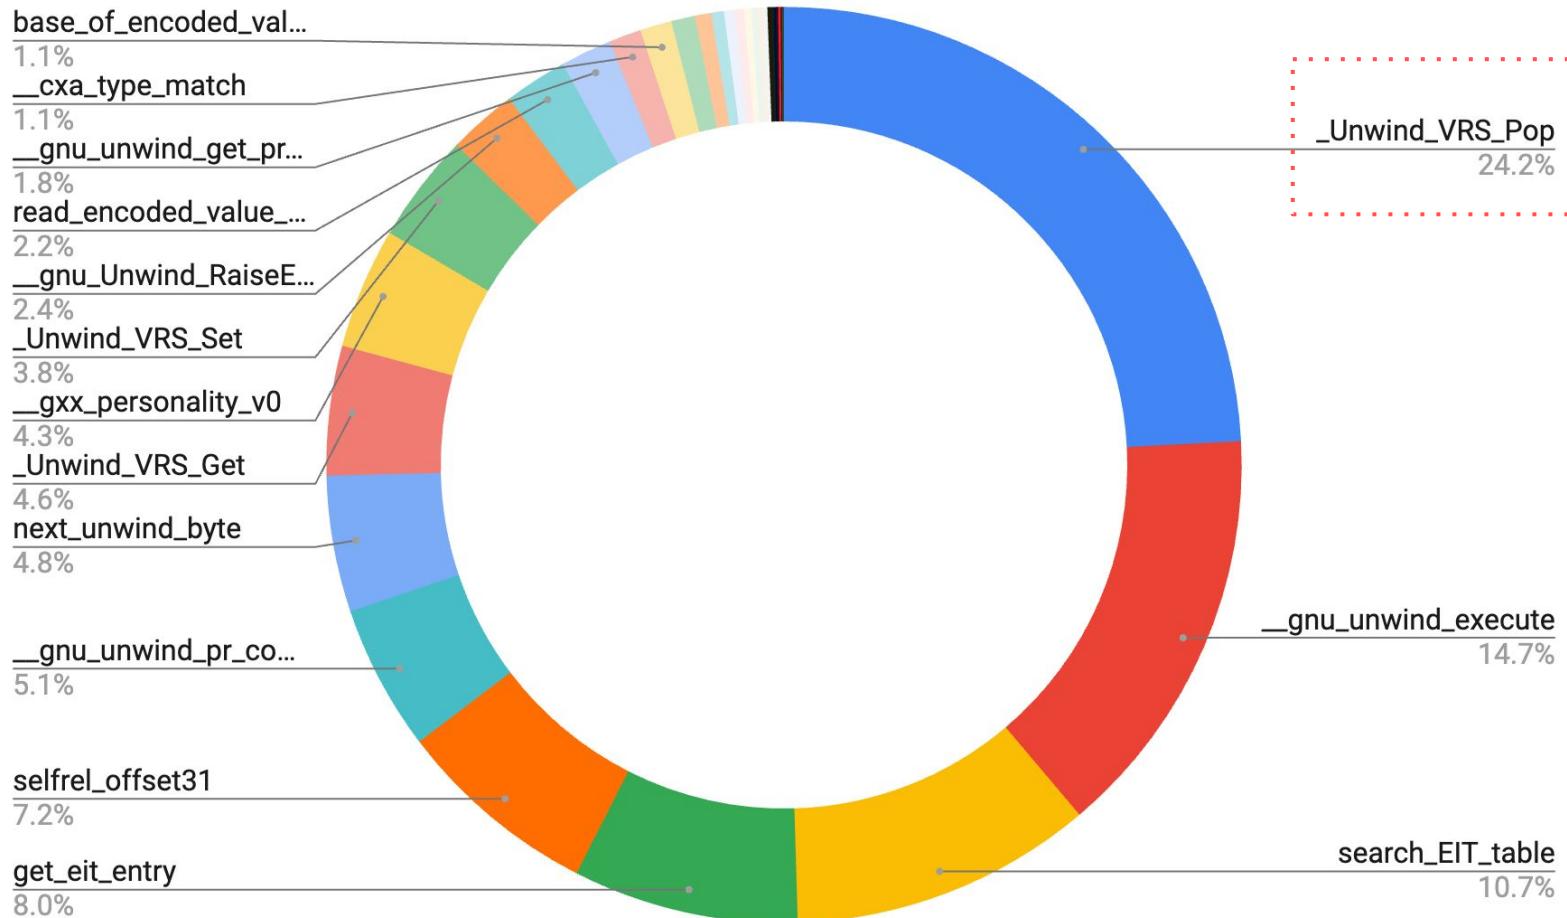
std::countz

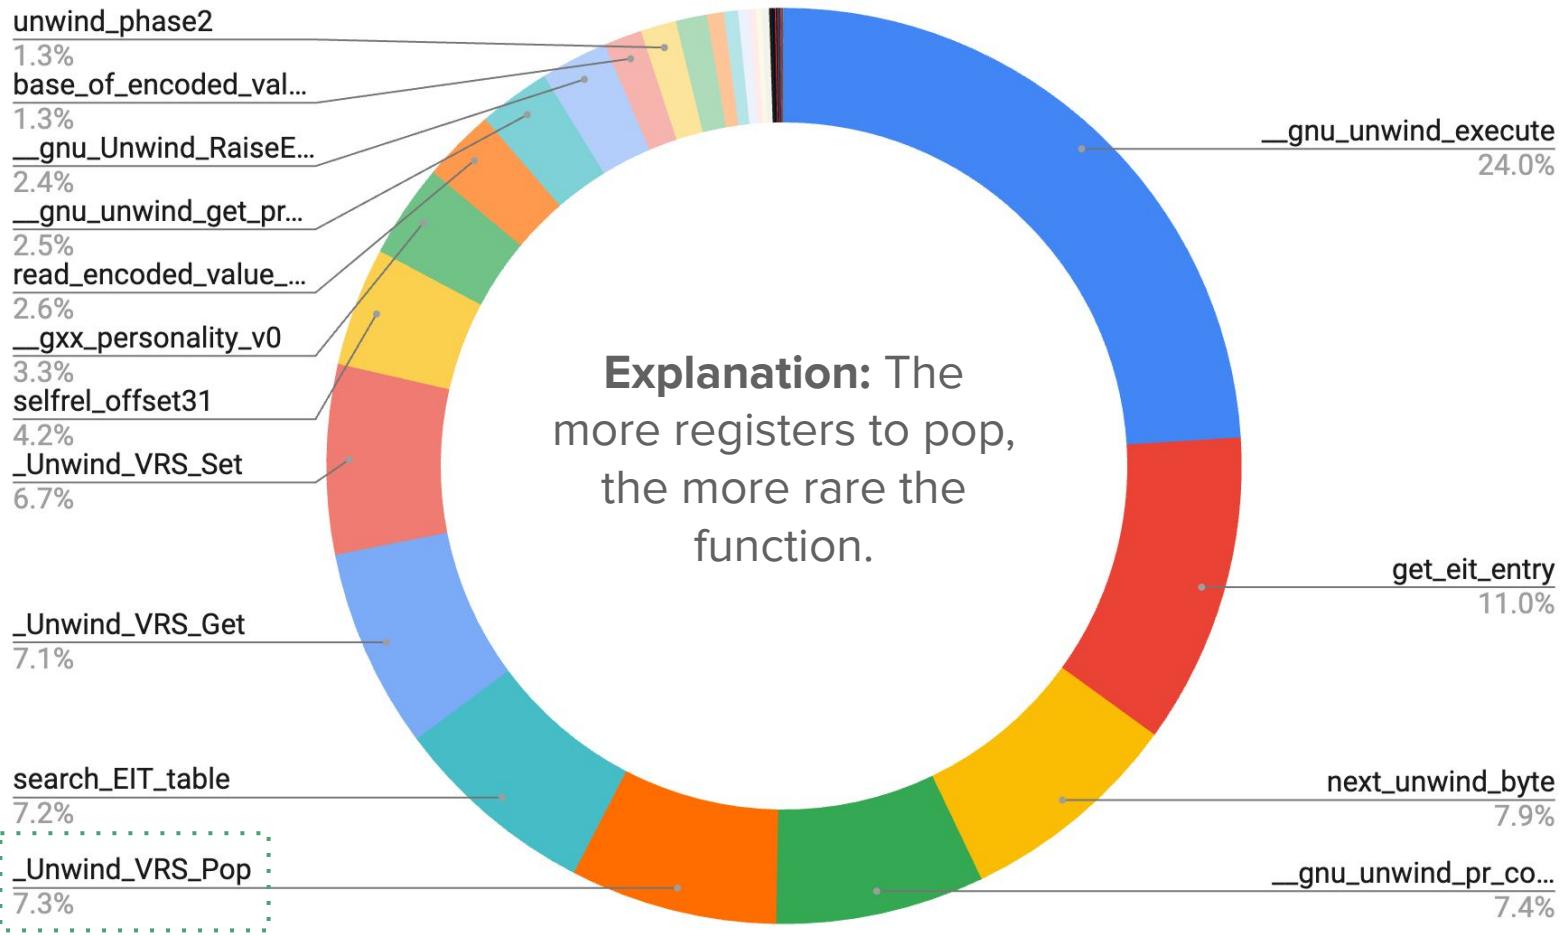
Defined in header `<bit>`

```
template< class T >  
constexpr int countz( T x ) noexcept; (since C++20)
```

Returns the number of consecutive `0` bits in the value of `x`, starting from the least significant bit ("right").

V1



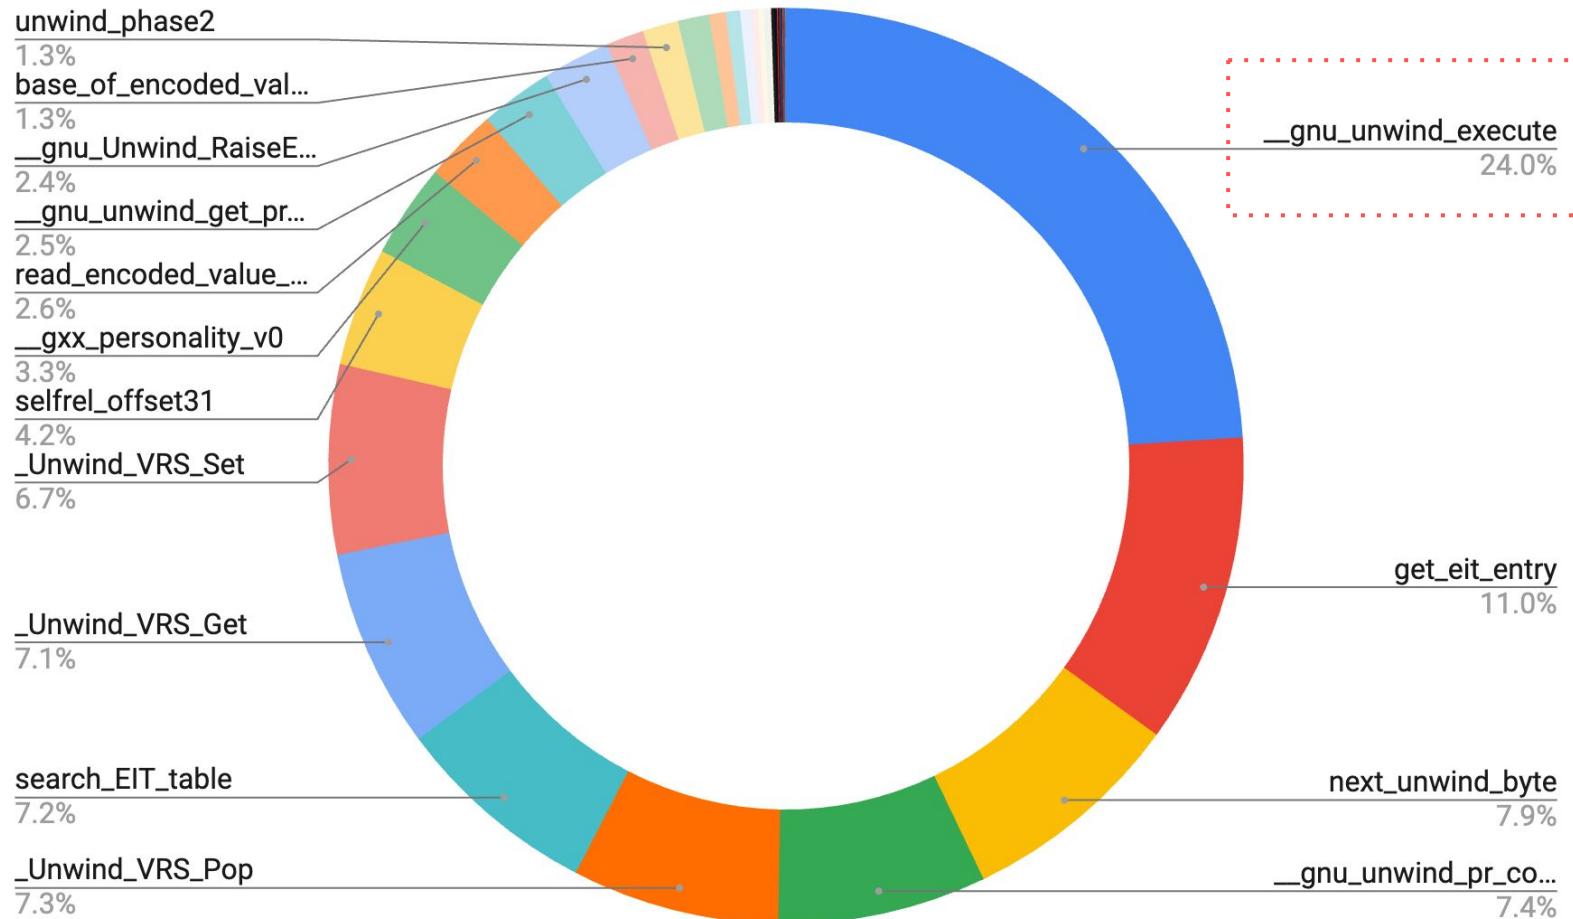


24%



7.3%

V2



```
/* Execute the unwinding instructions described by UWS.  */
_Unwind_Reason_Code __gnu_unwind_execute (_Unwind_Context * context,
                                         __gnu_unwind_state * uws)
{
    _uw op;
    int set_pc;
    _uw reg;
    _uw sp;
    set_pc = 0;
    for (;;) {
        op = [next_unwind_byte (uws)];
        if (op == CODE_FINISH) {
            if (!set_pc)
            {
                [_Unwind_VRS_Get] (context, _UVRSC_CORE, R_LR, _UVRSD_UINT32,
                                    &reg);
                [_Unwind_VRS_Set] (context, _UVRSC_CORE, R_PC, _UVRSD_UINT32,
                                    &reg);
                set_pc = 1;
            }
            /* Drop out of the loop.  */
            break;
        }
    }
}
```

```
/* Execute the unwinding instructions described by UWS.  */
_Undertake_Reason_Code __gnu_unwind_execute (_Unwind_Context * context,
                                              __gnu_unwind_state * uws)
{
    _uw op;
    int set_pc;
    _uw reg;
    _uw sp;
    set_pc = 0;
    for (;;) {
        op = next_unwind_byte (uws);
        if (op == CODE_FINISH) {
            if (!set_pc)
            {
                _Unwind_VRS_Get (context, _UVRSC_CORE, R_LR, _UVRSD_UINT32,
                                 &reg);
                _Unwind_VRS_Set (context, _UVRSC_CORE, R_PC, _UVRSD_UINT32,
                                 &reg);
                set_pc = 1;
            }
            /* Drop out of the loop.  */
            break;
        }
    }
}
```

```
static inline _uw8 next_unwind_byte (
    __gnu_unwind_state * uws)
{
    _uw8 b;

    if (uws->bytes_left == 0)
    {
        /* Load another word */
        if (uws->words_left == 0)
            return CODE_FINISH; /* Nothing left.  */
        uws->words_left--;
        uws->data = *(uws->next++);
        uws->bytes_left = 3;
    }
    else
        uws->bytes_left--;

    /* Extract the most significant byte.  */
    b = (uws->data >> 24) & 0xff;
    uws->data <<= 8;
    return b;
}
```

```
/* Execute the unwinding instructions described by UWS. */
_Unwind_Reason_Code __gnu_unwind_execute (_Unwind_Context * context,
                                         __gnu_unwind_state
{
    _uw op;
    int set_pc;
    _uw reg;
    _uw sp;
    set_pc = 0;
    for (;;) {
        op = next_unwind_byte (uws);
        if (op == CODE_FINISH) {
            if (!set_pc)
            {
                __Unwind_VRS_Get (context, _UVRSC_CORE, R_LR, _UVRSD_UI
                    &reg);
                __Unwind_VRS_Set (context, _UVRSC_CORE, R_PC, _UVRSD_UI
                    &reg);
                set_pc = 1;
            }
            /* Drop out of the loop. */
            break;
        }
    }
}
```

```
_Unwind_VRS_Result __Unwind_VRS_Get (_Unwind_Context *context,
                                       _Unwind_VRS_RegClass regclass, _uw regno,
                                       _Unwind_VRS_DataRepresentation representation, void *valuep)
{
    phase1_vrs *vrs = (phase1_vrs *) context;
    switch (regclass) {
        case _UVRSC_CORE:
            if (representation != _UVRSD_UINT32
                || regno > 15)
                return _UVRSR_FAILED;
            *(_uw *) valuep = vrs->core.r[regno];
            return _UVRSR_OK;
        case _UVRSC_VFP:
        case _UVRSC_WMMXD:
        case _UVRSC_WMMXC:
            return _UVRSR_NOT_IMPLEMENTED;
        case _UVRSC_PAC:
            *(_uw *) valuep = vrs->pseudo.pac;
            return _UVRSR_OK;
        default:
            return _UVRSR_FAILED;
    }
}
```

```
/* Execute the unwinding instructions described by UWS.  */
_Unwind_Reason_Code __gnu_unwind_execute (_Unwind_Context * context,
                                         _gnu_unwind_state
{
    _uw op;
    int set_pc;
    _uw reg;
    _uw sp;
    set_pc = 0;
    for (;;) {
        op = next_unwind_byte (uws);
        if (op == CODE_FINISH) {
            if (!set_pc)
            {
                _Unwind_VRS_Get (context, _UVRSC_CORE, R_LR, _UVRSD_UI
                                  &reg);
                _Unwind_VRS_Set (context, _UVRSC_CORE, R_PC, _UVRSD_UI
                                  &reg);
                set_pc = 1;
            }
            /* Drop out of the loop. */
            break;
        }
    }
}
```

```
_Unwind_VRS_Result _Unwind_VRS_Set (_Unwind_Context *context,
                                     _Unwind_VRS_RegClass regclass, _uw regno,
                                     _Unwind_VRS_DataRepresentation representation, void *valueep)
{
    phase1_vrs *vrs = (phase1_vrs *) context;
    switch (regclass) {
        case _UVRSC_CORE:
            if (representation != _UVRSD_UINT32
                || regno > 15)
                return _UVRSR_FAILED;
            vrs->core.r[regno] = *(_uw *) valueep;
            return _UVRSR_OK;
        case _UVRSC_VFP:
        case _UVRSC_WMMXD:
        case _UVRSC_WMMCX:
            return _UVRSR_NOT_IMPLEMENTED;
        case _UVRSC_PAC:
            vrs->pseudo.pac = *(_uw *) valueep;
            return _UVRSR_OK;
        default:
            return _UVRSR_FAILED;
    }
}
```

```
_Unwind_VRS_Result _Unwind_VRS_Set (_Unwind_Context *context,
_Unwind_VRS_RegClass regclass, _uw regno,
_Unwind_VRS_DataRepresentation representation, void *valuep)
{
phasel_vrs *vrs = (phasel_vrs *) context;
switch (regclass) {
    case _UVRSC_CORE:
        if (representation != _UVRSD_UINT32
            || regno > 15)
            return _UVRSR_FAILED;
        vrs->core.r[regno] = *(_uw *) valuep;
        return _UVRSR_OK;
    case _UVRSC_VFP:
    case _UVRSC_WMMXD:
    case _UVRSC_WMMXC:
        return _UVRSR_NOT_IMPLEMENTED;
    case _UVRSC_PAC:
        vrs->pseudo.pac = *(_uw *) valuep;
        return _UVRSR_OK;
    default:
        return _UVRSR_FAILED;
}
}
```



```
[[gnu::always_inline]] _Unwind_VRS_Result
My_Unwind_VRS_Set(
    _Unwind_Context* context,
    _Unwind_VRS_RegClass regclass,
    _uw regno,
    _Unwind_VRS_DataRepresentation representation,
    void* valuep)
{
    auto* vrs = reinterpret_cast<phasel_vrs*>(context);
    vrs->core.r[regno] = *(_uw*)valuep;
    return _UVRSR_OK;
}
```

```
_Unwind_VRS_Result _Unwind_VRS_Get (_Unwind_Context *context,
_Unwind_VRS_RegClass regclass, _uw regno,
_Unwind_VRS_DataRepresentation representation, void *valuep)
{
phasel_vrs *vrs = (phasel_vrs *) context;
switch (regclass) {
    case _UVRSC_CORE:
        if (representation != _UVRSD_UINT32
            || regno > 15)
            return _UVRSR_FAILED;
        *_uw * valuep = vrs->core.r[regno];
        return _UVRSR_OK;
    case _UVRSC_VFP:
    case _UVRSC_WMMXD:
    case _UVRSC_WMMXC:
        return _UVRSR_NOT_IMPLEMENTED;
    case _UVRSC_PAC:
        *_uw * valuep = vrs->pseudo.pac;
        return _UVRSR_OK;
    default:
        return _UVRSR_FAILED;
}
}
```

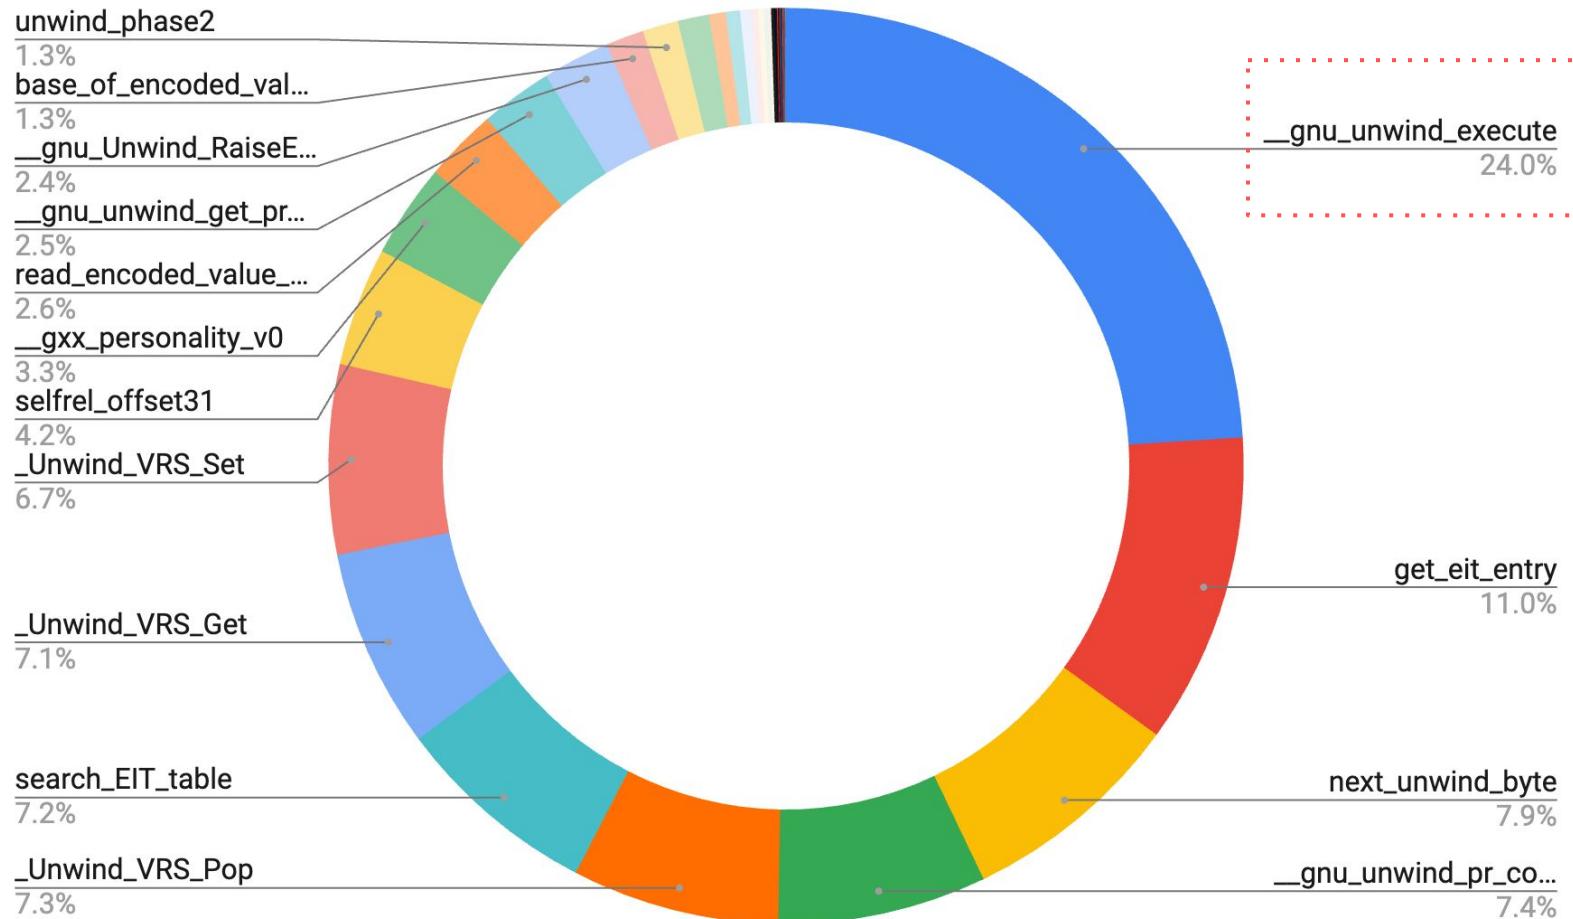


```
[[gnu::always_inline]] _Unwind_VRS_Result
My_Unwind_VRS_Get(
    _Unwind_Context* context,
    _Unwind_VRS_RegClass regclass,
    _uw regno,
    _Unwind_VRS_DataRepresentation representation,
    void* valuep)
{
    auto* vrs = reinterpret_cast<phasel_vrs*>(context);
    *_uw* valuep = vrs->core.r[regno];
    return _UVRSR_OK;
}
```

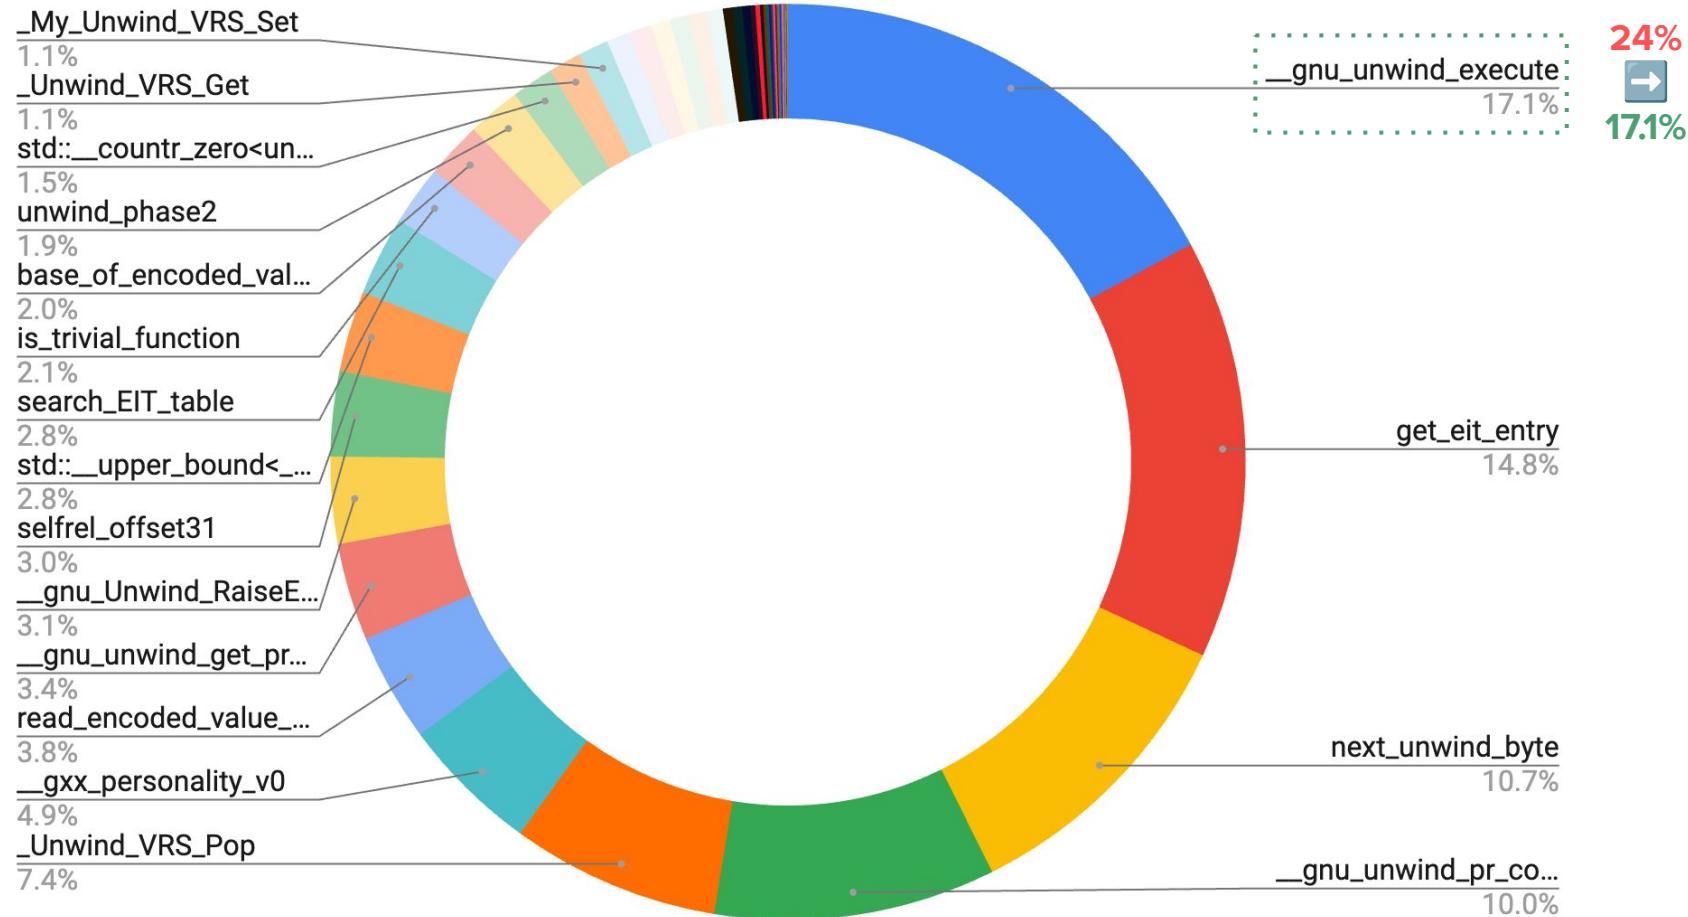
```
_uw8 next_unwind_byte(__gnu_unwind_state* uws)
{
    _uw8 b;
    if (uws->bytes_left == 0) {
        /* Load another word */
        if (uws->words_left == 0)
            return 0xB0; /* Nothing left. */
        uws->words_left--;
        uws->data = *(uws->next++);
        uws->bytes_left = 3;
    } else
        uws->bytes_left--;
    /* Extract the most significant byte. */
    b = (uws->data >> 24) & 0xff;
    uws->data <= 8;
    return b;
}
```

```
[ [ gnu::always_inline ] ] _uw8
next_unwind_byte(__gnu_unwind_state* uws)
{
    _uw8 b;
    if (uws->bytes_left == 0) {
        /* Load another word */
        if (uws->words_left == 0)
            return 0xB0; /* Nothing left. */
        uws->words_left--;
        uws->data = *(uws->next++);
        uws->bytes_left = 3;
    } else
        uws->bytes_left--;
    /* Extract the most significant byte. */
    b = (uws->data >> 24) & 0xff;
    uws->data <= 8;
    return b;
}
```

V2



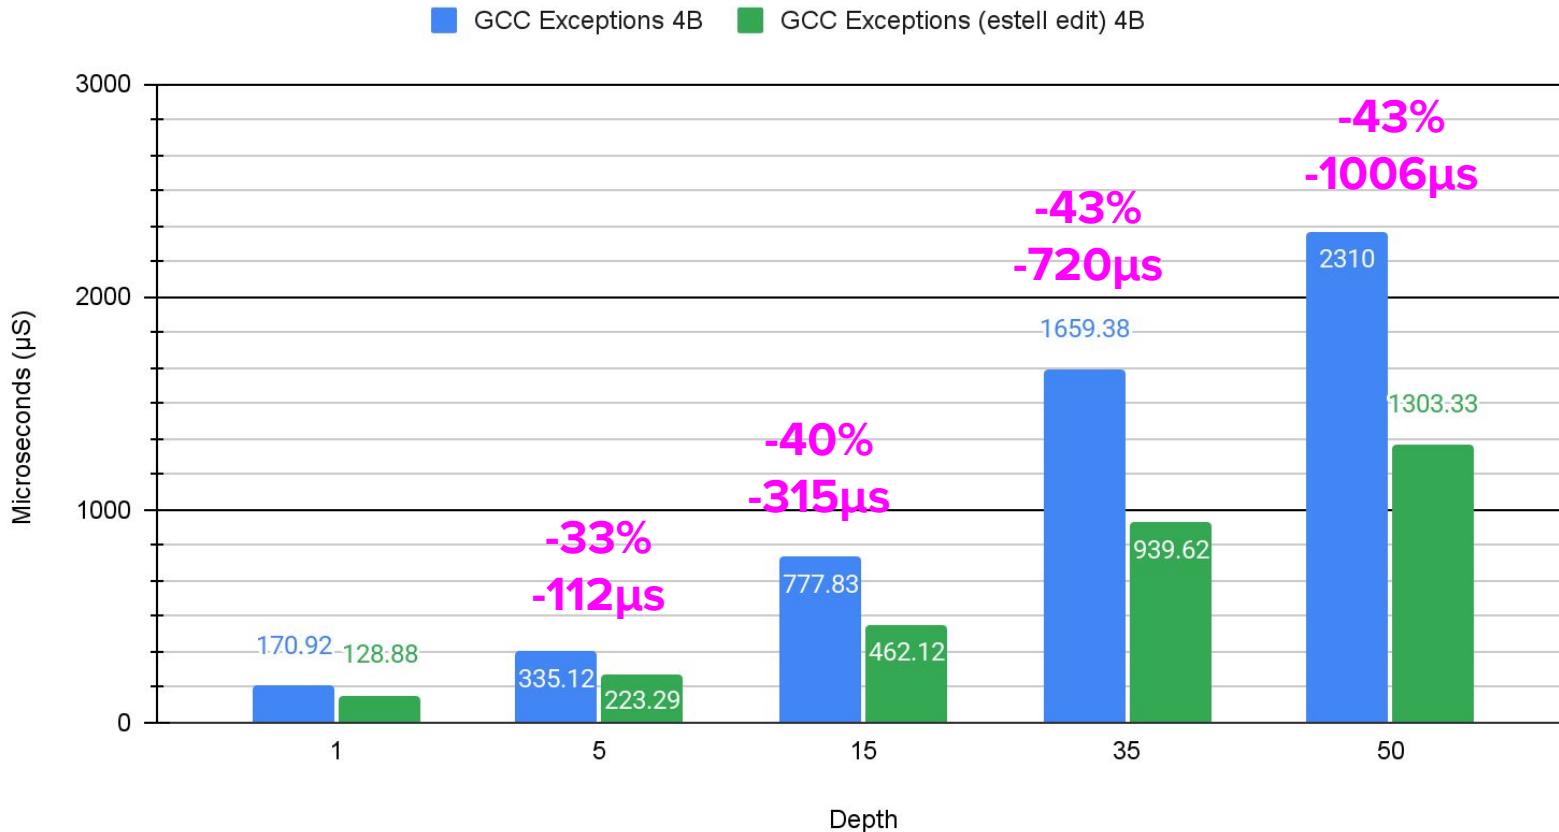
V3



Results!!



Exception Performance (Estell Edit)





Telling people about it!



October 2023



Jason Turner

So after that change the cycle count dropped down to ~270 cycles!

That's fascinating... I wonder if you could reduce the space you need to check by doing some sort of linear interpolation.

Ah, kinda like assuming "if the program counter is the middle of the binary then maybe the function's entry is also near the middle?

Kinda of. I just think there may be some structure here you can take advantage of. But maybe not.



Nearpoint Search Algorithm



Madeline Schneider

... well exceptions could never be faster than result types.

Oh yes, I don't expect to beat result types... But now that you say that
I'm going to try!

How is that even going to be possible?

"Consider this madeline..."

Depth	Cycles
	std::expected<T,E>
6	556 cycles*
1	92 cycles*

Did some napkin math and got ~65 cycles

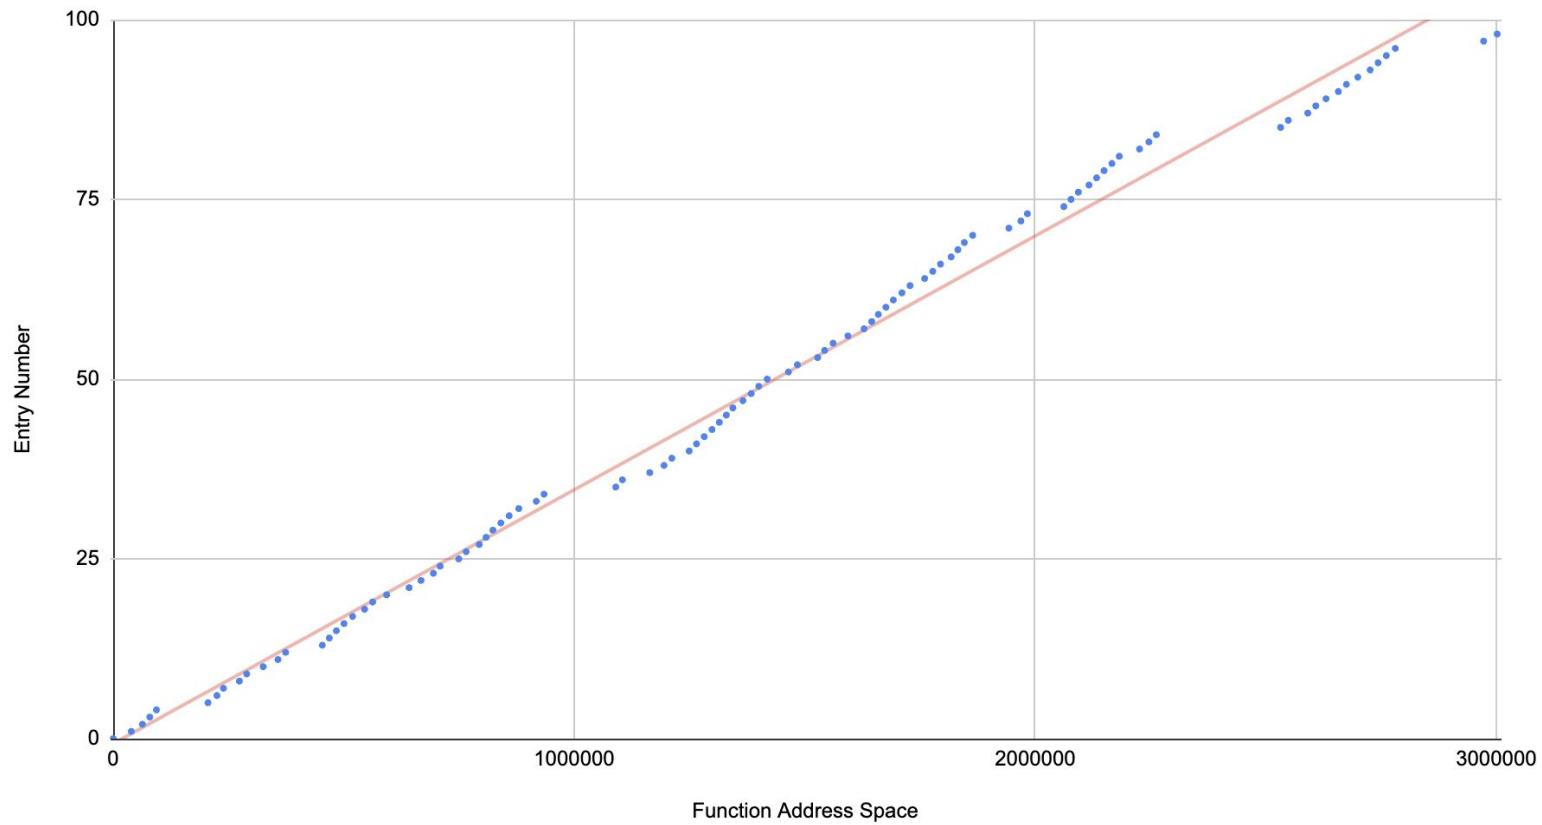
Search time: 270 cycles  27 cycles

Nearpoint Search Algorithm Requirements

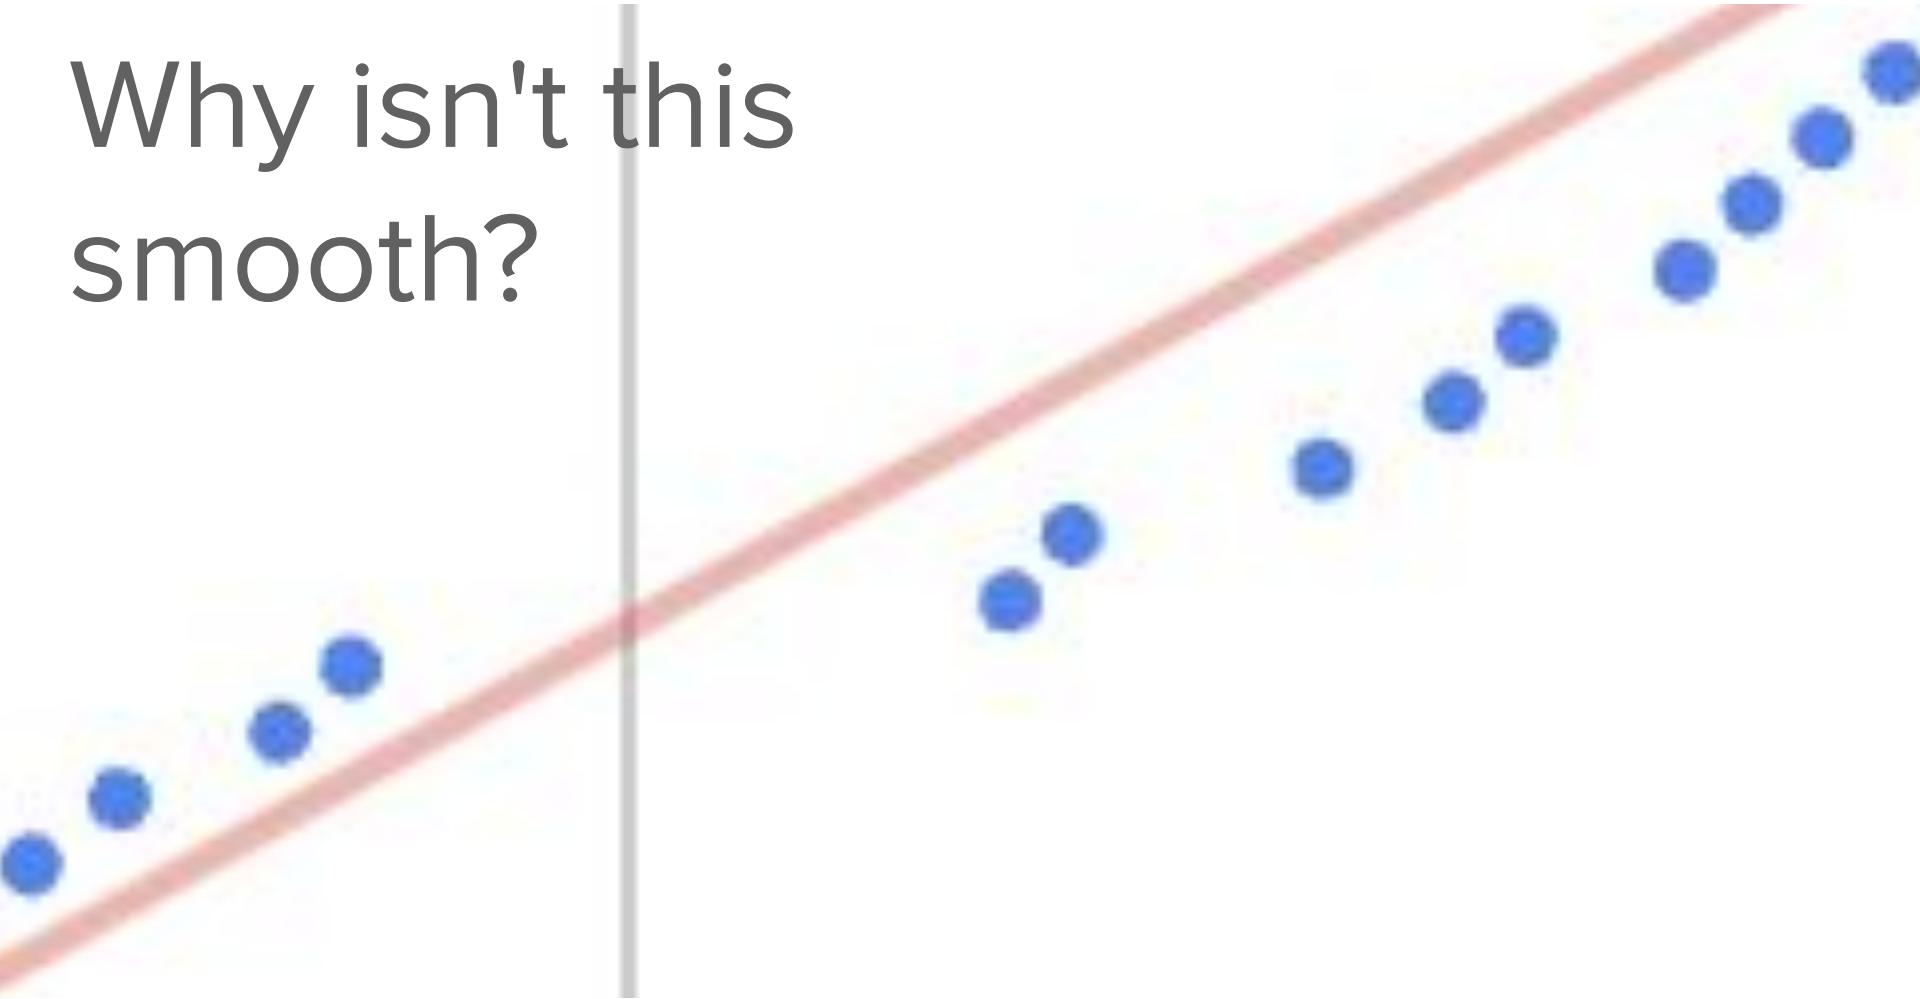
-  Order of magnitude faster than binary search
-  Try and preserve exception space saving

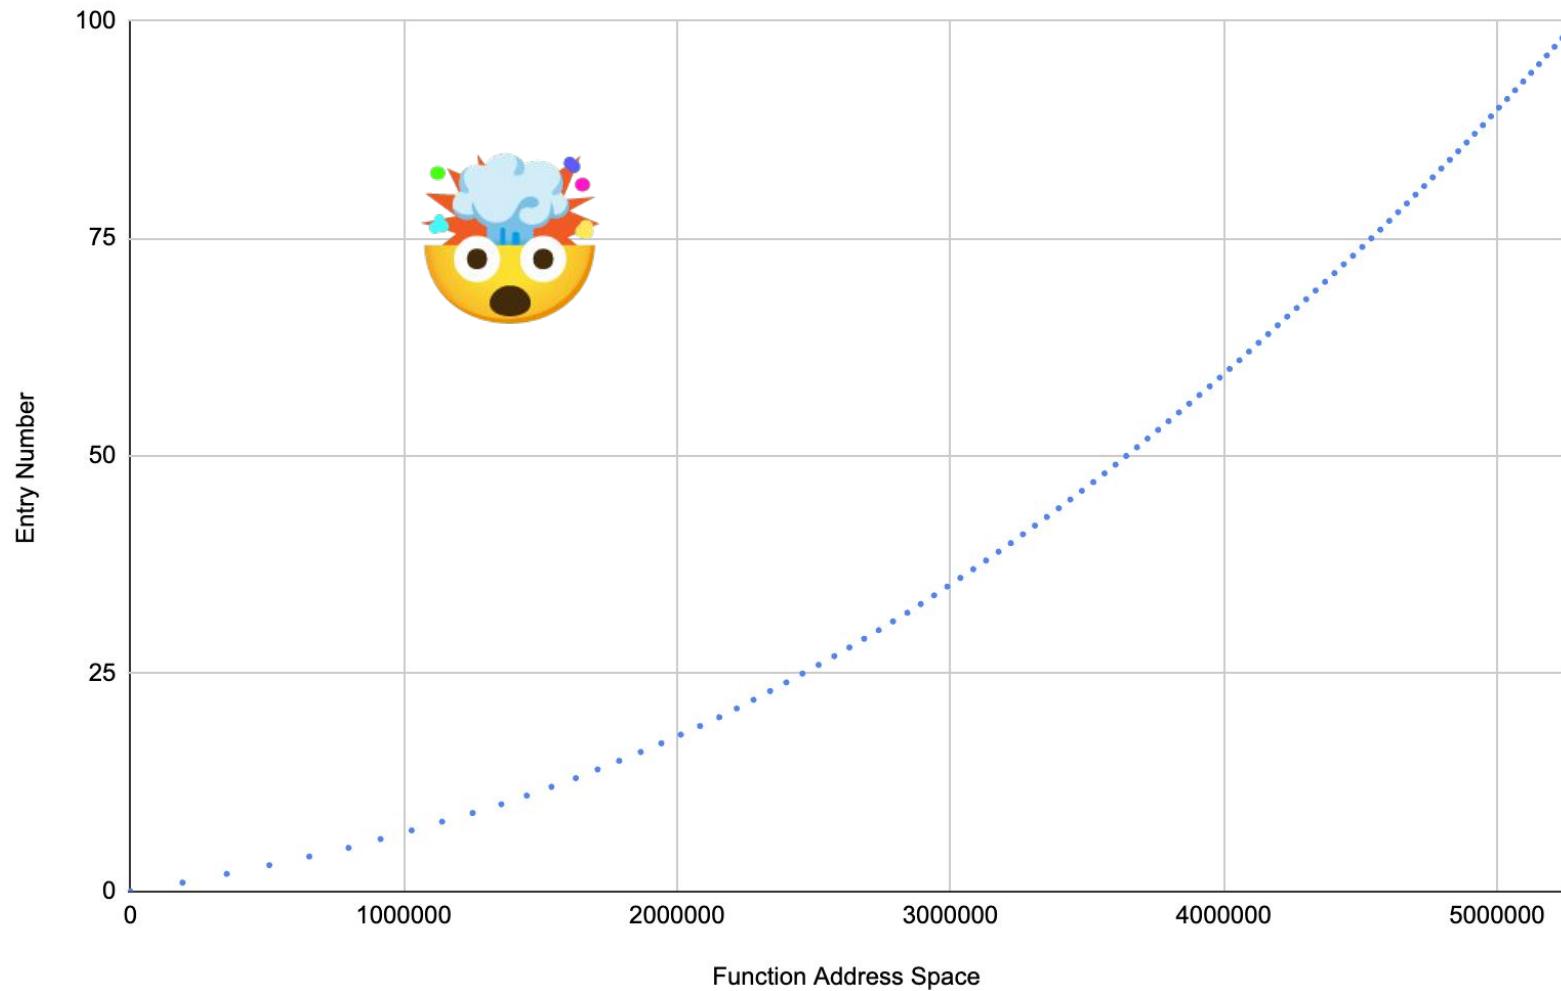
December 2023

On a flight home
from a business
trip...



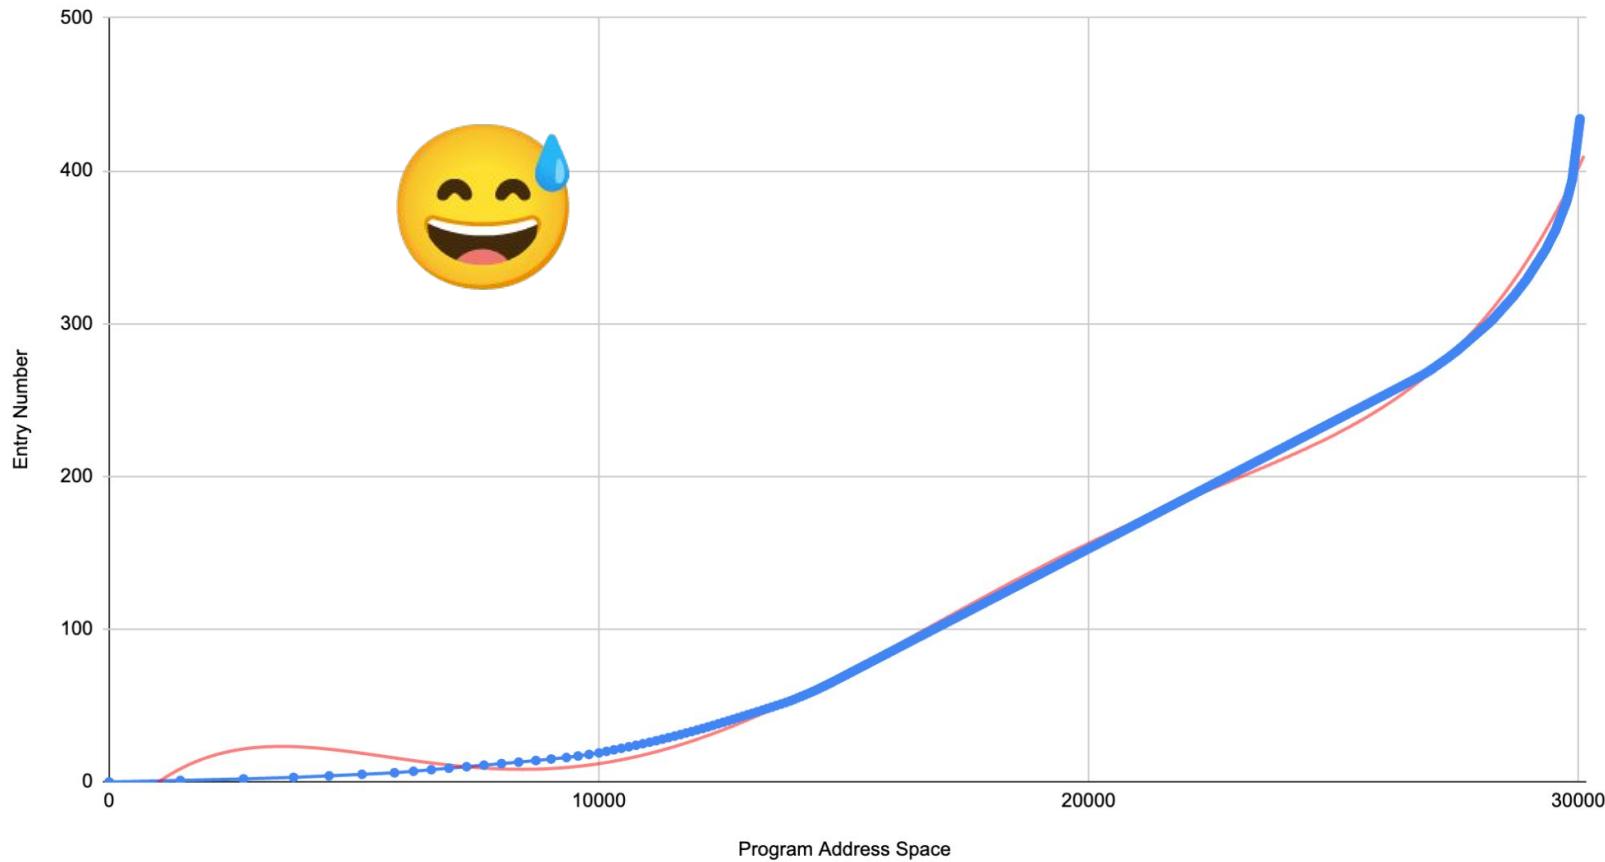
Why isn't this
smooth?





Index Entry and Decimal Addr

● Index Entry -30.4 + 0.0386x + -9.32E-06x² + 8.7E-10x³ + -3.25E-14x⁴ + 4.32E-19x⁵ R² = 0.996



ISO C++ Meeting

Japan 

Host: Woven By Toyota

March 2024



Ben Craig

Oh come on, floating point arithmetic isn't that bad. It performs quite well actually!

The problem is the space budget for embedded devices **without** an **FPU**. The floating point operations would **exceed our ROM**.

Hmmm, alright, then maybe desktop/servers environments could opt for this?

There is a problem there too. In the kernel, using floating point operations **corrupts user mode state**.

Nearpoint Search Algorithm Requirements

- ✓ Order of magnitude faster than binary search
- ✓ Try and preserve exception space saving
- ✓ Integer math ONLY ✨ NEW ✨



Robin Rowe

Yammers on about exceptions on embedded...



Robin Rowe

Yammers on about exceptions on embedded...

Huh... You know. I agree with your premise, but I don't know if I agree with your conclusion.



Robin Rowe

... 5 hours of discussion ...



Robin Rowe

... 5 hours of discussion ...

TrapC: Memory Safe C Programming with No UB

Pr. Robin Rowe <Robin.Rowe@trasec.com>

Number: n3423 - TrapC

Date: 2025-01-07

For: ISO WG14 C Committee meeting at Graz University of Technology, Graz, Austria,
February 24 to February 28, 2025



ACCU
conference
2025

Keynote

**C++ Exceptions are
Code Compression**

Months go by...

ISO C++ Meeting

St. Louis 

June 2024



Herb Sutter

... Shows Herb my current runtime results ...

Very interesting. Now, how do you handle catching exceptions?

My runtime doesn't support polymorphic catch yet. I can only catch types **directly**. I plan to write my own "**dynamic-cast-like**" algorithm for this.

Make sure your dynamic-cast-like implementation has deterministic performance. It should have an **upper bound** on **space** and **time** cost.

Chapter 3. ⚡

Estell Exception Runtime

Inefficiencies spotted

- Searches index **2x** per frame
- dynamic_cast tree traversal
- Multilang support
- TODOs in source

```
// FILE: gcc/libstdc++-v3/libsupc++/unwind-cxx.h
bool __is_gxx_exception_class(_Unwind_Exception_Class c)
{
    // TODO: Take advantage of the fact that C will always be
    //        word aligned.

    return c[0] == 'G'
        && c[1] == 'N'
        && c[2] == 'U'
        && c[3] == 'C'
        && c[4] == 'C'
        && c[5] == '+'
        && c[6] == '+'
        && (c[7] == '\0' || c[7] == '\x01');
}

// FILE: gcc/libstdc++-v3/libsupc++/eh_arm.cc
__cxa_type_match_result __cxa_type_match() {
    // Check if exception came from GNU C++
    bool foreign_exception
        = !__is_gxx_exception_class(ue_header->exception_class);
}
```

A TODO almost old enough to drink



20 years ago Makefile.in: Set and use UNWI... | ⚙️ | ...

```
239     static inline bool
240     __is_gxx_exception_class(_Unwind_Exception_Class c)
241     {
242         // TODO: Take advantage of the fact that c will always be word aligned
243         return c[0] == 'G'
244             && c[1] == 'N'
245             && c[2] == 'U'
246             && c[3] == 'C'
247             && c[4] == 'C'
248             && c[5] == '+'
249             && c[6] == '+'
250             && (c[7] == '\0' || c[7] == '\x01');
251     }
252
253     // Only checks for primary or dependent, but not that it is a C++ except
254     // all.
255     static inline bool
256     __is_dependent_exception(_Unwind_Exception_Class c)
257     {
258         return c[7] == '\x01';

```

17 years ago [multiple changes] | ⚙️

Optimization #0

Rewrote from scratch in C++

Optimization #1

Single Phase Exceptions

Single Phase Exceptions

Two Phase Exceptions:

1. **Phase 1:** Determine if a suitable catch handler exists. If non exist call `std::terminate`
2. **Phase 2:** Otherwise, unwind stack, perform cleanup, and enter catch handler.



Single Phase

Single Phase Exceptions

Two Phase Exceptions:

1. **Phase 1:** Determine if a suitable catch handler exists. If non exist call `std::terminate`
2. **Phase 2:** Otherwise, unwind stack, perform cleanup, and enter catch handler.



Single Phase

One Phase Exceptions:

1. **Phase 1:** Unwind stack and perform cleanup until you reach a suitable catch handler or a point of termination (noexcept, main, thread)

Single Phase Exceptions

Consequences?

1. With two phases, if a suitable catch block is not found, your terminate handler can do a **core dump**.
2. Applications may depend on destructors only being called if a handler exists.



Single Phase

Positives?

Single Phase Exceptions

Consequences?

1. With two phases, if a suitable catch block is not found, your terminate handler can do a **core dump**.
2. Applications may depend on destructors only being called if a handler exists.



Single Phase

Positives?

1. Faster, without using additional storage
2. **Same behaviour** as using `std::expected`

Optimization #2

TLS Control Block

thread local storage

Exception Header

(all the things the exception runtime needs to remember during unwinding)

Thrown Object

```
struct exception_control_block
{
    cortex_m_cpu cpu{};

    flattened_hierarchy<4> type_info{};

    std::size_t catch_type_offset = 0;

    void* thrown_object;

    destructor_t* destructor = nullptr;

    exception_cache cache{};

};
```

Returned address

Control Block Size:
128B per thread

assuming max hierarchy size of 4

vs

128B per exception

```
thread_local exception_control_block
control block{};
```

Optimization #3

Use PMR Allocation

Polymorphic memory resource

```
#include <memory_resource>

/***
 * @brief Set the global exception allocator function
 *
 * @param p_allocator - the global exception memory resource. The lifetime of
 * the memory resource should be static.
 */
void set_exception_allocator(std::pmr::memory_resource* p_allocator) noexcept;

/***
 * @brief Get the global exception allocator function
 *
 * @returns the global exception memory allocator
 */
std::pmr::memory_resource* get_exception_allocator() noexcept;
```

Exception Header

```
struct exception_header
{
    std::pmr::memory_resource* allocator = nullptr;
    std::size_t size = 0uz;
    alignas(std::max_align_t) std::byte thrown_object;
};
```



```
void* __wrap___cxa_allocate_exception(std::size_t p_thrown_size)
{
    auto const total_memory = sizeof(exception_header) + p_thrown_size;

    auto* const resource = get_exception_allocator();
    auto* allocated_memory = resource->allocate(total_memory);
    auto* header = static_cast<exception_header*>(allocated_memory);

    header->size = total_memory;
    header->allocator = resource;

    return & header->data;
}
```

Optimization #4

Skipping the first search

```
void __wrap___cxa_throw (void* p_thrown_exception ,
                        std::type_info* p_type_info ,
                        ke::destructor_t p_destructor) noexcept (false)
{
    auto& control_block = ke::control_block; // Load TLS reference
    control_block.thrown_object = p_thrown_exception;
    control_block.destructor = p_destructor;
    ke::capture_cpu_core (control_block.cpu);
    ke::flatten_rtti (p_thrown_exception , control_block.type_info , p_type_info);
    // Use the expanded unwind instructions to quickly unwind cxa_throw.
    ke::unwind_frame (ke::cxa_throw_unwind_instructions , control_block.cpu);
    // Raise exception returns when an error or call to terminate has been found
    ke::raise_exception (control_block);
    // If raise exception returns it means it failed and we should terminate
    control_block.cache.state (ke::runtime_state::handled_state);
    std::terminate ();
}
```

```
void __wrap___cxa_throw (void* p_thrown_exception ,
                        std::type_info* p_type_info ,
                        ke::destructor_t p_destructor) noexcept (false)
{
    auto& control_block = ke::control_block; // Load TLS reference
    control_block.thrown_object = p_thrown_exception;
    control_block.destructor = p_destructor;
    ke::capture_cpu_core (control_block.cpu);
    ke::flatten_rtti (p_thrown_exception , control_block.type_info , p_type_info);
    // Use the expanded unwind instructions to quickly unwind cxa_throw.
    ke::unwind_frame (ke::cxa_throw_unwind_instructions , control_block.cpu);
    // Raise exception returns when an error or call to terminate has been found
    ke::raise_exception (control_block);
    // If raise exception returns it means it failed and we should terminate
    control_block.cache.state (ke::runtime_state::handled_state);
    std::terminate ();
}
```

```
template<typename F>
instructions_t cache(F* p_function_to_be_cached)
{
    auto const& entry = ke::get_index_entry(function_address);
    if (entry.has_inlined_personality()) {
        return personality_to_unwind_instructions(&entry.personality_offset);
    } else {
        return spare_instruction();
    }
}

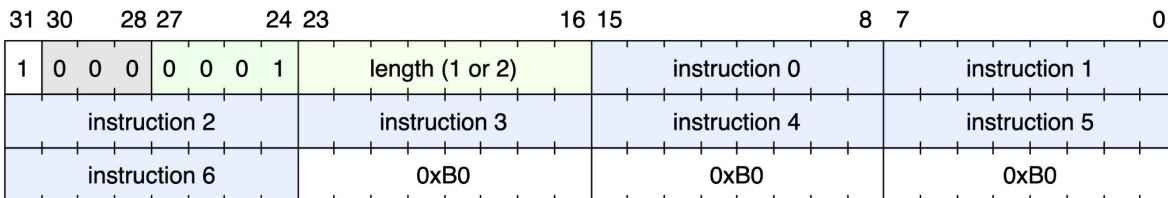
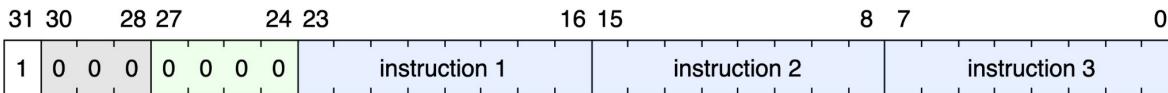
auto const cxa_throw_unwind_instructions = cache(&__wrap_cxa_throw);
auto const cxa_rethrow_unwind_instructions = cache(&__wrap_cxa_rethrow);
```

Computed at **static initialization** time.

Optimization #5

Prearrange unwind instructions

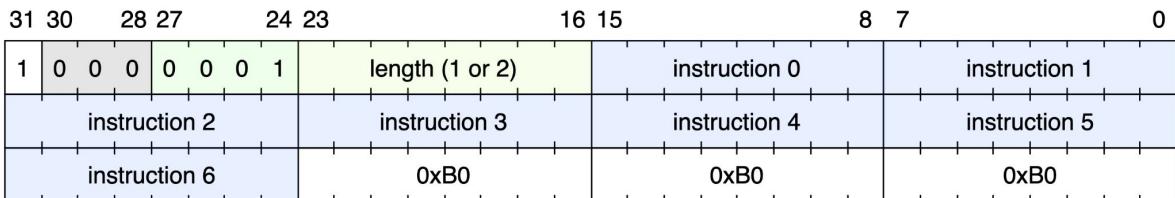
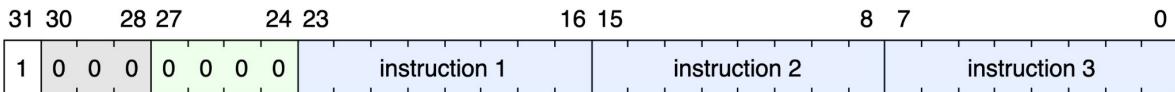
ARM short & long unwind formats



```
case runtime_state::unwind_frame: {
    auto const instructions =
        create_instructions_from_entry(p_exception_object);
    unwind_frame(instructions, p_exception_object.cpu);
    p_exception_object.cache.state(runtime_state::get_next_frame);
}
```

```
[[gnu::always_inline]] _uw8
next_unwind_byte(__gnu_unwind_state* uws)
{
    _uw8 b;
    if (uws->bytes_left == 0) {
        /* Load another word */
        if (uws->words_left == 0)
            return 0xB0; /* Nothing left. */
        uws->words_left--;
        uws->data = *(uws->next++);
        uws->bytes_left = 3;
    } else
        uws->bytes_left--;
    /* Extract the most significant byte. */
    b = (uws->data >> 24) & 0xff;
    uws->data <= 8;
    return b;
}
```

ARM short & long unwind formats



Optimization #6

JIT unwinding

just in time

```
void unwind_frame(instructions_t const& p_instructions,  
                  cortex_m_cpu& p_cpu)  
{  
    static constexpr std::array<void*, 256> jump_table{  
        &&vsp_add_0, // [0]  
        &&vsp_add_1, // [1]  
        &&vsp_add_2, // [2]  
        &&vsp_add_3, // [3]  
        &&vsp_add_4, // [4]  
        &&vsp_add_5, // [5]  
        &&vsp_add_6, // [6]  
        &&vsp_add_7, // [7]  
        &&vsp_add_8, // [8]  
        &&vsp_add_9, // [9]  
        &&vsp_add_10, // [10]  
        &&vsp_add_11, // [11]  
        ....  
    };
```

ARM Unwind
instructions are
8-bits in size.

$$\text{Cost} = 256 * 4\text{B} = \mathbf{1024\text{ B}}$$

GCC label address extension

&&vsp_add_11, // [11]

```
auto instruction_ptr = p_instructions.data.begin();
auto const* sp_ptr = *p_cpu.sp;

while (true) {
    auto const* instruction_handler = jump_table[*instruction_ptr++];
    goto *instruction_handler;

finish_unwind:
    p_cpu.sp = sp_ptr;
    [[likely]] if (move_lr_to_pc) {
        p_cpu.pc = p_cpu.lr;
    }
    break;

reserved_or_spare_thus_terminate:
    std::terminate();
    break;
```

```
//  
+-----+  
// | Sequentially Pop Registers + LR  
|  
//  
+-----+  
pop_off_stack_r4_to_r11_and_lr :  
    sp_ptr = pop_register_range<7, pop_lr::do_it>(sp_ptr, p_cpu);  
    continue;  
pop_off_stack_r4_to_r10_and_lr :  
    sp_ptr = pop_register_range<6, pop_lr::do_it>(sp_ptr, p_cpu);  
    continue;  
pop_off_stack_r4_to_r9_and_lr :  
    sp_ptr = pop_register_range<5, pop_lr::do_it>(sp_ptr, p_cpu);  
    continue;  
pop off stack r4 to r8 and lr :  
    sp_ptr = pop_register_range<4, pop_lr::do_it>(sp_ptr, p_cpu);  
    continue;
```

```
template<size_t PopCount, pop_lr PopLinkRegister = pop_lr::skip>
[[nodiscard("The result of this function MUST be assigned to reg SP")]]
std::uint32_t const* pop_register_range(
    std::uint32_t const* sp_ptr,
    cortex_m_cpu& p_cpu) {
    auto* r4_pointer = &p_cpu.r4.data;
    static_assert(PopCount <= 7, "Pop Count cannot be above 7");
    if constexpr (PopCount == 0) {
        *r4_pointer = *sp_ptr;
    } else {
        for (std::size_t i = 0; i < PopCount + 1; i++) {
            r4_pointer[i] = sp_ptr[i];
        }
    }
    if constexpr (PopLinkRegister == pop_lr::do_it) {
        p_cpu.lr = sp_ptr[PopCount + 1];
    }
    return sp_ptr + PopCount + 1 +
        unsigned{ PopLinkRegister == pop_lr::do_it };
}
```

The only runtime
conditional

```
template<size_t PopCount, pop_lr PopLinkRegister = pop_lr::skip>
[[nodiscard("The result of this function MUST be assigned to reg SP")]]
std::uint32_t const* pop_register_range(
    std::uint32_t const* sp_ptr,
    cortex_m_cpu& p_cpu) {
    auto* r4_pointer = &p_cpu.r4.data;
    static_assert(PopCount <= 7, "Pop Count cannot be above 7");
    if constexpr (PopCount == 0) {
        *r4_pointer = *sp_ptr;
    } else {
        for (std::size_t i = 0; i < PopCount + 1; i++) {
            r4_pointer[i] = sp_ptr[i];
        }
    }
    if constexpr (PopLinkRegister == pop_lr::do_it) {
        p_cpu.lr = sp_ptr[PopCount + 1];
    }
    return sp_ptr + PopCount + 1 +
        unsigned{ PopLinkRegister == pop_lr::do_it };
}
```

The only runtime
conditional

```

/Users/kammce/.conan2/p/b/libhaae3c10bd31de3/b/src/arm_cortex/estell/exception.o
823     if constexpr (PopCount == 0) {
824         *r4_pointer = *sp_ptr;
825     } else {
> 827         for (std::size_t i = 0; i < PopCount + 1; i++) {
828             r4_pointer[i] = sp_ptr[i];
829         }
830     }
831
832     if constexpr (PopLinkRegister == pop_lr::do_it) {
833         p_cpu.lr = sp_ptr[PopCount + 1];

```

```

0x80005e6 <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+278>
0x80005ea <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+282>
0x80005ec <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+284>
0x80005f0 <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+288>
0x80005f2 <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+290>
>0x80005f4 <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+292>
0x80005f6 <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+294>
0x80005f8 <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+296>
0x80005fa <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+298>
0x80005fc <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+300>
0x80005fe <_ZN2ke12unwind_frameERKNS_14instructions_tERNS_12cortex_m_cpuE+302>

```

remote Thread 1 (src) In: ke::unwind_frame

```

(gdb) nexti
ke::register_t::operator* (this=0x200005fc <ke::control_block+52>)
  at /Users/kammce/.conan2/p/b/libhaae3c10bd31de3/b/src/arm_cortex/estell/intern
ke::unwind_frame (p_instructions=..., p_cpu=...) at /Users/kammce/.conan2/p/b/lib
ke::pop_register_range<7u, (ke::pop_lr)1> (sp_ptr=0x2000233c, p_cpu=...)
  at /Users/kammce/.conan2/p/b/libhaae3c10bd31de3/b/src/arm_cortex/estell/excep
(gdb) ■

```

<pre> 80005f4: 6823 ldr r3, [r4, #0] 80005f6: 4610 mov r0, r2 80005f8: 610b str r3, [r1, #16] 80005fa: 6863 ldr r3, [r4, #4] 80005fc: 3424 adds r4, #36 @ 0x24 80005fe: 614b str r3, [r1, #20] 8000600: f854 3c1c ldr.w r3, [r4, #-28] 8000604: 618b str r3, [r1, #24] 8000606: f854 3c18 ldr.w r3, [r4, #-24] 800060a: 61cb str r3, [r1, #28] 800060c: f854 3c14 ldr.w r3, [r4, #-20] 8000610: 620b str r3, [r1, #32] 8000612: f854 3c10 ldr.w r3, [r4, #-16] 8000616: 624b str r3, [r1, #36] @ 0x24 8000618: f854 3c0c ldr.w r3, [r4, #-12] 800061c: 628b str r3, [r1, #40] @ 0x28 800061e: f854 3c08 ldr.w r3, [r4, #-8] 8000622: 62cb str r3, [r1, #44] @ 0x2c 8000624: f854 3c04 ldr.w r3, [r4, #-4] 8000628: 638b str r3, [r1, #56] @ 0x38 800062a: e755 b.n 80004d8 </pre>	<pre> <ke::unwind_frame(ke::instructions_t const&, ke::cortex_m_cpu&)+0x8> </pre>
--	--

Optimization #7

Flatten RTTI hierarchy

run time type information

Possible implementation

```
template<class P, class From, class To = remove_pointer_t<P>>
To *dynamic_cast(From *p) {
    if constexpr (is_same_v<From, To>) {
        return p;
    } else if constexpr (is_base_of_v<To, From>) {
        return (To *)(p);
    } else if constexpr (is_void_v<To>) {
        return truly_dynamic_to_mdo(p);
    } else if constexpr (is_base_of_v<From, To>) {
        return truly_dynamic_from_base_to_derived<To>(p);
    } else {
        return truly_dynamic_between_unrelated_classes<To>(p);
    }
}
```

Technically we should also verify that To is an *accessible* base of From in the current lexical scope...

59



dynamic_cast
From Scratch

CppCon.org

RTTI for classes hierarchies

Class type with no inheritance

```
class std_exception_rtti :  
public __class_type_info {  
    const char* name;  
};
```

8 B

Class type with one parent

```
class std_logic_error_rtti :  
public __si_class_type_info {  
    const char* name;  
    std::type_info* parent;  
};
```

12 B

Class type with multiple parents

```
class multi_inherit_error_rtti :  
public __vmi_class_type_info {  
    const char* name;  
    long flags;  
    size_t length;  
    type_info_and_offset parents[];  
};
```

20 B + (length * 8B)

$$\text{Cost} = 8\text{B} + 12\text{B} + 12\text{B} = \mathbf{32\text{ B}}$$

RTTI for classes hierarchies

Ignoring the size of strings
Ignoring the size of base classes



std::exception

std::logic_error

std::out_of_range

```
class std_exception_rtti :  
public __class_type_info {  
    const char* name;  
};
```

8B

```
class std_logic_error_rtti :  
public __si_class_type_info {  
    const char* name;  
    std::type_info* parent;  
};
```

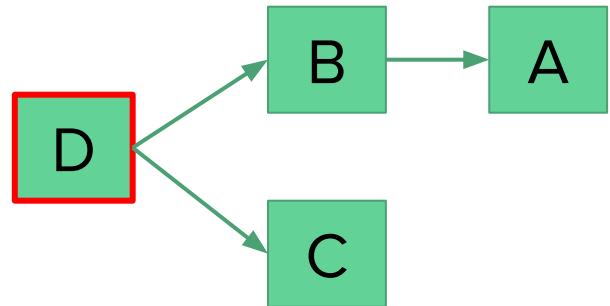
12B

```
class std_out_of_range_rtti :  
public __si_class_type_info {  
    const char* name;  
    std::type_info* parent;  
};
```

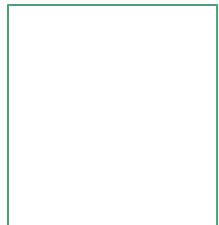
12B

`std::fixed_vector<void*, 7>`

Flattening Algorithm



`size`



`capacity`



`evaluator`



```
std::fixed_vector<void*,7>
```

Flattening Algorithm

```
class multi_inherit_error_rtti :  
public __vmi_class_type_info {  
    const char* name;  
    long flags;  
    size_t length;  
    type_info_and_offset parents[];  
};
```

size

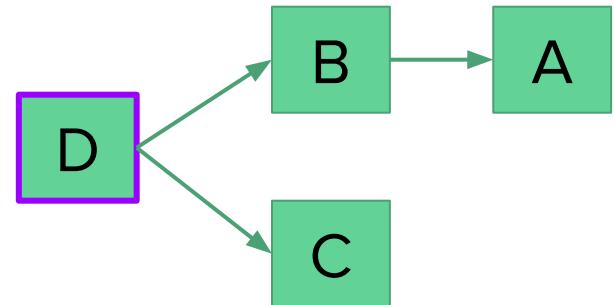


&D



evaluator

```
for(parent : parents) { types.push_back(parent); }
```

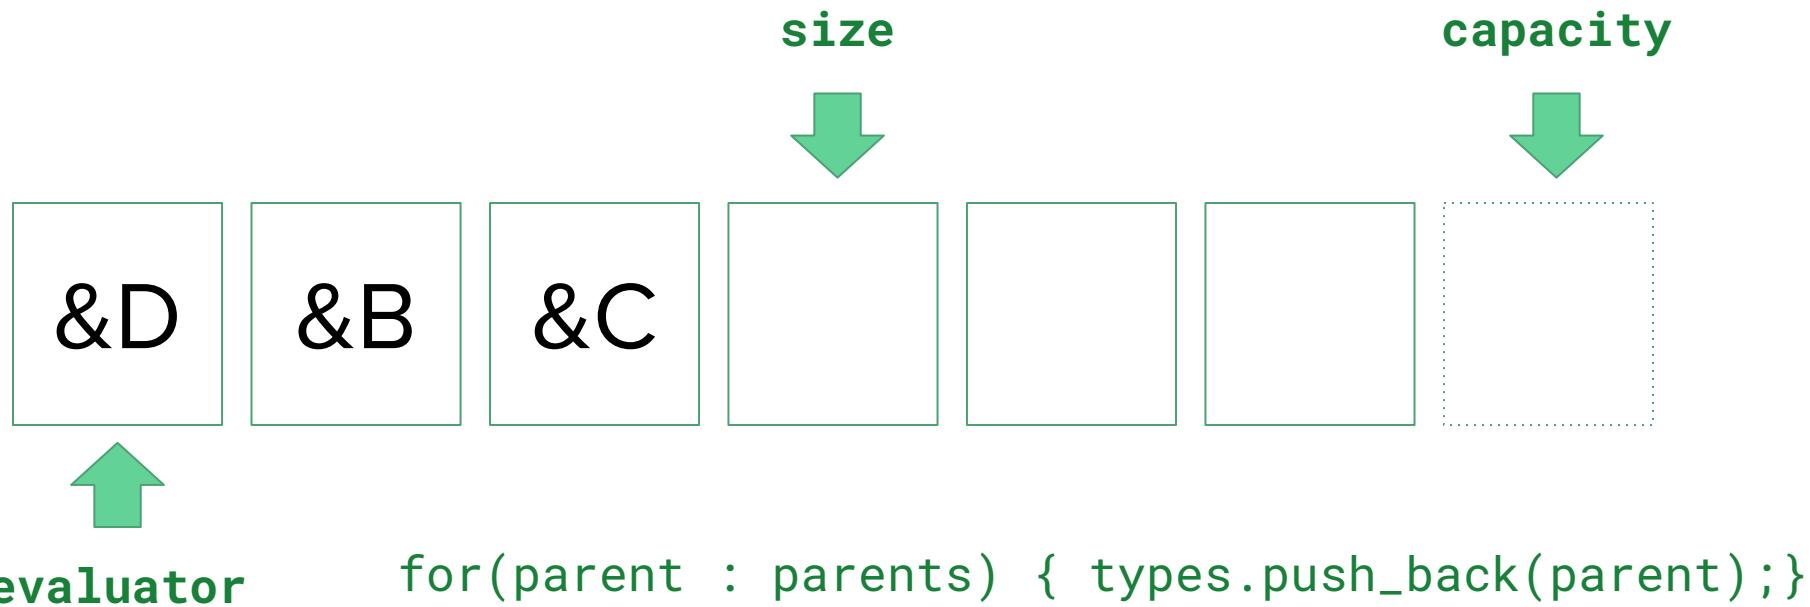


capacity



```
std::fixed_vector<void*, 7>
```

Flattening Algorithm

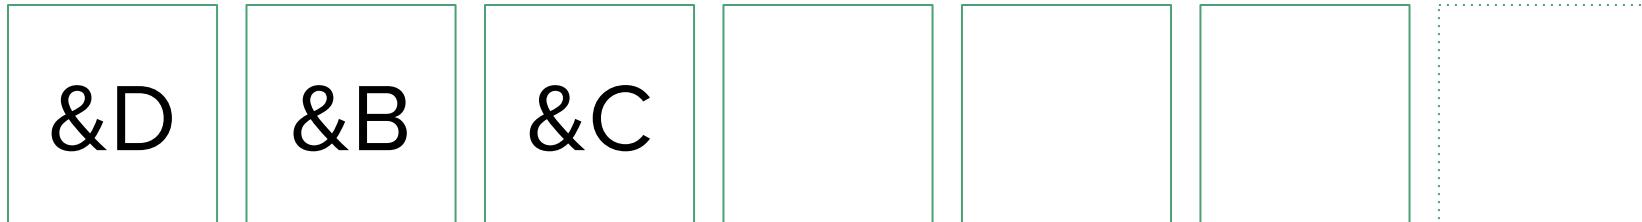


```
std::fixed_vector<void*,7>
```

Flattening Algorithm

```
class std_logic_error_rtti :  
public __si_class_type_info {  
    const char* name;  
    std::type_info* parent;  
};
```

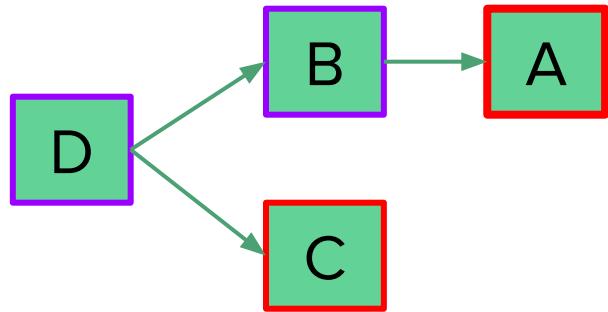
size



capacity



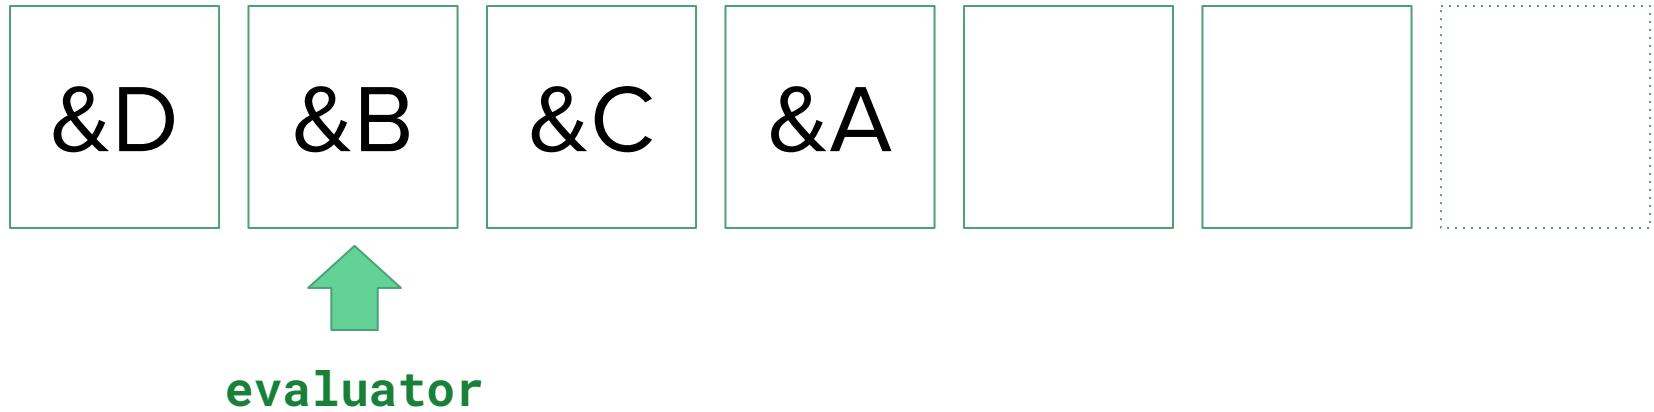
evaluator



```
std::fixed_vector<void*,7>
```

Flattening Algorithm

```
class std_logic_error_rtti :  
public __si_class_type_info {  
    const char* name;  
    std::type_info* parent;  
};
```



`std::fixed_vector<void*, 7>`

Flattening Algorithm

```
class std_exception_rtti :  
public __class_type_info {  
    const char* name;  
};
```

&D

&B

&C

&A

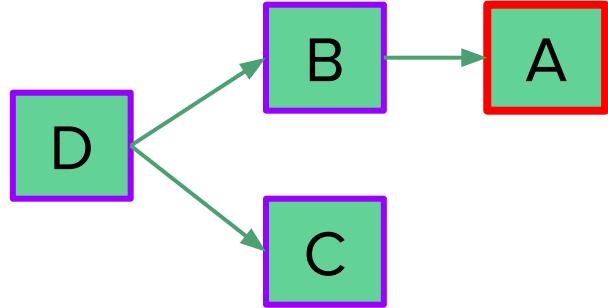


evaluator

size



capacity



`std::fixed_vector<void*, 7>`

Flattening Algorithm

```
class std_exception_rtti :  
public __class_type_info {  
    const char* name;  
};
```

&D

&B

&C

&A

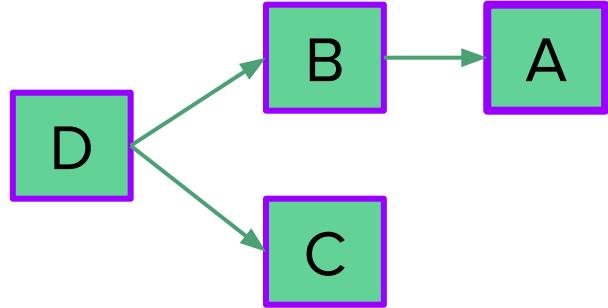


evaluator

size

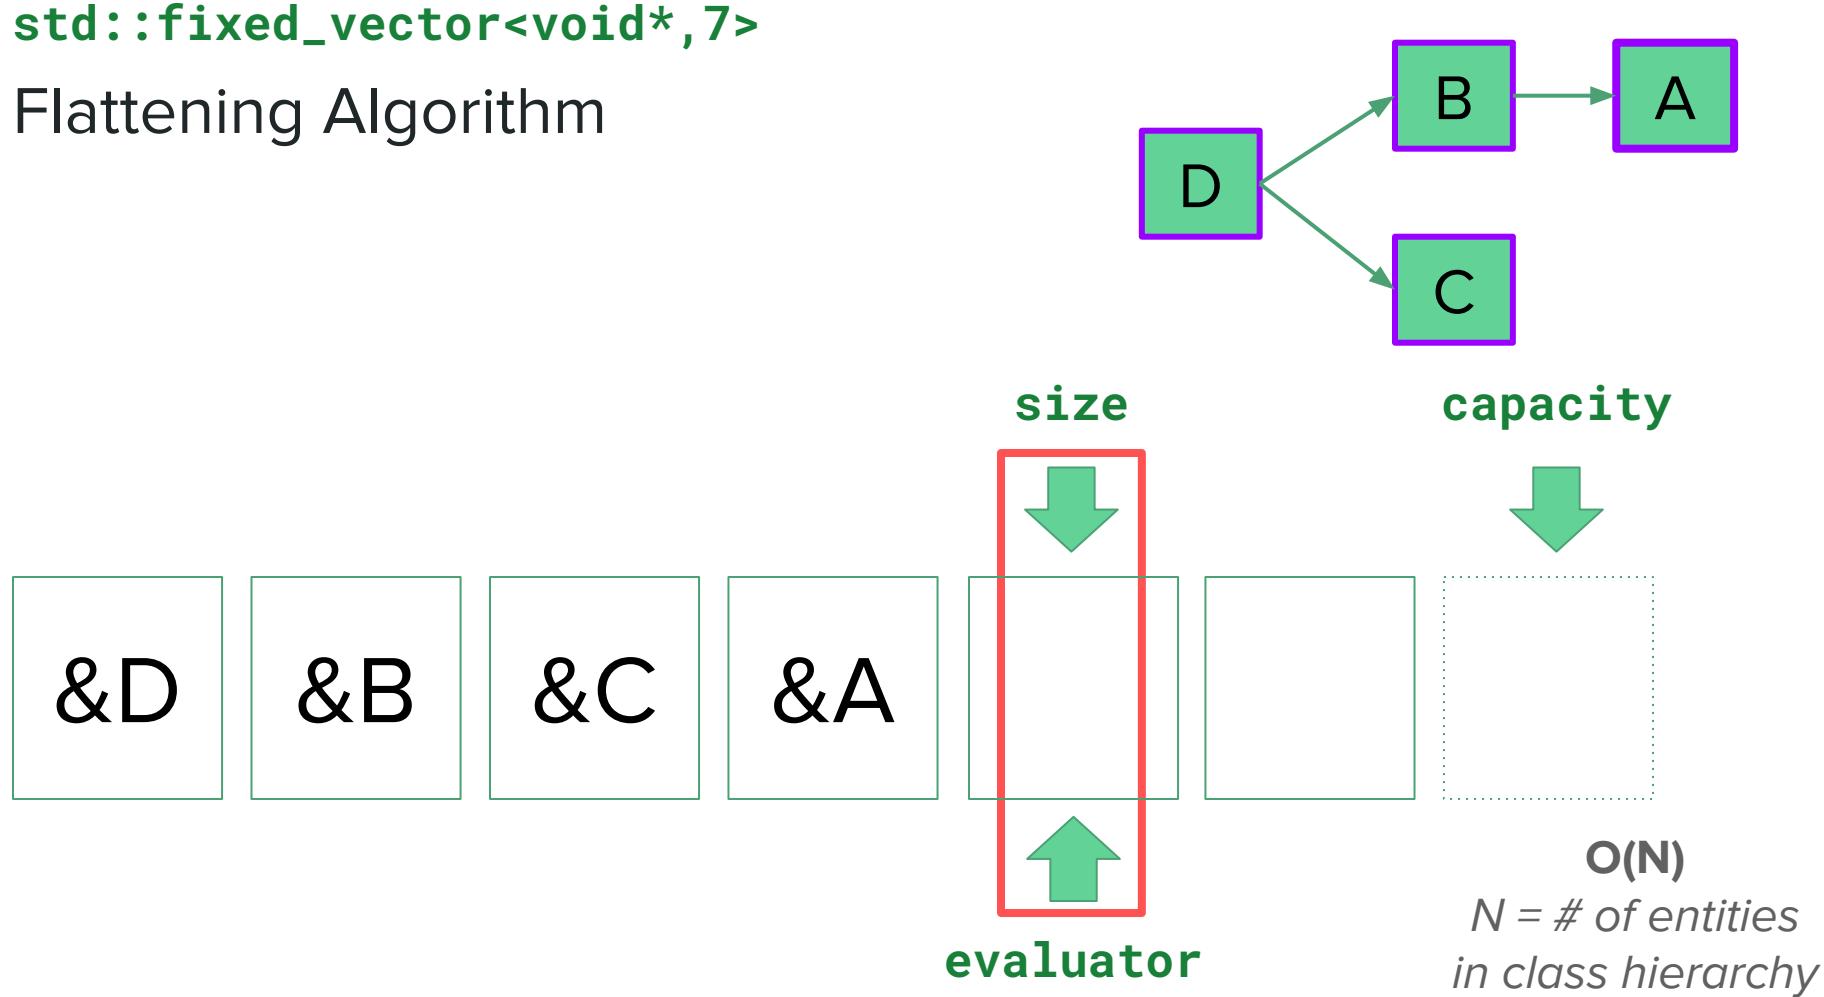


capacity

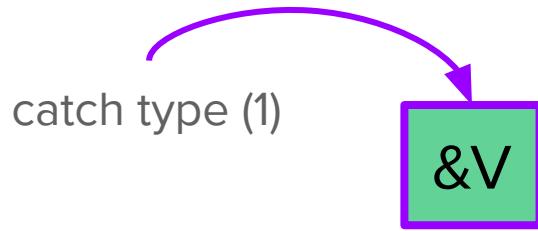


```
std::fixed_vector<void*,7>
```

Flattening Algorithm



Dynamic_cast Algorithm



```
try {  
    // stuff...  
} catch (V const&) {  
} catch (C const&) {  
}
```

size



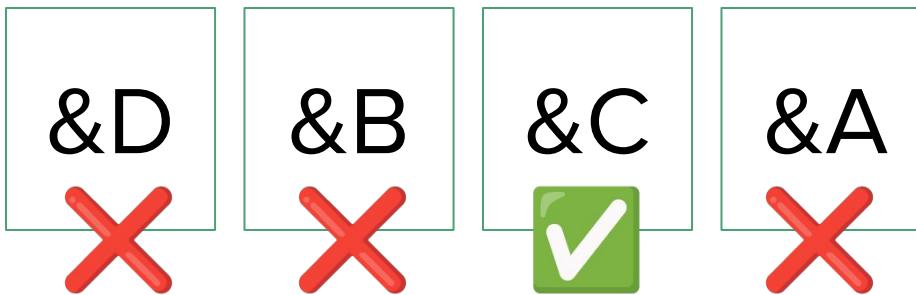
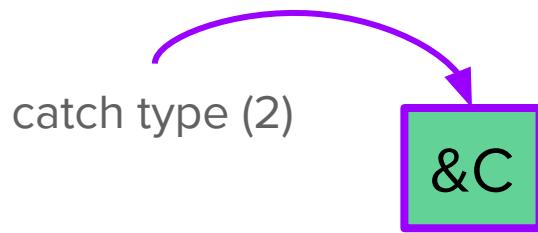
capacity



$O(N)$

$N = \# \text{ of entities}$
in class hierarchy

Dynamic_cast Algorithm



```
try {  
    // stuff...  
} catch (V const&) {  
} catch (C const&) {  
}
```

size



capacity



$O(N)$

$N = \# \text{ of entities}$
in class hierarchy

"dynamic_cast" Algorithm

O(N*M)

N = # of entities in class hierarchy

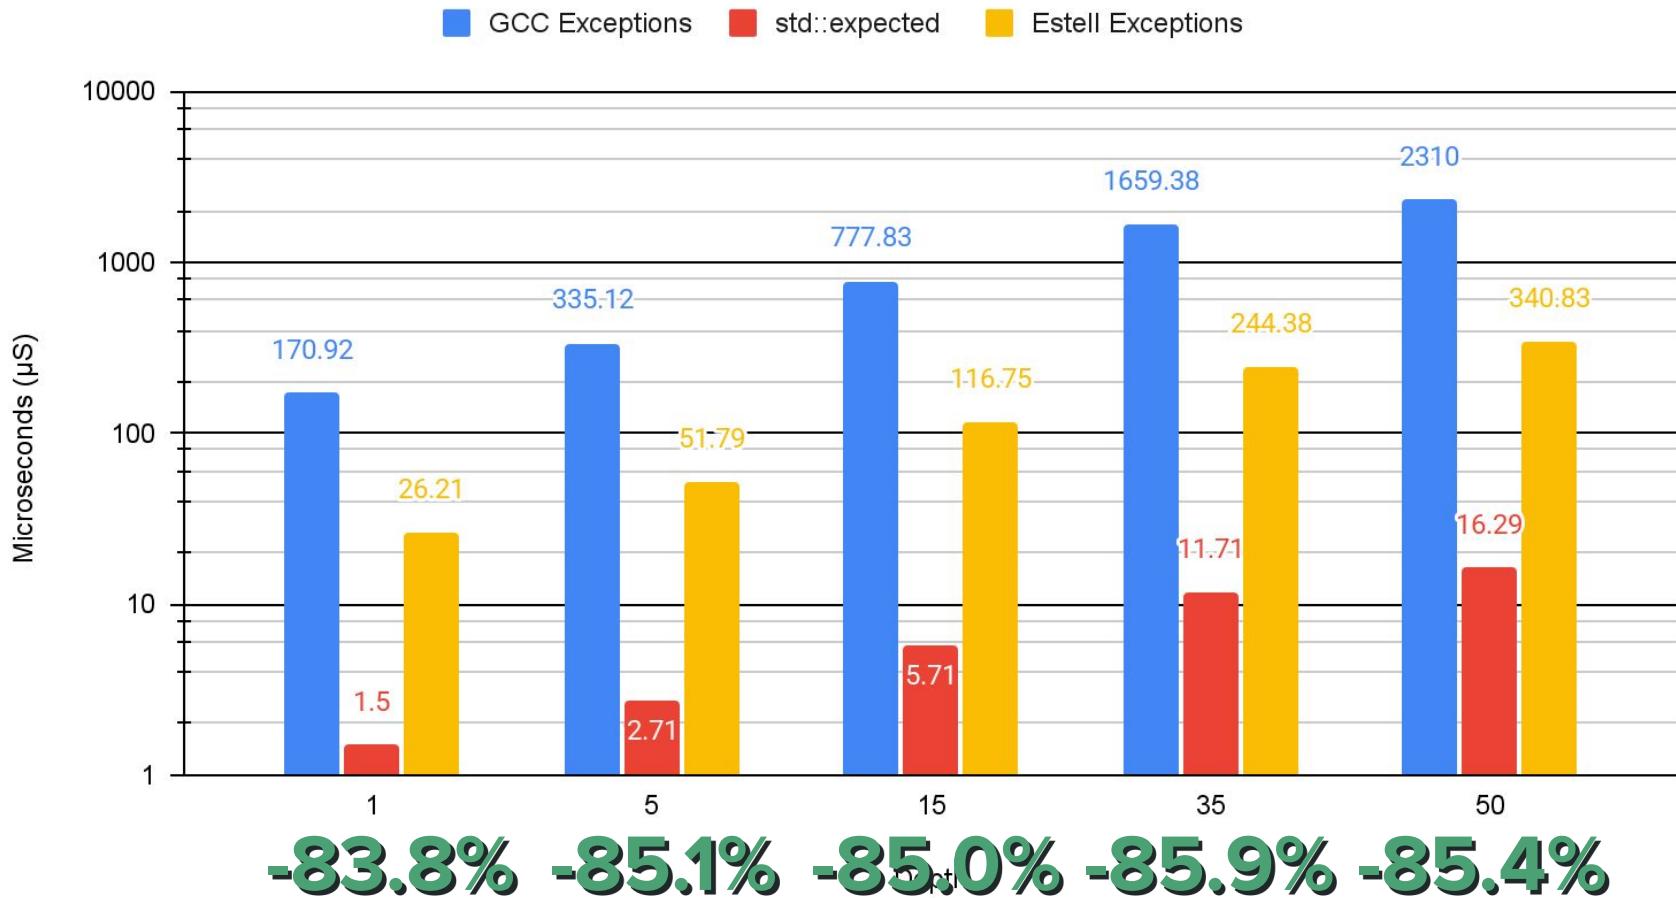
M = # of catches until handler found

```
for (auto const& type_entry : p_exception_object.type_info) {  
    if (type_entry.type_info == catch_type) {  
        p_exception_object.choose_type_offset = type_entry.offset;  
        // Prepare for context installation...  
        restore_cpu_core(cpu);  
    }  
}
```

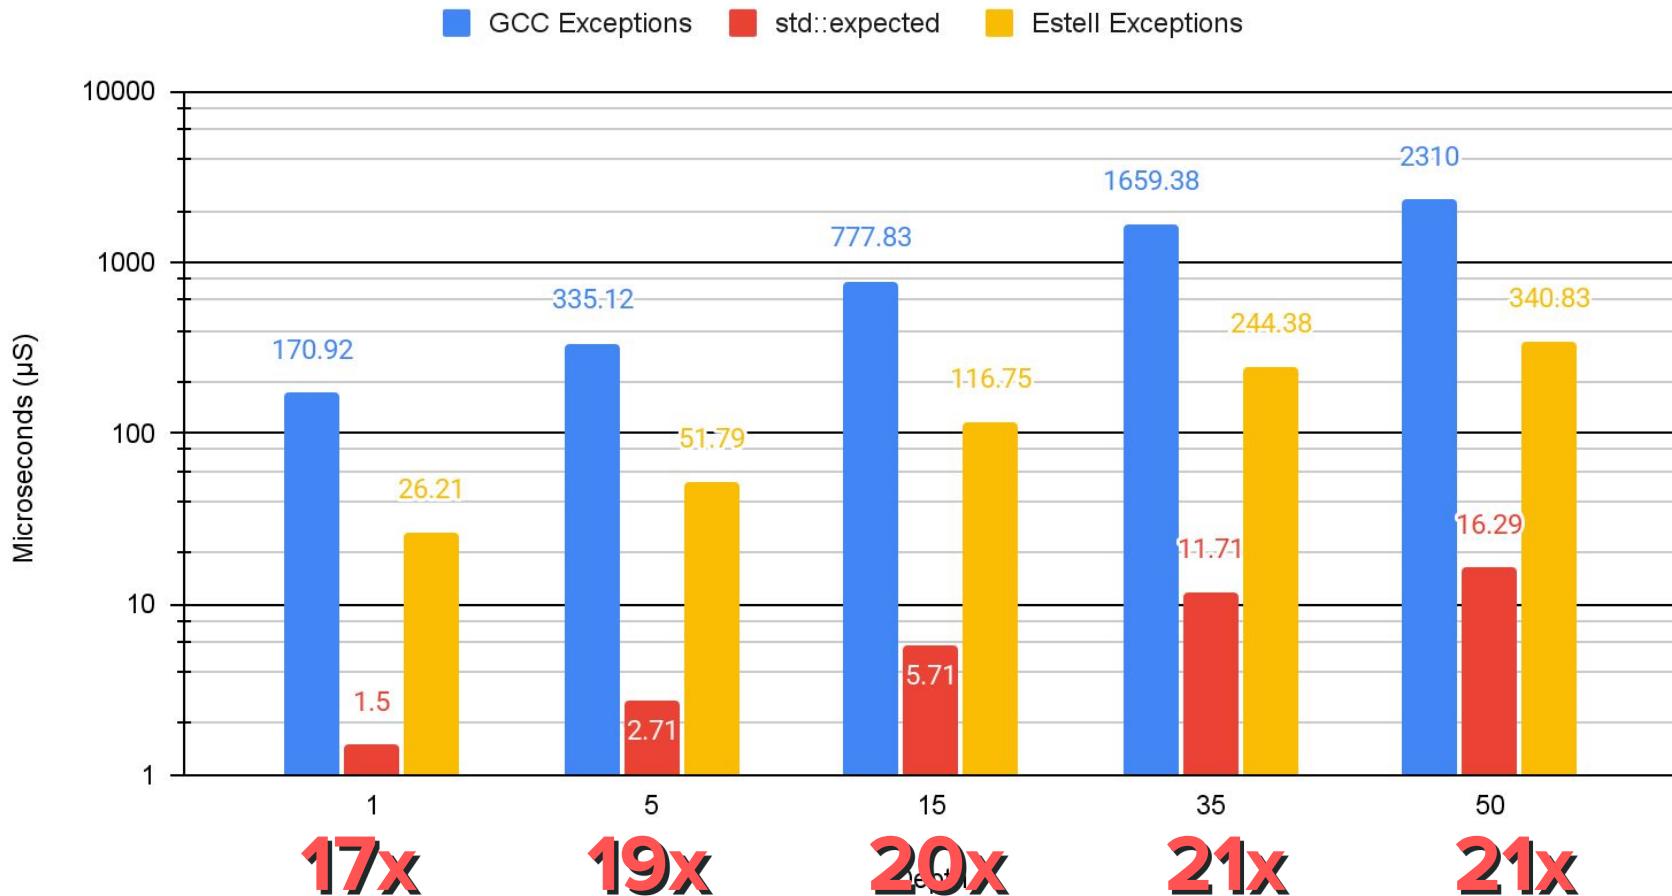


RESULTS!!

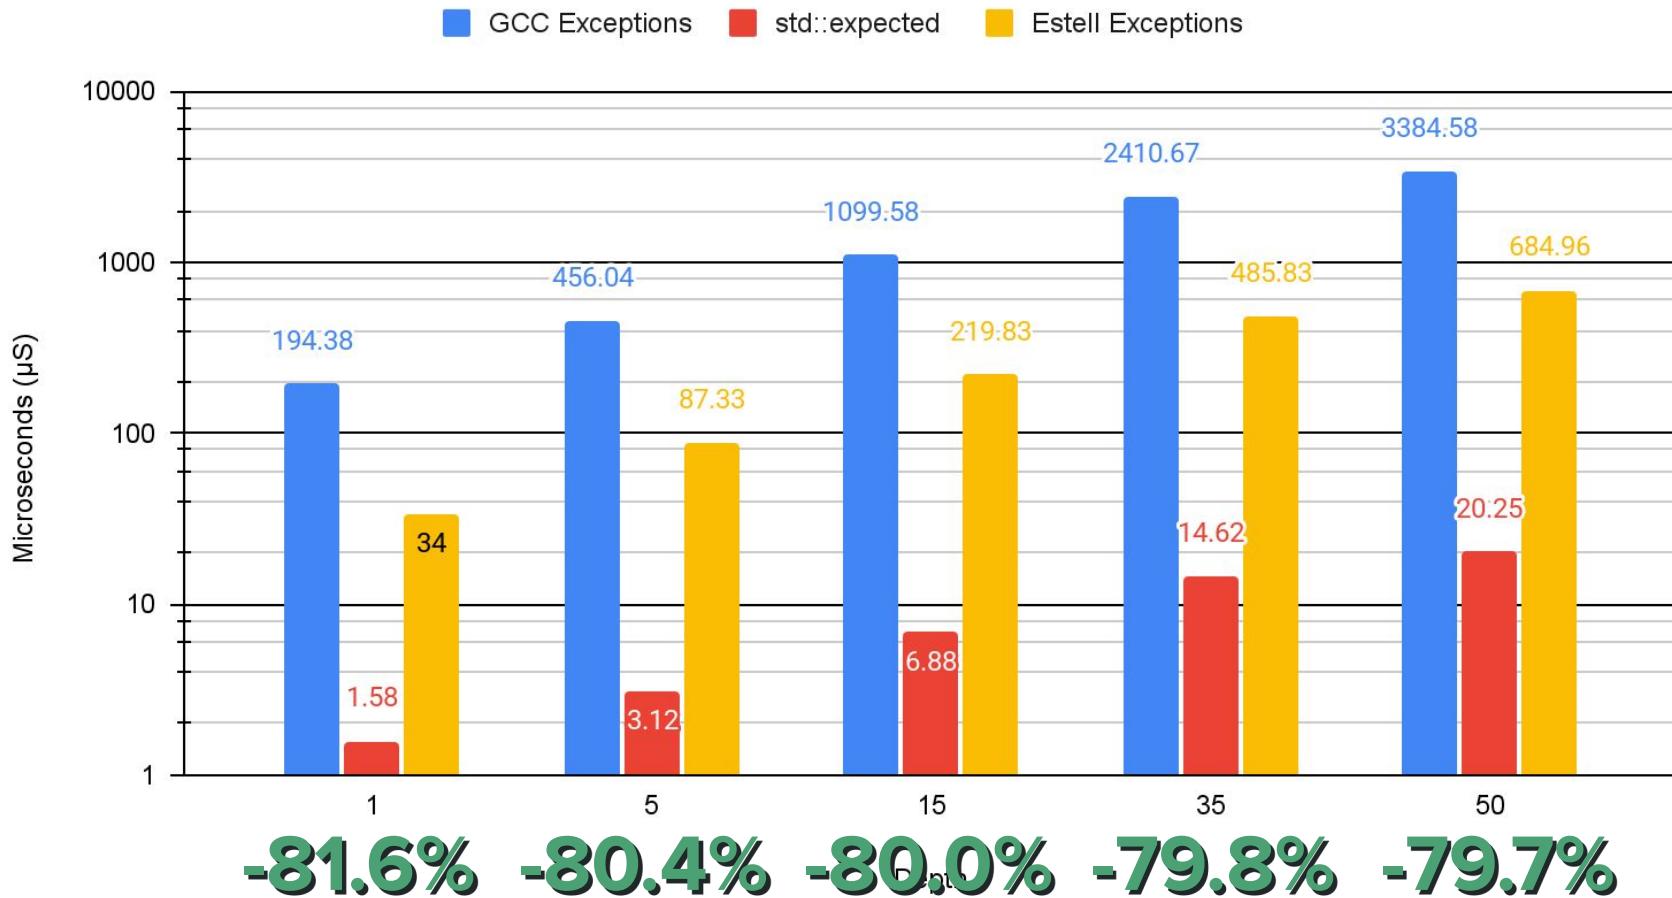
Error Propagation Time - 0% Cleanup, Error Size 4B



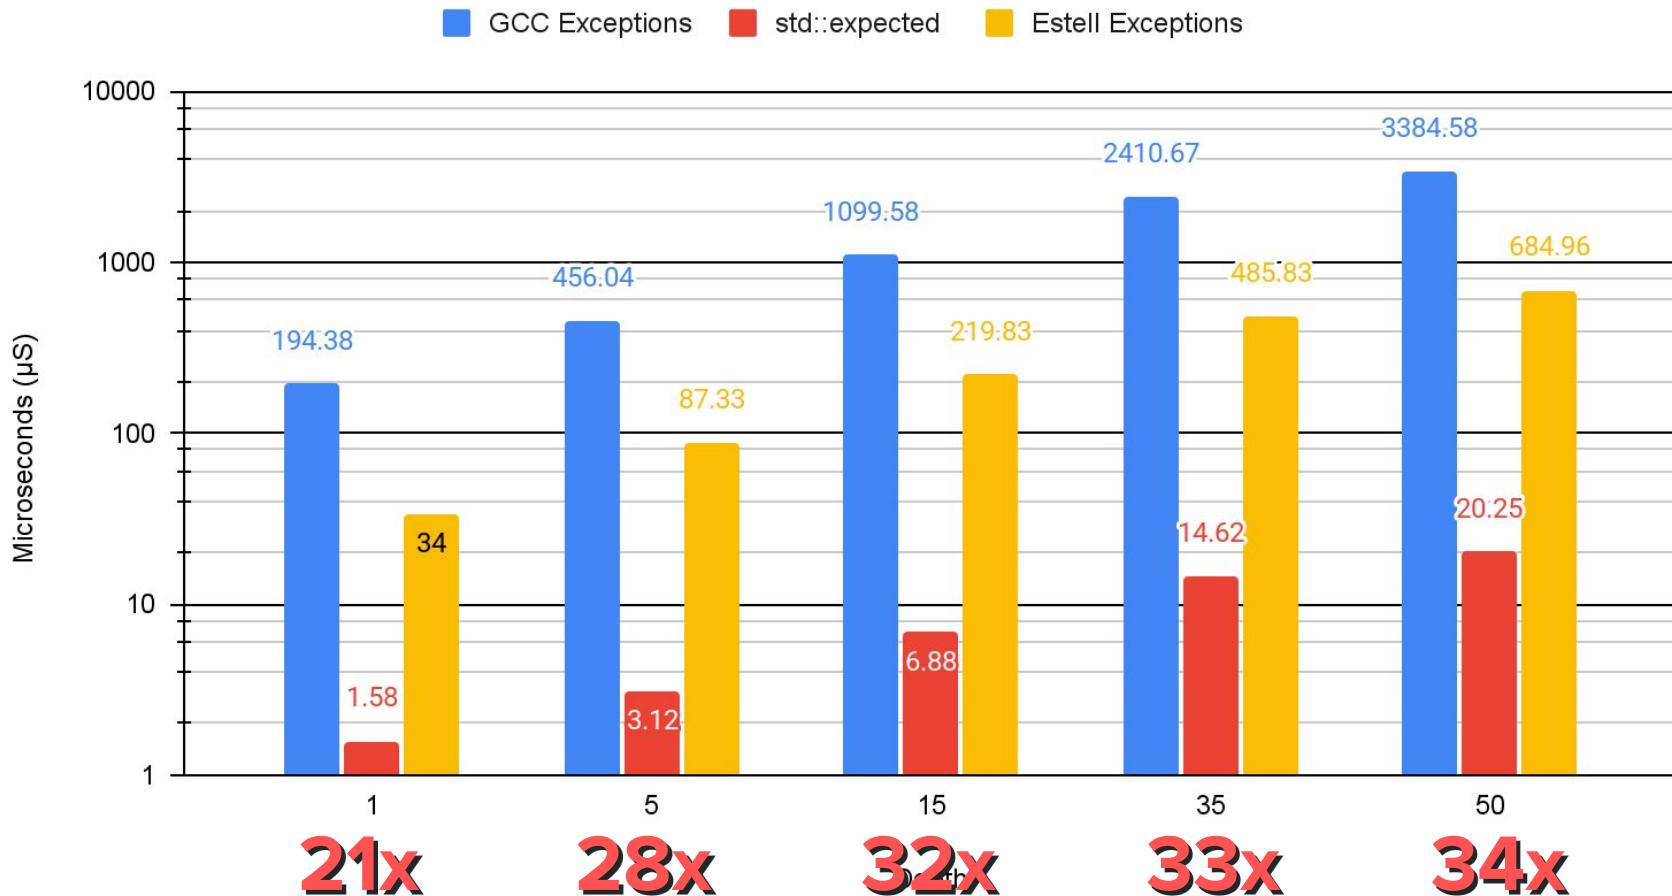
Error Propagation Time - 0% Cleanup, Error Size 4B



Error Propagation Time - 100% Cleanup, Error Size 4B



Error Propagation Time - 100% Cleanup, Error Size 4B



Wait... where is the



non-determinism???

I couldn't find any



On my microcontroller...



Use shared libraries

So why do firmware developers avoid C++ exceptions?

SPACE

Requires Dynamic Memory/Heap

Requires Malloc

Increases Binary Size

Requires whole
~~C++ STL~~

Run Time Type Info
(RTTI)

Unbounded Memory Usage

Exception Tables

Exception Code

TIME

Nondeterministic

Type Comparison (dynamic_cast)



Slow

Binary Search

Frame Unwinding

Frame Evaluation

So why do firmware developers avoid C++ exceptions?

SPACE

Requires Dynamic Memory/Heap

Requires Malloc

Increases Binary Size

Requires whole
C++ STL

Run Time Type Info
(RTTI)

Unbounded Memory Usage

Exception Tables

Exception Code

TIME

Nondeterministic

Type Comparison (dynamic_cast)

Slow

Binary Search

Frame Unwinding

 **Shared libraries**



Frame Evaluation

[Bug 71744](#) - Concurrently throwing exceptions is not scalable

Status: RESOLVED FIXED

Alias: None

Product: gcc

Component: libgcc ([show other bugs](#))

Version: 5.3.1

Importance: P3 normal

Target Milestone: ---

Assignee: Florian Weimer

URL:

Keywords:

Depends on:

Blocks:

Reported: 2016-07-03 22:11 UTC by Nadav Har'El

Modified: 2023-01-09 07:30 UTC ([History](#))

CC List: 5 users ([show](#))

See Also:

Host:

Target:

Build:

Known to work:

Known to fail:

Last reconfirmed: 2018-11-19 00:00:00

Florian Weimer 2023-01-09 07:30:59 UTC

Fixed via [r12-6210-g790854ea7670f1](#) for dynamically linked code:

commit 790854ea7670f11c14d431c102a49181d2915965
 Author: Florian Weimer <fweimer@redhat.com>
 Date: Tue Jan 4 15:47:30 2022 +0100

libgcc: Use _dl_find_object in _Unwind_Find_FDE

libgcc/ChangeLog:

- * unwind-dw2-fde-dip.c (_Unwind_Find_FDE): Call _dl_find_object if available.

And [r13-2706-g6e80a1d164d1f9](#) for the run-time code registration interface:

commit 6e80a1d164d1f996ad08a512c000025a7c2ca893
 Author: Thomas Neumann <tneumann@users.sourceforge.net>
 Date: Tue Mar 1 21:57:35 2022 +0100

eliminate mutex in fast path of __register_frame

The __register_frame/_deregister_frame functions are used to register unwinding frames from JITed code in a sorted list. That list itself is protected by object_mutex, which leads to terrible performance in multi-threaded code and is somewhat expensive even if single-threaded. There was already a fast-path that avoided taking the mutex if no frame was registered at all.

This commit eliminates both the mutex and the sorted list from the atomic fast path, and replaces it with a btree that uses optimistic lock coupling during lookup. This allows for fully parallel unwinding and is essential to scale exception handling to large core counts.

libgcc/ChangeLog:

- * unwind-dw2-fde.c (release_registered_frames): Cleanup at shutdown. (__register_frame_info_table_bases): Use btree in atomic fast path. (_deregister_frame_info_bases): Likewise. (_Unwind_Find_FDE): Likewise. (base_from_object): Make parameter const. (classify_object_over_fdes): Add query-only mode. (get_pc_range): Compute PC range for lookup.
- * unwind-dw2-fde.h (last_fde): Make parameter const.
- * unwind-dw2-btree.h: New file.

Bug 71744 - Concurrently throwing exceptions is not scalable

Status: RESOLVED FIXED

Alias: None

Product: gcc

Component: libgcc ([show other bugs](#))

Version: 5.3.1

Importance: P3 normal

Target Milestone: ---

Assignee: Florian Weimer

URL:

Keywords:

Depends on:

Blocks:

So why do firmware developers avoid C++ exceptions?

SPACE

Requires Dynamic Memory/Heap

~~Requires Malloc~~

Increases Binary Size

~~Requires whole C++ STL~~

Run Time Type Info (RTTI)

~~Exception Tables~~

~~Exception Code~~

Unbounded Memory Usage

TIME

Nondeterministic

~~Type Comparison (dynamic_cast)~~



Shared libraries



Slow

Binary Search

Frame Unwinding

Frame Evaluation

Chapter 4.

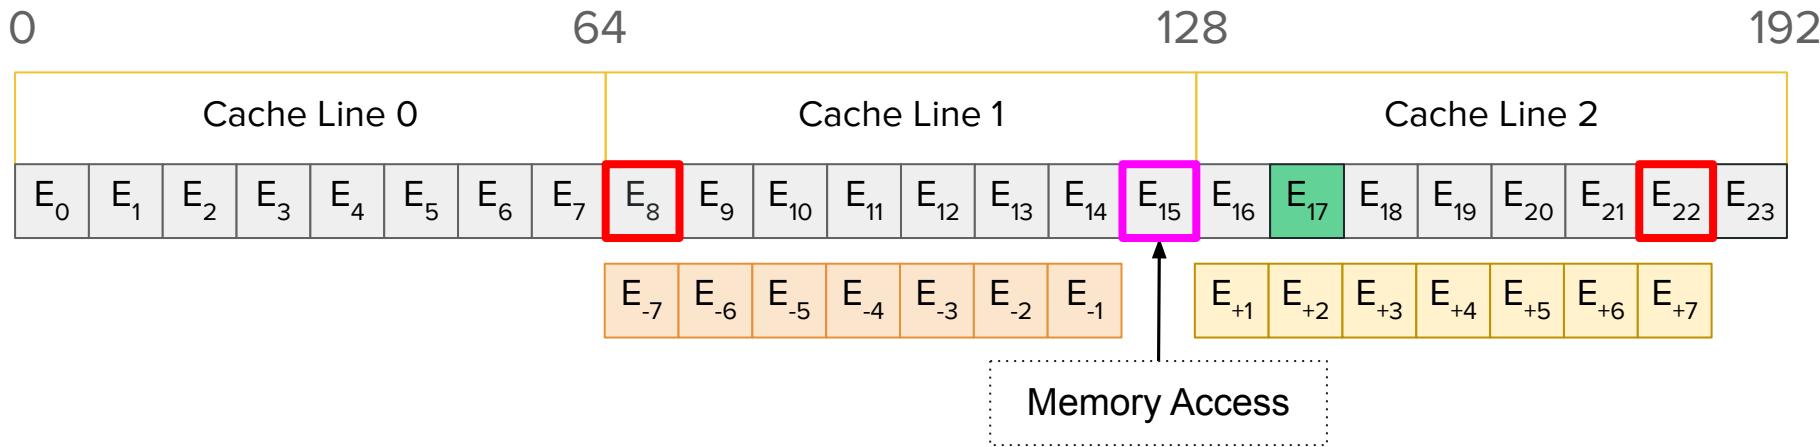


Making Nearpoint Search Work

Nearpoint Search Algorithm Requirements

- ✓ Order of magnitude faster than binary search
- ✓ Try and preserve exception space saving
- ✓ Integer math ONLY
- ✓ Max guess error < 8 ✨ NEW ✨

Why limit error to less than 8? (Max 2 cache misses)



64 bytes can hold
8 exception entries



Entry Coalescence



P3313R0

Impacts of noexcept on ARM table based exception metadata

Draft Technical Report, 2024-05-22

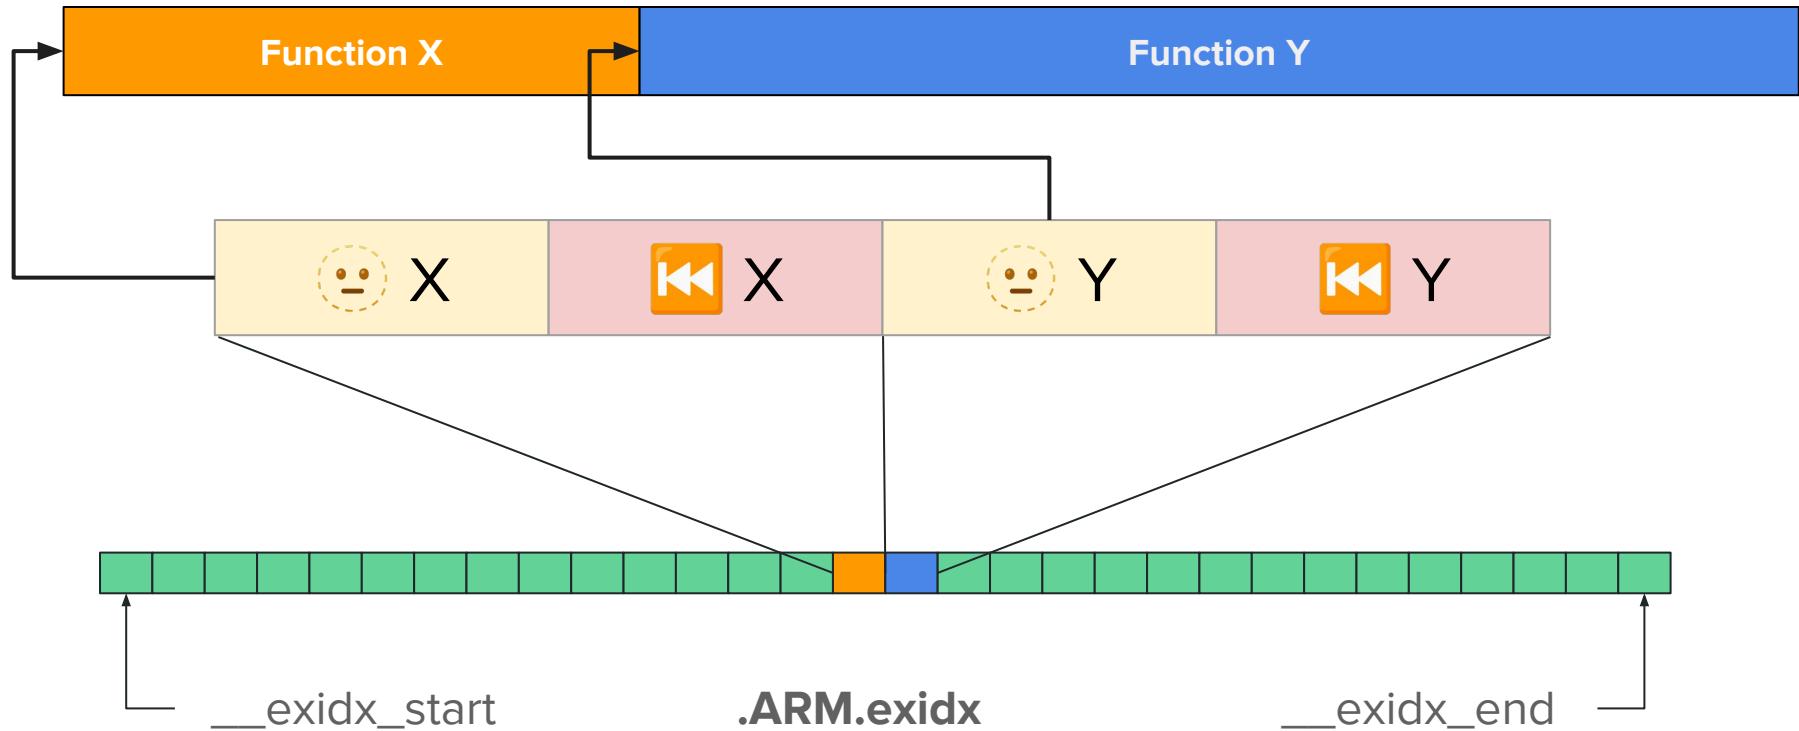
This version:

<https://wg21.link/P3313R0>

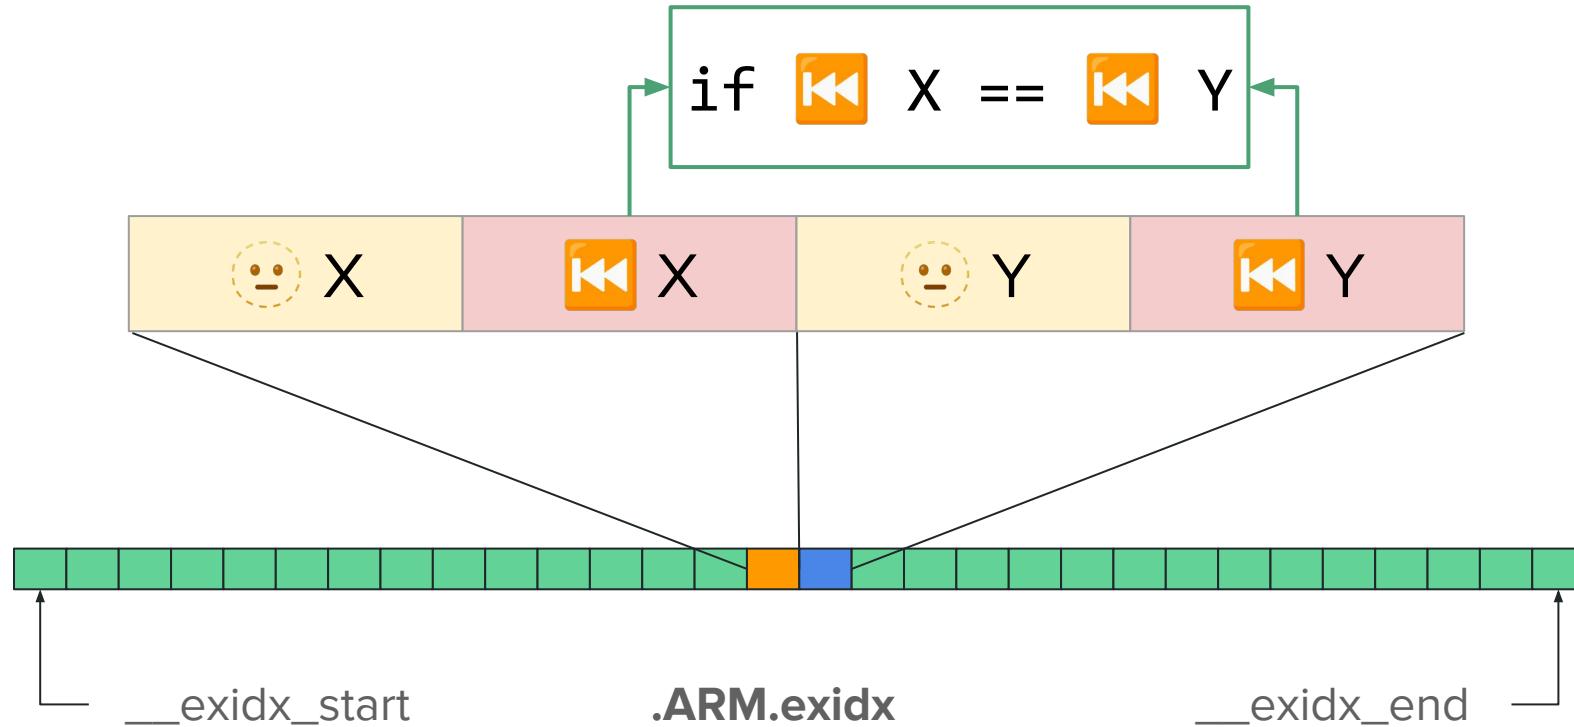
Author:

Khalil Estell

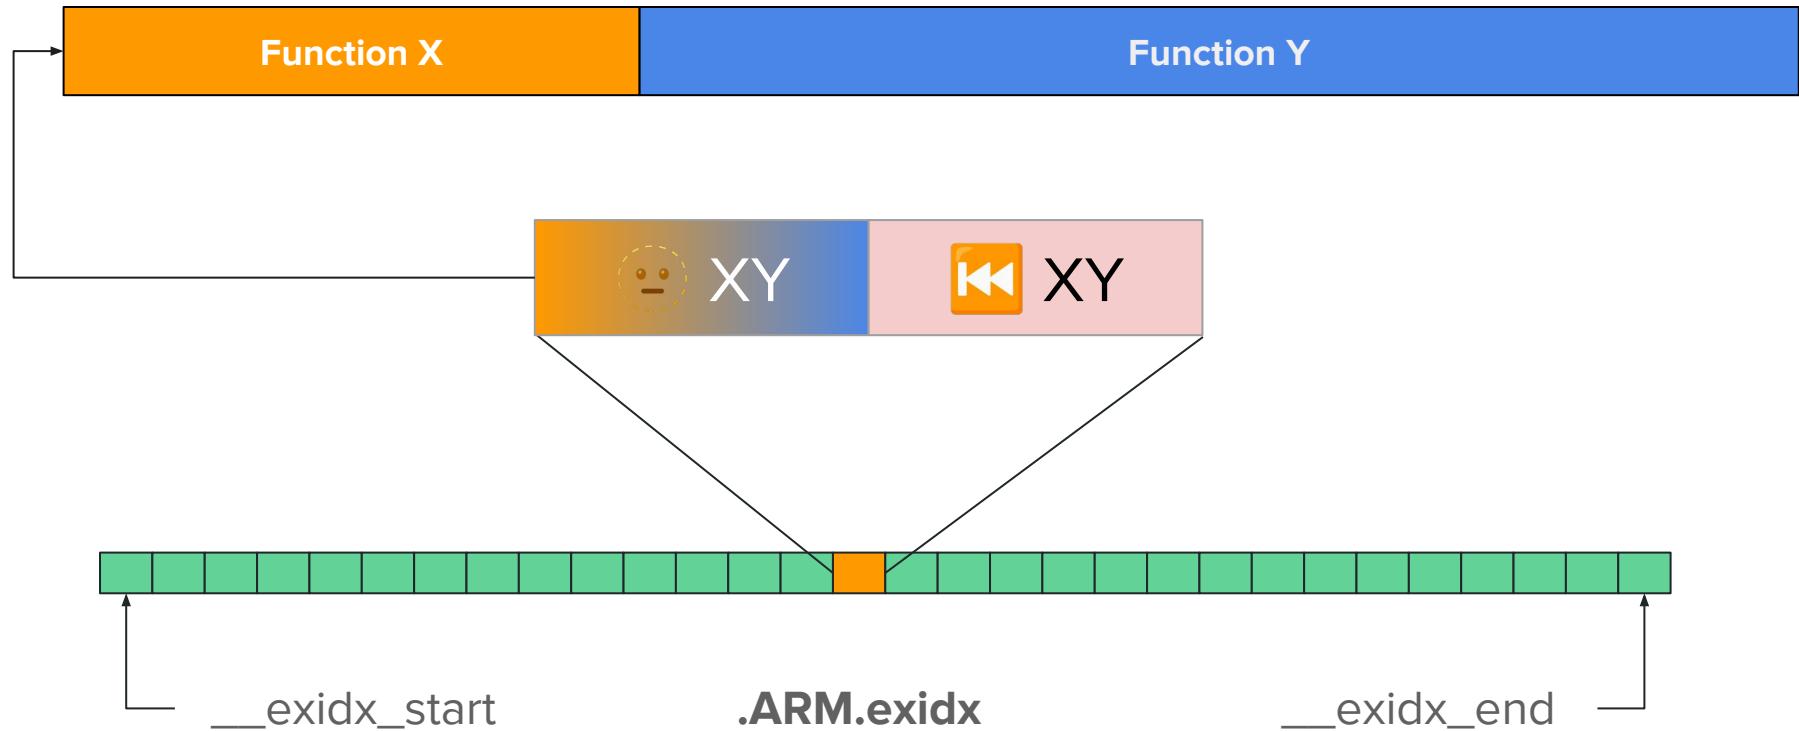
GCC's Entry Coalescence



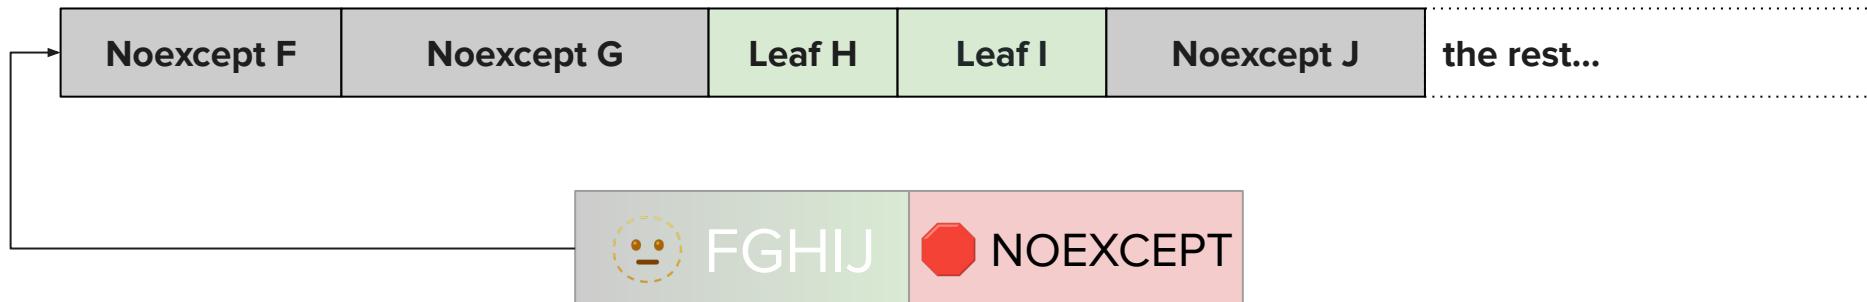
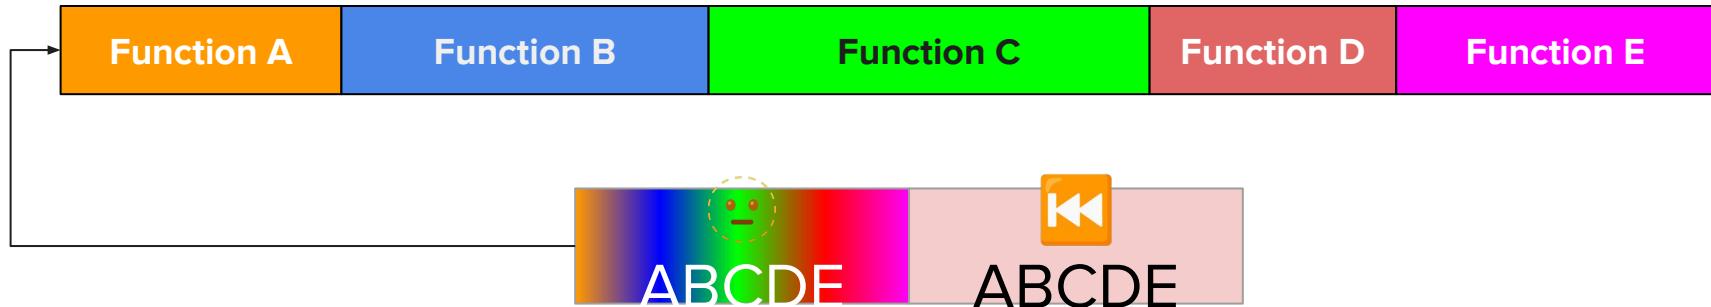
GCC's Entry Coalescence



GCC's Entry Coalescence



Function Grouping



Nearpoint Function Group Sort by Size



entry #0

ABCDEF	ABCDEF
F	F
G	G
HIJ	HIJ

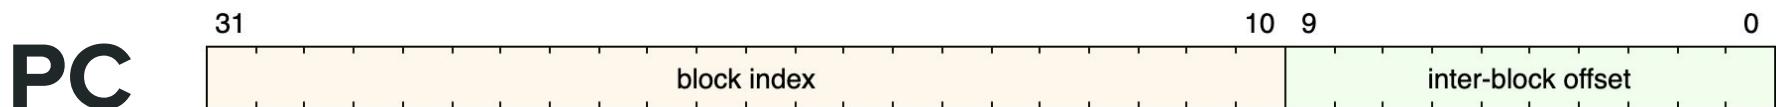
entry #1

entry #2

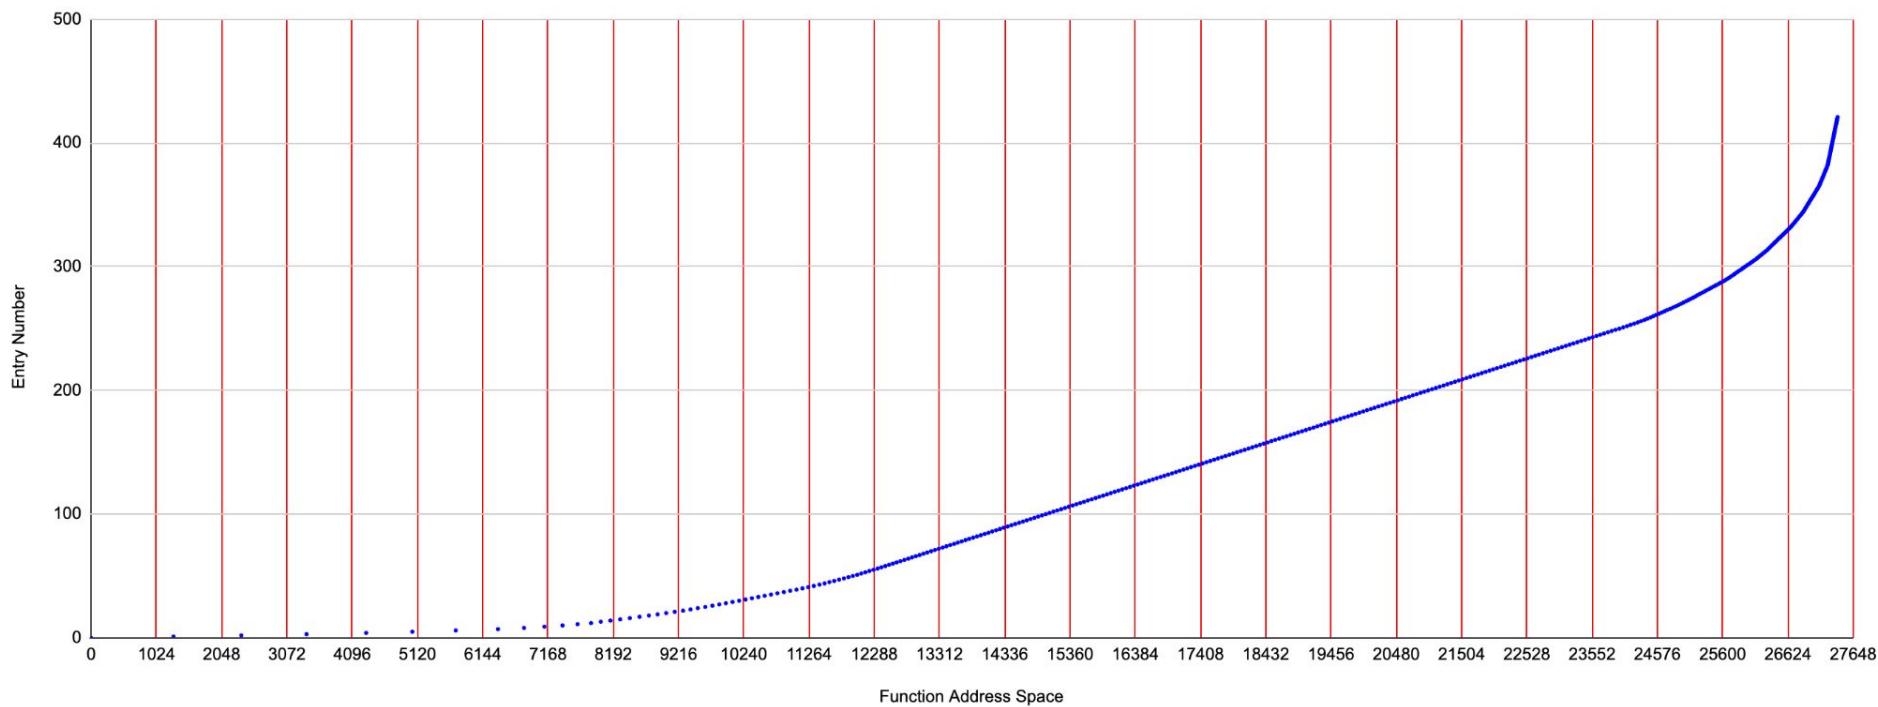
entry #3

•
•
•
•

Block size must be a power of 2 (1024 in this example)



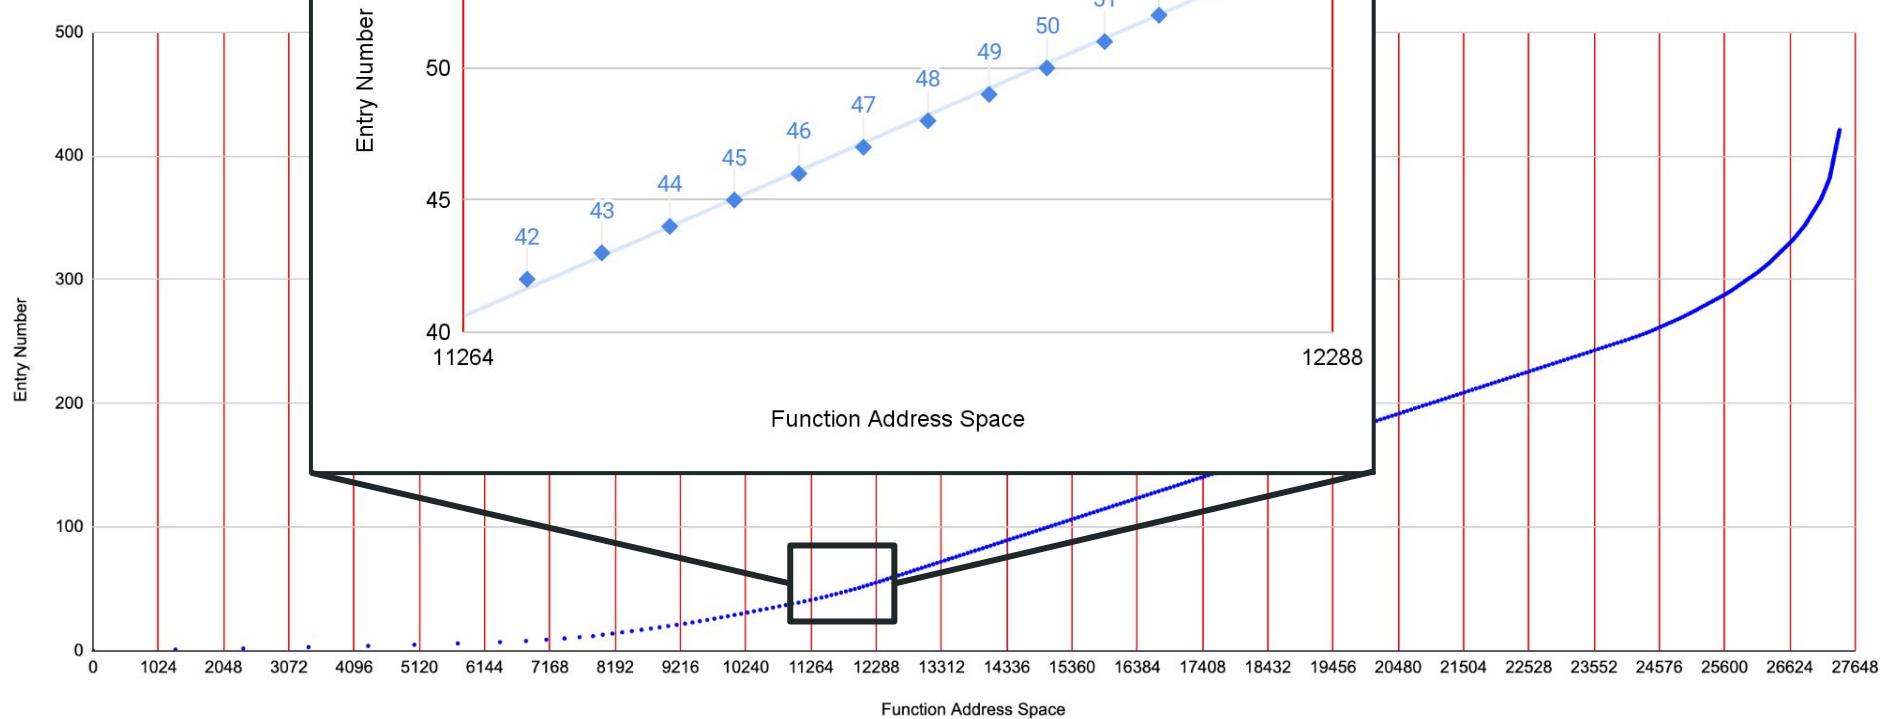
Final Address vs. Index Entry



Final Address vs. Index Entry

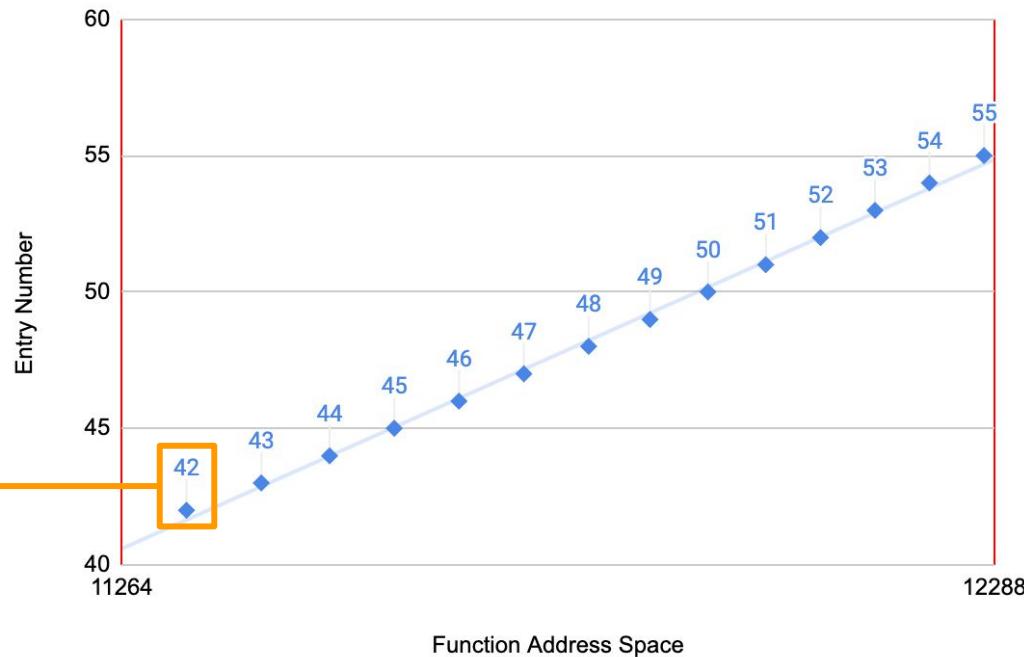
$$y = m * (x - x_0) + b$$

Final Address vs. Index



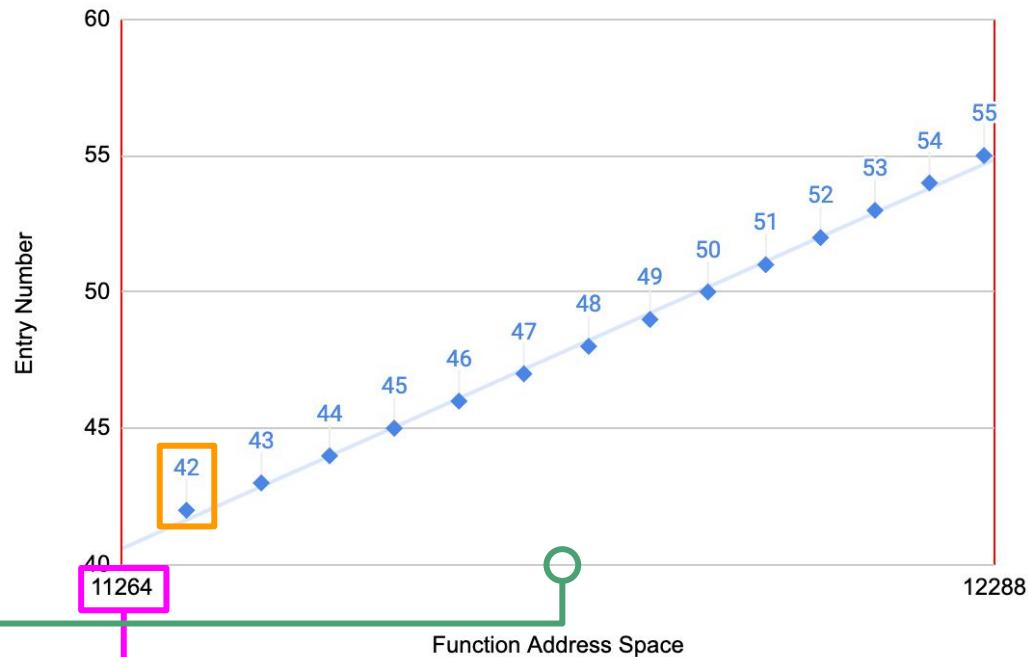
$$Y = m * (x - x_0) + b$$

Final Address vs. Index Entry



$$Y = m * (x - x_0) + b$$

Final Address vs. Index Entry

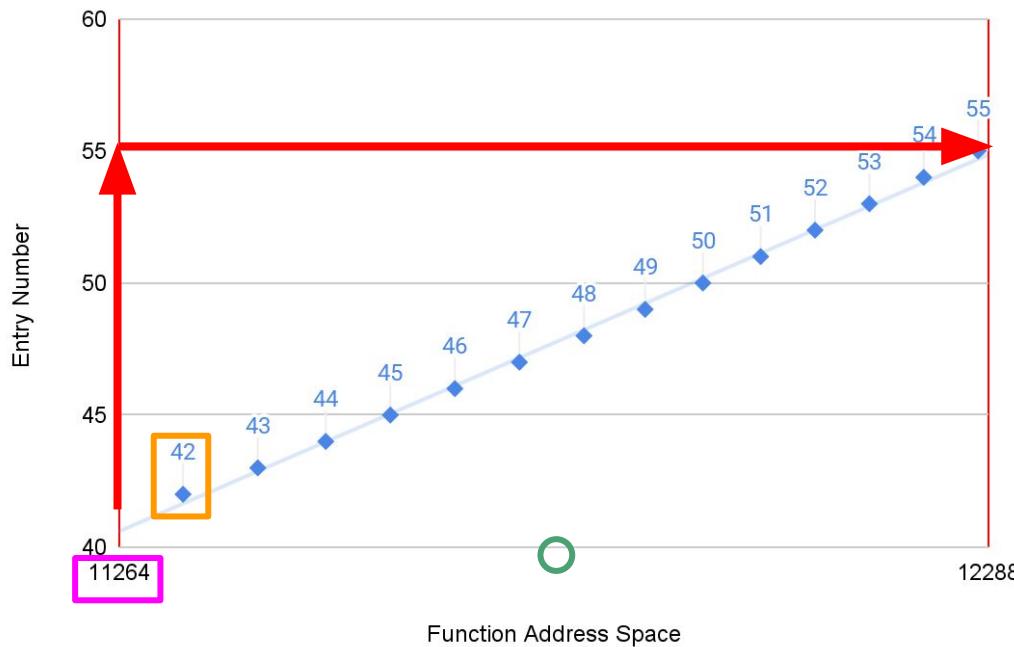


$$y = \boxed{m} * \boxed{(x - x_0)} + \boxed{b}$$

Rise
13

Run
1024
(same for all blocks)

Final Address vs. Index Entry



31

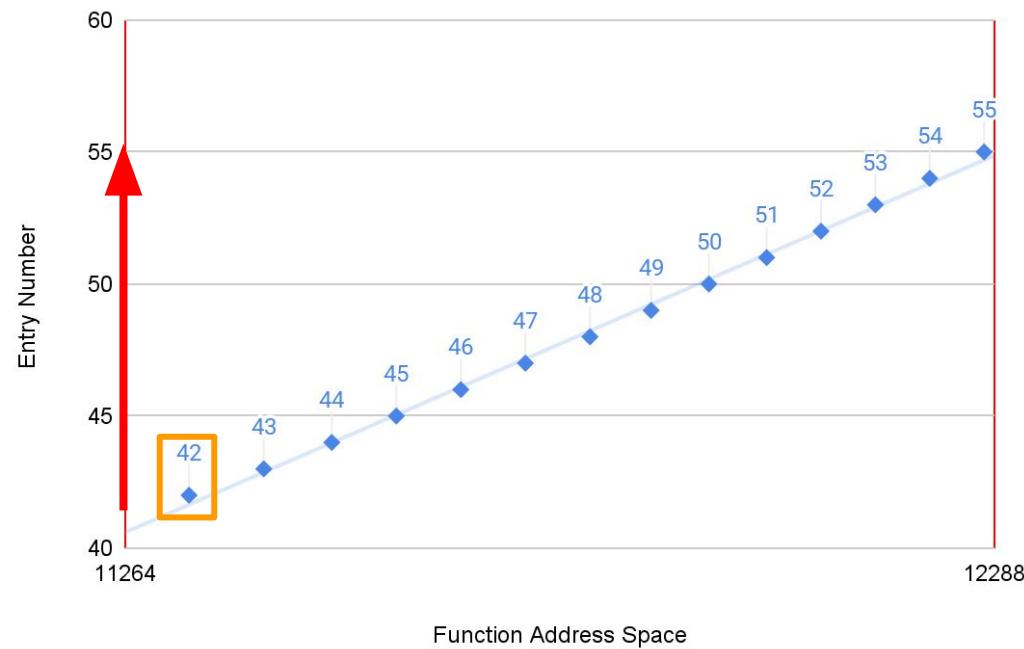
22 21

0

entry count

starting entry number

Final Address vs. Index Entry



$$Y = m * (x - x_0) + b$$

```
struct nearpoint_descriptor
{
    std::uint32_t block_power = 0;
    std::uint32_t text_address = 0;
};

std::array<std::uint32_t, 2> const
_near_point_descriptor_data = {
    0x0000000b,
    0x08000044,
};
```

```
std::array<std::uint32_t, 18> const _normal_table_data = {
    (0 << 11) | 1, // Block(start=0, count=1)
    (1 << 11) | 1, // Block(start=1, count=1)
    (1 << 11) | 1, // Block(start=1, count=1)
    (2 << 11) | 2, // Block(start=2, count=2)
    (3 << 11) | 3, // Block(start=3, count=3)
    (5 << 11) | 3, // Block(start=5, count=3)
    (7 << 11) | 6, // Block(start=7, count=6)
    (12 << 11) | 14, // Block(start=12, count=14)
    (25 << 11) | 20, // Block(start=25, count=20)
    (44 << 11) | 27, // Block(start=44, count=27)
    (70 << 11) | 28, // Block(start=70, count=28)
    (97 << 11) | 32, // Block(start=97, count=32)
    (128 << 11) | 26, // Block(start=128, count=26)
};
```

```
std::uintptr_t near_point_guess_index (std::uintptr_t p_program_counter)
{
    auto const block_power = __except_abi::near_point_descriptor.block_power;
    auto const program_offset = __except_abi::near_point_descriptor.text_address;
    auto const inter_block_mask = (1U << block_power) - 1U;

    auto const pc = p_program_counter - program_offset;
    auto const inter_block_location = pc & inter_block_mask;
    auto const block_index = pc >> block_power;
    auto const block_info = __except_abi::normal_table[block_index];

    auto const entry_start = block_info >> block_power;
    auto const entry_count = block_info & inter_block_mask;

    auto const guess_offset = (inter_block_location * entry_count) >> block_power;
    auto const guess_location = entry_start + guess_offset
    return guess_location;
}
```

Pull info from descriptor

Break apart PC

Lookup block_info

Break up block_info

Calculate linear equation

```
index_entry_t const& get_index_entry_near_point (std::uint32_t p_program_counter)
{
    auto const index_table = get_arm_exception_index ();
    auto const initial_guess = near_point_guess_index (p_program_counter);
    auto it = index_table.begin () + initial_guess;

    if (p_program_counter < it->function ()) {
        do {
            --it;
        } while (it->function () > p_program_counter);
    } else {
        do {
            ++it;
        } while (it->function () <= p_program_counter);
        --it;
    }

    return *it;
}
```

Guess

Scan Left

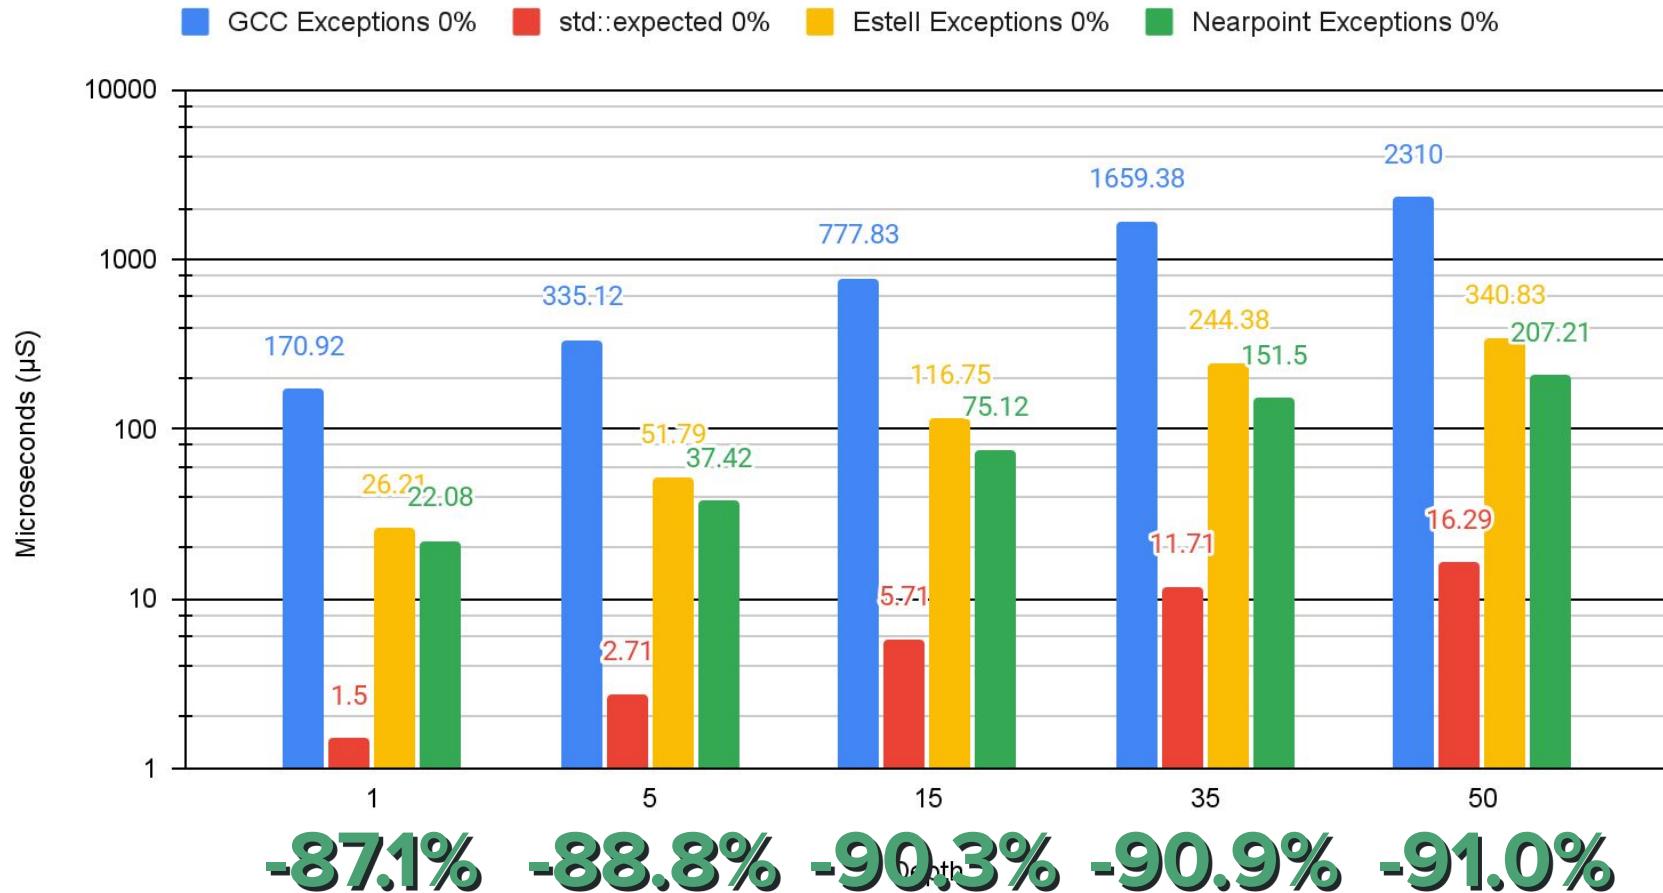
Scan Right



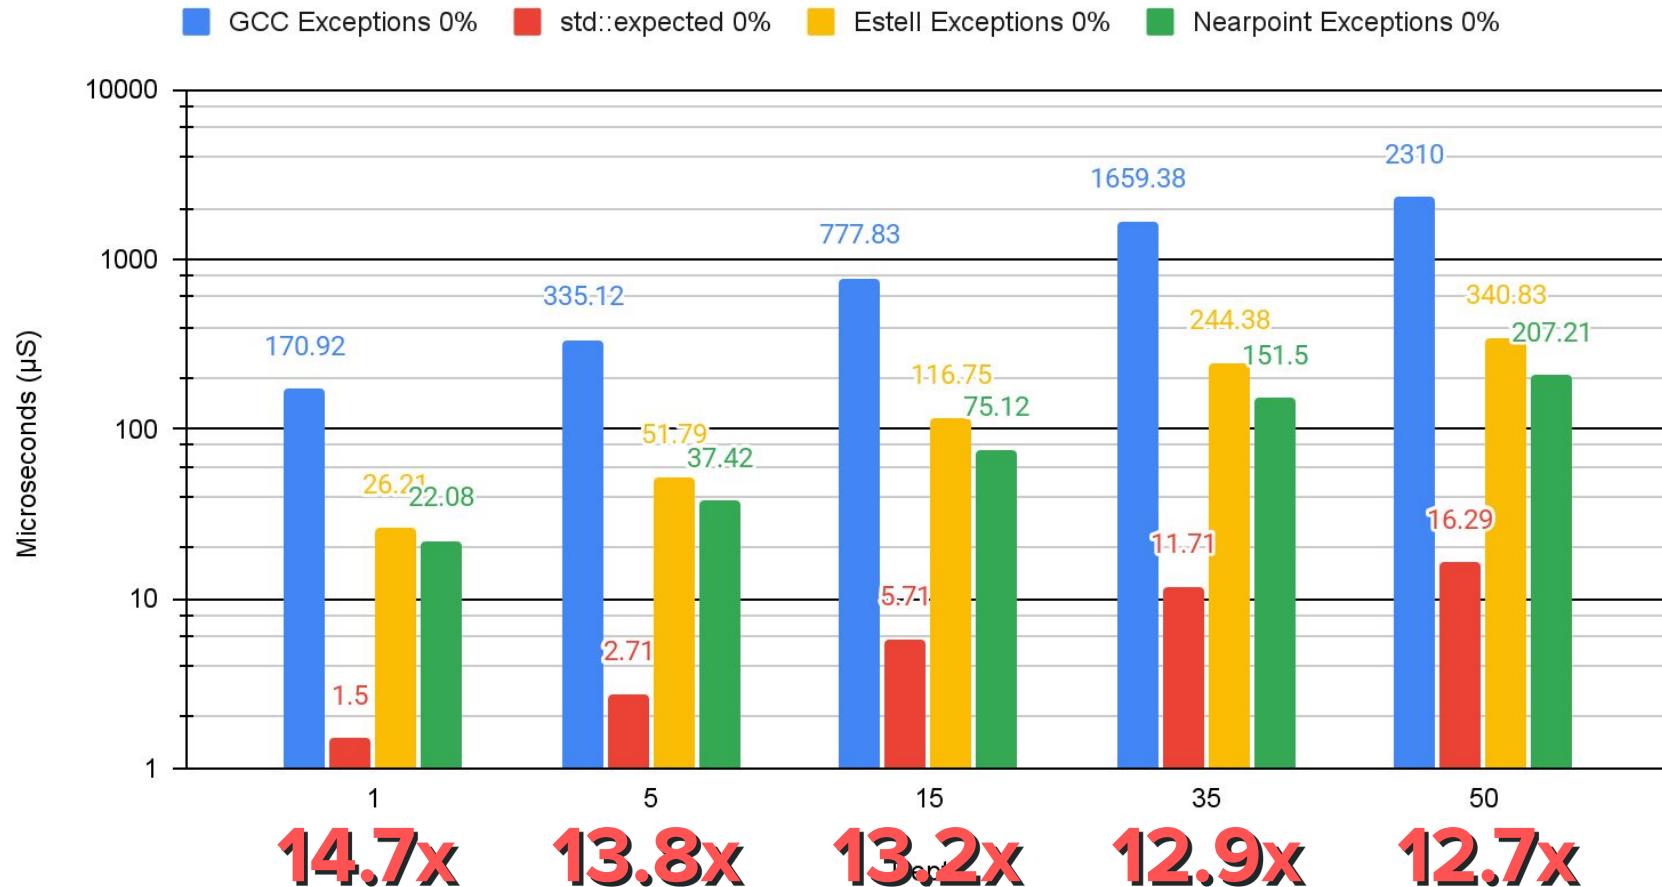
RESULTS!!

>90% @ >=11 frames

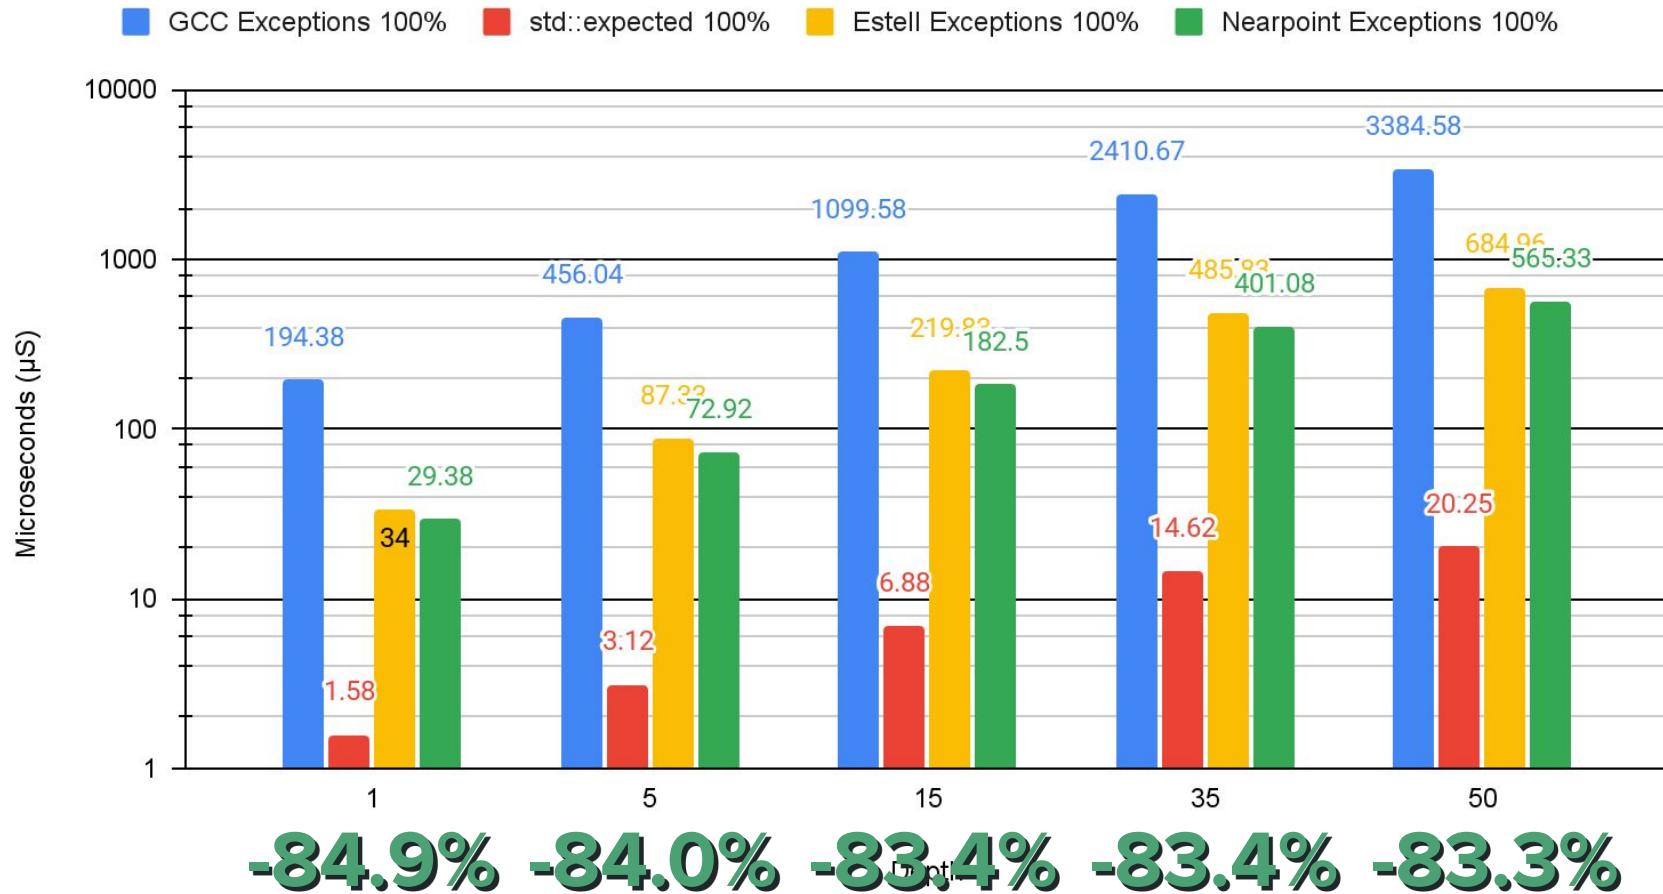
Error Propagation Time - 0% Cleanup, Error Size 4B



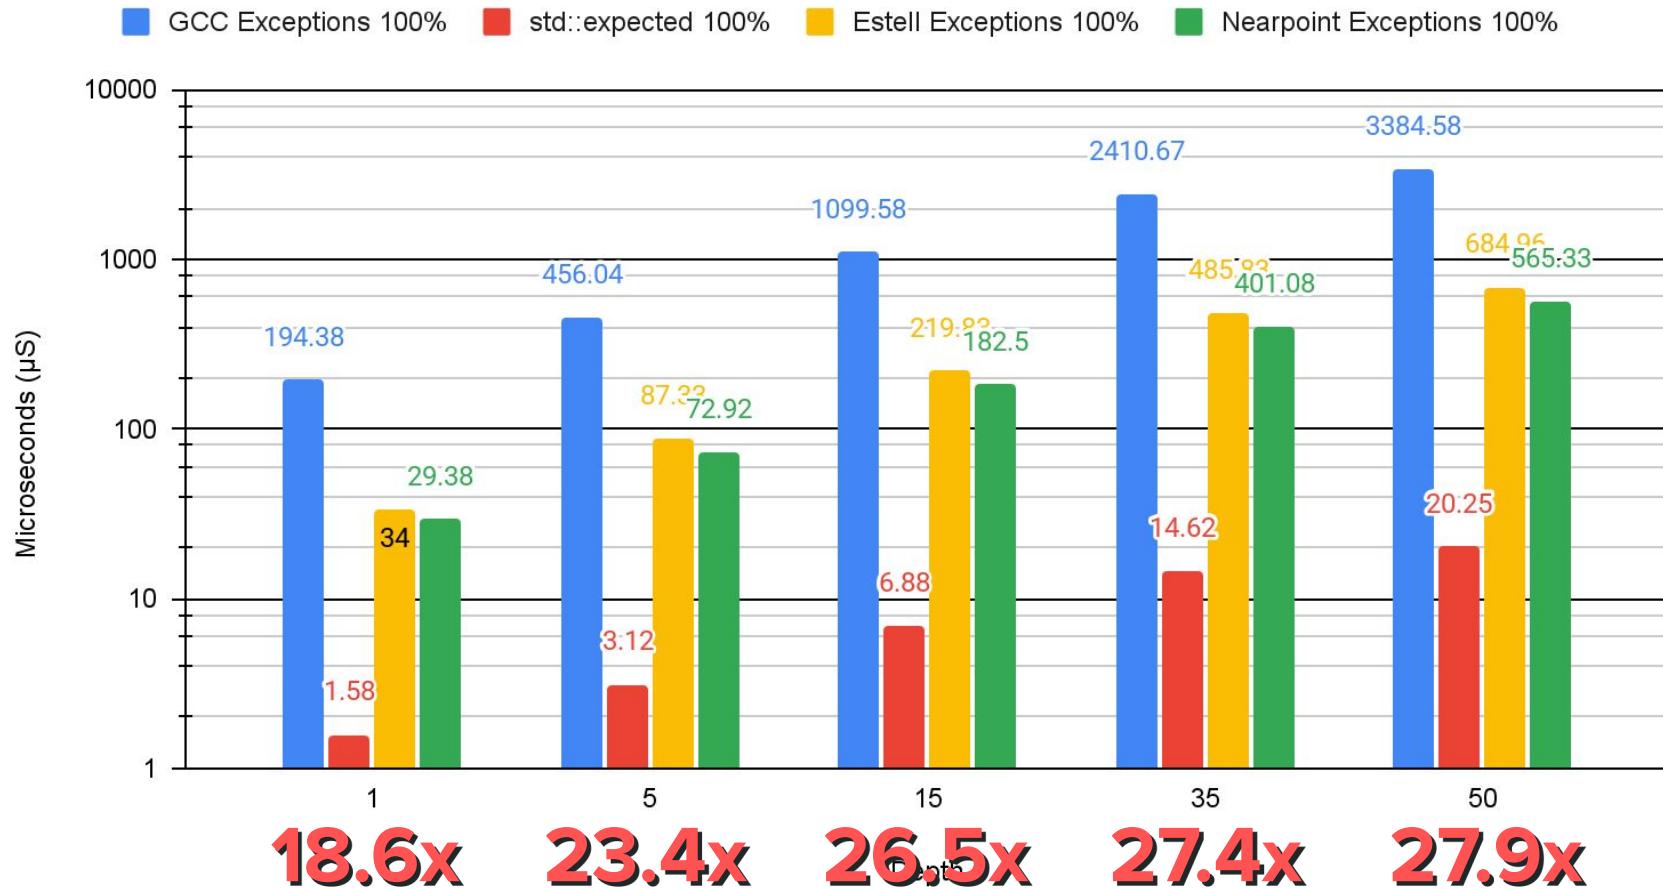
Error Propagation Time - 0% Cleanup, Error Size 4B



Error Propagation Time - 100% Cleanup, Error Size 4B



Error Propagation Time - 100% Cleanup, Error Size 4B



Binary Size Difference

text	data	bss	dec	hex	filename
47575	240	2648	50463	c51f	stm32f103c8/estell/Release/except.cpp.elf

Binary Size Difference

text	data	bss	dec	hex	filename
47575	240	2648	50463	c51f	stm32f103c8/estell/Release/except.cpp.elf
46987	268	2648	49791	c27f	stm32f103c8/estell/Release/nearpoint.cpp.elf



-588 bytes

**Block power = 11*

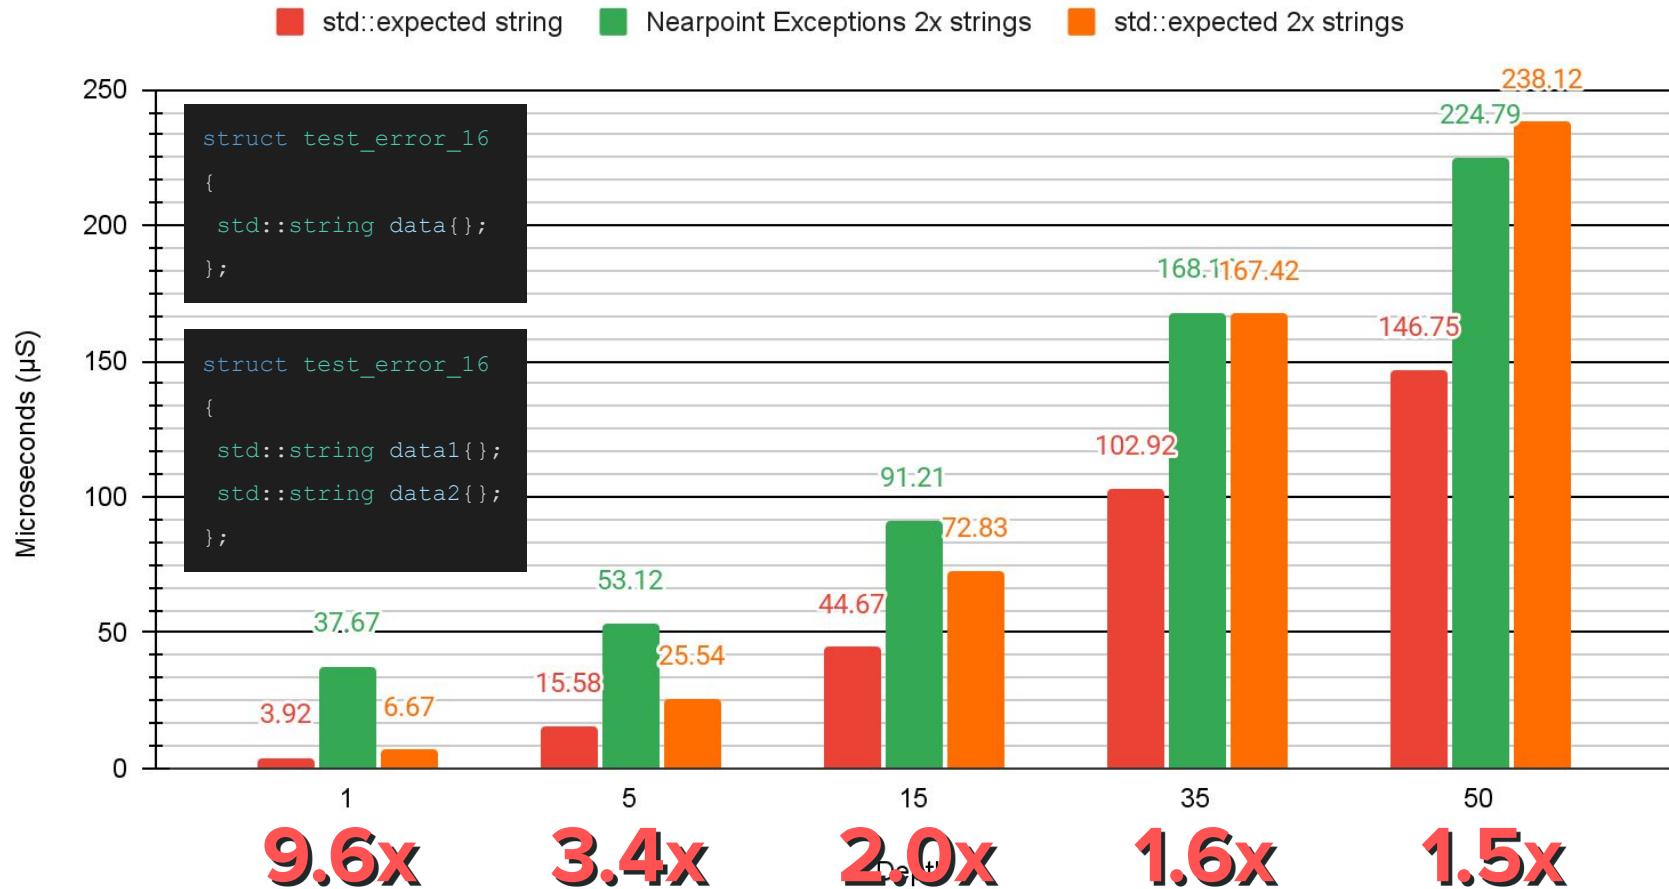
 Let's compare to an `std::string{}; // empty`

- Default initialized to an empty string
 - No dynamic memory allocation
- Has a nontrivial move constructor
- A stand-in for a **complex error object**

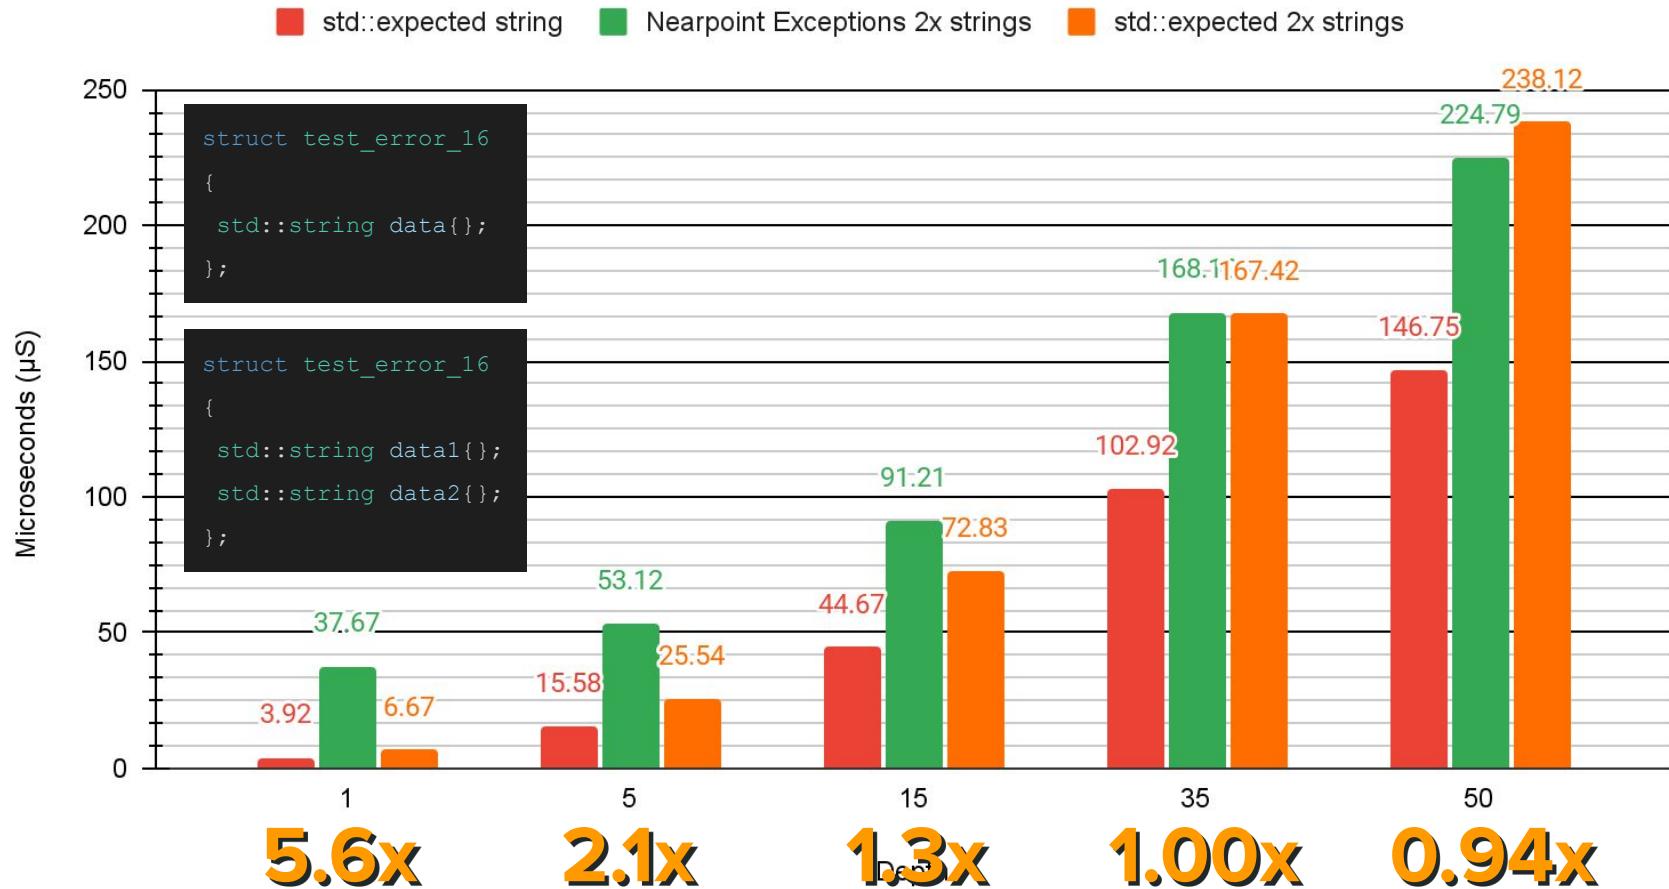
```
struct test_error_16
{
    std::string data{};
};
```

```
struct test_error_16
{
    std::string data1{};
    std::string data2{};
};
```

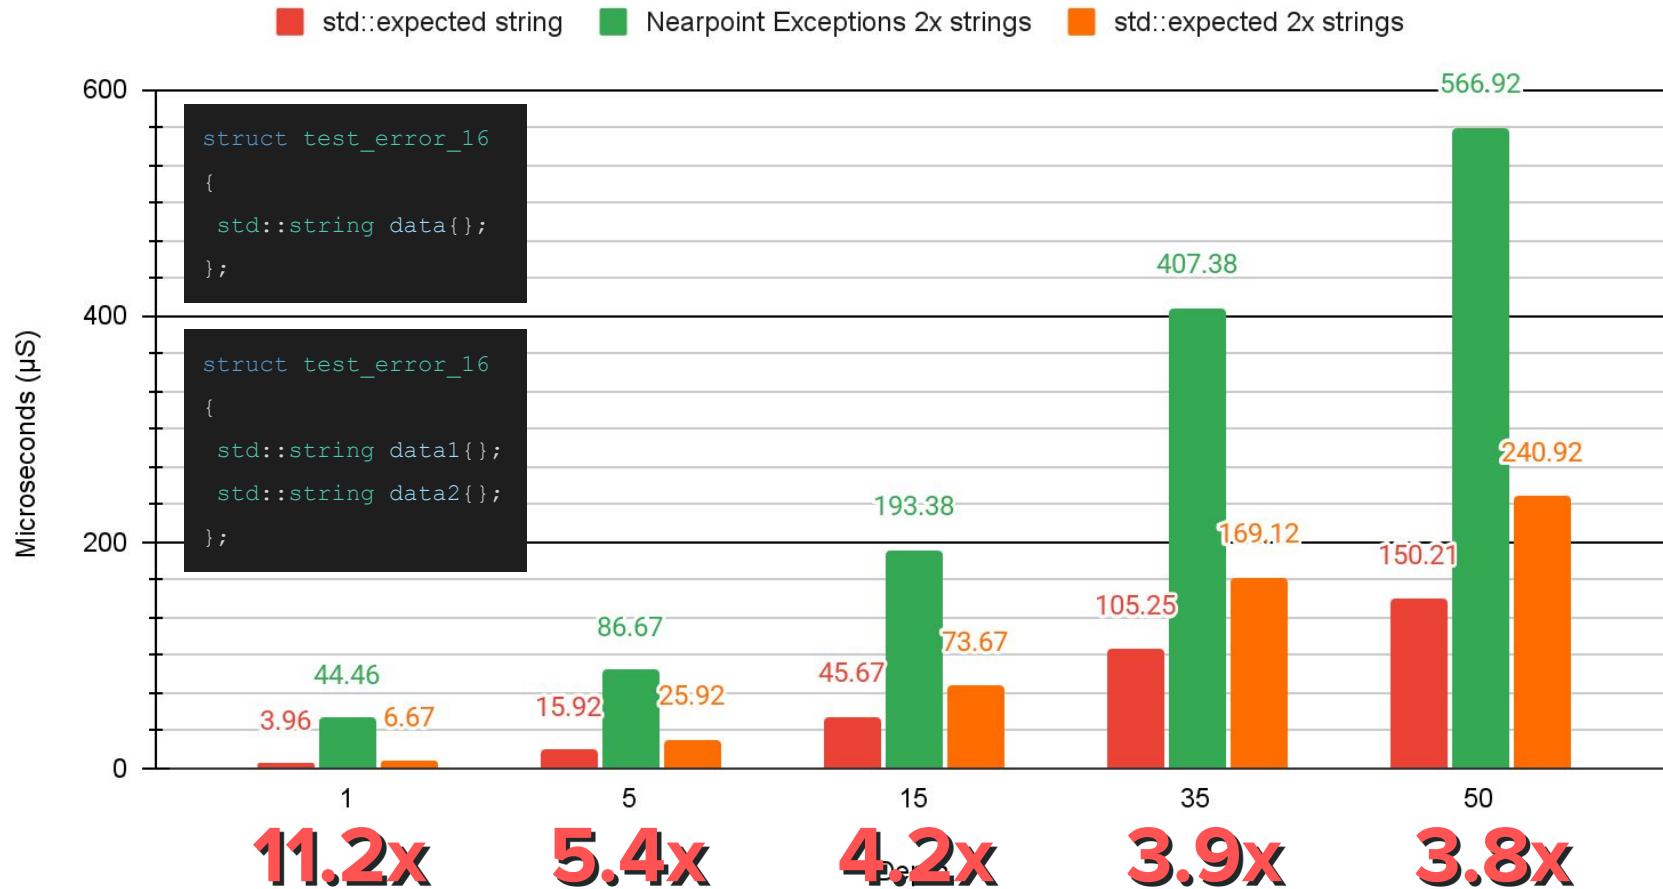
Error Propagation Time - 0% Cleanup, Error == empty std::string



Error Propagation Time - 0% Cleanup, Error == empty std::string



Error Propagation Time - 100% Cleanup, Error == empty std::string



Why is `std::expected` impacted so badly?

Result type propagation must:

1. $O(M)$: Determine a suitable handler
2. $O(N)$: Unwind frame
3. **$O(N)$:** Move/copy error up a frame

Exception propagation must:

1. $O(M)$: Determine a suitable handler
 2. $O(N)$: Unwind frames
 3. **$O(1)$:** Allocate & construct error
-

RECAP: How did I reduce C++ exceptions via GCC on an ARM Cortex-M by 90%?

- Used C++ standard libraries when appropriate
- Moved as much of the code to compile time as possible
- Applied data driven design to exception index & unwind instructions
- Cached & pre-formatted data when appropriate.
- Rewrote runtime in C++
- **I thought it was possible**

Chapter 5.



Future of Exceptions (as Khalil sees it)

Exceptional Software Design & Exception Insights Tool

Part 3

Unique Exception Format

(De-duplicate Exception Data Table)

First step to exceptions in freestanding C++

```
void set_exception_allocator(std::pmr::memory_resource*) noexcept;  
std::pmr::memory_resource* get_exception_allocator() noexcept;
```

Async (coro) Exceptions Unwinding Async Destructors

Vendoring this Technology

```
-Wl,--plugin=estell_exceptions.so  
-Wl,--plugin=exception_insight.so
```

-  No need to update the standard
-  No need to update the compiler
-  No need to break the ABI
-  No need for recompile your
binaries*
-  Might need to update the linker
-  Exceptions become modular

* *-ffunction-sections (or similar) is needed for nearpoint ordering*

Special Thanks

- Jason Turner
- Patrice Roy
- Guy Davidson
- Inbal Levi
- Ben Craig
- Lewis Baker
- Herb Sutter
- Bjarne Stroustrup
- Robin Rowe

Special Thanks

- Jason Turner
- Patrice Roy
- **Guy Davidson**
- Inbal Levi
- Ben Craig
- Lewis Baker
- Herb Sutter
- Bjarne Stroustrup
- Robin Rowe



Possible implementation

```
template<class P, class From, class To = remove_pointer_t<P>>
To *dynamic_cast<From *p> {
    if constexpr (is_same_v<From, To>) {
        return p;
    } else if constexpr (is_base_of_v<To, From>) {
        return (To *)p;
    } else if constexpr (is_void_v<To>) {
        return truly_dynamic_to_nullopt(p);
    } else if constexpr (is_base_of_v<From, To>) {
        return truly_dynamic_from_base_to_derived<To>(p);
    } else {
        return truly_dynamic_between_unrelated_classes<To>(p);
    }
}
```

Technically we should also verify that To is an accessible base of From in the current lexical scope...

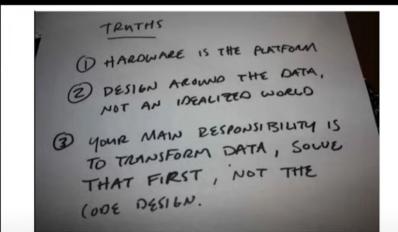


ARTHUR O'DWYER
dynamic_cast
From Scratch

CppCon.org

59

0:53 / 58:48



DATA-ORIENTED DESIGN AND C++
Mike Acton

1:09:54 / 1:27:45



TO EXIT FULL SCREEN PRESS ESC



Andrei Alexandrescu

Robots Are after Your Job
(or At Least the Boring Parts
of It): Exploring Generative
AI for C++

13:10 / 1:33:08

Commonly Taught binary_search

The "violet eyes" version

- First comparison in the middle of the range
- Each pass divides searched space in half
- 1-2 tests per pass (some consider it 1)
- May return early (as soon as element found)
- Not found decision made after range collapsed

```
template <typename I, typename T>
bool binary_search(I b, I e, const T& v) {
    while (b < e) {
        auto m = b + (e - b) / 2;
        if (v == *m) {
            e = m;
        } else if (*m < v) {
            b = m + 1;
        } else {
            return true;
        }
    }
    return false;
}
```

Video Sponsorship Provided By:
ansatz think-cell

CC BY-NC-SA



Cppcon.org

Video Sponsorship Provided By
think-cell



push_back with an $\mathcal{O}(n)$ budget



Budget:

v.push_back(2);

Cost:

0	0	1	2
---	---	---	---

Amortized $O(1)$ Complexity

Andreas Weis



How to support this research?

Consider making a 501(c)(3) eligible donation to my research fund!

San Jose State University Research Fund Name:

Exceptional Software Research Fund

Objective

Research & improve exception performance, size efficiency and tooling.

SJSU Tower Account #034-1300-1248

Want To learn More?

Technical Research Briefings Available!

Interactive sessions where I share my current findings, methodologies, and future roadmap while gathering feedback on your team's specific needs.

What you'll get

- Deep dive into exception performance breakthroughs
- Early insights into the Exception Insights tool development
- Understanding of how your codebase could benefit
- Opportunity to influence the research direction

Questions?

Contact Info

University Email:

khalil.estell@sjsu.edu

Inquiries:

estell.exceptions@gmail.com

Email me for benchmark source

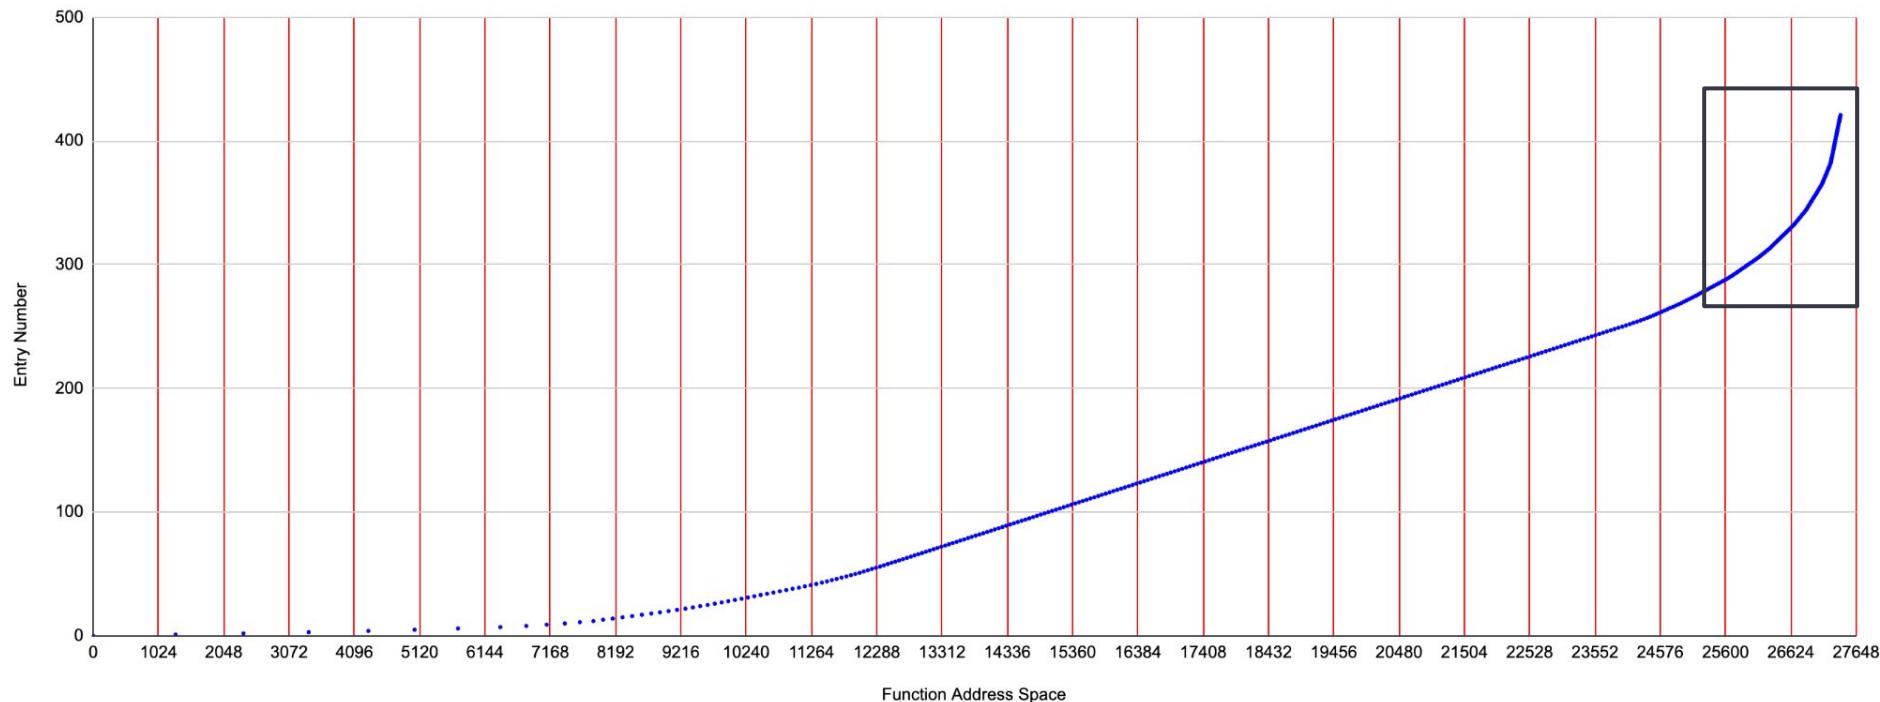


Interested in HALbORD?

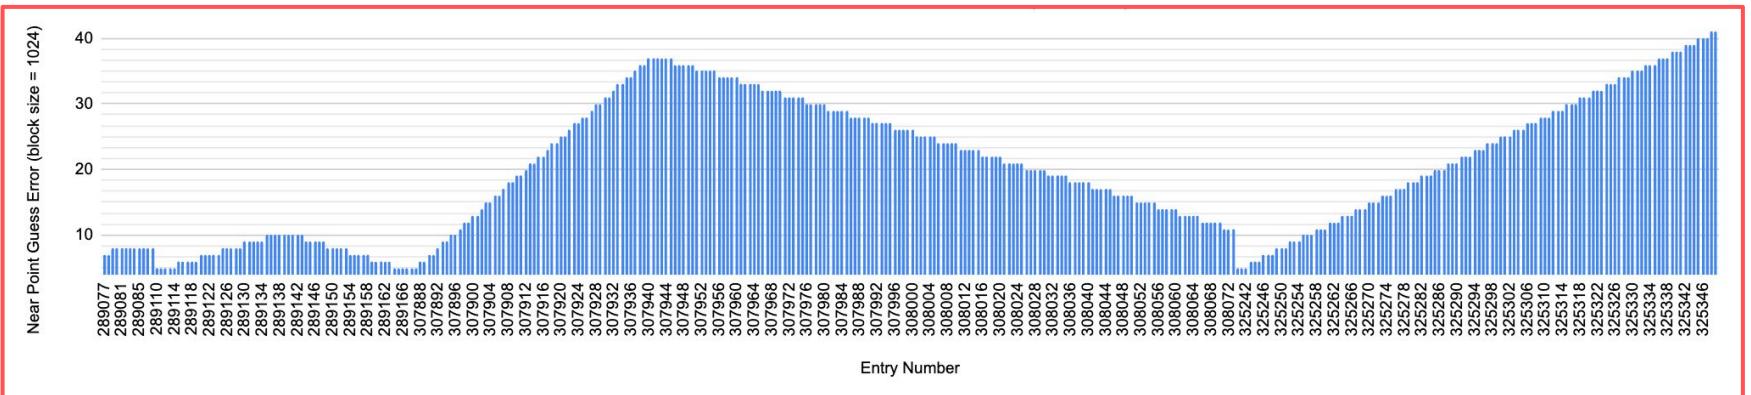
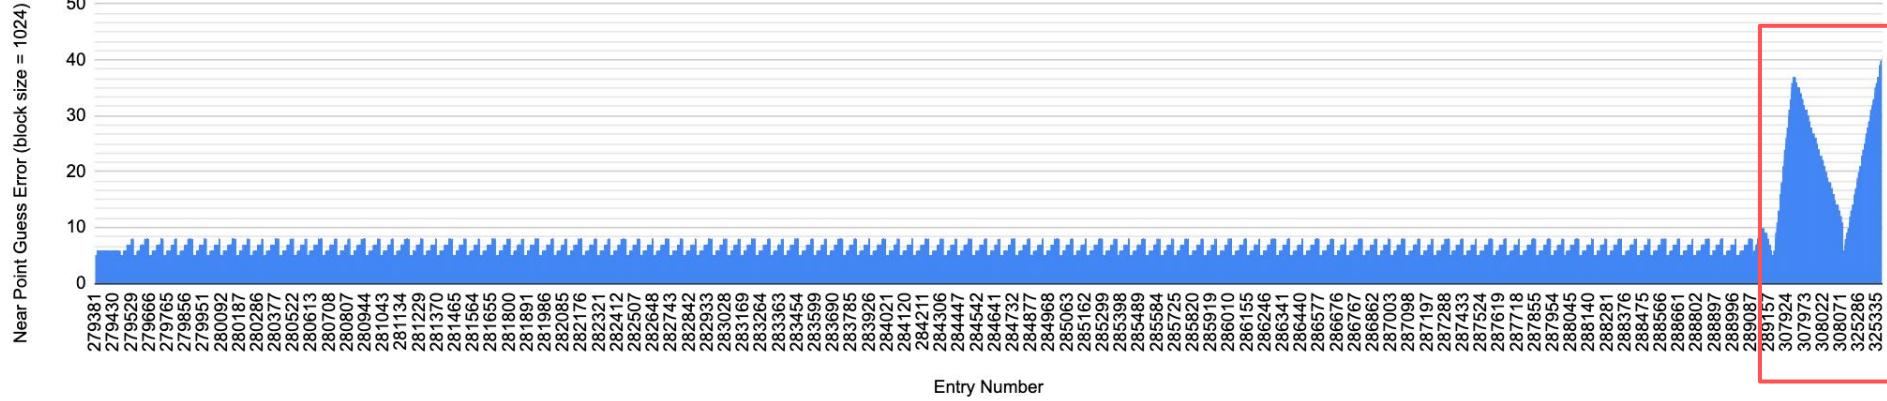
Scan the QR code to learn more!



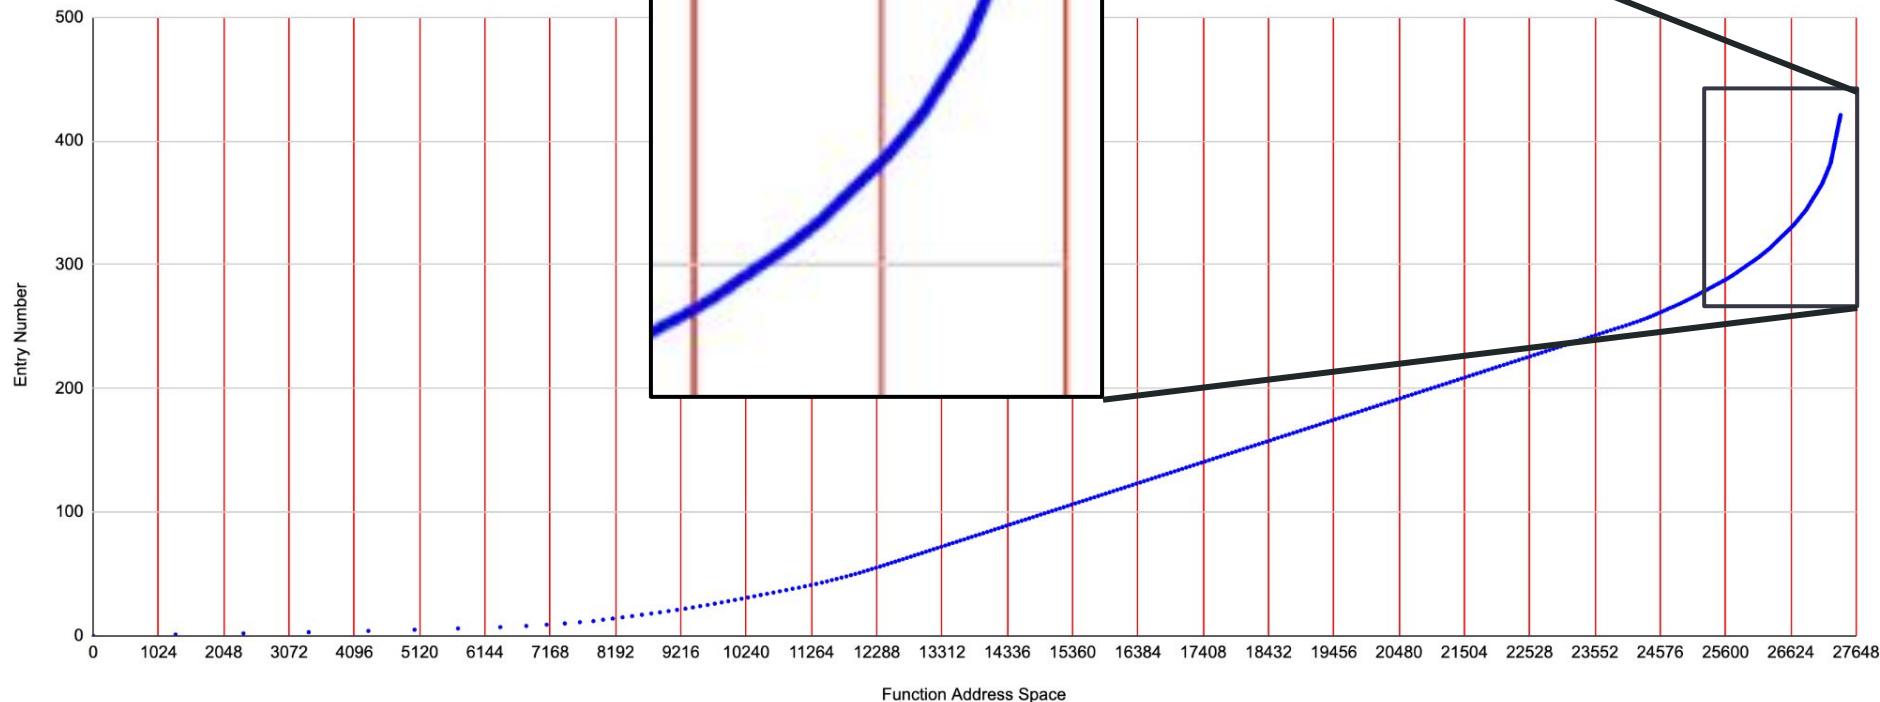
Final Address vs. Index Entry



Nearpoint Simulated Error Testing

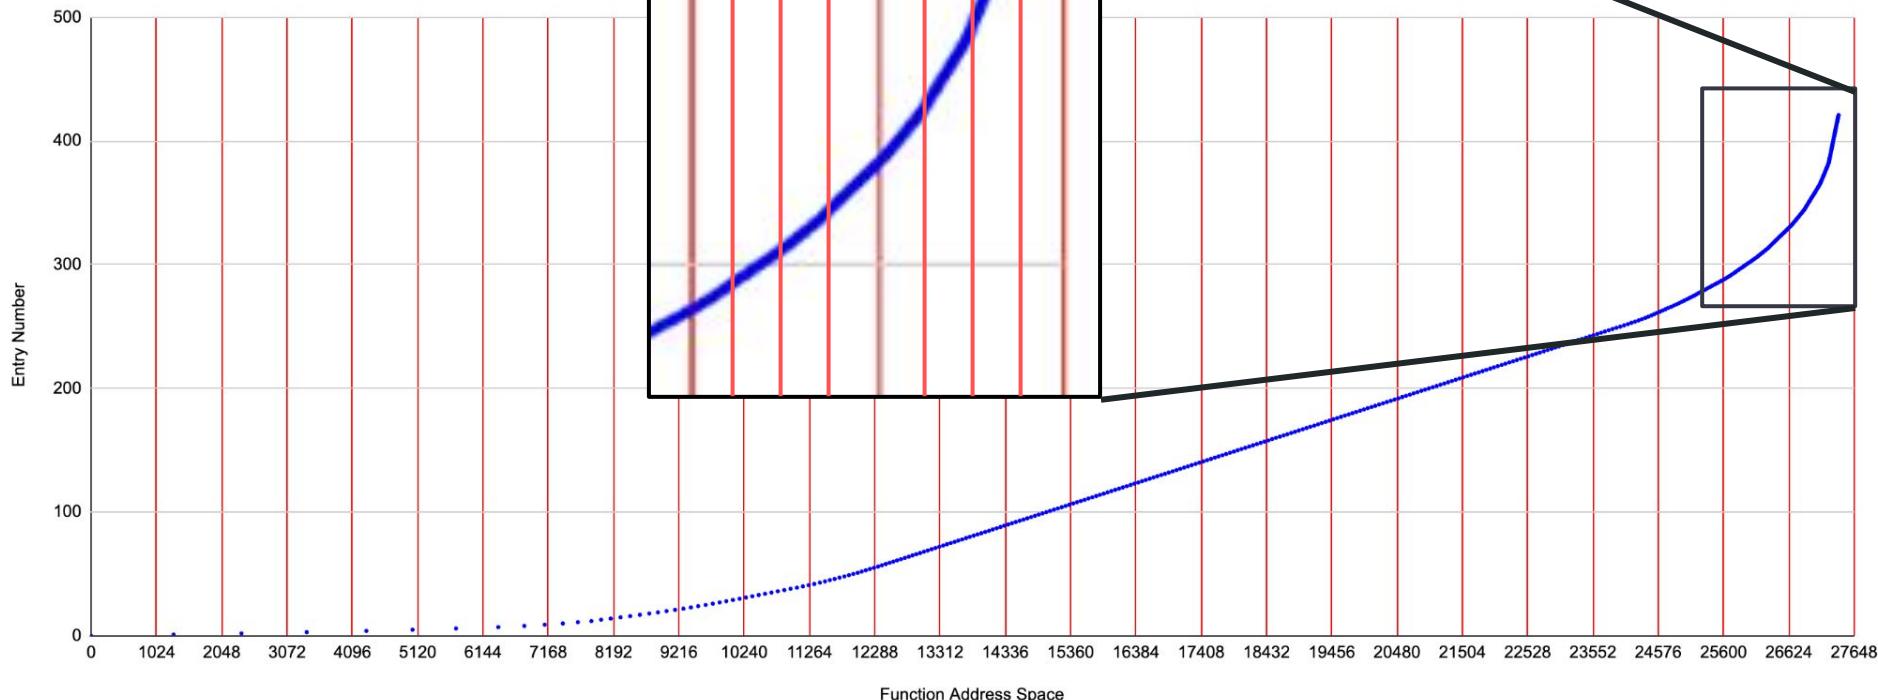


Final Address vs. Index Entry



The fix? **Zoom in!**

Final Address vs. Index Entry



Near Point Linear Approx. Table



```
struct nearpoint_descriptor
{
    std::uint32_t normal_block_size = 0;
    std::uint32_t text_starting_address = 0;
    std::uint32_t small_block_size = 0;
    std::uint32_t small_starting_address = 0;
};
```