



Debugging C++ Coroutines in GDB

ZARTAJ MAJEED



How hard is it to debug C++ coroutines in GDB?

- Debugging is hard
- C++ coroutines are hard
- GDB is hard
- Debugging C++ coroutines in GDB is 3 times as hard

GCC Bug 99215 - Coroutines: debugging with gdb

Open and unassigned since 23 Feb 2021

"I am itching to get into C++20 coroutines (and very grateful for their implementation) but am somewhat put off by the apparent inability to inspect them from within a debugger currently.

While looking for existing related GCC specific issues, discussions or commits (none of which I found) the following paper [Debugging C++ coroutines] did come up:
<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2020/p2073r0.pdf>

This seems to at least confirm the current state that I was seeing

I can not tell if support for this is missing in GCC or GDB or both but I figured I'd try finding out here first."

Is it still really that bad?

No.

Things have improved a lot since the survey in P2073

You can set breakpoints in coroutines, view local variables, view coroutine parameters and examine a coroutine's promise type

Some features of coroutines that are unusual for functions are still difficult to work with in GDB

These include examining suspended coroutines and figuring out async call stacks

A problem Knuth used to motivate coroutines

String Decoding Problem

- From The Art of Computer Programming by Donald Knuth
- A text parsing and transformation problem
- Use coroutines as loosely-coupled functions collaborating to solve a problem

String Decoding Problem

The Art of Computer Programming, Vol 1 & Fascicle 1

Donald Knuth

String Decoding Problem - Input

-- Fascicle 1, TAOCP, p.67

Suppose we want to write a program that translates one code into another. The input code to be translated is a sequence of 8-bit characters terminated by a period, such as

a2b5e3426fg0zyw3210pq89r. (1)

This code appears on the standard input file, interspersed with whitespace characters in an arbitrary fashion. For our purposes a "whitespace character" will be any byte whose value is less than or equal to #20, the ASCII code for ''. All whitespace characters in the input are ignored

String Decoding Problem - Output

-- Fascicle 1, TAOCP, p.67

The other characters should be interpreted as follows, when they are read in sequence:

- (1) If the next character is one of the decimal digits 0 or 1 or ... or 9, say n, it indicates $(n + 1)$ repetitions of the following character, whether the following character is a digit or not.
- (2) A nondigit simply denotes itself.

The output of our program is to consist of the resulting sequence separated into groups of three characters each, until a period appears; the last group may have fewer than three characters. For example,

a2b5e3426fg0zyw3210pq89r. (1)

should be translated into

abb bee eee e44 446 66f gzy w22 220 0pq 999 999 999 r. (2)

String Decoding Problem - Notes

-- Fascicle 1, TAOCP, p.67

Notice that **3426f** does not mean 3427 repetitions of the letter **f**; it means 4 fours and 3 sixes followed by **f**.

If the input sequence is '**1 .**', the output is simply '.', not '..', because the first period terminates the output.

The goal of our program is to produce a sequence of lines on the standard output file, with 16 three-character groups per line (except, of course, that the final line might be shorter).

The three-character groups should be separated by blank spaces, and each line should end as usual with the ASCII newline character **#a**.

String Decoding Problem

Original TAOCP Solution with Coroutines

What is a coroutine to Knuth?

-- Fascicle 1, TAOCP, pp.66-67

Subroutines are special cases of more general program components, called coroutines. In contrast to the unsymmetric relationship between a main routine and a subroutine, there is complete symmetry between coroutines, which call on each other.

A subroutine is always initiated *at its beginning*, which is usually a fixed place; a coroutine is always initiated *at the place following* where it last terminated.

Such coroutine linkage is easy to achieve with MMIX if we set aside two global registers, **a** and **b**. In coroutine **A**, the instruction **GO a, b, 0** is used to activate coroutine **B**; in coroutine **B**, the instruction **GO b, a, 0** is used to activate coroutine **A**.

String decoding - A solution using coroutines

Split problem into three functions along input, output and processing of single character items

- **NextChar**, to manage the raw input buffer
- **In**, to parse input and provide a character item for further processing
- **Out**, to format a character item and print output

In and **Out** are coroutines working in tandem

In is a generator that yields one character

Out fills a group of three characters awaiting each in sequence from **In**

NextChar is a subroutine called by **In**

String Decoding Problem

C++ Coroutines Port of TAOCP MMIX Solution

In() coroutine

```
InR0 In() {  
    char inchr;  
    int count;  
  
    for(;;) {  
        inchr = NextChar();  
        if(inchr > '9') {  
            co_yield inchr;  
            continue;  
        }  
        count = inchr - '0';  
        if(count < 0) {  
            co_yield inchr;  
            continue;  
        }  
        inchr = NextChar();  
        co_yield inchr;  
        for(--count; count >= 0; --count) {  
            co_yield inchr;  
        }  
    }  
}
```

NextChar function

```
char NextChar() {  
    static char InBuf[1000];  
    static char* inptr = InBuf;  
    char inchar;  
  
    for(;;) {  
        if(*inptr == '\0') {  
            inptr = InBuf;  
            if(fgets(InBuf, sizeof InBuf,  
                     stdin) == NULL)  
                *inptr = '.';  
        }  
        inchar = *inptr;  
        ++inptr;  
        if(!isspace(inchar) &&  
            !iscntrl(inchar))  
            break;  
    }  
    return inchar;  
}
```

Out() coroutine, 1/3

```
OutR0 Out() {  
  
    char OutBuf[] = {  
        0, 0, 0, ' ',  
        /* group repeated 14 times */,  
        0, 0, 0, '\n',  
        0  
    };  
    char* outptr = OutBuf;  
  
    InR0&& in = In();  
  
    for(;;) {  
  
        char outchr = co_await in;  
        outptr[0] = outchr;  
  
        if(outchr == '.') {  
            finishLine(OutBuf,  
                (uint64_t)(outptr + 3 - OutBuf));  
            break;  
        }  
    }  
}
```

Out() coroutine, 2/3

```
outchr = co_await in;
outptr[1] = outchr;

if(outchr == '.') {
    ++outptr;
    finishLine(OutBuf,
        (uint64_t)(outptr + 3 - OutBuf));
    break;
}
```

```
outchr = co_await in;
outptr[2] = outchr;

if(outchr == '.') {
    outptr += 2;
    finishLine(OutBuf,
        (uint64_t)(outptr + 3 - OutBuf));
    break;
}
```

Out() coroutine, 3/3

```
outptr += 4;

if(outptr != OutBuf + sizeof(OutBuf) - 1) {
    continue;
}

fputs(OutBuf, stdout);
outptr = OutBuf;
}
```

```
void finishLine(char* line, uint64_t len) {
    line[len - 2] = '\n';
    line[len - 1] = '\0';
    fputs(line, stdout);
}
```

First among equals?

- Unlike subroutines, coroutines don't have hierarchical caller-callee relationship
- Which coroutine should run first and what should its initial state be?
- Better to start with **Out** for decoding since it needs to await a character from **In**
- Run **Out** right away to suspend at first **co_await**
- Start **In** suspended and resume when **Out** needs a value

Program start - Decoding problem

- Call **Out** coroutine to kick off processing
- Program does not exit till **Out** coroutine suspends or returns and control is transferred to **main**

```
int main() {  
    Out();  
}
```

String Decoding Problem

Compiler Machinery for C++ Coroutines Solution

In coroutine - Return object - Compiler machinery

- **InRO** is return type of **In** coroutine
- Forward declare nested **promise_type**
- Has **coroutine_handle** to In coroutine frame
- Constructor runs when compiler calls **get_return_object** method of **promise_type**
- Destructor destroys coroutine frame

```
struct InRO {  
    struct promise_type;  
    coroutine_handle<promise_type> coro;  
  
    InRO(coroutine_handle<promise_type> h)  
        : coro(h) {}  
  
    ~InRO() {  
        coro.destroy();  
    }  
    // omitted ...
```

In coroutine - Promise type - Compiler machinery

- No need for `coroutine_traits` if promise type is inside return object
- Compiler calls `get_return_object` to construct return object instance using coroutine handle
- `return_void` required for coroutine that does not `co_return` a value
- `unhandled_exception` also required
 - ours simply terminates program

```
// struct InR0 {  
struct promise_type {  
    InR0 get_return_object() {  
        return InR0{  
            coroutine_handle<promise_type>::  
            from_promise(*this)};  
    }  
  
    void return_void() {}  
    void unhandled_exception() noexcept {  
        terminate();  
    }  
};  
// omitted ...  
};
```

In coroutine - Initial/Final suspend - Compiler machinery

- Previously decided to suspend **In** coroutine at start
- **In** never actually breaks out of its main loop
- So **final_suspend** is never called and it can return any **awaitable**
- We arbitrarily return **suspend_always**

```
// struct promise_type {  
  
    auto initial_suspend() noexcept {  
        return suspend_always{};  
    }  
  
    auto final_suspend() noexcept {  
        return suspend_always{};  
    }  
  
};
```

In coroutine - co_yield handler - Compiler machinery

- `yield_value` is called with `co_yield` argument
- Allows **In** promise to update `val` member that can be retrieved later through promise
- Returns `awaitable` because `co_yield` is just a form of `co_await`
- We suspend **In** coroutine to return control to where **In** was resumed

```
// struct promise_type {  
    char val;  
  
    auto yield_value(char x) {  
        val = x;  
        return suspend_always{};  
    }  
};
```

In coroutine - Make InRO awaitable - Compiler machinery

- The return object of **In()** is made **awaitable** by adding **await_ready**, **await_suspend** and **await_resume** methods
- **In() await_ready** returns **true** because **In()** is always ready to return next character
- Simplest **void** form of **await_suspend** is used since it's never called

```
// struct InRO {  
  
    bool await_ready() const noexcept {  
        return true;  
    }  
  
    void await_suspend(coroutine_handle<>) {}  
  
};
```

In coroutine - Orchestrate awaitable - Compiler machinery

- `InR0 await_resume` resumes its own `In()` coroutine
- This causes `In()` to run till it yields
 - Recall `In()` promise `yield_value` sets `val` member of promise then suspends
- Then `await_resume` can return promise `val` member to awaiting coroutine

```
// struct InR0 {  
  
    char await_resume() const noexcept {  
        coro.resume();  
        return coro.promise().val;  
    }  
  
};
```

Out coroutine - Return object - Compiler machinery

- **OutRO** is return type of **Out()** coroutine
- Forward declare nested **promise_type**
- Has **coroutine_handle** to **Out()** coroutine frame
- Constructor runs when compiler calls **get_return_object** method of **promise_type**
- Destructor destroys coroutine frame

```
struct OutRO {  
    struct promise_type;  
    coroutine_handle<promise_type> coro;  
  
    OutRO(coroutine_handle<promise_type> h)  
        : coro(h) {}  
  
    ~OutRO() {  
        coro.destroy();  
    }  
    // omitted ...
```

Out coroutine - Promise type - Compiler machinery

- Easier to define promise type inside return object than using **coroutine_traits**
- Compiler calls **get_return_object** to construct return object instance with coroutine handle
- **return_void** required for coroutine that does not **co_return** a value
- **unhandled_exception** also required
 - ours simply terminates program

```
// struct OutRO {  
struct promise_type {  
    OutRO get_return_object() {  
        return OutRO{  
            coroutine_handle<promise_type>::  
            from_promise(*this)};  
    }  
  
    void return_void() {}  
    void unhandled_exception() noexcept {  
        terminate();  
    }  
};  
// omitted ...  
};
```

Out coroutine - Initial/Final suspend - Compiler machinery

- Previously decided to immediately run `Out()` coroutine at start
- Want `final_suspend` to suspend `Out()` to let its return object destructor destroy coroutine frame

```
// struct promise_type {  
  
    auto initial_suspend() noexcept {  
        return suspend_never{};  
    }  
  
    auto final_suspend() noexcept {  
        return suspend_always{};  
    }  
  
};
```

Demo