

25

From Pure ISO C++20 To Compute Shaders

KOEN SAMYN



Cppcon
The C++ Conference

20
25



September 13 - 19

About me



- Koen Samyn
- Program Coordinator Game Development at DAE



<https://digitalartsandentertainment.be/>

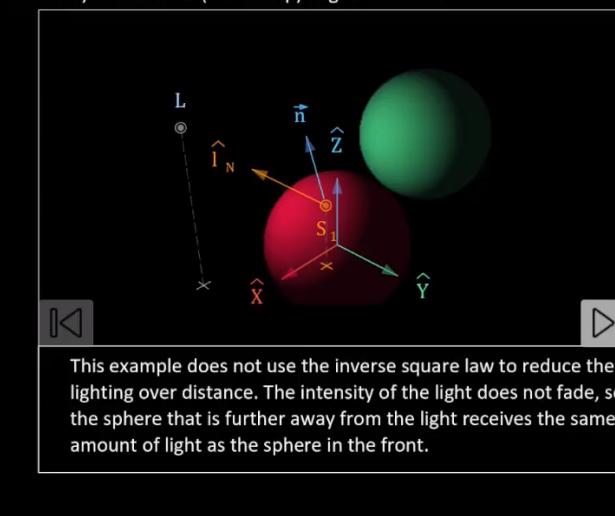
About me



- Koen Samyn
- Program Coordinator Game Development at DAE

Diffuse lighting scene

In the following example, the distance is not taking into account. For each point on the sphere, the **dot product** of the normal \hat{n} at that point is taken with the light vector \hat{l}_N . It is possible that this dot product is negative, so when calculating lighting (or when implementing them in shaders), it is necessary to check for (and clamp) negative values.



Calculation of diffuse intensity

For each pixel on the screen you calculate the light vector from that pixel in **world space** to the light source:

$$\hat{l} = L - S_1$$

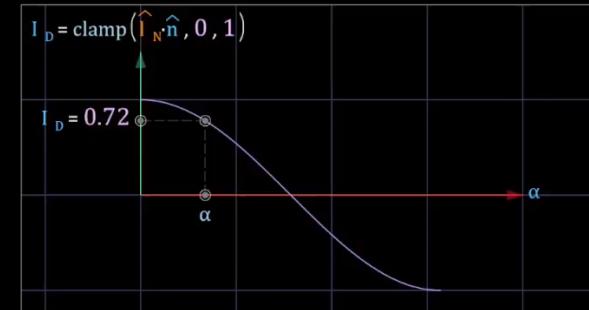
For lighting calculations it is important to work with normalized vectors:

$$\hat{l}_N = \frac{\hat{l}}{\|\hat{l}\|}$$

The diffuse intensity is the the dot product of the normal of the sphere at the pixel in world space with the normalize light vector:

$$I_D = \text{clamp}(\hat{l}_N \cdot \hat{n}, 0, 1)$$

The dot product of two normalized vectors is the cosine function of the angle between the two vectors:



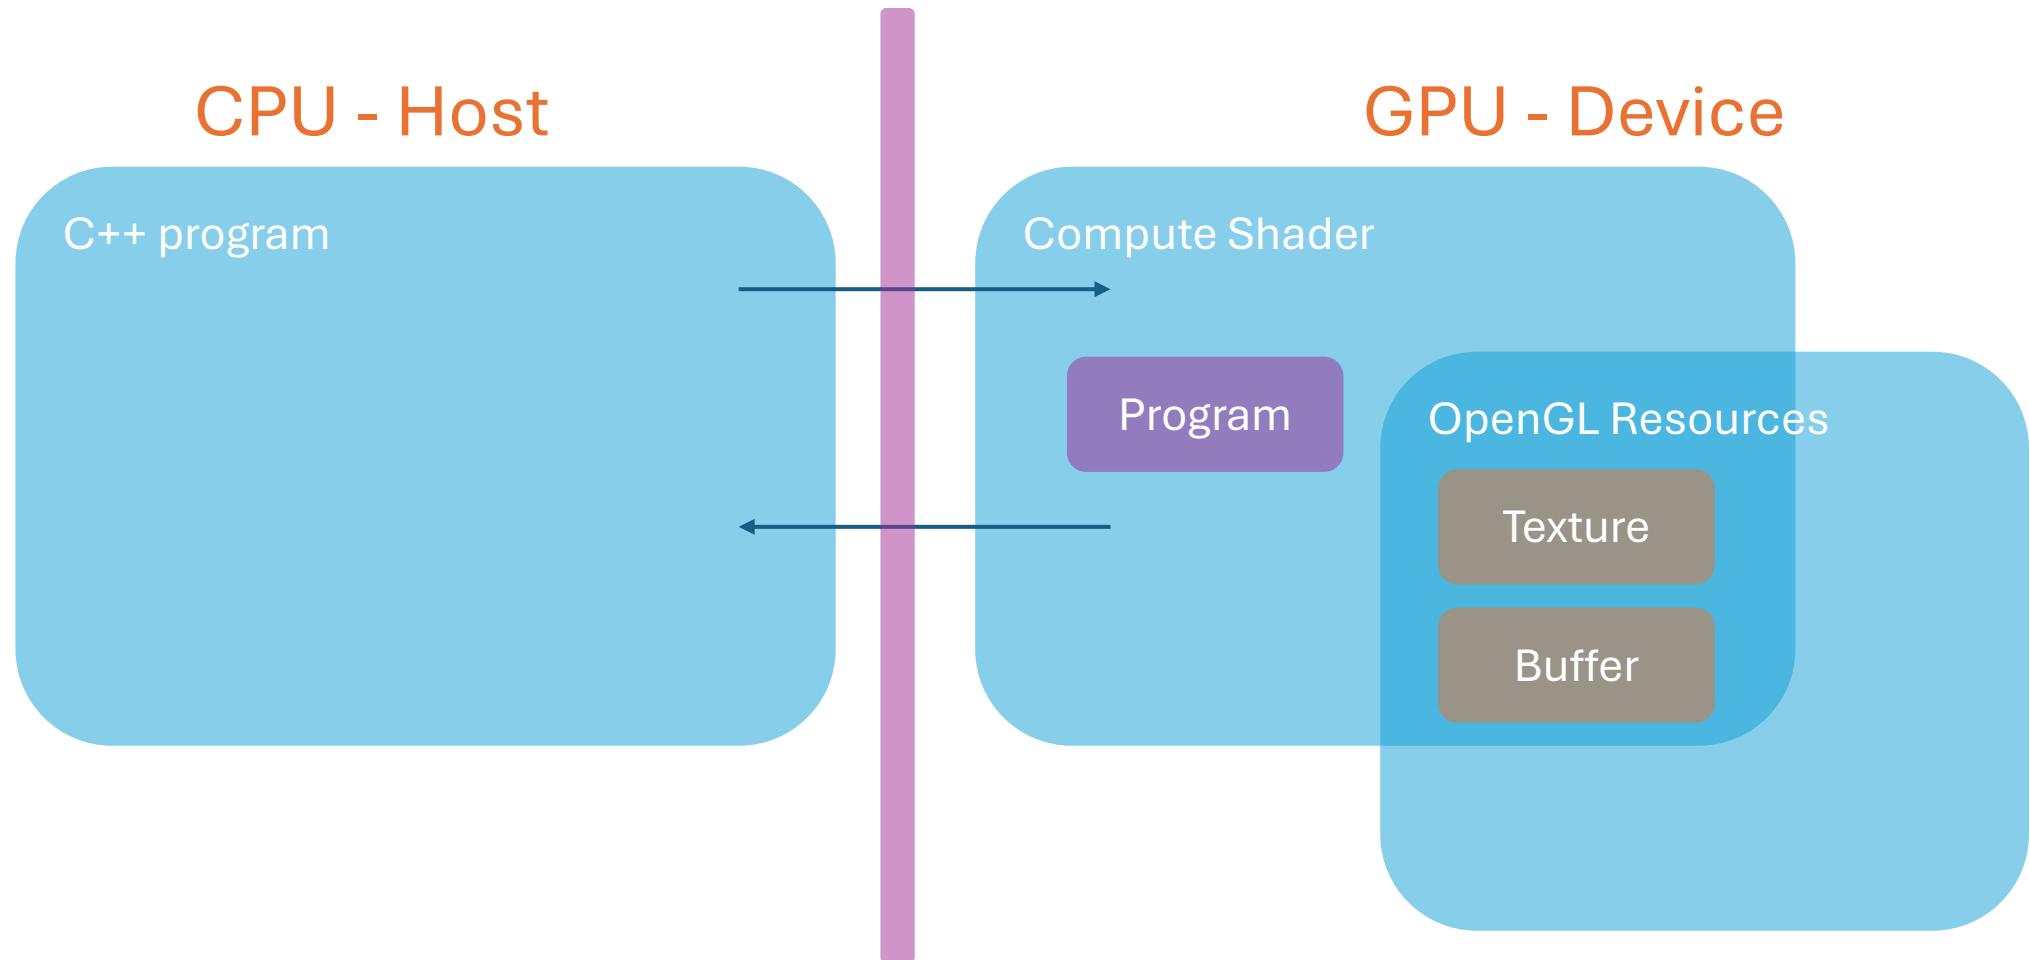
Why this project?

- Mental load to develop compute shaders is heavy
 - Still developing C++ skills.
 - Management layer needed for GPU resources.
 - Debugging GPU programs is hard.
- Programming model is different
 - Workgroups / threads
 - Homogeneous parallelism: every thread performs the same task
 - Existing C++ code can be hard to convert

Is this new?

- A lot of movement in this area and existing frameworks
 - CUDA /HIP
 - SyCL
 - CompuShady (Python)
 - Slang
 - OpenMP
- Niche I am trying to fill:
 - Non-proprietary
 - Lightweight
 - **Educational focus:** teach transferable skills.
 - Focus on techniques that are transferable to graphics programming shaders: vertex/fragment/mesh shaders

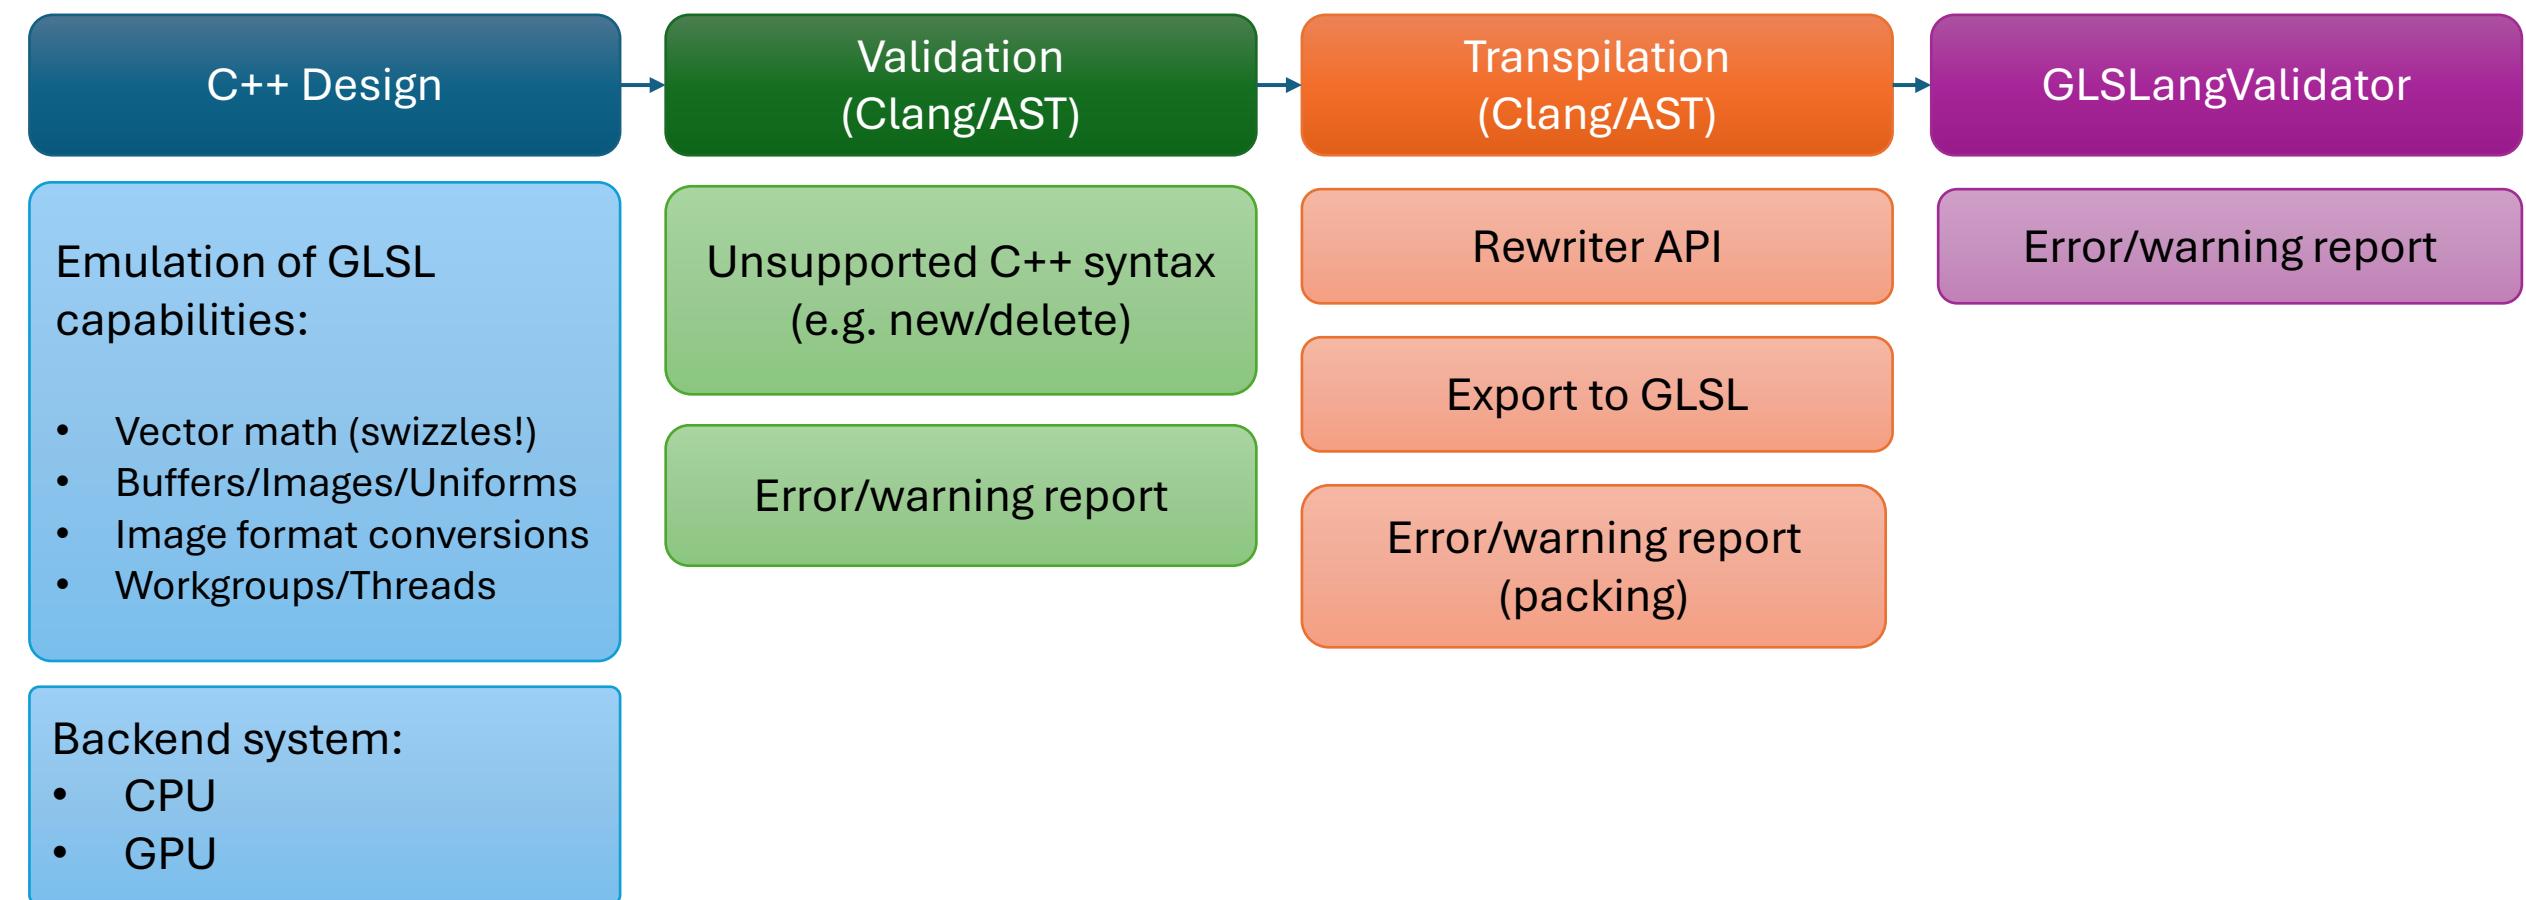
What is a compute shader?



Design goals

- Start with C++20 → no new keywords
- Compact code
- Close mapping with GPU parallelism
 - Workgroups/threads
- Equivalence between CPU and GPU execution

Overview



Cases

Adding two vectors
of floats

1	2	3	4	5	...
---	---	---	---	---	-----

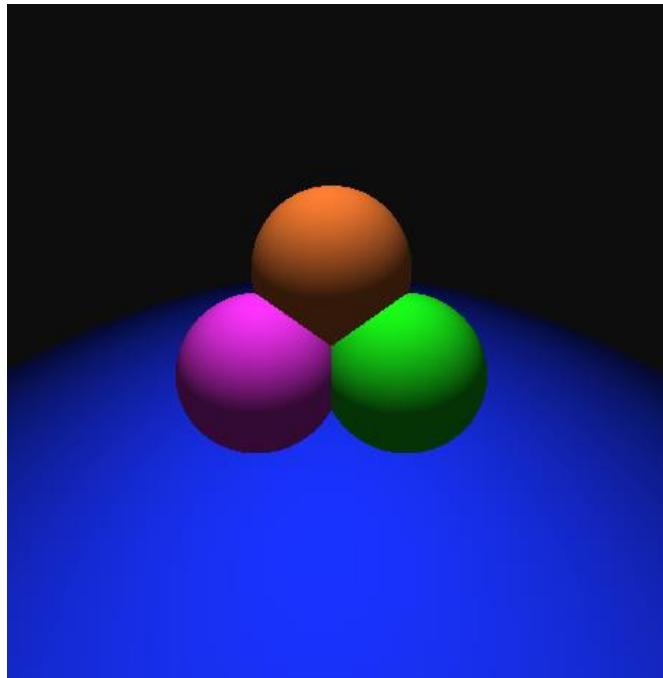
+

7	8	9	10	11	...
---	---	---	----	----	-----

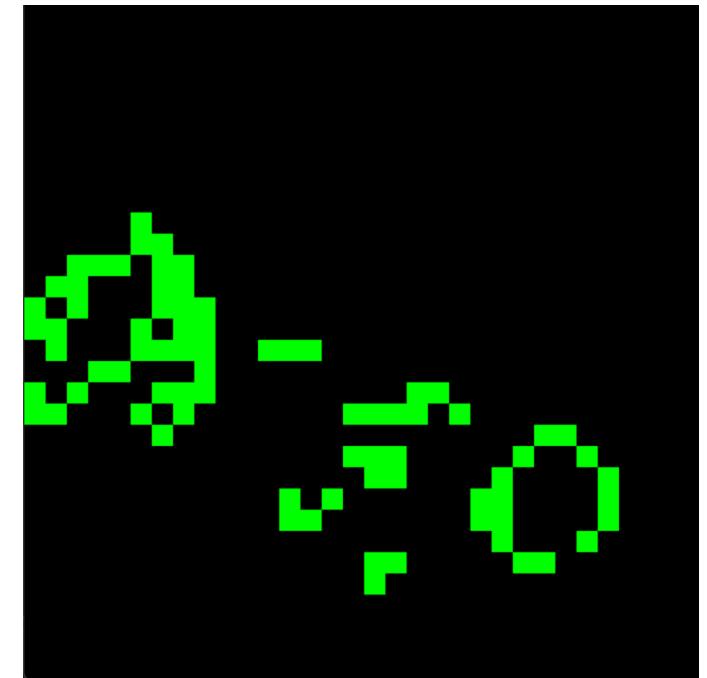
=

8	10	12	14	16	...
---	----	----	----	----	-----

Raytracing



Game of life



Case 1: Adding floats – Compute Shader

```
1 #version 430  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15
```

Compute shader compatible with
OpenGL version 4.3

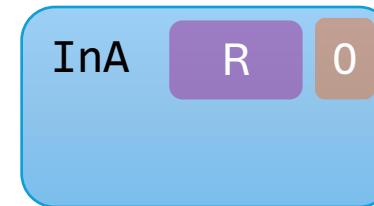
Case 1: Adding floats – Compute Shader

```
1 #version 430  
2 layout(local_size_x = 256) in;  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15
```

Workgroup size → defines the number of threads that will be created for one **workgroup**.

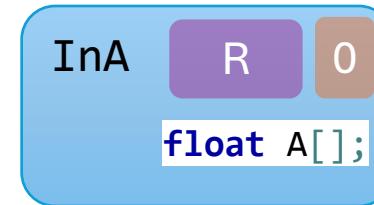
Case 1: Adding floats – Compute Shader

```
1 #version 430
2
3 layout(local_size_x = 256) in;
4
5 layout( binding = 0 ) readonly buffer InA {
6
7
8
9
10
11
12
13
14
15 }
```



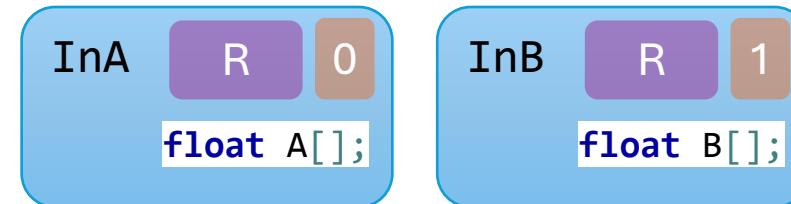
Case 1: Adding floats – Compute Shader

```
1 #version 430
2
3 layout(local_size_x = 256) in;
4
5 layout( binding = 0 ) readonly buffer InA {
6     float A[];
7 }
8
9
10
11
12
13
14
15
```



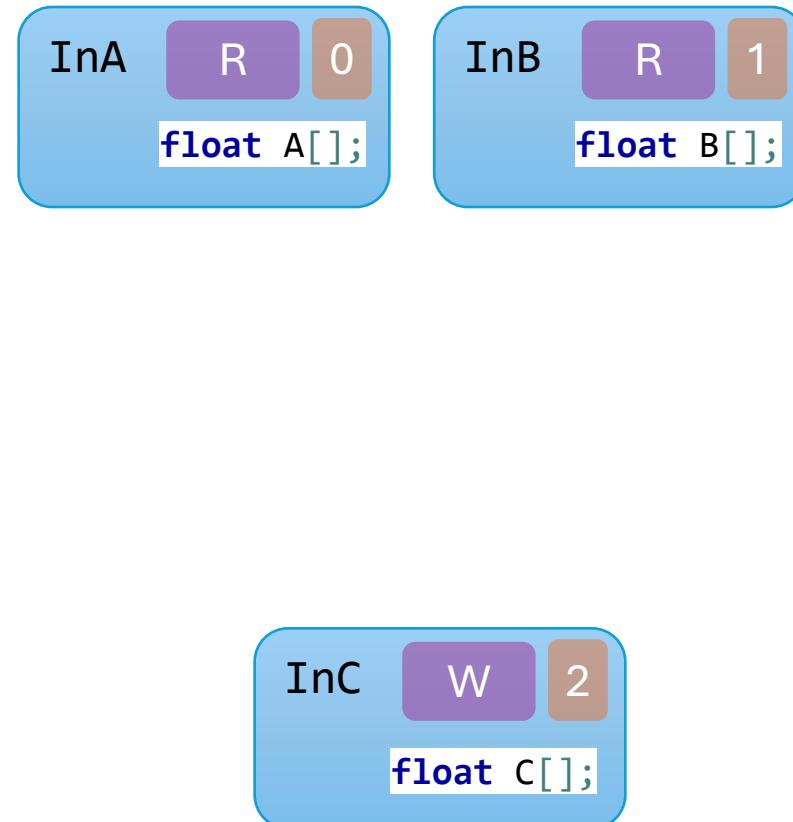
Case 1: Adding floats – Compute Shader

```
1 #version 430
2
3 layout(local_size_x = 256) in;
4
5 layout( binding = 0 ) readonly buffer InA {
6     float A[];
7 }
8
9
10
11
12
13
14
15
```



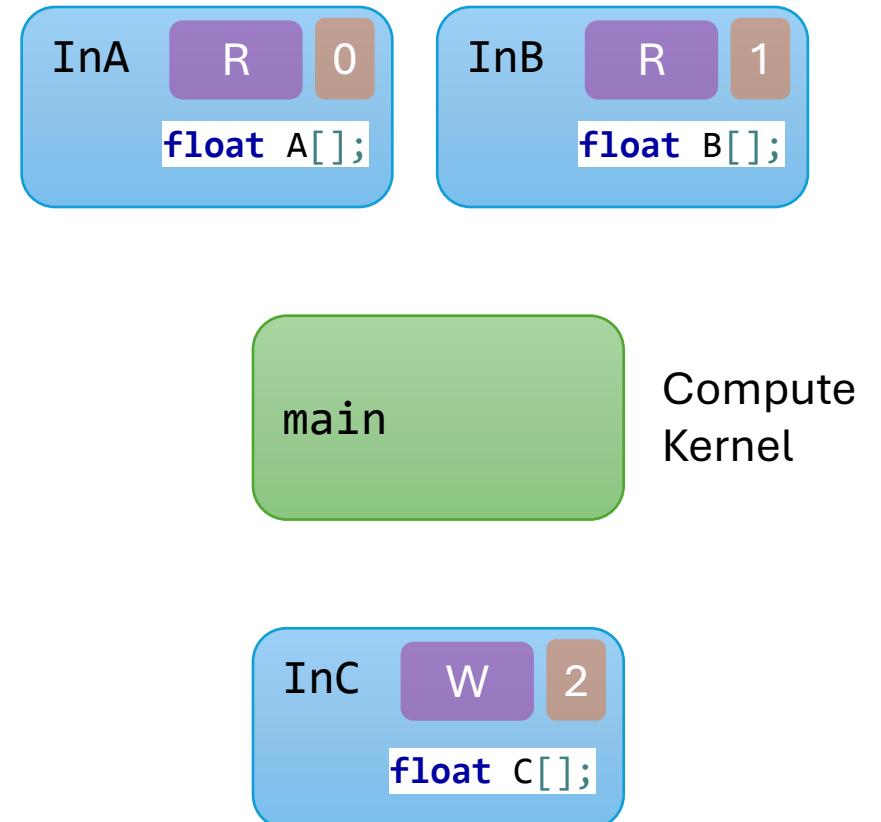
Case 1: Adding floats – Compute Shader

```
1 #version 430
2
3 layout(local_size_x = 256) in;
4
5 layout( binding = 0 ) readonly buffer InA {
6     float A[];
7 }
8
9 layout( binding = 1 ) readonly buffer InB {
10    float B[];
11 }
12
13 layout( binding = 2 ) writeonly buffer OutC
14 {   float C[];
15 }
```



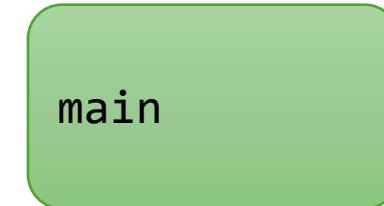
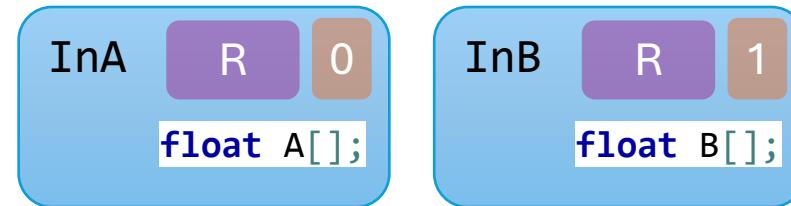
Case 1: Adding floats – Compute Shader

```
1 #version 430
2
3 layout(local_size_x = 256) in;
4
5 layout( binding = 0 ) readonly buffer InA {
6     float A[];
7 }
8
9 layout( binding = 1 ) readonly buffer InB {
10    float B[];
11 }
12
13 void main() {
14 }
```

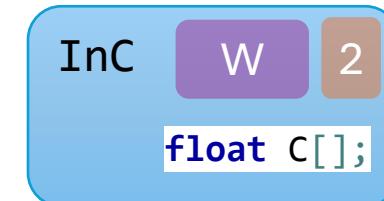


Case 1: Adding floats – Compute Shader

```
1 #version 430
2
3 layout(local_size_x = 256) in;
4
5 layout( binding = 0 ) readonly buffer InA {
6     float A[];
7 }
8
9 layout( binding = 1 ) readonly buffer InB {
10    float B[];
11 }
12
13 void main() {
14     uint i = gl_GlobalInvocationID.x;
15 }
```

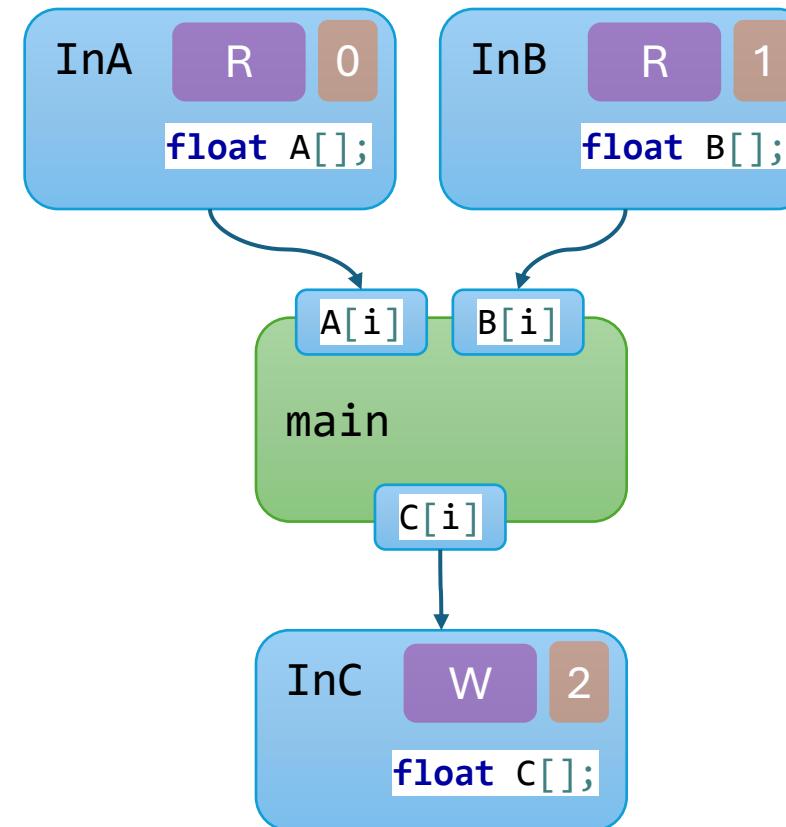


Thread ID: can be used as **index** into arrays A,B and C



Case 1: Adding floats – Compute Shader

```
1 #version 430
2
3 layout(local_size_x = 256) in;
4
5 layout( binding = 0 ) readonly buffer InA {
6     float A[];
7 };
8
9 layout( binding = 1 ) readonly buffer InB {
10    float B[];
11 };
12
13 layout( binding = 2 ) writeonly buffer OutC {
14     float C[];
15 };
16
17 void main() {
18     uint i = gl_GlobalInvocationID.x;
19     C[i] = A[i] + B[i];
20 }
```



Case 1: C++ equivalent

```
1 #include "vec.hpp"
2
3 struct [[clang::annotate("kernel")]] FloatAdder
4 {
5
6
7
8
9
10
11
12
13
14 };
```

vec.hpp : for vector types: **uvecn**, **ivecn**,
vecn, **dvecn**, **nvecn**

Kernel annotation: clang annotation to
find kernels in the AST (see further)

Case 1: C++ equivalent

```
1 #include "vec.hpp"
2
3 struct [[clang::annotate("kernel")]] FloatAdder
4 {
5     static constexpr char fileLocation[] = "floatadder";
6
7
8
9
10
11
12
13
14 };
```

Where to store the generated glsl file.

Case 1: C++ equivalent

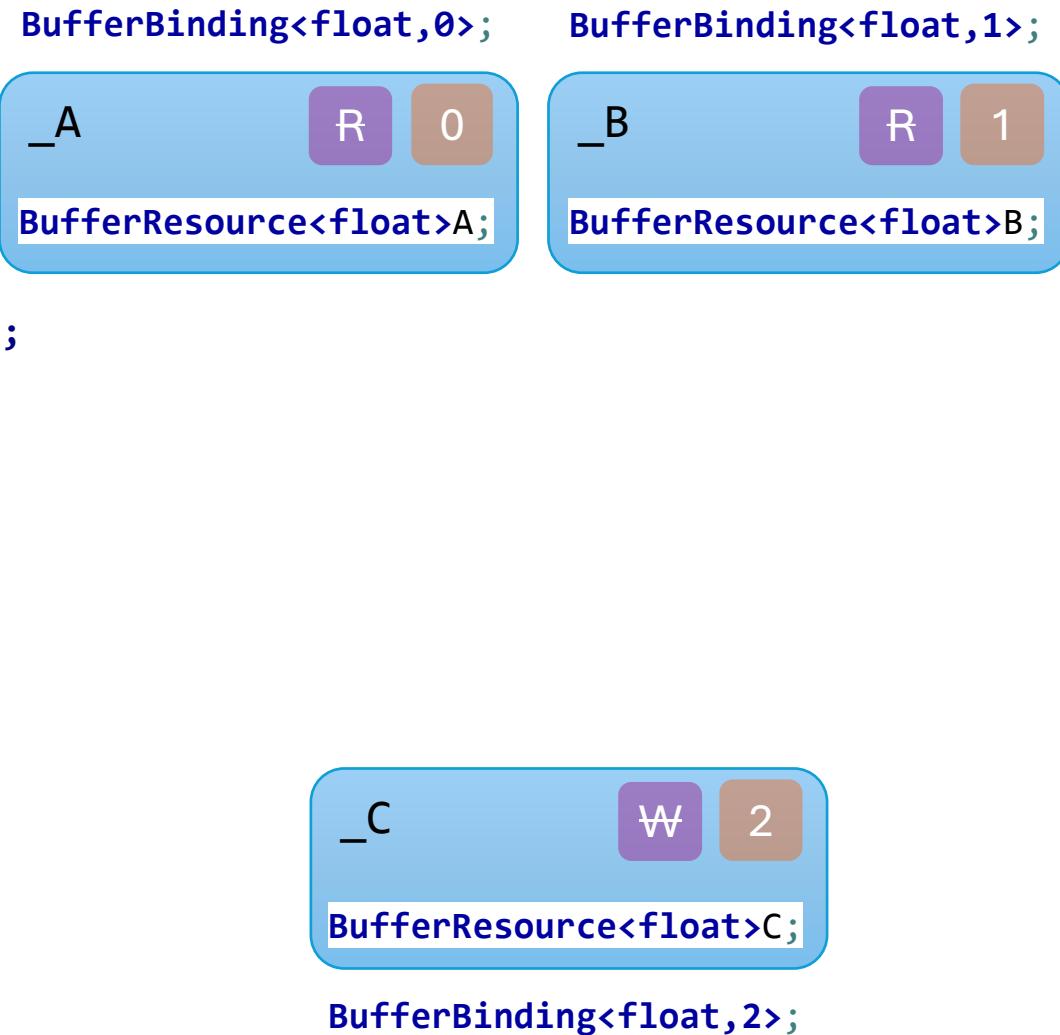
```
1 #include "vec.hpp"
2
3 struct [[clang::annotate("kernel")]] FloatAdder
4 {
5     static constexpr char fileLocation[] = "floatadder";
6     uvec3 local_size{ 256, 1, 1 };
7
8
9
10
11
12
13
14 };
```

Corresponds with :

```
layout(local_size_x = 256) in;
```

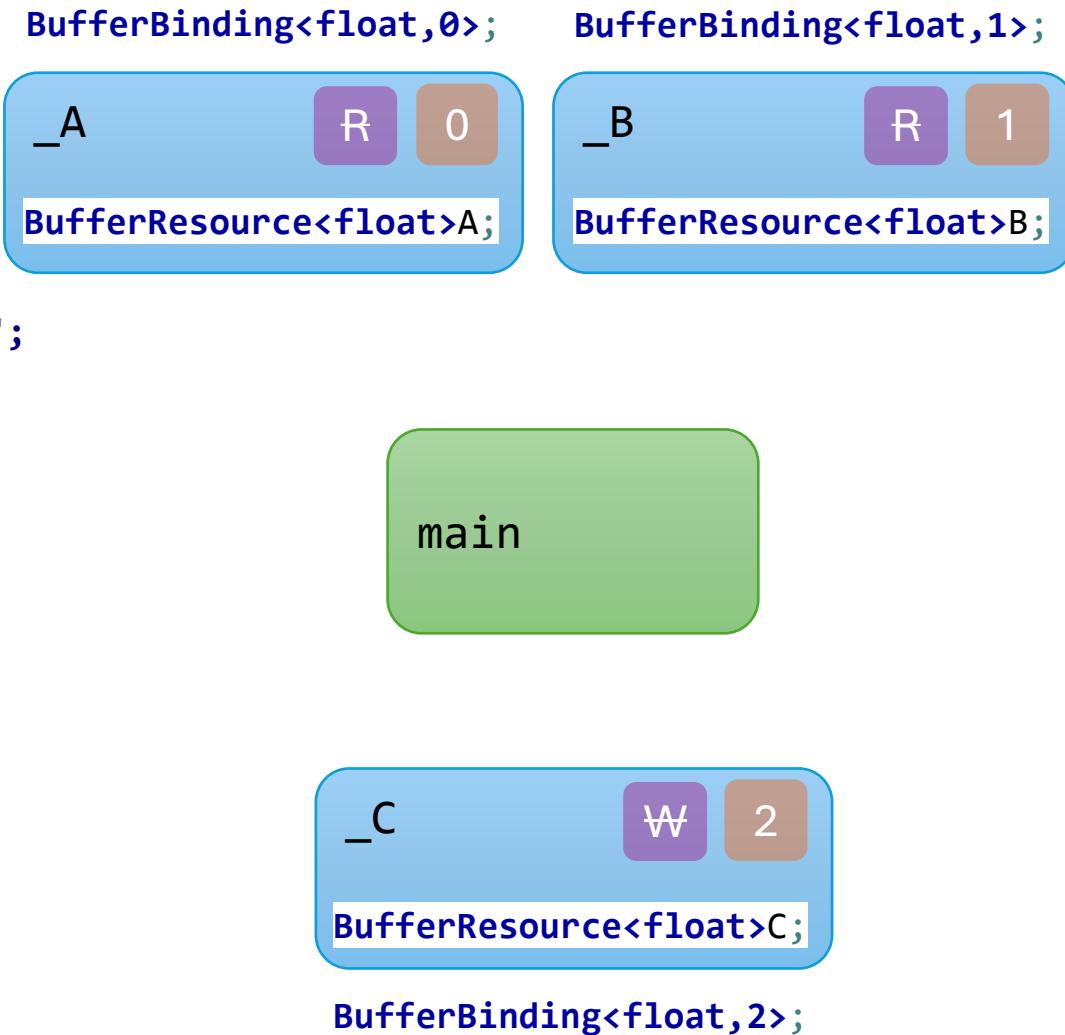
Case 1: C++ equivalent

```
1 #include "vec.hpp"
2 #include "computebackend.hpp"
3 struct [[clang::annotate("kernel")]] FloatAdder
4 {
5     static constexpr char fileLocation[] = "floatadder";
6     uvec3 local_size{ 256, 1, 1 };
7     BufferBinding< float , 0> A;
8     BufferBinding< float , 1> B;
9     BufferBinding< float , 2> C;
10
11
12
13
14 };
```



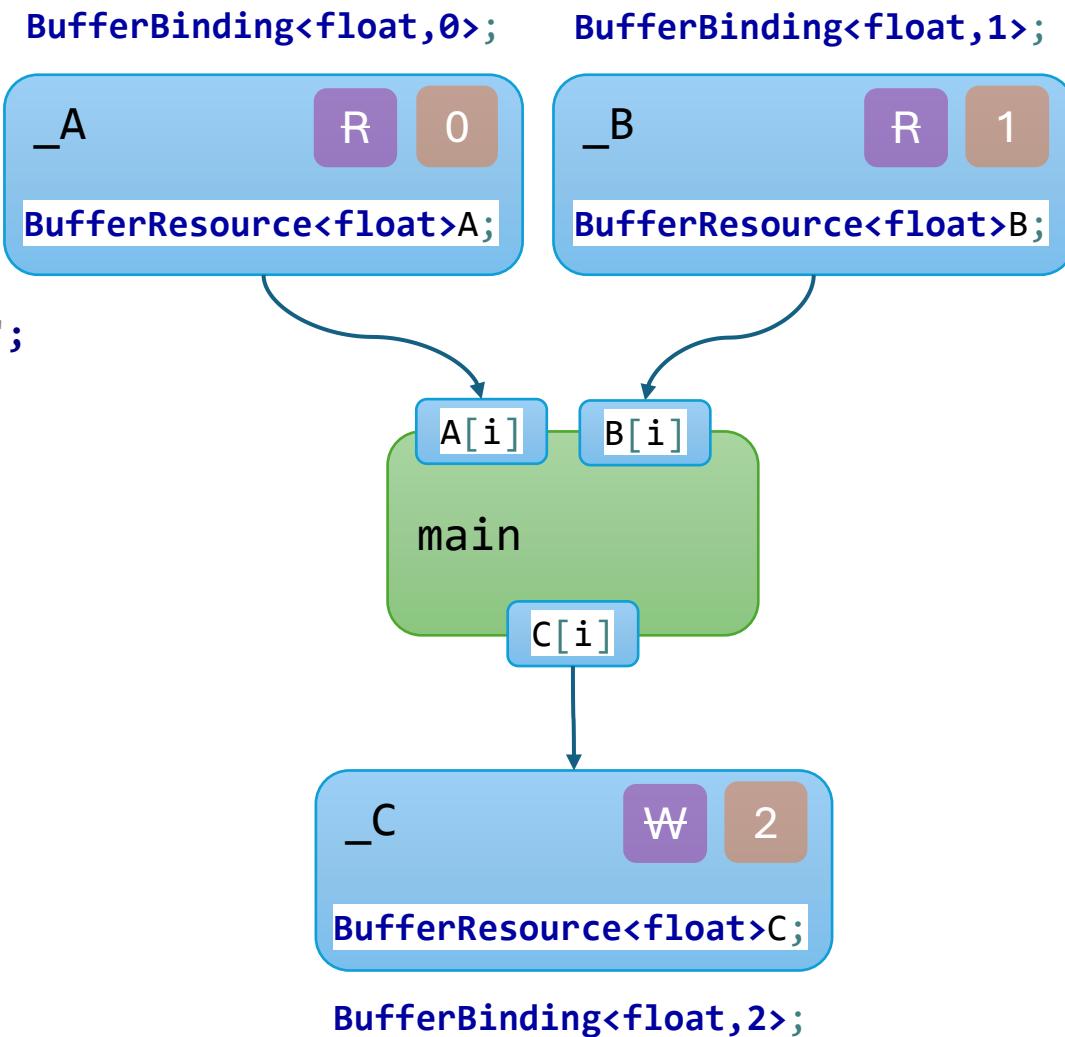
Case 1: C++ equivalent

```
1 #include "vec.hpp"
2 #include "computebackend.hpp"
3 struct [[clang::annotate("kernel")]] FloatAdder
4 {
5     static constexpr char fileLocation[] = "floatadder";
6     uvec3 local_size{ 256, 1, 1 };
7     BufferBinding< float , 0> A;
8     BufferBinding< float , 1> B;
9     BufferBinding< float , 2> C;
10    void main() {
11
12    }
13
14};
```



Case 1: C++ equivalent

```
1 #include "vec.hpp"
2 #include "computebackend.hpp"
3 struct [[clang::annotate("kernel")]] FloatAdder
4 {
5     static constexpr char fileLocation[] = "floatadder";
6     uvec3 local_size{ 256, 1, 1 };
7     BufferBinding< float , 0> A;
8     BufferBinding< float , 1> B;
9     BufferBinding< float , 2> C;
10    void main() {
11        tc::uint i = tc::gl_GlobalInvocationID.x;
12        C[i] = A[i] + B[i];
13    }
14};
```



Case 1 : Validation

- Clang 22.0 : *matchers* and *visitors*
- Custom C++ tool called via CMake

```
-CXXRecordDecl 0x26d24bde080 <line:24:1, line:104:1> line:24:38 struct GameOfLifeKernel definition
|-DefinitionData pass_in_registers aggregate standard_layout trivially_copyable has_constexpr_non_copy_move_ctor
| |-DefaultConstructor exists non_trivial constexpr needs_implicit defaulted_is_constexpr
| |-CopyConstructor simple trivial has_const_param needs_implicit implicit_has_const_param
| |-MoveConstructor exists simple trivial needs_implicit
| |-CopyAssignment simple trivial has_const_param needs_implicit implicit_has_const_param
| |-MoveAssignment exists simple trivial needs_implicit
`-Destructor simple irrelevant trivial constexpr needs_implicit
-AnnotateAttr 0x26d24bde1b8 <col:10, col:34> "kernel"
-CXXRecordDecl 0x26d24bde258 <col:1, col:38> col:38 implicit struct GameOfLifeKernel
-VarDecl 0x26d24bde358 <line:26:2, col:41> col:24 fileLocation 'const char[11]' static inline constexpr cinit
|-value: Array size=11
| |-elements: Int 103, Int 97, Int 109, Int 101
| |-elements: Int 111, Int 102, Int 108, Int 105
| |-elements: Int 102, Int 101
`-filler: 1 x Int 0
-StringLiteral 0x26d24bde440 <col:41> 'const char[11]' "gameoflife"
-FieldDecl 0x26d24bde558 <line:27:2, col:32> col:12 referenced local_size 'sf::uvec3'
`-CXXConstructExpr 0x26d24bded38 <col:22, col:32> 'sf::uvec3' 'void (uint32_t, uint32_t, uint32_t)' list
|-ImplicitCastExpr 0x26d24bdecf0 <col:24> 'uint32_t':'unsigned int' <IntegralCast>
| |-IntegerLiteral 0x26d24bdec20 <col:24> 'int' 18
|-ImplicitCastExpr 0x26d24bded08 <col:27> 'uint32_t':'unsigned int' <IntegralCast>
| |-IntegerLiteral 0x26d24bdec48 <col:27> 'int' 18
`-ImplicitCastExpr 0x26d24bded20 <col:30> 'uint32_t':'unsigned int' <IntegralCast>
`-IntegerLiteral 0x26d24bdec70 <col:30> 'int' 1
-FieldDecl 0x26d24bdea90 <line:28:2, col:16> col:16 t 'Test<float, 3>'
-CXXMethodDecl 0x26d24bdeb48 <line:29:2, line:32:2> line:29:7 main 'void ()' implicit-inline
`-CompoundStmt 0x26d24be96a0 <col:13, line:32:2>
`-DeclStmt 0x26d24be9688 <line:31:3, col:30>
`-VarDecl 0x26d24bded98 <col:3, col:29> col:12 idx 'unsigned int' cinit
```

Find kernel structs via matchers

```
1 #include "vec.hpp"
2 #include "computebackend.hpp"
3 struct [[clang::annotate("kernel")]] FloatAdder
4 {
5     static constexpr char fileLocation[] = "floatadder";
6     uvec3 local_size{ 256, 1, 1 };
7     BufferBinding< float , 0> A;
8     BufferBinding< float , 1> B;
9     BufferBinding< float , 2> C;
10    void main() {
11        tc::uint i = tc::gl_GlobalInvocationID.x;
12        C[i] = A[i] + B[i];
13    }
14};
```

Validation
(Clang/AST)

```
auto kernelStructMatcher = cxxRecordDecl(
    isDefinition(),
    hasAttr(attr::Annotate),
    hasMethod(
        cxxMethodDecl(hasName("main"))
    )
    hasDescendant(
        varDecl(
            hasName("fileLocation"),
            isConstexpr()
        ).bind("fileLocation")
    ),
    hasDescendant(
        fieldDecl(
            hasName("local_size")
        ).bind("localSizeVar")
    ),
).bind("kernelStruct");
```

Clang validation

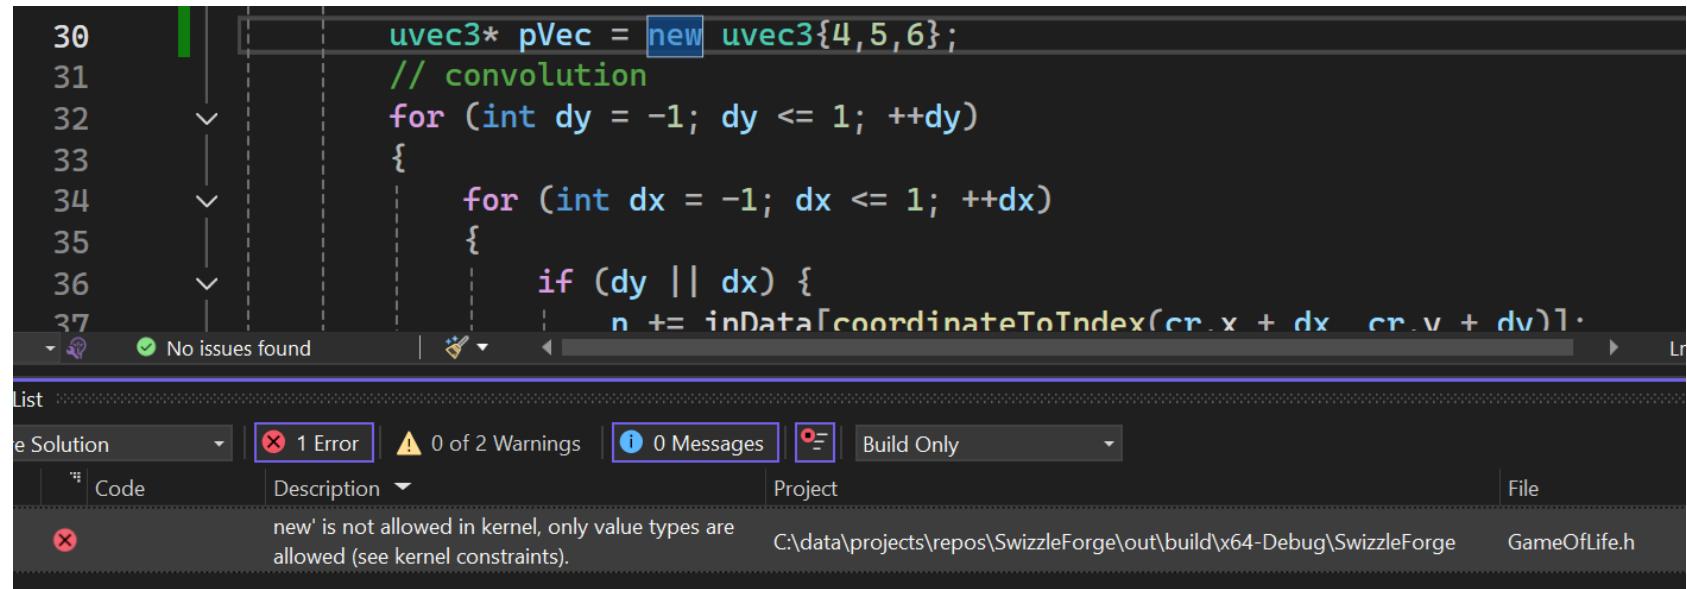
- C++ AST library
- Visit methods
 - VisitFieldDecl(clang::FieldDecl* FD)
 - VisitCXXNewExpr(clang::CXXNewExpr* E)
 - ...
- Report errors back to the IDE

Note: only look
within the found
kernel structs

Clang validation – example 1

```
bool VisitCXXNewExpr(clang::CXXNewExpr* E){
    llvm::errs()
        << E->getBeginLoc().printToString(Context.getSourceManager())
        << " Error: 'new' is not allowed in kernel, only value types are allowed.\n";

    reportError(E->getBeginLoc(), dynamicMemory);
    Valid = false;
    return true;
}
```



The screenshot shows a code editor with a dark theme. A vertical line of numbers on the left indicates line 30. The code itself is a C++ snippet involving nested loops and a conditional statement. At the bottom of the editor, a status bar displays 'No issues found'. Below the editor is a 'List' pane with a toolbar above it. The toolbar includes buttons for 'Solution' (selected), 'Code' (highlighted in blue), 'Description' (highlighted in blue), 'Project' (disabled), and 'File' (disabled). The 'Code' tab shows one error message: 'new' is not allowed in kernel, only value types are allowed (see kernel constraints). The 'Description' tab shows the same message. The 'Project' and 'File' tabs are disabled.

```
30     uvec3* pVec = new uvec3{4,5,6};
31     // convolution
32     for (int dy = -1; dy <= 1; ++dy)
33     {
34         for (int dx = -1; dx <= 1; ++dx)
35         {
36             if (dy || dx) {
37                 n += inData[coordinateToIndex(cr.x + dx, cr.y + dy)];
```

No issues found

1 Error | 0 of 2 Warnings | 0 Messages | Build Only

Code | Description | Project | File

new' is not allowed in kernel, only value types are allowed (see kernel constraints).

Case 1: Transpilation

```
#include "vec.hpp"
#include "computebackend.hpp"
struct [[clang::annotate("kernel")]] FloatAdder
{
    static constexpr char fileLocation[] = "...";
    uvec3 local_size{ 256, 1, 1 };
    BufferBinding< float , 0> A;
    BufferBinding< float , 1> B;
    BufferBinding< float , 2> C;
    void main() {
        tc::uint i = tc::gl_GlobalInvocationID.x;
        C[i] = A[i] + B[i];
    }
};
```

Transpilation
(Clang/AST)

```
#version 430
layout (
    local_size_x = 256,
    local_size_y = 1,
    local_size_z = 1)
in;
layout(set = 0, binding = 0) buffer _ALayout {
    float A[];
};
layout(set = 0, binding = 1) buffer _BLayout {
    float B[];
};
layout(set = 0, binding = 2) buffer _CLayout {
    float C[];
};
void main() {
    uint i = gl_GlobalInvocationID.x;
    C[i] = A[i] + B[i];
}
```

From Pure ISO C++20 To Compute Shaders

29

Case 1: Usage

```
using Backend = tc::gpu::GPUBackend;  
Backend compute;
```

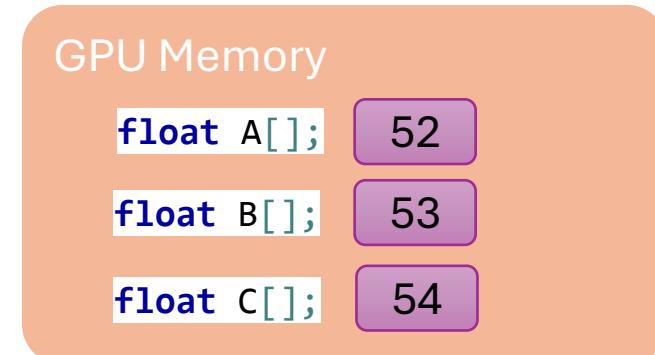
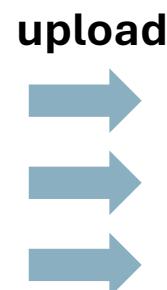
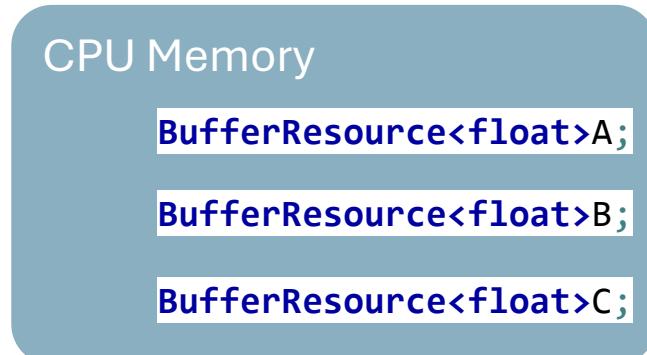
Case 1: Usage

```
using Backend = tc::gpu::GPUBackend;
Backend compute;

compute.uploadBuffer(m_pA.get());
compute.uploadBuffer(m_pB.get());
compute.uploadBuffer(m_pC.get());
```

With the definitions for `m_pA`, `m_pB` and `m_pC` given as:

```
std::unique_ptr<tc::BufferResource<float>> m_pA;
std::unique_ptr<tc::BufferResource<float>> m_pB;
std::unique_ptr<tc::BufferResource<float>> m_pC;
```



Note :
use renderdoc

Case 1: Usage

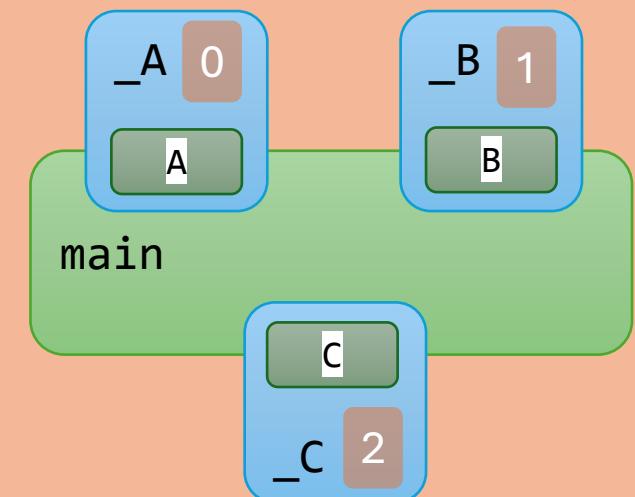
```
using Backend = tc::gpu::GPUBackend;
Backend compute;

compute.uploadBuffer(m_pA.get());
compute.uploadBuffer(m_pB.get());
compute.uploadBuffer(m_pC.get());

compute.useKernel(m_FloatAdder);
```

GPU Memory

float A[];	52
float B[];	53
float C[];	54



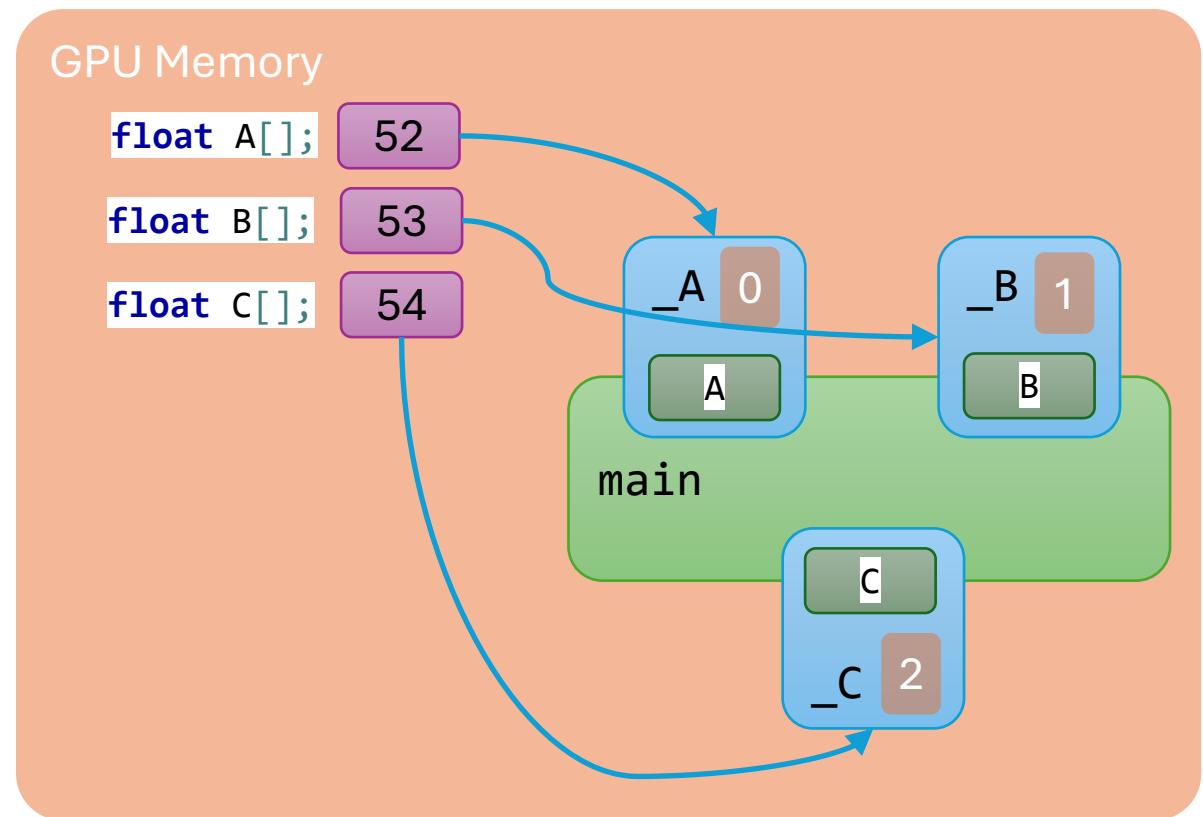
Case 1: Usage

```
using Backend = tc::gpu::GPUBackend;
Backend compute;

compute.uploadBuffer(m_pA.get());
compute.uploadBuffer(m_pB.get());
compute.uploadBuffer(m_pC.get());

compute.useKernel(m_FloatAdder);

compute.bindBuffer(m_FloatAdder.A);
compute.bindBuffer(m_FloatAdder.B);
compute.bindBuffer(m_FloatAdder.C);
```



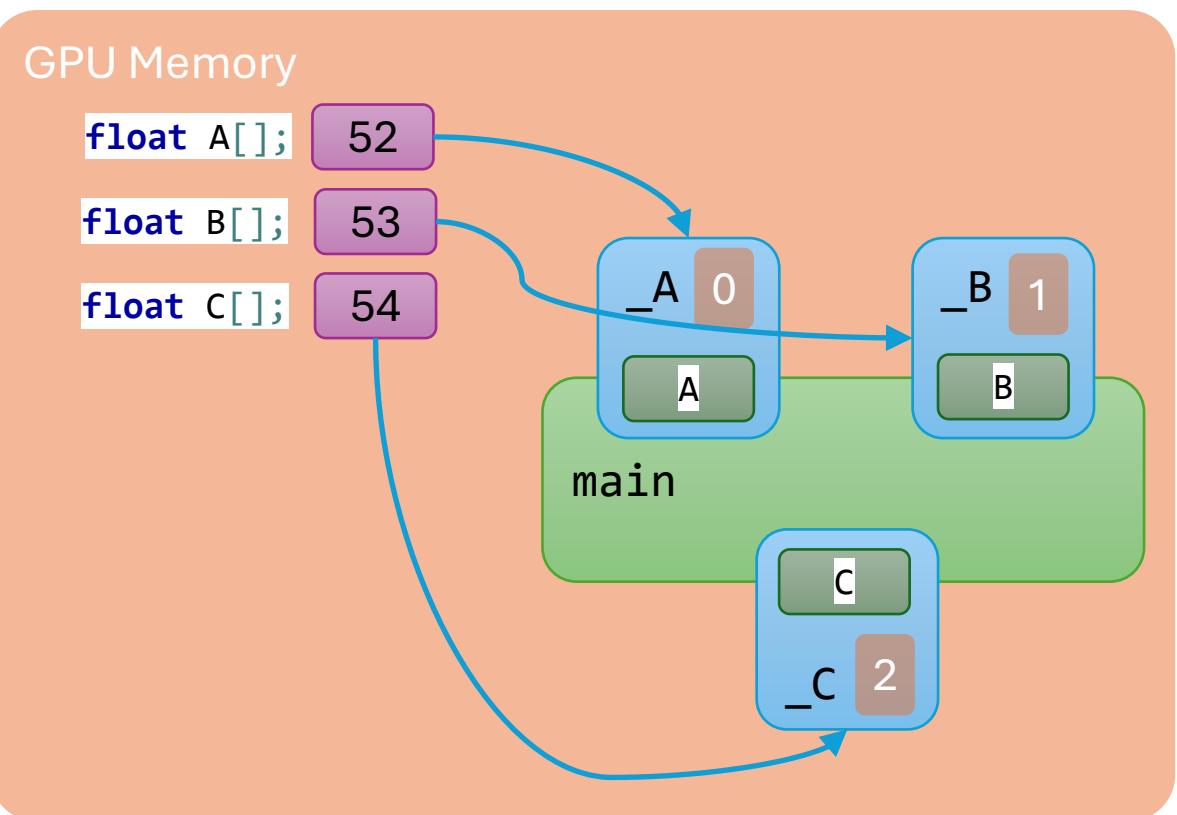
Case 1: Usage

```
using Backend = tc::gpu::GPUBackend;
Backend compute;

compute.uploadBuffer(m_pA.get());
compute.uploadBuffer(m_pB.get());
compute.uploadBuffer(m_pC.get());

compute.useKernel(m_FloatAdder);

compute.bindBuffer(m_FloatAdder.A);
compute.bindBuffer(m_FloatAdder.B);
compute.bindBuffer(m_FloatAdder.C);
compute.execute(m_FloatAdder, tc::uvec3{ N, 1, 1 });
```



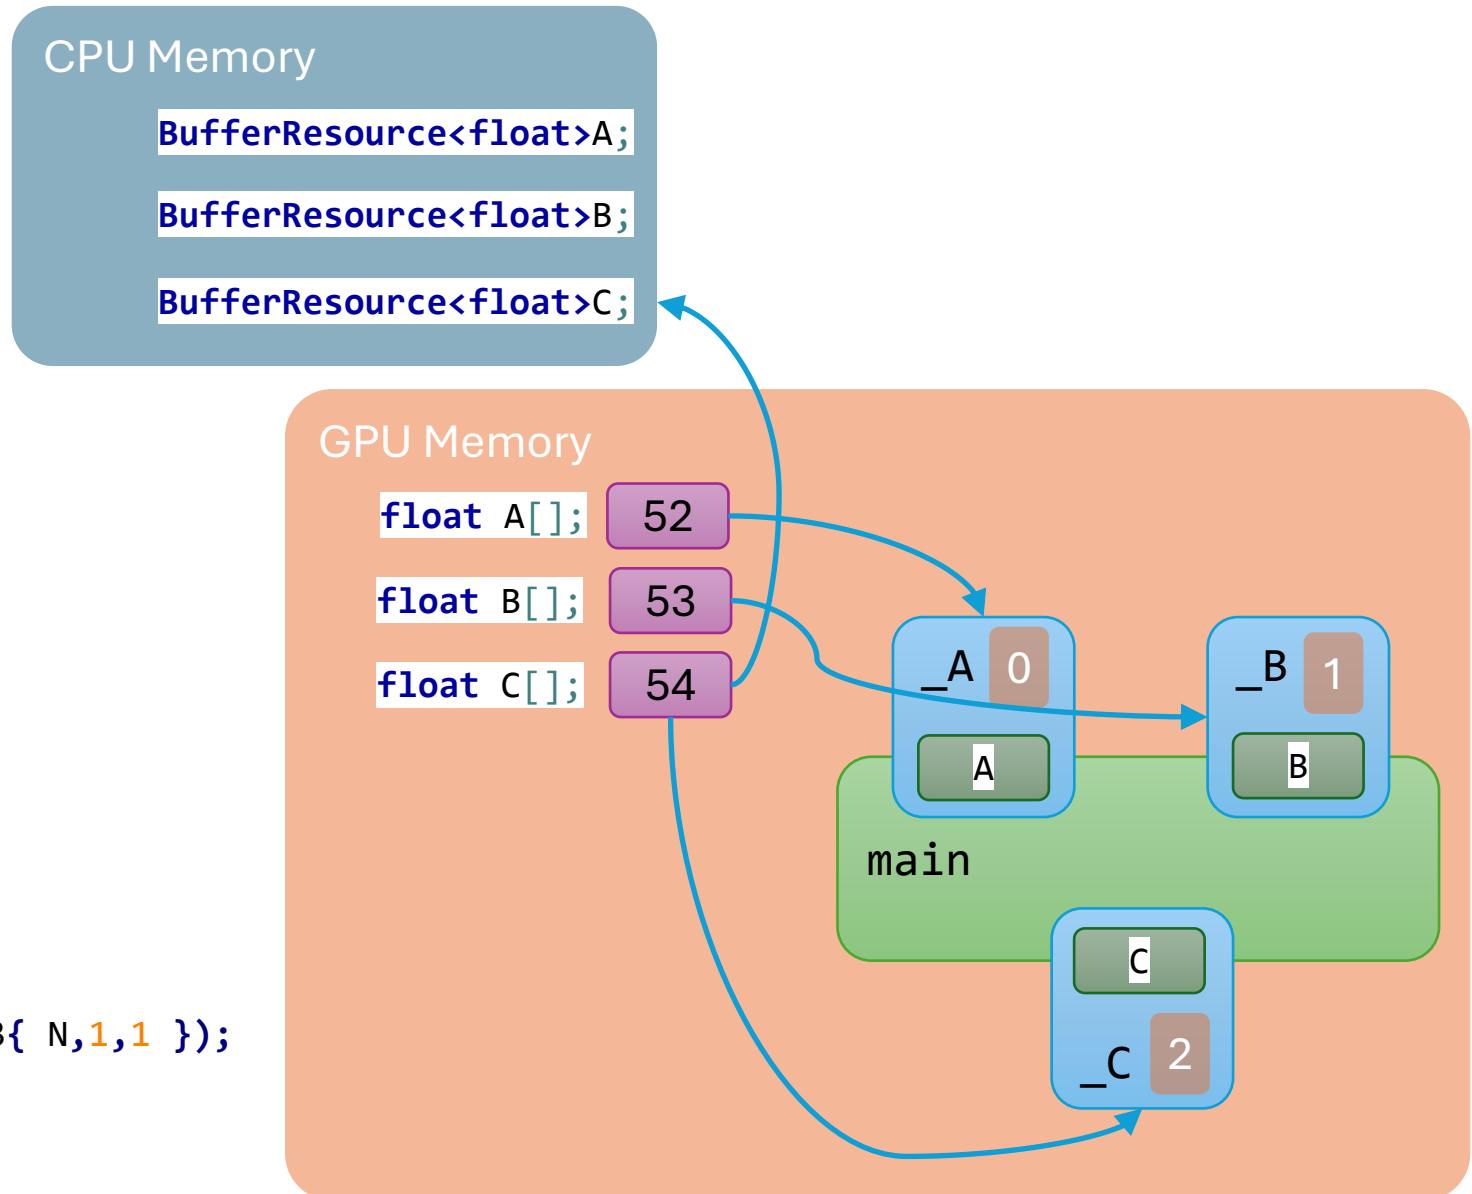
Case 1: Usage

```
using Backend = tc::gpu::GPUBackend;
Backend compute;

compute.uploadBuffer(m_pA.get());
compute.uploadBuffer(m_pB.get());
compute.uploadBuffer(m_pC.get());

compute.useKernel(m_FloatAdder);

compute.bindBuffer(m_FloatAdder.A);
compute.bindBuffer(m_FloatAdder.B);
compute.bindBuffer(m_FloatAdder.C);
compute.execute(m_FloatAdder, tc::uvec3{ N, 1, 1 });
compute.downloadBuffer(m_pC.get());
```



Case 1: Usage

```
using Backend = tc::gpu::GPUBackend;
Backend compute;

compute.uploadBuffer(m_pA.get());
compute.uploadBuffer(m_pB.get());
compute.uploadBuffer(m_pC.get());

compute.useKernel(m_FloatAdder);

compute.bindBuffer(m_FloatAdder.A);
compute.bindBuffer(m_FloatAdder.B);
compute.bindBuffer(m_FloatAdder.C);
compute.execute(m_FloatAdder, tc::uvec3{ N, 1, 1 });
compute.downloadBuffer(m_pC.get());
```

CPU Memory

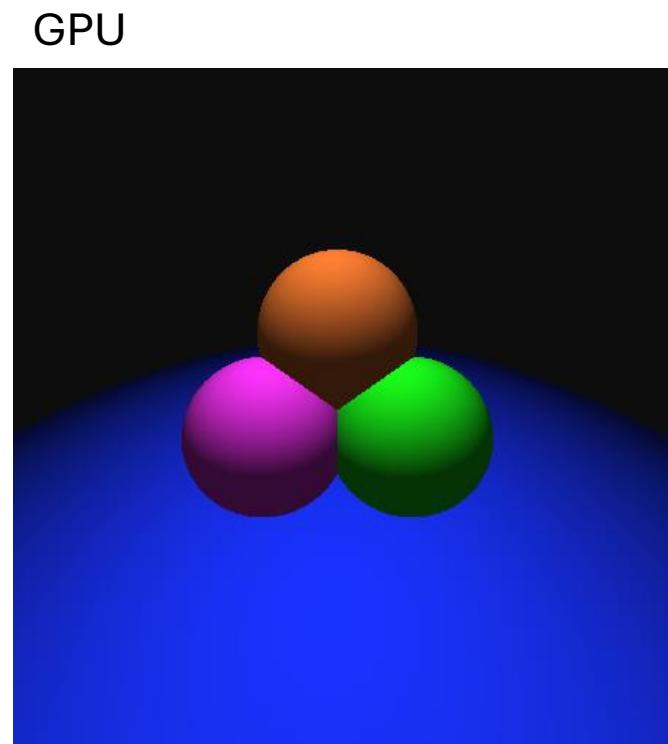
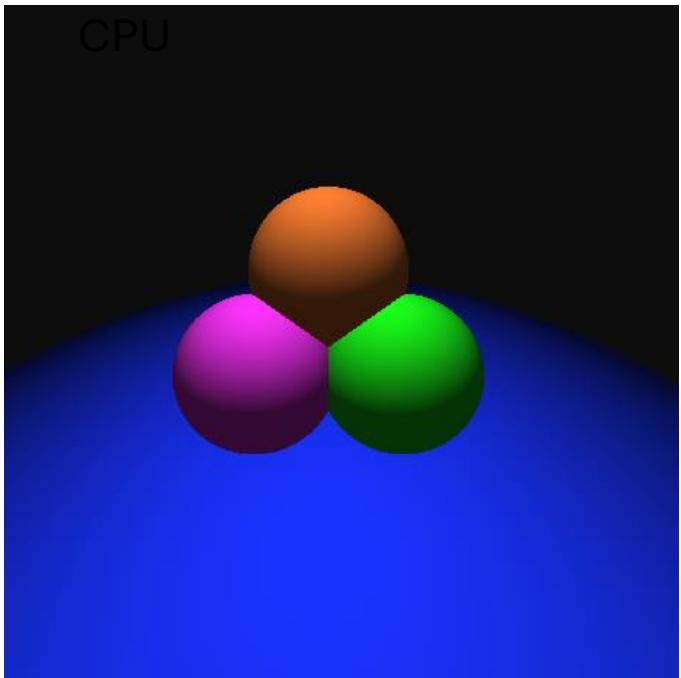
BufferResource<float>A;

BufferResource<float>B;

BufferResource<float>C;

Everything is already in the
correct memory space →
log error if buffer not bound

Raytracer - demo



File Edit View Git Project Build Debug Test Analyze Tools Extensions Window Help Search SwizzleForge

Local Machine x64-Debug RaytracerCPU.exe (Projects\Raytracer\RaytracerCpu\RaytracerCPU.exe)

RayTracerWindow.cpp RayTracerCPU.hpp

RaytracerCPU.exe (Projects\Raytracer\RaytracerCpu\RaytracerCPU.e SphereRayTracer

32 struct [[clang::annotate("kernel")]] SphereRayTracer
33 {
34 static constexpr char fileLocation[] = "sphere_raytracer";
35 tc::uvec3 local_size{ 16,16,1 };
36 tc::ImageBinding<tc::InternalFormat::RGBA32F, tc::Dim::D2, tc::cpu::RGBA8, 0> rays;
37 tc::BufferBinding<float, 1> tBuffer;
38
39 struct Sphere {
40 tc::vec3 loc;
41 float R;
42 tc::vec4 color;
43 };
44 tc::BufferBinding<Sphere, 2> spheres;
45 tc::Uniform<int, 0> nrOfSpheres;
46
47 tc::ImageBinding<tc::InternalFormat::RGBA8, tc::Dim::D2, tc::cpu::RGBA8UI, 1> outputTexture;
48 tc::Uniform<tc::vec3, 1> lightPos{ tc::vec3{-0.25,3,0.2} };
49
50 void main()
51 {
52 using namespace tc;
53 }

136% No issues found

Output

Show output from: Debug

The thread 31436 has exited with code 0 (0x0).
The thread 19328 has exited with code 0 (0x0).
The thread 31592 has exited with code 0 (0x0).
The thread 29872 has exited with code 0 (0x0).
The thread 30772 has exited with code 0 (0x0).
The thread 32440 has exited with code 0 (0x0).
The program '[31084] RaytracerCPU.exe' has exited with code 0 (0x0).

Error List Output

Solution Explorer - CMake Targets View

Search Solution Explorer - CMake Targets View (Ctrl+.)

SwizzleForge (C:\data\projects\repos\SwizzleForge)
Pinned Targets
TinyCompute Project
01-Core
AssetLib (static library)
ComputeLibOpenGL (static library)
TinyCompute (interface library)
02-Tools
03-Examples
A-AddFloats
B-Raytracer
CameraRays (executable)
RaytracerCPU (executable)
References
CMakeLists.txt
main.cpp
RayTracerCPU.hpp
RayTracerWindow.hpp
RaytracerGPU (executable)
C-GameOfLife
03-Examples
Project01 (executable)
Project02 (executable)
Project03 (executable)
swizzleforge_demo (executable)
04-Tests
3rd-party
cmake
CMakeLists.txt
CMakeSettings.json
util.cmake

Unity Project Explorer GitHub Copilot Chat Solution Explorer Git Changes

Properties

Ready

File Edit View Git Project Build Debug Test Analyze Tools Extensions Window Help Search SwizzleForge

Local Machine x64-Debug RaytracerCPU.exe (Projects\Raytracer\RaytracerCpu\RaytracerCPU.exe)

RayTracerWindow.cpp RayTracerCPU.hpp

```
38 tc::BufferBinding<float, 1> tBuffer;
39
40 struct Sphere {
41     tc::vec3 loc;
42     float R;
43     tc::vec4 color;
44 };
45 tc::BufferBinding<Sphere, 2> spheres;
46 tc::Uniform<int, 0> nrOfSpheres;
47
48 tc::ImageBinding<tc::InternalFormat::RGBA8, tc::Dim::D2, tc::cpu::RGBA8UI, 1> outputTexture;
49 tc::Uniform<tc::vec3, 1> lightPos{ tc::vec3{-0.25, 3, 0.2} };
50
51 void main()
52 {
53     using namespace tc;
54     uvec2 gid = gl_GlobalInvocationID["xy"_sw];
55     ivec2 coordinate = ivec2(gid.x, gid.y);
56     vec3 rayOrigin = vec3(0, 0, 0);
57     vec3 rayDirection = imageLoad(rays, coordinate)["xyz"_sw];
58
59     ivec2 imgSize = imageSize(rays);
```

No issues found

Output

```
Show output from: Debug
The thread 5952 has exited with code 0 (0x0).
The thread 30060 has exited with code 0 (0x0).
The thread 15264 has exited with code 0 (0x0).
The thread 14564 has exited with code 0 (0x0).
The thread 31440 has exited with code 0 (0x0).
The thread 25772 has exited with code 0 (0x0).
The program '[31608] RaytracerCPU.exe' has exited with code 0 (0x0).
```

Error List Output

Solution Explorer - CMake Targets View

Search Solution Explorer - CMake Targets View (Ctrl+.)

- SwizzleForge (C:\data\projects\repos\SwizzleForge)
 - Pinned Targets
- TinyCompute Project
 - 01-Core
 - AssetLib (static library)
 - ComputeLibOpenGL (static library)
 - TinyCompute (interface library)
 - 02-Tools
 - 03-Examples
 - A-AddFloats
 - B-Raytracer
 - CameraRays (executable)
 - RaytracerCPU (executable)
 - References
 - CMakeLists.txt
 - main.cpp
 - RayTracerCPU.hpp
 - RayTracerWindow.hpp
 - RaytracerGPU (executable)
 - C-GameOfLife
 - 03-Examples
 - Project01 (executable)
 - Project02 (executable)
 - Project03 (executable)
 - swizzleforge_demo (executable)
 - 04-Tests
 - 3rd-party
 - cmake
 - CMakeLists.txt
 - CMakeSettings.json
 - util.cmake

Unity Project Explorer GitHub Copilot Chat Solution Explorer Git Changes

Properties

Ready

File Edit View Git Project Build Debug Test Analyze Tools Extensions Window Help Search SwizzleForge

Local Machine x64-Debug RaytracerCPU.exe (Projects\Raytracer\RaytracerCpu\RaytracerCPU.exe)

RayTracerWindow.cpp RayTracerCPU.hpp

RaytracerCPU.exe (Projects\Raytracer\RaytracerCpu\RaytracerCPU.e SphereRayTracer::Sphere

32 struct [[clang::annotate("kernel")]] SphereRayTracer
33 {
34 static constexpr char fileLocation[] = "sphere_raytracer";
35 tc::uvec3 local_size{ 16,16,1 };
36 tc::ImageBinding<tc::InternalFormat::RGBA32F, tc::Dim::D2, tc::cpu::RGBA8, 0> rays;
37 tc::BufferBinding<float, 1> tBuffer;
38
39 struct Sphere {
40 tc::vec4 locR;
41 tc::vec4 color;
42 };
43 tc::BufferBinding<Sphere, 2> spheres;
44 tc::Uniform<int, 0> nrOfSpheres;
45
46 tc::ImageBinding<tc::InternalFormat::RGBA8, tc::Dim::D2, tc::cpu::RGBA8UI, 1> outputTexture;
47 tc::Uniform<tc::vec3, 1> lightPos{ tc::vec3{-0.25,3,0.2} };
48
49 void main()
50 {
51 using namespace tc;
52 uvec2 gid = gl_GlobalInvocationID["xy"_sw];
53 }
54 }

No issues found

Output

Show output from: Debug

The thread 5952 has exited with code 0 (0x0).
The thread 30060 has exited with code 0 (0x0).
The thread 15264 has exited with code 0 (0x0).
The thread 14564 has exited with code 0 (0x0).
The thread 31440 has exited with code 0 (0x0).
The thread 25772 has exited with code 0 (0x0).
The program '[31608] RaytracerCPU.exe' has exited with code 0 (0x0).

Solution Explorer - CMake Targets View

Search Solution Explorer - CMake Targets View (Ctrl+.)

SwizzleForge (C:\data\projects\repos\SwizzleForge)
Pinned Targets
TinyCompute Project
01-Core
AssetLib (static library)
ComputeLibOpenGL (static library)
TinyCompute (interface library)
02-Tools
03-Examples
A-AddFloats
B-Raytracer
CameraRays (executable)
RaytracerCPU (executable)
References
CMakeLists.txt
main.cpp
RayTracerCPU.hpp
RayTracerWindow.hpp
RaytracerGPU (executable)
C-GameOfLife
03-Examples
Project01 (executable)
Project02 (executable)
Project03 (executable)
swizzleforge_demo (executable)
04-Tests
3rd-party
cmake
CMakeLists.txt
CMakeSettings.json
util.cmake

Unity Project Explorer GitHub Copilot Chat Solution Explorer Git Changes

Error List Output

Ready

File Edit View Git Project Build Debug Test Analyze Tools Extensions Window Help Search SwizzleForge

Local Machine x64-Debug RaytracerCPU.exe (Projects\Raytracer\RaytracerCpu\RaytracerCPU.exe)

RayTracerWindow.cpp RayTracerCPU.hpp

RaytracerCPU.exe (Projects\Raytracer\RaytracerCpu\RaytracerCPU.e SphereRayTracer main()

50 void main()
51 {
52 using namespace tc;
53 uvec2 gid = gl_GlobalInvocationID["xy"_sw];
54 ivec2 coordinate = ivec2(gid.x, gid.y);
55 vec3 rayOrigin = vec3(0, 0, 0);
56 vec3 rayDirection = imageLoad(rays, coordinate)["xyz"_sw];
57
58 ivec2 imgSize = imageSize(rays);
59 int index = imgSize.x * coordinate.y + coordinate.x;
60
61 bool write = false;
62 vec4 color = vec4(0.05, 0.05, 0.75, 1);
63 for (int si = 0; si < nrOfSpheres; ++si)
64 {
65 vec3 sphereLoc = spheres[si].locR["xyz"_sw];
66 float r = spheres[si].locR.w; I
67
68 vec3 raySphereDiff = sphereLoc - rayOrigin;
69 float L2 = dot(raySphereDiff, raySphereDiff);
70 float tca = dot(raySphereDiff, rayDirection);
71 float od2 = L2 - tca * tca;

136% No issues found Ln: 66 Ch: 32 Col: 41 TABS CRLF

Output Show output from: Debug
The thread 5952 has exited with code 0 (0x0).
The thread 30060 has exited with code 0 (0x0).
The thread 15264 has exited with code 0 (0x0).
The thread 14564 has exited with code 0 (0x0).
The thread 31440 has exited with code 0 (0x0).
The thread 25772 has exited with code 0 (0x0).
The program '[31608] RaytracerCPU.exe' has exited with code 0 (0x0).

Error List Output

Solution Explorer - CMake Targets View

Search Solution Explorer - CMake Targets View (Ctrl+.)

- SwizzleForge (C:\data\projects\repos\SwizzleForge)
 - Pinned Targets
- TinyCompute Project
 - 01-Core
 - AssetLib (static library)
 - ComputeLibOpenGL (static library)
 - TinyCompute (interface library)
 - 02-Tools
 - 03-Examples
 - A-AddFloats
 - B-Raytracer
 - CameraRays (executable)
 - RaytracerCPU (executable)
 - References
 - CMakeLists.txt
 - main.cpp
 - RayTracerCPU.hpp
 - RayTracerWindow.hpp
 - RaytracerGPU (executable)
 - C-GameOfLife
 - 03-Examples
 - Project01 (executable)
 - Project02 (executable)
 - Project03 (executable)
 - swizzleforge_demo (executable)
 - 04-Tests
 - 3rd-party
 - cmake
 - CMakeLists.txt
 - CMakeSettings.json
 - util.cmake

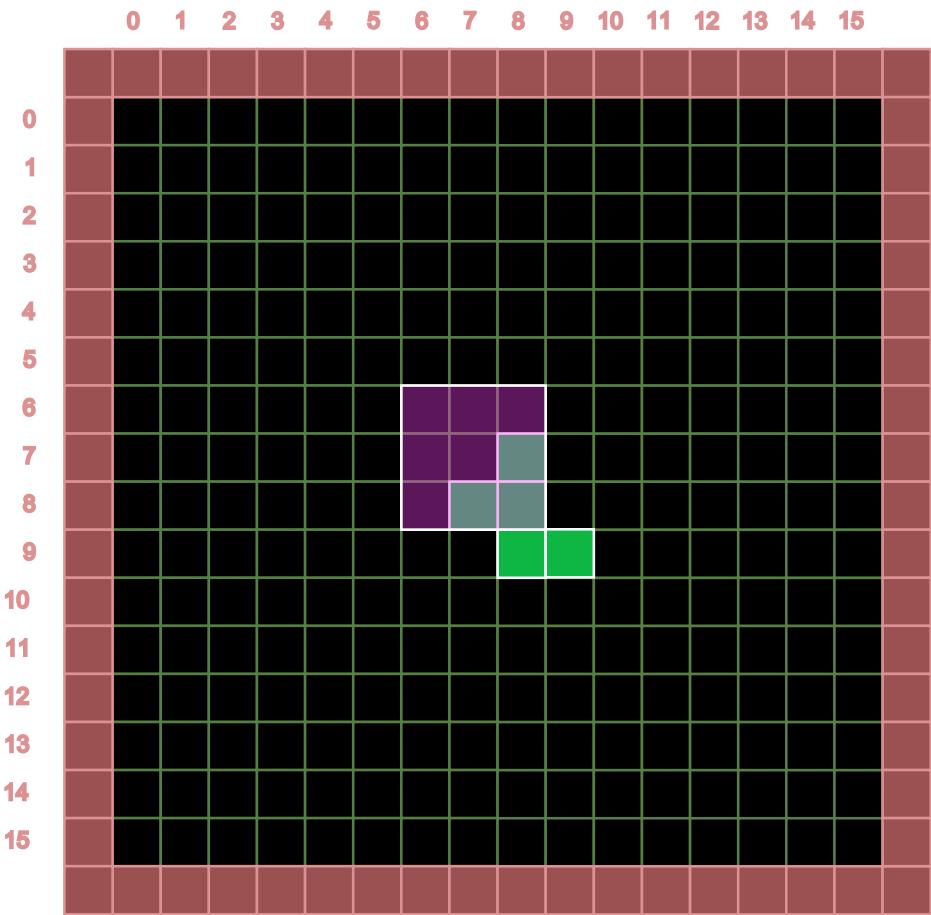
Unity Project Explorer GitHub Copilot Chat Solution Explorer Git Changes

Properties

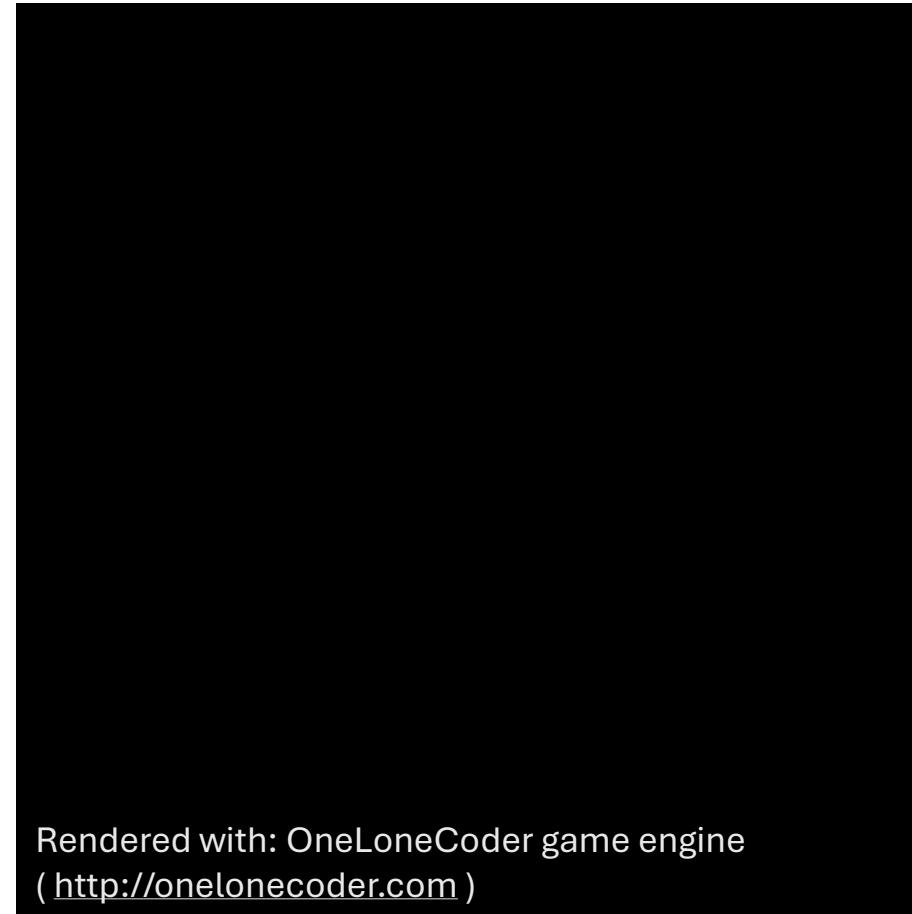
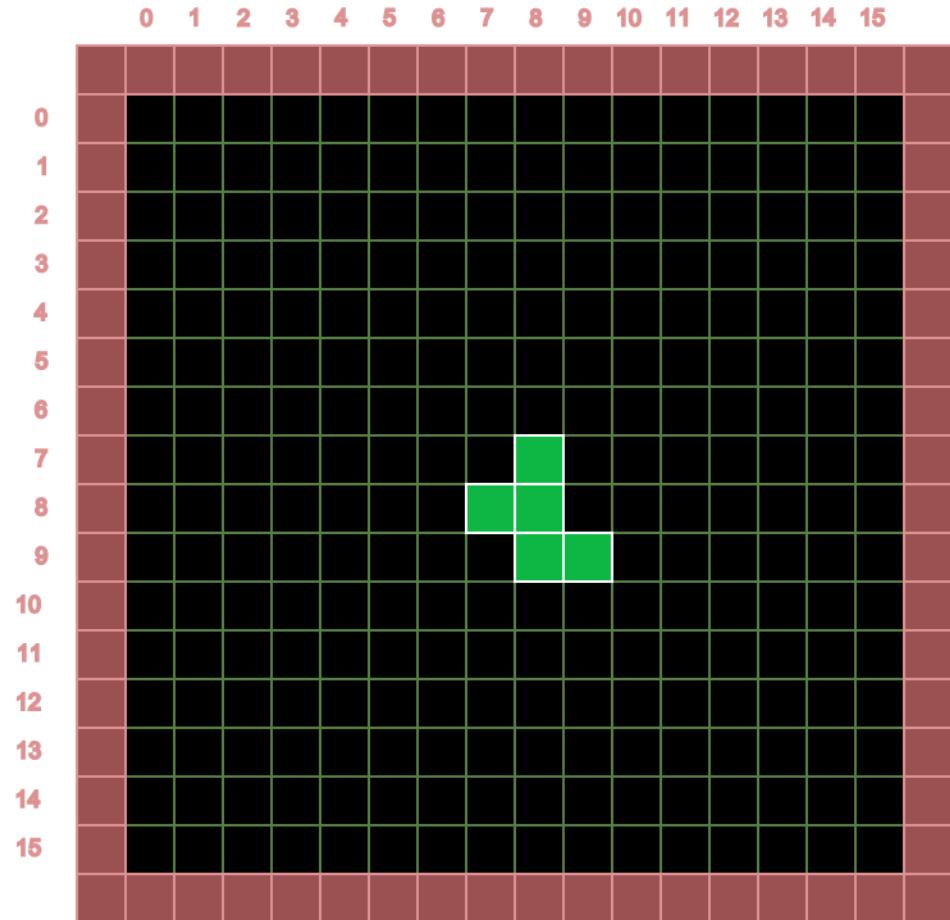
Ready 0 / 0 17 main SwizzleForge

Case 3 : Game of life

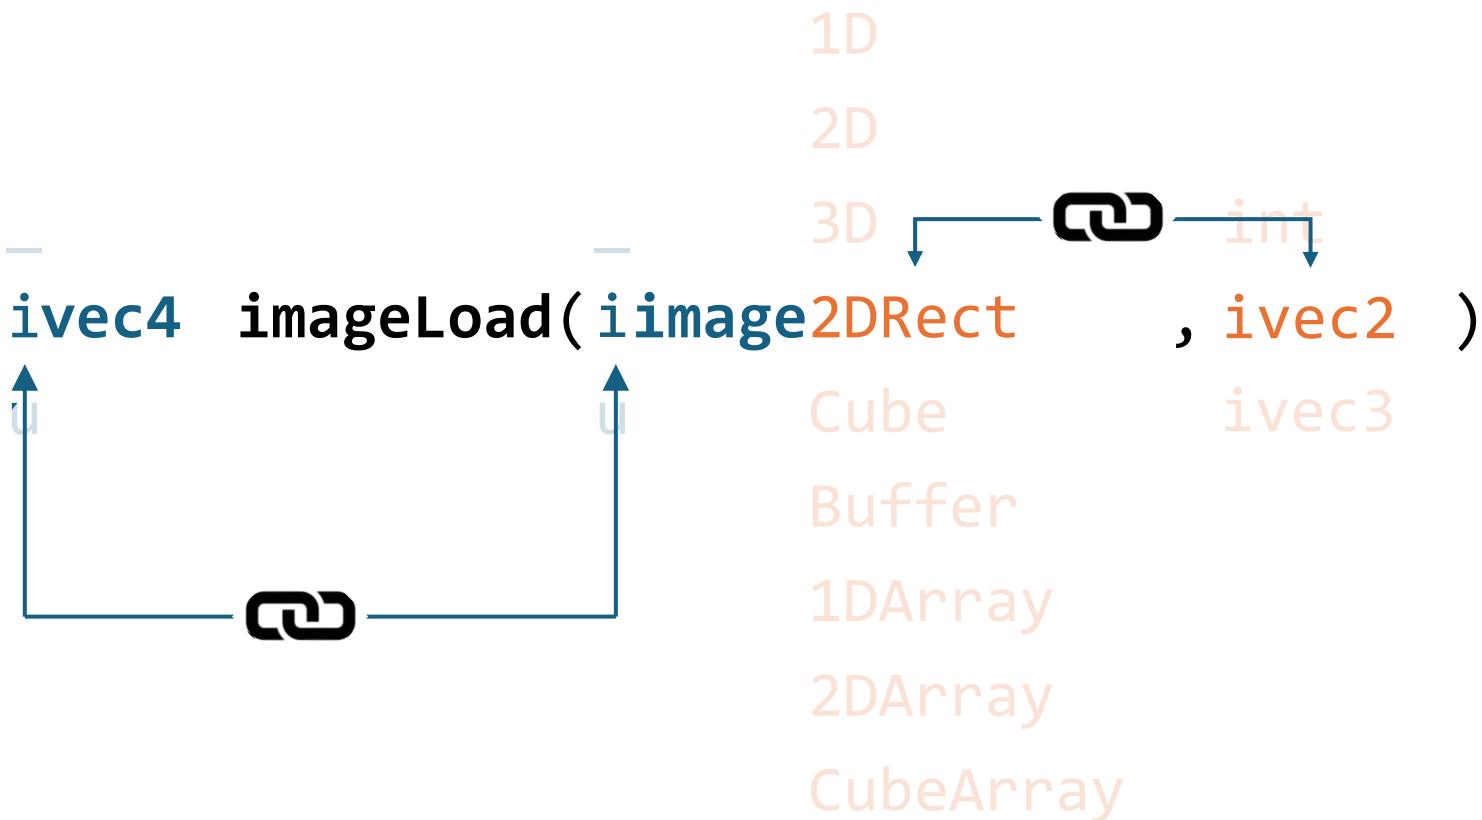
- Image
 - Source Formats
 - External type
 - Destination Formats
 - Internal type
- GLSL Image operations
 - `imageLoad`
 - `ImageStore`



Case 3: Game of life

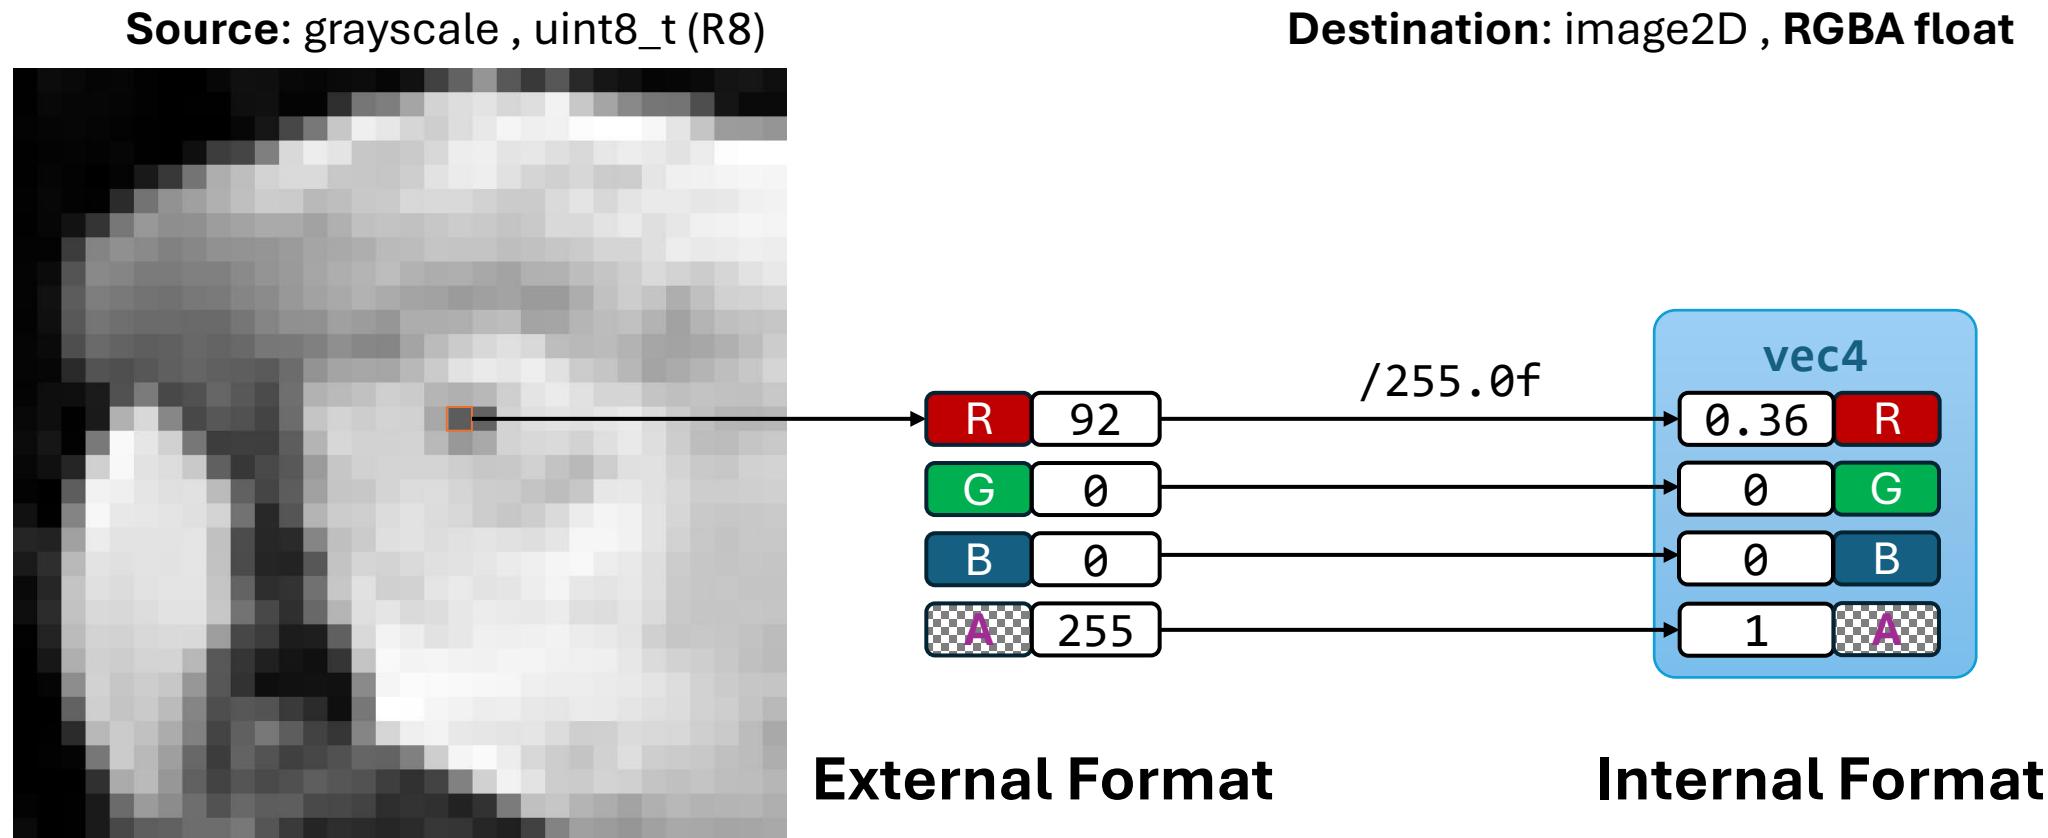


ImageLoad - GLSL



27
combinations!

But wait! There's more



Images in GLSL

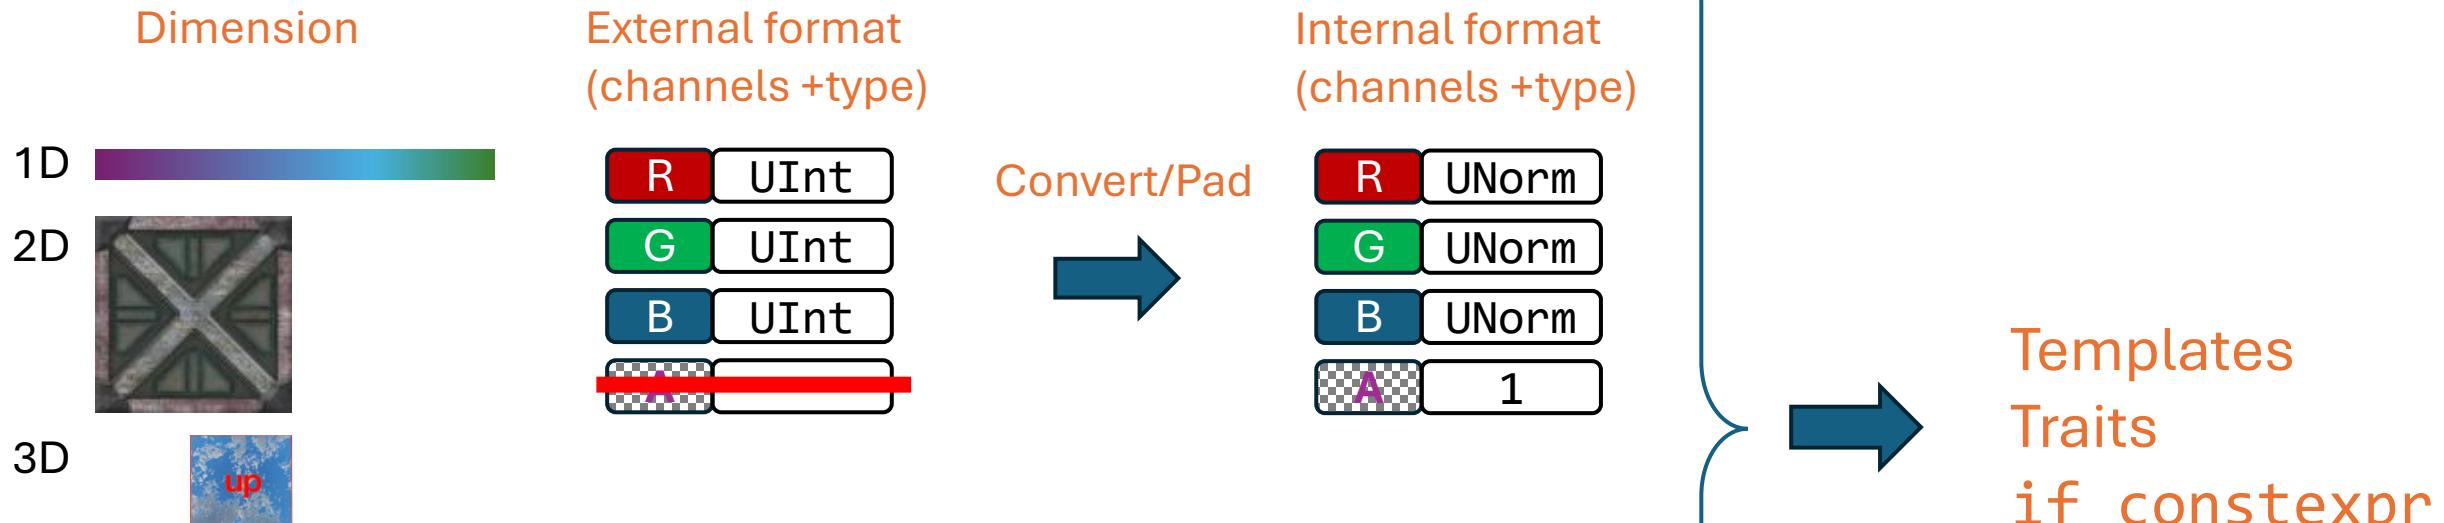
```
1 #version 430

2 layout (local_size_x = 16, local_size_y = 16, local_size_z = 1) in;

3 layout(binding=0, r8ui ) uniform uimage2D inData;           Internal format
4 layout(binding=1, r8ui ) uniform uimage2D outData;

5 void main()
{
    6
    7     uvec4 pixelIn = imageLoad( inData, coordinate );
    8     // convolution
    9     imageStore( outData, coordinate, uvec4(1,0,0,0) );
10 }
```

imageLoad – C++



Can we create **one** version of the **imageLoad** free function?

Without runtime branching ?

imageLoad - BufferResource

```
template<typename T, Dim D = Dim::D1>
```

BufferResource

```
    BufferResource(dimType bufferSize)
    :m_Data(/* bufferSize (nD) → size */)
    {
    }

    const T& operator[](dimType index) const
    {
        return m_Data[ /* index (nD) → 1D */ ];
    }
```

private:

```
    dimType m_BufferSize;
    std::vector<T> m_Data; —————→ 1D
```

D	dimType
D1	int32_t
D2	ivec2
D3	ivec3

GLSL: *signed int* is type used in **imageLoad**

Dimension - Traits

```
enum class Dim : uint8_t {
    D1 = 1,
    D2 = 2,
    D3 = 3
};

1 template<>
2 struct DimTraits< Dim::D1 > {
3     using IndexType = int32_t;
4     static constexpr int32_t flatten( IndexType coord, IndexType /*size*/ ) {
5         return coord;
6     }
7     static constexpr int32_t product( IndexType d ) {
8         return d;
9     }
10};
```

Dimension – 2D

```
enum class Dim : uint8_t {
    D1 = 1,
    D2 = 2,
    D3 = 3
};

1 template<>
2 struct DimTraits< Dim::D2 > {
3     using IndexType = ivec2;
4     static constexpr int32_t flatten( IndexType c, IndexType size ) {
5         return c.y * size.x + c.x;
6     }
7     static constexpr int32_t product( IndexType d ) {
8         return d.x * d.y;
9     }
10};
```

Dimension – 3D

```
enum class Dim : uint8_t {
    D1 = 1,
    D2 = 2,
    D3 = 3
};

1 template<>
2 struct DimTraits< Dim::D3 > {
3     using IndexType = ivec3;
4     static constexpr int32_t flatten( IndexType c, IndexType size ) {
5         return (coord.z * size.y + coord.y) * size.x + coord.x;
6     }
7     static constexpr int32_t product( IndexType d ) {
8         return d.x * d.y * d.z;
9     }
10};
```

imageLoad - BufferResource

```
template<typename T, Dim D = Dim::D1>

class BufferResource
public:
    using dimType = typename Traits::IndexType;

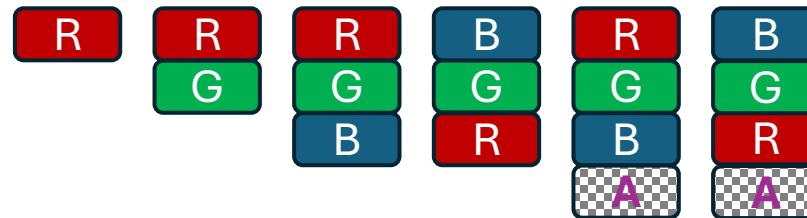
    BufferResource(dimType bufferSize)
        :m_Data( Traits::product(bufferSize) )
    {
    }

    const T& operator[](dimType index) const
    {
        return m_Data[ Traits::flatten(index) ];
    }

private:
    using Traits = DimTraits<D>;
    dimType m_BufferSize;
    std::vector<T> m_Data;
```

imageLoad – external format

Channels



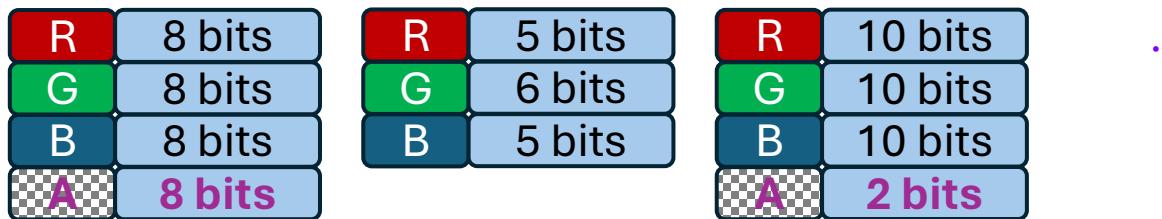
Semantics

Literal float or integer	Signed normal [-1,1]	Unsigned normal [0,1]	_INTEGER
-----------------------------	-------------------------	--------------------------	----------

Representation

uint8_t	int8_t	uint16_t	int16_t	float16_t	...
					C++23!

Packing



Pixel concept

```
1 using namespace tc;  
2 enum class Channel{ R=0, G=1, B=2, A=3 };  
3  
4 RGB8 px{ 0.2, 0.3, 0.5 };  
5 float r = px.getChannel<Channel::R>();  
6 float g = px.getChannel<Channel::G>();  
7 float b = px.getChannel<Channel::B>();  
8 float a = px.getChannel<Channel::A>();  
9  
10 EXPECT_EQ(r, 0.2);  
11 EXPECT_EQ(g, 0.3);  
12 EXPECT_EQ(b, 0.5);  
13 EXPECT_EQ(a, 1.0);
```

R	0.2
G	0.3
B	0.5



Padded value

Simple pixel: **r**, **g** and **b** have the same type → **float**

Pixel concept

```
using r_t = tc::cpu::RGB8::ChannelType<Channel::R>;  
using g_t = tc::cpu::RGB8::ChannelType<Channel::G>;  
using b_t = tc::cpu::RGB8::ChannelType<Channel::B>;  
using a_t = tc::cpu::RGB8::ChannelType<Channel::A>;
```

R	0.2
G	0.3
B	0.5

It is possible to have a different type for each channel.

Pixel Concept

```
template<class P, Channel C>
concept ChannelConcept = requires(P p) {

    typename P::template ChannelType<C>;
    { p.template get<C>() }
        -> std::convertible_to<typename P::template ChannelType<C>>;
    { p.template set<C>( p.template get<C>() ) };

};

template<class P>
concept PixelConcept =
    ChannelConcept<P, Channel::R> &&
    ChannelConcept<P, Channel::G> &&
    ChannelConcept<P, Channel::B> &&
    ChannelConcept<P, Channel::A>;
```

Conversions

```
template<class Src, class Dst>
struct ChannelConverter;
```

```
template<class T>
struct ChannelConverter<T, T> {
    static constexpr T apply(T v) noexcept { return v; }
};

template<>
struct ChannelConverter<std::uint8_t, float> {
    static constexpr float apply(std::uint8_t v) noexcept {
        return float(v) * (1.0f / 255.0f);
    }
};
```

External format
(channels +type)

R	UInt
G	UInt
B	UInt
A	

Convert/Pad



Internal format
(channels +type)

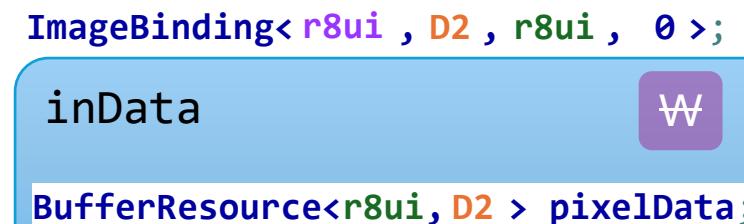
R	UNorm
G	UNorm
B	UNorm
A	1

ImageBinding

```
1 #version 430  
2  
3 layout (local_size_x = 16, local_size_y = 16, local_size_z = 1) in;  
4 layout(binding=0, r8ui ) uniform uimage2D inData;  
5 layout(binding=1, r8ui ) uniform uimage2D outData;  
6  
7 void main()  
8 {  
9     uvec4 pixelIn = imageLoad( inData, coordinate );  
10    // convolution  
11    imageStore( outData, uvec4(1,0,0,0) );  
12 }
```

Reminder:


`BufferResource<float>C;`
`BufferBinding<float,2>;`


`ImageBinding< r8ui , D2 , r8ui , 0 >;`
`inData`
`BufferResource<r8ui, D2 > pixelData;`

Internal Format

External format
(channels +type)

R	UInt
G	UInt
B	UInt
A	UInt

Convert/Pad



R8

R	0-1
G	0
B	0
A	1

Pad with internal
format **defaults**

Internal Format

External format
(channels + type)

R	UInt
G	UInt
B	UInt
A	UInt

Convert/Pad



R8

R	UInt
G	0
B	0
A	1

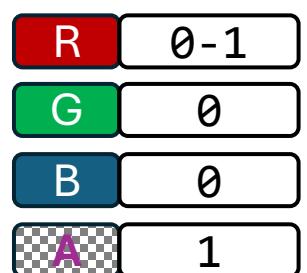
Convert



vec4

Internal Format

```
template<>
1 struct GPUFormatTraits<tc::gpu::R32F> {
2     using ChannelType = float;
3     using VectorType = tc::vec4;
4     static inline constexpr std::array<tc::Channel, 4> channels{           R8
5         Channel::R,
6         Channel::Min,
7         Channel::Min,
8         Channel::Max
9     };
10};
```



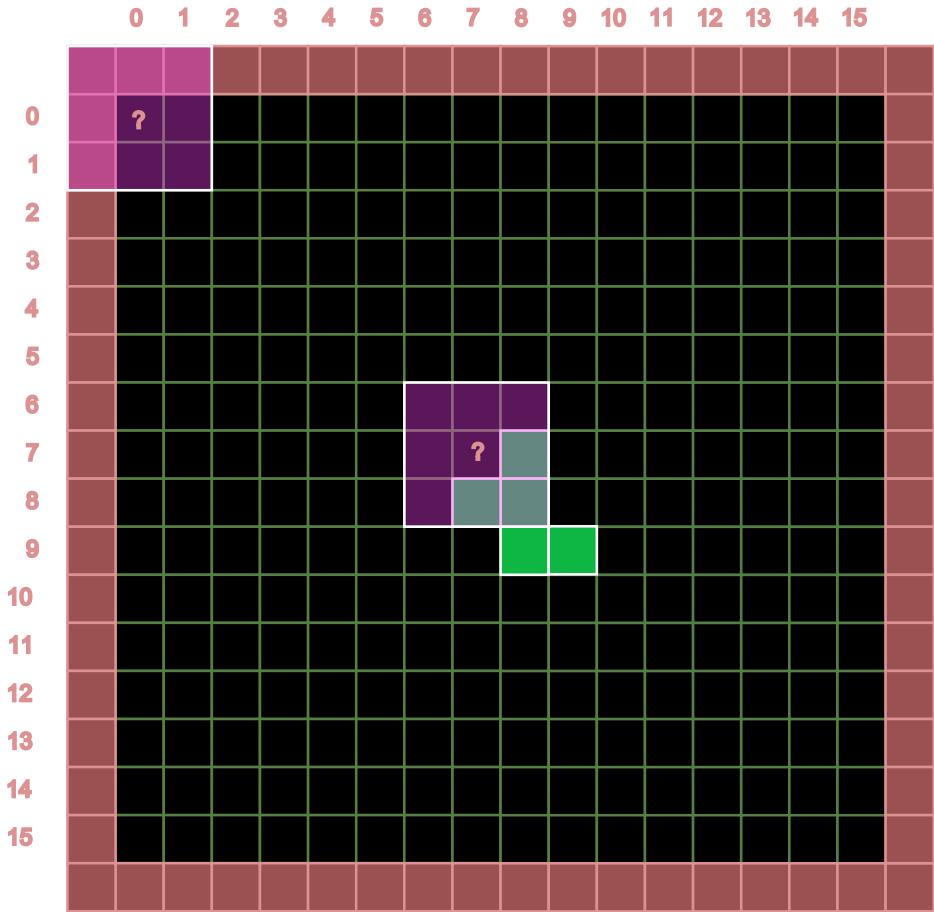
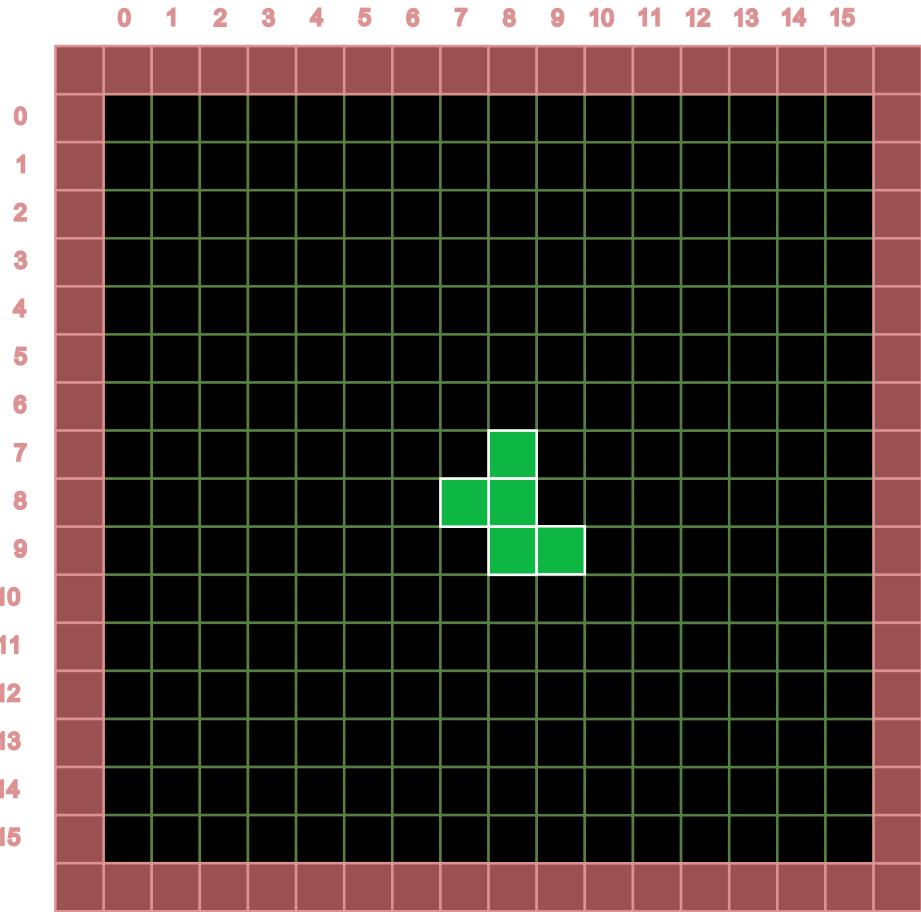
R8

R	0-1
G	0
B	0
A	1

ImageLoad

```
1 template<tc::InternalFormat G, tc::Dim D, tc::cpu::PixelConcept P>
2 auto imageLoad(const ImageBinding<G, D, P>& image, tcVec<D> texCoord)
3 {
4     const auto* buf = image.getBufferData();
5     const P& px = (*buf)[texCoord];
6     using gpuTraits = GPUFormatTraits<G>;
7     using dst_t = typename gpuTraits::ChannelType;
8     using vec_t = typename gpuTraits::VectorType;
9
10    dst_t r = loadChannel<gpuTraits::channels[0], dst_t, P>(px);
11    dst_t g = loadChannel<gpuTraits::channels[1], dst_t, P>(px);
12    dst_t b = loadChannel<gpuTraits::channels[2], dst_t, P>(px);
13    dst_t a = loadChannel<gpuTraits::channels[3], dst_t, P>(px);
14    return vec_t{ r,g,b,a };
15 }
```

Conway's game of life



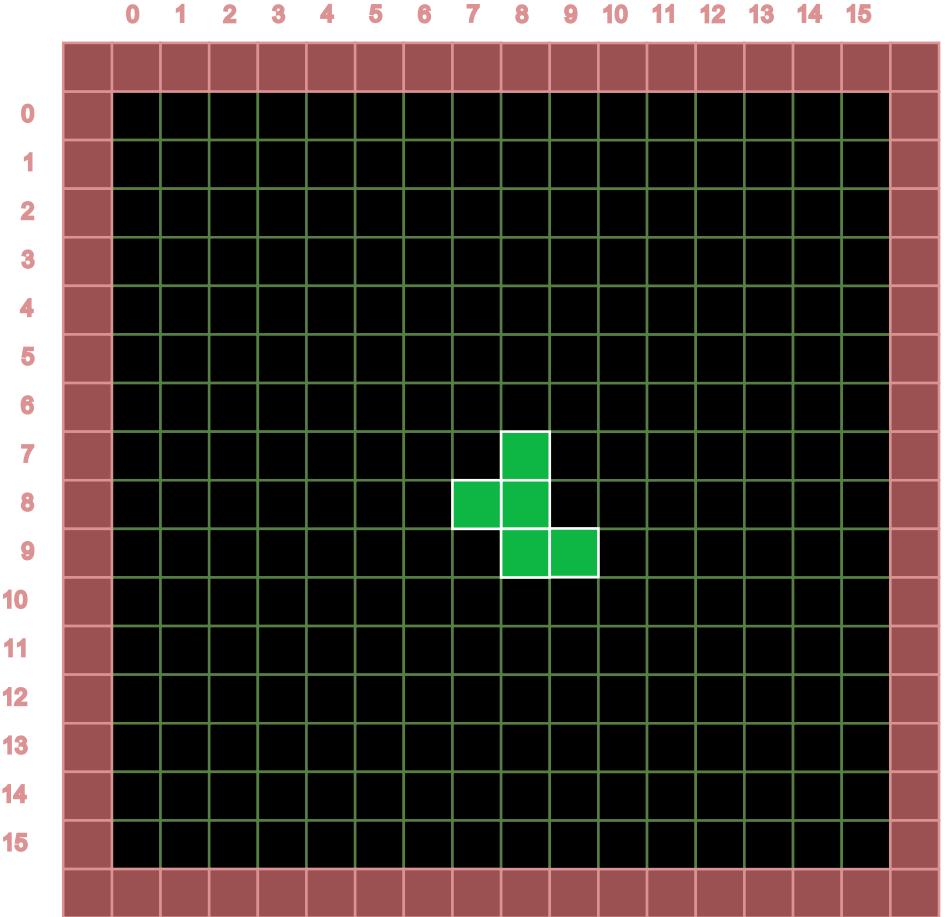
Compute Shader – C++

```
1 struct [[clang::annotate("kernel")]] GameOfLifeKernel
2 {
3     ImageBinding<gpu::R8UI, tc::Dim::D2, cpu::R8UI, 0> inData;
4     ImageBinding<gpu::R8UI, tc::Dim::D2, cpu::R8UI, 1> outData;
5     Uniform<int32_t, 1> pad{ 1 };
6     std::array<ivec2, 8> kIndices {
7         ivec2{-1,-1}, ivec2{0,-1}, ivec2{1,-1},
8         ivec2{-1, 0},           ivec2{1, 0},
9         ivec2{-1, 1}, ivec2{0, 1}, ivec2{1, 1}
10    };
11    void main() {
12        uvec2 gId = tc::gl_GlobalInvocationID["xy"_sw];
13        ivec2 coordinate = ivec2(gId.x+pad, gId.y+pad);
14        uint n = 0;
15        bool alive = imageLoad(inData, coordinate).x;
16        for (int ki = 0; ki < kernelIndices.size(); ++ki) {
17            n += imageLoad(inData, coordinate + kIndices[ki]).x;
18        }
19        bool newState = (n == 3) || (alive && n == 2);
20        imageStore(outData, coordinate, uvec4(newState));
21    }
};
```

Compute Shader – GLSL

```
1 struct [[clang::annotate("kernel")]] GameOfLifeKernel
2 {
3     layout(binding=0,r8ui) uniform uimage2D inData; 0> inData;
4     layout(binding=1,r8ui) uniform uimage2D outData;1> outData;
5     layout(location=1) uniform int pad;
6     const ivec2 kIndices[8] = ivec2[8] (
7         ivec2(-1,-1), ivec2(0,-1), ivec2(1,-1),,
8         ivec2(-1, 0),           ivec2(1, 0),,
9         ivec2(-1, 1), ivec2(0, 1), ivec2(1, 1})
10    );
11
12    void main() {
13        uvec2 gId = gl_GlobalInvocationID.xy;"xy"_sw];
14        ivec2 coordinate = ivec2(gId.x+pad, gId.y+pad);
15        uint n = 0;
16        bool alive = bool( imageLoad(inData, coordinate).x );
17        for (int ki = 0; ki < kIndices.size(); ++ki) {
18            n += imageLoad(inData, coordinate + kIndices[ki]).x;
19        }
20        bool newState = (n == 3) || (alive && n == 2);
21        imageStore(outData, coordinate, uvec4(newState));
22    }
};
```

Game of life GPU



$$C^{\infty} \rightarrow C^{\infty}$$

Conclusions

- 3 vertical slices → now go horizontal
- Generic programming reduces code:
 - Traits
 - Constexpr
 - Concepts
 - ...
- Compile time programming → no performance penalty for very flexible code
- I need to upgrade to C++26
 - `float16_t`
 - Dimensional spans
 - Reflection!

