



Implement Standard Library

Design Decisions, Optimisations and Testing in
Implementing Libc++

HUI XIE



20
25 | 
September 13

What is libc++

- A Standard Library Implementation
- Part of LLVM Project
- Ships with clang

Contents

- Various Optimisations in the Library
 - `std::expected`
 - `stop_token`
 - `ranges::for_each`, `ranges::copy`
 - `flat_map`
- Testing

About Me

- Hui Xie (GitHub @huixie90)
- Software Developer @Qube Research & Technologies
- libc++ contributor
- BSI (WG21 UK National Body) member
- WG21 (C++ Standard Committee) member
- Active in SG9 (Ranges Study Group)
- Write C++ Proposals

std::expected

```
1 std::expected<MyData, MyError> compute();
2
3 int main() {
4     std::expected<MyData, MyError> res = compute();
5
6     if (res.has_value()) {
7         // some code that uses res.value();
8     } else {
9         // some code that handles res.error();
10    }
11 }
```

std::expected

```
class Foo {
    int i;
    char c;
    bool b;
};

enum class ErrCode : int { Err1, Err2, Err3 };

static_assert(sizeof(Foo) == 8); // on most implementations

static_assert(sizeof(std::expected<Foo, ErrCode>) == ?);
```

std::expected in libstdc++

```
// gcc libstdc++
template <class Val, class Err>
class expected {
    union {
        Val val_;
        Err err_;
    };
    bool has_val_;
};

static_assert(sizeof(std::expected<Foo, ErrorCode>) == ?);
```

std::expected in libstdc++

```
1 // gcc libstdc++  
2 static_assert(sizeof(std::expected<Foo, ErrCode>) == 12);
```

```
1 int3 | int2 | int1 | int0 | char | bool | xxxx | xxxx | has_val | xxxx | xxxx | xxxx  
2 <----- Foo Data ----->  
3                                     <- Foo pad. ->  
4 <----- std::expected<Foo, ErrCode> Data ----->  
5                                     <- expected pad. ->  
6 <----- std::expected<Foo, ErrCode> ----->
```

std::expected in libc++

```
// clang libc++ simplified version
template <class Val, class Err>
class expected {
    union U {
        [[no_unique_address]] Val val_;
        [[no_unique_address]] Err err_;
    };
    [[no_unique_address]] U u_;
    bool has_val_;
};

static_assert(sizeof(std::expected<Foo, ErrorCode>) == ?);
```

The attribute-token no_unique_address specifies that a non-static data member is a potentially-overlapping subobject.

std::expected in libc++

```
1 int3 | int2 | int1 | int0 | char | bool | has_val | xxxx
2 <----- Foo Data ----->
3                                     <--  Foo pad. -->
4 <---- std::expected<Foo, ErrCode> Data ----->
5                                     <ex pd>
6 <----- std::expected<Foo, ErrCode> ----->
```

```
// clang libc++
static_assert(sizeof(std::expected<Foo, ErrCode>) == 8);
```

- What are the benefits?

sizeof(std::expected) Smaller, Better ?

- Smaller Memory Footprint
- Better Cache Locality
- std::expected is likely to be used as a return type

std::expected return Type

```
std::expected<Foo, ErrCode> compute() { return Foo{}; }
```

```
# gcc libstdc++
compute():
    mov    DWORD PTR [rsp-24], 0
    xor    eax, eax
    mov    BYTE PTR [rsp-24], 1
    mov    ecx, DWORD PTR [rsp-24]
    mov    rdx, rcx
    ret
```

```
# clang libc++
compute():
    movabs  rax, 281474976710656
    ret
```

Takeaway 1

- Reuse tail padding with `[[no_unique_address]]`
 - Including the padding of an empty type

A Nasty Bug

```
1 template <class Val, class Err>
2 class expected {
3     union U {
4         [[no_unique_address]] Val val_;
5         [[no_unique_address]] Err err_;
6     };
7     [[no_unique_address]] U u_;
8     bool has_val_;
9
10 public:
11     expected(const expected& other)
12         : has_val_(other.has_val_) {
13         if (has_val_) {
14             std::construct_at(std::addressof(u_.val_), other.u_.val_);
15         } else {
16             std::construct_at(std::addressof(u_.err_), other.u_.err_);
17         }
18     }
19 }
```

A Nasty Bug

```
1 int main() {
2     std::expected<Foo, int> e1(Foo{});
3     std::expected e2(e1);
4     assert(e2.has_value());
5 }
```

```
output.s: /app/example.cpp:15: int main(): Assertion `e2.has_value()' failed.
Program terminated with signal: SIGSEGV
```

Zero-Initialisation

To zero-initialize an object or reference of type T means: ... if T is a (possibly cv-qualified) non-union class type, its padding bits ([basic.types.general]) are initialized to zero bits and ...

A Nasty Bug

```
1 template <class Val, class Err>
2 class expected {
3     union U {
4         [[no_unique_address]] Val val_;
5         [[no_unique_address]] Err err_;
6     };
7     [[no_unique_address]] U u_;
8     bool has_val_;
9
10 public:
11     expected(const expected& other)
12         : has_val_(other.has_val_) {
13         if (has_val_) {
14             std::construct_at(std::addressof(u_.val_), other.u_.val_);
15         } else {
16             std::construct_at(std::addressof(u_.err_), other.u_.err_);
17         }
18     }
19 }
```

Takeaway 2

- Don't mix [[no_unique_address]] with manual lifetime management (union, construct_at, placement-new).

stop_source, stop_token, stop_callback

```
1 std::stop_source stop_src;
2
3 // Thread 1
4 std::stop_token token = stop_src.get_token();
5 while(!token.stop_requested()) {
6     // do some work
7 }
8
9 // Thread 2
10 std::stop_token token = stop_src.get_token();
11 std::stop_callback cb(token, [] { /*clean up*/ });
12
13 // Thread 3
14 std::stop_token token = stop_src.get_token();
15 std::stop_callback cb(token, [] { /* do stuff */ });
16
17 // Thread 4
18 stop_src.request_stop();
```

Under The Hood

```
1 class stop_token {
2     std::shared_ptr<__stop_state> state_;
3 };
4
5 class stop_source {
6     std::shared_ptr<__stop_state> state_;
7 };
8
9 template <class Callback>
10 class stop_callback {
11     [[no_unique_address]] Callback callback_;
12     std::shared_ptr<__stop_state> state_;
13 };
```

But What About The Shared State?

- `stop_requested()`: needs a flag to hold whether a stop was requested
- `stop_possible()`: needs to count how many `stop_sources` exist
- `stop_callback`: needs to store a list of refs to all the `stop_callbacks`
- `stop_callback`: needs to synchronise the list of callbacks
 - major requirement is that the whole thing is `noexcept`, ruling out `mutex`
- The state needs a ref count to manage its lifetime

A Naive Implementation

```
1 class __stop_state {
2     std::atomic<bool> stop_requested_;
3     std::atomic<unsigned> stop_source_count_; // for stop_possible()
4     std::list<stop_callback*> stop_callbacks_;
5     std::mutex list_mutex_;
6 };
```

```
static_assert(sizeof(__stop_state) == 72);
```

```
static_assert(sizeof(stop_token) == 16);
```

libc++'s Implementation

```
1 class __stop_state {
2     // The "callback list locked" bit implements a 1-bit lock to guard
3     // operations on the callback list
4     //
5     //      31 - 2      |  1      |  0
6     // stop_source counter  | callback list locked | stop_requested |
7     atomic<uint32_t> state_ = 0;
8
9     // Reference count for stop_token + stop_callback + stop_source
10    atomic<uint32_t> ref_count_ = 0;
11
12    // Lightweight intrusive non-owning list of callbacks
13    // Only stores a pointer to the root node
14    __intrusive_list_view<stop_callback_base> callback_list_;
15};
```

```
static_assert(sizeof(__stop_state) == 16);
```

```
static_assert(sizeof(stop_token) == 8);
```

Takeaway 3

- Sometimes we can reuse unused bits to save space
- Look for existing padding bytes for free storage

Segmented Iterators for_each


```
1 std::deque<int> d = ...
2
3 // 1
4 for (int& i : d) {
5     i = std::clamp(i, 200, 500);
6 }
7
8 // 2
9 std::ranges::for_each(d, [](int& i) {
10     i = std::clamp(i, 200, 500);
11});
```

Benchmarking Iteration

Benchmark	for loop	for_each	speed up
[for_loop vs. for_each]/32	12.5 ns	3.69 ns	3.4x
[for_loop vs. for_each]/8192	2973 ns	259 ns	11.5x
[for_loop vs. for_each]/65536	24221 ns	3327 ns	7.3x

clang20 -O3 cpu: M4 MAX 16 cores, memory: 48GB

`std::deque`



std::deque for Loop

```
1 deque_iterator begin = d.begin();
2 deque_iterator end = d.end();
3 for ( ; begin != end; ++begin ) {
4     int& i = *begin;
5     i = std::clamp(i, 200, 500);
6 }
```

```
1 deque_iterator& operator++() {
2     if (++elem_ptr_ - *block_iter_ == block_size) {
3         ++block_iter_;
4         elem_ptr_ = *block_iter_;
5     }
6     return *this;
7 }
```

Segmented Iterators Concept

- Segmented Iterators and Hierarchical Algorithms, Matt Austern
- We introduced this Segmented Iterator concept into libc++
- Iterators over ranges of sub-ranges
- Allows algorithms to operate over these multi-level iterators natively

for_each_segment

```
1 template <class SegmentedIterator, class Func>
2 void __for_each_segment(SegmentedIterator first, SegmentedIterator last, Func func) {
3     using Traits = SegmentedIteratorTraits<SegmentedIterator>;
4     auto seg_first = Traits::segment(first);
5     auto seg_last = Traits::segment(last);
6     // some code handles the first segment ...
7     // iterate over the segments
8     while (seg_first != seg_last) {
9         func(Traits::begin(seg_first), Traits::end(seg_first));
10        ++seg_first;
11    }
12    // some code handles the last segment ...
13 }
```

```
1 template <class SegmentedIterator, class Func>
2     requires is_segmented_iterator<SegmentedIterator>
3 void for_each(SegmentedIterator first, SegmentedIterator last, Func func) {
4     std::__for_each_segment(first, last, [&](auto inner_first, auto inner_last) {
5         std::for_each(inner_first, inner_last, func);
6     });
7 }
```

Code Generation Comparison

for loop

ranges::for_each

x86-64 clang (trunk) -O3 -stdlib=libc++ -std=c++26

```
25     cmp    rcx,  rax
26     je     .LBB0_6
27 BB0_3:
28     mov    r8d, dword ptr [rcx]
29     cmp    r8d, 201
30     cmovl r8d, esi
31     cmp    r8d, 500
32     cmovge r8d, edi
33     mov    dword ptr [rcx], r8d
34     add    rcx, 4
35     mov    r8, rcx
36     sub    r8, qword ptr [rdx]
37     cmp    r8, 4096
38     jne    .LBB0_5
39     mov    rcx, qword ptr [rdx +
40     add    rdx, 8
41     jmp    .LBB0_5
42 BB0_6:
43     ret
```

x86-64 clang (trunk) -O3 -stdlib=libc++ -std=c++26

```
51 BB0_10:
52     movdqu xmm2, xmmword ptr [rcx]
53     movdqu xmm3, xmmword ptr [rcx]
54     movdqa  xmm4, xmm2
55     pcmppgtd xmm4, xmm0
56     pand   xmm2, xmm4
57     pandn  xmm4, xmm0
58     por    xmm4, xmm2
59     movdqa  xmm2, xmm3
60     pcmppgtd xmm2, xmm0
61     pand   xmm3, xmm2
62     pandn  xmm2, xmm0
63     por    xmm2, xmm3
64     movdqa  xmm3, xmm1
65     pcmppgtd xmm3, xmm4
66     pand   xmm4, xmm3
67     pandn  xmm3, xmm1
68     por    xmm3, xmm4
```

Optimizing ranges::copy

```
1 std::vector<std::vector<int>> v = ...;
2 std::vector<int> out;
3
4 // 1
5 out.reserve(total_size);
6 for (const auto& inner : v) {
7     for (int i: inner) {
8         out.push_back(i);
9     }
10 }
11
12 // 2
13 out.resize(total_size);
14 std::ranges::copy(v | std::views::join, out.begin());
```

Benchmarking ranges::copy

Benchmark	for loop	copy	speed up
[for_loop vs. copy]/32	490 ns	322 ns	1.5x
[for_loop vs. copy]/8192	160632 ns	63372 ns	2.5x
[for_loop vs. copy]/65536	1066403 ns	208083 ns	5.1x

join_view::iterator is a Segmented Iterator

```
1 template <class Iter, class OutIter>
2   requires is_segmented_iterator<Iter>
3 pair<Iter, OutIter> __copy(Iter first, Iter last, OutIter result) {
4   std::__for_each_segment(first, last, [&](auto inner_first, auto inner_last) {
5     result = std::__copy(inner_first, inner_last, std::move(result)).second;
6   });
7   return std::make_pair(last, std::move(result));
8 }
```

```
1 template <class In, class Out>
2   requires can_lower_copy_assignment_to_memmove<In, Out>
3 pair<In*, Out*> __copy(In* first, In* last, Out* result) {
4   std::memmove(result, first, last - first);
5   return std::make_pair(last, result + n);
6 }
```

Takeaway 4

- We constantly add optimisations to algorithms. Please use them
- We use general concepts to capture optimisation opportunities
- Optimisations are often generic and can compose with each other

flat_map Insertion

```
// flat_map<int, double>
class flat_map {
    std::vector<int> keys_; // always sorted
    std::vector<double> values_;
    [[no_unique_address]] std::less<int> compare_;
};
```

```
1 std::flat_map<int, double> m1 = ...;
2 std::flat_map<int, double> m2 = ...;
3
4 // Insert the elements of m2 into m1
5 for (const auto& [key, val] : m2) {
6     m1.emplace(key, val);
7 }
```

What is the time complexity?

flat_map::insert_range

```
template<container-compatible-range<value_type> R>
constexpr void insert_range(R&& rg);
```

- Complexity: $N + M \log M$, where N is `size()` before the operation and M is `ranges::distance(rg)`.

```
1 std::flat_map<int, double> m1 = ...;
2 std::flat_map<int, double> m2 = ...;
3
4 // Insert the elements of m2 into m1
5 m1.insert_range(m2);
```

P3567 insert_range(sorted_unique, rg)

```
template<container-compatible-range<value_type> R>
void insert_range(sorted_unique_t, R&& rg);
```

- Complexity: Linear in N, where N is size() after the operation.

```
1 std::flat_map<int, double> m1 = ...;
2 std::flat_map<int, double> m2 = ...;
3
4 // Insert the elements of m2 into m1
5 m1.insert_range(std::sorted_unique, m2);
```

flat_map::insert_range Implementation

```
1 template<container-compatible-range<value_type> R>
2 constexpr void insert_range(R&& rg) {
3
4     __append(ranges::begin(rg), ranges::end(rg)); // O(M)
5
6     auto zv = ranges::views::zip(keys_, values_);
7     ranges::sort(zv.begin() + old_size, zv.end()); // O(MLogM)
8
9     ranges::inplace_merge(zv.begin(), zv.begin() + old_size, zv.end()); // O(M+N)
10
11    auto dup_start = ranges::unique(zv).begin(); // O(M+N)
12    __erase(dup_start); // O(M+N)
13 }
```

How to append ?

```
1 template <class InputIterator, class Sentinel>
2 void __append(InputIterator first, Sentinel last) {
3     for (; first != last; ++first) {
4         std::pair<Key, Val> kv = *first;
5         keys_.insert(keys_.end(), std::move(kv.first));
6         values_.insert(values_.end(), std::move(kv.second));
7     }
8 }
```

- This misses existing optimisations in `vector::insert(pos, first, last)`
- But we are given a range of “pairs”, not two ranges
- Can we reuse these optimizations in some cases?

Introducing a Concept `product_iterator`

- iterators that *aggregate* multiple underlying iterators
- `flat_map::iterator` is `product_iterator`

```
1 template <class Iter>
2   requires is_product_iterator_of_size<Iter, 2>
3 void __append(Iter first, Iter last)
4 {
5   using Traits = product_iterator_traits<Iter>;
6
7   keys_.insert(keys_.end(),
8     Traits::template get_iterator<0>(first),
9     Traits::template get_iterator<0>(last));
10
11  values_.insert(values_.end(),
12    Traits::template get_iterator<1>(first),
13    Traits::template get_iterator<1>(last));
14 }
```

Benchmarking insert_range

Benchmark	insert_pair	product_iterator	speed up
[insert_range]/32	149 ns	74 ns	2.0x
[insert_range]/8192	26682 ns	2995 ns	8.9x
[insert_range]/65536	226235 ns	27844 ns	8.1x

Are there any other product_iterator?

```
1 std::flat_map<int, double> m = ...;
2 std::vector<int> newKeys = ...;
3 std::vector<double> newValues = ...;
4
5 // Insert newKeys and newValues into m
```



```
m.insert_range(std::views::zip(newKeys, newValues));
```

`zip_view::iterator` is also a `product_iterator`

Takeaway 5

- Use the most precise API for what you're trying to achieve
 - `insert_range` instead of `insert` in a loop
 - Use `sorted_unique` if the inputs are already sorted
- Use library facilities (e.g `views::zip`) to benefit from concept-based optimisations

Testing in libc++

- We test almost every single word in the standard

What Shall We Test?

```
constexpr expected(expected&& rhs) noexcept(see below);
```

- Constraints:
 - `is_move_constructible_v<T>` is true and
 - `is_move_constructible_v<E>` is true.
- Effects: If `rhs.has_value()` is true, direct-non-list-initializes `val` with `std::move(*rhs)`. Otherwise, direct-non-list-initializes `unex` with `std::move(rhs.error())`.
- Postconditions: `rhs.has_value()` is unchanged; `rhs.has_value() == this->has_value()` is true.
- Throws: Any exception thrown by the initialization of `val` or `unex`.
- Remarks: The exception specification is equivalent to `is_nothrow_move_constructible_v<T> && is_nothrow_move_constructible_v<E>`.
- This constructor is trivial if
 - `is_trivially_move_constructible_v<T>` is true and
 - `is_trivially_move_constructible_v<E>` is true.

Testing constexpr

```
1 constexpr bool test() {
2     // Test point 1
3     {
4         std::expected<MoveOnly, int> e1 = ...;
5         std::expected<MoveOnly, int> e2 = std::move(e1);
6         assert(e2.has_value());
7         assert(e2.value() == ... );
8     }
9     // more test points
10
11    return true;
12 }
13
14 int main() {
15     test();
16     static_assert(test());
17 }
```

- Share compile time and run time tests

Testing Constraints

```
constexpr expected(expected&& rhs) noexcept(see below);
```

- Constraints: `is_move_constructible_v<T>` is true and `is_move_constructible_v<E>` is true.
- Constraints translate to requires

```
1 struct Foo { Foo(Foo&&) = delete;};
2
3 static_assert(std::is_move_constructible_v<std::expected<int, int>>);
4 static_assert(!std::is_move_constructible_v<std::expected<Foo, int>>);
5 static_assert(!std::is_move_constructible_v<std::expected<int, Foo>>);
6 static_assert(!std::is_move_constructible_v<std::expected<Foo, Foo>>);
```

Testing Mandates

```
template<class U = remove_cv_t<T>> constexpr T value_or(U&& v) const &;
```

- Mandates: `is_copy_constructible_v<T>` is true and `is_convertible_v<U, T>` is true.
- Mandates translate to `static_assert`
- Test that usage that violates the mandate should not compile

```
const std::expected<NonCopyable, int> f1{5};  
f1.value_or(5); // expected-note{{in instantiation of function template ...}}  
// expected-error-re@*: {{static assertion failed {{.*}}value_type has to be copy ...}}
```

```
-Xclang -verify
```

Testing {Hardened} Preconditions

```
constexpr const T* operator->() const noexcept;
constexpr T* operator->() noexcept;
```

- Hardened preconditions: `has_value()` is true.
- Preconditions/Hardened preconditions translate to `assert/contract_assert/other assert` macros

```
std::expected<int, int> e{std::unexpected, 5};
TEST_LIBCPP_ASSERT_FAILURE(e.operator->(),
    "expected::operator-> requires the expected to contain a value");
```

- `TEST_LIBCPP_ASSERT_FAILURE` installs assertion handler, forks the process and matches the child process `stderr` with the expected errors

Takeaway 6

- Share the same test code for `constexpr` and runtime
- Use negative tests to ensure that things fail **as expected**
- `-Xclang -verify` is very useful to test `static_asserts`

Contribute to libc++

- Getting Started
- Github Issues

QRT is Hiring

- <https://www.qube-rt.com/careers/>



If You Do Mix `[no_unique_address]` with `union`, `std::construct_at`, or `placement new`

- Be ready for bugs

Dude, Where is my_char ?

```
1 struct MyStruct {  
2     [[no_unique_address]] std::expected<Foo, ErrorCode> my_foo;  
3     char my_char;  
4 };  
5  
6 MyStruct s{.my_foo = Foo{}, .my_char = 'x'};  
7  
8 s.my_foo.emplace(Foo{});
```

```
1 // Assertion failure!  
2 assert(s.my_char == 'x');
```

Wait, is my_char in the Padding?

```
template <class Val, class Err>
class expected {
    union U {
        [[no_unique_address]] Val val_;
        [[no_unique_address]] Err err_;
    };
    [[no_unique_address]] U u_;
    bool has_val_;
};
```

```
1 int3 | int2 | int1 | int0 | char | bool | has_val | my_char
2 <----- Foo Data ----->
3                                     <---- Foo pad. ---->
4 <----- std::expected<Foo, ErrorCode> Data ----->
5                                     <ex pad.>
6 <----- std::expected<Foo, ErrorCode> ----->
7 <----- MyStruct ----->
```

```
1 s.my_foo.emplace(Foo{}); // It calls construct_at and clears the padding!
```

Stop Users from Reusing expected Paddings

```
1 template <class Val, class Err>
2 class expected {
3     struct repr {
4         union U {
5             [[no_unique_address]] Val val_;
6             [[no_unique_address]] Err err_;
7         };
8         [[no_unique_address]] U u_;
9         bool has_val_;
10    };
11
12    repr repr_; // No [[no_unique_address]] here
13};
```

- `[[no_unique_address]]` is not transitive

It Works

```
1 struct MyStruct {  
2     [[no_unique_address]] std::expected<Foo, ErrCode> my_foo;  
3     char my_char;  
4 };
```

```
1 int3 | int2 | int1 | int0 | char | bool | has_val | xxxx | my_char | xxxx | xxxx | xxxx  
2 <----- Foo Data ----->                                <--- Foo pad. --->  
3                                         <----- repr Data ----->          <repr pad.>  
4                                         <----- std::expected<Foo, ErrCode> Data ----->  
5                                         <----- std::expected<Foo, ErrCode> ----->  
6                                         <----- MyStruct Data ----->           <-MyStruct padding ->  
7                                         <----- MyStruct ----->  
8  
9  
10
```

But Can We do Better?

```
1 struct MyStruct {  
2     [[no_unique_address]] std::expected<int, ErrCode> my_int;  
3     char my_char;  
4 };
```

```
1 int3 | int2 | int1 | int0 | has_val | xxxx | xxxx | xxxx | my_char | xxxx | xxxx | xxxx  
2 <----- repr Data ----->  
3                                     <--- repr padding --->  
4 <----- std::expected<int, ErrCode> Data ----->  
5 <----- std::expected<int, ErrCode> ----->  
6 <----- MyStruct Data ----->  
7                                     <-MyStruct padding ->  
8 <----- MyStruct ----->
```

```
1 s.my_int.emplace(42); // construct_at(int*, int) is very safe! int has no padding
```

Final Fix

```
1 template <class Val, class Err>
2 struct repr {
3     union U {
4         [[no_unique_address]] Val val_;
5         [[no_unique_address]] Err err_;
6     };
7     [[no_unique_address]] U u_;
8     bool has_val_;
9 };
10
11 template <class Val, class Err>
12 struct expected_base {
13     repr<Val, Err> repr_; // no [[no_unique_address]]
14 };
15 template <class Val, class Err> requires bool_is_not_in_padding
16 struct expected_base {
17     [[no_unique_address]] repr<Val, Err> repr_;
18 };
19
20 template <class Val, class Err>
21 class expected : expected_base<Val, Err> {};
```