

+ 25

Instrumenting the Stack

Strategies for End-to-end Sanitizer Adoption

DAMIEN BUHL



20
25 | A graphic of two white mountain peaks with a yellow square at the top of the taller one.
September 13 - 19

EngFlow

- Co-Founded by the lead developer of Bazel, Google's open-source build system
- Acquired tipi.build the team behind CMake RE
- Co-founded by Boost Contributors, CMake lovers and Standard C++ ISO members

Texas-based company acquires tipi.build

13.03.2025



Google Cloud
Partner

CMake RE

Drop-in for CMake

```
cmake \
-G Ninja -S . -B build/ \
-DCMAKE_TOOLCHAIN_FILE=toolchains/linux-ubuntu24.04-clang20.cmake \
-DCMAKE_BUILD_TYPE=Release

cmake --build build/
```

CMake RE

Drop-in for CMake

```
cmake \
  -G Ninja -S . -B build/ \
  -DCMAKE_TOOLCHAIN_FILE=toolchains/linux-ubuntu24.04-clang20.cmake \
  -DCMAKE_BUILD_TYPE=Release

cmake --build build/
```

CMake with Remote, Distributed, Containerized Execution and Caching

```
cmake[--distributed] [--remote|--host] \
  -G Ninja -S . -B build/ \
  -DCMAKE_TOOLCHAIN_FILE=toolchains/linux-ubuntu24.04-clang20.cmake \
  -DCMAKE_BUILD_TYPE=Release

cmake[--re [--jobs 1000+]] --build build/
```

CppCon 2023 : Delivering Safe C++

``Are we there yet?''

- We have come a long way
 - From “classic C”
 - From “C with Classes”
 - From C++11
- We can write type-and-resource safe C++
 - Use contemporary C++
 - Avoid C-style and 1980s style C++
- C++ Core Guidelines
 - Directions, rules, and some enforcement
 - But not standardized
 - Uneven enforcement across implementations
- We need to standardize “Profiles”
 - And get implementations deployed
 - Finally reach type-and-resource safe C++!

Stroustrup - C++ safety - CppCon - October 2023



Are we safe yet ?

- Not fully yet, but in practice it's doable.
- Good tooling exists, it is just waiting for teams to use it.
- Requires more build types and test pipelines.

Practical Ways of getting Safety in C++

Productivity vs Safety

- Setting up tools
- Good test code coverage
- Requires more build types, means longer iterations in getting code merged

Safety Toolbox

Static Analysis : clang-tidy

1

Using Sanitizers : MSan, ASan + LSan + UBSan, TSan

2

TSA (Thread Safety Analysis)

3

clang-tidy : Static Analysis

Default set of checks

- clang-analyzer-*: Clang Static Analyzer detect potential bugs and vulnerabilities
- cppcoreguidelines-*: Enforces the C++ Core Guidelines
- bugprone-*: Identify common programming mistakes that often lead to bugs

→ <https://clang.llvm.org/extra/clang-tidy/checks/list.html>

→ Parses code twice

- Once for clang-tidy
- Once for compilation

→ Checks can be silenced in C++ code with `// NOLINT`

clang-tidy : Static Analysis



```
# Enable clang-tidy checks
set(CMAKE_CXX_CLANG_TIDY clang-tidy)
set(CMAKE_C_CLANG_TIDY clang-tidy)
```

MSan : MemorySanitizer

Detector of Uninitialized Memory

→ Shadow Memory

1 corresponding bit for each application's memory bits.

Mark Poisoned (1) on allocation, unpoisoned (0) on initialization.

→ Compiler Instrumentation on

- Memory allocation
- Initialization (e.g. assignments, memset, memcpy)
- Destructors to catch use-after-destruction
- Memory Read, dereference, function call

→ Doubles the execution time

- Faster than Valgrind which slows down execution by a factor 20
- Linux only

MSan : How to get full-instrumentation



```
RUN mkdir -p /llvm-project && \
    git clone --branch llvmorg-20.1.2 --depth 1 https://github.com/llvm/llvm-
project.git /llvm-project && \
    cmake \
    -DCMAKE_BUILD_TYPE=Release \
    -G Ninja \
    -S /llvm-project/runtimes \
    -B /llvm-project/build-msan \
    -DCMAKE_INSTALL_PREFIX=/usr/local/instrumented/msan \
    -DLLVM_USE_SANITIZER=MemoryWithOrigins \
    -DLLVM_ENABLE_RUNTIMES='libcxx;libcxxabi' \
    -DLIBCXXABI_USE_LLVM_UNWINDER=Off && \
    cmake --build /llvm-project/build-msan --target install && \
    rm -rf /llvm-project/
```

ENV MSAN_SYMBOLIZER_PATH=/usr/bin/llvm-symbolizer

MSan : CMAKE_TOOLCHAIN_FILE



```
# Link to MSAN instrumented libc++
add_compile_options($<$<COMPILE_LANGUAGE:C,CXX>:-nostdlib++>)
add_link_options($<$<COMPILE_LANGUAGE:C,CXX>:-nostdlib++>)
link_libraries("/usr/local/instrumented/msan/lib/libc++.so.1.0")
link_libraries("/usr/local/instrumented/msan/lib/libc++abi.so.1.0")
add_compile_options($<$<COMPILE_LANGUAGE:C,CXX>:-stdlib=libc++>)
add_link_options($<$<COMPILE_LANGUAGE:C,CXX>:-stdlib=libc++>)

# Compilation and linker flags to ensure proper sanitizer support
set(CMAKE_POSITION_INDEPENDENT_CODE ON)
set(MSAN_FLAGS
    -fsanitize=memory
    -fsanitize-memory-track-origins
)
add_compile_options(
    $<$<COMPILE_LANGUAGE:C,CXX>:-fno-omit-frame-pointer>
    ${MSAN_FLAGS})
add_link_options(
    $<$<COMPILE_LANGUAGE:C,CXX>:-fno-omit-frame-pointer>
    ${MSAN_FLAGS})
```

MSan : Let's take an example



```
int compute_array_sum(const std::array<int, 10> arr) {
    int sum=0;
    for (auto v : arr) { sum+=v; }
    return sum;
}
```



```
FROM tipibuild/tipi-ubuntu-2404-clang20@sha256:...
```

```
# Preinstall dependencies
RUN mkdir -p /compute_array_sum && \
    git clone ...compute_array_sum.git /compute_array_sum && \
    cmake -DCMAKE_BUILD_TYPE=Release \
        -G Ninja \
        -S /compute_array_sum \
        -B /compute_array_sum/build \
        -DCMAKE_INSTALL_PREFIX=/usr/ && \
    cmake --build /compute_array_sum/build --target install
```

MSan : Let's take an example



```
find_package(compute_array_sum REQUIRED)

add_executable(main test/main.cpp)
target_link_libraries(main compute_array_sum::compute_array_sum)

add_test(NAME main COMMAND $<TARGET_FILE:main>

C++ #include <print>
#include <array>

#include <compute_array_sum.hpp>

int main() {
    std::array<int,10> arr;
    arr[5] = 0;
    std::println("Array sum: {} \n", compute_array_sum(arr));
    return 0;
}
```



MSan : Let's take an example



```
find_package(compute_array_sum REQUIRED)

add_executable(main test/main.cpp)
target_link_libraries(main compute_array_sum::compute_array_sum)

add_test(NAME main COMMAND $<TARGET_FILE:main>

C++ #include <print>
      #include <array>

      #include <compute_array_sum.hpp>

      int main() {
          std::array<int,10> arr;
          arr[5] = 0;
          std::println("Array sum: {} \n", compute_array_sum(arr));
          return 0;
      }
```

No instrumentation

MSan : Let's get instrumentation everywhere



```
- find_package(compute_array_sum REQUIRED)

+ FetchContent_Declare(
+ compute_array_sum
+ GIT_REPOSITORY .../unittest-compute_array_sum.git
+ GIT_TAG 9cbe670)
+ FetchContent_MakeHermetic(compute_array_sum)
+ HermeticFetchContent_MakeAvailableAtBuildTime(compute_array_sum)

add_executable(main test/main.cpp)
target_link_libraries(main compute_array_sum::compute_array_sum)

add_test(NAME main COMMAND $<TARGET_FILE:main>)
```

Instrumentation 

MSan : Let's get instrumentation everywhere



```
# HFC to the rescue
FetchContent_MakeHermetic(<name>

    HERMETIC_BUILD_SYSTEM cmake | autotools | openssl

    HERMETIC_FIND_PACKAGES <list of exposed dependencies>

    HERMETIC_TOOLCHAIN_EXTENSION <cmake code>

    HERMETIC_CMAKE_EXPORT_LIBRARY_DECLARATION <cmake code>

    SBOM_LICENSE "<SPDX-Identifier>"
    SBOM_SUPPLIER "<Author>"

)
```

★ Please Star it on Github : <https://github.com/tipi-build/hfc>

MSan Release vs Debug Symbols

```
==222900==WARNING: MemorySanitizer: use-of-uninitialized-value
#0 0xf in compute_array_sum(std::__1::array<int, 10ul>) (/main+0x1082bd)compute_array_sum.cpp:6:3
#1 0xf in main (/main+0xce5ed)test/main.cpp:12:36
#2 0xf in __libc_start_call_main csu/../sysdeps/nptl/libc_start_call_main.h:58:16
#3 0xf in __libc_start_main csu/../csu/libc-start.c:360:3
#4 0xf in _start (/main+0x33494)
```

Uninitialized value was created by an allocation of 'agg.tmp1' in the stack frame

```
#0 0xf in main (/main+0xce577)test/main.cpp:6
```

```
SUMMARY: MemorySanitizer: use-of-uninitialized-value (/main+0x1082bd)compute_array_sum.cpp:6:3 in
compute_array_sum(std::__1::array<int, 10ul>)
```

Choosing CMAKE_BUILD_TYPE

CMAKE_BUILD_TYPE

Specifies the build type on single-configuration generators (e.g. [Makefile Generators](#) or [Ninja](#)). Typical values include `Debug`, `Release`, `RelWithDebInfo` and `MinSizeRel`, but custom build types can also be defined.

This variable is initialized by the first `project()` or `enable_language()` command called in a project when a new build tree is first created. If the `CMAKE_BUILD_TYPE` environment variable is set, its value is used. Otherwise, a toolchain-specific default is chosen when a language is enabled. [The default value is often an empty string, but this is usually not desirable and one of the other standard build types is usually more appropriate.](#)

Choosing CMAKE_BUILD_TYPE

"": No Optimizations, No Debug Symbols

Release: Optimizations, No Debug Symbols

Debug: No Optimizations, Debug Symbols

MinSizeRel: Optimizations for Size, No Debug Symbols

RelWithDebInfo: Optimizations, Debug Symbols

→ Rationales

- Release & MinSizeRel have no Debug Symbols to ensure binaries only contain necessary data for operations
- Debug has no optimizations because optimizations can make debugging surprising

→ Typical Pitfalls when considering only Debug & Release

- Developers test most of their time in Debug
- Release crashes are undetectable

→ **Hot Take** : RelWithDebInfo is the only sane option

Handling Debug Symbols Bloat

Problem with many very long template<> symbols

```
GNU LD: relocation truncated to fit: R_X86_64_32 against `._debug_str'  
mold: (.debug_str): mergeable section too large
```

→ Can be avoided with :

Either `-g -gsplit-dwarf`

Or `-gsimple-template-names`

Debug Symbols Impacts

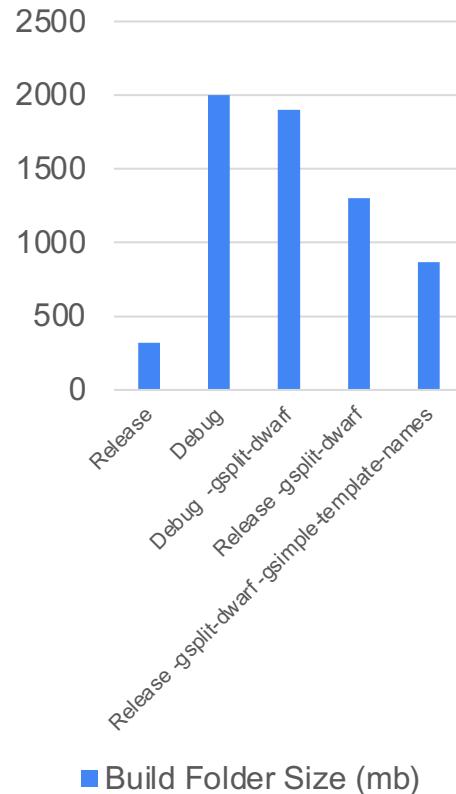
Building Boost 1.85.0 and Unit Tests with `-gsplit-dwarf`, generates a .o and a .dwo

.o

- .debug_info links to the .dwo
- Index for symbols availability in the .dwo
- .debug_gnu_pubnames
- .debug_gnu_pubtypes

.dwo

- .debug_info.dwo
- .debug_str.dwo
- .debug_rnglists.dwo
- .debug_loclists.dwo
- .debug_str_offsets.dwo
- .debug_abbrev.dwo



Precompiled : Valgrind as last resort

When **precompiled dependencies** are **unavoidable**

- Runtime Instrumentation via Dynamic Binary Instrumentation
- Slows down execution by 20x
- Complex Diagnostics



```
# Enable valgrind on CTest
set(CMAKE_TEST_LAUNCHER /usr/bin/valgrind;--track-origins=yes)
```

ASan + LSan + UBSan

AddressSanitizer + LeakSanitizer

- Overrides malloc() and free() at runtime
Same shadowing strategy than MSAN, checked on Memory Read.
- Slows down execution only by 2x
- Supported on Windows (Asan only), Linux, macOS

UndefinedBehaviourSanitizer

- Compiler Instrumentation
e.g. Adds non-null checks before any pointer dereference, check integer overflow on addition...
- Supported on Windows, Linux, macOS

ASan + LSan + UBSan : CMAKE_TOOLCHAIN_FILE



```
# Compilation and linker flags to ensure proper sanitizer support
set(CMAKE_POSITION_INDEPENDENT_CODE ON)

set(ASAN_FLAGS
    -fsanitize=address
)

set(UBSAN_FLAGS
    -fsanitize=undefined
    -fno-sanitize-merge
)

add_compile_options(
    ${${<}<COMPILE_LANGUAGE:C,CXX>:-fno-omit-frame-pointer>
    ${ASAN_FLAGS}
    ${UBSAN_FLAGS})
add_link_options(
    ${${<}<COMPILE_LANGUAGE:C,CXX>:-fno-omit-frame-pointer>
    ${ASAN_FLAGS}
    ${UBSAN_FLAGS})
```

TSan

ThreadSanitizer

→ Code Instrumentation

For every memory access (read or write) + synchronization primitive

→ Each thread is assigned a vector clock, on communication (e.g. via a mutex) the vector clock are synchronized

→ A data race is reported when 2 accesses to the same memory location have no happens-before relation, and one of them is a write

→ Performance slow down of 5x-15x

Setting up TSAN : CMAKE_TOOLCHAIN_FILE



```
# Compilation and linker flags to ensure proper sanitizer support
set (CMAKE_POSITION_INDEPENDENT_CODE ON)

add_compile_options($<$<COMPILE_LANGUAGE:C,CXX>:-fsanitize=thread>
add_link_options($<$<COMPILE_LANGUAGE:C,CXX>:-fsanitize=thread>

add_compile_options(
    $<$<COMPILE_LANGUAGE:C,CXX>:-fno-omit-frame-pointer>
)
add_link_options(
    $<$<COMPILE_LANGUAGE:C,CXX>:-fno-omit-frame-pointer>
)
```

TSan : Let's take an example



```
class Writer {
public:
    void write(std::string word) {
        words_.push_back(std::move(word));
    }
private:
    std::vector<std::string> words_;
};

void threadedwork(int thread_id, Writer *writer) {
    for (auto count : std::ranges::iota_view(0, 100)) {
        writer->write(std::format("Thread {} count {}", thread_id, count));
    }
}

Writer w;
std::thread t1{[&](){threadedwork(1,&w);}};
std::thread t2{[&](){threadedwork(2,&w);}};
...
```

TSan : Detects the data race

WARNING: ThreadSanitizer: data race (pid=11586)

```
Read of size 8 at 0x00016b07b128 by thread T2:
#0 std::__1::vector<std::__1::basic_string<char, std::__1::char_traits<char>, std::__1::allocator<char>>, std::__1::allocator<std::__1::basic_string<char, std::__1::char_traits<char>, std::__1::allocator<char>>>::push_back[abi:ne190102](std::__1::basic_string<char, std::__1::char_traits<char>, std::__1::allocator<char>>&&) <null> (exe1_tsan:arm64+0x1000020b8)
#1 writer::write(std::__1::basic_string<char, std::__1::char_traits<char>, std::__1::allocator<char>>)
<null> (exe1_tsan:arm64+0x1000018fc)
#2 threadedwork(int, Writer*) <null> (exe1_tsan:arm64+0x1000015f4)
#3 decltype(std::declval<void (*)(int, Writer*)>()(&std::declval<int>(), std::declval<Writer*>()))
std::__1::__invoke[abi:ne190102]<void (*)(int, Writer*), int, Writer*>(void (*&&)(int, Writer*), int&&, Writer*&&) <null> (exe1_tsan:arm64+0x100006788)
#4 void std::__1::__thread_execute[abi:ne190102]<std::__1::unique_ptr<std::__1::__thread_struct, std::__1::default_delete<std::__1::__thread_struct>>, void (*)(int, Writer*), int, Writer*, 2ul, 3ul>(std::__1::tuple<std::__1::unique_ptr<std::__1::__thread_struct, std::__1::default_delete<std::__1::__thread_struct>>, void (*)(int, Writer*), int, Writer*>&, std::__1::__tuple_indices<2ul, 3ul>) <null> (exe1_tsan:arm64+0x100006654)
#5 void*
std::__1::__thread_proxy[abi:ne190102]<std::__1::tuple<std::__1::unique_ptr<std::__1::__thread_struct, std::__1::default_delete<std::__1::__thread_struct>>, void (*)(int, Writer*), int, Writer*>>(&void*) <null> (exe1_tsan:arm64+0x100005868)
```

(+10 more screens)

Productivity & Safety

Really want to avoid safety bugs ?

Sanitizers requires you to add at least 3 new build types to a project existing targets

→ **New Build Types for :**

1. Msan
2. ASan + LSan + UBSan
3. TSan

+ Adding clang-tidy on top of your existing build

We optimize this, that's one of the main reasons why we developed our remote-execution engine in CMake RE.

Clang TSA

Thead Safety Analysis

- Annotations allowing compile-time checks for data races
- No Runtime Performance Impact
- Works on all platform with libc++

TSA : Improve our example



```
#include "tsa.h"

class writer {
public:
    void write(std::string word) {
        words_.push_back(std::move(word));
    }
private:
    std::mutex mutex_;
    std::vector<std::string> words_ GUARDED_BY(mutex_);
};

void threadedwork(int thread_id, writer *writer) {
    for (auto count : std::ranges::iota_view(0, 100)) {
        writer->write(std::format("Thread {} count {}", thread_id, count));
    }
}

Writer w;
std::thread t1{[](){threadedwork(1,&w);}}; ...
```

TSA : Improve our example

```
/usr/bin/c++ -D_LIBCPP_ENABLE_THREAD_SAFETY_ANNOTATIONS -Wthread-safety -std=c++20  
exe2.cpp -o exe2
```

```
exe2.cpp:13:9: warning: reading variable 'words_' requires holding mutex 'mutex_' [-  
Wthread-safety-analysis]
```

```
13 |         words_.push_back(std::move(word));  
|         ^
```

Clang TSA

Thead Safety Analysis

- Requires Clang and Compilation warning: `-wthread-safety`
- std:: types support: `-D_LIBCPP_ENABLE_THREAD_SAFETY_ANNOTATIONS`
- boost:: types support: `-DBOOST_THREAD_ENABLE_THREAD_SAFETY_ANALYSIS`

★ For Windows FetchContent libcxx :

- <https://github.com/Orphis/LibcxxFetchContentTest>

- ➡ Learn more about Clang TSA with Florent Castelli (@Orphis) at Stockholm Cpp Meetup
 - <https://www.youtube.com/watch?v=VAnIVjrDouA>

Getting all the sanitizers

→ Enhance your codebase safety today : CMake RE Local Containerized builds **is free** to use.

→ Launch your sanitizers builds from macOS, Windows, Linux

```
cmake-re -S . -B build/msan \
-DCMAKE_TOOLCHAIN_FILE=environments/linux-ubuntu-2404-clang20-msan.cmake
```

```
cmake-re --build build/msan --run-test main --monitor
```



Get the free cmake-re sanitizers environments and example :
<https://github.com/tipi-build/example-cmake-re-sanitizers>

→ For faster builds+tests with **-j1000+** cores, sign-up on tipi.build & contact us : hello@tipi.build !