



Modern C++ for Robust Bundle Adjustment

From Outliers to Optimization

VISHNU SUDHEER MENON



Outline of the talk

- What is Visual Reconstruction?
- Bundle adjustment in the large
- SFM pipeline
- Data Parser
- Outlier Rejection
- Optimization
- Results



Problem Statement

How do we design a visual reconstruction library that is adaptable, robust, scalable across datasets and optimization challenges?

- Dataset Diversity
Works across different dataset formats and sensor types
- Flexible Filtering
Plug-and-play support for methods like RANSAC, outlier rejection
- Optimizable Core
Bundle Adjustment that converges reliably
Efficient for both small and large datasets
- Scalability Requirement
Avoids runtime explosion as dataset size grows



What is Visual Reconstruction?

Recovering 3D structure of a scene from 2D images

- Infer geometry and structure of the world.
- Using primarily camera inputs.

Inputs/Outputs

- Inputs: single or multiple images, video sequences
- Outputs: 3D reconstruction (point clouds, meshes, dense models)

Applications

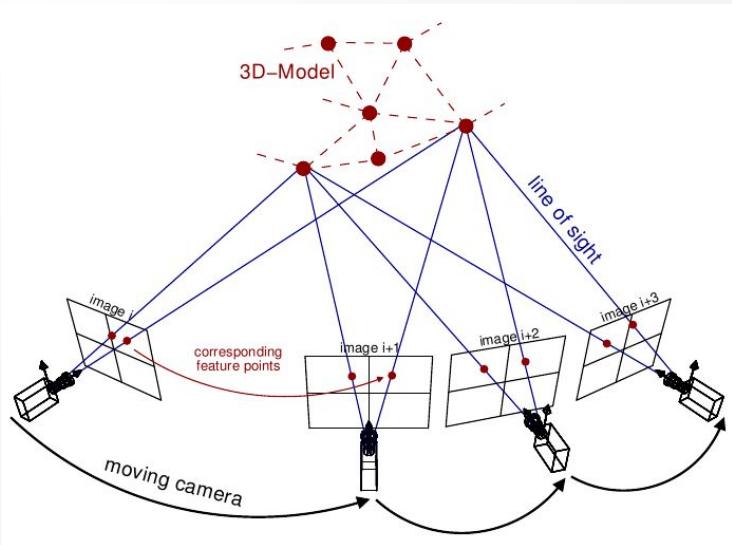
- Robotics & autonomous navigation
- Digital twins & mapping

Challenges

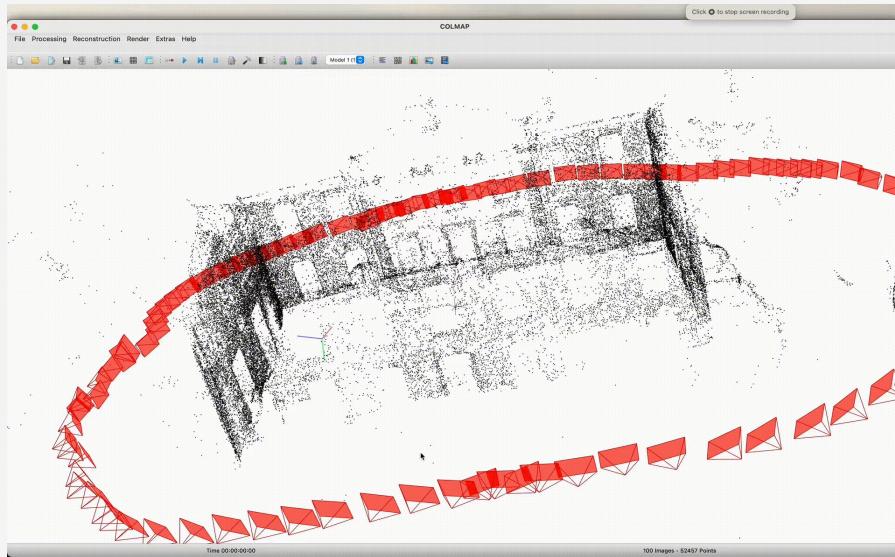
- Occlusions & missing data
- Noise and outliers in features & correspondences
- Large-scale optimization can be computationally expensive



What is Visual Reconstruction?



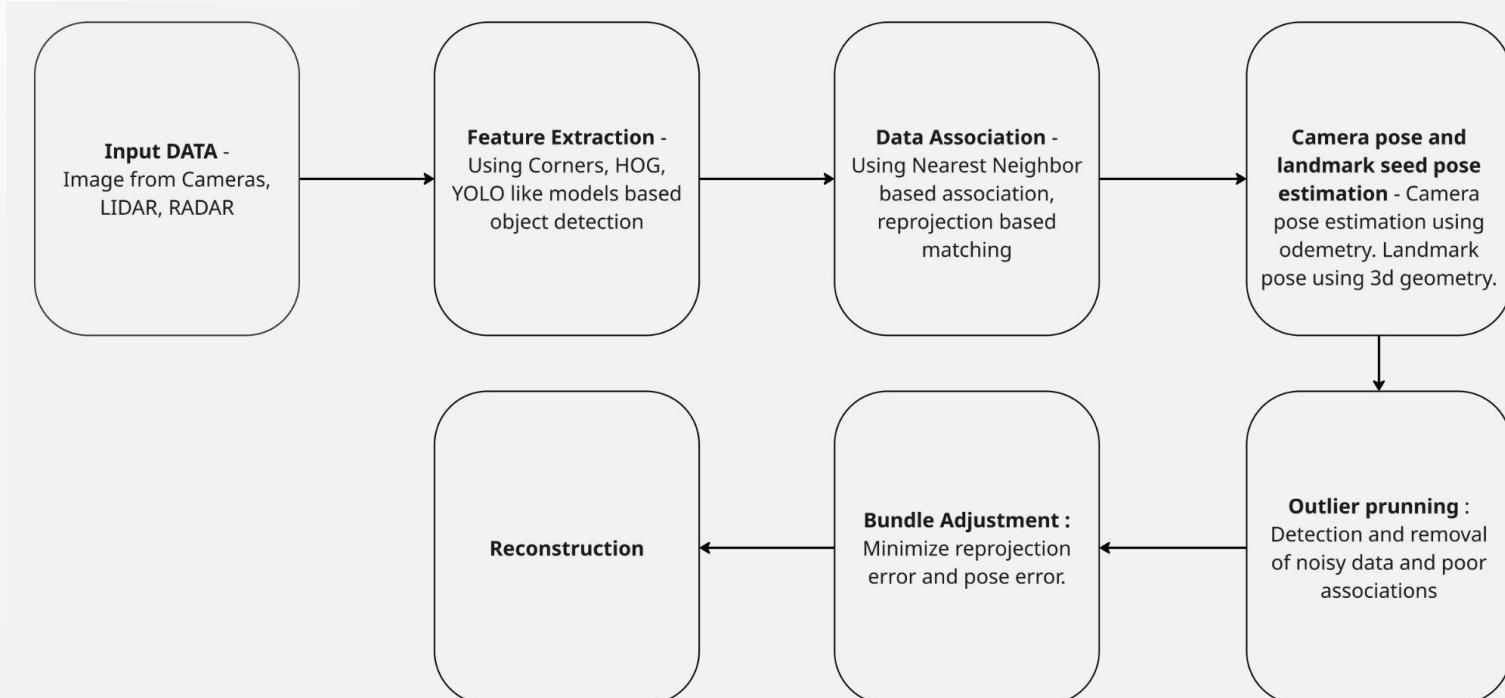
<http://www.theia-sfm.org/sfm.html>



Garrard-Hall using GLOMAP
<https://medium.com/@sugunsegu/understanding-glomap-global-structure-from-motion-64b052be7927>



What is Visual Reconstruction?



Bundle Adjustment in the Large (BAL)

- Benchmark for bundle adjustment
- Provides realistic 3D reconstruction problems from images
- Includes ground-truth data association for fair comparisons
- Designed for mid to large-scale optimization problems

Data Format

```
<num_cameras> <num_points> <num_observations>
<camera_index_1> <point_index_1> <x_1> <y_1>
...
<camera_index_num_observations> <point_index_num_observations> <x_num_observations> <y_num_observations>
<camera_1>
...
<camera_num_cameras>
<point_1>
...
<point_num_points>
```



Bundle Adjustment in the Large (BAL)



Data Format

- Observations are associated with landmarks.
 - Single observation per landmark per camera.
- Seed camera poses with intrinsics.
- Seed landmark poses.

```
<num_cameras> <num_points> <num_observations>
<camera_index_1> <point_index_1> <x_1> <y_1>
...
<camera_index_num_observations> <point_index_num_observations> <x_num_observations> <y_num_observations>
<camera_1>
...
<camera_num_cameras>
<point_1>
...
<point_num_points>
```



Bundle Adjustment in the Large (BAL)

Ladybug

Cameras : 1723

Points : 156502

Observations : 678718

Reprojection error

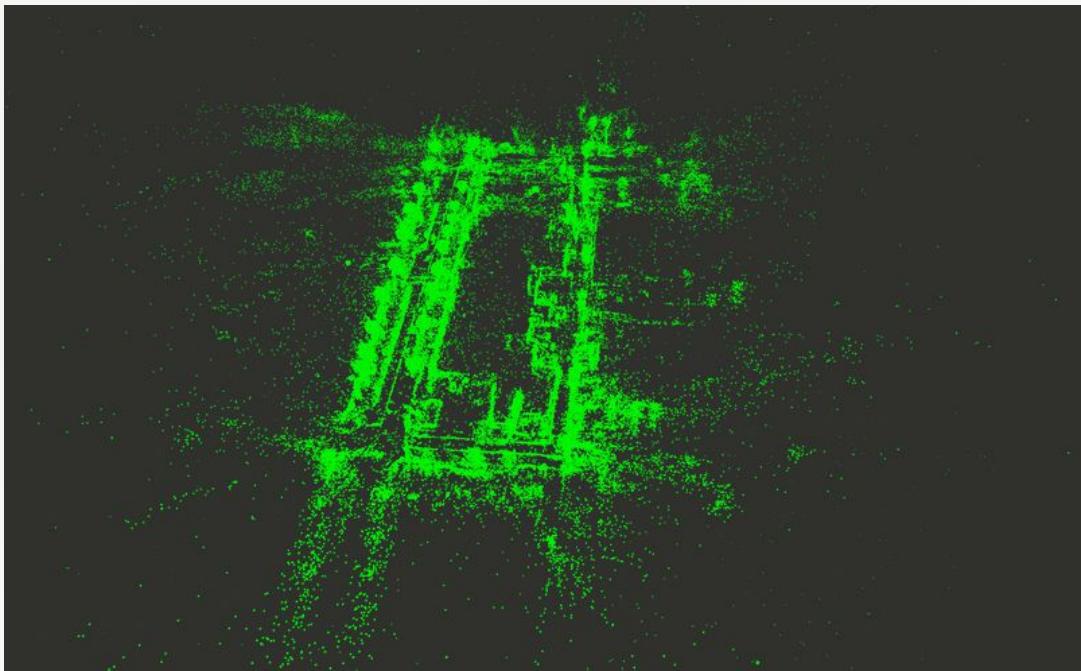
- $\mu=3.86, \sigma=8.61, \text{med}=1.67, \text{max} = 1693.285$

Observations per Camera

- $\mu=393.6, \sigma=195.87, \text{med}=395, \text{max}=970$

Observations per Point

- $\mu=4.34, \sigma=4.99, \text{med}=3, r=[2.000, 145.00]$



Bundle Adjustment in the Large (BAL)

Trafalgar Square dataset

Cameras : 257

Points : 65132

Observations : 225911

Reprojection error

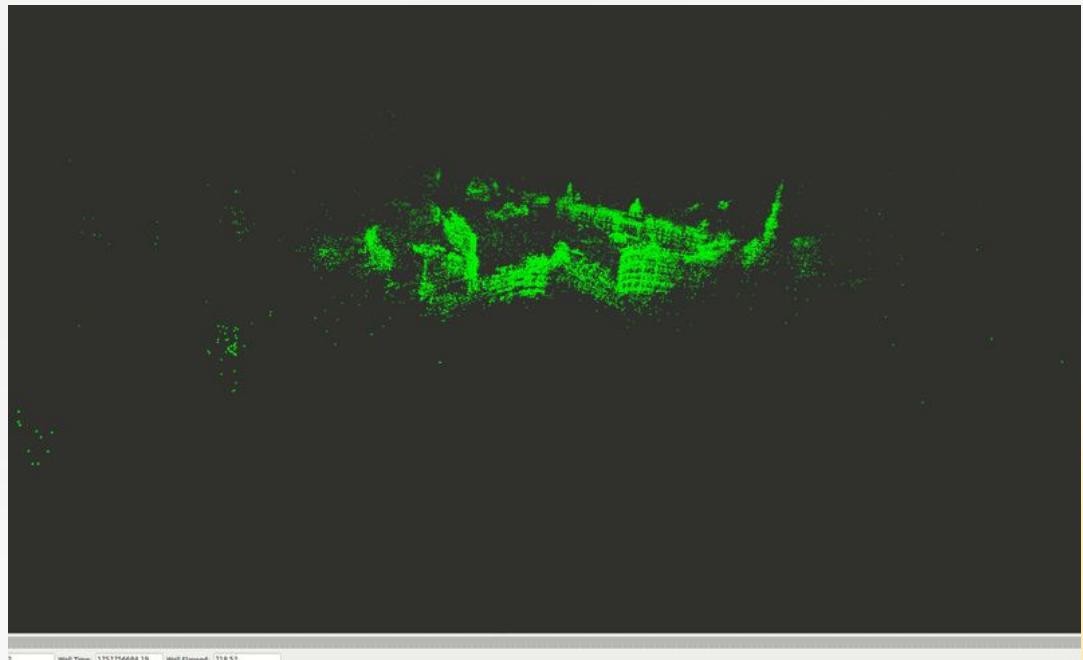
- $\mu=8.3$, $\sigma=12.2$, med=3.5, max=456.01

Observations per Camera

- $\mu=876.5$, $\sigma=792.29$, med=625, max= 4173

Observations per Point

- $\mu=3.47$, $\sigma=3.2$, med=2, range=[2, 42]



Bundle Adjustment in the Large (BAL)



Challenges

- Large Dataset, risks high processing time
 - Ladybug has 1723 images with 678718 observations
 - Trafalgar Square has 257 images with 225911 observations
- Noisy data and outliers can lead to non-convex optimization.

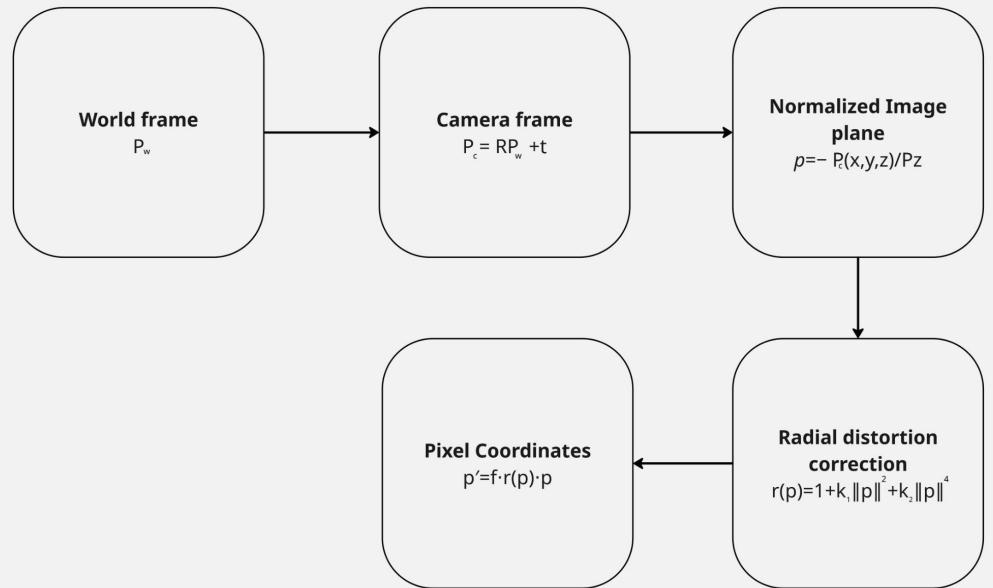


BAL Camera Model

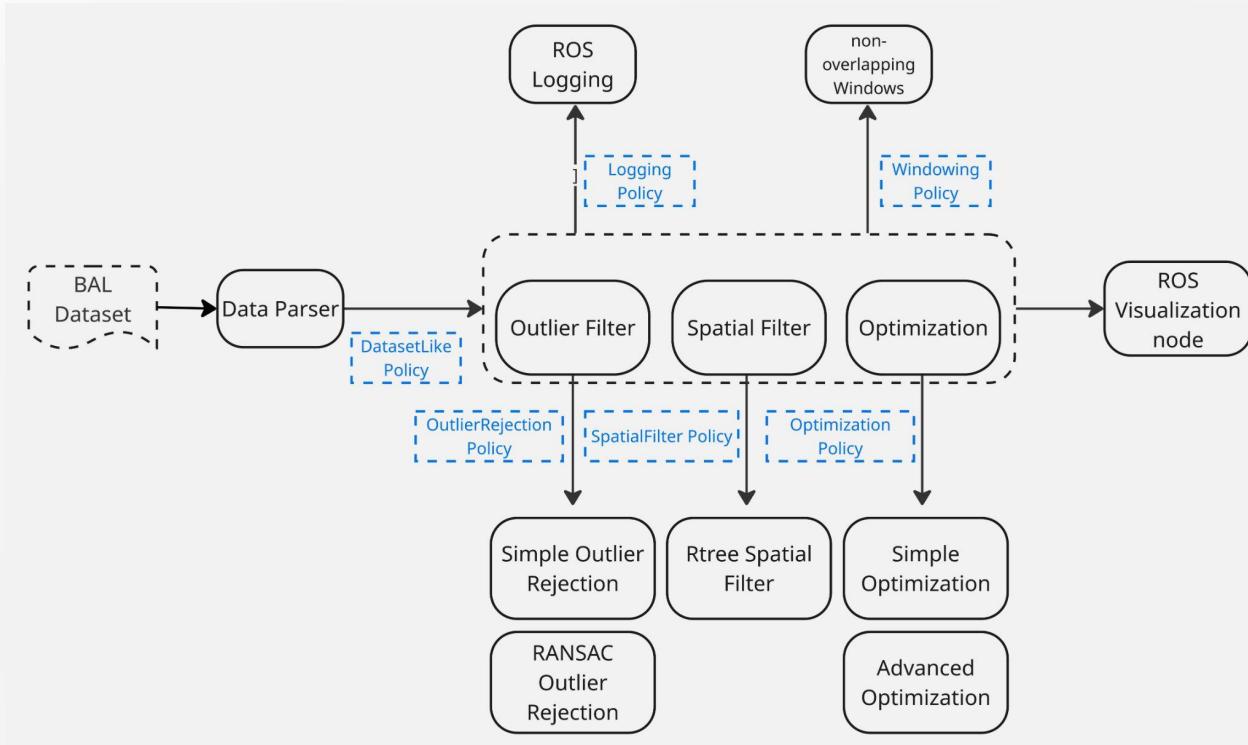
Parameters Estimated per Camera

- Rotation R, Translation t
- Focal length f
- Radial distortion k_1, k_2

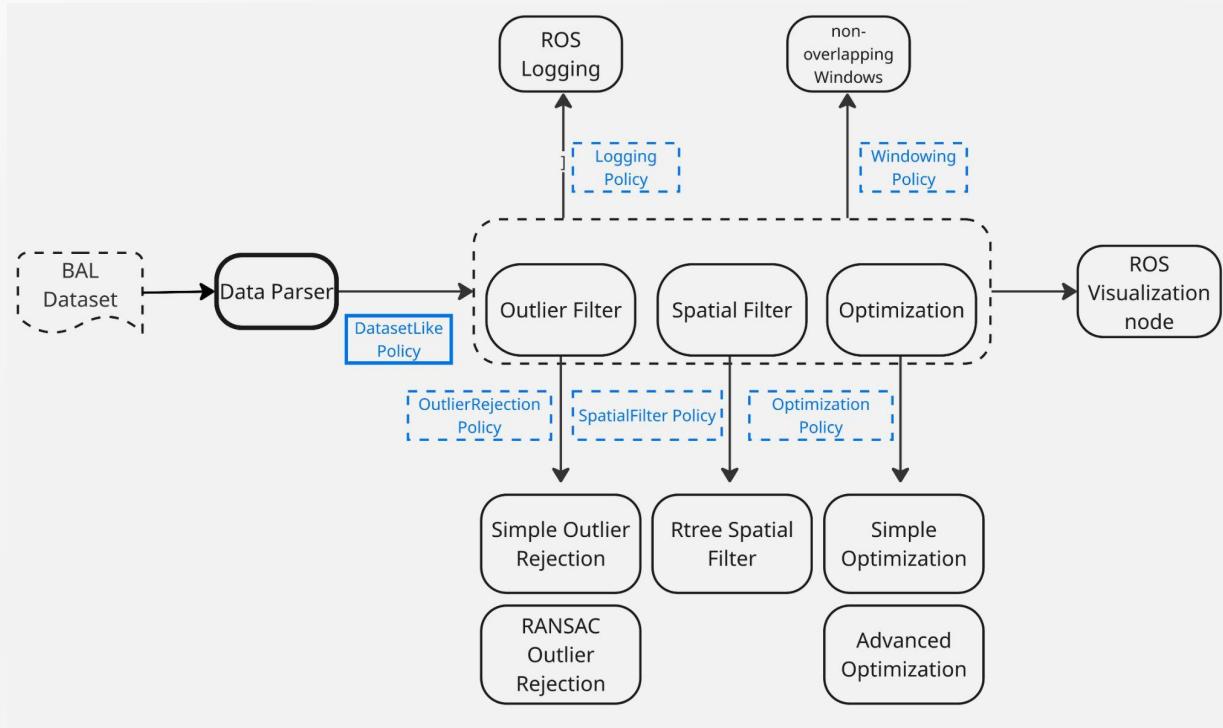
Projection Pipeline



SFM Pipeline



SFM Pipeline



Data Parser

```
[[nodiscard]] std::expected<BALDataType, ParseError> extract_data()
{
    try
    {
        auto cameras_landmarks = load_data();
        if (!cameras_landmarks)
        {
            return std::unexpected{cameras_landmarks.error()};
        }
        auto [cameras, landmarks] = std::move(*cameras_landmarks);

        if (config_.perform_filtering)
        {
            filter_observations(cameras, landmarks);
        }

        if (config_.normalize_data)
        {
            normalize(cameras, landmarks);
        }

        if (cameras.empty() || landmarks.empty())
        {
            return std::unexpected{ParseError::InsufficientData};
        }

        auto result = convert_to_policy_format(cameras, landmarks);
        return result;
    } catch (const std::exception&)
    {
        return std::unexpected{ParseError::CorruptedData};
    }
}
```

- `nodiscard` used to prevent accidental call to `extract_data()`
- Return type wrapped with `std::expected` to accommodate exception handling
- Performs filtering, normalization and conversion to policy based datatype.



Data Parser: preprocessing

```
void filter_observations(CameraVector& cameras, LandmarkVector& landmarks) noexcept
{
    if (config_.depth_threshold <= 0.0) return;
    auto observation_filter = [this, &cameras](auto& landmark)
    {
        std::unordered_map<int, typename Traits::Vec2> valid_obs;
        for (const auto& [cam_idx, obs] : landmark.obs)
        {
            if (cam_idx < 0 || cam_idx >= static_cast<int>(cameras.size()))
            {
                continue;
            }

            const auto& camera = cameras[cam_idx];
            auto p_cam = camera.T_c_w * landmark.p_w;

            if (is_valid_observation(p_cam, config_.depth_threshold))
            {
                valid_obs.emplace(cam_idx, obs);
            }
        }

        landmark.obs = std::move(valid_obs);
    };

    std::for_each(std::execution::par, landmarks.begin(), landmarks.end(), observation_filter);

    auto insufficient_observations = [this](const auto& landmark)
    {
        return landmark.observation_count() < config_.min_observations_per_landmark;
    };

    std::erase_if(landmarks, insufficient_observations);
}
```

- Parallel execution using `for::each` to speed up the processing (c++ 17).
 - Check if the depth of the landmark in camera frame is positive.
- `std::erase_if` used to remove any empty landmarks.

Data Parser : preprocessing

```
void normalize(CameraVector& cameras, LandmarkVector& landmarks) noexcept {
    if (landmarks.empty()) return;

    typename Traits::Vec3 median = Traits::Vec3::Zero();

    for (int dim = 0; dim < 3; ++dim) {
        auto coords = landmarks | std::views::transform([dim](const auto& landmark) { return landmark.p_w(dim); });
        median(dim) = compute_median(coords);
    }

    auto deviations = landmarks
        | std::views::transform([&median](const auto& landmark) {
            return (landmark.p_w - median).template lpNorm<1>();
        });

    Scalar median_dev = compute_median(deviations);
    Scalar scale_factor = static_cast<Scalar>(config_.scale) / median_dev;

    auto landmark_normalizer = [scale_factor, &median](auto& landmark) {
        landmark.p_w = scale_factor * (landmark.p_w - median);
    };

    auto camera_normalizer = [scale_factor, &median](auto& camera) {
        auto T_w_c = camera.T_c_w.inverse();
        T_w_c.translation() = scale_factor * (T_w_c.translation() - median);
        camera.T_c_w = T_w_c.inverse();
    };

    std::for_each(std::execution::par, landmarks.begin(), landmarks.end(), landmark_normalizer);
    std::for_each(std::execution::par, cameras.begin(), cameras.end(), camera_normalizer);
}
```

- Normalizes using Median absolute deviation(MAD).
 - MAD uses median which is more resilient to outliers.
 - Value becomes a factor of median deviation.
- Parallelized std::for_each using intel TBB.

Concepts : DatasetLike Policy

```

template<typename T>
concept CameraModelLike = requires(T model, const Eigen::Vector3d& point)
{
    {model.project(point)} -> std::convertible_to<Eigen::Vector2d>;
    {model.focal_length()} -> std::convertible_to<Eigen::Vector2d>;
    {model.principal_point()} -> std::convertible_to<Eigen::Vector2d>;
    {model.distortion_coefficients()} -> std::convertible_to<Eigen::Vector2d>;
};

template<typename T>
concept CameraLike = requires(T camera)
{
    typename T::Traits;
    typename T::CameraModelType;

    { camera.intrinsics } -> CameraModelLike;
    { camera.translation() } -> std::convertible_to<typename T::Traits::Vec3>;
    { camera.rotation() } -> std::convertible_to<typename T::Traits::SO3>;
    { camera.id } -> std::convertible_to<int>;
};

template<typename T>
concept ObservationLike = requires(T obs)
{
    typename T::Traits;
    { obs.camera_index } -> std::convertible_to<int>;
    { obs.point_index } -> std::convertible_to<int>;
    { obs.pixel_coordinates() } -> std::convertible_to<typename T::Traits::Vec2>;

    { obs.outlier } -> std::convertible_to<bool>;
    { obs.outlier_reason } -> std::convertible_to<std::string>;
};

```

CameraModelLike
function :
 project()

ObservationLike
variable :
 camera_index
 point_index
 outlier
 outlier_reason
function :
 pixel_coordinates()

CameraLike
variable :
 intrinsics
 id
function :
 translation()
 rotation()



Concepts : DatasetLike Policy

```
template<typename T>
concept DatasetLike = requires(T dataset)
{
    typename T::ScalarType;
    typename T::ObservationType;
    typename T::CameraType;
    typename T::Traits;

    // Functions for modifying data
    { dataset.observations() } -> std::same_as<std::vector<typename T::ObservationType>&>;
    { dataset.cameras() } -> std::same_as<std::vector<typename T::CameraType>&>;
    { dataset.points() } -> std::same_as<std::vector<typename T::Vec3>&>;

    // Functions for analysis
    { dataset.num_observations() } -> std::convertible_to<size_t>;
    { dataset.num_cameras() } -> std::convertible_to<size_t>;
    { dataset.num_points() } -> std::convertible_to<size_t>;
    { dataset.empty() } -> std::convertible_to<bool>;

    // maps for easy access
    { dataset.camera_to_observations_map() } -> std::same_as<std::unordered_map<int, std::vector<int>>&>;
    { dataset.point_to_observations_map() } -> std::same_as<std::unordered_map<int, std::vector<int>>&>;

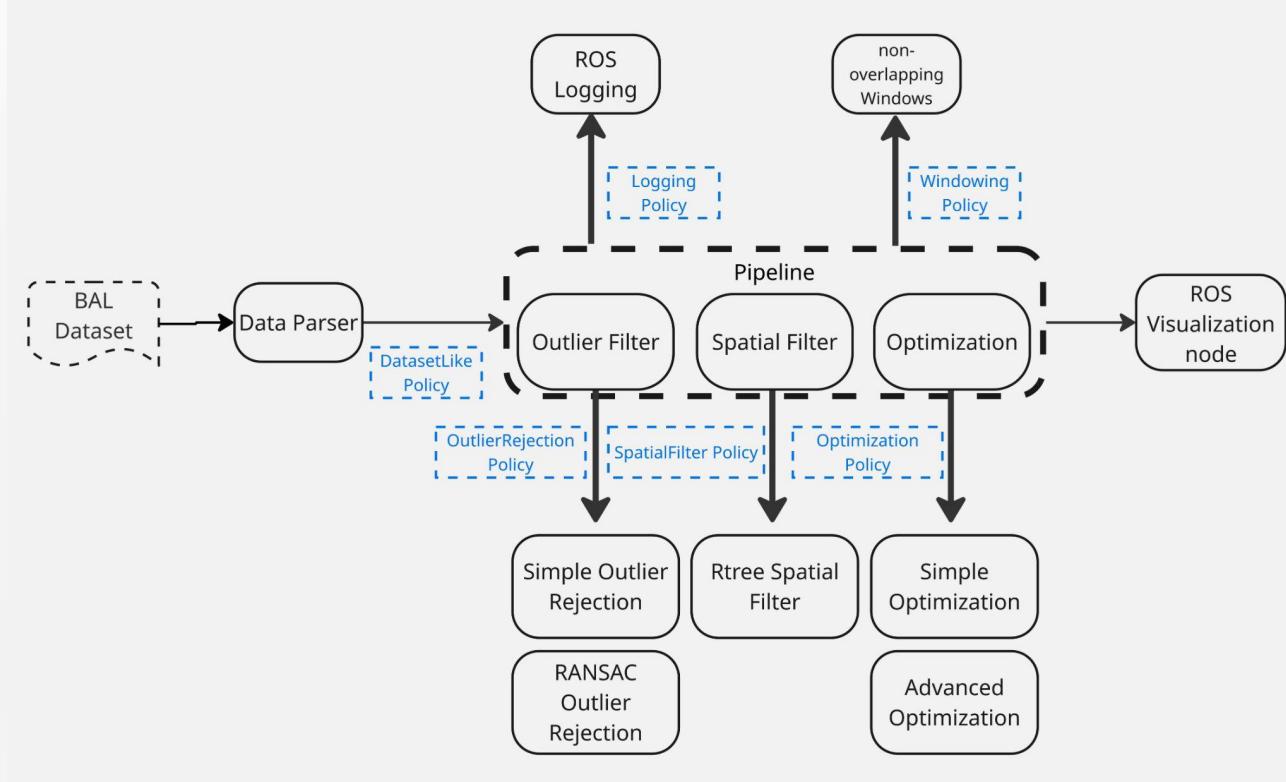
    // concepts for the data types
    requires ObservationLike<typename T::ObservationType>;
    requires CameraLike<typename T::CameraType>;
};

}
```

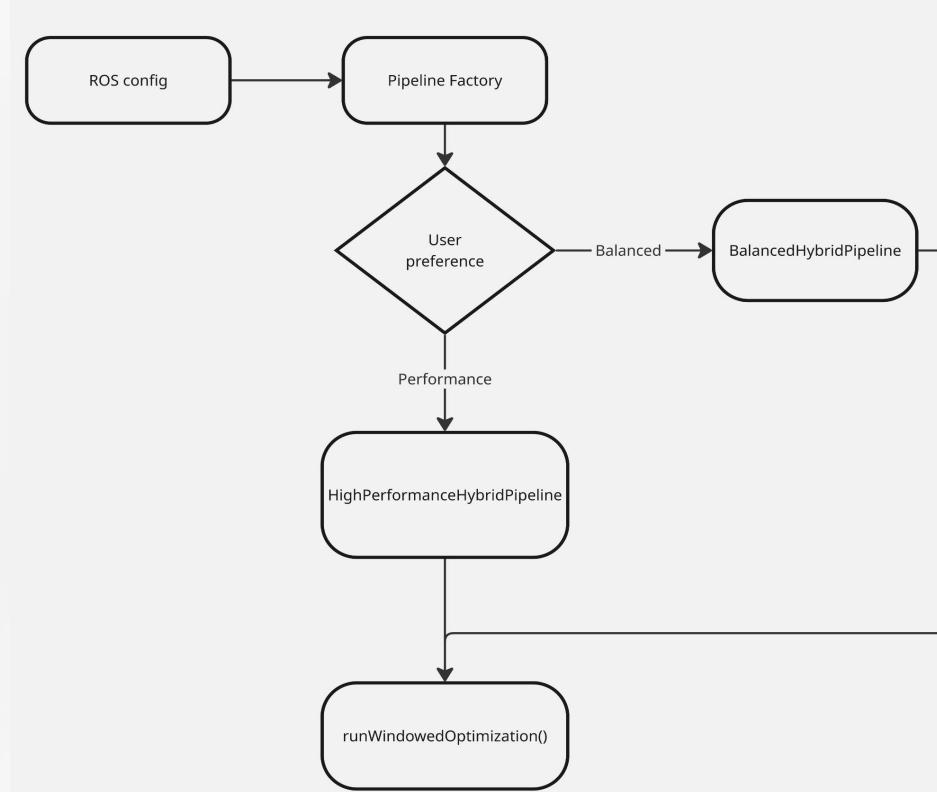
- `std::same_as` is used as a strict requirement.
- `std::convertible_to` is more flexible, just needs the return type to be convertible to the specified type.
- Use *requires* to specify the datatype.



Pipeline : Setup



Pipeline : Factory



Pipeline : Factory

```
40 /**
41 * @brief Hybrid Pipeline class for bundle adjustment
42 */
43 template<
44     typename Scalar,
45     typename DatasetType,
46     typename WindowingPolicyType,
47     typename SpatialFilterPolicyType,
48     typename OutlierRejectionPolicyType,
49     typename OptimizationPolicyType>
50 requires
51     DatasetLike<DatasetType> &&
52     policies::WindowingPolicy<WindowingPolicyType, DatasetType> &&
53     policies::SpatialFilterPolicy<SpatialFilterPolicyType, DatasetType> &&
54     policies::OutlierRejectionPolicy<OutlierRejectionPolicyType, DatasetType> &&
55     policies::OptimizationPolicy<OptimizationPolicyType, DatasetType>
56 class HybridPipeline {
57 public:
58     /**
59      * @brief Constructor to initialize the pipeline
60      */
61     explicit HybridPipeline(std::shared_ptr<rclcpp::Node> node, DatasetType data) noexcept
62     : node_(node),
63     optimized_data_(data),
64     data_(std::move(data)),
65     config_(loadHybridPipelineConfig(node))
66     {
67         data_stats_ = std::make_unique<DataStats<DatasetType>>(node_, config_.outlier_threshold, config_.verbose_logging);
68         policies::ROSLoggingPolicy logging_policy(node_);
69         windowing_policy_ = std::make_unique<WindowingPolicyType>(node_, logging_policy);
70
71         spatial_filter_policy_ = std::make_unique<SpatialFilterPolicyType>(node_, logging_policy);
72         outlier_rejection_policy_ = std::make_unique<OutlierRejectionPolicyType>(node_, logging_policy);
73         optimization_policy_ = std::make_unique<OptimizationPolicyType>(node_, logging_policy);
74     }
75
76     virtual ~HybridPipeline() = default;
```

- Concepts used to specify requirements on templated types.



Pipeline : Factory

```

256 /**
257  * @brief Balanced pipeline configuration
258 */
259 template<typename Scalar = double>
260 class BalancedHybridPipeline : public HybridPipeline<
261     Scalar,
262     BALData<Scalar>,
263     policies::implementations::NonOverlappingWindows<policies::ROSLoggingPolicy>,
264     policies::implementations::RTreeSpatialFilter<policies::ROSLoggingPolicy>,
265     policies::implementations::SimpleOutlierRejection<policies::ROSLoggingPolicy>,
266     policies::implementations::CeresOptimization<policies::ROSLoggingPolicy>>
267 {
268 public:
269     explicit BalancedHybridPipeline(std::shared_ptr<rclcpp::Node> node, BALData<Scalar> data) noexcept
270         : HybridPipeline<Scalar>(
271             BALData<Scalar>,
272             policies::implementations::NonOverlappingWindows<policies::ROSLoggingPolicy>,
273             policies::implementations::RTreeSpatialFilter<policies::ROSLoggingPolicy>,
274             policies::implementations::SimpleOutlierRejection<policies::ROSLoggingPolicy>,
275             policies::implementations::CeresOptimization<policies::ROSLoggingPolicy>>(node, std::move(data))
276     {}
277
278     std::string getPipelineTypeName() const {
279         return "Balanced HybridPipeline";
280     }
281 };
282
283 /**
284  * @brief High-performance pipeline configuration
285 */
286 template<typename Scalar = double>
287 class HighPerformanceHybridPipeline : public HybridPipeline<
288     Scalar,
289     BALData<Scalar>,
290     policies::implementations::NonOverlappingWindows<policies::ROSLoggingPolicy>,
291     policies::implementations::RTreeSpatialFilter<policies::ROSLoggingPolicy>,
292     policies::implementations::SimpleOutlierRejection<policies::ROSLoggingPolicy>,
293     policies::implementations::AdvancedOptimization<policies::ROSLoggingPolicy>>
294 {
295 public:
296     explicit HighPerformanceHybridPipeline(std::shared_ptr<rclcpp::Node> node, BALData<Scalar> data) noexcept
297         : HybridPipeline<Scalar>(
298             BALData<Scalar>,
299             policies::implementations::NonOverlappingWindows<policies::ROSLoggingPolicy>,
300             policies::implementations::RTreeSpatialFilter<policies::ROSLoggingPolicy>,
301             policies::implementations::SimpleOutlierRejection<policies::ROSLoggingPolicy>,
302             policies::implementations::AdvancedOptimization<policies::ROSLoggingPolicy>>(node, std::move(data))
303     {}
304
305     std::string getPipelineTypeName() const {
306         return "High-Performance HybridPipeline";
307     }
308 };

```

- Derived classes with pre-set template arguments.
- User can select at runtime using ROS params.

```

58     if (pipeline_type == "balanced")
59     {
60         auto pipeline = bal_cppcon::HybridPipelineFactory<double>::createBalancedPipeline(node, data);
61         runPipeline(node, pipeline, data);
62     } else if (pipeline_type == "high_performance")
63     {
64         auto pipeline = bal_cppcon::HybridPipelineFactory<double>::createHighPerformancePipeline(node, data);
65     }
66 }

```

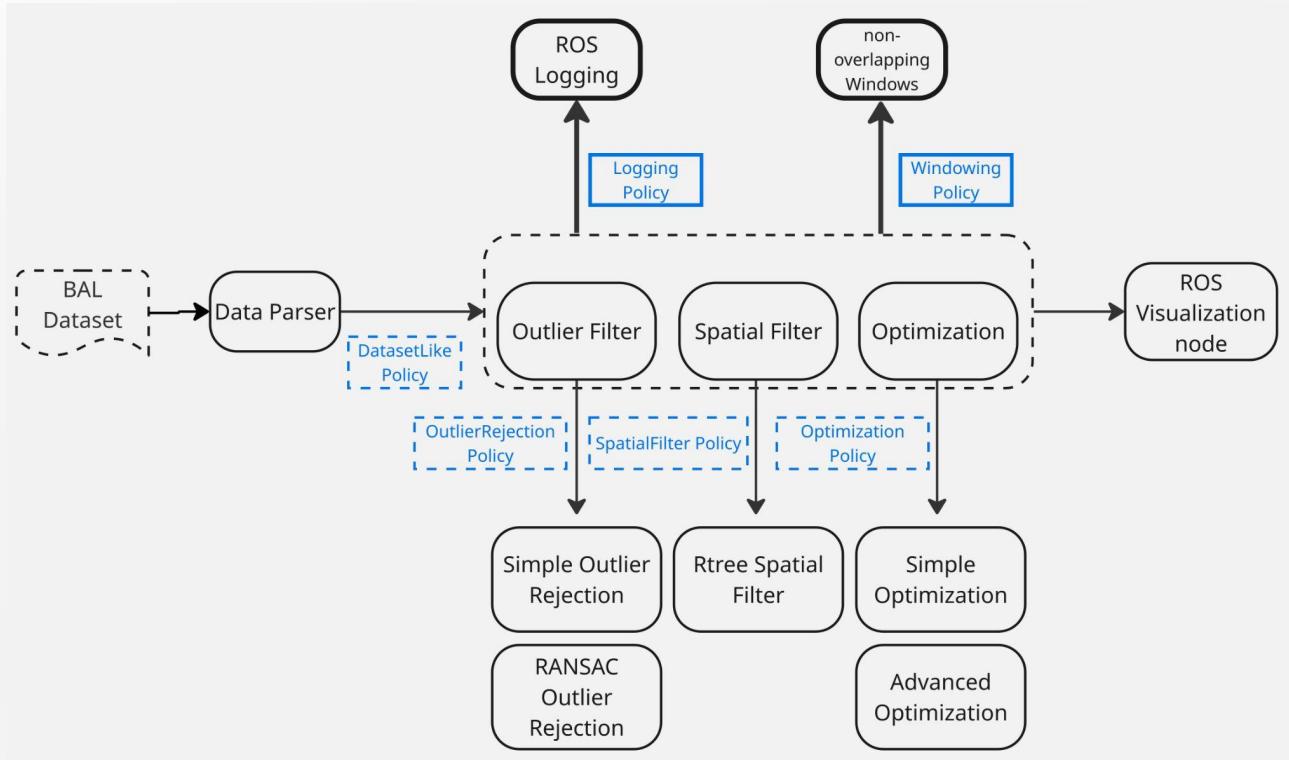
Pipeline : main



```
81 void runWindowedOptimization()
82 {
83     // prepare windows based on camera poses
84     auto windowing_result = windowing_policy_->generateWindows(optimized_data_);
85     const auto& windows = windowing_result.windows;
86
87     for (const auto& window : windows)
88     {
89         // filter observations in the window using outlier and spatial filter policies
90         std::span<const int> camera_window_span(window.camera_indices);
91         auto outlier_result = outlier_rejection_policy_->rejectOutliers(optimized_data_, camera_window_span);
92         auto spatial_result = spatial_filter_policy_->filterSpatially(optimized_data_, camera_window_span);
93     }
94
95     for (const auto& window : windows)
96     {
97         // optimize the window using the optimization policy
98         std::span<const int> camera_window_span(window.camera_indices);
99         auto optimization_result = optimization_policy_->optimize(optimized_data_, camera_window_span);
```



Logging and window policy



Logging and window policy

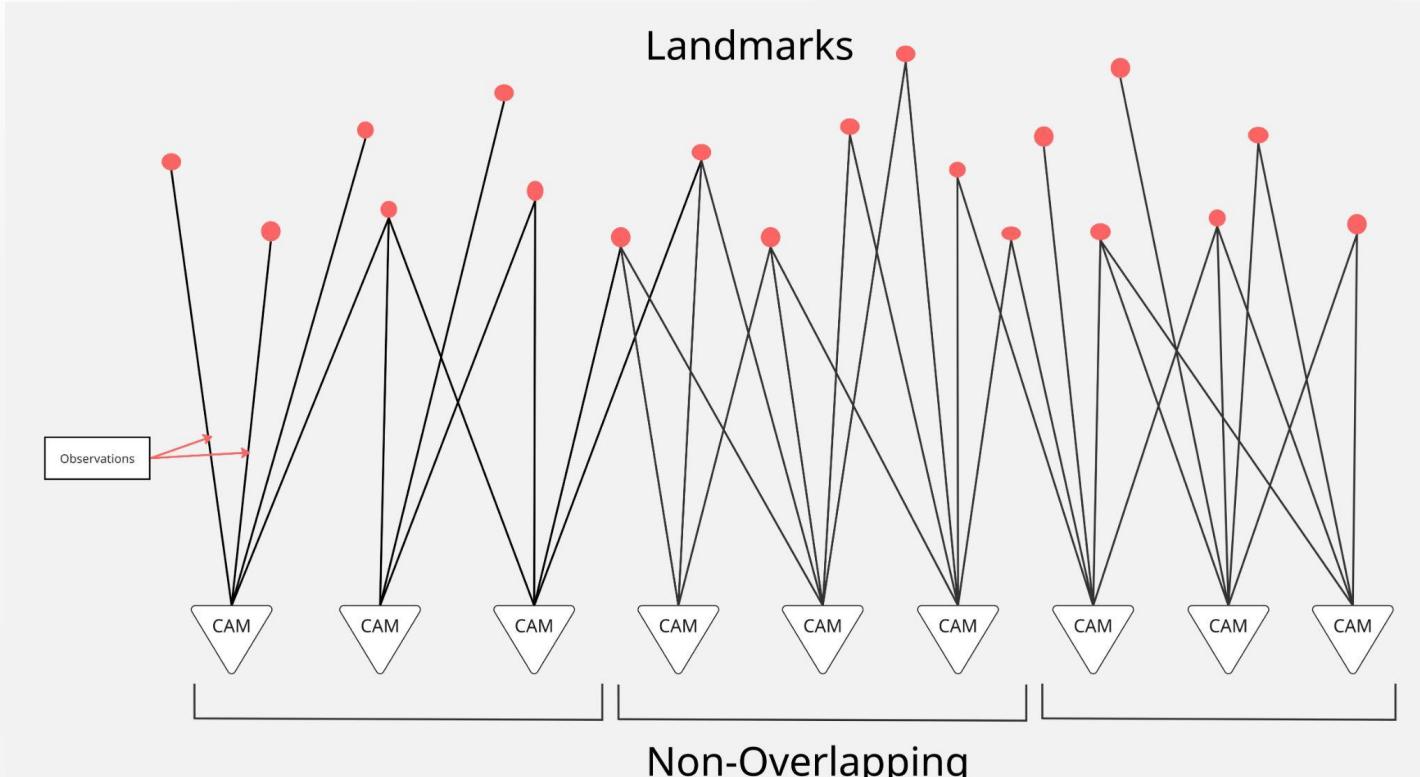
```
enum class LogLevel : std::uint8_t
{
    DEBUG = 0,
    INFO = 1,
    WARN = 2,
    ERROR = 3
};

template<typename Policy>
concept LoggingPolicy = requires(Policy policy, LogLevel level, std::string_view message)
{
    { policy.log(level, message) } -> std::same_as<void>;
};
```

```
36 // ROS Logging Policy
37 class ROSLoggingPolicy
38 {
39 public:
40     explicit ROSLoggingPolicy(std::shared_ptr<rclcpp::Node> node) noexcept
41         : node_(std::move(node)) {}
42
43     // Satisfy LoggingPolicy concept
44     void log(LogLevel level, std::string_view message) noexcept
45     {
46         if (!node_) return;
47
48         switch (level) {
49             case LogLevel::DEBUG:
50                 RCLCPP_DEBUG(node_->get_logger(), "%s", message.data());
51                 break;
52             case LogLevel::INFO:
53                 RCLCPP_INFO(node_->get_logger(), "%s", message.data());
54                 break;
55             case LogLevel::WARN:
56                 RCLCPP_WARN(node_->get_logger(), "%s", message.data());
57                 break;
58             case LogLevel::ERROR:
59                 RCLCPP_ERROR(node_->get_logger(), "%s", message.data());
60                 break;
61         }
62     }
63 private:
64     std::shared_ptr<rclcpp::Node> node_;
65 };
66
67 // Verify that ROSLoggingPolicy satisfies the concept
68 static_assert(LoggingPolicy<ROSLoggingPolicy>);
```



Logging and window policy



Logging and window policy

```
11 struct CameraWindow
12 {
13     std::vector<int> camera_indices;
14     int start_camera_id = -1;
15     size_t window_size = 0;
16
17     // Keep essential methods
18     [[nodiscard]] bool empty() const noexcept { return camera_indices.empty(); }
19     [[nodiscard]] size_t size() const noexcept { return camera_indices.size(); }
20
21     // Iterator support for range-based loops
22     [[nodiscard]] auto begin() const noexcept { return camera_indices.begin(); }
23     [[nodiscard]] auto end() const noexcept { return camera_indices.end(); }
24 };
25
26
27 struct WindowingResult
28 {
29     std::vector<CameraWindow> windows;
30     size_t total_cameras_covered = 0;
31     size_t overlap_count = 0;
32
33     [[nodiscard]] bool empty() const noexcept { return windows.empty(); }
34     [[nodiscard]] size_t num_windows() const noexcept { return windows.size(); }
35 };
36
37 template<typename Policy, typename Dataset>
38 concept WindowingPolicy = requires(
39     Policy policy,
40     const Dataset& dataset)
41 {
42     requires DatasetLike<Dataset>;
43     { policy.generate_windows(dataset) }-> std::same_as<WindowingResult>;
44     requires std::destructible<Policy>;
45 };
46
```

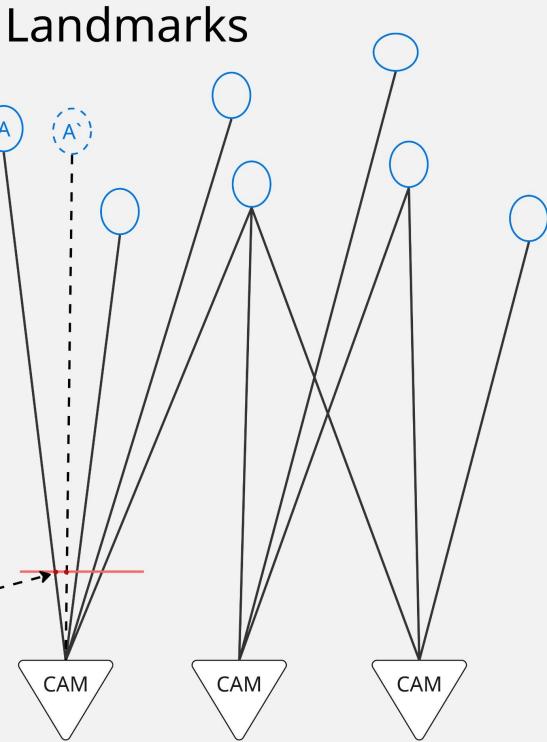


Pipeline - outlier filtering model

```
14
15     struct OutlierRejectionResult
16 {
17     size_t total_observations_processed{0};
18     size_t total_inliers_found{0};
19
20     // Default constructor
21     OutlierRejectionResult() = default;
22
23     // Statistics
24     [[nodiscard]] double inlier_ratio() const noexcept {
25         return total_observations_processed > 0 ?
26             static_cast<double>(total_inliers_found) / total_observations_processed : 0.0;
27     }
28 };
29
30 // outlier rejection policy
31 template<typename Policy, typename Dataset>
32 concept OutlierRejectionPolicy = requires(
33     Policy policy,
34     Dataset& dataset,
35     std::span<const int> camera_window)
36 {
37     requires DatasetLike<Dataset>;
38     // Main outlier rejection method
39     { policy.rejectOutliers(dataset, camera_window) } -> std::same_as<OutlierRejectionResult>;
40
41     // Stats
42     { policy.get_algorithm_name() } -> std::convertible_to<std::string_view>;
43     { policy.get_statistics() } -> std::convertible_to<std::string>;
44
45     requires std::destructible<Policy>;
46 };
47
```



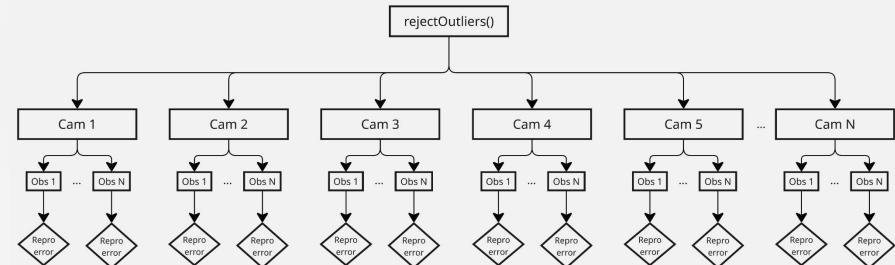
Pipeline - Simple outlier filtering



Pipeline - Simple outlier filtering

```

55 // OutlierRejectionPolicy interface
56 OutlierRejectionResult rejectOutliers(DatasetLike auto& dataset,
57                                     std::span<const int> camera_window)
58 {
59     bal_cppcon::policies::concepts::OutlierRejectionResultInternal result_internal;
60
61     // Process each camera in the window
62     processCamerasParallel(dataset, camera_window, result_internal);
63     OutlierRejectionResult result = result_internal.to_result();
64     update_statistics(result);
65     return result;
66 }
67
68 void processCamerasParallel(
69     DatasetLike auto& dataset,
70     std::span<const int> camera_window,
71     bal_cppcon::policies::concepts::OutlierRejectionResultInternal& result)
72 {
73     auto window_cameras = camera_window |
74         std::views::transform([&dataset](int idx) {
75             return dataset.cameras()[idx];});
76
77     // Process each camera in parallel
78     std::for_each(std::execution::par, window_cameras.begin(),
79                  window_cameras.end(),
80                  [this, &dataset, &result](const auto& camera) {
81                     processSingleCamera(dataset, camera, result);});
82 }
```

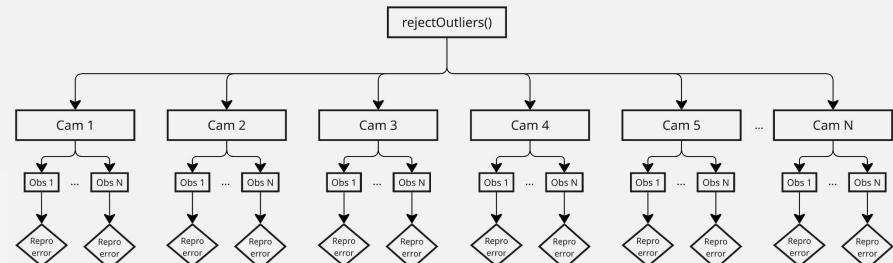


Pipeline - Simple outlier filtering

```

141 // Get mutable access to observations to update outlier flags
142 auto& observations = dataset.observations();
143
144 // Compute reprojection errors for all observations
145 std::vector<std::pair<int, double>> error_pairs(obs_indices_set.size());
146
147 std::transform(std::execution::par, obs_indices_set.begin(),
148               obs_indices_set.end(), error_pairs.begin(),
149               [this, &dataset, &camera](int obs_idx) {
150                 return compute_reprojection_error(dataset, camera, obs_idx);});
151
152 // Update outlier flags in observations and count inliers
153 std::atomic<size_t> inlier_count{0};
154
155 std::for_each(std::execution::par, error_pairs.begin(), error_pairs.end(),
156               [this, &observations, &inlier_count](const auto& pair) {
157                 const auto& [obs_idx, error] = pair;
158                 if (obs_idx >= 0 &&
159                     obs_idx < static_cast<int>(observations.size())) {
160                   auto& obs = observations[obs_idx];
161
162                   // Determine if observation is an outlier based on reprojection error
163                   bool is_outlier = error >= config_.reprojection_threshold;
164
165                   // Update the observation's outlier status
166                   obs.outlier = is_outlier;
167                   if (is_outlier) {
168                     obs.outlier_reason = "simple_outlier_rejection";
169                   } else {
170                     obs.outlier_reason = "";
171                     inlier_count.fetch_add(1);
172                   }
173               });
174 };
175

```

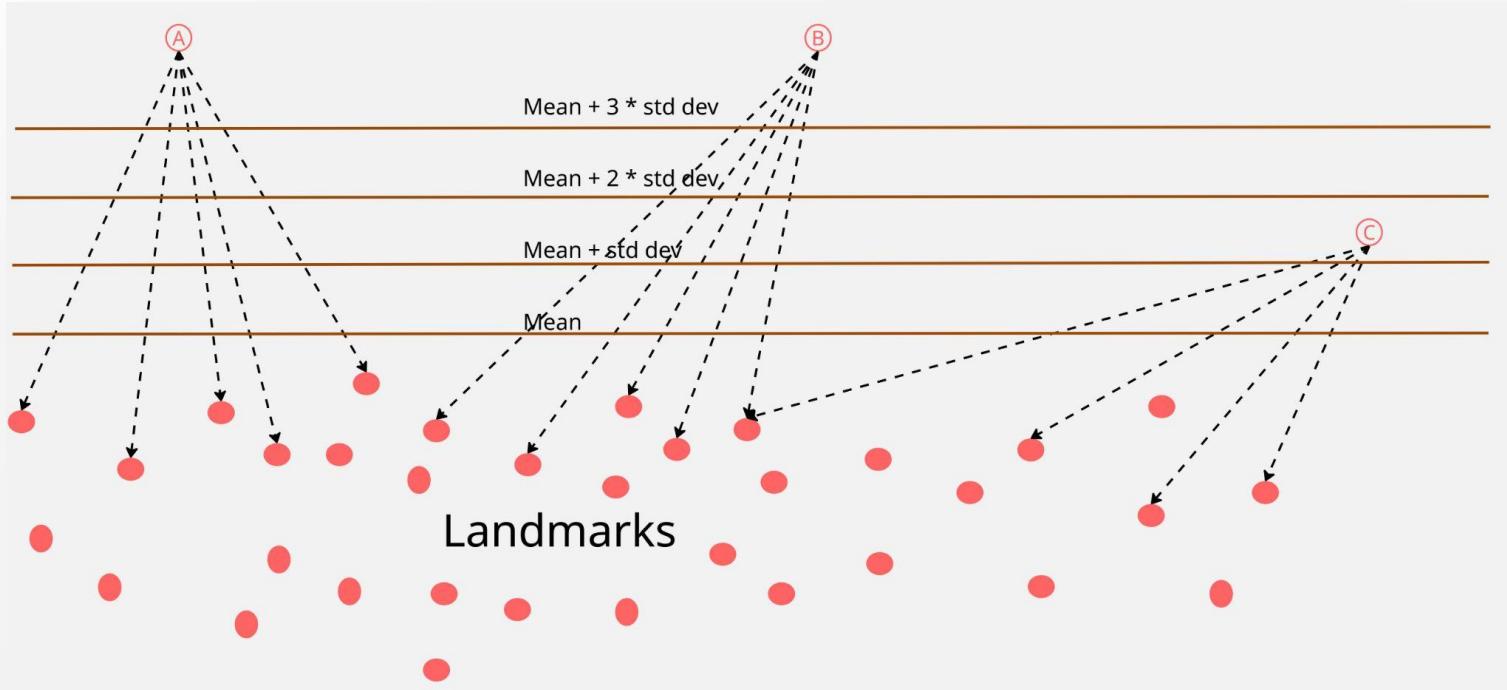


Pipeline - Spatial Filtering

```
14 // spatial filter result
15 struct SpatialFilterResult {
16     size_t total_observations_processed{0};
17     size_t total_observations_kept{0};
18
19     SpatialFilterResult() = default;
20
21     // Statistics
22     [[nodiscard]] double retention_ratio() const noexcept {
23         return total_observations_processed > 0 ?
24             static_cast<double>(total_observations_kept) / total_observations_processed : 0.0;
25     }
26 };
27
28 //spatial filter policy concept
29 template<typename Policy, typename Dataset>
30 concept SpatialFilterPolicy = requires(
31     Policy policy,
32     Dataset& dataset,
33     std::span<const int> camera_window)
34 {
35     requires DatasetLike<Dataset>;
36
37     // Main spatial filtering method
38     { policy.filterSpatially(dataset, camera_window) }
39         -> std::same_as<SpatialFilterResult>;
40
41     // stats
42     { policy.getLastStatistics() } -> std::convertible_to<std::string>;
43
44     requires std::destructible<Policy>;
45
46 };
```



Pipeline - Spatial Filtering



Pipeline - Spatial Filtering

```
116 template <typename Dataset>
117 requires DatasetLike<Dataset>
118 void buildSpatialIndex(Dataset& dataset)
119 {
120     std::lock_guard<std::mutex> lock(mutex_);
121
122     if (config_.verbose_logging)
123     {
124         logging_policy_.log(LogLevel::INFO, "Building R-tree spatial index");
125     }
126
127     rtree_.clear();
128
129     // build rtree with points as spatial index
130     size_t count = 0;
131     const auto& points = dataset.points();
132     for (const auto& point : points)
133     {
134         if (point.size() >= 3)
135         {
136             rtree_.insert(std::make_pair(Point3D(point[0], point[1], point[2]), count++));
137         }
138     }
139
140     rtree_built_ = true;
141
142     if (config_.verbose_logging)
143     {
144         logging_policy_.log(LogLevel::INFO,
145             "R-tree built with " + std::to_string(count) + " points");
146     }
147 }
```

- Rtree used to store 3d points for quick nearest neighbour query.
- Groups nearby points into bounding boxes, forming a tree structure.
- Complexity for querying is $O(\log n)$.



Pipeline - Spatial Filtering

```
223     for (int obs_idx : obs_indices_set)
224     {
225         if (obs_idx >= 0 && obs_idx < static_cast<int>(observations.size()))
226         {
227             const auto& obs = observations[obs_idx];
228             if (!obs.outlier) { // Only process inliers
229                 inlier_obs_indices.push_back(obs_idx);
230             }
231         }
232     }
233
234     // Process each inlier observation for this camera in parallel and update outlier flags
235     std::atomic<size_t> kept_count{0};
236     std::for_each(std::execution::par, inlier_obs_indices.begin(), inlier_obs_indices.end(),
237                  [this, &dataset, &kept_count](int obs_idx)
238     {
239         bool is_valid = processSingleObservation(dataset, obs_idx);
240         if (is_valid)
241         {
242             kept_count.fetch_add(1);
243         }
244     });
245 }
```

- Only process inliers. Reject points that were detected as outliers by an earlier filter.
- `std::for_each` with parallel execution is used to speed up the filtering



Pipeline - Spatial Filtering

```

307 // Query R-tree for k-nearest neighbors
308 std::vector<Value> neighbors;
309 rtree_.query(boost::geometry::index::nearest(query_pt, config_.k_neighbors + 1),
310               std::back_inserter(neighbors));
311
312 // Filter out the query point itself and calculate distances
313 auto distances = neighbors
314   | std::views::filter([point_idx = obs.point_index](const auto& neighbor) {
315     return static_cast<int>(neighbor.second) != point_idx;
316   })
317   | std::views::transform([&query_pt](const auto& neighbor) {
318     return boost::geometry::distance(query_pt, neighbor.first);
319   })
320   | std::ranges::to<std::vector>();
321
322 // Calculate mean and standard deviation
323 auto [distance_sum, squared_distance_sum, num_distances] = std::ranges::fold_left(
324   distances,
325   std::tuple{0.0, 0.0, 0u},
326   [] (auto previous_result, double distance_value)
327   {
328     auto& [total_distance_sum, total_squared_distance_sum, distance_count] = previous_result;
329     total_distance_sum           += distance_value;
330     total_squared_distance_sum  += distance_value * distance_value;
331     ++distance_count;
332     return previous_result;
333   }
334 );
335
336 double mean = distance_sum / num_distances;
337 double variance = (squared_distance_sum - num_distances * mean * mean) / num_distances;
338 double stddev = std::sqrt(variance);
339
340 // Check if the nearest neighbor distance is reasonable (not an outlier)
341 double nearest_distance = std::ranges::min(distances);
342 bool is_valid = stddev > 0.0 &&
343             std::abs(nearest_distance - mean) <= config_.std_dev_multiplier * stddev;
344

```

- Query k neighbors nearest to the query point and find its distance from them.
- Fold::left is used to calculate sum of distances, squared sum and count.
- Calculate mean and variance using those values.
- Outlier check is a variant of z-score. Checks if the nearest distance to a neighbour is within a multiple of the standard deviation.



Pipeline - Filtering

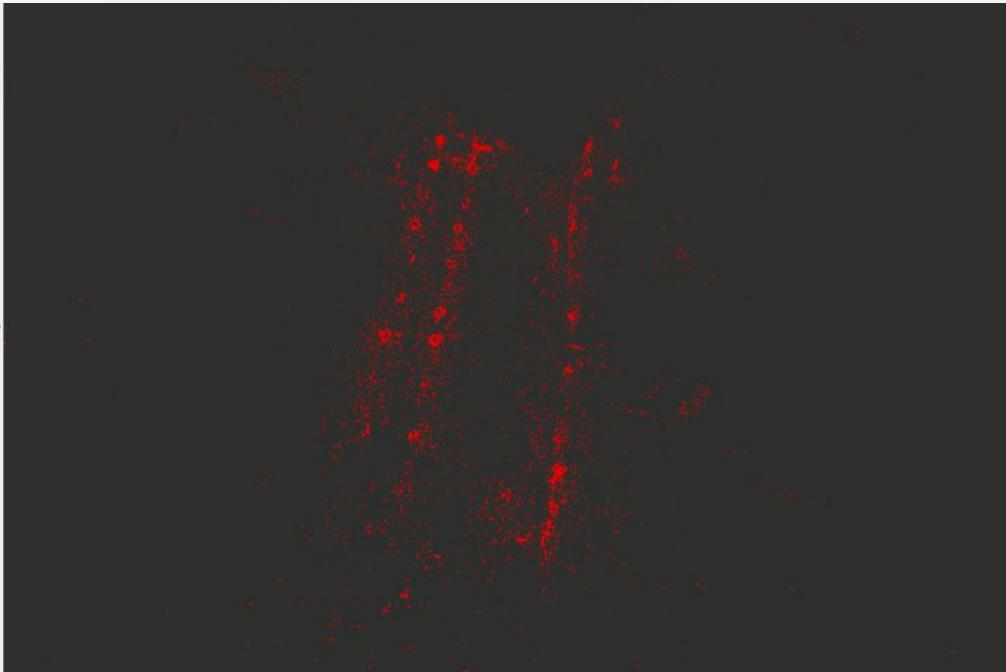
Ladybug dataset

Thresholds

- Reprojection error threshold : 25.0pix
- Neighbours for spatial filter : 100
- std_dev multiplier : 4.5

Outliers

- Count=22866/678155 (3.371796%),
- ReprojError:
 - Mean=21.980313px,
 - StdDev=36.999753px,
 - Median=19.146612px



Pipeline - Filtering

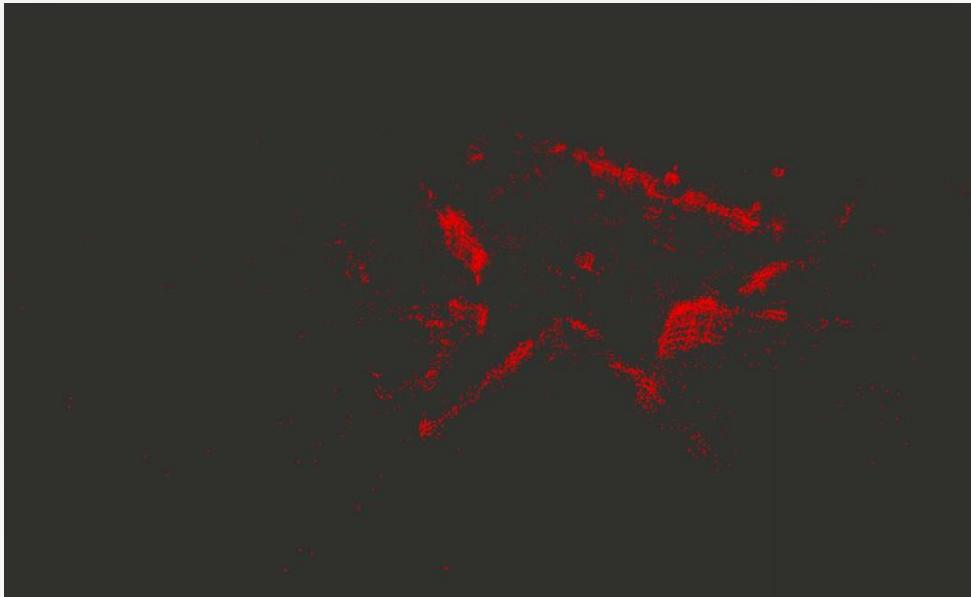
Trafalgar Square dataset

Thresholds

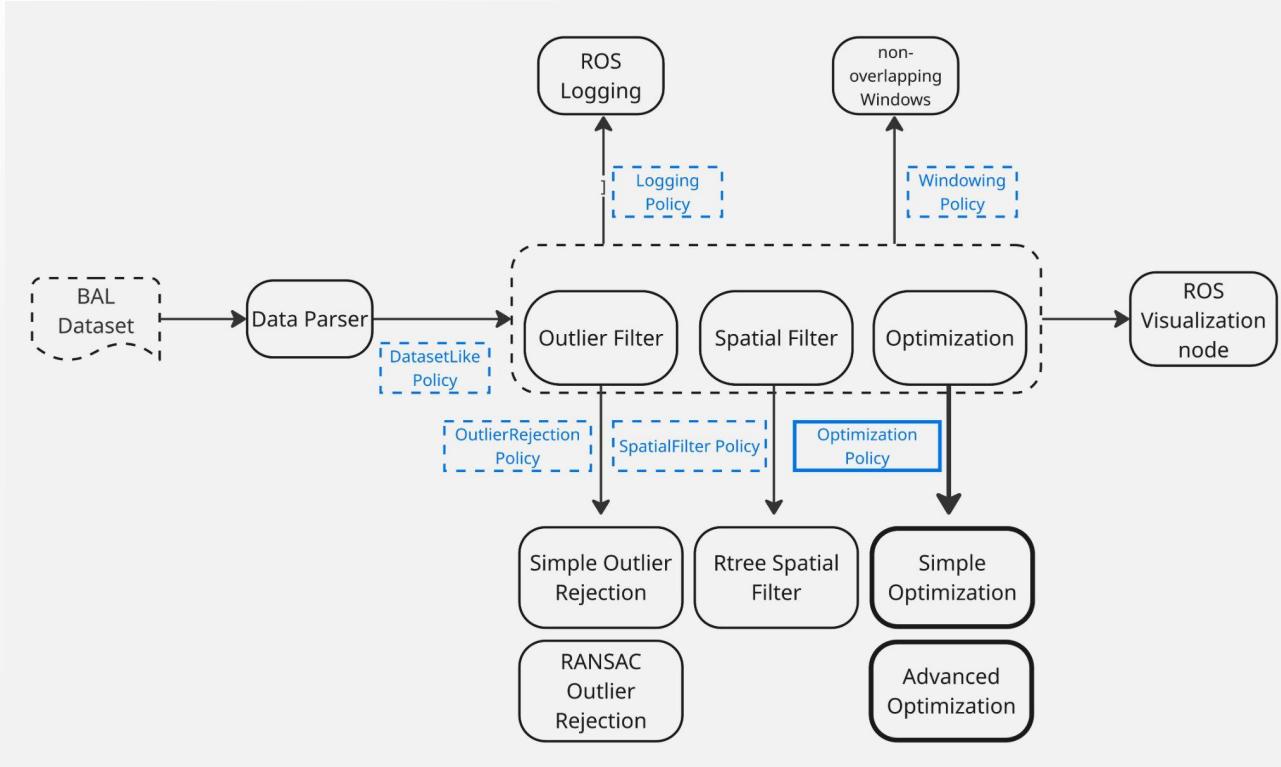
- Reprojection error threshold : 25.0pix
- Neighbours for spatial filter : 100
- std_dev multiplier : 4.5

Outliers

- Count=20114/225267 (8.928960%),
- ReprojError:
 - Mean=37.676941px,
 - StdDev=19.108898px,
 - Median=34.209722px



Pipeline - Optimization



Optimization policy

```
struct IterationSummary
{
    double cost{0.0};
    double cost_change{0.0};
    double cost_reduction_percentage{0.0};
    bool step_is_successful{false};
};

bool success{false};
double initial_cost{0.0};
double final_cost{0.0};
double cost_reduction_percentage{0.0};

std::vector<IterationSummary> iteration_summaries;
std::string solver_summary;
};

template<typename Policy, typename Dataset>
concept OptimizationPolicy = requires(
    Policy policy,
    Dataset& dataset,
    std::span<const int> camera_window)
{
    requires DatasetLike<Dataset>;
    { policy.optimize(dataset, camera_window) } -> std::same_as<OptimizationResult>;
    { policy.getLastStatistics() } -> std::convertible_to<std::string>;
    { policy.reset() } -> std::same_as<void>;
    requires std::destructible<Policy>;
};
```



Pipeline - Optimization

Problem: Refine camera poses and 3D points from noisy observations.

Solution: Minimize some cost functions that describe the reprojection error of the features

Given (Input Data)

- Observations: 2D points (u_{ij}, v_{ij}) where camera i sees 3D point j .
- Initial estimates: Camera parameters - (R, T, f, k_1, k_2) and 3d point (x, y, z) .

Optimize

- 9 Camera parameters - Extrinsic(R, t) and intrinsic(f, k_1, k_2).
- 3 Landmark parameters - Pose (x, y, z)

Optimization Problem

$$\min_{\{c_i\}, \{X_j\}} \sum_{(i,j) \in \mathcal{O}} \rho(\|p_{ij} - \pi(c_i, X_j)\|^2)$$

- For each observation, solve
- This equation can be expressed as $\min \|r(x)\|^2$



Pipeline - Optimization



- **Linearization**

- At the current estimate, compute Jacobians of residuals wrt parameters.
 - Using a first-order Taylor expansion, we can rephrase the problem as $f(x + \delta) \sim f(x) + J\delta$
 - So the reprojection error becomes this Normal Equation :
 - $\min \|r(x)\|^2 \rightarrow \min \|r(x + \delta)\|^2 \rightarrow \min \|f(x) + J\delta\|^2 \rightarrow (J^T J)\delta = -J^T r$

- **Solve Linear Subproblem**

- Use **Schur complement** for a BA problem. Exploit sparsity in the jacobian.
 - Schur complement eliminates point variables, turning the large system into a smaller one in camera parameters.
 - Point variables are solved for by back-substituting the solved camera poses.

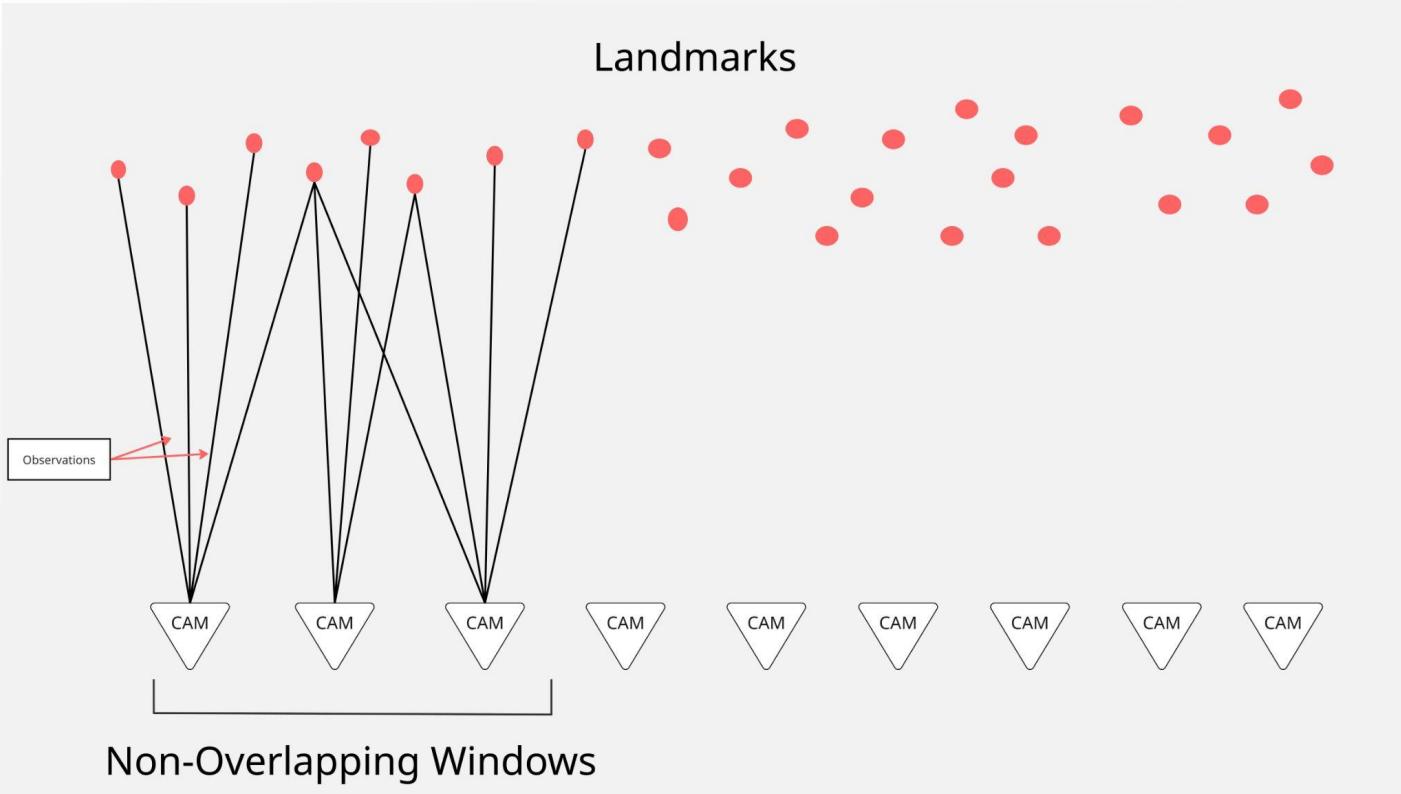
- **Trust-Region Strategy**

- Use **Levenberg–Marquardt** or **Dogleg** to decide step size.
 - Prevents overshooting of estimates in cases of poorly conditioned or highly nonlinear problems.
 - Balances speed of Gauss–Newton with stability of gradient descent in flat or narrow directions.

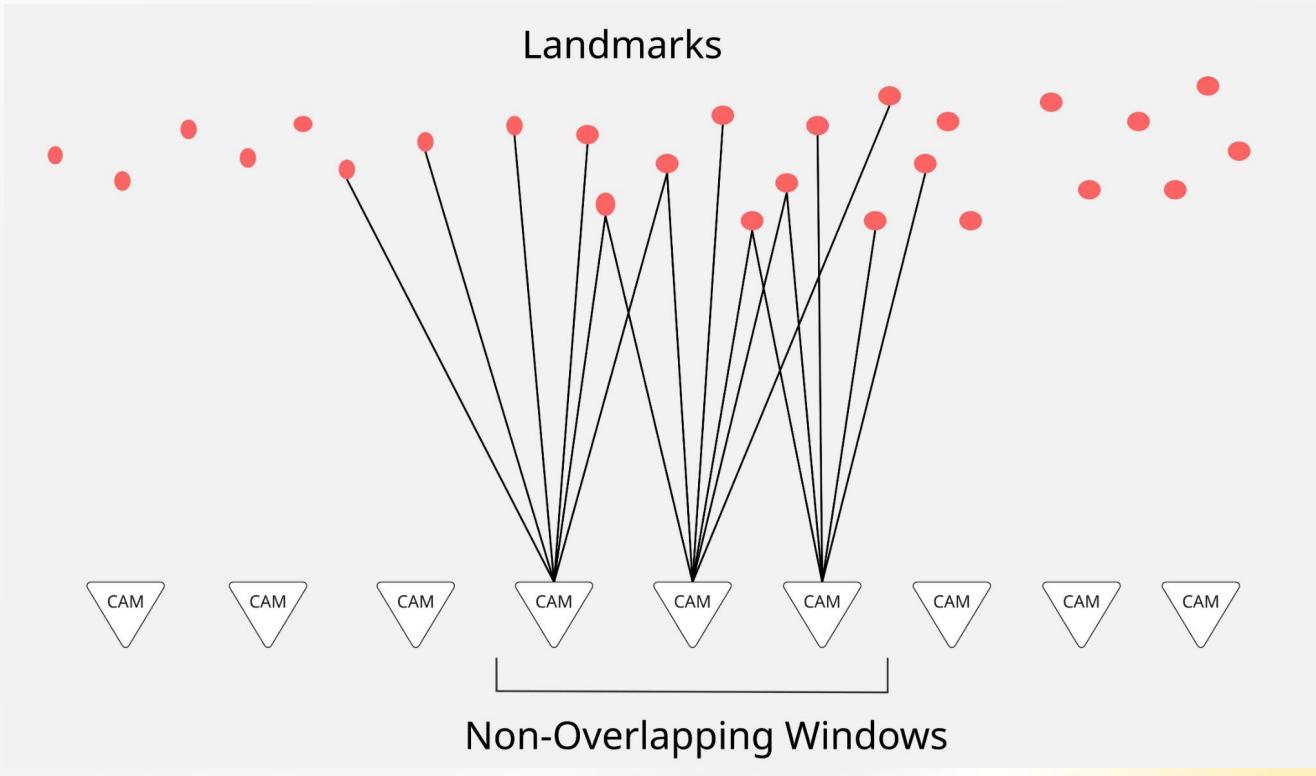
- **Update Parameters and re-iterate**



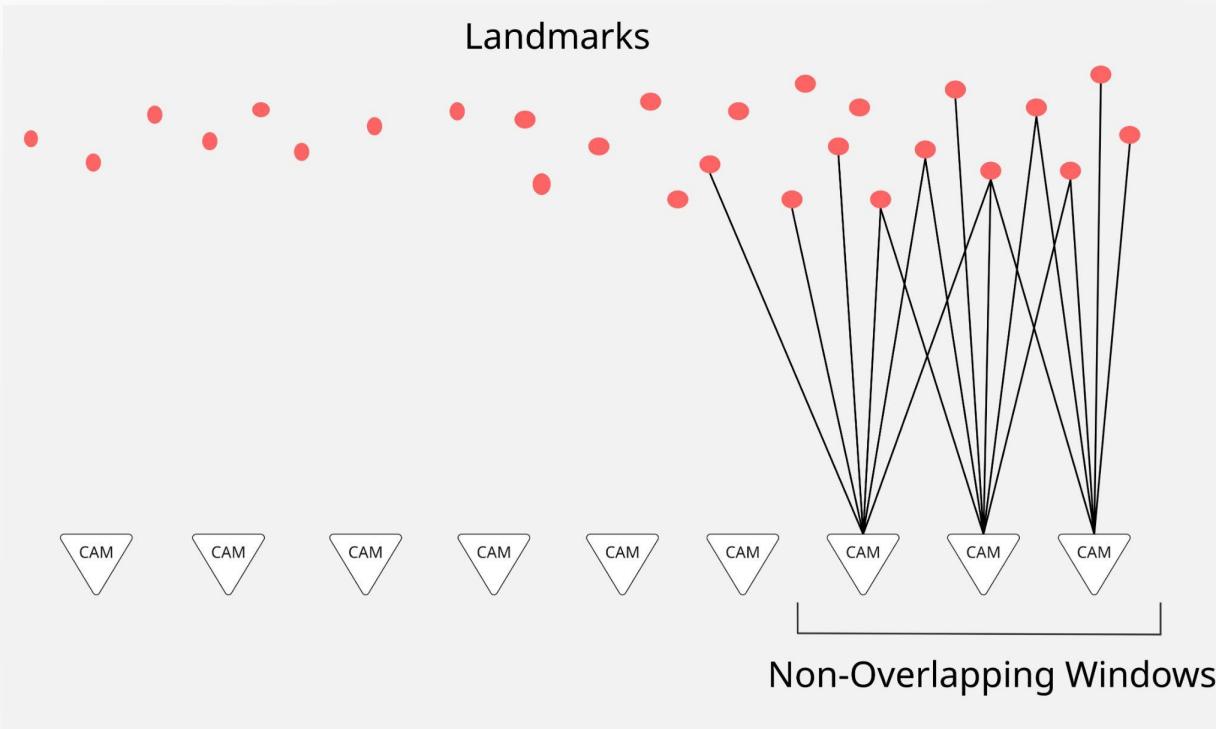
Pipeline - Optimization



Pipeline - Optimization



Pipeline - Optimization



Pipeline - Optimization



Optimizer settings

```
if (config_.fix_first_camera && storage_info.cam_count > 0
    && !camera_params_storage_.empty())
{
    problem.SetParameterBlockConstant(camera_params_storage_[0].data());
}

ceres::Solver::Options options;
options.minimizer_type = ceres::TRUST_REGION;
options.trust_region_strategy_type = ceres::LEVENBERG_MARQUARDT;
options.linear_solver_type = ceres::SPARSE_SCHUR;
options.preconditioner_type = ceres::SCHUR_JACOBI;
options.num_threads = config_.num_threads;
options.minimizer_progress_to_stdout = config_.minimizer_progress_to_stdout;
options.max_num_iterations = config_.max_num_iterations;
options.function_tolerance = config_.function_tolerance;
options.gradient_tolerance = config_.gradient_tolerance;
options.parameter_tolerance = config_.parameter_tolerance;
options.logging_type = ceres::PER_MINIMIZER_ITERATION;
```



Pipeline - Optimization

```
auto cameras_with_observations = camera_window_span  
| std::views::filter([&](int camera_index) {return camera_index >= 0 && camera_index < static_cast<int>(in_dataset.cameras().size());})  
| std::views::transform[&](int camera_index) {return in_dataset.cameras()[camera_index].id;})  
| std::views::filter([&](int camera_id) {  
    auto obs_it = camera_to_observations.find(camera_id);  
    return obs_it != camera_to_observations.end() && !obs_it->second.empty();});  
  
for (int camera_id : cameras_with_observations)  
{  
    // Find camera in input dataset by ID with bounds checking  
    auto cam_it = std::ranges::find_if(in_dataset.cameras(),  
        [camera_id](const auto& cam) { return cam.id == camera_id; });  
    if (cam_it == in_dataset.cameras().end()) {continue;}  
  
    // Extract camera parameters  
    const auto& cam = *cam_it;  
    auto q = cam.rotation().unit_quaternion();  
    auto t = cam.translation();  
    auto k = cam.intrinsics.as_vec3();  
  
    // Create camera parameter vector  
    std::vector<double> cam_params =  
    {  
        static_cast<double>(q.x()),  
        static_cast<double>(q.y()),  
        static_cast<double>(q.z()),  
        static_cast<double>(q.w()),  
        static_cast<double>(t.x()),  
        static_cast<double>(t.y()),  
        static_cast<double>(t.z()),  
        static_cast<double>(k[0]),  
        static_cast<double>(k[1]),  
        static_cast<double>(k[2])  
    };  
  
    // Validate parameters are finite  
    bool valid = std::ranges::all_of(cam_params, [](double p) { return std::isfinite(p); });  
    if (!valid) continue;  
  
    // Store vector and get pointer  
    camera_id_to_storage_idx[camera_id] = camera_params_storage_.size();  
    camera_params_storage_.push_back(std::move(cam_params));  
    double* cam_ptr = camera_params_storage_.back().data();  
    problem.AddParameterBlock(cam_ptr, 10);  
}
```

- Filter and transform used to reject invalid camera indices and extract camera ids
- Camera broken down into 10 parameters - extrinsic + intrinsic
- Only add valid(non Nan) blocks
- Use quaternion parameterization for rotation

Pipeline - Optimization

```
// Process each observation for this camera
auto valid_observations = obs_indices
| std::views::filter([&](int idx) {return idx >= 0 && idx < static_cast<int>(observations.size());})
| std::views::filter([&](int idx) {const auto& obs = observations[idx];
    if (obs.outlier) {
        outlier_filtered++;
        return false;
    }
    return true;
})
| std::views::filter([&](int idx) {return observations[idx].camera_index == camera_id;})
| std::views::filter([&](int idx) {const auto& obs = observations[idx];
    return obs.point_index >= 0 && obs.point_index < static_cast<int>(points.size());});

size_t camera_obs_count = 0;
for (int obs_idx : valid_observations)
{
    total_observations_processed++;
    const auto& obs = observations[obs_idx];

    // Create or reuse point parameter block
    size_t pt_storage_idx;
    auto pit = point_index_to_storage_idx.find(obs.point_index);
    if (pit == point_index_to_storage_idx.end())
    {
        // Create new point parameter block
        const auto& p = points[obs.point_index];
        // Validate point coordinates
        if (!std::ranges::all_of(p, [](double coord) { return std::isfinite(coord); })) {continue;}

        // Store vector and get pointer
        pt_storage_idx = point_params_storage_.size();
        point_index_to_storage_idx[obs.point_index] = pt_storage_idx;
        point_params_storage_.push_back({
            static_cast<double>(p[0]),
            static_cast<double>(p[1]),
            static_cast<double>(p[2])});

        double* pt_ptr = point_params_storage_.back().data();
        problem.AddParameterBlock(pt_ptr, 3);
    } else
    {
        pt_storage_idx = pit->second;
    }
}
```

- Select only valid observations using `views::filter`
- Add the 3 dimensions of the point as a parameter block.

Pipeline - Optimization

```
for (const auto& [camera_id, storage_idx] : valid_camera_updates)
{
    auto cam_it = std::ranges::find_if(dataset.cameras(), [camera_id](const auto& cam) { return cam.id == camera_id; });
    if (cam_it == dataset.cameras().end()) continue;

    std::span<const double, 10> params{camera_params_storage_[storage_idx]};
    Eigen::Quaterniond q(params[3], params[0], params[1], params[2]);
    q.normalize();
    Eigen::Vector3d t(params[4], params[5], params[6]);
    cam_it->T_c_w = typename Dataset::Traits::SE3(typename Dataset::Traits::SO3(q), t);

    // Update intrinsics
    cam_it->intrinsics.focal_length = params[7];
    cam_it->intrinsics.k1 = params[8];
    cam_it->intrinsics.k2 = params[9];
}

auto valid_point_updates = point_index_to_storage_idx
    | std::views::filter([&](const auto& pair) {
        const auto& [point_index, storage_idx] = pair;
        return storage_idx < point_params_storage_.size() &&
            point_index >= 0 &&
            point_index < static_cast<int>(dataset.points().size());
    });
for (const auto& [point_index, storage_idx] : valid_point_updates)
{
    std::span<const double, 3> point_params{point_params_storage_[storage_idx]};
    dataset.points()[point_index] = typename Dataset::Traits::Vec3(
        point_params[0], point_params[1], point_params[2]);
}
```

- For each camera in the window, find its optimized block and update its values.
- For each point that was updated, access the value and update it.



Pipeline - Results



Optimization results in balanced mode

Dataset	Avg Repro Error before	Avg Repro Error after
Ladybug	3.859 pixels	202.43 pixels
Trafalgar Square	8.309 pixels	294.173 pixels



Pipeline - Optimization

```
// cameras out of the current window become constant
auto prev_window_cameras = prev_window_cameras | std::views::filter([&](int cam_id){ return !window_cameras.contains(cam_id); });
std::ranges::for_each(prev_window_cameras, [&](int cam_id)
{
    if (auto it = camera_id_to_index_.find(cam_id); it != camera_id_to_index_.end())
    {
        auto& camera = cameras_[it->second];
        problem_->SetParameterBlockConstant(camera.data());
        camera.is_variable = false;
    }
});

// Entering cameras become variable
auto current_window_cameras = window_cameras | std::views::filter([&](int cam_id){ return !prev_window_cameras.contains(cam_id); });
std::ranges::for_each(current_window_cameras, [&](int cam_id)
{
    if (auto it = camera_id_to_index_.find(cam_id); it != camera_id_to_index_.end())
    {
        auto& camera = cameras_[it->second];
        problem_->SetParameterBlockVariable(camera.data());
        camera.is_variable = true;
    }
});
```

```
// camera.parameters = [qx,qy,qz,qw, tx,ty,tz, fx, k1,k2]
for (int i = 4; i <= 6; ++i)
{
    problem_->SetParameterLowerBound(camera.data(), i, config_.cam_translation_min);
    problem_->SetParameterUpperBound(camera.data(), i, config_.cam_translation_max);
}
problem_->SetParameterLowerBound(camera.data(), 7, config_.focal_min);
problem_->SetParameterUpperBound(camera.data(), 7, config_.focal_max);
problem_->SetParameterLowerBound(camera.data(), 8, config_.k_min);
problem_->SetParameterUpperBound(camera.data(), 8, config_.k_max);
problem_->SetParameterLowerBound(camera.data(), 9, config_.k_min);
problem_->SetParameterUpperBound(camera.data(), 9, config_.k_max);
}
```

- Add all cameras to the optimization problem
- Only set cameras that in the window as variable, set all other cameras as constants.
- Use parameterLowerBound and parameterUpperBound to restrict the optimization space.



Pipeline - Optimization

```
// Adjust in-window support counts based on entering/leaving cameras
std::unordered_set<int> visited_points; // points whose support may have changed
std::ranges::for_each(curr_camera, [&](int cam_id)
{
    auto cam_obs_it = camera_id_to_obs_indices_.find(cam_id);
    std::ranges::for_each(cam_obs_it->second, [&](int obs_idx)
    {
        const auto& obs = observations_[static_cast<size_t>(obs_idx)];
        if (auto pt_it = point_id_to_index_.find(obs.point_id); pt_it != point_id_to_index_.end())
        {
            points_[pt_it->second].in_window_support++;
            visited_points.insert(obs.point_id);
        }
    });
});
std::ranges::for_each(prev_camera, [&](int cam_id)
{
    auto cam_obs_it = camera_id_to_obs_indices_.find(cam_id);
    std::ranges::for_each(cam_obs_it->second, [&](int obs_idx)
    {
        const auto& obs = observations_[static_cast<size_t>(obs_idx)];
        if (auto pt_it = point_id_to_index_.find(obs.point_id); pt_it != point_id_to_index_.end())
        {
            points_[pt_it->second].in_window_support--;
            visited_points.insert(obs.point_id);
        }
    });
});
```

- Keep points observed by any in-window cameras, remove those with zero in-window support
- Freeze points with exactly 1 in-window camera. Optimize those with 2+ in-window cameras.
- Add residuals for ALL observations of kept points, including cross-window observations



Pipeline - Optimization

```
if (!point.is_in_problem)
{
    problem_->AddParameterBlock(point.data(), 3);
    point.is_in_problem = true;
    if (config_.use_parameter_bounds) {
        for (int i = 0; i < 3; ++i)
        {
            problem_->SetParameterLowerBound(point.data(), i, config_.point_min);
            problem_->SetParameterUpperBound(point.data(), i, config_.point_max);
        }
    }

    if (point.in_window_support == 1 && config_.freeze_weak_points)
    {
        problem_->SetParameterBlockConstant(point.data());
        point.is_variable = false;
    } else
    {
        problem_->SetParameterBlockVariable(point.data());
        point.is_variable = true;
    }
}
```

- Keep points observed by any in-window cameras, remove those with zero in-window support
- Freeze points with exactly 1 in-window camera. Optimize those with 2+ in-window cameras.
- Add residuals for ALL observations of kept points, including cross-window observations



Pipeline - Optimization



```
if (!config_.use_schur_ordering) return;

parameter_ordering_ = std::make_shared<ceres::ParameterBlockOrdering>();

// Group 0: Points (to be eliminated via Schur complement)
for (auto& point : points_)
{
    if (point.is_variable && point.is_in_problem)
    {
        parameter_ordering_->AddElementToGroup(point.data(), 0);
    }
}

// Group 1: Cameras (retained in the reduced system)
for (auto& camera : cameras_)
{
    if (camera.is_variable && camera.is_in_problem)
    {
        parameter_ordering_->AddElementToGroup(camera.data(), 1);
    }
}
solver_options_.linear_solver_ordering = parameter_ordering_;
```

Parameter ordering

- For schur complement, ceres wants to eliminate the points block and solve for the camera poses.
- Groups parameters into two categories to explicitly tag the blocks associated with cameras and points.
- Elements added to group 0 are eliminated first while group 1 is retained.
- Ceres does automatically group terms but this guarantees desired performance .



Pipeline - Optimization

```
if (!config_.use_inner_iterations) return;

inner_iteration_ordering_ = std::make_shared<ceres::ParameterBlockOrdering>();

// Group 0: Points (optimize first in inner iterations)
// This allows points to adjust while cameras are held fixed
for (auto& point : points_)
{
    if (point.is_variable && problem_->HasParameterBlock(point.data()))
    {
        inner_iteration_ordering_->AddElementToGroup(point.data(), 0);
    }
}

// Group 1: Cameras (optimize second in inner iterations)
// This allows cameras to adjust while points are held fixed
for (auto& camera : cameras_)
{
    if (camera.is_variable && problem_->HasParameterBlock(camera.data()))
    {
        inner_iteration_ordering_->AddElementToGroup(camera.data(), 1);
    }
}

solver_options_.inner_iteration_ordering = inner_iteration_ordering_;
```

Inner iterations

- Instead of solving the whole problem in one go, inner iterations solves it one group at a time. Decouples the problem.
- This solve is done in addition to the main solve.
- Group 0 - Fix all cameras, optimize only points
- Group 1 - Fix all points, optimize only the cameras
- Each subproblem is simpler and more focused
- Drawback : Adds more iterations and not guaranteed to give u better solution than joint optimization



Pipeline - Optimization



```
9  
10 class AdaptiveHuberLoss : public ceres::LossFunction {  
11 public:  
12     explicit AdaptiveHuberLoss(double initial_delta = 1.0)  
13         : delta_(initial_delta)  
14     {}  
15  
16     void Evaluate(double s, double rho[3]) const override  
17     {}  
18 }
```

Loss functions

- Cauchy and huber loss functions
- Huber : Quadratic loss till threshold and liner beyond.
- Cauchy : Quadratic loss till threshold and log beyond.
- Cauchy is more tolerant to large outlier but can slow down optimization if not tuned.
- Custom loss function usually an overkill.



Pipeline - Results

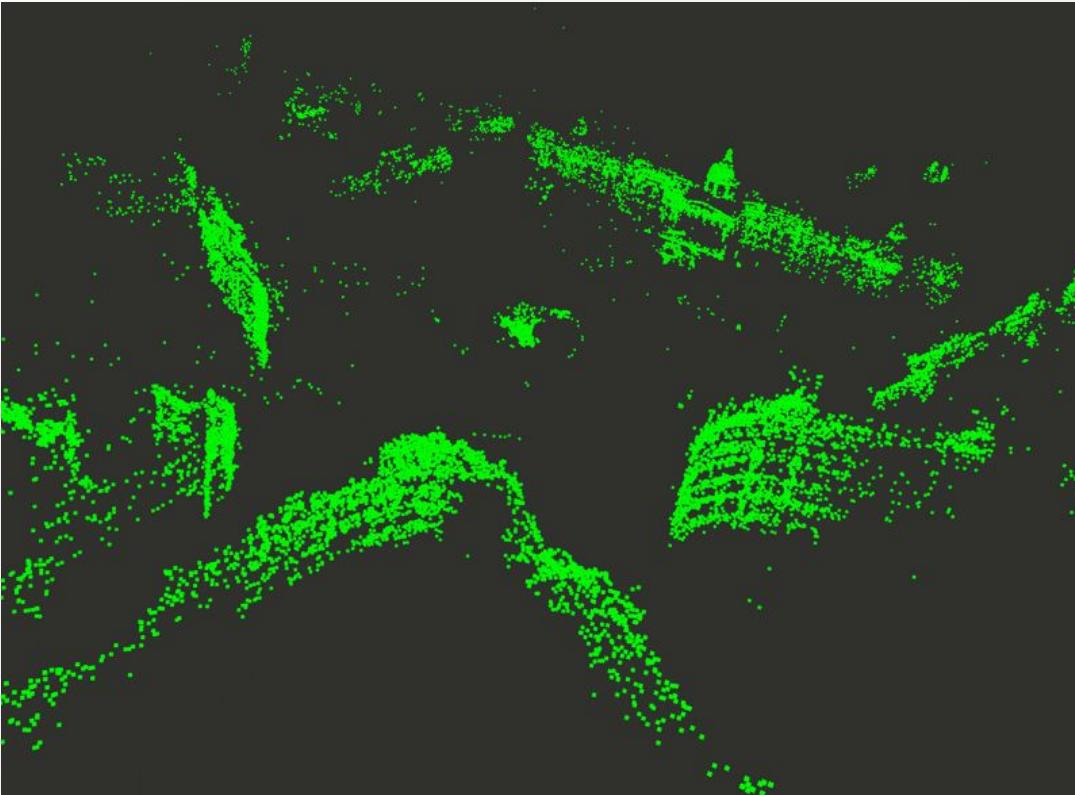


Optimization results

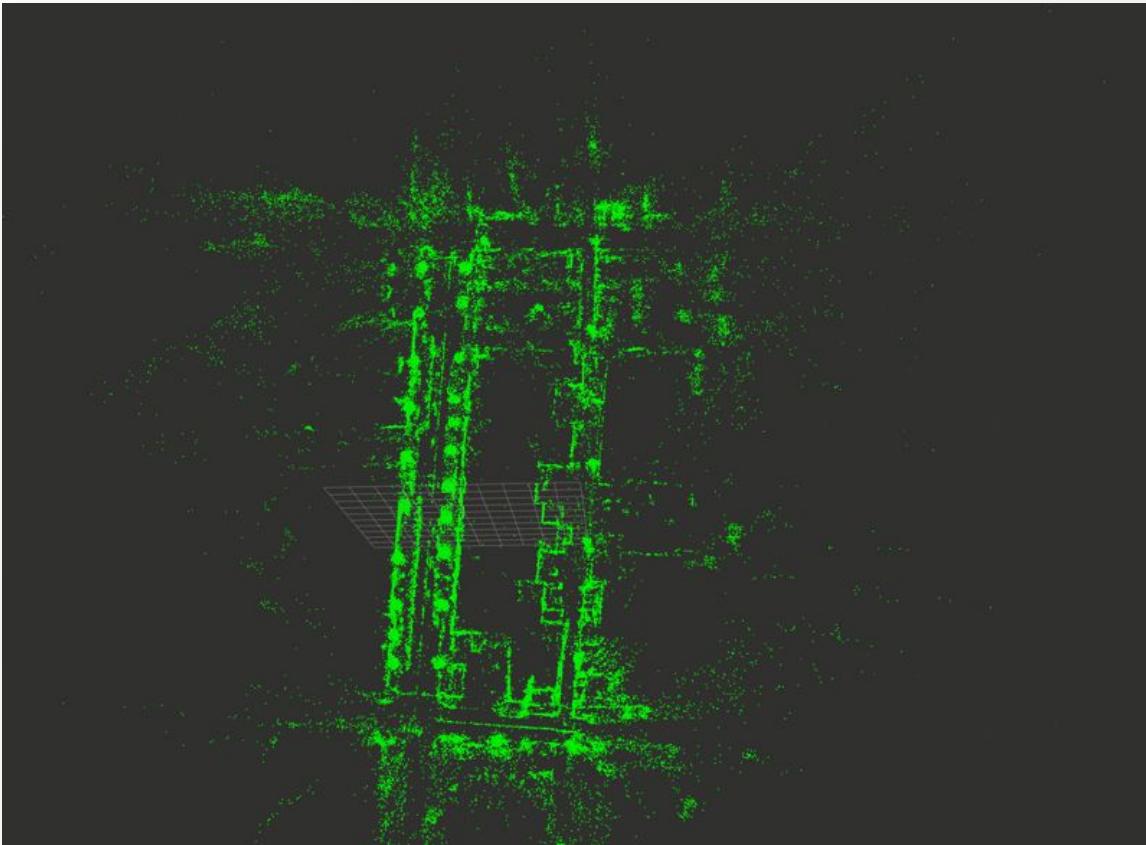
Dataset	Avg Repro Error before	Avg Repro Error after
Ladybug	3.859 pixels	1.060 pixels
Trafalgar Square	8.309 pixels	0.748 pixels



Pipeline - Results



Pipeline - Results



References

Back to basics - concepts by Nicolai Josuttis ([link](#))

Implementing particle filter with C++ ranges - Nahuel Espinosa ([link](#))

Linear algebra and optimization Seminar - Sameer Agarwal([link](#))

Ceres solver - documentation([link](#))



Pipeline - Questions

