

+ 25

Performance of std::expected with Monadic Operations

VITALY FANASKOV



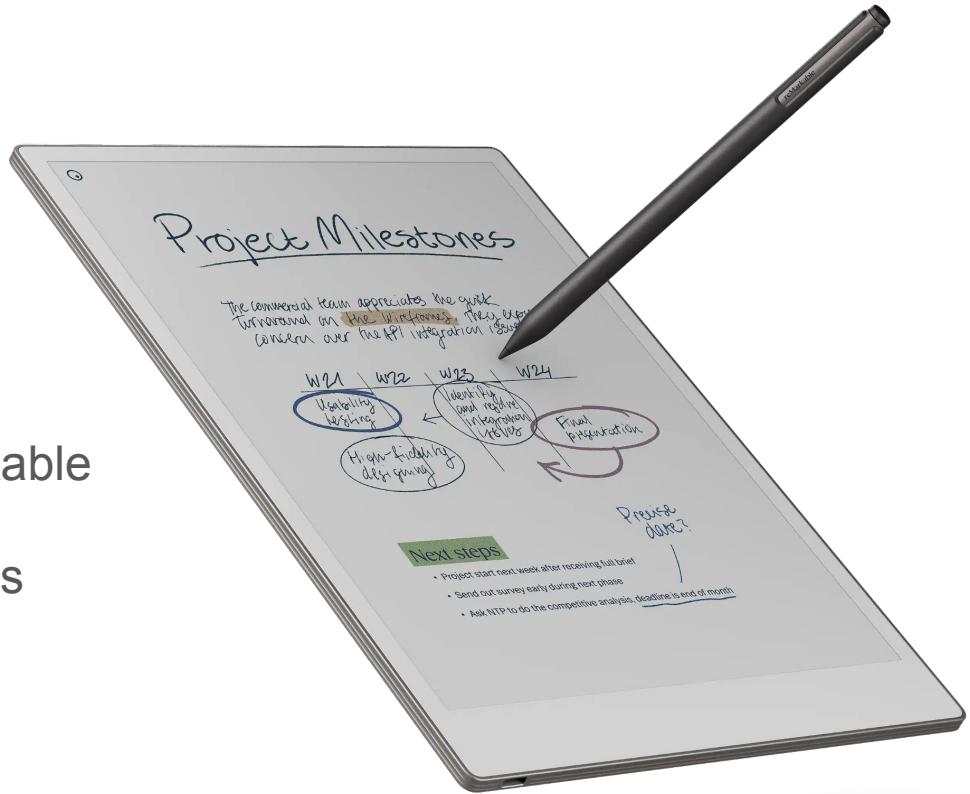
September 13 - 19

About me

Vitaly Fanaskov

Principal software engineer at reMarkable

Work on C++ frameworks and libraries



Agenda

- Error handling in general
- Briefly about exceptions and std::expected
- Important implementation details of monadic operations
- Common use cases and benchmarks
- Practical recommendations

Error handling

Error handling methods

- Return codes
- Function arguments
- Exceptions
- Optional and Result/Either/Expected

Different error handling methods



Error handling in
C++ - Vitaly
Fanaskov - NDC
Techtown 2022

The image shows a video player interface. The main video frame displays a presentation slide with the title "Error handling in C++" and the subtitle "NDC TechTown, 2022". To the right of the video frame, a man with long hair tied back is standing behind a white podium, speaking. The podium has a microphone and a small sign that reads "QUALITY BY DESIGN" and "GRAND KONGSBERG BY KIRKJU COFFEE". A large digital screen at the bottom right of the stage displays the number "1" and the text "NDC [TechTown]". The video player interface includes a play button, volume control, and a progress bar indicating the video is at 0:08 of 44:01. The subtitle "Intro >" is also visible. The overall background is dark, suggesting a conference setting.

What and why?

- Exceptions
- std::expected
- Similar flexibility
- Similar use cases
- The closest alternatives

Briefly about exceptions

Definition

- Unplanned event
 - Occurs during program execution
 - Creates an exception object
-
- The object is handed to the runtime system (*throw*)
 - The runtime system is looking for exception handler
 - Exception handler is found (*catch*) or terminate/crash if not

Exception safety

- No-throw (always succeeds)
- Strong (keeps the same state as before an exception thrown)
- Basic (no resources are leaked, objects' invariants are intact)
- No exception safety (memory corruption, leaks, inconsistency, etc.)

Try-catch example

```
const std::vector vec{1, 2, 3};

try {
    std::println("The fourth element: {}", vec.at(3));
} catch (const std::out_of_range& e) {
    std::println("Error: {}", e.what());
}

// Error: vector::_M_range_check: __n (which is 3) ≥ this->size() (which is 3)
```

Throw an exception

```
template <class E>
E elementByIndex(const std::vector<E>& vec, std::size_t index)
{
    if (index ≥ vec.size()) {
        throw std::out_of_range(
            std::format(
                "Error: index (which is {}) out of range (which is [0, {})).",
                index,
                vec.size()));
    }

    return vec[index];
}
```

Briefly about std::expected

Basic example of using std::expected

```
#include <expected>

std::expected<int, Error> expectedBox;
expectedBox = 42;

std::println("The value is: {}", expectedBox.value());
```

std::expected as a return value

When an operation can fail and we need to know why:

```
std::expected<Widget, WidgetError> LoadWidget()
{
    // If error
    return std::unexpected(WidgetError{ /* ... */ });

    // Actual result
    return Widget{ /* ... */ };
}
```

Process std::expected

```
void loadWidget()
{
    if (const auto widgetBox = getNewWidget(); widgetBox.has_value()) {
        const auto widget = widgetBox.value();
        // Do something with the widget ...
    } else {
        const auto error = widgetBox.error();
        // Handle the error ...
    }
}
```

Process std::expected with monadic operations

```
getWidget()
    .and_then([](const auto &widget) → std::expected<Widget, WidgetError> {
        // Do something with the widget ...
        return widget;
})
    .transform([](const auto &widget) → ID { return widget.id(); })
    .or_else([](const auto& error) → std::expected<ID, WidgetError> {
        log(error);
        // Possibly recover and/or cleanup ...
        // Return value or error ...
})
}
```

Available monadic operations

`transform(F f)`

f: Value → ValueOut

`transform_error(F f)`

f: Error → ErrorOut

`and_then(F f)`

f: Value → expected<ValueOut, Error>

`or_else(F f)`

f: Error → expected<Value, ErrorOut>

For more information



[Monadic Operations in Modern C++: A Practical Approach - Vitaly Fanaskov - CppCon 2024](#)

A slide from CppCon 2024 featuring a presentation titled "Monadic Operations in Modern C++: A Practical Approach" by Vitaly Fanaskov. The slide includes the CppCon logo and the date "2024".

**Monadic Operations in Modern C++:
A Practical Approach**

VITALY FANASKOV

Cppcon
The C++ Conference

2024 | September 15 - 20

Important implementation details of std::expected

Types

- `template <class E> class unexpected`
- `template <class E> class bad_expected_access`
- `template <class T, class E> class expected`

Data structure

```
union {
    T val;
    E unex; // NOTE: not unexpected<E>!
};

bool has_val;
```

Data structure

```
union {  
    T val;  
    E unex; // NOTE: not unexpected<E>!  
};  
(bool has_val;
```

Simplified “and_then”

```
template <class F>
auto and_then(F&& f) &
{
    using U = std::invoke_result_t<F, T&>; // + compile-time checks
    if (has_val) {
        return std::invoke(std::forward<F>(f), val);
    } else {
        return U{unex};
    }
}
```

Simplified “and_then”

```
template <class F>
auto and_then(F&& f) &
{
    using U = std::invoke_result_t<F, T&>; // + compile-time checks
    if (has_val) {
        return std::invoke(std::forward<F>(f), val);
    } else {
        return U{unex};
    }
}
```

Simplified “transform”

```
template <class F>
auto transform(F&& f) &
{
    using U = std::invoke_result_t<F, T&>;
    using R = expected<U, E>;

    if (has_val) {
        return R{std::invoke(std::forward<F>(f), val)};
    } else {
        return R{unex};
    }
}
```

Simplified “transform”

```
template <class F>
auto transform(F&& f) &
{
    using U = std::invoke_result_t<F, T&>;
    using R = expected<U, E>;
    if (has_val) {
        return R{std::invoke(std::forward<F>(f), val)};
    } else {
        return R{unex};
    }
}
```

Simplified “or_else”

```
template <class F>
auto or_else(F&& f) &
{
    using U = std::invoke_result_t<F, E&>; // + compile-time checks
    if (has_val) {
        return U{val};
    } else {
        return std::invoke(std::forward<F>(f), unex);
    }
}
```

Simplified “or_else”

```
template <class F>
auto or_else(F& f) &
{
    using U = std::invoke_result_t<F, E&>; // + compile-time checks
    if (has_val) {
        return U{val};
    } else {
        return std::invoke(std::forward<F>(f), unex);
    }
}
```

Simplified “transfrom_error”

```
template <class F>
auto transform_error(F&& f) &
{
    using U = std::invoke_result_t<F, E&>;
    using R = expected<T, U>;

    if (has_val) {
        return R{val};
    } else {
        return R{std::invoke(std::forward<F>(f), unex)};
    }
}
```

Simplified “transfrom_error”

```
template <class F>
auto transform_error(F&& f) &
{
    using U = std::invoke_result_t<F, E&>;
    using R = expected<T, U>;
    if (has_val) {
        return R{val};
    } else {
        return R{std::invoke(std::forward<F>(f), unex)};
    }
}
```

Potential issues

- Extra “has_val” field
- “if” inside all operations
- Potentially: copies of value or error
- Potentially: result type changes

Compare use cases

Compare implementation

```
double sqrt(double v)
{
    if (v < 0.0) {
        throw std::invalid_argument(
            "Negative number");
    } else {
        return std::sqrt(v);
    }
}
```

```
using Result =
    std::expected<double, std::string>;
```



```
Result sqrt(double v) noexcept
{
    if (v < 0.0) {
        return std::unexpected(
            "Negative number");
    } else {
        return std::sqrt(v);
    }
}
```

Some computation

```
double calc(double v)
{
    return log(sqrt(log(v)));
}
```

```
Result calc(double v) noexcept
{
    return log(v)
        .and_then(&sqrt)
        .and_then(&log);
}
```

//+ two hidden "ifs"

Handle errors

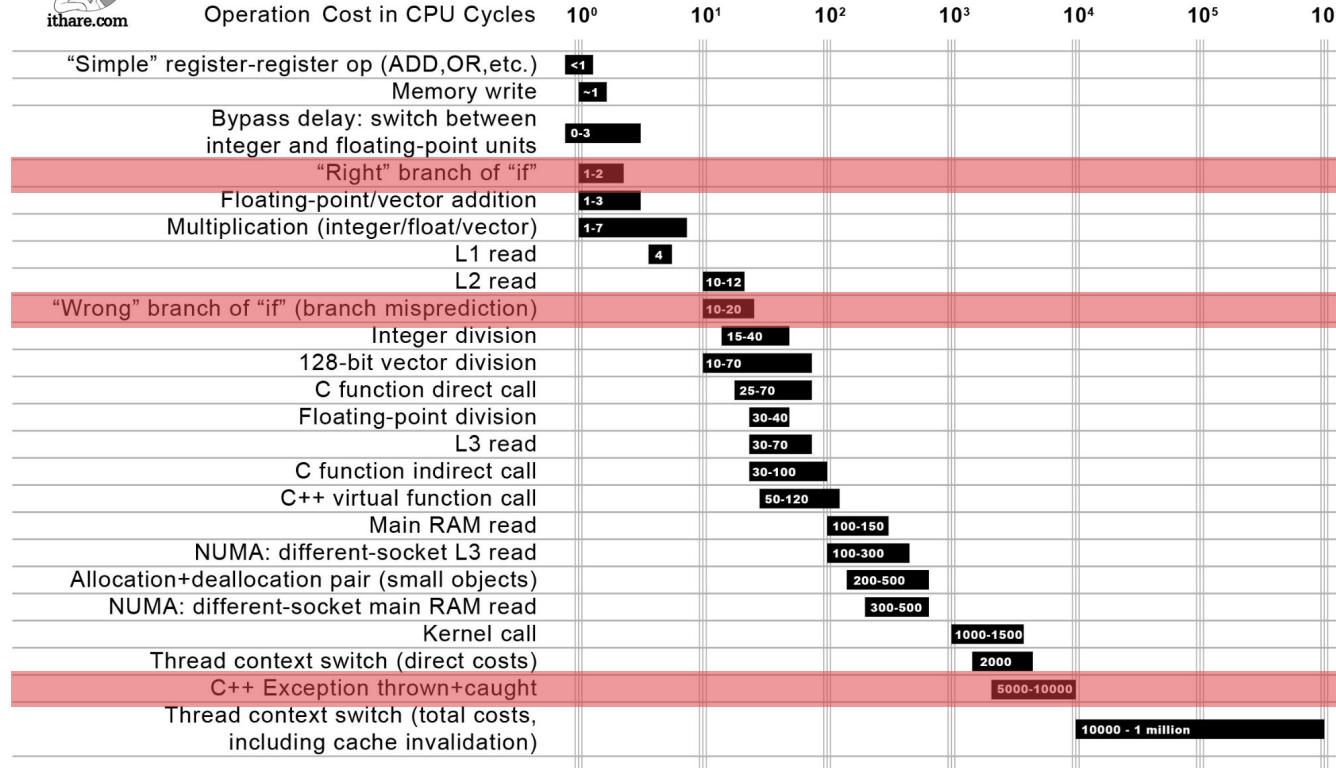
```
try {
    const auto result =
        with_exceptions::calc(125);
    std::println("Result: {}", result);
} catch (const std::invalid_argument& e) {
    std::println("Error: {}", e.what());
}
```

```
const auto result =
    with_expected::calc(125);
if (result.has_value()) {
    std::println(
        "Result: {}", result.value());
} else {
    std::println(
        "Error: {}", result.error());
}
```

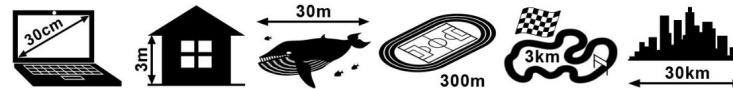
Some reference data



Not all CPU operations are created equal



Distance which light travels while the operation is performed



Operations cost

Operation	Cost in CPU cycles
“Right” branch of “if”	1-2
“Wrong” branch of “if” (branch misprediction)	10-20
C++ exception thrown+caught	5000-10000

Why are exceptions so slow?



- The exceptions are allocated in dynamic memory.
- The table driven unwinder logic used by modern C++ compilers grabs a global mutex to protect the tables from concurrent changes.

[P2544R0: C++ exceptions are becoming more and more problematic](#)

Let's do some benchmarks!

Do your own benchmarks!

- Build type (release or release with debug info mode)
- Certain toolchain
- Hardware
- Use cases
- Data types
- Number of elements/operations

Hardware config

CPU:

```
model name      : Intel(R) Core(TM) i9-9900K CPU @ 3.60GHz  
cpu MHz        : 3600.000  
cache size     : 16384 KB
```

RAM (x4):

```
Size: 16 GB  
Speed: 3200 MT/s  
Configured Voltage: 1.2 V
```

RAM timings:

AA-RCD-RP-RAS (cycles) as DDR4-1866	13-13-13-31
AA-RCD-RP-RAS (cycles) as DDR4-1600	11-11-11-27

How to compare

- Pick a relatively "light" operation
- Perform several operations in sequence (10, 100, 1000, 10000)
- Run without errors
- Run with an error (at the beginning, middle, and end)
- Check `std::expected` or catch an exception

Operation code

```
Data operation(Data obj)
{
    if (wrongStep(obj)) {
        throw /* ... */;
    } else {
        incOp(obj.counter);
        return obj;
    }
}
```

```
/* ... */obj = operation(obj);
/* ... */
```

```
std::expected/* ... */ operation(Data obj) noexcept
{
    if (wrongStep(obj)) {
        return /* ... */;
    } else {
        incOp(obj.counter);
        return obj;
    }
}
```

```
operation(obj).and_then(&operation).
/* ... */
```

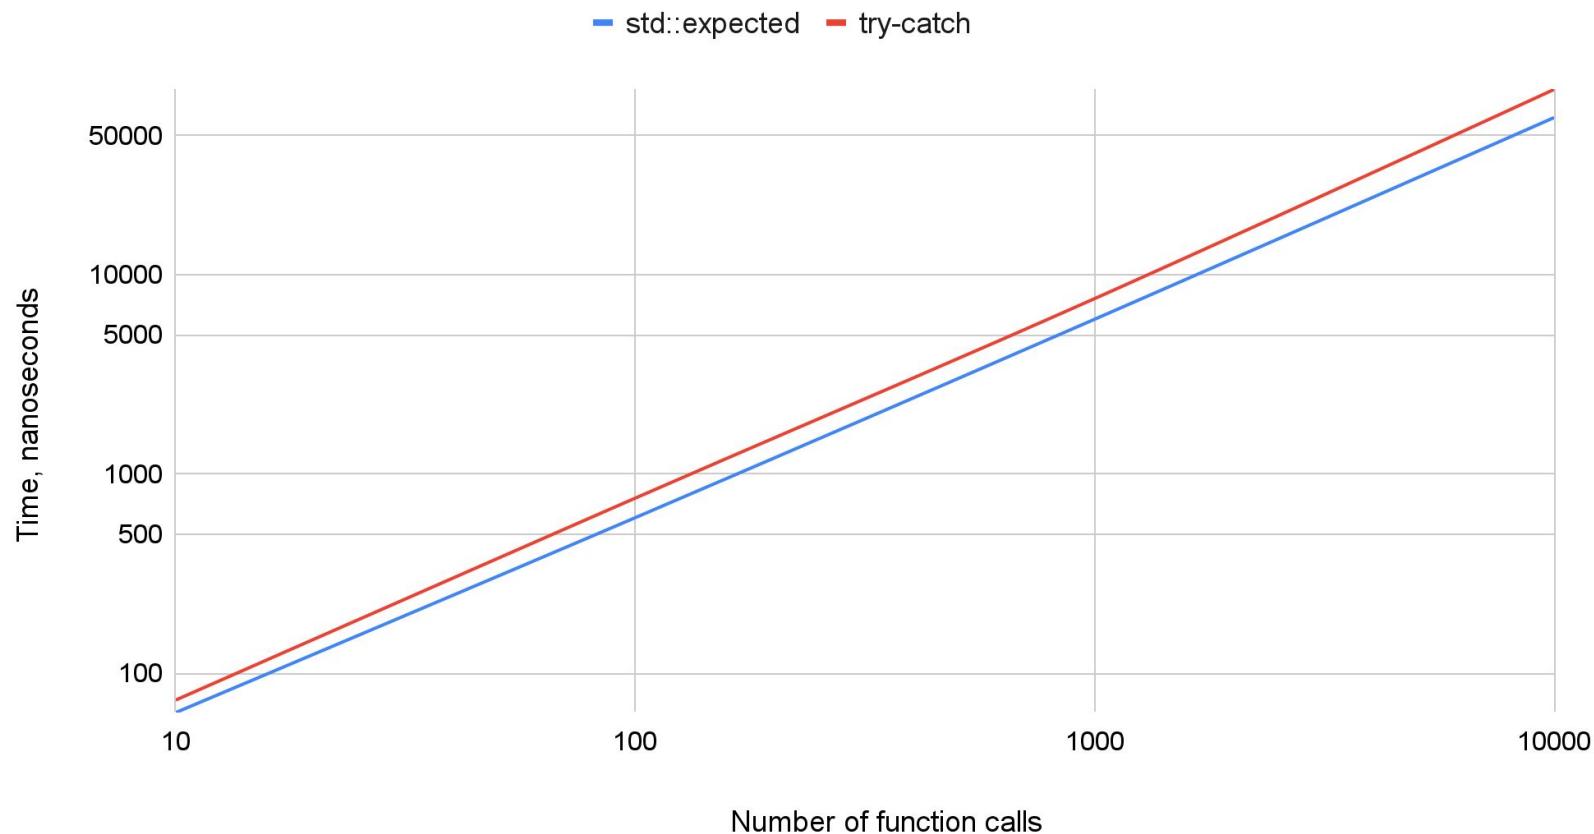
Close to declarative software design

- N monadic operations

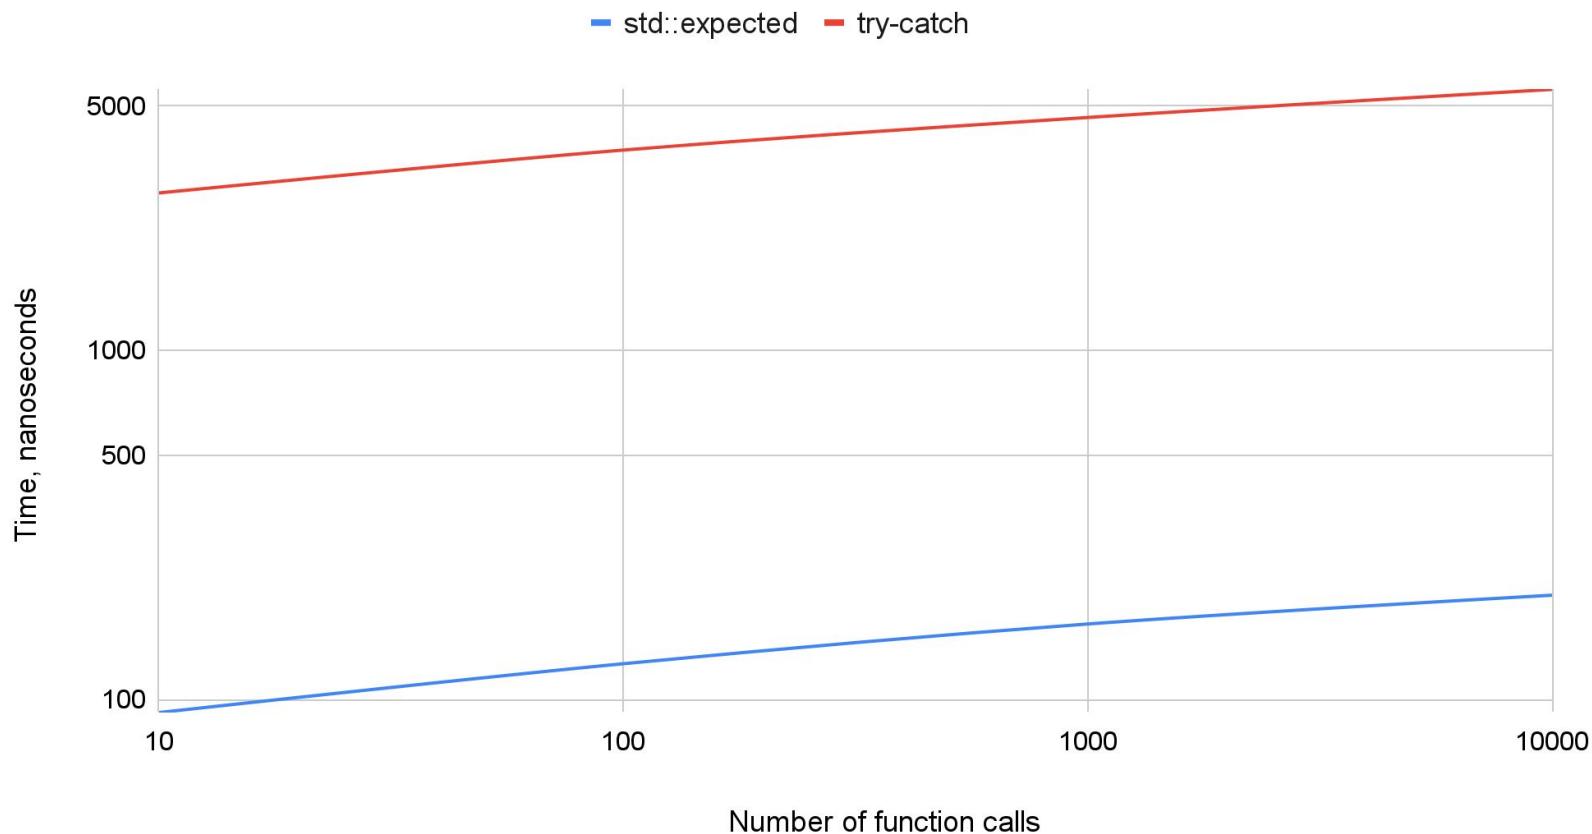
Compared to:

- 1 throw
- 1 catch

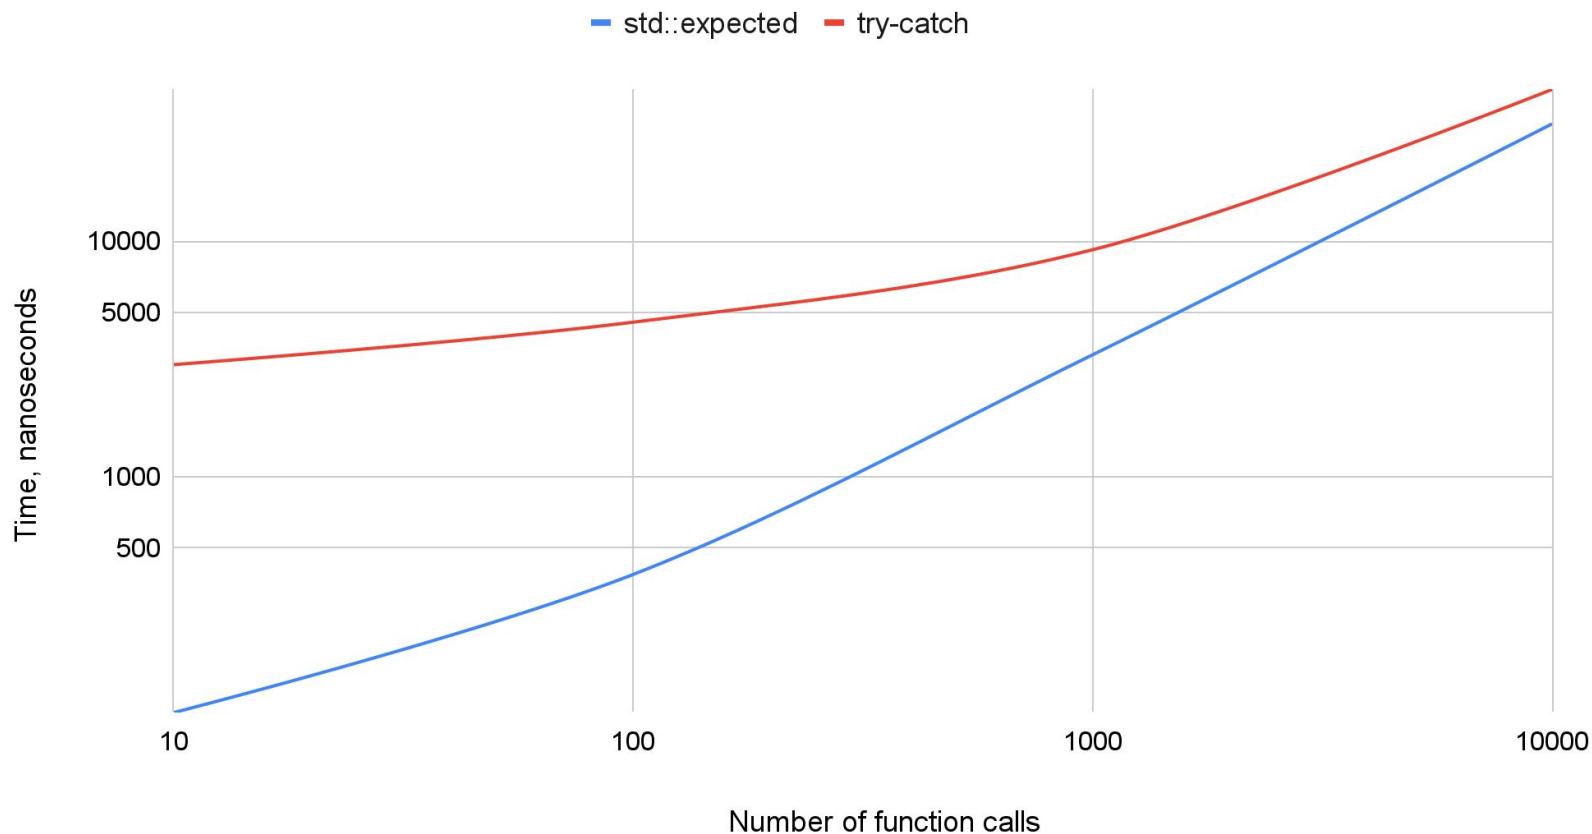
Simple run (no errors)



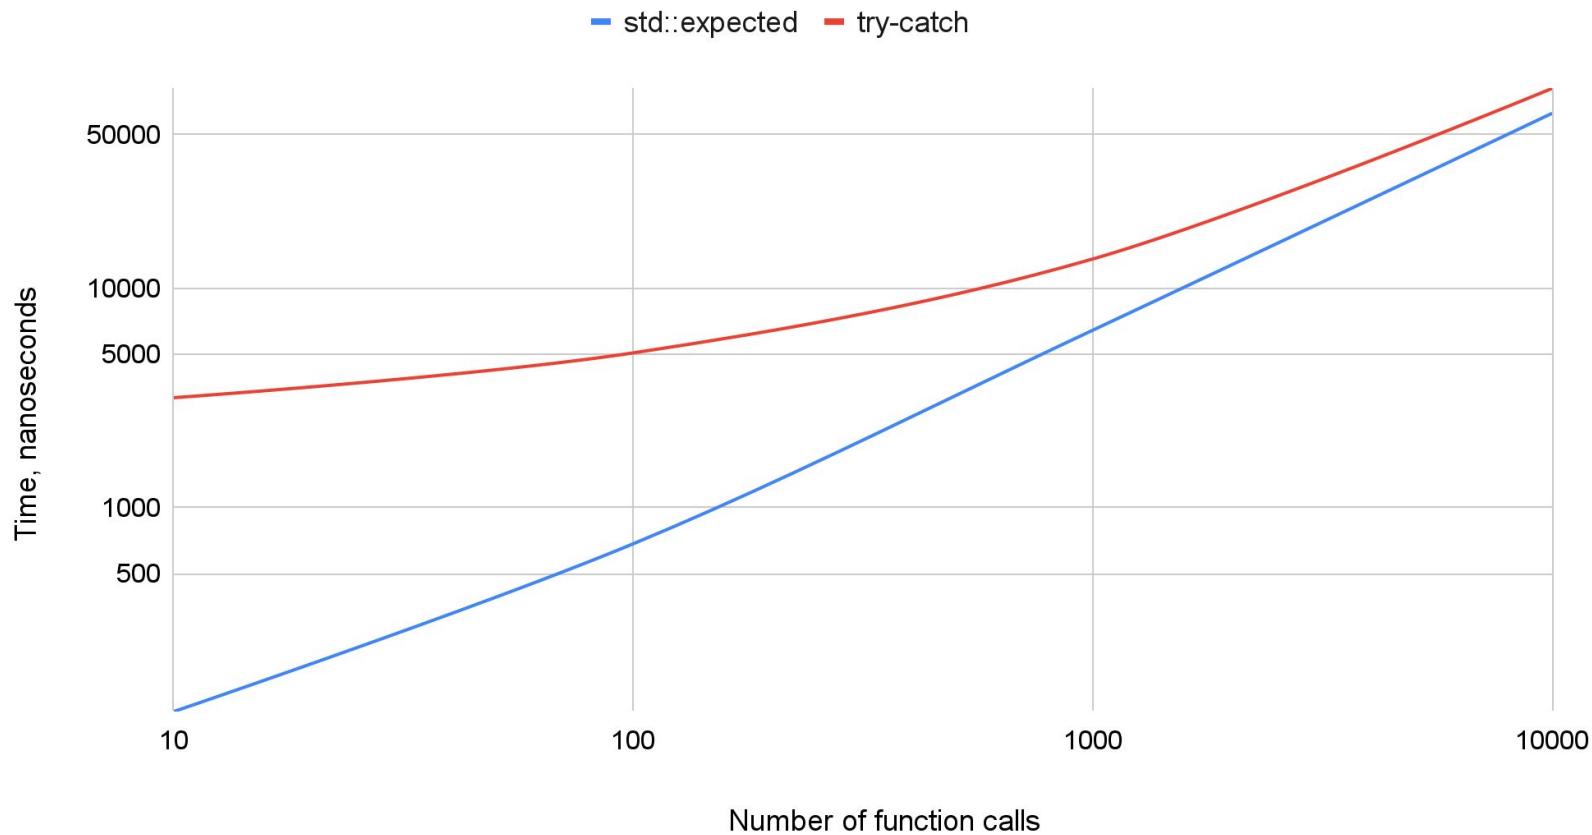
An error at the very beginning



An error in the middle



An error at the end



Conclusions

- In general, code with monadic operations is faster than with exceptions
- If an error happens at the beginning of the calculations, then code with monadic operations is significantly faster
- For a call chain of approx. less than 500 operations, monadic operations work at least two times faster than exceptions

What do we do with the extra copy/move operations?

A simple operation

```
using Result =
    std::expected<Noisy, std::string>;  
  
Result operation()
{
    return Noisy{};
}  
  
// ...  
  
const auto result = operation();
```

Ctor: 0x7ffc99fddda7
Move ctor: 0x7ffc99fdddc8
Dtor: 0x7ffc99fddda7
Dtor: 0x7ffc99fdddc8

and_then/transform operation

```
std::ignore = result
.and_then(
    [](const Noisy& noisy) → Result {
        doSmth(noisy);
        return noisy;
})
.and_then(
    [](const Noisy& noisy) → Result {
        doSmth(noisy);
        return noisy;
});
```

Do smth: 0x7fffb9c65828
Copy ctor: 0x7fffb9c657d8
Do smth: 0x7fffb9c657d8
Copy ctor: 0x7fffb9c65800
Dtor: 0x7fffb9c65800
Dtor: 0x7fffb9c657d8

Can we just move?

```
std::ignore = result.and_then(Compilation error!
    [](Noisy&& noisy) → Result {
        doSmth(noisy);
        return noisy;
});
```

Explicit move or a temporary object

```
std::ignore = std::move(result).and_then(  
    [](Noisy&& noisy) → Result {  
        doSmth(noisy);  
        return noisy;  
    }).and_then(  
    [](Noisy&& noisy) → Result {  
        doSmth(noisy);  
        return noisy;  
    })
```

Do smth: 0x7fff8da77a68
Move ctor: 0x7fff8da77a18
Do smth: 0x7fff8da77a18
Move ctor: 0x7fff8da77a40
Dtor: 0x7fff8da77a40
Dtor: 0x7fff8da77a18

Transfer an error

```
using AnotherResult =
    std::expected<std::string, Noisy>;  
  
AnotherResult result = /* error */;  
  
std::ignore = std::move(result)
    .and_then([](const std::string& o) → AnotherResult {
        return o;
})
    .and_then([](const std::string& o) → AnotherResult {
        return o;
})
    .and_then([](const std::string& o) → AnotherResult {
        return o;
});
```

Move ctor: 0x7fff55df36b8
Move ctor: 0x7fff55df36e0
Move ctor: 0x7fff55df3708

Conclusions

- Copy/move operations for value/error will be performed on each step
- Highly unlikely will they be optimised out

Practical recommendations

Use "in-place" constructors

```
Result operation() Ctor: 0x7ffc82707b0
{
    return Result{std::in_place,};
}
```

```
Result operation()
{
    return Result{std::unexpected,};
}
```

Use smart-pointers or references

```
std::expected<std::unique_ptr<Noisy>, std::string> operation()
{
    return std::make_unique<Noisy>();
}
```

```
std::expected<std::reference_wrapper<Noisy>, std::string> operation(Noisy& noisy)
{
    doSmth(noisy);
    return std::ref(noisy);
}
```

Use simple data structures

```
struct Noisy
{
    int field1{};
    double field2{};
    std::string field3{};
    Foo field4{};
};
```

- Just a structure
- With “regular” fields only
- You can simply pass by value in many cases

Thank you!