

+ 25

The Joy of C++26 Contracts (and Some Mythconceptions)

HERB SUTTER
≡ CITADEL | Securities



20
25 | September 13 - 19

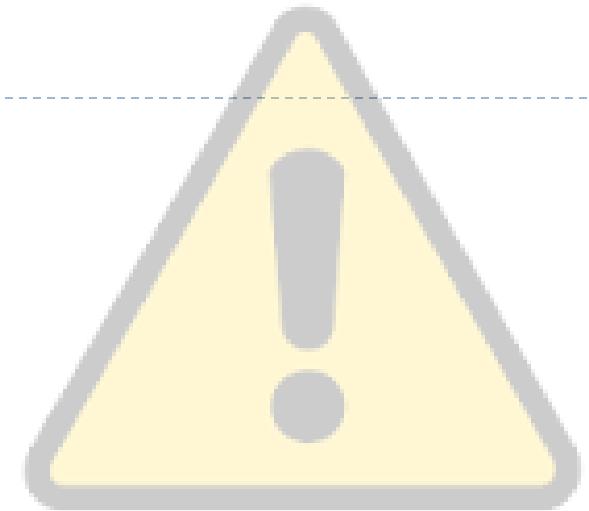
Disclaimer

Current status

Contracts are part of draft C++26, with consensus in February

We have two more meetings of technical fit-and-finish,
including objections to contracts

Every concern will continue to get a full hearing



This talk summarizes my best understanding of the design,
expected best practices, FAQs, and major concerns

But (a) I could be wrong (and am happy to be wrong and learn),
and (b) you should treat the answers herein as provisional

Let me make myself a target...

Some experts strongly believe C++26 contracts are
beautiful and ready to go
not ready and shouldn't ship

Q: How many of you think that I can offend both groups in two sentences?

C++26 contracts have open questions and don't yet have widespread deployment experience.

We “likely” won’t get a substantially better design than C++26 contracts, without wishing for magic linker improvements.



Roadmap

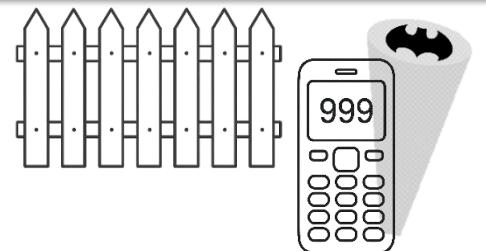
Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,
::handle_contractViolation()`



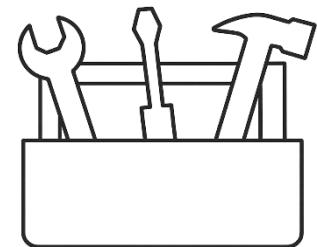
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

Including open questions



Functional safety (vs. memory safety)

Functional safety: **Incremental confidence** that the system **meets its requirements**

Primary aim of contracts

Write one contract \Rightarrow get one check

Incremental confidence: Depends on how many contracts you have (your own + in libraries)

Memory safety: **Guarantees** that the system is **free of certain classes of bugs**

Build a certain way \Rightarrow get a guarantee (e.g., systematic language rules, systematic analysis rules)

Contracts can also be mechanically generated for this, as an implementation detail

Potential example: “Dereferencing a null pointer is checked by `contract_assert(pointer)`”

Build-time guarantee: Automatically enforced by language or other build step (e.g., analyzer)

Example

Consider this code:

```
void my_function( Widget* w ) {
    if( is_gnarly(w) ) {
        write_to_gnarly_file( data.header(), w->to_string() );
    }
}
```

What could be better about this function? Discuss...

Example

Consider this code:

```
void my_function( Widget* w ) {
    assert(w);
    if( is_gnarly(w) ) {
        write_to_gnarly_file( data.header(), w->to_string() );
    }
}
```

Assert liberally

68. Assert liberally to document internal assumptions and invariants.

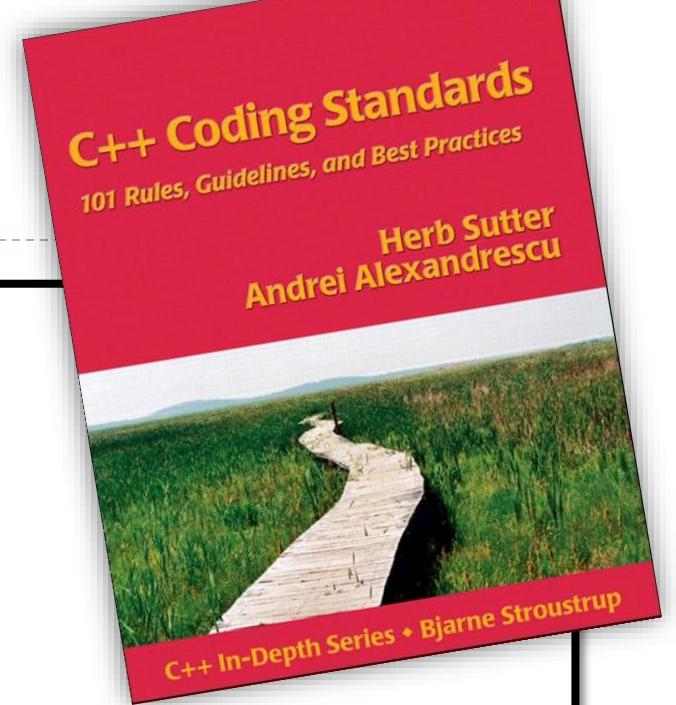
Summary

Be assertive! Use `assert` or an equivalent liberally to document assumptions internal to a module (i.e., where the caller and callee are maintained by the same person or team) that must always be true and otherwise represent programming errors (e.g., violations of a function's postconditions detected by the caller of the function). (See also Item 70.) Ensure that assertions don't perform side effects.

Discussion

It's hard enough to find an error in your code when you're looking for it; it's even harder when you've assumed your code is error-free. —Steve McConnell

It is hard to overestimate the power of assertions. The `assert` macro and alternatives



It is hard to overestimate the power of assertions. The `assert` macro and alternatives such as compile-time (and, less preferably, run-time) assertion templates are invaluable tools for detecting and debugging programming errors during a project's development. Of all such tools, they arguably have the best complexity/effectiveness ratio. The success of a project can be conditioned at least in part by the effectiveness with which developers use assertions in their code.

Assertions commonly generate code in debug mode only (when the `NDEBUG` macro is not defined), so they can be made “free” in release builds. Be generous with what you check. Never write expressions with side effects in `assert` statements. In release mode, when the `NDEBUG` macro is defined, `asserts` don’t generate any code at all:

```
assert( ++i < limit );                                // bad: i is incremented in debug mode only
```

According to information theory, the quantity of information in an event is inversely proportional to the likelihood of that event happening. Thus, the less likely some `assert` is to fire, the more information it will bring to you when it does fire.

Avoid `assert(false)`, and prefer `assert(!"informational message")`. Most compilers will helpfully emit the string in their error output. Also consider adding `&&` “informational message” to more complex assertions, especially instead of a comment.

A YouTube short I recently sent to the C++ committee...

```
6 import { assert } from "node:console";
5
4 type Foo = {
3     bar?: number
2 }
1
7 function morphFoo(foo: Foo): number {
1     return foo.bar * 5;
2 }
3
```

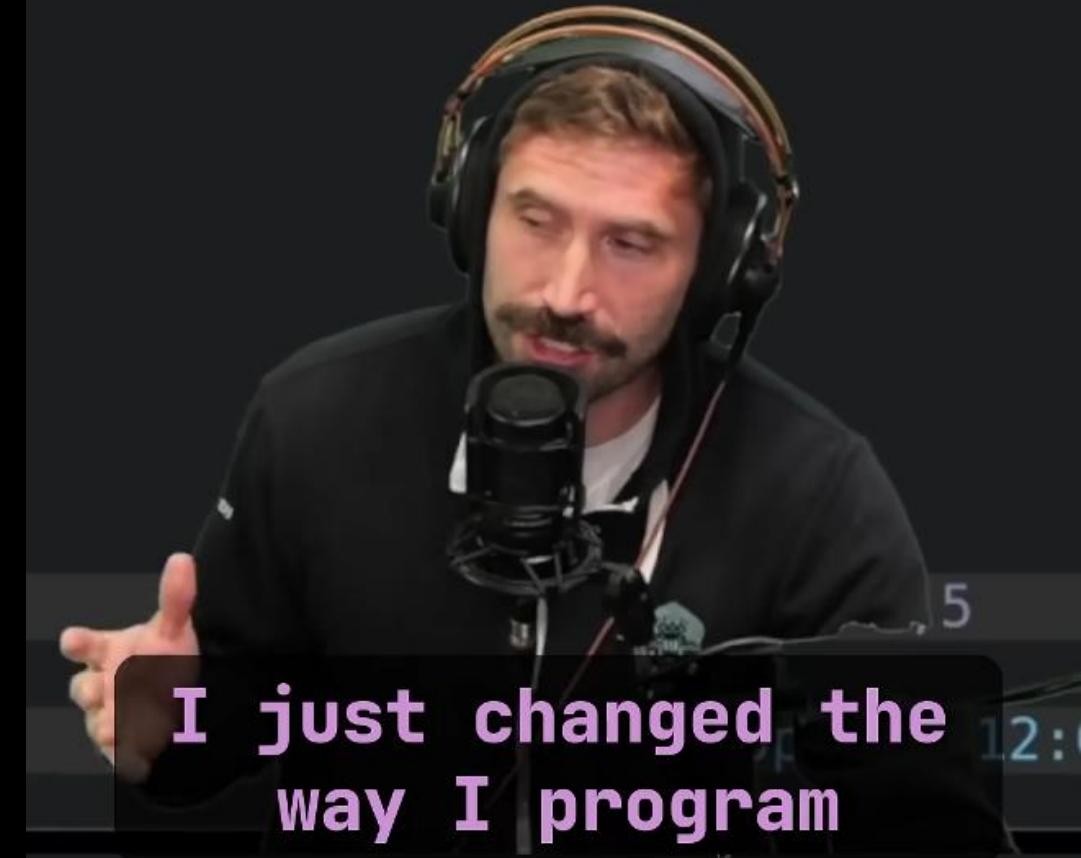
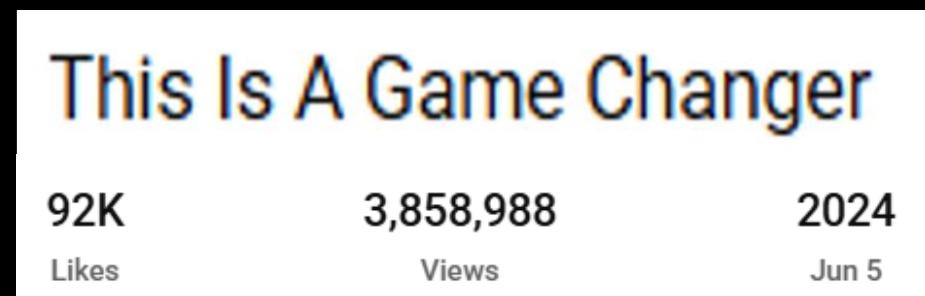
This Is A Game Changer

92K
Likes

3,858,988
Views

2024
Jun 5

A YouTube short I recently sent to the C++ committee...



```
6 import { assert } from "node:console";
5
4 type Foo = {
3   bar?: number
2 }
1
7 function morphFoo(foo: Foo): number {
1   return foo.bar * 5;
2 }
3
```

Roadmap

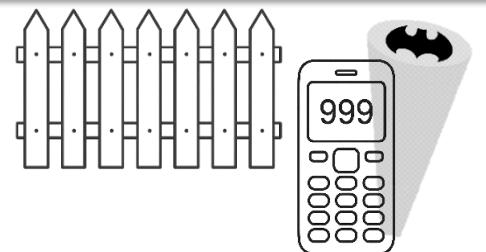
Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,
::handle_contractViolation()`



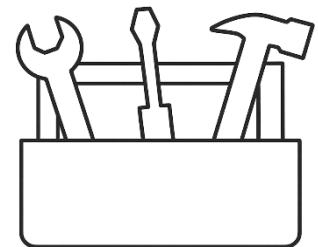
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

Including open questions



Roadmap

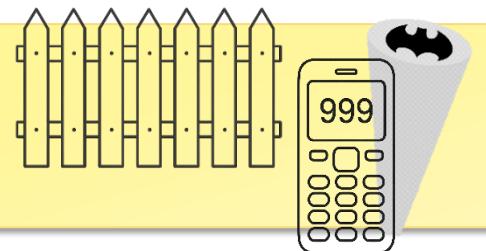
Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,`
`::handle_contractViolation()`



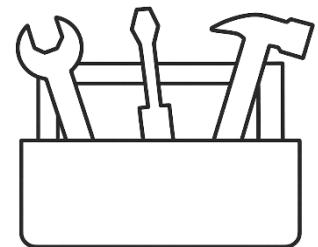
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

Including open questions



A C***GPT prompt

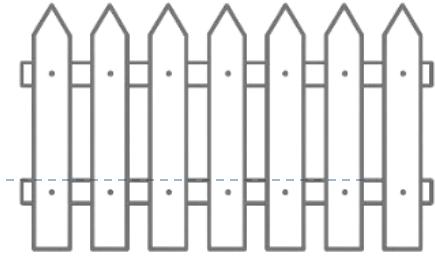
prompt #2: "show me the
chatgpt logo with amusing
little red devil horns"



"Give an example C++ function that should use assert() in three different ways"

```
double calculateAverage(const std::vector<int>& numbers) {  
    assert(!numbers.empty() && "Input vector must not be empty.");  
    for (int num : numbers) {  
        assert(num >= 0 && "All numbers must be non-negative.");  
        sum += num;  
    }  
    double average = static_cast<double>(sum) / numbers.size();  
    assert(average <= 100.0 && "Average exceeds maximum expected value.");  
    return average;  
}
```

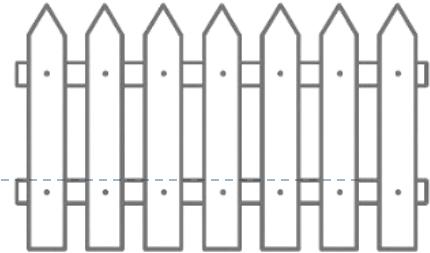
Corrected + using C++26 contracts



In current draft C++26, we can write:

```
auto calculateAverage(std::span<int const> numbers) -> double
    pre(!numbers.empty() && "Input vector must not be empty.")
    post(ret: 0.0 <= ret && ret <= 100.0 && "Average outside expected range.")
{
    auto sum = 0;
    for (int num : numbers) {
        contract_assert(0 <= num && num <= 100 && "Numbers must be in (0, 100).");
        sum += num;
    }
    return static_cast<double>(sum) / numbers.size();
}
```

The four horsemen (semantics)



Standard semantic	Check (evaluate)?	Call violation handler?	Terminate?
ignore	—	—	—
observe	✓	✓	—
enforce	✓	✓	✓
quick_enforce	✓	—	✓

Covers all
four basic
combinations

`::handle_contractViolation`

Optionally, you can supply a global handler (invoked only for `observe` or `enforce`)

```
::handle_contractViolation( std::contracts::contractViolation const& ) noexceptopt;
```

Parameter information

`location()`

QoI: for a precondition, should be the call site source_location

`comment()`

QoI: could be the pretty-printed predicate

`kind()`

`pre`, `post`, `contract_assert`

`semantic()`

`ignore`, `observe`, `enforce`, `quick_enforce` (but open-ended)

`detection_mode()`

predicate was false vs. predicate threw an exception

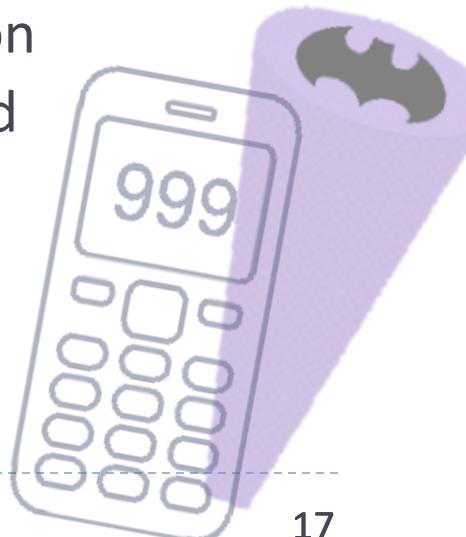
`is_terminating()`

whether a *standard* terminating semantic was used

Like all global replacement functions, provided:

By the whole program owner

At link time (avoiding security issues)



Roadmap

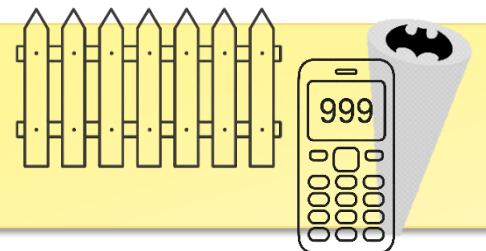
Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,`
`::handle_contractViolation()`



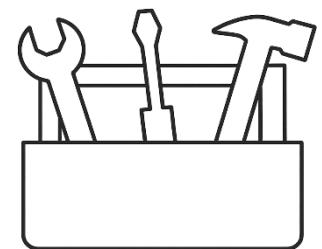
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

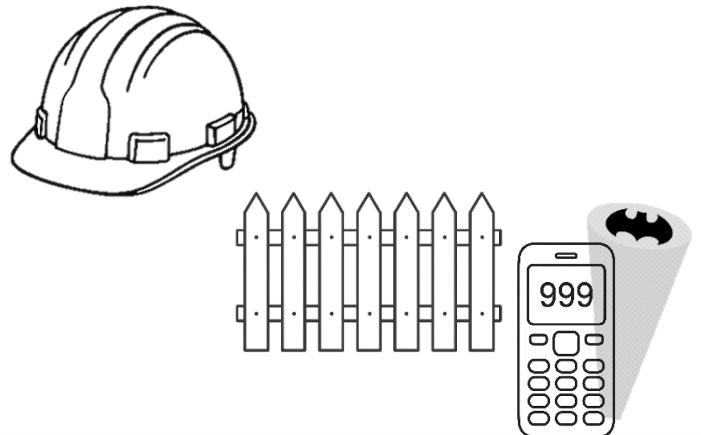
Including open questions



Roadmap

Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,
::handle_contractViolation()`

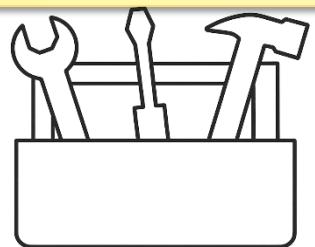
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

Including open questions



Interlude: Program bugs vs run-time errors

Say we're about to perform `vec[index]`

If index is **computed by the program**, a bad index is a **bug in the program code**

```
contract_assert(  
    0 <= index && index < vec.size()  
    && "our code computed a bad index"  
);
```

The caller generally can't handle this or retry or anything sensible (short of a restart, which we'll discuss soon...)

The program is just wrong and needs to be fixed

If index is **read from input**, a bad index is a **user/data error**

```
if (index < 0 || index >= vec.size()) {  
    /* error: "dear user, go fix your data" +  
       report error/exception to the caller */  
}
```

The caller can handle it (back out, retry, whatever is sensible)

► assertions are not for run-time input validation

Item 0: Never use contracts/assertions for program logic

Example: Never use contracts for run-time input validation

Example: Never use contracts for normal program behavior

If you need control flow, use “if”

“The first thing to know about assertions ... is they are redundant.

If you’re writing something that’s not redundant, don’t use an assertion.

This is a basic fact of information theory:

There is no error [here, program bug] detection without redundancy.”

— Lisa Lippincott (CppCon 2024)



Item 1: Avoid side effects in contract predicates

(Corollary of Item 0)

Bad: `contract_assert(++x > 10); // don't do this`

Just like all assertions

C assert: Predicate will be executed zero or one times

C++26: Predicate will be **executed zero or more times**

Good news:

✓ Most of these will be compile-time errors!

Contract predicates treat variables as const — similar to lambda capture



Item 2: Avoid splitting compound conditions For short-circuit eval you need to write the `&&` !

Good: `contract_assert(p && p->foo()); // do this`

Bad: `contract_assert(p);
contract_assert(p->foo()); // don't do this`

With `observe` semantic:

The second test is not protected by the first, and might be UB

Not too surprising: You only get short-circuit evaluation for `&&` when you write `&&`

Good news:

- ✓ Even in `observe`, the first check will be checked and tell you what you did wrong



Item 3: Use throwing violation handlers “judiciously”

Q: Can it be reasonable to install an exception-throwing violation handler?

- A1 (default) Throwing handlers are mainly to **optimize testing contracts themselves**
- A2 (advanced) Yes in a **specific scenario**, but be aware of **tradeoffs**

Scenario	We want to “enforce” contracts in production, but can’t tolerate termination Typical idea: “Failed subcomponent” unwind-and-restart
Beware	Injecting an exception during stack unwinding (or in noexcept fn) causes termination Nearly all code could be transitively called during stack unwinding (or in noexcept fn)

Suggestion Run with **observe** + a violation handler like this (*disclaimer: experimental!*)

```
    ::handle_contractViolation( /*...*/ ) {  
        /// anything else ///  
        if (!uncaughtException()) {  
            throw restartComponentException();  
        }  
    }
```

Prior art: Ada, Eiffel, Java, and other languages throw for assertion violations, but don't have C++'s unwinding-termination



Item 4: Understand whether/how build modes combine

Recall existing issue: “Debug” and “Release” builds aren’t link-compatible

NDEBUG on/off is often a “benign” ODR violation, but...

... `_DEBUG` (which affects ABI) and NDEBUG often vary together in real builds

Existing best practices

Best, no ODR issues: Build everything in the same mode (including NDEBUG)

“Benign” ODR violation: Test NDEBUG (incl. assert) in .cpp files only,
not in inline/template .h code

Contracts offer an improvement, while still in today’s “TUs + linkers” world

Deliberately designed to allow mixing semantics (per-TU) without ABI/link incompat

The current GCC and Clang contracts implementations allow linking TUs built with any configuration (without “benign” ODR violations = better than C assert)

The standard allows platforms to document restrictions on if/when/how semantics mix



Item 4: Understand whether/how build modes combine

Consider this example (P3835)

You can guarantee local semantic per-TU

So lib.cpp can use contract_assert inside .cpp bodies (or ensure inlining) to get guaranteed enforcement...

lib.cpp – **quick_enforce**

```
#include "lib.h"
void g() {
    f(0); // expected to fail
} // quick_enforce guaranteed
// if f(0) is inlined
```

lib.h

```
inline void f(int x) {
    contract_assert(x > 0);
}
```

main.cpp – **ignore**

```
#include "lib.h"
extern void g();
int main() {
    f(1); // expected to pass
    g();
}
```

... and main.cpp gets guaranteed ignore with zero run-time cost

if whole program has two copies of f(), linker picks one

Item 4: Understand whether/how build modes combine

To know whether you can have a program that combines TUs built with different semantics, check your compiler/platform docs

Expectation: You probably can, but you still have to know about inlining/templates

If yes, a translation unit author (library or mainline) can get guaranteed semantics (checking, ignoring) for contract assertions **in their translation unit**:

Code that appears in the .cpp file

Code that appears in a .h file that you ensure is inlined (incl. templates)

That's an important restriction you can't always meet, so be aware



Item 5: Understand C++26's v1.0 limitations

No custom error messages

For `static_assert` similarly, we had to add them later

No contracts on virtual functions or function pointers

Some people really want this urgently

No contract groups/labels – e.g., a group for bounds violations, and a way to always enforce those

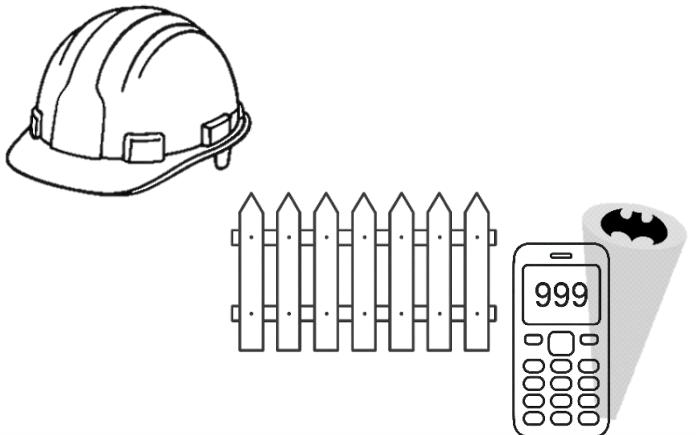
Today: Per-TU semantics today (see Item 4), single global violation handler



Roadmap

Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,
::handle_contractViolation()`

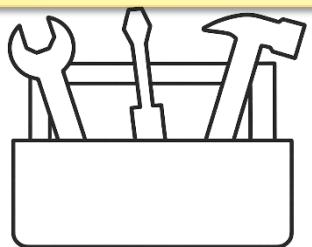
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

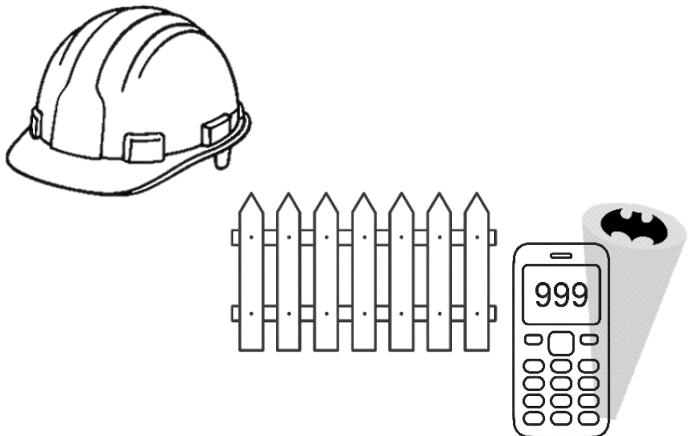
Including open questions



Roadmap

Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,
::handle_contractViolation()`

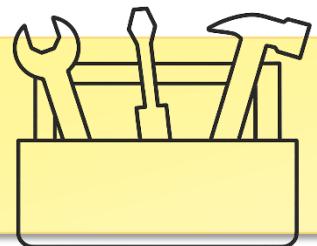
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

Including open questions



Minimal: Addresses key requirements that assert() doesn't

Contracts

`pre / post`

Major advance: Express assertions **on declarations**, lights up analysis

`contract_assert`

Better than a macro... 'nuff said

Semantics

`ignore`

Essential

`observe`

Enables **checking in production**, and testing new predicates

`enforce`

Essential, for testing

`quick_enforce`

Enables **enforcing in production** with absolute min. size & perf impact

Violation handler

`default`

Typically prints source file/line/col, message, and similar

and replaceable

Enables hooking into your **existing logging/alerting** mechanisms

... at link time

Avoids security issues of run-time replacement



Belt and suspenders...



Shift left

Bring more bug detection
to test time, pre-merge

Safety net

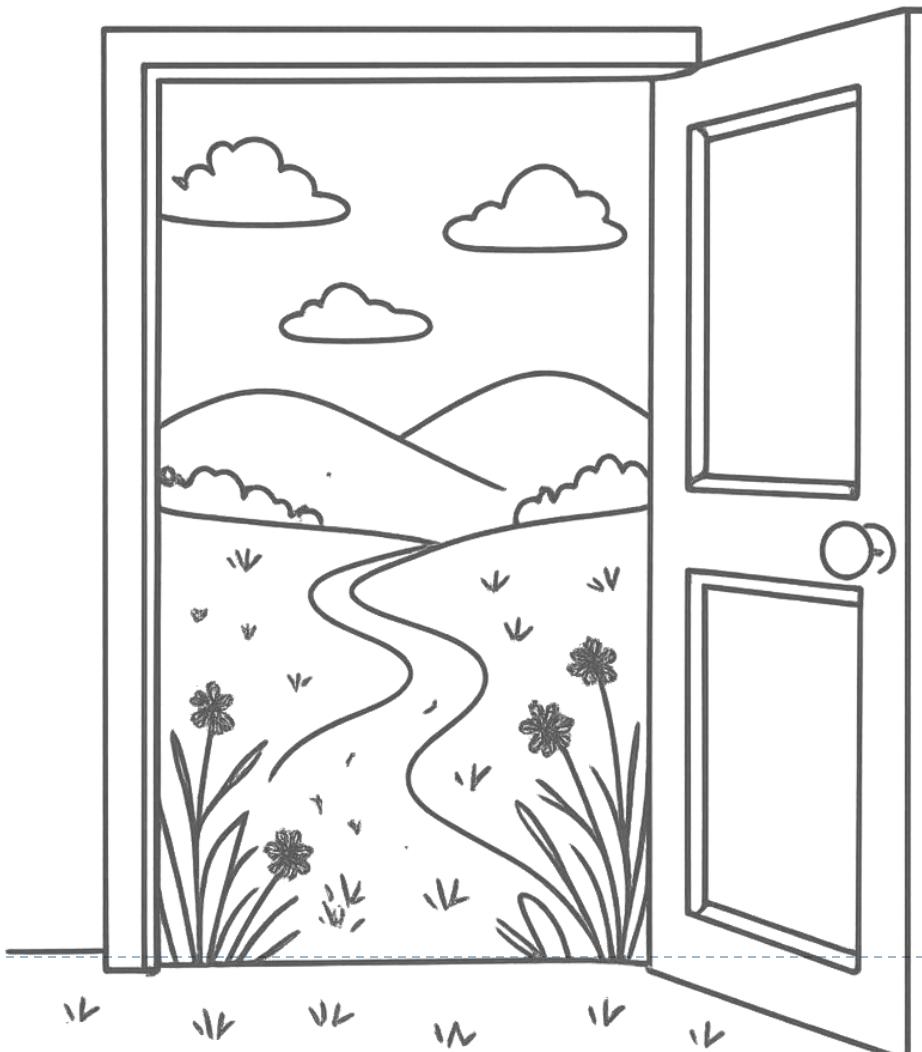
Also enable bug detection
always, including production

Viable: Works, and doesn't close doors

Guiding principle: Don't close doors

Open doors include:

- Custom error messages
- Virtual functions
- Function pointers
- Labels/groups



Example: `noexcept_enforce/observe` (in GCC dev branch)

As a future extension example:

A semantic that makes throwing violation handlers nonthrowing

"The breakthrough realization, the epiphany, and the stroke of genius there is that instead of trying to cater to two different crowds by making the violation handler `noexcept` or not, we can serve two crowds with the same violation handler.

If it throws, the authors of components that don't want exceptions to escape get what they want. The authors of components that want exceptions to escape get what they want. The different decisions can be mixed in the same binary, and one violation handler does what both audiences want.

Author A compiles with `noexcept_enforce`, author B compiles with `enforce`. The same violation handler works for both. No linker magic, no tricks, just different compilation of the violation handler calls. Done in the wrapper functions we already have."

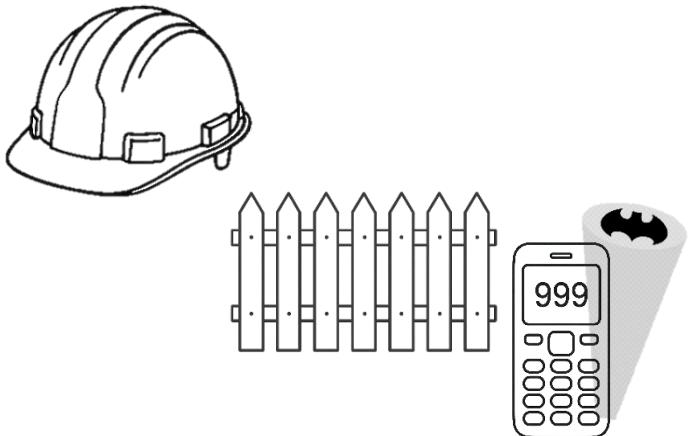
— Ville Voutilainen



Roadmap

Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,
::handle_contractViolation()`

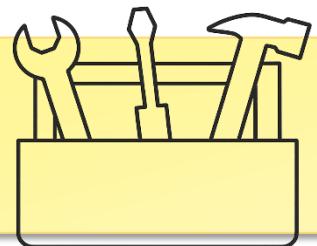
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

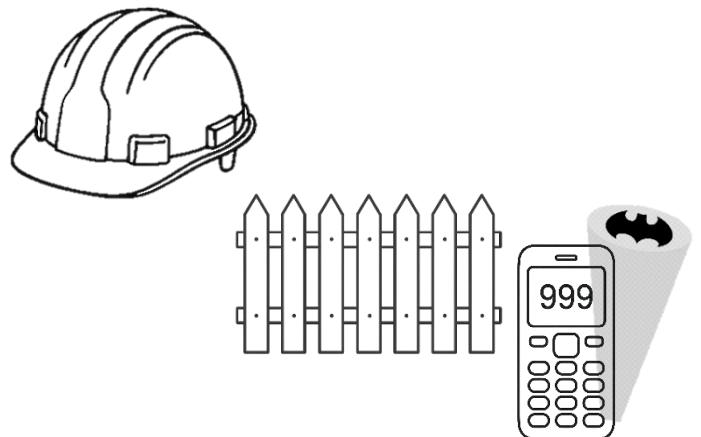
Including open questions



Roadmap

Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,
::handle_contractViolation()`

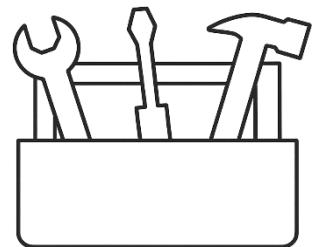
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

Including open questions



Some quick FAQs

Q Do C++26 contracts have lots of implementation-defined behavior?

A They're things we usually make implementation-defined
(e.g., the standard doesn't specify /Compiler --switches, or build modes)

Q Are C++26 contracts hard to implement?

A GCC and Clang are complete, in code review, targeting next versions

Q Is it true that C++26 contracts don't help static analysis?

A  Catching Bugs Early: Validating C++ Contracts with Static Analysis

Thursday @ 15:15 MDT

Mike Fairhurst, GitHub CodeQL (Microsoft)

Peter Martin, Product Security (Bloomberg)



More quick FAQs

Q Is there a portable default behavior?

A C++26 says... ‘Enforce by default, print contract_violation information by default’ + “rate-limited for potentially repeated violations of observed contract assertions”

Current GCC message:

*contract violation in function T ret(int) [with T = double] at mypath/myfilename:76: a > 1
[assertion_kind: pre, semantic: observe, mode: predicate_false, terminating: no]*

Q Do contracts have a flaw where writing UB in code protected by a contract can silently eliminate another contract check by time-travel optimization?

A Not anymore, fixed for C++26

In a nutshell: “observable checkpoint,” treat a contract predicate like I/O



S'mores quick FAQs

Q What if I try to write a `pre/post` on a virtual? ... a function pointer?

A Compile-time error in C++26; likely future extension post-C++26
C++11: `->auto` for lambdas ... C++14: `->auto` for all functions

Q Isn't making objects implicitly const in predicate checks weird?

A So far: Found actual bugs, and no realistic problem examples
Not novel in C++: Similar to implicit const in C++ lambda capture

Q Why aren't C++26 contracts a drop-in replacement for [my assertion library]?

A1 A language feature beats macros, and won't always work exactly like a macro
A2 Some examples break because they are *already const-incorrect* (see previous)



Open questions

1a A predicate that throws is converted to a contract violation

1b (related) Implementation and deployment experience on non-Itanium ABIs

Q: Are those good semantics?

Q: Is the run-time cost acceptable? Need to measure.

(The hopeful conjecture: If most predicates are trivially non-throwing, we can elide most notional injected try/catches?)

Suggested extension: Evaluation semantics that don't translate predicate exceptions into contract violations

2 We can't yet write functions that don't exhibit undefined behavior when preconditions are violated

Suggested extension: A way to express a precondition check that can never be ignored, no matter which compiler flags are used (example: `pre<quick_enforce> (index < size())`)

These suggested extensions can likely be added post-26, e.g., using an additional semantic or a “labels” grouping feature

More open questions

3 Should we hold v1.0 from the standard till we can add { groups/labels, virtuals } ?

Groups: to allow grouping of contracts with a similar purpose, or tied to a component

Virtuals: with substitutability?

Note the tension:

“v1.0 must be strictly minimal” +
“but add one feature I need”

Clang is providing groups as a vendor extension →

Examples of hierarchical matching:

```
// Group hierarchy: mylib -> mylib.debug -> mylib.debug.verbose
contract_assert [[clang::contract_group("mylib")]] (basic_check());
contract_assert [[clang::contract_group("mylib.debug")]] (debug_check());
contract_assert [[clang::contract_group("mylib.debug.verbose")]] (detailed_check());
contract_assert [[clang::contract_group("mylib.performance")]] (perf_check());
```

With these compiler flags:

```
# Set base mylib group to observe
-fcontract-group-evaluation-semantic=mylib=observe
# Override debug subgroup to enforce
-fcontract-group-evaluation-semantic=mylib.debug=enforce
# Override specific verbose group to ignore
-fcontract-group-evaluation-semantic=mylib.debug.verbose=ignore
```



The evaluation semantics would be:

- mylib contracts: **observe** (exact match)
- mylib.debug contracts: **enforce** (exact match, overrides parent)
- mylib.debug.verbose contracts: **ignore** (exact match, overrides parent)
- mylib.performance contracts: **observe** (inherits from parent mylib)

Roadmap

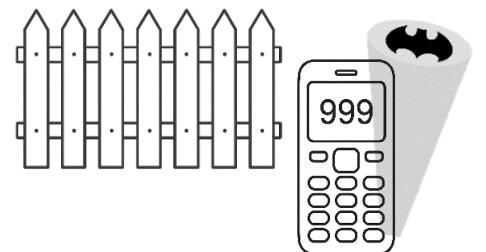
Motivation

Why defensive programming?



Overview

`pre, post, contract_assert,
::handle_contractViolation()`



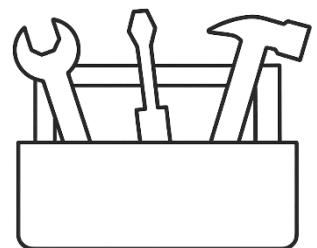
Best practices

“Effective C++26 Contracts”



Why ____?

Viable, ~minimal
What each part enables



FAQs

Including open questions



+ 25

The Joy of C++26 Contracts (and Some Mythconceptions)

HERB SUTTER
≡ CITADEL | Securities



20
25 | September 13 - 19