

+ 25

# Where'd All That RAM Go?

Measuring Objects With Ownership-Aware  
Memory Profiling in C++

ALECTO IRENE PEREZ



20  
25



# Introduction

- My name is Alecto
- I'm a Senior Software Engineer at Vola Dynamics, The Volatility Company
- I wrote a memory profiler called `mem_profile`
- This is a personal project
- [https://github.com/codeinred/mem\\_profile](https://github.com/codeinred/mem_profile)

# Overview

## This presentation will cover:


- Why we care about memory profiling (+background)
- What `mem_profile` supports (direct object-level memory statistics)
- How all of this works

```
> uv run python -m mp_reader type_stats malloc_stats.json --type 'my_object'

// Totals for my_object sizeof=72 bytes
// └─ 398 bytes across 6 allocs and 1 instance
struct my_object
{
    bytes      a; // 10 bytes across 1 allocs : std::vector<std::byte>
    bytes      b; // 100 bytes across 1 allocs : std::vector<std::byte>
    string_map m; // 288 bytes across 4 allocs : std::map<std::string, std::string>
};
```

# Why do we care about memory profiling?

- Memory on a system is a finite resource
- Swap is sometimes available, but it is orders of magnitude slower than RAM
- Out-of-memory is often unrecoverable



# Why do we care about memory profiling?

- You can buy more RAM, you can get bigger servers...
- These things are *expensive*.
- If you ship to clients, they may not be able to get more RAM!
- **You need to be able to measure what you optimize.**

# Measuring Memory Usage - An Introduction

# Measuring Memory Usage

The simplest way to measure memory usage is with built-in utilities such as `time`.

Consider this program:

```
#include <cstdint>
#include <map>

int main() {
    std::map<int64_t, int64_t> m;
    for (size_t i = 0; i < 1'000'000; i++) {
        m[i] = i * i;
    }
}
```

# Measuring Memory Usage

`time` can measure peak memory footprint with either `-l` or `-v`:

```
> \time -l ./a.out
0.06 real      0.06 user      0.00 sys
 49332224 maximum resident set size
         0 average shared memory size
         0 average unshared data size
         0 average unshared stack size
    3210 page reclaims
         1 page faults
         0 swaps
         0 block input operations
         0 block output operations
         0 messages sent
         0 messages received
         0 signals received
         0 voluntary context switches
         3 involuntary context switches
816149941 instructions retired
237223462 cycles elapsed
 49217896 peak memory footprint
```



# Measuring Memory Usage - **time**

## Pros:

- Quick
- Easy
- You get one number: peak memory footprint

## Cons:

- No visibility into what was using memory
- You only get one number: peak memory footprint

# Measuring Memory Usage - `valgrind massif`

- `massif` can measure memory usage over time
- You get more granularity with regards to when peak memory usage occurs
- Can detect leaks!



## 9. Massif: a heap profiler

### Table of Contents

#### [9.1. Overview](#)

#### [9.2. Using Massif and ms\\_print](#)

##### [9.2.1. An Example Program](#)

##### [9.2.2. Running Massif](#)


##### [9.2.3. Running ms\\_print](#)

##### [9.2.4. The Output Preamble](#)

##### [9.2.5. The Output Graph](#)

##### [9.2.6. The Snapshot Details](#)

##### [9.2.7. The Summary](#)

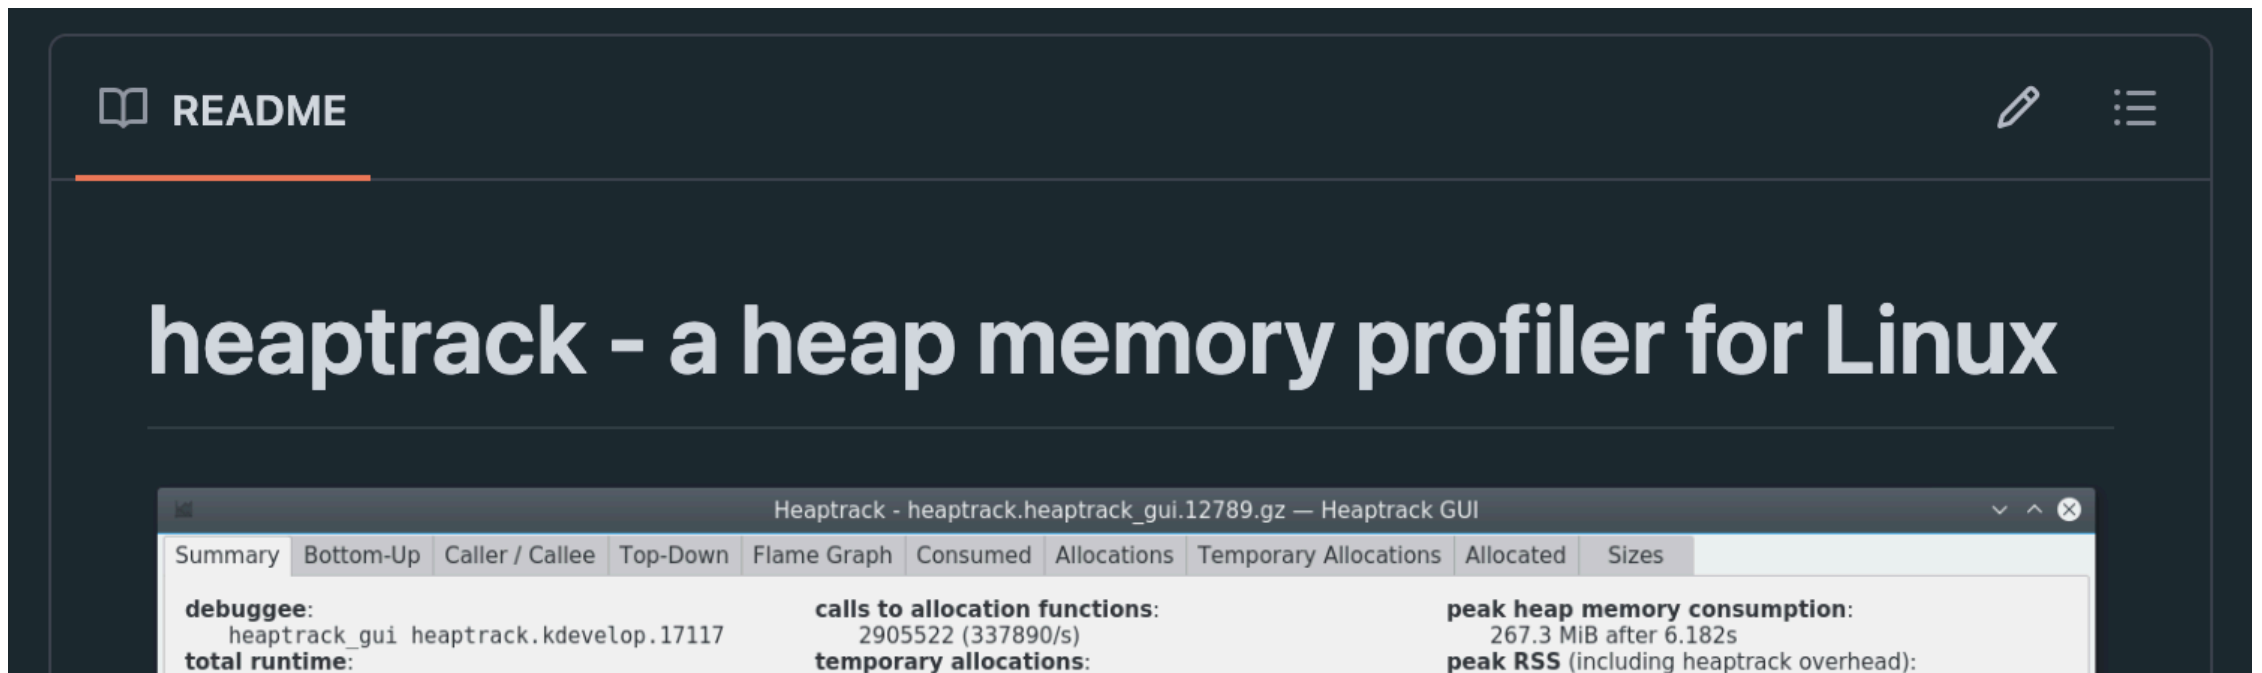



Number of snapshots: 63

Detailed snapshots: [3, 4, 10, 11, 15, 16, 29, 33, 34, 36, 39, 41,  
42, 43, 44, 49, 50, 51, 53, 55, 56, 57 (peak)]

# Measuring Memory Usage - KDE **heaptrack**

- Can see where allocations occur
- Can get flamegraph of memory allocations
- Can identify temporary allocations





**So why another profiler?**

# So why another profiler?

We write code in terms of data structures, containers, and objects. I want to be able to measure memory in terms of these things!

- Which objects are consuming memory?
- Why isn't memory released?
- What members/subobjects are using that memory?

`mem_profile` can answer these questions!

# We resolved limitations of similar efforts

mem\_profile's RAll-based approach resolves many limitations faced by prior attempts at profiling via object introspection.

- `mem_profile` can profile user-defined containers
- `mem_profile` handles template specialization of containers
- `mem_profile` handles virtual inheritance
- `mem_profile` handles c-style unions, AND ALSO subobjects allocated with placement-new.
- `mem_profile` resolves shared ownership
- if a pointer is freed in the destructor, `mem_profile` handles it



## Can Your Profiler Do This?

**mem\_profile embeds type and layout information in a binary to provide comprehensive statistics, even for release builds.**

```

6  using bytes      = std::vector<std::byte>;
7  using string_map = std::map<std::string, std::string>;
8
9  struct my_object {
10     bytes      a;
11     bytes      b;
12     string_map m;
13 };
14
15 int main() {
16     my_object m{
17         bytes(10),
18         bytes(100),
19         string_map{
20             {"key1", "Hello, world"},
21             {"key2", "The quick brown fox jumps over the lazy dogs"},
22             {"key3", "Have a great CppCon!"},
23         },
24     };
25 }

```

mem\_profile/stats\_reader on  main at v0.1.0 via  v3.13.2

```
> uv run python -m mp_reader type_stats test_files/macos/objects.json \  
    --type 'my_object'
```

```
// Totals for my_object sizeof=72 bytes  
// └─ 398 bytes across 6 allocs and 1 instance
```

```
struct my_object  
{  
    bytes      a; // 10 bytes across 1 allocs : std::vector<std::byte>  
    bytes      b; // 100 bytes across 1 allocs : std::vector<std::byte>  
    string_map m; // 288 bytes across 4 allocs : std::map<std::string, std::string>  
};
```

mem\_profile/stats\_reader on  main at v0.1.0 via  v3.13.2

```
>
```

mem\_profile/stats\_reader on  main at v0.1.0 via  v3.13.2


```
> uv run python -m mp_reader type_stats test_files/macos/objects.json \  
    --type 'my_object' --show-offsets
```

```
// Totals for my_object sizeof=72 bytes  
// └─ 398 bytes across 6 allocs and 1 instance
```

```
struct my_object  
{  
    bytes      a; //  0.. 24  10 bytes across 1 allocs : std::vector<std::byte>  
    bytes      b; // 24.. 48 100 bytes across 1 allocs : std::vector<std::byte>  
    string_map m; // 48.. 72 288 bytes across 4 allocs : std::map<std::string, std::string>  
};
```

mem\_profile/stats\_reader on  main at v0.1.0 via  v3.13.2

```
> |
```


examples > src >  objects\_in\_lambda.cpp > ...

```
1  #include <cstdint>
2  #include <vector>
3
4  using bytes    = std::vector<std::byte>;
5  using floats   = std::vector<float>;
6  using doubles  = std::vector<double>;
7
8  int main() {
9      auto myLambda = [a = bytes(10),    //
10                      b = floats(100),  //
11                      c = doubles(1000)]() {};
12 }
13
```

```
mem_profile/stats_reader on 🐱🐹 main at v0.1.0 via 🐍 v3.13.2
> uv run python -m mp_reader type_stats test_files/macos/objects_in_lambda.json --type 'lambda'

// Totals for (lambda at /opt/pages/mem_profile/examples/src/objects_in_lambda.cpp:9:21) sizeof=72 bytes
// └─ 8,410 bytes across 3 allocs and 1 instance
struct (lambda at /opt/pages/mem_profile/examples/src/objects_in_lambda.cpp:9:21)
{
    bytes    (unnamed); // 10 bytes across 1 allocs : std::vector<std::byte>
    floats   (unnamed); // 400 bytes across 1 allocs : std::vector<float>
    doubles  (unnamed); // 8,000 bytes across 1 allocs : std::vector<double>
};

mem_profile/stats_reader on 🐱🐹 main at v0.1.0 via 🐍 v3.13.2
>
```

examples > src >  objects\_in\_array.cpp > ...

```
1  #include <cstdint>
2  #include <vector>
3
4  using bytes = std::vector<std::byte>;
5
6  struct array3 {
7      bytes elems[3];
8  };
9
10 int main() {
11     array3 values{
12         bytes(10),
13         bytes(100),
14         bytes(1000),
15     };
16 }
17
```

```
mem_profile/stats_reader on main +26 -10 !2 at v0.1.0 via v3.13.2
> uv run python -m mp_reader type_stats test_files/macos/objects_in_array.json --type 'array3'
```

```
// Totals for array3 sizeof=72 bytes
// └─ 1,110 bytes across 3 allocs and 1 instance
```

```
struct array3
{
    bytes[3] elems; // [0] 10 bytes across 1 allocs : std::vector<std::byte>
                  // [1] 100 bytes across 1 allocs : std::vector<std::byte>
                  // [2] 1,000 bytes across 1 allocs : std::vector<std::byte>
};
```

```
mem_profile/stats_reader on main +26 -10 !2 at v0.1.0 via v3.13.2
> |
```



```

8  template <class T, size_t Cap>
9  struct svector {
10     size_t count = 0;
11     alignas(T) std::byte buffer[sizeof(T) * Cap];
12
13     void push_back(T&& value) { ...
14
15     }
16     T* data() noexcept { return reinterpret_cast<T*>(buffer); }
17     ~svector() {
18         for (size_t i = 0; i < count; i++) data()[i].~T();
19     }
20 };
21
22
23 int main() {
24     svector<bytes, 10> v;
25     v.push_back(bytes(10)); // [0]
26     v.push_back(bytes(100)); // [1]
27     v.push_back(bytes(1000)); // [2]
28     v.push_back(bytes()); // [3] - empty: no allocation will be recorded here
29     v.push_back(bytes(12345)); // [4]
30 }

```

mem\_profile/stats\_reader on  main at v0.1.0 via  v3.13.2

> uv run python -m mp\_reader type\_stats test\_files/macos/objects\_in\_svector.json --type 'svector'

```
// Totals for svector<std::vector<std::byte>, 10> sizeof=248 bytes
```

```
// └─ 13,455 bytes across 4 allocs and 1 instance
```

```
struct svector<std::vector<std::byte>, 10>
```

```
{
```

```
    size_t        count;
```


```
    std::byte[240] buffer; // [0] 10 bytes across 1 allocs : std::vector<std::byte>
```

```
                        // [1] 100 bytes across 1 allocs : std::vector<std::byte>
```

```
                        // [2] 1,000 bytes across 1 allocs : std::vector<std::byte>
```

```
                        // [4] 12,345 bytes across 1 allocs : std::vector<std::byte>
```

```
};
```

mem\_profile/stats\_reader on  main +4 -9 !1 at v0.1.0 via  v3.13.2

> |

```
1  #include <ankerl/svector.h>
2  #include <cstdint>
3  #include <vector>
4
5  using bytes = std::vector<std::byte>;
6
7
8  ∨ int main() {
9      ankerl::svector<bytes, 16> sv;
10     sv.push_back(bytes(10));
11     sv.push_back(bytes(100));
12     sv.push_back(bytes(1000));
13 }
14
```

```

mem_profile/stats_reader on 🐱🔑 main +26 -10 !2 at v0.1.0 via 🐍 v3.13.2
> uv run python -m mp_reader type_stats test_files/macos/objects_in_svector_martinus.json \
  --type 'svector' \
  --replace 'detail::size_of_svector<vector<byte, allocator<byte>>>(16UL)' \
  --replace-with '...'

// Totals for ankerl::svector<std::vector<std::byte>, 16> sizeof=392 bytes
// └─ 1,110 bytes across 3 allocs and 1 instance
struct ankerl::svector<std::vector<std::byte>, 16>
{
    std::array<uint8_t, ...> m_data; // +8 10 bytes across 1 allocs : std::vector<std::byte>
                                   // +32 100 bytes across 1 allocs : std::vector<std::byte>
                                   // +56 1,000 bytes across 1 allocs : std::vector<std::byte>
};

mem_profile/stats_reader on 🐱🔑 main +26 -10 !2 at v0.1.0 via 🐍 v3.13.2
> |

```

```
1  #include <cstdint>
2  #include <vector>
3
4  using bytes = std::vector<std::byte>;
5
6
7  int main() {
8      std::vector values{
9          bytes(10),
10         bytes(100),
11         bytes(1000),
12     };
13 }
14
```

```

mem_profile/stats_reader on 🐱 main at v0.1.0 via 🐍 v3.13.2
> uv run python -m mp_reader type_stats test_files/macos/objects_in_vector.json --type 'std::vector' \
  --clean-members

// Totals for std::vector<std::byte> sizeof=24 bytes
// └─ 2,220 bytes across 6 allocs and 6 instances
struct std::vector<std::byte>
{
    // (6 non-owning fields cleaned)

    ~std::vector<std::byte>();
    // directly owned (or unannotated):
    // └─ 2,220 bytes across 6 allocs
};

// Totals for std::vector<std::vector<std::byte>> sizeof=24 bytes
// └─ 1,182 bytes across 4 allocs and 1 instance
struct std::vector<std::vector<std::byte>>
{
    // (6 non-owning fields cleaned)

    ~std::vector<std::vector<std::byte>>();
    // directly owned (or unannotated):
    // └─ 72 bytes across 1 allocs
    // children on heap:
    // └─ 1,110 bytes across 3 allocs : std::vector<std::byte>
};

```

**How big is an object?**

**Measuring Objects and the Perils of Shared Ownership**

# Measuring Objects - Simple Example

```
#include <memory>

struct A {
    std::shared_ptr<char[]> data;

    static A alloc( size_t bytes ) {
        return A{std::shared_ptr<char[]>(new char[bytes])};
    }
};

int main() {
    // How big is this object?
    A a = A::alloc(1'000'000);
}
```



# Measuring objects - Simple Example

```
int main() {  
    // How big is this object?  
    A a = A::alloc(1'000'000);  
}
```

- `A` is 16 bytes (assuming 8-byte pointers)
- `a.data` holds 1,000,000 bytes, allocated on the heap

Therefore:

- `a` takes up 1,000,016 bytes total (1,000,000 on the heap)

# Measuring Objects - Simple Example

Suppose we had a magic function, `measure_heap(object)`, that tells you how much memory is reachable from an object:

```
int main() {  
    // How big is this object?  
    A a = A::alloc(1'000'000);  
  
    fmt::println("total size: {}", measure_heap(a));  
}
```

This prints `total size: 1000000`, or 1,000,000 bytes.

**Unfortunately, not everything is simple.**

# Measuring Objects -

*Sometimes, objects have shared custody.*

Q: Is this correct? (output: `total bytes = 2000000` )

```
int main() {  
    // How big is this object?  
    A a = A::alloc(1'000'000);  
    A b = A::alloc(1'000'000);  
  
    size_t total_bytes = measure_heap(a) + measure_heap(b);  
  
    fmt::println("total bytes = {}", total_bytes);  
}
```

# Measuring Objects -

*Sometimes, objects have shared custody.*

Q: Is this correct? (output: `total bytes = 2000000` )

```
int main() {  
    // How big is this object?  
    A a = A::alloc(1'000'000);  
    A b = A::alloc(1'000'000);  
  
    size_t total_bytes = measure_heap(a) + measure_heap(b);  
  
    fmt::println("total bytes = {}", total_bytes);  
}
```

Yes, this is fine. `a` and `b` don't share any memory.

# Measuring Objects -

*Sometimes, objects have shared custody.*

Q: Is this correct? (output: `total bytes = 2000000` )

```
int main() {  
    // How big is this object?  
    A a = A::alloc(1'000'000);  
    A b = a;  
  
    size_t total_bytes = measure_heap(a) + measure_heap(b);  
  
    fmt::println("total bytes = {}", total_bytes);  
}
```

# Measuring Objects - 🌟

*Sometimes, objects have shared custody.*

Q: Is this correct? (output: `total bytes = 2000000` )

```
int main() {  
    // How big is this object?  
    A a = A::alloc(1'000'000);  
    A b = a;  
  
    size_t total_bytes = measure_heap(a) + measure_heap(b);  
  
    fmt::println("total bytes = {}", total_bytes);  
}
```

NO! This is not correct! Even with a magic `measure_heap` function, these objects share memory!

# Measuring Objects - How do we fix this?



# Measuring Objects - How do we fix this?

Let's take some simple code, and desugar it:

```
A a = A::alloc(1'000'000); // use count: 1
A b = a;                  // use count: 1 -> 2

// ... (stuff happens) ...

b.~A(); // b's lifetime ends (IMPLICIT)
        // use count: 2 -> 1
a.~A(); // a's lifetime ends (IMPLICIT)
        // use count: 1 -> 0 (deallocation occurs)
```

- `b` is destroyed first, and performs no deallocations
- `a` is destroyed last, and `a`'s destructor cleans up the memory.

# Measuring Objects - How do we fix this?

- Ownership is shared, but deallocation is unique.
- In a leak-free program, any block of memory should be deallocated precisely once.
- In our example, *only* `a`'s destructor frees memory
- Therefore: **attribute the memory to `a` !**

```
int main() {  
    // How big is this object?  
    A a = A::alloc(1'000'000);  
    A b = a;  
  
    // ...  
}
```

# Measuring Objects - A Destructor-Based Approach

## These are the rules:

- We attribute memory to the objects that actually deallocate that memory. `a` performed the deallocation, so the memory held by the shared pointer is deallocated to `a`
- Any memory deallocated in a destructor call was owned by the object being destroyed
- If no deallocations occur on destruction, there's nothing to attribute
- Because each deallocation is unique, this *also* correctly assigns an owner, *even when the memory is shared*

# Is This Correct?

## I claim that it is!

- This approach is *generic* - any object that owns memory *must free that memory*
- This approach ignores non-owning types - we don't have to care about `std::span` or `string_view`, because these don't deallocate!
- This approach handles both unique ownership and shared ownership
- Memory is *always* attributed to the *last remaining object* which owns a piece of memory.
- This also allows us to discover why memory wasn't freed sooner!

# The Profiler

**Let's Measure Some Objects!**

# Measuring Objects - Test Code

```
#include <memory> // shared.cpp, compiled as `build/shared`
#include <cstdio>

struct A {
    std::shared_ptr<char[]> data;

    static A alloc( size_t bytes ) {
        return A{std::shared_ptr<char[]>(new char[bytes])};
    }
};

int main() {
    A a = A::alloc(1'000'000);
    A b = a;
    printf("a.this = %p\n", &a);
    printf("b.this = %p\n", &b);
}
```

# Measuring Objects - Compiling

You can build your project with `mem_profile` by adding these lines near the top of your `CMakeLists.txt`.

```
find_package(mem_profile REQUIRED)
link_libraries(mp::mp_build_with_plugin)
```

Use `llvm clang` when compiling:

```
env CC=clang CXX=clang++ cmake -B build -G Ninja
cmake --build build
```

*Note: on macos, this should be `/opt/homebrew/opt/llvm/bin/clang++`*

# Measuring Objects - Compiling

- `mp_build_with_plugin` is an INTERFACE target
- Linking it adds compile flags to inject the plugin when compiling with clang.
- Use `link_libraries` to ensures the plugin is linked against all targets in the binary.
- The plugin injects information into destructor calls to enable introspection and ownership tracing



# Compile Output with Printing Enabled

```
Rewrote my_object::~~my_object @ /opt/pages/mem_profile/examples/src/objects.cpp:9:8
inline ~my_object() noexcept __noinline__ {
    static constexpr const char *__MP_FIELD_NAMES[3] = {"a", "b", "m"};
    static constexpr const char *__MP_FIELD_TYPES[3] = {"bytes", "bytes", "string_map"};
    static constexpr unsigned long __MP_FIELD_SIZES[3] = {24UL, 24UL, 24UL};
    static constexpr unsigned long __MP_FIELD_OFFSETS[3] = {0UL, 24UL, 48UL};
    static constexpr const char *__MP_BASE_TYPES[0] = {};
    static constexpr unsigned long __MP_BASE_SIZES[0] = {};
    static constexpr unsigned long __MP_BASE_OFFSETS[0] = {};
    static constexpr _mp_type_data __MP_TYPE_DATA = {sizeof(my_object), "my_object", 0UL,
BASE_SIZES, __MP_BASE_OFFSETS};
    save_state(this, __builtin_alloca(40UL), __MP_TYPE_DATA);
}
[2/2] Linking CXX executable objects
```

You can use `MEM_PROFILE_PRINT_BODY=1` to view instrumented destructor calls.

# Measuring Objects - Running the Profiler

Let's profile our executable. We can use either `LD_PRELOAD` (on linux) or `DYLD_INSERT_LIBRARIES` to inject the profiler:

```
# linux
env LD_PRELOAD=/usr/local/lib/libmp_runtime.so build/shared

# macos
env DYLD_INSERT_LIBRARIES=/usr/local/lib/libmp_runtime.dylib build/shared
```

The profiler will be initialized on load. `libmp_runtime` provides definitions for `malloc`, `free`, `new`, `delete`, etc that record allocations and deallocations as they occur.

# Measuring Objects - Output

When we run our executable with the profiler, it generates a file `malloc_stats.json`. This contains all of the information necessary for analysis.

```
$ env LD_PRELOAD=/usr/local/lib/libmp_runtime.so build/shared  
a.this = 0x16d6fe980  
b.this = 0x16d6fe970  
  
$ ll malloc_stats.json  
-rw-r--r--@ 1 alecto  staff   3.1K Sep 11 13:44 malloc_stats.json
```

You can configure the output filename by setting `MEM_PROFILE_OUT=<filename>`.

```
$ env MEM_PROFILE_OUT=output.json LD_PRELOAD=... /path/to/executable
```

# Measuring Objects - Output

If we dump an event trace for the file, we can see that there are two free events:

- a 1MB allocation for the `char[]` array held by the shared pointer, and
- a 32-byte allocation for the metadata the shared pointer uses to track the object.

Note that the `this` pointer recorded in `malloc_stats` matches the address displayed by the executable.

```
#           module      subcommand  input file
#           v           v           v
$ uv run python -m mp_reader dump_events malloc_stats.json
```

```

type: FREE
addr: 128008000
size: 1000000
— /opt/homebrew/Cellar/llvm/20.1.7/bin/./include/c++/v1/__memory/shared_count.h:121:7 (inlined)
  std::__1::__shared_weak_count::__release_shared[abi:ne200100]()
— /opt/homebrew/Cellar/llvm/20.1.7/bin/./include/c++/v1/__memory/shared_ptr.h:558:17
  std::__1::shared_ptr<char []>::~~shared_ptr[abi:ne200100]()
  OBJECT this=0x16d6fe980 offset=0 size=16 type=std::shared_ptr<char[]>
— /opt/pages/cppcon/example/shared.cpp:0:8
  A::~~A()
  OBJECT this=0x16d6fe980 size=16 type=A
— /opt/pages/cppcon/example/shared.cpp:17:1
  main

type: FREE
addr: 125605f50
size: 32
— /opt/homebrew/Cellar/llvm/20.1.7/bin/./include/c++/v1/__memory/shared_count.h:0:7 (inlined)
  std::__1::__shared_weak_count::__release_shared[abi:ne200100]()
— /opt/homebrew/Cellar/llvm/20.1.7/bin/./include/c++/v1/__memory/shared_ptr.h:558:17
  std::__1::shared_ptr<char []>::~~shared_ptr[abi:ne200100]()
  OBJECT this=0x16d6fe980 offset=0 size=16 type=std::shared_ptr<char[]>
— /opt/pages/cppcon/example/shared.cpp:0:8
  A::~~A()
  OBJECT this=0x16d6fe980 size=16 type=A
— /opt/pages/cppcon/example/shared.cpp:17:1
  main

```

## Output (for reference)

```
$ env LD_PRELOAD=/usr/local/lib/libbmp_runtime.so build/shared  
a.this = 0x16d6fe980  
b.this = 0x16d6fe970
```

# Measuring Objects - The Trace

- Each entry in the trace corresponds to a stack frame.
- Some entries correspond to destructor calls
- These destructor calls are annotated with information about the object being destroyed, namely,
  - the `this` pointer of the object,
  - it's size,
  - and it's typename.
- Additional type information is also recorded and stored in `malloc_stats.json`, although this is not displayed in the trace.

# Measuring Objects - Statistics

We can use the `mp_reader stats` command to generate stats for the program. These statistics are on a per-type basis.

```
> uv run python -m mp_reader stats /opt/3rd/FTXUI/malloc_stats.json
```

Allocation Statistics by Type:

140 objects	1,438,080 bytes	<code>std::vector&lt;ftxui::Pixel&gt;</code>
10 objects	1,338,720 bytes	<code>ftxui::Screen</code>
10 objects	1,338,480 bytes	<code>std::vector&lt;std::vector&lt;ftxui::Pixel&gt;&gt;</code>
10 objects	1,338,480 bytes	<code>ftxui::Image</code>
600 objects	708,000 bytes	<code>ftxui::flexbox_helper::Global</code>
600 objects	536,640 bytes	<code>std::vector&lt;ftxui::flexbox_helper::Block&gt;</code>
1,930 objects	500,400 bytes	<code>std::shared_ptr&lt;ftxui::Node&gt;</code>
290 objects	498,880 bytes	<code>std::vector&lt;std::shared_ptr&lt;ftxui::Node&gt;&gt;</code>
40 objects	498,880 bytes	<code>ftxui::(anonymous namespace)::Border</code>
100 objects	495,840 bytes	<code>ftxui::(anonymous namespace)::Flexbox</code>
110 objects	442,200 bytes	<code>ftxui::(anonymous namespace)::VBox</code>
2,950 objects	307,200 bytes	<code>std::vector&lt;ftxui::box_helper::Element&gt;</code>



# Measuring Objects - Peak Usage

By default, mem\_profile shows accumulated stats over the full runtime of the program, but you can limit statistics to objects alive at peak usage.

```
> uv run python -m mp_reader stats /opt/3rd/FTXUI/malloc_stats.json --filter-peak --count 6
Allocation Statistics by Type:
```

14 objects	143,808 bytes	std::vector<ftxui::Pixel>
1 object	133,848 bytes	std::vector<std::vector<ftxui::Pixel>>
1 object	133,848 bytes	ftxui::Image
1 object	133,848 bytes	ftxui::Screen
193 objects	38,240 bytes	std::shared_ptr<ftxui::Node>
29 objects	38,088 bytes	std::vector<std::shared_ptr<ftxui::Node>>

...

10 entries filtered

9,640 allocs	21 bytes	<untyped>
--------------	----------	-----------

302 objects	194,216 bytes	<total>
-------------	---------------	---------

# Measuring Objects - Breakdown by Member

Stats by type are nice, but we can do a lot more!

- We know the `this` pointer of every object
- We know the size of every object
- *this allows us to compute the offset of subobjects within a greater whole!*

Additionally:

- We already have a compiler plugin
- The compiler knows every field in an object, it's name, it's type, and it's offset
- We can use this to embed any and all desired type info, directly in the binary!

```
> uv run python -m mp_reader type_stats /opt/3rd/FTXUI/malloc_stats.json \
--type '^ftxui' --filter-mode=REGEX --clean-members
```

```
// Totals for ftxui::Screen sizeof=128 bytes
// └─ 1,338,720 bytes across 150 allocs and 10 instances
```

```
struct ftxui::Screen
{
    : Image // 1,338,480 bytes across 140 allocs : ftxui::Image
{
    // (2 non-owning fields cleaned)
    std::vector<std::string> hyperlinks_; // 240 bytes across 10 allocs : std::vector<std::string>
};
```

```
// Totals for ftxui::Image sizeof=56 bytes
// └─ 1,338,480 bytes across 140 allocs and 10 instances
```

```
struct ftxui::Image
{
    // (3 non-owning fields cleaned)
    std::vector<std::vector<Pixel>> pixels_; // 1,338,480 bytes across 140 allocs : std::vector<std::vector<ftxui::P
};
```

```
// Totals for ftxui::flexbox_helper::Global sizeof=88 bytes
// └─ 708,000 bytes across 2100 allocs and 600 instances
```

```
struct ftxui::flexbox_helper::Global
{
    // (3 non-owning fields cleaned)
    std::vector<Block> blocks; // 536,640 bytes across 600 allocs : std::vector<ftxui::flexbox_helper::Block>
    std::vector<Line> lines; // 171,360 bytes across 1,500 allocs : std::vector<ftxui::flexbox_helper::Line>
};
```

# Testing The Profiler on a Non-Trivial Codebase

We're going to profile CMake

# Profiling CMake

- 935 translation units, a million lines of code, multiple dependencies...
- We can build the *whole thing* with the compiler plugin, and profile it with

mem\_profile

```
// Totals for cmState sizeof=600 bytes
// └─ 3,474,745 bytes across 36738 allocs and 3 instances
struct cmState
{
    cmPropertyDefinitionMap      PropertyDefinitions; // 2,592 bytes across 18 allocs : cmPropertyDefinitionMap
    std::vector<std::string>      EnabledLanguages; // 144 bytes across 3 allocs : std::vector<std::string>
    std::unordered_map<std::string, Command> BuiltinCommands; // 34,569 bytes across 468 allocs : std::unordered_map<std::string, std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>>
    std::unordered_map<std::string, Command> ScriptedCommands; // 1,973,623 bytes across 18,180 allocs : std::unordered_map<std::string, std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>>
    std::unordered_set<std::string> FlowControlCommands; // 2,592 bytes across 54 allocs : std::unordered_set<std::string>
    cmPropertyMap                GlobalProperties; // 3,049 bytes across 55 allocs : cmPropertyMap
    std::unique_ptr<cmCacheManager> CacheManager; // 78,980 bytes across 1,207 allocs : std::unique_ptr<cmCacheManager>
    std::unique_ptr<cmGlobVerificationManager> GlobVerificationManager; // 216 bytes across 3 allocs : std::unique_ptr<cmGlobVerificationManager>
    cmLinkedTree<cmStateDetail::BuildsystemDirectoryStateType> BuildsystemDirectory; // 8,192 bytes across 43 allocs : cmLinkedTree<cmStateDetail::BuildsystemDirectoryStateType>
    cmLinkedTree<std::string>      ExecutionListFiles; // 40,088 bytes across 315 allocs : cmLinkedTree<std::string>
    cmLinkedTree<cmStateDetail::PolicyStackEntry> PolicyStack; // 3,072 bytes across 6 allocs : cmLinkedTree<cmStateDetail::PolicyStackEntry>
    cmLinkedTree<cmStateDetail::SnapshotDataType> SnapshotData; // 129,024 bytes across 6 allocs : cmLinkedTree<cmStateDetail::SnapshotDataType>
    cmLinkedTree<cmDefinitions> VarTree; // 1,198,188 bytes across 16,374 allocs : cmLinkedTree<cmDefinitions>
    std::string                  SourceDirectory;
    std::string                  BinaryDirectory;
    bool                         IsGeneratorMultiConfig;
    bool                         WindowsShell;
    bool                         WindowsVSIIDE;
    bool                         CMakeMultiIDE;
```

> ll

total 99941168

drwxr-xr-x@	7	alecto	staff	224B	Aug	28	20:09	Auxiliary
-rwxr-xr-x@	1	alecto	staff	65K	Aug	28	20:09	bootstrap
drwxr-xr-x@	27	alecto	staff	864B	Sep	11	10:06	build
drwxr-xr-x@	4	alecto	staff	128B	Sep	1	21:54	build_mp
-rw-r--r--@	1	alecto	staff	15G	Aug	28	22:20	cmake_stats.json
-rw-r--r--@	1	alecto	staff	794B	Aug	28	20:09	cmake_uninstall.cmake.in
-rw-r--r--@	1	alecto	staff	9.8K	Aug	28	20:09	CMakeCPack.cmake
-rw-r--r--@	1	alecto	staff	12K	Aug	28	20:09	CMakeCPackOptions.cmake.in
-rw-r--r--@	1	alecto	staff	153B	Aug	28	20:09	CMakeGraphVizOptions.cmake
-rw-r--r--@	1	alecto	staff	22K	Aug	28	20:19	CMakeLists.txt
-rw-r--r--@	1	alecto	staff	4.4K	Aug	28	20:09	CMakeLogo.gif
-rw-r--r--@	1	alecto	staff	5.2K	Aug	28	20:09	CompileFlags.cmake
-rwxr-xr-x@	1	alecto	staff	99B	Aug	28	20:09	configure
-rw-r--r--@	1	alecto	staff	3.2K	Aug	28	20:09	CONTRIBUTING.rst
-rw-r--r--@	1	alecto	staff	4.1K	Aug	28	20:09	CONTRIBUTORS.rst
-rw-r--r--@	1	alecto	staff	564B	Aug	28	20:09	CTestConfig.cmake
-rw-r--r--@	1	alecto	staff	7.7K	Aug	28	20:09	CTestCustom.cmake.in
-rw-r--r--@	1	alecto	staff	372B	Aug	28	20:09	DartConfig.cmake
-rw-r--r--@	1	alecto	staff	31G	Aug	29	14:08	event_table.json



```
> uv run python -m mp_reader stats --top-n-layouts=100 /opt/3rd/CMake/cmake_stats.json
```

```
29m 6.0s to load json...
```

```
7m 55.0s to create output record...
```

## Allocation Statistics by Type:

14,215 objects	544,249,813 bytes	std::vector<char>
5,719 objects	375,678,464 bytes	std::unique_ptr<cmUVStreamReadHandle>
5,719 objects	375,175,192 bytes	cmUVStreamReadHandle
116,042 objects	202,228,957 bytes	std::vector<cmListFileFunction>
764,374 objects	180,782,606 bytes	std::shared_ptr<const cmListFileFunction::Implementation>
764,374 objects	180,782,606 bytes	cmListFileFunction
2,109,345 objects	139,738,246 bytes	cmListFileContext
326 objects	137,039,317 bytes	std::unique_ptr<cmState>
660,459 objects	104,075,801 bytes	std::vector<cmListFileArgument>
780,757 objects	96,498,522 bytes	std::vector<std::string>
575,310 objects	95,172,718 bytes	cmListFileFunction::Implementation
20,832 objects	90,850,328 bytes	std::__function::__value_func<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>
20,832 objects	90,850,328 bytes	std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>
9,206 objects	88,932,201 bytes	cmListFile
652 objects	87,755,955 bytes	std::__hash_table<std::__hash_value_type<std::string, std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>>, std::__unordered_map_hasher<std::string, std::__hash_value_type<std::string, std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>>, std::hash<std::string>, std::equal_to<std::string>>, std::__unordered_map_equal<std::string, std::__hash_value_type<std::string, std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>>, std::equal_to<std::string>, std::hash<std::string>>, std::allocator<std::__hash_value_type<std::string, std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>>>>
652 objects	87,755,955 bytes	std::unordered_map<std::string, std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>>
17,239 objects	83,871,867 bytes	std::pair<const std::string, std::function<bool (const std::vector<cmListFileArgument> &, cmExecutionStatus &)>>
363,885 objects	70,733,073 bytes	std::shared_ptr<const cmStackEntry<const cmListFileContext, cmStackType::Const>>
8,982 objects	63,414,569 bytes	(anonymous namespace)::cmMacroHelperCommand
323,312 objects	62,795,168 bytes	cmMakefile::CallScope
697,026 objects	54,859,728 bytes	std::shared_ptr<const std::string>
4,104 objects	46,640,909 bytes	std::__hash_table<std::__hash_value_type<cm::String, cmDefinitions::Def>, std::__unordered_map_hasher<cm::String, std::__hash_value_type<cm::String, cmDefinitions::Def>, std::hash<cm::String>, std::equal_to<cm::String>>, std::__unordered_map_equal<cm::String, std::__hash_value_type<cm::String, cmDefinitions::Def>, std::equal_to<cm::String>, std::hash<cm::String>>, std::allocator<std::__hash_value_type<cm::String, cmDefinitions::Def>>>>
4,104 objects	46,640,909 bytes	std::unordered_map<cm::String, cmDefinitions::Def>

```
// Totals for cmCacheManager::CacheEntry sizeof=80 bytes
```

```
// └─ 4,282,520 bytes across 82949 allocs
```

```
struct cmCacheManager::CacheEntry
```

```
{
```

```
    std::string          Value;
```

```
    cmStateEnums::CacheEntryType Type;
```

```
    cmPropertyMap        Properties; // 4,022,680 bytes across 78,851 allocs : cmPropertyMap
```

```
    bool                 Initialized;
```

```
~cmCacheManager::CacheEntry();
```

```
// directly owned (or unannotated):
```

```
// └─ 259,840 bytes across 4,098 allocs
```

```
};
```

```
// Totals for cmNinjaBuild sizeof=240 bytes
```

```
// └─ 4,234,664 bytes across 61284 allocs
```

```
struct cmNinjaBuild
```

```
{
```

```
    std::string Comment;
```

```
    std::string Rule;
```

```
    cmNinjaDeps Outputs; // 331,034 bytes across 9,107 allocs : std::vector<std::string>
```

```
    cmNinjaDeps ImplicitOuts;
```

```
    cmNinjaDeps WorkDirOuts; // 1,539 bytes across 56 allocs : std::vector<std::string>
```

```
    cmNinjaDeps ExplicitDeps; // 317,001 bytes across 6,768 allocs : std::vector<std::string>
```

```
    cmNinjaDeps ImplicitDeps; // 45,657 bytes across 270 allocs : std::vector<std::string>
```

```
    cmNinjaDeps OrderOnlyDeps; // 131,077 bytes across 3,398 allocs : std::vector<std::string>
```

```
    cmNinjaVars Variables; // 3,045,588 bytes across 35,244 allocs : std::map<std::string, std::string>
```

```
    std::string RspFile;
```

```
~cmNinjaBuild();
```

```
// directly owned (or unannotated):
```

```
// └─ 362,768 bytes across 6,441 allocs
```

```
};
```

```
// Totals for std::__tree<std::__value_type<std::string, std::unique_ptr<cmComputeLinkInformation>>, std::__map_value_compare<std::string, std::__value_type<std::string, std::unique_ptr<cmComputeLinkInformation>>, std::less<std::string>>, std::allocator<std::__value_type<std::string, std::unique_ptr<cmComputeLinkInformation>>>> sizeof=24 bytes
```

```
unique_ptr<cmComputeLinkInformation>>>> sizeof=24 bytes
```



# Caveats

- You need to use clang
- The plugin currently has issues annotating non-inline destructor calls (a fix is in the works)
- If a shared library provides a definition for a destructor (such as `std::string`'s destructor, which is provided by the standard library), this can result in missing annotations.

Another fix is in the works here.

# Final Remarks

- This tool provides increased visibility to memory usage on a per-type and per-object basis
- It makes it a lot more obvious which types are consuming a lot of memory
- It shows a breakdown by *field*, so that you can see which fields are taking up space.

All together, this provides the data necessary to identify inefficiencies, and optimize code for memory usage.

**Thank You!**